

sercon

User Manual

Version 0.49.1

2026-06-11

Repository · <https://github.com/codedeviate/sercon>

License · MIT

sercon manual

Long-form reference for the `pkg/scriptengine` library, the `sercon` CLI, the built-in script API, and the JavaScript runtime surface that `scriptengine` exposes via `goja` and `goja_nodejs`. Examples are runnable — save snippets as `.ts` files and execute with `sercon file.ts` unless otherwise noted.

Table of contents

1. [Overview](#)
 2. [Quickstart](#)
 3. [Library API](#) — `pkg/scriptengine`
 4. [CLI](#) — `sercon`
 5. [Reserved globals](#) (script surface)
 6. [Servers](#)
 7. [JavaScript runtime built-ins](#) (`goja`)
 8. [Async runtime additions](#) (`goja_nodejs`)
 9. [TypeScript support](#)
 10. [Top-level](#) `await`
 11. [Module resolution](#)
 12. [Timeouts and cancellation](#)
 13. [Error semantics](#)
 14. [Type generation](#) (`.d.ts`)
 15. [Limitations and gotchas](#)
 16. [Binding reference](#) (generated)
-

1. Overview

`sercon` is a Go program that runs TypeScript files. It's built from two parts:

- `pkg/scriptengine` — a library you embed in your own Go code. You register Go-callable bindings, then call `Run` / `RunFile` to execute a `.ts` script that talks to them.
- `cmd/sercon` — a thin CLI built on the library, useful for ad-hoc test scripts and as a working example of every binding kind.

The runtime is pure Go: `goja` is the JS engine, `esbuild` (used as a Go library) is the TS→JS transpiler, and `goja_nodejs` provides Promises, `setTimeout`, `console`, and CommonJS `require`. No cgo, no Node.

2. Quickstart

Install:

```
go install github.com/codedeviate/sercon/cmd/sercon@latest
```

Run one of the bundled examples:

```
sercon examples/scripts/smoke.ts examples/scripts/async.ts
```

...or run them all via `make demo`. `examples/scripts/` ships a runnable `.ts` (or `.tsx`) file per feature area — `hash.ts`, `strings.ts`, `path-and-time.ts`, `default-export.ts`, `tsx-demo.ts`, and the original `smoke.ts` / `async.ts` / `hang.ts` (`hang.ts` is the timeout demo and intentionally exits non-zero).

Generate a declaration file for editor autocomplete:

```
sercon --emit-dts sercon.d.ts
```

Embed in your own Go program:

```
eng := scriptengine.New(scriptengine.Options{
    Timeout:    5 * time.Second,
    ScriptRoot: "./scripts",
})
eng.Register("upper", strings.ToUpper)
val, err := eng.RunFile(ctx, "scripts/main.ts")
```

3. Library API — `pkg/scriptengine`

`Engine`, `Options`, `New`

```
type Options struct {
    Timeout      time.Duration // wall-clock limit per Run; 0 disables
    ScriptRoot   string        // base for require/import resolution
    DisableConsole bool          // turn off the goja_nodejs console module
    Verbose      io.Writer     // diagnostic traces, prefixed "[sercon] "
    ModuleLoader func(candidatePath string) (source string, found bool, err error)
}

func New(opts Options) *Engine
```

`ModuleLoader`, when non-nil, is consulted for every require/import candidate path **before** the filesystem — the hook for embedders that want to serve modules from somewhere other than disk (an in-memory FS, a network source, an embedded bundle). goja probes several candidates per specifier (`./x`, `./x.ts`, `./x/index.ts`, ...); the engine also tries the bare path plus the usual extension fallbacks (`.ts` / `.tsx` / `.js` / `.mjs` / `.cjs` / `.json`) against the loader, so a loader can match on

a plain suffix. Return `(source, true, nil)` to serve the module (transpiled when the path ends in `.ts` / `.tsx`, exactly like a disk read); `( "", false, nil)` to fall through to the filesystem; `( "", false, err)` to abort resolution.

```
eng := scriptengine.New(scriptengine.Options{
    ModuleLoader: func(path string) (string, bool, error) {
        if strings.HasSuffix(path, "greeting.ts") {
            return `export const hi = (n: string) => "hi " + n;`, true, nil
        }
        return "", false, nil // fall through to disk
    },
})
```

`Engine` is the host of all registered bindings. Construct it once, then call `Run` / `RunFile` as many times as needed — each call gets a *fresh* `*goja.Runtime` so script state never leaks between invocations.

`ScriptRoot` defaults to the working directory at `New` time if left empty. Both `Timeout` and `ScriptRoot` can be overridden on a per-Run basis only by constructing a fresh `Engine`; see [Out-of-scope](#) for the per-Run override on the roadmap.

Registering bindings

Function	Use when...
<code>Register(name, value)</code>	The binding is a pure function, struct, or primitive that doesn't need access to the runtime.
<code>RegisterNamespace(name, members map[string]any)</code>	You want to expose <code>name.x</code> , <code>name.y</code> without defining a Go struct.
<code>RegisterConstructor(name, ctor)</code>	The binding is a Go function whose return value should be treated as a class in the emitted <code>.d.ts</code> . (Runtime semantics today are identical to <code>Register</code> ; the <code>d.ts</code> emitter is the only differentiator.)
<code>RegisterFactory(name, factory func(vm, loop) any)</code>	The binding needs the per-Run <code>*goja.Runtime</code> or <code>*eventloop.EventLoop</code> in scope (typical for Promise-returning I/O).
<code>RegisterNamespaceFactory(name, factory)</code>	Same, but the factory returns the full member map.

Example — synchronous function binding:

```
eng.Register("greet", func(name string) string { return "hi " + name })
```

Example — namespace with a sync member:

```
eng.RegisterNamespace("math2", map[string]any{
    "square": func(n float64) float64 { return n * n },
    "pi":     3.14159,
})
```

Example — async (Promise-returning) binding via `PromisifyAsync`:

```
eng.RegisterFactory("httpGet", func(vm *goja.Runtime, loop *eventloop.EventLoop)
any {
    return scriptengine.PromisifyAsync(vm, loop,
        func(ctx context.Context, call goja.FunctionCall) (map[string]any,
error) {
            url := call.Argument(0).String()
            resp, err := http.Get(url)
            if err != nil {
                return nil, err
            }
            defer resp.Body.Close()
            body, _ := io.ReadAll(resp.Body)
            return map[string]any{"status": resp.StatusCode, "body":
string(body)}, nil
        })
})
```

Run, RunFile

```
func (e *Engine) Run(ctx context.Context, name, source string, opts
...RunOption) (goja.Value, error)
func (e *Engine) RunFile(ctx context.Context, path string, opts ...RunOption)
(goja.Value, error)
```

`Run` executes `source` as an entry-script TS. `name` is used in stack traces. The returned value is currently always `undefined` — top-level expression capture is on the backlog.

Both methods respect `Options.Timeout` and `ctx` cancellation, whichever fires first; the resulting error is either `ErrScriptTimeout`, `ctx.Err()`, or the underlying JS exception.

Per-Run options

```
type RunOption func(*runConfig)

func WithScriptRoot(dir string) RunOption
```

Reuse a single Engine across many scripts that each live in their own directory by passing `WithScriptRoot(dir)` to `Run` / `RunFile`. The override applies only to this call; `Options.ScriptRoot` is used for every other `Run`.

```
_, _ = eng.Run(ctx, "main.ts", source,
scriptengine.WithScriptRoot("/path/to/run42"))
```

```
func WithArgs(args []string) RunOption
```

Set the per-script argument vector. The script reads it as `runtime.argv`, a global with Node/Bun layout: `argv[0]` is `Options.ProgramName` (defaults to `filepath.Base(os.Args[0])`); the CLI sets it to `"sercon"`), `argv[1]` is the running script's path, and the `WithArgs` values follow from index 2. The global is always present — with no args it is just `[programName, scriptPath]`.

```
_, _ = eng.Run(ctx, "main.ts", source, scriptengine.WithArgs([]string{"--port",
"8080"}))
// inside the script: runtime.argv == [ProgramName, "/abs/main.ts", "--port",
"8080"]
```

Reset

```
func (e *Engine) Reset()
```

Clears every registered binding. Useful when a long-lived Engine is reused across unrelated batches of scripts that each want a clean global namespace. **Not safe to call concurrently with Run / RunFile.**

WriteTypes

```
func (e *Engine) WriteTypes(w io.Writer) error
```

Emits a `.d.ts` describing the registered surface. See section [14. Type generation](#) for what the mapping looks like.

SetDocs and SetMemberDocs

```
func (e *Engine) SetDocs(path, doc string)
func (e *Engine) SetMemberDocs(namespace string, docs map[string]string)
```

Attach JSDoc strings to registered bindings so the emitted `.d.ts` grows readable editor hover.

`SetDocs` takes a dotted path — either a bare top-level name (`"greet"`, `"http"`) or `"namespace.member"` for namespace members (`"http.get"`, `"exec.shell"`); `SetMemberDocs` is the convenience for documenting many members of a namespace at once (the namespace itself can still be documented separately via `SetDocs`).

Multi-line docs are written as `\n`-separated strings; the emitter expands them into a standard `* -` prefixed JSDoc block. Single-line docs collapse to `/** ... */`. Calling `SetDocs` with an empty string removes any previous doc for that path. Bindings without a doc entry get no JSDoc block at all — the emitter doesn't insert empty `/** */` placeholders.

`SetDocs` may be called before or after the matching `Register...` call; docs are looked up by name at emit time. Like the registration methods, `SetDocs` must not race with a `Run` / `WriteTypes` / `Reset` on the same engine.

```
eng.Register("greet", func(name string) string { return "hi " + name })
eng.SetDocs("greet", "Greet someone by name.")

eng.RegisterNamespace("math2", map[string]any{
    "pi":      3.14159,
    "squared": func(n float64) float64 { return n * n },
})
eng.SetDocs("math2", "Tiny math helpers.")
eng.SetMemberDocs("math2", map[string]string{
    "pi":      "Circumference / diameter ratio.",
    "squared": "Multiply a number by itself.",
})
```

PromisifyAsync[T] and AsyncBinding

```
type AsyncBinding struct {
    Func      func(goja.FunctionCall) goja.Value
    TSReturnType string
}

func PromisifyAsync[T any](vm *goja.Runtime, loop *eventloop.EventLoop,
    work func(ctx context.Context, call goja.FunctionCall) (T, error),
) AsyncBinding
```

`PromisifyAsync` turns blocking Go work into a JS Promise. It launches the work in a goroutine, parks a `SetTimeout` sentinel so `loop.Run` doesn't return early, and schedules the resolve/reject back onto the event loop. **Required** for any Promise-returning binding — `RunOnLoop` alone is not counted as a live job by the event loop.

`PromisifyAsync` returns an `AsyncBinding` carrier rather than the bare callback. The engine unwraps it to `.Func` at `vm.Set` time so goja's special-case detection of `func(goja.FunctionCall) goja.Value` host-callbacks still fires; the `.d.ts` emitter reads `.TSReturnType` to emit `Promise<T>` (mapped from the generic `T` via reflect at construction time) instead of the previous `unknown`. Hosts assigning the result into a `map[string]any` for `RegisterNamespace` / `RegisterNamespaceFactory` get the unwrap recursively too — no manual work required.

Version

```
import "github.com/codedeviate/sercon/pkg/scriptengine"

fmt.Println(scriptengine.Version) // "0.1.0"
```

Bumped in lockstep with the git tag.

4. CLI — sercon

```
sercon [flags] <script.ts> [script.ts ...]
sercon [flags] -                # read entry script from stdin
sercon --examples | --help | --version
```

Flag	Purpose
<code>-timeout DURATION</code>	Wall-clock limit per script (default <code>10s</code> ; <code>0</code> disables).
<code>-root DIR</code>	Override the script root for <code>require</code> / <code>import</code> resolution.
<code>-emit-dts PATH</code>	Write the example bindings' <code>.d.ts</code> to <code>PATH</code> and exit.
<code>-v</code>	Verbose: trace the rewritten entry-script JS and each module resolution to stderr; also print duration on failure.
<code>-h</code> , <code>--help</code>	In-depth colored help.
<code>--examples</code>	In-depth colored walkthrough of every feature.
<code>--no-pager</code>	Don't page <code>--help</code> / <code>--examples</code> . By default, when stdout is a terminal they pipe through <code>\$PAGER</code> (falling back to <code>less</code> with <code>LESS=FRX</code> , color preserved); a pipe/redirect, <code>--no-pager</code> , or <code>PAGER=cat</code> renders directly.
<code>--version</code>	Print the engine version (plus goja/esbuild build-info versions).
<code>--watch</code>	Re-run on file change under the script root. Debounced 150 ms. Ctrl-C exits.
<code>--secrets-prefix=P</code>	Namespace prefix for <code>runtime.secrets</code> keystore items. Overrides <code>\$SERCON_SECRETS_PREFIX</code> ; default <code>sercon/</code> .

Each positional argument is either a path to a `.ts` / `.tsx` file or `-` to read an entry script from standard input:

```
echo 'runtime.log(1 + 2);' | sercon -
sercon prelude.ts -                # prelude then stdin
```

Passing arguments: `--` and `runtime.argv`

Everything after a standalone `--` is the script argument vector, exposed to every script via `runtime.argv`:

```
sercon run.ts -- --port 8080 hello
```

`runtime.argv` uses the Node/Bun layout `[programName, scriptPath, ...userArgs]` — `argv[0]` is `"sercon"`, `argv[1]` is the path of the script currently executing, and the post-`--` arguments start at index 2 (so `runtime.argv.slice(2)` is just your args). All scripts in one invocation share the same `--` tail, but each sees its own path at `argv[1]`. The global is always present; with no `--` it is just `[programName, scriptPath]`.

```
const args = runtime.argv.slice(2); // ["--port", "8080", "hello"]
```

Exit codes are distinct per failure type so shells can react sensibly. When several scripts run, the highest applicable code wins:

Code	Meaning
0	every script passed.
1	CLI usage error (unknown flag, missing script argument, ...).
2	at least one script failed to transpile — never ran.
3	at least one script timed out (<code>-timeout</code>) or was context-cancelled.
4	at least one script ran and threw a JS exception.

`-v` writes lines prefixed with `[sercon]` to stderr. The traces include the full rewritten entry-script JS (the form goja actually runs, after the ESM→CJS rewrite + async IIFE wrapper) and every module-resolution event, so debugging an unexpected resolve target or a mis-rewritten import is straightforward.

The CLI registers ten reserved top-level globals; see the next section.

--watch : re-run on file change

`sercon --watch <script.ts>` runs the script once and then blocks, re-running it every time a watched file changes under the script root. Useful for iterating on a script body or on the binding surface during development.

```
sercon --watch examples/scripts/smoke.ts
```

What's watched:

- The script root (resolved the same way as the one-shot mode — `-root DIR` if set, else `dirname` of the first script argument).
- Recursively. Symlinks aren't followed.
- New directories that appear during the session are picked up automatically.

What's filtered:

- **File extensions:** `.ts` / `.tsx` / `.js` / `.jsx` / `.json` trigger a re-run; `.d.ts` files match via suffix too. Anything else (Markdown, images, editor swap files, build artefacts) is ignored.
- **Directories:** hidden directories (`.git`, `.vscode`, anything starting with `.`) and `node_modules` are excluded — both because they generate floods of irrelevant events and because recursing into them inflates the watcher's directory count for no gain.

Re-runs are **debounced 150 ms**. Editors typically fire several events per save (write → rename → chmod); collapsing them into a single re-run keeps the loop responsive without thrashing. The debounce window resets on every event, so a rapid burst of saves becomes one re-run after the burst settles.

Each run is delimited by a line like `--- sercon re-run @ HH:MM:SS ---` so the output is visually distinct from the previous run.

`Ctrl-C` (SIGINT) or `SIGTERM` exit cleanly. Per-script failures inside a watch session log as `FAIL` but don't terminate the loop — a syntax error you're trying to fix shouldn't kick you out of watch mode.

The watcher reuses the same `Engine` across runs. Each `Run` already gets a fresh `*goja.Runtime`, so re-running is just re-invoking `runOne` on the existing engine — there's no per-iteration setup cost beyond the transpile + module-resolution work.

Module-graph invalidation. Watch mode tracks each entry script's import graph (its own file plus every module it resolves, captured via `Engine.SetResolveHook`) during the run. On a file change it re-runs **only the entries whose graph includes the changed file** — editing a helper that just one of three entry scripts imports re-runs that one entry, not all three. An entry's own file counts (editing it re-runs it); stdin entries and any entry whose graph isn't known yet re-run unconditionally (conservative). Before each re-run the module cache is busted (`Engine.ResetModuleCache`) so an edited import's new source actually takes effect — the registry otherwise caches compiled bytecode across runs.

Editor autocomplete (`sercon init`)

```
sercon init [dir]
```

Drops two files into `dir` (default the current directory) so any editor backed by the TypeScript language server (VSCode, Zed, Neovim+coc, Sublime LSP, ...) gives completion and hover docs for the reserved globals with no plugin or manual config:

- **`sercon.d.ts`** — the binding declarations (same content as `sercon -emit-dts`): ambient `declare const` blocks for `runtime`, `crypto`, `text`, `codec`, `fs`, `net`, `db`, `server`, `services`, `tui`, and `console`, each with JSDoc.
- **`jsconfig.json`** — points the language server at `sercon.d.ts` and sets `moduleResolution: "Bundler"` (so the extensionless relative imports sercon scripts use `resolve`) and `types: []` (sercon isn't Node, so no stray `@types/*`).

Existing files are left untouched unless `--force` is given. The `examples/scripts/` directory ships with both files already, so the bundled demos autocomplete out of the box. For a single file without a config, the no-setup fallback is a leading `/// <reference path="./sercon.d.ts" />`.

Shebang lines and executable scripts (`sercon run`)

A script may begin with a `#!` shebang line. It is stripped before transpile (the line is blanked in place, so transpile/syntax error line numbers still match the source), which means a `.ts` / `.tsx` / `.js` file can be made directly executable:

```
#!/usr/bin/env sercon
runtime.log("hello from an executable script");
```

```
chmod +x hello.ts
./hello.ts
```

The kernel launches a shebang script as `sercon <script> <args...>`, and in the default mode every positional is treated as a *separate script* — so positional arguments to a shebang script would be mistaken for additional script paths. For an executable script that takes arguments, use the **run subcommand** in the shebang:

```
sercon run [flags] <script.ts> [args...]
```

```
#!/usr/bin/env -S sercon run
runtime.log("args:", JSON.stringify(runtime.argv.slice(2)));
```

`sercon run` executes exactly one script and hands every token after the script path to it as `runtime.argv[2:]` (Node/Bun layout: `[program, script, ...args]`) — no standalone `--` separator is needed (and a shebang line can't inject one). The `env -S` form splits the single shebang argument so the kernel runs `sercon run /abs/path/script.ts arg1 arg2 ...`. Flags (`-timeout`, `-root`, `-v`) are accepted before the script path; everything from the script path onward is positional and becomes script args. `env -S` is supported by GNU coreutils, macOS, and the BSDs.

```
sercon run script.ts --port 8080 alice    # argv[2:] = ["--port","8080","alice"]
```

sercon serve : long-running scripts with production niceties

```
sercon serve [flags] <script.ts> [-- args...]
```

`sercon serve` is a sibling subcommand that wraps the same engine but adds the conveniences a process supervisor (systemd, docker, launchd) usually wants from a long-lived script. Use it whenever a script binds an HTTP/HTTPS listener (or any future `server.*` protocol) and is expected to keep running. Vanilla `sercon script.ts` also keeps the loop alive while listeners are bound, but gives you none of the deltas below.

Flag	Purpose
<code>--shutdown-timeout DURATION</code>	Graceful-shutdown deadline on SIGTERM/SIGINT (default <code>30s</code>). After the deadline the engine hard-cancels.
<code>--port-override N</code>	If non-zero, replaces every literal <code>port:</code> in <code>server.*.listen({...})</code> calls with <code>N</code> . Useful for spinning up the same script on a different port without editing source. Only literal port values are rewritten — computed expressions (<code>{port: getPort()}</code>) are left alone.
<code>-timeout DURATION</code>	Per-script wall-clock limit. Defaults to <code>0</code> (disabled) under <code>serve</code> so listeners don't get cut off; opt back in if you want one.
<code>-root DIR</code>	Script root for <code>require / import</code> resolution (same semantics as vanilla <code>sercon</code>).
<code>-v</code>	Verbose engine tracing to stderr.

Behavioural deltas vs vanilla `sercon` :

- **Access log to stderr.** Each request emits one line in the format `ts remote method path status dur_μs`, e.g. `2026-05-29T10:23:14Z 127.0.0.1:54321 GET /users/42 200 1843μs`. WebSocket upgrades log only the upgrade line; per-frame traffic is suppressed.
- **Readiness signal to stdout.** When each listener calls back into the binding with `bound`, the binding prints `READY` listening on `tcp/0.0.0.0:8080` to stdout. Multiple listeners produce one `READY` line each. Supervisors that read-line on startup (`systemd-notify`, `docker-compose` healthchecks) can wait on it.
- **Graceful shutdown.** `SIGTERM` / `SIGINT` triggers `srv.close()` on every active listener concurrently, then waits up to `--shutdown-timeout` for the script to drain. Clean exit is exit code `0`. Past the deadline, `loop.Terminate()` fires and the process exits with the usual classification.
- **No `--watch`.** Re-running while a listener is bound would race port rebinding; `sercon serve --watch ...` exits with a usage error. Use vanilla `sercon --watch script.ts` for the hot-reload development loop.

```
sercon serve examples/scripts/server-http.ts
sercon serve --port-override 9090 server.ts
sercon serve --shutdown-timeout 10s server.ts
```

5. Reserved globals (script surface)

`sercon` scripts get ten reserved top-level globals. Use them directly — there is no enclosing namespace. The full per-binding reference is the generated `examples/scripts/sercon.d.ts`; this section is the at-a-glance prose.

Reserved names: `runtime`, `crypto`, `text`, `codec`, `fs`, `net`, `db`, `server`, `services`, `tui`. User code can shadow these with a local `let` / `const` / `var` per normal JavaScript scoping — `sercon` does not intervene at runtime. `server` (added in v0.10.0) covers inbound listeners — see [section 6](#) for the long-form treatment.

Deterministic key order. Objects returned from bindings enumerate their keys in a stable order (declaration order for fixed-shape results; source / column / insertion order for dynamic ones), so `JSON.stringify(result)` is byte-stable run-to-run — safe to hash for canonical serialization (payment signing, webhook signatures). The lone exception is `text.jq`, whose underlying engine discards key order internally.

`console` (browser/Node compatibility)

Beyond the ten reserved globals, `sercon` provides a `console` global so scripts pasted from a browser or Node run unchanged:

Method	Stream
<code>console.log</code> / <code>console.info</code> / <code>console.debug</code>	stdout
<code>console.warn</code> / <code>console.error</code>	stderr

Each prints one space-joined line — clean and prefix-free (no timestamp), unlike the goja default console. Primitives print raw; objects and arrays render as JSON (`console.log({a:1})` → `{"a":1}`), with circular references falling back to `[object Object]` rather than throwing. The CLI disables the engine's built-in console so this stream-correct shim is the only `console` a script sees. `runtime.log` is the native equivalent of `console.log` (`stdout`) and formats the same way. Library embedders get the `goja_nodejs` console by default instead; toggle it with `Options.DisableConsole`.

Migration from v0.8.0

v0.8.0	v0.9.0
<code>api.runtime.log</code>	<code>runtime.log</code>
<code>api.crypto.hash.sha256</code>	<code>crypto.hash.sha256</code>
<code>api.text.str.trim</code>	<code>text.str.trim</code>
<code>api.format.barcode.encode</code>	<code>codec.barcode.encode</code>
<code>api.fs.path.basename</code>	<code>fs.path.basename</code>
<code>api.net.http.get</code>	<code>net.http.get</code>
<code>api.db.sqlite.open</code>	<code>db.sqlite.open</code>
<code>api.tools.exec.shell</code>	<code>services.exec.shell</code>
<code>api.tools.git.commit</code>	<code>services.git.commit</code>
<code>api.ui.tui.layout</code>	<code>tui.layout</code>
<code>Sercon.argv</code>	<code>runtime.argv</code>

The `api` global no longer exists; reading it throws a `ReferenceError`. The `Sercon` global is also gone.

A mechanical sed pass over your scripts handles the prefix drop and the three renames — see the `CHANGELOG.md` `[0.9.0]` "Changed" entry for a sketch.

`runtime`

Script-host scaffolding. Members:

- `runtime.log(...args)` — print a space-joined line (primitives raw, objects/arrays as JSON).
- `runtime.assert.equal(actual, expected, msg?)` / `runtime.assert.ok(cond, msg?)` — throw on mismatch / falsy. `assert.equal` is **deep**: distinct objects/arrays with identical contents compare equal (structural, recursive — key order irrelevant).
- `runtime.time.nowMs()` — current wall-clock in milliseconds.
- `runtime.time.sleep(ms)` — Promise that resolves after `ms` on the event loop.
- `runtime.time.format(unixMs, layout, tz?)` — Go-style time layout (e.g. `"2006-01-02 15:04:05"`).

- `runtime.env.get(name)` — process environment variable; returns `undefined` when unset (never throws).
- `runtime.argv` — `[programName, scriptPath, ...userArgs]`. User args come from the CLI args after a `--` separator. The layout mirrors Node/Bun: `argv[0]` is `"sercon"`, `argv[1]` is the path of the currently-executing script (so each script in a multi-arg invocation sees its own path), and `argv.slice(2)` is just your args. The property is always present; with no `--` it is just `[programName, scriptPath]`.
- **`runtime.secrets`** — read/write string credentials in the OS keystore (macOS Keychain, Linux Secret Service / `libsecret`, Windows Credential Manager). Pure-Go (no cgo) via `zaland/go-keyring`.
 - `runtime.secrets.available` — `boolean`, advisory: is a keystore backend reachable this run? Gate on it to self-skip on headless boxes.
 - `runtime.secrets.get(name, account)` — `Promise<string | null>`; `null` when absent.
 - `runtime.secrets.set(name, account, secret)` — `Promise<void>`.
 - `runtime.secrets.delete(name, account)` — `Promise<boolean>` (`true` if removed).

All operations are confined to a **prefix namespace**: the keystore service is `PREFIX + name`, so a script can neither read nor clobber secrets outside it. `PREFIX` resolves from `--secrets-prefix > SERCON_SECRETS_PREFIX > the default sercon/`. Store a secret once with your OS tools, e.g. macOS `security add-generic-password -s 'sercon/devshop' -a tess -w`, Linux `secret-tool store --label='sercon devshop' service 'sercon/devshop' account tess`. Operations are async and bounded by a 10s timeout (a blocking macOS consent prompt rejects rather than hangs).

```
runtime.log("hello");
runtime.assert.equal(2 + 2, 4);
const args = runtime.argv.slice(2); // post-`--` user args
```

crypto

Hashing, JWT, and age-style encryption. Members:

- `crypto.hash.*` — `md5`, `sha1`, `sha256`, `sha384`, `sha512`, `sha3_256`, `sha3_512`, `blake3`, `crc32`. All take a `string` and return the hex digest.
- `crypto.jwt.sign(claims, key, opts?)` / `crypto.jwt.view(token)` / `crypto.jwt.validate(token, key, opts?)` — HS256/RS256/ES256 JWTs.
- `crypto.encrypt.{keygen, keygenPgp, encrypt, decrypt, rekey, detectBackend}` — age-format symmetric encryption with PGP-compatible key wrapping.

text

String, regex, charset, and data manipulation. Members:

- `text.str.*` — `trim`, `ltrim`, `rtrim`, `reverse`, `stripHtml`, `nl2br`, `br2nl`, `base64Encode`, `base64Decode`, `urlencode`, `urlDecode`, `htmlEntityDecode`, `pad` / `lpad` / `rpadd`, `sprintf`, `printf`, `normalizeNewlines`.
- `text.preg.*` and `text.preg2.*` — PCRE-style regex (PHP-compatible named-capture, lookbehind, etc.).
- `text.charset.detect(data)` / `decode(data, charset)` / `encode(text, charset)` — character-set detection and conversion.
- `text.jq.query(json, expr)` / `queryAll(json, expr)` — jq filtering against parsed JSON values.
- `text.diff.compare(a, b, opts?)` — unified / patience / inline diffs between two strings.

codec

Binary-format codecs (was `format` in v0.8.0). Members:

- `codec.compression.{algos, compress, decompress}` — `gzip`, `deflate`, `zlib`, `bzip2`, `zstd`, `brotli`, `lz4`, `xz`, `snappy`.
- `codec.barcode.{formats, decodableFormats, encode, decode}` — `QR`, `DataMatrix`, `Aztec`, `PDF417`, `Code128`, `Code39`, `Codabar`, `EAN-13/8`, `UPC-A`, `ITF`. `encode` returns PNG bytes; `decode` reads image bytes.
- `codec.checkdigit.{algos, validate, compute, inspect}` — `Luhn`, `ISBN-10/13`, `EAN-13/8`, `UPC-A`.
- `codec.php.{serialize, unserialize, varExport, parseVarExport, varDump, parseVarDump}` — read and write PHP data dumps.
- `codec.perl.{dumper, parseDumper}` — read and write Perl `Data::Dumper` output.
- `codec.xml.encode(value, opts?)` / `codec.xml.decode(xml)` — `value` ↔ XML. Attributes are `@`-prefixed keys, text is `#text`, other keys are child elements, and an array value becomes repeated sibling elements; a text-only element collapses to a bare string and an empty/self-closing element to `null`. The value must be a single-key object naming the root (or pass `opts.rootName`). `opts.indent` pretty-prints and `opts.declaration` prepends the XML declaration. Object key order and namespace prefixes are preserved. Decoded values are **strings** (XML is untyped — a number round-trips as its string form); surrounding whitespace is trimmed and mixed text/element ordering is not preserved. Cycles, mismatched tags, multiple roots, and malformed XML throw.

codec.php — PHP dump formats

Three PHP textual formats, each with an encoder and a matching parser:

- `codec.php.serialize(value, opts?): string` — PHP `serialize()`. Emits the canonical wire form (`s:`, `i:`, `d:`, `b:`, `N;`, `a:`, `0:`).
- `codec.php.unserialize(text, opts?): value` — inverse of `serialize`.
- `codec.php.varExport(value, opts?): string` — PHP `var_export()`: valid PHP source (`array ( ... ), \CIs::__set_state( ... )`).
- `codec.php.parseVarExport(text, opts?): value` — read a `var_export` literal back to a value.
- `codec.php.varDump(value, opts?): string` — PHP `var_dump()`: human-readable debug output with type/length annotations.
- `codec.php.parseVarDump(text, opts?): value` — **best-effort** read of `var_dump` output.

Array heuristic. A JS array, or a JS object whose keys are exactly the contiguous integer sequence `0..n-1`, maps to a PHP list array. Any other object maps to a PHP associative array (string/int keys preserved).

The `__class` sentinel. An object carrying a `__class` string key (configurable via `opts.classKey`, default `"__class"`) round-trips as a PHP *object* — `serialize` emits `0:...`, `var_export` emits `\Cls::__set_state(...)`, `var_dump` emits `object(Cls)#...`. The remaining keys become the object's properties. Decoding restores the `__class` key. Without the sentinel a value is treated as an array.

Shared references and cycles. `serialize/unserialize` model PHP's `r: / R:` reference markers: when the same object appears more than once, the decoder resolves both occurrences to the *same* JS object (a DAG is preserved). A true cycle (an object reachable from itself) **throws** on encode — there is no JS value that can represent it losslessly here.

`var_dump` parse is lossy. PHP's `var_dump` is a debug view, not a serialization format; `parseVarDump` is best-effort and **throws** when it hits something it cannot faithfully reconstruct: the `*RECURSION*` marker, a length-truncated string (its declared byte length disagrees with the bytes present), or visibility-annotated property names (`["x":"Cls":private]`).

`codec.perl — Perl Data::Dumper`

- `codec.perl.dumper(value, opts?): string` — emit a `Data::Dumper`-style dump (`$VAR1 = ... ;`) with normalized indentation.
- `codec.perl.parseDumper(text, opts?): value` — read `Data::Dumper` output back to a value.

The JSON bool convention. Perl has no native boolean, so JSON modules represent one as a blessed scalar reference (e.g. `bless( do{\(my $o = 1)}, 'JSON::XS::Boolean' )`). `dumper` emits JS booleans in this form, and `parseDumper` decodes a blessed scalar ref in the JSON-bool family (`JSON::XS::Boolean`, `JSON::PP::Boolean`, ...) back to a JS boolean. A bare `1 / 0` stays a number. The class used on encode is `opts.perlBoolClass` (default `"JSON::XS::Boolean"`). Blessed hashes carry the `__class` sentinel, mirroring the PHP side. Cycles **throw**.

`opts` and shared caveats

All `codec.php.* / codec.perl.*` functions accept an optional `opts` bag:

- `classKey` (default `"__class"`) — the key used as the class sentinel.
- `perlBoolClass` (default `"JSON::XS::Boolean"`) — class for Perl booleans.
- `indent` — override the default 2-space indentation step (`varExport` , `dumper`).

Caveats shared with JSON: strings are UTF-8 and lengths are reported in **bytes** (multibyte chars count as more than one); JS has a single number type, so `int` and `float` collapse the same way `JSON.stringify` does (e.g. `1.0` may re-emit as `1`). Decoded objects preserve a stable key order, so re-encoding is byte-stable — safe for canonical-JSON / payment hashing.

`fs`

Filesystem operations:

- `fs.path.dirname(p) / fs.path.basename(p, suffix?)` — POSIX-style.
- `fs.archive.create(srcPath, opts?) / fs.archive.extract(archivePath, destDir, opts?)` — tar/zip with optional compression.

net

Network clients and probes:

- `net.http.{get, post, request}` — fetch-style HTTP client; `request` accepts headers, body, timeout, retry, follow, and basic auth.
- `net.probe.{tcp, dns, tls, ntp, whois, ping, traceroute, smtp, wss}` — one-shot reachability / handshake probes. Each returns a structured result; failures surface as `Error` objects with details.
- `net.probe.traceroute(host, opts?)` — trace the network path: send probes with increasing TTL and report each responding router. **Needs root / CAP_NET_RAW** (intermediate hops appear only via ICMP `time-exceeded`). `opts.protocol` is `"icmp"` (default), `"udp"` (to an incrementing high port), or `"tcp"` (SYN via a TTL-limited connect); `port` sets the udp/tcp target; `maxHops` (30), `timeout` ms per probe (2000), and `probes` per hop (3) bound the trace. Resolves to one `{ ttl, address, rttMs, reached }` per hop (`address` is `null` for a timed-out hop). IPv4 only.
- `net.probe.ping` also accepts `mode: "udp"` — a UDP datagram to a (closed) port whose ICMP `port-unreachable` proves reachability (needs root / `CAP_NET_RAW`), alongside the existing `"tcp"` (default) and `"icmp"` modes.
- `net.netstatus.check(host)` — combined reachability summary.
- `net.email.*` — `spf`, `dmARC`, `mtasts`, `tlsrpt`, `bimi`, and `all(domain)` which runs every probe in parallel and aggregates; plus `send({...})` — outbound SMTP sender (see §6.7).
- `net.browser.open(url)` — launches the OS default browser.
- `net.tcp.connect(host, port, opts?)` — open a TCP client socket.
- `net.udp.open(opts)` — open a UDP socket (connected or bound).
- `net.icmp.open(opts?)` — open a raw ICMP socket (needs privileges).
- `net.capture.{interfaces, open, openFile, toFile}` — packet capture (live + pcap file I/O), pure-Go `gopacket`.
- `net.raw.{open, tcp}` — craft & send raw IPv4 packets (TCP flags / UDP / arbitrary IP protocol) and receive replies (needs privileges).

Raw sockets (`net.tcp` / `net.udp` / `net.icmp`) are long-lived, bidirectional client sockets with a *push / callback* read model — unlike the one-shot `net.probe.*` helpers, they stay open and deliver inbound data through callbacks until you `close()` them. Each constructor returns a `Promise` for a handle object:

- `net.tcp.connect(host, port, opts?)` — dial a TCP peer. `opts { timeout? (ms, default 10000), readBuffer? (inbound channel capacity, default 64) }`. The handle has `write(data)` (`string` → UTF-8 bytes, `Uint8Array` → raw bytes; returns a `Promise`), the `remote / local` address strings, plus the shared callbacks below. Inbound chunks arrive via `onData`.
- `net.udp.open(opts)` — two modes selected by `opts`. **Connected** `{ host, port, readBuffer? }` resolves the peer up front and exposes `send(data)`. **Bound** `{ bind: "127.0.0.1:0", readBuffer? }` listens on a local address and exposes `sendTo(data, host, port)`; its inbound events also carry `address / port` of the sender. Either way the

handle has `local` (the bound address — read the chosen port back from it when you bind to `:0`) and delivers datagrams via `onMessage`.

- `net.icmp.open(opts?)` — open a raw ICMP socket. **Requires root / CAP_NET_RAW**; without the privilege `open()` rejects with an error naming that requirement. `opts { network?: "ip4" | "ip6" (default "ip4"), readBuffer? }`. `send(opts)` writes a message in one of two modes: **Echo mode** `{ to, type?, code?, id?, seq?, payload? }` builds an Echo-shaped body (`type` defaults to the network's echo request), or **raw mode** `{ to, type, code?, body }` marshals `body (Uint8Array | string)` verbatim — for hand-built non-Echo messages such as a crafted destination-unreachable. In raw mode `type` is required and `body` is mutually exclusive with `id / seq / payload`. The handle has `network / local`; inbound events carry `address / type / code` and arrive via `onMessage`.

All three handles share the same callback surface and lifecycle:

- `onData(cb)` (TCP) / `onMessage(cb)` (UDP, ICMP) — register a listener for inbound data. The event object carries `bytes` (a `Uint8Array`) and `text` (the bytes decoded as UTF-8); UDP-bound / ICMP events add the sender metadata noted above.
- `onClose(cb)` — fires once when the socket closes (locally or remotely).
- `onError(cb)` — fires on a read / connection error (benign teardown closes are reported through `onClose`, not `onError`).
- `close()` — shut the socket down; returns a `Promise`.

The socket keeps the engine's event loop alive while open (via the internal `HoldRun` refcount), so a script that only listens won't exit until it `close()`s. `sercon` scripts can't *bind* a listening TCP/UDP server socket yet — these are client sockets — so a TCP example needs a peer to dial; a UDP loopback pair is fully self-contained.

Packet capture (`net.capture`)

Packet capture and pcap file I/O, powered by **pure-Go gopacket** (no `libpcap`, no `cgo`). Five bindings:

- `net.capture.interfaces()` — **synchronous**; returns an array of `{ name, addresses: string[], up, loopback }` for the host's network interfaces. Pure-Go (`net.Interfaces`); no privileges, all platforms.
- `net.capture.routes()` — **synchronous**; returns an array of `{ destination, gateway, interface, family, metric }` for the host's IP routing table. `destination` is a CIDR (`0.0.0.0/0 / ::/0` are the default routes); `gateway` is the next-hop IP or `""` for a directly-connected route; `family` is `"ip"` or `"ip6"`; `metric` is `0` where the platform doesn't report one. Pure-Go, unprivileged: Linux parses `/proc/net/route + /proc/net/ipv6_route`; macOS/BSD read the routing socket via `golang.org/x/net/route`. **Windows is unsupported** (throws).
- `net.capture.open({ iface, promisc?, snaplen?, filter? }, pkt => {...})` — start a **live** capture on `iface`. Resolves to a handle `{ iface, link, close() }`; the per-frame handler receives a decoded packet object (see below). `promisc` defaults `true`; `snaplen` defaults `262144`. `filter` is an optional tcpdump-like expression string (see **Filtering** below). **Linux + macOS only** — Linux uses `AF_PACKET`, macOS uses `BPF (/dev/bpf)`; **Windows is unsupported** and `open()` rejects there. Requires **root / CAP_NET_RAW (Linux)** or `/dev/bpf`

read access (macOS); without the privilege `open()` rejects naming the requirement. `close()` returns `Promise<void>`.

- `net.capture.openFile(path, pkt => {...}, opts?)` — read a `.pcap` / `.pcapng` file (format auto-detected from the magic bytes), calling the handler once per decoded packet. Returns a `Promise` that resolves at EOF. Offline; no privileges. `opts` is an optional trailing argument `{ filter? }` (the two-argument form still works); `filter` is the same tcpdump-like expression string as `open` (see **Filtering** below).
- `net.capture.toFile(path, { linkType?, snaplen? })` — open/create a `.pcap` file and return `{ write(bytes, { ts? }), close() }`. `write` appends a raw frame (a `Uint8Array`); `opts.ts` (ms) overrides the per-frame timestamp (defaults to now). `close()` flushes and returns `Promise<void>`. `linkType` defaults to Ethernet, `snaplen` to 262144. Offline; no privileges.

The **decoded packet object** passed to the `open` / `openFile` handlers:

```
{
  ts, length, captureLength, link,      // always present
  eth?: { src, dst, type },
  vlan?: { id, priority, drop, type }, // 802.1Q tag; `type` is the inner
  ethertype
  arp?: { operation, senderMac, senderIp, targetMac, targetIp }, // operation:
  "request"|"reply"
  ip?: { version, src, dst, protocol, ttl },
  tcp?: { srcPort, dstPort, seq, ack, window, checksum,
    flags: { syn, ack, fin, rst, psh, urg },
    options?: { mss?, windowScale?, sackPermitted?, timestamps?: { val,
  ecr } } },
  udp?: { srcPort, dstPort, length },
  icmp?: { type, code },
  dns?: { id, qr, opcode, rcode,
    questions: { name, type }[],
    answers: { name, type, data }[] }, // data rendered by type
  (A/AAAA→IP, CNAME→name, ...)
  payload?: Uint8Array,                // application-layer bytes, if any
  bytes: Uint8Array,                   // the full raw frame (always present)
}
```

Layer keys (`eth` / `vlan` / `arp` / `ip` / `tcp` / `udp` / `icmp` / `dns` / `payload`) are present **only when that layer decodes** — a truncated or unrecognised frame still yields the always-present `ts` / `length` / `captureLength` / `link` / `bytes`. `tcp.options` is present only when the segment carries recognised options.

Filtering. Both `open` and `openFile` accept an optional `filter` string — a **subset of tcpdump syntax**. Supported grammar:

- protocols: `tcp`, `udp`, `icmp`, `ip`, `ip6`;
- hosts: `host X`, `src host X`, `dst host X` (IPv4 or IPv6 address);
- networks: `net X/Y`, `src net X/Y`, `dst net X/Y` (IPv4 or IPv6 CIDR prefix, e.g. `net 10.0.0.0/8` or `net 2001:db8::/32`);
- ports: `port N`, `src port N`, `dst port N`;

- port ranges: `portrange A-B`, `src portrange A-B`, `dst portrange A-B` (inclusive, e.g. `portrange 8000-8100`);
- boolean composition: `and`, `or`, `not`, and parentheses;
- **implicit-and** between juxtaposed primaries — `tcp port 80` is exactly `tcp and port 80`.

Example: `net.capture.open({ iface: "en0", filter: "tcp and port 80" }, pkt => {...})`, or `net.capture.openFile("cap.pcap", pkt => {...}, { filter: "udp or icmp" })`.

The filter is a **userspace, post-decode** predicate — it is **not** compiled to a kernel BPF program. It runs on each frame *after* gopacket has decoded it, so it saves the JS-callback dispatch and object-conversion cost for non-matching packets but does **not** avoid the kernel→userspace copy (a real kernel filter would). A malformed expression makes `open` / `openFile` **reject**.

Limitations. No live capture on Windows. The `filter` grammar is the tcpdump subset above; for anything beyond it, filter inside the callback on the decoded fields. Common-layer decode only (Ethernet / IPv4 / IPv6 / TCP / UDP / ICMP); exotic protocols surface only as raw bytes. At high packet rates frames backpressure and drop at the kernel, exactly like `tcpdump`.

The pcap file API round-trips: write raw frames with `toFile`, read them back decoded with `openFile`, no privileges required (see `examples/scripts/capture-file.ts`).

Raw packets (`net.raw`)

Craft and send raw IPv4 packets and receive the replies — for SYN scans, RST/port-state inference, custom-TTL probes, and exotic-packet testing. **Linux + macOS only; needs root / CAP_NET_RAW** (Windows rejects). Two bindings:

- `net.raw.open({ iface?, filter?, readBuffer? })` — open a raw packet engine. `send(spec | Uint8Array)` crafts and fires a packet: a structured spec `{ dst, dstPort?, srcPort?, src?, proto?: "tcp"|"udp"|"ip", protocol?, flags?: string[], seq?, ack?, window?, ttl?, ipId?, payload? }` (full source-IP / TTL / IP-ID control; `sercon` computes checksums), or a `Uint8Array` holding a complete IPv4 packet sent verbatim. `onPacket(cb)` delivers decoded replies (the **same decoded packet object** as `net.capture`); `onClose` / `onError` / `close()` (returns a `Promise`) round out the handle. `iface` defaults to the auto-detected default-route interface — set it explicitly on multi-homed hosts; `filter` is the same tcpdump-like expression as `net.capture`. Default `flags ["SYN"]`, `ttl 64`, `window 65535`, `src` = the egress interface IP, `srcPort` = a random high port.
- `net.raw.tcp(host, port, opts?)` — one-shot: send a single crafted TCP segment (default a SYN) and resolve with the first reply correlated by the 4-tuple, or `null` on timeout. SYN → SYN/ACK = open, RST = closed, null = filtered/no answer. `opts { flags?, srcPort?, src?, seq?, ttl?, payload?, timeout? (ms, default 2000), iface? }`.

Kernel-RST caveat: because the host kernel does not own the crafted connection, it may answer the peer's SYN/ACK with its own RST. The probe still captures the reply; only the target's connection state is perturbed. To probe silently, add a host firewall rule dropping the outbound RST (`iptables -A OUTPUT -p tcp --tcp-flags RST RST -j DROP`, or the `pf` equivalent). `sercon` does not modify your firewall.

db

Database / KV / directory clients:

- `db.sqlite.open(path)` — embedded SQLite (cgo-free driver).
- `db.postgres.open(dsn | opts)` — PostgreSQL via pure-Go `pgx`.
- `db.mysql.open(dsn | opts)` — MySQL / MariaDB via pure-Go `go-sql-driver`.
- `db.mssql.open(dsn | opts)` — Microsoft SQL Server via pure-Go `go-mssqldb`.
- `db.clickhouse.open(dsn | opts)` — ClickHouse via pure-Go `clickhouse-go v2` (`opts.secure` for TLS).
- `db.oracle.open(dsn | opts)` — Oracle via pure-Go `go-ora` (`opts.database` is the service name).
- `db.redis.open(addr, opts?)` — Redis client.
- `db.valkey.open(url, opts?)` — Valkey client (the RESP-compatible Redis fork; same `{do, ping, close}` handle, accepts `valkey://` / `valkeys://` as well as `redis://`).
- `db.memcached.open(addr)` — Memcached client.
- `db.ldap.open(url, opts?)` — LDAP client (search, bind).
- `db.dict.{define, match}` — local dictionary lookup / fuzzy match.

SQL engines (`sqlite` / `postgres` / `mysql` / `mssql`) share one handle: `open()` resolves to `{exec, query, queryValue, begin, prepare, close}` (transactions via `begin()` → `{exec, query, queryValue, commit, rollback}`; prepared statements via `prepare(sql)` → `{exec, query, queryValue, close}`). The server engines accept either a driver **DSN string** or a connection **options object** (`{host, port, user, password, database}`, plus `sslmode` for postgres); credentials in the assembled URL DSNs are percent-escaped. The connection is pinged on `open()` so a bad DSN or unreachable server fails there, not at the first query. Bind parameters are positional and passed after the SQL string — write your engine's placeholder syntax: `?` (sqlite, mysql), `$1` (postgres), `@p1` (mssql). All four drivers are pure Go (no cgo). Scripts must `close()` the handle (and commit/rollback transactions, close prepared statements) — there is no GC finalizer.

Integration testing. The `db.*` tests above are unit-level. To exercise the networked bindings against real servers, point the `SERCON_TEST_*` env vars at a running database and run the `Integration` tests; unset vars skip, so `make test` / CI stay green without any server. `make test-integration` does it all against the [dbplayground](#) fleet: it brings the fleet up, runs the tests (postgres / mysql / mariadb / clickhouse / redis / valkey / memcached / ldap), and tears it down. Requires Docker + Compose v2.

services

Subprocess and external-CLI / service wrappers (was `tools` in v0.8.0). Members:

- `services.exec.shell(cmd, opts?)` — run a shell command; captures stdout/stderr or streams into a `tui` pane when `opts.pane` is set.
- `services.exec.stream(cmd, onLine, opts?)` — like `exec.shell` but streams the subprocess's output to `onLine(line, stream)` **line by line as it arrives** instead of buffering it. `cmd` and `opts` (`cwd` / `env` / `stdin` / `timeout`) match `exec.shell`, except `timeout` has **no default** (0 / absent = run until exit), since streaming targets long-running output. `stream` is `"stdout"` or `"stderr"`. Returns a Promise resolving to `{exitCode, success, durationMs}` on exit; a non-zero exit resolves with `success: false`, while spawn failures

and timeouts reject. Useful for processing large or incremental output without holding it all in memory.

- `services.exec.http(method, url, opts?)` — shell-level `curl` -style HTTP (separate from `net.http`, intended for raw protocol use).
- `services.git.*` — git porcelain wrappers (clone, status, commit, log, ...) invoking the local `git` binary.
- `services.gh.{authStatus, prList, repoView}` — thin `gh` CLI wrappers.
- `services.ai.{providers, send}` — provider-agnostic AI chat client.
- `services.agentBrowser.*` — headless-Chrome automation via the `agent-browser` CLI. See the subsection below.
- `services.webdriver.*` — W3C WebDriver client (via `tebeka/selenium`). See the subsection below.

`services.agentBrowser`

Bridge to the `agent-browser` headless-Chrome CLI. Every call shells out to `agent-browser --json --session <name> <command>` and returns the parsed JSON response.

Availability gate. `services.agentBrowser.available` is a sync boolean — check it before calling anything else. Every binding throws a clean error if the CLI is absent.

Launch model. `services.agentBrowser.launch(opts?)` is **synchronous**; no browser starts until the first command. It returns a handle whose I/O methods are all async (Promises). Sessions the script doesn't `close()` are best-effort closed when the Run ends.

Launch options (all optional):

Key	CLI flag	Description
<code>session</code>	<code>--session</code>	Name the session (auto-generated when omitted).
<code>headed</code>	<code>--headed</code>	Run with visible browser window.
<code>profile</code>	<code>--profile</code>	Browser profile directory.
<code>proxy</code>	<code>--proxy</code>	Proxy URL.
<code>userAgent</code>	<code>--user-agent</code>	Custom user-agent string.
<code>device</code>	<code>--device</code>	Emulate a device (e.g. "iPhone 12").
<code>colorScheme</code>	<code>--color-scheme</code>	"light" or "dark".
<code>ignoreHttpsErrors</code>	<code>--ignore-https-errors</code>	Skip TLS verification.
<code>engine</code>	<code>--engine</code>	Browser engine to use.
<code>executablePath</code>	<code>--executable-path</code>	Path to the browser binary.
<code>enable</code>	<code>--enable</code>	Enable a browser feature flag.
<code>args</code>	<code>--args</code>	Extra browser arguments.

Key	CLI flag	Description
<code>timeout</code>	<i>(sercon)</i>	Per-call subprocess timeout in ms (default 30 000, 0 disables). When the limit is reached the Promise rejects with a clear "timed out" message instead of hanging. Session <code>close()</code> is independently bounded at 10 s regardless of this value.

Response envelope. Every async method returns a parsed JSON object with the envelope `{ success: boolean, data: {...}, error: string | null }`. Drill into `.data` for the actual values (e.g. `data.title`, `data.value`, `data.visible`).

Handle method groups:

- **Navigation:** `open(url)`, `back()`, `forward()`, `reload()`, `wait(msOrSelector)`, `connect(portOrUrl)`
- **Interaction:** `click(sel)`, `dblclick(sel)`, `hover(sel)`, `focus(sel)`, `fill(sel, text)`, `type(sel, text)`, `press(key)`, `check(sel)`, `uncheck(sel)`, `select(sel, ...values)`, `scroll(dir, px?)`, `scrollIntoView(sel)`, `drag(src, dst)`, `upload(sel, files)`, `download(sel, path)`, `keyboard.{type, insertText}`, `mouse.{move, down, up, wheel}`
- **Inspection:** `get(what, sel?)`, `isVisible(sel)`, `isEnabled(sel)`, `isChecked(sel)`, `eval(jsCode)`, `snapshot(opts?)`, `console(opts?)`, `errors(opts?)`, `highlight(sel)`
- **Locators:** `find(locator, value, { action, text? })`, `locator(spec)` — returns a chainable handle bound to a `(locator, value)` spec.
- **Settings (Phase 2):** `set.viewport(w,h,scale?)`, `set.device(name)`, `set.geo(lat,lng)`, `set.offline(on?)`, `set.headers(obj)`, `set.credentials(user,pass)`, `set.media(scheme?,reducedMotion?)`
- **Recording (Phase 2):** `record.start(path.webm, url?)`, `record.stop()`
- **Capture (Phase 2):** `screenshot(path?,opts?)`, `pdf(path?)`
- **Network (Phase 3):** `network.route(url, opts?)`, `network.unroute(url?)`, `network.requests(opts?)`, `network.request(id)`, `network.har.start(path?)`, `network.har.stop(path?)`
- **Cookies (Phase 3):** `cookies.get()`, `cookies.set(name, value, opts?)`, `cookies.clear()`
- **Storage (Phase 3):** `storage.local.get(key?)`, `storage.local.set(key, value)`, `storage.local.clear()` and matching `storage.session.*`
- **Tabs (Phase 3):** `tabs.list()`, `tabs.new(url?, opts?)`, `tabs.close(ref?)`, `tabs.select(ref)`
- **Diff (Phase 3):** `diff.snapshot(opts?)`, `diff.screenshot(opts)`, `diff.url(url1, url2)`
- **Debug/Perf (Phase 4):** `trace.start()`, `trace.stop(path?)`, `profiler.start(opts?)`, `profiler.stop(path?)`, `inspect()`, `clipboard(op, text?)`, `vitals(url?)`, `pushstate(url)`
- **React DevTools (Phase 4):** `react.tree()`, `react.inspect(id)`, `react.renderers.start()`, `react.renderers.stop()`, `react.suspense(opts?)` — requires `launch({ enable: "react-devtools" })`
- **Streaming (Phase 4):** `stream.enable(opts?)`, `stream.disable()`, `stream.status()`
- **AI (Phase 4):** `chat(message, opts?)` — requires an AI gateway configured in agent-browser
- **Escape hatch (Phase 4):** `cmd(command, ...args)`, `batch(cmds, opts?)`

- **Auth login (Phase 4):** `auth.login(name)` — uses a saved auth profile on the current session
- **Lifecycle:** `session` (read-only name), `close()`

Argument vocabularies (the agent-browser surface). The string arguments to `get`, the `is*` checks, and `find` are passed through to the `agent-browser` CLI verbatim, so the accepted values are agent-browser's, not sercon's. The set below is current for **agent-browser 0.27.1** (the version this manual was written against — see the authoritative-reference note at the end of this section):

- **`get(what, sel?)`** — `what` is one of `text`, `html`, `value`, `attr` (the attribute name goes in the `sel` slot, e.g. `get("attr", "href")`), `title`, `url`, `count`, `box` (bounding box), `styles`, or `cdp-url`. When `sel` is a plain CSS selector it scopes the query; a snapshot `@ref` (see below) targets a specific element; omit it to query the page or the active element. The result is in the `.data` field of the envelope (e.g. `(await b.get("title")).data.title`).
- **`isVisible(sel)` / `isEnabled(sel)` / `isChecked(sel)`** — map to agent-browser's `isVisible|enabled|checked <selector>`. `sel` is required. The boolean lands in `.data` (e.g. `.data.visible`).
- **`find(locator, value, { action, text? })`** — `locator` is one of `role`, `text`, `label`, `placeholder`, `alt`, `title`, `testid`, `first`, `last`, or `nth`; `value` is the locator's argument (the role/text/etc. to match). `action` is **required** and is an act-style verb — `click`, `dblclick`, `hover`, `focus`, `check`, `uncheck`, `fill`, or `type` (the last two consume `text`). agent-browser's `find` cannot return a handle without acting, so there is no read-only `find`; for read-only matching use `snapshot()`. The chainable `locator(spec)` handle exposes the same verbs as methods.
- **`snapshot(opts?)`** — `opts` is `{ interactive?: boolean, compact?: boolean, depth?: number, selector?: string }`, mapping to agent-browser's `-i / -c / -d <n> / -s <sel>`. The snapshot is an accessibility tree whose interactive nodes carry refs like `@e2`; **those refs are valid selectors** for `get`, `click`, `fill`, etc. on the same session — the snapshot → act-by-ref loop is agent-browser's intended workflow for AI agents.
- **`console(opts?)` / `errors(opts?)`** — `opts` is `{ clear?: boolean }` (passes `--clear`); returns captured console logs / page errors.

Reaching the rest of agent-browser. sercon binds the commonly-used subset; agent-browser ships more (sessions, dashboard, confirmation prompts, iOS-simulator and cloud-provider backends, init scripts, ...). Any subcommand that isn't a first-class binding is reachable through the `cmd(command, ...args)` escape hatch — `cmd` maps 1:1 onto an `agent-browser <command> [args]` invocation. For the **authoritative, version-matched** command and flag reference — including data shapes returned in `.data`, which evolve with the CLI — run `agent-browser --help`, or the agent-oriented guide `agent-browser skills get core --full`. Treat the tables in this manual as a convenience snapshot, not a contract: agent-browser owns its own surface.

Capture (Phase 2). `screenshot` and `pdf` are path-first with opt-in in-memory bytes:

- **`b.screenshot(path?, opts?)`** — when `path` is given the image is written to that file and the result is `{ path: string, size: number, format: string }`. When omitted the raw bytes are returned as `{ bytes: number[], format: string }` — a plain JS `number[]` (byte-value array). Convert to a typed array with `new Uint8Array(bytes)`.
- **`b.pdf(path?)`** — same path-first model; format is always `"pdf"`.
- **`opts` for `screenshot`:** `{ selector?, full?, annotate?, format?: "png"|"jpeg", quality? }`.

Defaults bag (Phase 2). A per-Run namespace-level defaults object is merged under per-call `launch()` opts so common flags don't need to be repeated:

```
services.agentBrowser.setDefaultOptions({ headed: false, proxy:
"http://proxy:3128" });
const b = services.agentBrowser.launch(); // inherits headed + proxy
```

- `setDefaultOptions(obj)` — replace the defaults map entirely.
- `defaultOptions()` — return a shallow copy of the current defaults.
- `clearDefaultOptions()` — reset to `{}`.

Defaults are scoped to the current Run; they do not persist across `Run` calls.

Flat one-shot shortcuts (Phase 2). Launch an ephemeral session, open a URL, perform one action, and close — without managing a handle:

Shortcut	Equivalent handle sequence
<code>agentBrowser.screenshot(url, path?, opts?)</code>	<code>launch()+open(url)+screenshot(path?,opts?)+close()</code>
<code>agentBrowser.pdf(url, path?)</code>	<code>launch()+open(url)+pdf(path?)+close()</code>
<code>agentBrowser.snapshot(url, opts?)</code>	<code>launch()+open(url)+snapshot(opts?)+close()</code>
<code>agentBrowser.eval(url, js)</code>	<code>launch()+open(url)+eval(js)+close()</code>

Example (Phase 1 + Phase 2):

```

if (!services.agentBrowser.available) {
  runtime.log("install agent-browser first");
} else {
  const html = "data:text/html," + encodeURIComponent(
    "<title>sercon demo</title><h1 id=hi>Hello</h1><input id=box>"
  );
  const b = services.agentBrowser.launch({ headed: false });
  try {
    await b.open(html);

    const titleRes = await b.get("title");
    runtime.log("title:", titleRes.data?.title);    // "sercon demo"

    await b.fill("#box", "typed by sercon");
    const valRes = await b.get("value", "#box");
    runtime.log("value:", valRes.data?.value);      // "typed by sercon"

    const visRes = await b.isVisible("#hi");
    runtime.log("visible:", visRes.data?.visible); // true

    const snap = await b.snapshot({ compact: true });
    runtime.log("snap keys:", Object.keys(snap).join(", "));

    // Phase 2 – settings + capture
    await b.setViewport(1280, 800);
    const shot = await b.screenshot({ full: true }); // bytes path
    runtime.log("bytes:", new Uint8Array(shot.bytes).length, shot.format);

    const file = await b.screenshot("/tmp/demo.png"); // file path
    runtime.log("file:", JSON.stringify(file));      // { path, size, format }
  } finally {
    await b.close();
  }

  // One-shot shortcut
  const r = await services.agentBrowser.eval(html, "document.title");
  runtime.log("title:", r.data?.result); // "sercon demo"
}

```

Phase 3: network, cookies, storage, tabs, diff.

All Phase-3 groups are handle methods (operate on an open session) and return the standard `{ success, data, error }` envelope.

Network interception (`network.*`).

```

// Block all XHR requests matching a glob pattern.
await b.network.route("**/api/*", { abort: true });

// Stub a JSON response body instead of blocking.
await b.network.route("**/data.json", { body: { mock: true } });

// Remove a specific route (or all routes when url is omitted).
await b.network.unroute("**/api/*");

// Inspect captured requests.
const reqs = await b.network.requests({ clear: true, method: "POST" });
runtime.log(JSON.stringify(reqs.data));

// Get a single request by id.
const req = await b.network.request("req-1");

// Record a HAR trace.
await b.network.har.start("/tmp/trace.har");
// ... navigate and interact ...
await b.network.har.stop();

```

Route opts: abort?: boolean, body?: unknown (JSON-serialised), resourceType?: string (e.g. "xhr").

Requests opts: clear?: boolean, filter?: string, type?: string, method?: string, status?: string.

Cookie management (cookies.*).

```

// Read the current cookie jar.
const jar = await b.cookies.get();
runtime.log(jar.data?.cookies); // array of cookie objects

// Set a cookie (requires a real HTTP origin – not available on data: URLs).
await b.cookies.set("session_id", "abc123", {
  domain: ".example.com",
  httpOnly: true,
  secure: true,
  sameSite: "Lax",
  expires: Math.floor(Date.now() / 1000) + 3600, // Unix seconds
});

// Clear all cookies.
await b.cookies.clear();

```

cookies.set opts: url?, domain?, path?, httpOnly?, secure?, sameSite?: "Strict" | "Lax" | "None", expires? (Unix seconds).

Note: cookies.set and storage.* require a real HTTP origin. They will fail on data: URLs —

wrap in `try-catch` when testing with data: URLs.

Web storage (`storage.*`).

```
// localStorage round-trip (requires a real HTTP origin).
await b.storage.local.set("theme", "dark");
const val = await b.storage.local.get("theme");
runtime.log(val.data);           // { theme: "dark" } or similar
await b.storage.local.clear();

// sessionStorage works identically under storage.session.
await b.storage.session.set("token", "xyz");
const tok = await b.storage.session.get("token");
await b.storage.session.clear();
```

`get(key?)` — omit `key` to get all entries; pass a key to get one.

Tab management (`tabs.*`).

```
// Open a new tab (optionally at a URL, with an optional label).
await b.tabs.new("https://docs.example.com", { label: "docs" });

// List all open tabs. Refs are t1/t2/... or the user-supplied label.
const list = await b.tabs.list();
runtime.log(list.data?.tabs); // [{ tabId, title, url, active, label, type }]

// Switch the active tab.
await b.tabs.select("docs"); // by label
await b.tabs.select("t1");   // by ref

// Close a specific tab (or the current tab when ref is omitted).
await b.tabs.close("docs");
```

Page diffing (`diff.*`).

```

// Snapshot the current DOM state (baseline snapshot when no -b given).
const snap = await b.diff.snapshot();

// Compare against a saved baseline snapshot.
const delta = await b.diff.snapshot({ baseline: "/tmp/base.json", compact: true
});
runtime.log(delta.data);

// Pixel-level screenshot comparison.
const px = await b.diff.screenshot({
  baseline: "/tmp/base.png",
  output: "/tmp/diff.png",
  threshold: 0.1,           // 0-1; pixels above this are flagged
});

// Compare two URLs side-by-side without an open session.
const cmp = await b.diff.url("https://staging.example.com",
"https://prod.example.com");

```

`diff.snapshot` opts: `baseline?: string` (path to a prior snapshot JSON), `selector?: string` (scope to a subtree), `compact?: boolean`, `depth?: number`.

Phase 4: debug/perf, React DevTools, streaming, AI chat, escape hatch, auth vault.

All Phase-4 groups are handle methods (operate on an open session) unless noted. They return the standard `{ success, data, error }` envelope; `batch` returns an array of per-command envelopes.

Debug / perf (`trace.*`, `profiler.*`, `inspect`, `clipboard`, `vitals`, `pushstate`).

```

// Chrome DevTools tracing.
await b.trace.start();
// ... navigate and interact ...
await b.trace.stop("/tmp/trace.json");

// V8 CPU profiling (optional category filter).
await b.profiler.start({ categories: "v8,blink" });
await b.profiler.stop("/tmp/profile.json");

// Open the Chrome DevTools inspector and get the DevTools URL.
const devtools = await b.inspect();
runtime.log("DevTools:", devtools.data?.url);

// Clipboard read/write.
await b.clipboard("write", "copied text");
const clip = await b.clipboard("read");
runtime.log(clip.data?.text);

// Core Web Vitals for the current page (or a given URL).
const v = await b.vitals();
runtime.log("LCP:", v.data?.lcp);

// SPA client-side navigation without a full page reload.
await b.pushstate("/app/dashboard");

```

`profiler.start` opts: { categories?: string } — a comma-separated list of V8/Blink profiling categories.

React DevTools (`react.*`).

Requires the session to have been launched with `launch({ enable: "react-devtools" })`. `agent-browser` returns a clear error (surfaced as a throw) when `react-devtools` was not enabled at launch time.

```

const b = services.agentBrowser.launch({ enable: "react-devtools" });
await b.open("https://my-react-app.example.com");

// Get the component tree.
const tree = await b.react.tree();
runtime.log(JSON.stringify(tree.data).slice(0, 200));

// Inspect a specific fiber node by id (id from tree.data).
const node = await b.react.inspect("fiber-123");

// Measure re-renders.
await b.react.renders.start();
// ... interact ...
const renders = await b.react.renders.stop();
runtime.log(JSON.stringify(renders.data));

// List Suspense boundaries (optionally only dynamic ones).
const suspense = await b.react.suspense({ onlyDynamic: true });

```

`react.suspense` opts: { `onlyDynamic?: boolean` } — when true, only boundaries that have actually suspended are included.

Live streaming (`stream.*`).

```

// Enable a browser-event stream on a chosen port.
await b.stream.enable({ port: 9229 });

const status = await b.stream.status();
runtime.log("streaming:", status.data?.enabled);

await b.stream.disable();

```

`stream.enable` opts: { `port?: number` } — selects the streaming port (agent-browser picks a default when omitted).

AI chat (`chat`).

Requires an AI gateway to be configured in agent-browser; errors cleanly if not configured.

```

// Single-shot natural-language instruction.
const r = await b.chat("Click the Login button");
runtime.log("chat ok:", r.success);

// Optionally specify the model.
const r2 = await b.chat("Fill out the form with test data", { model: "gpt-4o" });

```

Escape hatch (`cmd` / `batch`).

`cmd` is a generic passthrough for any agent-browser subcommand. `batch` runs multiple commands in a single round-trip and returns an **array** of per-command result envelopes (not the usual single

envelope).

```
// cmd: call any agent-browser subcommand.
const title = await b.cmd("get", "title");
runtime.log("title:", title.data?.title);

// batch: multiple commands, one round-trip.
const results = await b.batch(["get title", "get url"], { bail: false });
for (const r of results) {
  runtime.log(r.success, JSON.stringify(r.data));
}
```

`batch` opts: { `bail?: boolean` } — stop at the first failing command and return results up to that point.

Auth vault — namespace (`auth.save` / `list` / `show` / `delete`) + handle (`auth.login`).

Vault CRUD is session-independent and lives on the namespace (`services.agentBrowser.auth.*`).
`login` fills a form on a live page and is a handle method (`b.auth.login(name)`).

Passwords are never placed in `argv`. `auth.save` feeds the password through the subprocess stdin using `--password-stdin`. This means the password never appears in `ps` output or shell history.

```
// Save a login profile (password is sent via stdin, not argv).
await services.agentBrowser.auth.save("prod", {
  url: "https://app.example.com/login",
  username: "admin",
  password: "s3cret",           // sent via stdin
  usernameSelector: "#user",
  passwordSelector: "#pass",
  submitSelector: "#submit",
});

// List all saved profiles (returns an envelope; data contains profile names).
const list = await services.agentBrowser.auth.list();
runtime.log(JSON.stringify(list.data));

// Show details of a saved profile.
const info = await services.agentBrowser.auth.show("prod");

// Delete a profile.
await services.agentBrowser.auth.delete("prod");

// Log in using a saved profile (handle method — requires an open session).
const b = services.agentBrowser.launch();
await b.open("https://app.example.com/login");
await b.auth.login("prod");    // fills the form and submits
runtime.log("logged in");
await b.close();
```

`auth.save opts: { url, username, password, usernameSelector?, passwordSelector?, submitSelector? }`. All string opts except `password` map to the corresponding agent-browser flags; `password` is stripped from the option object and fed via stdin.

`services.webdriver`

W3C WebDriver client backed by github.com/tebeka/selenium (pure Go — no cgo). Drives a real browser via an installed `chromedriver` or `geckodriver`; no driver binary is bundled or downloaded.

Availability gate. `services.webdriver.available` is a sync boolean — true when `chromedriver` or `geckodriver` is found on PATH. Gate all calls on this value.

Connect-or-start model. `connect(opts?)` is async. When `opts.url` is given, it dials the running driver endpoint at that URL. When `opts.url` is omitted, it locates the driver binary on PATH, picks an ephemeral port, starts the driver as a subprocess, and dials it. The resolved session handle has the same surface in both cases. Sessions quit (and any started subprocess is stopped) when the Run ends; call `session.quit()` explicitly for deterministic teardown.

Probe. `probe({ url })` checks whether a WebDriver endpoint is ready without opening a session. Returns `{ ready: boolean, status?: number }` on HTTP success or `{ ready: false, error: string }` on transport failure; never throws on network errors.

Connect options (all optional):

Key	Type	Default	Description
<code>browser</code>	<code>"chrome" "firefox"</code>	<code>"chrome"</code>	Selects the driver binary (<code>chromedriver</code> or <code>geckodriver</code>).
<code>headless</code>	<code>boolean</code>	<code>true</code>	Run the browser in headless mode.
<code>url</code>	<code>string</code>	—	Dial an already-running WebDriver server instead of starting one.
<code>args</code>	<code>string[]</code>	—	Extra browser command-line flags appended to the default set.
<code>capabilities</code>	<code>object</code>	—	Raw W3C capability overrides merged last (escape hatch).
<code>commandTimeout</code>	<code>number</code>	<code>30000</code>	Per low-level WebDriver request timeout (ms). Bounds each command so a driver blocked behind an open alert or an unreachable endpoint can't hang the call. <code>0</code> /negative disables it. <code>quit()</code> (and Run-end cleanup) also cancels any in-flight command.

Browser flags (`args`) and capabilities (the WebDriver surface). Unlike `agentBrowser`, this binding speaks the W3C WebDriver protocol directly (via `tebeka/selenium`) rather than wrapping a CLI, so the "external surface" is the W3C capability set and the browser's own command-line flags:

- **`args`** are the browser's launch flags, appended *after* the headless flag `sercon` adds for you (`--headless=new` for Chrome, `-headless` for Firefox when `headless` is true). They land in the

vendor capability block (`goog:chromeOptions.args` / `moz:firefoxOptions.args`). Common Chrome flags: `--no-sandbox` , `--disable-gpu` , `--disable-dev-shm-usage` , `--window-size=1280,800` , `--lang=en-US` , `--proxy-server=host:port` . Firefox takes `-width 1280` , `-height 800` , `-private` , etc.

- **capabilities** is an object of top-level **W3C standard capabilities**, merged into the session request *last* (so it overrides sercon's defaults). Standard keys: `browserName` , `browserVersion` , `platformName` , `acceptInsecureCerts` (boolean — accept self-signed/expired TLS), `pageLoadStrategy` ("normal" | "eager" | "none"), `unhandledPromptBehavior` ("dismiss" | "accept" | "ignore" | ...), `timeouts` ({ `script` , `pageLoad` , `implicit` } in ms), and `proxy` ({ `proxyType` , `httpProxy` , `sslProxy` , ... }). Vendor blocks go here too — `goog:chromeOptions` and `moz:firefoxOptions` (each { `args?` , `binary?` , `prefs?` , ... }). Because the merge is by top-level key, supplying `goog:chromeOptions` wholesale *replaces* the one built from `args` / `headless` — set `args` for additive flags, `capabilities` only when you need full control. See the [W3C WebDriver capabilities](#) spec for the complete list.

Session handle methods (all async):

Method	Returns	Description
<code>get(url)</code>	{ <code>ok: true</code> }	Navigate to a URL.
<code>url()</code>	string	Current URL.
<code>title()</code>	string	Current page title.
<code>back()</code> / <code>forward()</code> / <code>refresh()</code>	{ <code>ok: true</code> }	Browser history controls.
<code>find(by, value)</code>	element handle	Find the first matching element.
<code>findAll(by, value)</code>	element handle[]	Find all matching elements.
<code>source()</code>	string	Current page source HTML.
<code>screenshot(path?)</code>	{ <code>path?</code> , <code>size?</code> , <code>bytes?</code> , <code>format</code> }	Capture the viewport. No path → <code>bytes: number[]</code> ; with path → { <code>path</code> , <code>size</code> , <code>format</code> } .
<code>executeScript(js, args?)</code>	unknown	Run a synchronous script in the page context.
<code>executeScriptAsync(js, args?)</code>	unknown	Run an async script (callback-resolved).
<code>cookies()</code>	object[]	Get all cookies.
<code>setCookie(c)</code>	{ <code>ok: true</code> }	Set a cookie ({ <code>name</code> , <code>value</code> , <code>path?</code> , <code>domain?</code> , <code>secure?</code> , <code>httpOnly?</code> , <code>expiry?</code> }).
<code>deleteCookie(name)</code>	{ <code>ok: true</code> }	Delete a named cookie.
<code>deleteAllCookies()</code>	{ <code>ok: true</code> }	Clear all cookies.
<code>setImplicitWait(ms)</code>	{ <code>ok: true</code> }	Set the implicit element-wait timeout.

Method	Returns	Description
<code>waitFor(by, value, opts?)</code>	element handle	Poll until an element appears (opts: <code>timeout? ms</code> , <code>visible?: boolean</code>).
<code>quit()</code>	<code>{ closed: true }</code>	Close the session (idempotent; also called by Run-end cleanup).

Locator strategies (by argument): `css`, `xpath`, `id`, `name`, `tag`, `className`, `linkText`, `partialLinkText`.

Element handle methods (all async):

Method	Description
<code>click()</code>	Click the element.
<code>sendKeys(text)</code>	Type text into the element.
<code>clear()</code>	Clear the element's value.
<code>submit()</code>	Submit the form.
<code>text()</code>	Inner text content.
<code>getAttribute(name)</code>	Get the named HTML attribute ; returns <code>null</code> when the attribute is absent (W3C semantics). This is the <i>attribute</i> , not the live DOM <i>property</i> — an <code><input></code> 's typed text lives in its <code>value</code> property, not a <code>value</code> attribute, so read it with <code>executeScript</code> (e.g. <code>executeScript("return document.querySelector('#box').value")</code>). Chrome's legacy fallback returns the property; Firefox correctly returns <code>null</code> .
<code>cssValue(name)</code>	Get a computed CSS property value.
<code>tagName()</code>	Tag name in lower case.
<code>isDisplayed()</code> / <code>isEnabled()</code> / <code>isSelected()</code>	Visibility / state checks.
<code>find(by, value)</code> / <code>findAll(by, value)</code>	Sub-tree element search.
<code>screenshot(path?)</code>	Screenshot the element (scrolls into view).

Example:

```

if (!services.webdriver.available) {
  runtime.log("no chromedriver/geckodriver on PATH – skip");
} else {
  const d = await services.webdriver.connect({ browser: "chrome", headless: true
});
  try {
    await d.get("data:text/html," + encodeURIComponent(
      "<title>wd demo</title><h1 id=hi>Hello</h1><input id=box>"));
    runtime.log("title:", await d.title());
    const h1 = await d.find("id", "hi");
    runtime.log("text:", await h1.text(), "visible:", await h1.isDisplayed());
    const box = await d.find("css", "#box");
    await box.sendKeys("typed by sercon");
    runtime.log("id attr:", await box.getAttribute("id")); // "box"
    // Live input value is a DOM property, not an attribute – read via script
    // so it works the same on Chrome and Firefox.
    runtime.log("value:", await d.executeScript("return
document.querySelector('#box').value", []));
    runtime.log("eval:", await d.executeScript("return 6*7", []));
    const shot = await d.screenshot();
    runtime.log("bytes:", new Uint8Array(shot.bytes).length, shot.format);
  } finally {
    await d.quit();
  }
}

```

Phase 2 (v0.41.0). The session handle gains the following additional async methods.

Phase 2 session methods:

Method	Returns	Description
<code>windowHandles()</code>	<code>string[]</code>	All open window/tab handles.
<code>currentWindow()</code>	<code>string</code>	Handle of the currently focused window/tab.
<code>switchToWindow(handle)</code>	<code>{ ok: true }</code>	Focus a window/tab by its handle.
<code>newWindow(type?)</code>	<code>{ handle, type }</code>	Open a new tab ("tab" , default) or window ("window"). Returns the new handle.
<code>closeWindow()</code>	<code>string[]</code>	Close the current window/tab; returns the remaining handles and auto-switches to a survivor.
<code>switchToFrame(target)</code>	<code>{ ok: true }</code>	Switch into a frame by index (number) or by passing an element handle for the <code><iframe></code> . A name/id string is not accepted — locate the iframe via

Method	Returns	Description
		<code>find</code> and pass the element handle.
<code>switchToParentFrame()</code>	<code>{ ok: true }</code>	Switch to the parent frame.
<code>switchToDefaultContent()</code>	<code>{ ok: true }</code>	Switch back to the top-level browsing context.
<code>acceptAlert()</code>	<code>{ ok: true }</code>	Accept (OK/confirm) the current alert/confirm/prompt dialog.
<code>dismissAlert()</code>	<code>{ ok: true }</code>	Dismiss (cancel) the current dialog.
<code>alertText()</code>	string	Get the message text of the current dialog.
<code>sendAlertText(text)</code>	<code>{ ok: true }</code>	Type into a prompt dialog.
<code>maximize()</code>	<code>{ x, y, width, height }</code>	Maximize the browser window; returns the new rect.
<code>minimize()</code>	<code>{ x, y, width, height }</code>	Minimize the browser window; returns the new rect.
<code>fullscreen()</code>	<code>{ x, y, width, height }</code>	Enter fullscreen; returns the new rect.
<code>setWindowRect({width?,height?,x?,y?})</code>	<code>{ x, y, width, height }</code>	Set position and/or size.
<code>getWindowRect()</code>	<code>{ x, y, width, height }</code>	Get the current window rect.
<code>hover(el)</code>	<code>{ ok: true }</code>	Move the mouse pointer over an element (W3C actions).
<code>dragAndDrop(src, dst)</code>	<code>{ ok: true }</code>	Drag from <code>src</code> element to <code>dst</code> element.
<code>keyChord(...keys)</code>	<code>{ ok: true }</code>	Send a key chord (e.g. <code>keyChord("Control","a")</code>).
<code>performActions(sequence)</code>	<code>{ ok: true }</code>	Send a raw W3C actions array (escape hatch).
<code>releaseActions()</code>	<code>{ ok: true }</code>	Release any held action state (buttons/keys).

Element handle additions (Phase 2): Every element handle now carries an `elementId` string property (the raw W3C element reference id). Two new async methods are available: `el.hover()` moves the mouse over the element; `el.dragTo(target)` drags the element to a target element handle.

executeScript / executeScriptAsync — element handle return values: When a script returns a single W3C element reference, the binding wraps it in an element handle automatically. A top-level array of element references is also wrapped (each entry becomes an element handle). Nested references (e.g. an object with element-reference values) are returned as plain objects — return them flat or unwrap manually.

Caveat — alert unhandledPromptBehavior: chromedriver's default `unhandledPromptBehavior` is "dismiss and notify", which may auto-dismiss dialogs before your script can interact with them. To inspect alerts manually, pass `capabilities: { unhandledPromptBehavior: "ignore" }` in `connect()`. With "ignore", the alert is kept open but the renderer is blocked by the modal, so any WebDriver command that executes script in the page (including `executeScript`) will hang until the alert is dismissed. Trigger alerts via `<body onload="alert(...)">` or a synchronous inline script so that `get()` completes with the dialog already open — then `alertText()` and `acceptAlert()` work without racing against the modal state.

Caveat — file upload: For local `<input type="file">`, use `el.sendKeys("/absolute/path/to/file")` — this works with both Chrome and Firefox. Remote-grid file upload (the `/session/{id}/file` endpoint) is not supported by the underlying `tebeka/selenium` library and is out of scope.

Phase 2 example:

```

if (!services.webdriver.available) {
  runtime.log("no chromedriver/geckodriver on PATH – skip");
} else {
  const d = await services.webdriver.connect({
    browser: "chrome",
    headless: true,
    capabilities: { unhandledPromptBehavior: "ignore" },
  });
  try {
    // Windows / tabs
    const nw = await d.newWindow("tab");
    await d.switchToWindow(nw.handle);
    runtime.log("open tabs:", (await d.windowHandles()).length);
    await d.closeWindow();

    // Window rect
    await d.setWindowRect({ width: 1024, height: 768 });
    const r = await d.getWindowRect();
    runtime.log("rect:", r.width + "x" + r.height);

    // Frames
    await d.get("data:text/html," + encodeURIComponent(
      "<title>outer</title><iframe srcdoc='<p id=inner>hi</p>'></iframe>"));
    await d.switchToFrame(0);
    runtime.log("frame text:", await (await d.find("id", "inner")).text());
    await d.switchToParentFrame();
    await d.switchToDefaultContent();

    // Alerts (requires unhandledPromptBehavior: "ignore")
    // Navigate to a page that fires alert() synchronously on load; get()
    // blocks until the load event, so the dialog is open when it returns.
    try {
      await d.get("data:text/html," + encodeURIComponent(
        "<body onload='alert('hello')'><title>a</title>"));
    } catch (_) {
      // chromedriver may surface an unexpected-alert-open error – the alert
      // is still live, carry on.
    }
    runtime.log("alert:", await d.alertText());
    await d.acceptAlert();

    // Hover + drag
    await d.get("about:blank");
    await d.executeScript(`
      document.body.innerHTML =
        '<div id=a
style="position:absolute;top:50px;left:50px;width:80px;height:80px"></div>' +
        '<div id=b
style="position:absolute;top:50px;left:200px;width:80px;height:80px"></div>';
      return 1;`, []);
    const a = await d.find("id", "a");
    const b = await d.find("id", "b");

```

```

    await a.hover();
    await d.dragAndDrop(a, b);

    // executeScript returning an element handle
    const el = await d.executeScript("return document.getElementById('b')", []);
    runtime.log("script element tag:", await el.tagName());
  } finally {
    await d.quit();
  }
}

```

tui

Multi-pane terminal UI (was `ui.tui` in v0.8.0). `tui.layout(tree)` declares panes; `tui.pane(name)` returns a pane handle with `write`, `writeln`, `clear`, `title`. `services.exec.shell({pane})` streams subprocess I/O into a pane.

Scripts that orchestrate multiple subprocesses (e.g. `brew`, `npm`, `cargo` updates running in parallel) can declare a multi-pane terminal layout and route each subprocess's stdout/stderr into its own pane. Activation is **script-driven**: the first call to `tui.layout(...)` enters TUI mode. If stdout is not a TTY (CI, pipes, `make demo`), the same calls fall back to prefixed plain-text lines (`[paneName] line`) so the same script runs in both contexts.

Layout

```

tui.layout({
  rows: [
    { name: "log", title: "Orchestrator", weight: 1 },
    { cols: [
      { name: "brew" },
      { name: "npm" },
    ],
      weight: 2,
    },
  ],
});

```

Each layout node is one of:

- `{ name: string, title?: string, weight?: number }` — a leaf pane.
- `{ rows: LayoutNode[], weight?: number }` — a vertical split.
- `{ cols: LayoutNode[], weight?: number }` — a horizontal split.

`weight` defaults to 1. Pane names must be unique across the whole tree. `layout` may be called **once** per Run; a second call throws.

Pane handles

```
const log = tui.pane("log");
log.writeln("Updating Homebrew...");
log.title("Done");
log.clear();
log.write("partial line ");
```

`tui.pane(name)` returns a handle with `write`, `writeln`, `clear`, `title`. Methods are synchronous from the script's perspective — they enqueue to the TUI goroutine.

Subprocess routing

`services.exec.shell(cmd, opts)` accepts `opts.pane` (a Pane handle or a pane name string):

```
await services.exec.shell("brew update && brew upgrade", { pane: "brew" });
```

When set, the subprocess's stdout **and** stderr stream into the pane line by line. The returned `{ exitCode, durationMs, success }` is the same as without `pane`: except `stdout` and `stderr` are empty (data was streamed, not captured). ANSI colors are translated to `tview` color tags and rendered natively; `\r` without `\n` overwrites the current line so progress spinners render cleanly.

Pass `{ pty: true }` to run the command under a pseudo-terminal (Unix only). The child then believes it is a terminal and emits color, progress bars, and spinners, which a `pane` renders (or, without a pane, are captured into `stdout`). This is the general alternative to per-tool force-color flags (`FORCE_COLOR=1`, `--color=always`): it works for any tool that gates output on a TTY. A pty merges `stdout` and `stderr` onto one stream, so `stderr` is empty in pty mode. On Windows `pty` is ignored (pipe fallback, no color).

Keybindings (TTY mode only)

Key	Action
Tab / Shift-Tab	Cycle focus through panes
PgUp / PgDn	Scroll focused pane one page
↑ / ↓	Scroll focused pane one line
Home / End	Jump to top / bottom
Ctrl-C	Abort the script (single press)

The focused pane has a yellow border; the bottom status bar shows its name and the active keys.

Panes follow the tail automatically; scroll up (keys or, when enabled, the mouse wheel) to pause following, and scroll back to the bottom (or press `End`) to resume. Set `{ autoscroll: false }` on a leaf to keep it pinned at the top. Pass `{ mouse: true }` at the layout root to enable mouse support: the wheel scrolls the pane **under the cursor** (without changing which pane is focused), and a left-click focuses the pane under the cursor. This disables the terminal's native click-drag selection while active; hold Shift/Option to select.

Per-leaf `{ wrap: "char" | "word" | "off" }` controls line wrapping (default `"char"`; `"off"` lets long lines scroll horizontally), and `{ color: false }` strips a pane's ANSI to plain text instead of

rendering it. Both affect TTY rendering only — the non-TTY fallback always emits plain prefixed lines.

`tui.waitKey()` resolves with the next key (`{ name, rune, ctrl, alt, shift }`); `tui.onKey(handler)` registers a persistent per-keypress callback and returns an unsubscribe function. Both see every key except `Ctrl-C` (which always aborts) and coexist with the built-in navigation keys. Both require a TTY: `waitKey` rejects and `onKey` is a no-op in non-TTY (fallback) mode.

Limitations (v1)

- `tui` is **incompatible with** `--watch`: calling `tui.layout()` under `--watch` throws. Use one or the other.
- No runtime layout mutation, no input panes (`tui.input`), no snapshot-to-normal-screen on exit. The alt screen is restored at script end and the usual `PASS / FAIL` line prints.

6. Servers

The `server` global (added in v0.10.0) hosts inbound listeners — scripts that *accept* connections rather than make them. It ships `server.http` and `server.https` (§6.2), `server.smtp` (§6.7), and the raw `server.tcp` / `server.udp` listeners (§6.8); future cycles will add IMAP, FTP, and POP3 against the same engine foundation.

6.1 Foundation: `HoldRun` + `LoopCallable`

Two engine primitives back every long-lived binding. They live on `*scriptengine.Engine` (CLI use of them is internal; the public contract is "use these if you write a similar binding yourself").

`Engine.HoldRun(reason string) (release func())` — keeps `loop.Run` from exiting while a binding has resources outstanding. The event loop normally exits as soon as its `jobCount` drops to zero; `setTimeout` / `setInterval` / `setImmediate` are the only things that count. `HoldRun` parks a 24-hour sentinel timer under the hood (same trick `PromisifyAsync` uses for one-shot async work), returns a `release` function that clears it, and increments a refcount so multiple concurrent holders compose. `release` is idempotent. Any sentinels still parked when `Run` returns are cleanup-drained automatically as a safety net — a binding that forgets to call `release` does not leak into the next `Run`.

`scriptengine.NewLoopCallable(loop, fn)` — wraps a captured `goja.Callable` (a JS function reference) so it can be invoked from *any* goroutine. Goja's runtime is single-threaded; calling a `Callable` directly from a non-loop goroutine corrupts engine state. `LoopCallable.Call(buildArgs)` enqueues a `loop.RunOnLoop` callback, builds the arg list on the loop (where `vm.ToValue` is valid), invokes the callable, and returns the result to the caller's goroutine. `LoopCallable.CallOnLoop(vm, args...)` is the already-on-the-loop variant — using `Call` from inside a loop callback deadlocks because the new `RunOnLoop` job can't run until the current one returns.

You use them together. A typical server binding holds one or two `LoopCallable`s for the user's handlers, parks one `HoldRun` sentinel per active listener, and releases on close. The CLI bindings under `cmd/sercon/api_server*.go` follow this pattern.

Concurrency model. Because every handler invocation marshals back onto the single goja loop, **handlers serialize**: only one JS handler runs at a time, across all listeners. This is a feature (no JS data

ances, no shared-state confusion) and a throughput ceiling. For CPU-bound work in a handler, offload to a Go binding that does the work in a goroutine and resolves a Promise.

6.2 `server.http.listen` / `server.https.listen`

```
// HTTP
const srv = await server.http.listen({
  port: 8080, // required
  host: "0.0.0.0", // default
  routes: { "GET /": (req, res) => res.text("hi") },
  use: [logger], // optional global middleware
});

// HTTPS – identical plus cert/key (file paths OR inline PEM strings)
const srv2 = await server.https.listen({
  port: 8443,
  cert: "/etc/ssl/server.pem",
  key: "/etc/ssl/server.key",
  routes,
});

// HTTPS dev shortcut – ephemeral self-signed cert, no key, no files
const srv3 = await server.https.listen({
  port: 8443,
  cert: "self-signed",
  routes,
});

srv.address; // "tcp/0.0.0.0:8080"
await srv.close(); // graceful: stop accepting, drain, release HoldRun
await srv.stopped; // Promise that resolves when the listener exits
```

Options:

Key	Type	Notes
port	number	Required. Bind port. <code>0</code> lets the kernel pick; read the chosen port off <code>srv.address</code> .
host	string	Default <code>"0.0.0.0"</code> . Pass <code>"127.0.0.1"</code> for loopback-only.
routes	Record<string, RouteValue>	Required (may be <code>{}</code>). Keys are stdlib <code>http.ServeMux</code> patterns.
use	Middleware[]	Optional global middleware; runs in array order before any per-route middleware.
onError	(err, req, res) => unknown	Optional error handler; renders a custom response when a handler/middleware throws or rejects. See below.
cert	string	HTTPS only. File path, inline PEM, or the literal <code>"self-signed"</code> to mint an ephemeral in-process dev cert.

Key	Type	Notes
key	string	HTTPS only. File path or inline PEM. Not needed (ignored) when cert is "self-signed".

Custom error pages (`onError`). When a handler or middleware throws (or returns a rejected Promise) and `onError` is set, it is invoked as `onError(err, req, res)` instead of sending the stock 500 Internal Server Error. `err` is the thrown value (an `Error` for a plain `throw new Error(...)`); render whatever you like via the usual `res.*` terminals (`res.status(500).json({...})`, `res.html(...)`, etc.). `onError` may be `async`. If it finishes without finalizing a response, or itself throws/rejects, `sercon` falls back to the stock 500 — a buggy error handler can never wedge the request. Without `onError`, behaviour is unchanged (stock 500 + a log line); the per-route `try/catch` pattern still works and takes precedence (an error you catch never reaches `onError`).

Self-signed dev cert. `cert: "self-signed"` generates a fresh P-256 certificate in-process (valid ~1 year), with SANs covering `localhost`, `127.0.0.1`, `::1`, and the `listen host`. The keypair lives only for the life of the run and is never written to disk. Self-signed certs fail normal client verification by design (`net.probe.tls` skips verification, so it still works against them) — this is a local-dev convenience, not a production path. For real deployments own the cert in your supervisor (`caddy` / `nginx` / `traefik`) in front.

Route patterns are `stdlib` Go 1.22+ `http.ServeMux` syntax: `"METHOD /path/{param}/{rest...}"`. The leading method is optional (omit to match any). `{name}` captures one path segment; `{name...}` captures the tail (wildcard). No new routing library — `stdlib` does the matching, including longest-prefix wins. Captured values land in `req.params`.

RouteValue is either a bare handler — `(req, res) => res.text("...")` — or an object `{ use?: Middleware[], handler: HandlerFn }` for per-route middleware (see §6.4).

Request shape (`req`):

```
type Request = {
  method: string;           // "GET"
  url: string;              // full URL incl. scheme + host
  path: string;             // "/users/42"
  query: Record<string, string[]>; // ?a=1&a=2 → {a: ["1","2"]}
  headers: Record<string, string[]>; // lowercase keys
  params: Record<string, string>; // path-pattern captures
  body: string;             // UTF-8 decoded request body
  bodyBytes: Uint8Array;    // same data as raw bytes
  remote: string;           // "1.2.3.4:54321"
  cookies: Record<string, string>; // parsed Cookie header
};
```

Response builder (`res`) — every method returns `res` for chaining; the terminal `.json` / `.text` / `.html` / `.bytes` / `.empty` / `.redirect` sets the body and flips an internal `finalized` flag:

```

type CookieOpts = {
  domain?: string;
  path?: string;
  maxAge?: number; // seconds; 0 = session; <0 = delete
  expires?: number; // unix ms
  secure?: boolean;
  httpOnly?: boolean;
  sameSite?: "strict" | "lax" | "none";
};

type Response = {
  status(code: number): Response;
  header(name: string, value: string): Response;
  cookie(name: string, value: string, opts?: CookieOpts): Response;
  // Terminals (all return res so they remain chainable but no further
  // write makes sense after one fires; a second terminal throws):
  json(value: unknown): Response; // → application/json
  text(s: string): Response; // → text/plain
  html(s: string): Response; // → text/html
  bytes(b: Uint8Array, ct?: string): Response; // default application/octet-
stream
  empty(): Response; // no body; pair with .status() for
204/304
  redirect(loc: string, code?: number): Response; // default 302
  // Upgrade (WebSocket):
  upgradeWebSocket(opts?: { readBuffer?: number }): Promise<WebSocket>;
  // Stream (Server-Sent Events):
  sse(opts?: { keepAlive?: number; retry?: number }): SSEStream;
};

type SSEStream = {
  send(data: string | { data: unknown; event?: string; id?: string; retry?:
number }): Promise<void>;
  close(): Promise<void>;
  readonly closed: Promise<void>; // resolves on close OR client disconnect
  readonly remote: string;
};

```

After the handler's returned Promise resolves, the engine checks `res.finalized`:

- **true** — buffered status / headers / body get written to the underlying `http.ResponseWriter`.
- **false** — engine sends `204 No Content`.
- **handler threw** (sync or async) — engine sends `500 Internal Server Error` and logs the error (`stderr` under vanilla `sercon`; access log under `sercon serve`). The script keeps running; only that one request is affected.

Calling a terminal twice on the same `res` throws a `TypeError` (the second call sees `finalized === true` and refuses).

Multiple listeners compose naturally. A script can run

```
const [a, b] = await Promise.all([
  server.http.listen({ port: 80, routes }),
  server.https.listen({ port: 443, routes, cert, key }),
]);
```

Each `listen` call holds its own `HoldRun` sentinel; closing one listener does not affect the other.

6.3 Static-file serving

```
"GET /assets/{rest...}": server.http.static({
  dir:          "/var/www/public",    // root directory on disk
  stripPrefix:  "/assets/",           // URL prefix to strip before disk lookup
  index?:       "index.html",         // accepted but currently unused (stdlib
  default applies)
  etag?:        true,                 // accepted but currently unused (stdlib
  default applies)
}),
```

`server.http.static({...})` returns an opaque handler the route compiler unwraps and registers directly on the mux. **The route pattern must include a `{rest...}` wildcard** — without it the match only ever produces the literal prefix and no subpath ever resolves on disk. Internally a `stdlib http.FileServer(http.Dir(dir))` wrapped in `http.StripPrefix(stripPrefix, ...)`; ETag and range requests work out of the box. Symlinks outside `dir` are blocked (stdlib `http.Dir`'s default). No directory listing customisation yet — `index` and `etag` options are reserved for future tuning; v0.10.0 inherits the stdlib defaults.

6.4 Middleware

Onion model. A middleware is `(req, res, next) => Promise<void>`. `next` is an async function that runs the rest of the chain; awaiting it lets the middleware post-process.

```
const logger = async (req, res, next) => {
  const start = runtime.time.nowMs();
  await next(); // run downstream chain
  runtime.log(req.method, req.path, "→", runtime.time.nowMs() - start, "ms");
};

await server.http.listen({
  port: 8080,
  use: [logger], // global middleware
  routes: {
    "GET /api/secure": { // per-route middleware
      use: [authCheck],
      handler: (req, res) => res.json({ ok: true }),
    },
  },
});
```

Attachment points:

- **Global** — `use`: array in `listen({...})` options. Runs for every request, in declaration order, before any per-route middleware.
- **Per-route** — when a route value is `{ use: [...], handler: fn }`, those middleware run after globals but before the handler.

For a request to `POST /api/secure` with global middleware `[A, B]` and per-route middleware `[C]`, the composition is:

```
A(req, res, () => B(req, res, () => C(req, res, () => handler(req, res))))
```

Each layer can `await next()` then mutate response headers or record timings, exactly like Express / Koa.

Short-circuit: a middleware that does *not* call `await next()` stops the chain. It must terminate `res` itself (e.g. `res.status(401).text("denied")`) — otherwise the engine sends `204 No Content` per the unfinalized-response rule.

Throwing: any throw mid-chain (sync or via rejected Promise) bubbles to the engine's 500 path. Wrap with `try / catch` if you want custom error handling.

6.5 WebSocket upgrade

```
"GET /ws": async (req, res) => {
  const ws = await res.upgradeWebSocket({ readBuffer: 64 });
  // ws is both an async iterator AND has .send / .close.
  for await (const msg of ws) {
    if (msg.type === "text") {
      await ws.send("echo:" + msg.text);
    } else {
      await ws.send(msg.bytes);           // msg.type === "binary"
    }
  }
  // iterator ends on close; ws.closeCode / ws.closeReason populated
},
```

Type:

```
type WSMMessage =
  | { type: "text"; text: string }
  | { type: "binary"; bytes: Uint8Array };

type WebSocket = AsyncIterable<WSMMessage> & {
  send(data: string | Uint8Array): Promise<void>;
  close(code?: number, reason?: string): Promise<void>;
  closeCode?: number;
  closeReason?: string;
  remote: string;
};
```

Library: `github.com/coder/websocket` (the modern successor to `nhooyr.io/websocket`; `gorilla/websocket` is in maintenance mode). Pure Go, zero deps.

Marshalling. Per upgraded connection, a Go goroutine reads frames into a buffered channel (capacity = `readBuffer`, default 64). Each `next()` on the async iterator pulls one message off the channel and resolves on the loop via `LoopCallee.ws.send()` schedules a write back through the connection-owning goroutine.

Backpressure. If the read buffer fills (slow JS consumer), back-pressure propagates into the websocket library's read loop — the client eventually sees `1009 message too big` after the library's own limit (1MB default). Bump `readBuffer` if your workload bursts and you want more in-flight queueing.

`ws.close()` closes the send channel, drains the receive channel, sends a close frame, and ends the async iterator on the next `next()` call. The script's `for await` exits cleanly; the handler's Promise resolves; the connection's goroutine returns. The HTTP server's `HoldRun` is **unaffected** — it stays alive for new connections until `srv.close()`.

`async function*` and `for await (...)` are not natively parsed by goja. esbuild's `Supported` flag lowers both during transpile (see `CHANGELOG`), and every Run installs `Symbol.asyncIterator = Symbol.for("@@asyncIterator")` so the lowered helper and user code agree on the same iteration key.

Server-Sent Events (`res.sse`)

`res.sse(opts?)` starts a one-way `text/event-stream` response: the handler holds the connection open and pushes events until it (or the client) closes the stream.

```
"GET /events": (req, res) => {
  const stream = res.sse({ keepAlive: 15000 });
  let n = 0;
  const t = setInterval(() => stream.send({ event: "tick", data: { n: n++ } } ),
    1000);
  stream.closed.then(() => clearInterval(t)); // stop on client disconnect
  return stream.closed;                      // keep the handler alive
},
```

Options. `keepAlive` (ms) sends : `ping` comment frames at that interval to defeat idle-proxy timeouts (off by default). `retry` (ms) emits a one-time `retry:` line telling the client its reconnect delay.

send(data) . `data` is a string (→ a single `data:` line) or an object `{ event?, data, id?, retry? }`. Object `data` is JSON-encoded; string `data` is passed through, and multi-line string `data` becomes one `data:` line per source line (per the SSE grammar). The returned Promise resolves once the frame is flushed to the socket, and rejects if the stream is already closed.

close() ends the stream and resolves once it is torn down. `closed` is a Promise that resolves on `close()` or when the client disconnects — use it to stop timers/producers.

Mechanics. Unlike `upgradeWebSocket`, SSE does **not** hijack the connection — it keeps writing to the same `http.ResponseWriter`. So the request's dispatcher goroutine parks until the stream closes

(otherwise net/http would finish the response). A dedicated pump goroutine owns the writer and flushes each event; the event loop never writes to the socket directly. The listener's `HoldRun` keeps the loop alive while streams are open, exactly like `WebSocket`.

6.6 Lifecycle

Vanilla `sercon script.ts`. Each `server.*.listen` call parks a `HoldRun` sentinel and the event loop stays alive while the listener is bound. `SIGINT` terminates abruptly: the engine's interrupt watcher calls `loop.Terminate()`, the cleanup drain runs (closing each listener), and the process exits. No access log, no readiness signal, no shutdown timeout.

`sercon serve script.ts`. Adds the production niceties documented in §4: structured access log to `stderr`, a `READY` listening on `tcp/...` line on `stdout` per listener, and graceful shutdown with `--shutdown-timeout` (default `30s`). Clean `SIGTERM` exits `0`. Choose `serve` whenever the script is supervised (`systemd`, `docker compose`, `launchd`).

```
# Quick local test:
sercon examples/scripts/server-http.ts

# Production-style supervised run:
sercon serve examples/scripts/server-http.ts
```

6.7 SMTP

SMTP is a **stateful, multi-stage transaction**: a client connects, optionally authenticates, declares a sender (`MAIL FROM`), one or more recipients (`RCPT TO`), then streams the message body (`DATA`). `sercon` maps each stage to a JavaScript callback so a script can accept or reject the transaction incrementally — refuse an unknown sender at `MAIL`, refuse an unroutable recipient at `RCPT`, or inspect the parsed message at `DATA`. The inbound listener is `server.smtp.listen({...})`; the outbound side is `net.email.send({...})` (covered at the end of this section).

`server.smtp.listen({...})`

Synchronously binds the listening socket (so a port already in use throws immediately), then serves in the background. Returns a handle once the loop schedules it.

```
const srv = await server.smtp.listen({
  port: 2525,
  hostname: "mx.example.com",
  handlers: {
    onMail: (env) => env.from.endsWith("@example.com"),
    onRcpt: (env, rcpt) => isLocalMailbox(rcpt),
    onData: (env, msg) => { store(msg); return true; },
  },
});

// srv.address → "tcp/0.0.0.0:2525"
// srv.close() → Promise<void> (graceful, 30s drain)
// srv.stopped → Promise<void> (resolves when the listener exits)
```

Options:

Field	Type	Default	Notes
port	number	—	Required. TCP port to bind.
host	string	"0.0.0.0"	Bind address.
hostname	string	os.Hostname()	EHLO greeting + Received: identity.
handlers	{onMail, onRcpt, onData}	—	Required. All three functions; see below.
auth	(user, pass, env) => bool Promise<bool>	—	Optional SASL verifier (PLAIN + LOGIN).
starttls	{cert, key}	—	Enables STARTTLS. File paths or inline PEM.
allowInsecureAuth	boolean	false	Permit AUTH on a plaintext connection (dev only).
maxMessageBytes	number	10485760 (10 MB)	DATA size cap; a non-positive value keeps the default.
maxRecipients	number	100	Max RCPT TO per transaction.
sessionTimeout	number	30000	Per-session idle timeout, milliseconds.

Handlers. `onMail(envelope)`, `onRcpt(envelope, recipient)`, and `onData(envelope, message)` are invoked at the matching protocol stage. All three are required. Each may be `async` — a returned Promise is awaited before the SMTP response is written. The return value (or a thrown exception) controls the SMTP reply, uniformly across all three:

Return value	SMTP response	Use case
true / undefined	250 OK	Accept the stage
false	550 Command refused	Generic permanent reject
string	550 <string>	Permanent reject with a reason
throw	451 Temporary failure	Transient error (client should retry)

Envelope (passed to all three handlers; `recipients` grows as RCPT TO accumulates):

```

type Envelope = {
  from:          string; // "" for the null sender (DSN bounces)
  recipients:    string[]; // accepted RCPT TO so far
  remote:        string; // "1.2.3.4:54321"
  helo:          string; // EHLO/HELO hostname the client gave
  authenticatedUser?: string; // set if AUTH succeeded
  tls?:          { version: string; cipher: string }; // set under STARTTLS
};

```

Message (only `onData`; parsed via `jhillierd/enmime`, with the exact original bytes always available as `raw`):

```

type Message = {
  from:          string; // From: header (may differ from
  envelope.from)
  to:            string[]; // To: header
  cc:            string[]; // Cc: header
  subject: string;
  headers: Record<string, string[]>; // lowercase keys; multi-valued
  body: {
    text: string; // text/plain part; "" if absent
    html: string; // text/html part; "" if absent
  };
  attachments: Array<{
    filename: string;
    contentType: string;
    bytes: Uint8Array;
  }>;
  raw: Uint8Array; // exact DATA bytes (post dot-unstuffing)
};

```

Forwarders and DKIM-style signature verifiers should work from `raw`; everyone else uses the parsed fields. `enmime` is lenient with malformed messages, so a script needing strict RFC 5322 conformance can re-parse `raw` itself.

AUTH. Supply an `auth` callback to enable SASL. Both **PLAIN** and **LOGIN** mechanisms are advertised; the callback receives the decoded `(username, password, envelope)` and returns a boolean (or a Promise of one). By default AUTH is refused on a plaintext connection — the client must complete STARTTLS first. Set `allowInsecureAuth: true` to accept credentials over plaintext; this is a **dev-only** escape hatch for trusted local networks and should never be set in production.

STARTTLS. Pass `starttls: {cert, key}` (file paths or inline PEM, same convention as `server.https`) to advertise STARTTLS. The upgrade happens mid-session on the same connection; once active, `envelope.tls` is populated. The listener itself is plaintext at bind time — there is no implicit-TLS (SMTPS / port 465) inbound mode.

Concurrency. Each connection runs on its own Go goroutine, but every handler invocation **serializes on the goja loop** — exactly one handler runs at a time (the same single-threaded model as the HTTP listener in §6.2). A slow `onData` therefore blocks other connections' handlers. For slow work,

acknowledge the stage (`return true`) and process in the background (start a Promise without awaiting it).

A short, self-contained round-trip (bind, send to self, assert receipt):

```
const port = 38095;
let captured: { subject: string; from: string } | null = null;

const srv = await server.smtp.listen({
  port,
  hostname: "test.local",
  handlers: {
    onMail: () => true,
    onRcpt: () => true,
    onData: (env, msg) => {
      captured = { subject: msg.subject, from: env.from };
      return true;
    },
  },
});

await net.email.send({
  to:      "alice@test.local",
  from:    "bob@test.local",
  subject: "round-trip demo",
  body:    "hello from the SMTP demo",
  server:  { host: "127.0.0.1", port, tls: "none" },
});

runtime.assert.equal(captured!.subject, "round-trip demo", "subject");
await srv.close();
```

`net.email.send({...})` — outbound

The outbound counterpart lives under `net.email` (next to the `spf` / `dmARC` / ... probes). It opens **one TCP connection per call** — no pooling, no queueing, no automatic retries — composes the MIME body in-tree, and returns a per-recipient outcome.

```

const result = await net.email.send({
  to:      ["alice@example.com", "bob@example.com"], // string OR string[]
  from:    "noreply@my.app",
  subject: "your verification code",
  body:    "your code is 123456",                    // plain text
  html:    "<p>your code is <b>123456</b></p>",        // optional HTML
  alternative
    attachments: [
      { filename: "qr.png", contentType: "image/png", bytes: qrBytes },
    ],
    headers: { "X-App": "myapp" },                    // optional extra headers
    server: {
      host: "smtp.example.com",
      port: 587,                                       // default 587
      (submission)
        auth: { username: "u", password: "p" },      // optional PLAIN auth
        tls: "starttls",                             // "starttls" (default) |
        "tls" | "none"
      },
      timeout: 30000,                                // ms; default 30s
    });

// result = { accepted: string[], rejected: Array<{ address, reason }> }

```

MIME layout. `text only` → `text/plain`; `charset=utf-8` (no multipart). `text + html` → `multipart/alternative`. Any attachments → `multipart/mixed` (each part `base64`, `Content-Disposition: attachment`). `Date`, `Message-ID`, and `MIME-Version` are auto-generated; headers are merged on top.

TLS modes. `"starttls"` (default) connects plaintext then upgrades, refusing AUTH until TLS is active. `"tls"` is implicit TLS from connect (SMTPS, typically port 465). `"none"` is plaintext throughout, with AUTH disabled.

Outcomes. Transport-level failures (connection refused, TLS handshake, AUTH, `MAIL FROM` rejected) **throw** a JS exception. The `rejected` array captures only individual `RCPT TO` addresses the server permitted us to attempt but then refused — the transaction continues for the rest, so a partial send still delivers to the accepted recipients.

6.8 Raw TCP / UDP

`server.tcp.listen` and `server.udp.listen` are the inbound counterparts to the `net.tcp.connect` / `net.udp.open` clients (§5 `net`). They bind a listening socket **synchronously** and return the server handle directly — a port already in use throws immediately — then serve in the background, keeping the event loop alive via the same `HoldRun` model as the HTTP and SMTP listeners. `srv.close()` returns a `Promise<void>` (resolving once the listener is closed and accepted connections are drained), matching the rest of the `server.*` family. Passing `port: 0` binds an OS-chosen ephemeral port; read it back from the returned handle's `address`.

Both return a server handle:

```

srv.address;      // "tcp/127.0.0.1:54321" or "udp/127.0.0.1:54321"
srv.close();      // Promise<void> – stop accepting, drain, release HoldRun

```

server.tcp.listen({...}, conn => {...})

The second argument is a **connection handler** invoked once per accepted socket. The `conn` it receives is the **same handle shape** as `net.tcp.connect` — there is no separate server-connection type to learn:

```

const srv = await server.tcp.listen({ port: 0 }, (conn) => {
  conn.onData(ev => conn.write(ev.bytes)); // ev = { bytes: Uint8Array, text }
  conn.onClose(() => runtime.log("peer disconnected"));
  conn.onError(e => runtime.log("conn error", String(e)));
  // conn.write(data) – string (UTF-8) or Uint8Array
  // conn.remote / conn.local – peer / local "host:port"
  // conn.close() – half/full close this connection
});

const port = Number(srv.address.split(":").pop()); // "tcp/127.0.0.1:PORT"
// ... net.tcp.connect("127.0.0.1", port) to talk to it ...
await srv.close();

```

Options: `port` (required; 0 for ephemeral), `host` (default 127.0.0.1), `readBuffer` (per-connection read chunk size, same meaning as `net.tcp.connect`).

server.udp.listen({...}, (msg, reply) => {...})

UDP is connectionless, so the handler fires **once per inbound datagram** rather than once per connection. `msg` describes the datagram and its sender; `reply` sends a datagram back to that same sender:

```

const srv = await server.udp.listen({ port: 0 }, (msg, reply) => {
  // msg = { bytes: Uint8Array, text, address, port } – the sender
  runtime.log("got", msg.text, "from", msg.address + ":" + msg.port);
  reply("ack:" + msg.text); // string or Uint8Array; returns a Promise
});

const port = Number(srv.address.split(":").pop()); // "udp/127.0.0.1:PORT"
await srv.close();

```

Options: `port` (required; 0 for ephemeral), `host` (default 127.0.0.1).

Lifecycle. Identical to §6.6: vanilla `sercon script.ts` keeps the loop alive while bound and terminates abruptly on SIGINT; `sercon serve script.ts` adds the access log, a `READY` listening on `tcp/...` (or `udp/...`) line on stdout per listener, and graceful shutdown. As with all `server.*` bindings, the connection / datagram handlers marshal back onto the single goja loop, so **handlers serialize** (one at a time across all listeners).

- `server.icmp.listen(opts?, (msg, reply) => ...)` — bind a raw ICMP listener. **Requires root / CAP_NET_RAW** (synchronous bind throws otherwise); raw ICMP has no ports, so the

socket receives **all** host ICMP traffic. `opts { network?: "ip4" | "ip6" (default "ip4"), readBuffer? }`. The handler runs per received packet with `msg { bytes, text, address, type, code }` and a `reply(opts?)` that sends an ICMP message back to the sender (or `opts.to`) — Echo mode `{ type?, code?, id?, seq?, payload? }` or raw mode `{ type, code?, body }`, the same options as `net.icmp send`. The handle is `{ address: "icmp/<addr>", close() }`; it emits a `READY` line under `sercon serve` and joins graceful shutdown.

7. JavaScript runtime built-ins (goja)

`scriptengine` runs on goja, which implements **ES5.1 + a large subset of ES6+**. The following are present and behave per spec:

Globals

- `Object`, `Array`, `String`, `Number`, `Boolean`, `BigInt`
- `Math`, `Date`, `JSON`, `RegExp`
- `Error`, `TypeError`, `RangeError`, `SyntaxError`, `ReferenceError`, `EvalError`, `URIError`
- `Promise` (provided via goja's native implementation; resolution microtasks are pumped by the event loop)
- `Symbol` (partial: well-known symbols `Symbol.iterator`, `Symbol.asyncIterator`, etc., are supported)
- `Proxy`, `Reflect` (partial)

```
runtime.log(Math.PI.toFixed(4), Math.max(1, 9, 3));
runtime.log(new Date().toISOString());
runtime.log(JSON.stringify({ a: 1, b: [2, 3] }));
runtime.log("abc".repeat(3), "abc".padStart(6, "_"));
```

Collections

- `Map`, `Set`, `WeakMap`, `WeakSet`

```
const counts = new Map<string, number>();
for (const ch of "banana") counts.set(ch, (counts.get(ch) ?? 0) + 1);
runtime.log([...counts]); // [{"b",1}, {"a",3}, {"n",2}]
```

Typed arrays

- `ArrayBuffer`, `DataView`
- `Int8Array`, `Uint8Array`, `Uint8ClampedArray`, `Int16Array`, `Uint16Array`, `Int32Array`, `Uint32Array`, `Float32Array`, `Float64Array`

```
const buf = new ArrayBuffer(4);
new DataView(buf).setUint32(0, 0xCAFEBAFE);
runtime.log(new Uint8Array(buf)[0].toString(16)); // ca
```

Iteration and generators

- `for...of`, `for...in`, spread, destructuring
- Generator functions (`function*`)
- Async generators (`async function*`) — supported via goja's event loop

```
function* fib() {
  let [a, b] = [0, 1];
  while (true) {
    yield a;
    [a, b] = [b, a + b];
  }
}
const it = fib();
const first10 = Array.from({ length: 10 }, () => it.next().value);
runtime.log(first10);
```

Not (yet) provided

- `fetch` — use `net.http.*` from the CLI, or register your own binding.
- `Worker`, threading primitives.
- DOM globals (`window`, `document`).
- Native ES modules (`import` / `export` are *transpiled* to CommonJS at load time — see [section 9](#)).

8. Async runtime additions (goja_nodejs)

The event loop bundles a small Node-compatible runtime on top of goja:

Timers

```
setTimeout(() => runtime.log("after 50ms"), 50);
setInterval(() => runtime.log("tick"), 100);
setImmediate(() => runtime.log("next tick"));
clearTimeout(t); clearInterval(i); clearImmediate(im);
```

The event loop holds the script alive while *any* timer is pending. If your script ends with only `setTimeout` callbacks scheduled, the run will wait for them to fire before returning.

console

```
console.log("info");
console.error("oops");
console.warn("careful");
console.info("fyi");
console.debug("trace");
```

Disable with `Options.DisableConsole = true` if you want a clean global namespace.

require

CommonJS `require`, with TS-aware extension fallback added by `sercon` (see [section 11](#)):

```
const helpers = require("./helpers");
helpers.doThing();
```

Both `require("./foo")` and `import { x } from "./foo"` resolve to the same module instance per Run; esbuild rewrites `import` to `require` at transpile time.

9. TypeScript support

TS is transpiled by esbuild in-process. There is no `tsc`, no `tsconfig.json`, no type-checking — types are erased and the resulting JS is what executes. Practical implications:

- All TS syntax esbuild supports is allowed: type annotations, `interface`, `type`, generics, enums (compiled to objects), namespaces, decorators (legacy & TC39 stage 3).
- `import type { X } from "y"` is stripped entirely (no runtime require).
- TS errors are *not* surfaced as build errors — only esbuild parse errors are. Run a separate `tsc -noEmit` in CI if you want type enforcement.
- `tsconfig.json` is **not** consulted.
- **Runtime errors report TypeScript positions.** Stack traces in thrown errors point at the original `.ts` line/column, not the transpiled JS — for both the entry script and imported `.ts` / `.tsx` modules, and for both synchronous throws and async (top-level- `await`) rejections. This is done with inline source maps that goja consumes natively; the entry script's ESM→CJS rewrite is line-shift-aware so its frames map correctly too. If a map is ever unavailable the trace falls back to transpiled-JS positions.

`.tsx` and `.jsx` are supported via esbuild's `LoaderTSX` / `LoaderJSX`, including for required modules resolved by the engine's extension fallback. JSX has no React runtime in scope by default — either set the factory at the file level with an `@jsx` pragma:

```
/** @jsx h */
function h(tag: string, props: any, ...children: any[]) {
  return { tag, props: props ?? {}, children };
}
export const Box = (label: string) => <div className="box">{label}</div>;
```

...or provide a `React.createElement`-compatible binding from Go and rely on esbuild's default pragma. ESM `export default <value>` also works through the entry-script rewriter's `__esModule ? .default : m.unwrap`, so `import answer from './mod'` resolves to the default export.

An **entry script's own `export default <expr>`** is captured as the value `Engine.Run` resolves to, and the `sercon` CLI prints that value as JSON to stdout (scripts without a default export resolve to `undefined` and print nothing; a value that doesn't JSON-encode, such as a function, is skipped). This is the supported way for a script to emit a structured result.

10. Top-level `await`

Top-level `await` works in entry scripts:

```
const r = await net.http.get("https://example.com");
runtime.log(r.status);
```

How it works under the hood: esbuild rejects top-level `await` with the `cjs` output format, so the engine transpiles the *entry* script with `format=esm`, then rewrites `import` statements to `require()` calls and wraps the rest of the body in:

```
;(async () => { /* your transpiled body */ })().then(__resolve, __reject);
```

`__resolve` and `__reject` are bindings the engine sets on the runtime to capture the final settlement. *Required-module* code (anything loaded through `require`) is transpiled with `format=cjs` and does **not** support top-level `await` — wrap in `(async () => ... )()` yourself if you need it inside a module.

11. Module resolution

The engine plugs a custom source loader into `goja_nodejs/require` that adds:

1. Direct path: if the requested file exists on disk, use it. `.ts` / `.tsx` files are transpiled on the fly.
2. **Extension fallback** when the request has no extension: try `.ts`, `.tsx`, `.js`, `.cjs`, `.mjs`, `.json` in that order.
3. **`.js` → `.ts` swap**: if a path ending in `.js` / `.cjs` / `.mjs` is requested but the literal file doesn't exist, the same path ending in `.ts` (or `.tsx`) is tried. Handles `package.json` `main` fields that point at compiled output where only the TypeScript source is on disk.
4. **`package.json` source preference**: when reading a `package.json`, if a `source` field is present and points at an existing `.ts` / `.tsx` file, `main` is rewritten to that path before the registry consumes it. Matches the convention used by parcel, microbundle, and similar bundlers to mark the TS source of truth.
5. **Directory index**: if the resolved path is a directory, try `index.ts`, `index.tsx`, `index.js`.

The base directory follows Node's CommonJS rules — relative imports resolve against the requiring module's directory, courtesy of `goja_nodejs`.

```
// All resolve to ./helpers/assert.ts:
import { check } from "./helpers/assert";
import { check } from "./helpers/assert.ts";
const { check } = require("./helpers/assert");
const { check } = require("./helpers/assert.js");
```

JSON imports work out of the box via either the ESM-style default import (esbuild rewrites it to a `require`) or a direct `require`:

```
import data from "./data.json";
const r = require("./data.json");
```

`node_modules` lookup *works* via the upstream registry, but the source loader doesn't currently transpile through that path — keep TS helpers under `ScriptRoot`.

12. Timeouts and cancellation

Two mechanisms interrupt a running script:

- `Options.Timeout` — wall-clock limit per `Run`. Exceeding it returns `scriptengine.ErrScriptTimeout`.
- `ctx` passed to `Run / RunFile` — cancelling the context returns `ctx.Err()` (typically `context.Canceled` or `context.DeadlineExceeded`).

Both call `vm.Interrupt(...)` and `loop.Terminate()` from a watcher goroutine. This works for **sync** JS — including `while(true){}` — because the goja VM checks the interrupt flag between bytecode steps. A host Go binding that's blocked in cgo or a long syscall isn't interruptible mid-call; if you write such a binding, plumb the engine `ctx` through to it yourself.

```
eng := scriptengine.New(scriptengine.Options{Timeout: 200 * time.Millisecond})
_, err := eng.Run(ctx, "loop.ts", `while (true) {}`)
// err is scriptengine.ErrScriptTimeout, returned ~200-300ms after Run.
```

13. Error semantics

- A Go binding returning `(T, error)` automatically throws as a JS exception when `err != nil`. The exception's `.message` is `err.Error()`.

```
try { mayFail(); } catch (e) { runtime.log(String(e)); }
```

- A Go binding that `panic`s with a `*goja.Object` (e.g. `vm.NewGoError(...)`) does the same.
- A script throwing a JS error out of the top level surfaces from `Run` as a Go `error` whose `message` is the JS error's `.message`.

- Timeout errors satisfy `errors.Is(err, scriptengine.ErrScriptTimeout)`. Cancellation errors satisfy `errors.Is(err, context.Canceled)` or `context.DeadlineExceeded`.

14. Type generation (.d.ts)

`Engine.WriteTypes(w)` walks the registered bindings and emits a TypeScript declaration file. The mapping table (abbreviated):

Go type	TS type
<code>string</code>	<code>string</code>
any numeric / <code>float64</code>	<code>number</code>
<code>bool</code>	<code>boolean</code>
<code>[]T</code>	<code>T[]; []byte → Uint8Array</code>
<code>map[string]T</code>	<code>Record<string, T></code>
<code>*T</code>	<code>T</code> (transparent unwrap)
<code>error</code> (single return)	omitted (collapses to <code>void</code>)
<code>(T, error)</code> (tuple return)	<code>T</code>
<code>interface{}</code> / any	<code>unknown</code>
<code>goja.FunctionCall</code> parameter	folded into <code>(...args: unknown[])</code>
<code>goja.Value</code> return	<code>unknown</code>
<code>scriptengine.AsyncBinding</code> value (from <code>PromisifyAsync[T]</code>)	<code>Promise<T></code>
self-referential types	<code>unknown</code> (cycle detection bails after depth 4)

```
sercon --emit-dts sercon.d.ts
```

In your editor, place the resulting `.d.ts` next to your scripts (or reference it via `/// <reference path="..." />`) for autocomplete.

JSDoc comments on declarations

The emitter pulls JSDoc strings from the engine's doc map (set via `Engine.SetDocs` / `Engine.SetMemberDocs`) and renders them above each declaration. Single-line docs collapse to `/** ... */`; multi-line docs expand to a standard `/*`-prefixed block. Bindings without a doc entry emit no JSDoc — the emitter never inserts empty placeholders. Docs are looked up by the same dotted path the binding was registered under, so namespace members get hover text too:

```
// AUTO-GENERATED by scriptengine.Engine.WriteTypes – do not edit by hand.

/** Hashing primitives. */
declare const crypto: {
  hash: {
    /** SHA-256 hex digest of a UTF-8 input. */
    sha256(...args: unknown[]): string;
    /** BLAKE3 hex digest (32-byte output, lukechampine.com/blake3). */
    blake3(...args: unknown[]): string;
    // ...
  };
  // ...
};
```

The CLI's reserved-global surface ships with a curated doc map; see `cmd/sercon/docs.go` for the source of truth and add new entries there when adding new bindings (the lockstep rule keeps `examples / MANUAL / sercon.d.ts` in sync, and the doc map is now the seventh artifact in that chain).

15. Limitations and gotchas

- **No native ES modules at runtime.** All `import` is rewritten to `require` by esbuild. Side-effect-only imports run the module body exactly once per Run.
- **No top-level await inside required modules.** Wrap with `(async () => ... )()` inside the module if you need it.
- **No `tsconfig.json` honoured.** Module resolution and type checking are independent from any project-level TS config.
- **Per-Run runtime.** Side effects on `globalThis` do not persist between calls to `Run` on the same `Engine`. The require *compile* cache is shared; module *exports* are not.
- **Multi-line ESM imports may stress the entry rewriter.** The current rewriter is line-based and handles the common forms; especially gnarly inputs (comments inside the spec list, very unusual whitespace) fall back to executing as-is, which may throw.
- **RegisterConstructor gives real new semantics.** `new Foo(...)` runs the registered Go constructor; JS arguments are coerced to its parameter types (an argument that can't be coerced becomes the zero value rather than throwing); the result's exported methods/fields are reachable (subject to the `json`-tag field mapper); and a non-nil trailing `error` result throws. A panic inside the constructor propagates like any Go binding (it is not converted to a catchable JS error), and `instanceof Foo` is not guaranteed (the instance is the Go-wrapped object). This is a library-side API — the CLI registers namespaces, not constructors.
- **HTTP bindings are real network calls.** `net.http.*` uses `net/http` with a 5s timeout. They are not mockable from JS.

See [OUT-OF-SCOPE.md](#) for the active backlog of deferred ideas.

16. Binding reference (generated)

The per-function reference below is generated from the structured binding docs (MemberDoc) by `sercon --emit-reference`, spliced in by `make reference`. **Do not edit between the markers** — regenerate instead. Coverage grows per release as each namespace adopts the structured doc model (parameters, return shapes, errors, examples); namespaces not yet migrated show their one-line summary and a collapsed signature.

codec

Binary-format codecs: compression, barcodes, check digits.

codec.barcode.decodableFormats

```
decodableFormats(): string[]
```

Available decode formats (qr / datamatrix / aztec / code128 / code39 / code93 / codabar / ean13 / ean8 / upca / upce / itf). PDF417 is encode-only.

Returns: `string[]` — the twelve symbology names `barcode.decode` can recognise. PDF417 is absent (gozxing has no PDF417 decoder).

Throws: Never throws.

```
const fmts = codec.barcode.decodableFormats();
```

codec.barcode.decode

```
decode(data: string | Uint8Array | ArrayBuffer, format?: string): Promise<{  
  format: string, text: string }>
```

Decode a PNG/JPEG/WebP image to { format, text } via gozxing. Optional format hint skips the auto-detect walk. EAN/UPC need a quiet zone in the input. Async.

Parameters

- `data` (*string* | *Uint8Array* | *ArrayBuffer*) — Image bytes (PNG / JPEG / WebP). A string is treated as its raw UTF-8 bytes.
- `format` (*string*, *optional*) — Symbology hint (case-insensitive) from `decodableFormats`. When given, only that reader runs; otherwise every decoder is tried in priority order and the first hit wins.

Returns: `Promise<{ format: string, text: string }>` — the detected symbology name and the decoded payload.

Throws: Throws if data is missing/empty or not a string/ArrayBuffer/Uint8Array, the image cannot be decoded, the format hint is unsupported, or no barcode is recognised.

```
const { format, text } = await codec.barcode.decode(png);
```

codec.barcode.encode

```
encode(format: string, data: string, opts?: { width?: number, height?: number, quietZone?: boolean | number }): Promise<Uint8Array>
```

Render data into a PNG of the chosen format. `opts.width` / `opts.height` default to 256x256 (2D) or 400x120 (1D). `opts.quietZone` (true or px count) pads a white margin — required for EAN/UPC to decode. Async.

Parameters

- `format` (*string*) — Symbology (case-insensitive): qr / datamatrix / aztec / pdf417 / code128 / code39 / codabar / ean13 / ean8 / upca.
- `data` (*string*) — Payload to encode. EAN/UPC require the exact digit count for the variant; the encoder validates content per symbology.
- `opts` (*{ width?: number, height?: number, quietZone?: boolean | number }, optional*) — `width` / `height` set the output pixel dimensions (default 256x256 for 2D qr/datamatrix/aztec, 400x120 otherwise). `quietZone` pads a white margin: true uses 10% of width (min 10px), a number uses that many pixels per side, false/0/absent adds none. EAN/UPC need a quiet zone to be decodable.

Returns: Promise — PNG image bytes.

Throws: Throws if the format is unknown, the data is invalid for that symbology, or scaling / PNG encoding fails.

```
const png = await codec.barcode.encode("qr", "https://example.com", { width: 320, height: 320 });
```

codec.barcode.formats

```
formats(): string[]
```

Available encode formats (qr / datamatrix / aztec / pdf417 / code128 / code39 / codabar / ean13 / ean8 / upca).

Returns: string[] — the ten symbology names accepted by `barcode.encode`.

Throws: Never throws.

```
const fmts = codec.barcode.formats(); // ["qr", "datamatrix", ...]
```

codec.checkdigit.algos

```
algos(): string[]
```

Supported algorithms (luhn / isbn10 / isbn13 / ean13 / ean8 / upca).

Returns: string[] — the six supported algorithm names. isbn13 is an alias for ean13 (same check-digit math).

Throws: Never throws.

```
const algos = codec.checkdigit.algos();
```

codec.checkdigit.compute

```
compute(algo: string, partial: string): string
```

Compute the missing trailing check digit for a partial input.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `partial` (*string*) — The number WITHOUT its check digit (whitespace trimmed). Fixed-length algorithms expect exactly length-1 digits (e.g. 12 for ean13, 9 for isbn10).

Returns: string — the single check digit ('0'–'9', or 'X' for isbn10 when the value is 10).

Throws: Throws if the algorithm is unknown, the input is empty / the wrong length, or contains a non-digit.

```
const cd = codec.checkdigit.compute("ean13", "123456789012"); // "8"
```

codec.checkdigit.inspect

```
inspect(algo: string, input: string): { algo: string, input: string, valid: boolean, given: string, computed: string }
```

Diagnostic combining validate + compute: { valid, given, computed, ... }.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `input` (*string*) — The full number including its trailing check digit (whitespace trimmed).

Returns: { algo, input, valid, given, computed } — algo/input echo the normalised arguments; given is the input's last character; computed is the recalculated check digit (empty when the input is too short or malformed to split); valid is true when given equals computed (case-insensitive).

Throws: Never throws; malformed input yields valid:false with an empty computed.

```
const r = codec.checkdigit.inspect("ean13", "1234567890128");  
// { algo: "ean13", input: "...", valid: true, given: "8", computed: "8" }
```

codec.checkdigit.validate

```
validate(algo: string, input: string): boolean
```

Return whether the input passes the named algorithm's check digit.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `input` (*string*) — The full number including its trailing check digit (whitespace trimmed). ISBN-10 may end in 'X'.

Returns: boolean — true when the input's check digit is valid for the algorithm. Returns false (does not throw) for wrong length, non-digit characters, or an unknown algorithm.

Throws: Never throws; any failure (unknown algorithm included) returns false.

```
const ok = codec.checkdigit.validate("luhn", "4539578763621486"); // true
```

`codec.compression.algos`

```
algos(): string[]
```

Available compression algorithm names (gzip / deflate / zlib / bzip2 / zstd / brotli / lz4 / xz / snappy).

Returns: string[] — the nine supported algorithm names, lowercase, in registration order.

Throws: Never throws.

```
const algos = codec.compression.algos(); // ["gzip", "deflate", ...]
```

`codec.compression.compress`

```
compress(algo: string, data: string | Uint8Array | ArrayBuffer):  
Promise<Uint8Array>
```

Compress data with the named algorithm. Returns Uint8Array. Async.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive): gzip / deflate / zlib / bzip2 / zstd / brotli / lz4 / xz / snappy.
- `data` (*string* | *Uint8Array* | *ArrayBuffer*) — Input bytes. Strings are interpreted as their UTF-8 byte sequence.

Returns: Promise — the compressed bytes.

Throws: Throws if data is undefined/null or an unsupported type, the algorithm name is unknown, or the underlying compressor errors.

```
const packed = await codec.compression.compress("gzip", "hello world");
```

`codec.compression.decompress`

```
decompress(algo: string, data: string | Uint8Array | ArrayBuffer):  
Promise<Uint8Array>
```

Decompress data previously produced by compress (same algorithm name required). Returns Uint8Array. Async.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive), matching the one used to compress: gzip / deflate / zlib / bzip2 / zstd / brotli / lz4 / xz / snappy.
- `data` (*string* | *Uint8Array* | *ArrayBuffer*) — Compressed input bytes.

Returns: Promise — the original decompressed bytes.

Throws: Throws if data is undefined/null or an unsupported type, the algorithm name is unknown, or the input is not valid for that algorithm.

```
const raw = await codec.compression.decompress("gzip", packed);  
const text = new TextDecoder().decode(raw);
```

codec.perl.dumper

```
dumper(value: unknown, opts?: { classKey?: string, perlBoolClass?: string,  
indent?: string }): string
```

Perl Data::Dumper-style dump (\$VAR1 = ... ;), normalized indentation. JS booleans emit the JSON::XS::Boolean blessed-ref form (opts.perlBoolClass).

Parameters

- `value` (*unknown*) — Any JSON-like value (see php.serialize). Arrays/objects emit as Perl array/hash refs; class-key objects emit as blessed refs.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `perlBoolClass` names the blessed class emitted for JS booleans (default "JSON::XS::Boolean"). `indent` overrides the indentation step. `classKey` overrides the class-name sentinel (default "__class").

Returns: string — a Data::Dumper-style dump (\$VAR1 = ... ;) with normalized indentation.

Throws: Throws on a circular reference or an unsupported value type.

```
const d = codec.perl.dumper({ ok: true });
```

codec.perl.parseDumper

```
parseDumper(input: string, opts?: { classKey?: string, perlBoolClass?: string,  
indent?: string }): unknown
```

Read Data::Dumper output back. Blessed scalar refs in the JSON bool family decode to booleans; bare 1/0 stay numbers; cycles throw.

Parameters

- `input` (*string*) — `Data::Dumper` output to parse back (a `$VARn = ... ; assignment` or a bare value).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded blessed refs (default `"__class"`). The JSON bool family (`JSON::XS::Boolean`, `JSON::PP::Boolean`, `Types::Serialiser::Boolean`) decodes to JS booleans regardless.

Returns: unknown — the decoded value. Blessed scalar refs in the JSON-bool family become JS booleans; bare 1/0 stay numbers; other blessed refs carry the `classKey` sentinel.

Throws: Throws on malformed input or a reference that would close a cycle.

```
const v = codec.perl.parseDumper("$VAR1 = [1, 2];"); // [1, 2]
```

`codec.php.parseVarDump`

```
parseVarDump(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

Best-effort read of `var_dump()` output. Throws on lossy markers (*RECURSION*, truncation, visibility-annotated props).

Parameters

- `input` (*string*) — PHP `var_dump()` output to parse back.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded objects (default `"__class"`).

Returns: unknown — the reconstructed value (best-effort; `var_dump` is a lossy format).

Throws: Throws on lossy markers it cannot faithfully reverse: *RECURSION*, truncation, or visibility-annotated (private/protected) properties.

```
const v = codec.php.parseVarDump('int(42)'); // 42
```

`codec.php.parseVarExport`

```
parseVarExport(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

Read a `var_export()` literal (arrays, scalars, NULL, `\Cls::__set_state`) back to a value.

Parameters

- `input` (*string*) — A PHP `var_export()` literal: `array(...)`, scalars, NULL, or `\Cls::__set_state(array(...))`.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded `__set_state` objects (default `"__class"`).

Returns: unknown — the decoded value; \Cls::__set_state objects become plain objects carrying the classKey sentinel.

Throws: Throws on input that is not a parseable var_export() literal.

```
const v = codec.php.parseVarExport("array (\n  0 => 1,\n)");
```

codec.php.serialize

```
serialize(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP serialize(): encode a value to PHP's canonical serialization string. Objects use the __class sentinel; cycles throw.

Parameters

- `value` (*unknown*) — Any JSON-like value: null, boolean, number, string, array, or plain object. An object carrying the class-key sentinel (opts.classKey, default "__class") encodes as a PHP object (O:).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — classKey overrides the class-name sentinel property (default "__class"). indent / perlBoolClass are unused by serialize.

Returns: string — the PHP serialize() string (e.g. a:1:{...}).

Throws: Throws on a circular reference or an unsupported value type (function, Symbol, BigInt, Date, Map, Set, RegExp).

```
const s = codec.php.serialize({ a: 1, b: [2, 3] });
```

codec.php.unserialize

```
unserialize(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

PHP unserialize(): decode a serialize() string back to a value. r:/R: references resolve to shared objects (DAGs); cycles throw.

Parameters

- `input` (*string*) — A PHP serialize() string.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — classKey sets the sentinel property used to tag decoded PHP objects (default "__class").

Returns: unknown — the decoded value. PHP objects become plain objects carrying the classKey sentinel; r:/R: references rebuild as shared object identities (DAGs).

Throws: Throws on malformed input or a reference that would close a cycle.

```
const v = codec.php.unserialize('a:2:{i:0;i:1;i:1;i:2;}'); // [1, 2]
```

codec.php.varDump

```
varDump(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP `var_dump()`: human-readable debug output. String lengths are byte counts.

Parameters

- `value` (*unknown*) — Any JSON-like value (see `php.serialize`).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `indent` overrides the default indentation step. `classKey` overrides the class-name sentinel (default `"__class"`).

Returns: string — `var_dump()`-style output. String lengths in the output are byte counts, matching PHP.

Throws: Throws on a circular reference or an unsupported value type.

```
const dump = codec.php.varDump({ name: "ok" });
```

codec.php.varExport

```
varExport(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP `var_export()`: emit valid PHP code for a value. `opts.indent` overrides the 2-space step.

Parameters

- `value` (*unknown*) — Any JSON-like value (see `php.serialize`). Objects with the class-key sentinel emit as `\Cls::__set_state(...)`.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `indent` overrides the default 2-space indentation step. `classKey` overrides the class-name sentinel (default `"__class"`).

Returns: string — valid PHP source, the kind `var_export()` prints.

Throws: Throws on a circular reference or an unsupported value type.

```
const code = codec.php.varExport({ x: 1 }, { indent: "  " });
```

codec.xml.decode

```
decode(xml: string): unknown
```

Parse an XML string to a value using the same `@-prefix + #text` convention as `xml.encode`. Attributes become `@-keys`, text becomes `#text` (or a bare string for a text-only element), child elements become

keys, and repeated same-name siblings become an array. Empty/self-closing elements decode to null. Namespace prefixes are kept literally; all values are strings (no type coercion). Mismatched tags, multiple roots, and malformed XML throw.

Parameters

- `xml` (*string*) — The XML document to parse.

Returns: A single-key object whose key is the root element name and whose value is the parsed content (key order follows document order; all leaf values are strings).

Throws: Throws on malformed XML, mismatched/mis-nested end tags, multiple root elements, or no root element.

```
const v = codec.xml.decode("<note id=\"5\">hi</note>");
// { note: { "@id": "5", "#text": "hi" } }
```

`codec.xml.encode`

```
encode(value: unknown, opts?: { rootName?: string, indent?: string,
  declaration?: boolean }): string
```

Serialize a value to an XML string. Convention: @-prefixed keys are attributes, #text is element text, other keys are child elements, and an array value becomes repeated sibling elements (a scalar key becomes a text-only element, null a self-closing tag). The value must be a single-key object naming the root element, or pass `opts.rootName` to wrap it. Scalars are stringified; object key order is preserved.

Parameters

- `value` (*unknown*) — A single-key object whose one key names the root element — e.g. `{ note: { "@id": "5", "#text": "hi", to: "alice" } }` → `hialice`. Or any value plus `opts.rootName` to wrap it. Cycles throw.
- `opts` (`{ rootName?: string, indent?: string, declaration?: boolean }, optional`) — `rootName` wraps the value under that root element. `indent` pretty-prints with the given unit per level (default compact). `declaration` prepends (default off).

Returns: The XML string.

Throws: Throws if the value has no single root element and no `opts.rootName`, if the root content is an array, if a non-scalar is used as an attribute or #text value, or if the value contains a cycle.

```
const xml = codec.xml.encode({ note: { "@id": "5", "#text": "hi" } });
// <note id="5">hi</note>
```

console

Browser/Node-style console shim: `log/info/debug` to `stdout`, `warn/error` to `stderr`. For porting scripts; `runtime.log` is the native equivalent.

console.debug

```
debug(args: ...unknown[]): void
```

Alias of console.log — stringified arguments, space-joined, to stdout.

Parameters

- `args (...unknown[])` — Values to print; identical formatting and stdout destination as console.log.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.debug("cache hit", key);
```

console.error

```
error(args: ...unknown[]): void
```

Like console.log but writes to stderr.

Parameters

- `args (...unknown[])` — Values to print; same space-joined / JSON formatting as console.log but routed to stderr.

Returns: void — output is written to stderr as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.error("request failed", { status: 500 });
```

console.info

```
info(args: ...unknown[]): void
```

Alias of console.log — stringified arguments, space-joined, to stdout.

Parameters

- `args (...unknown[])` — Values to print; identical formatting and stdout destination as console.log.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.info("listening on", 8080);
```

console.log

```
log(args: ...unknown[]): void
```

Print a space-joined line of the arguments to stdout. Primitives print raw; objects/arrays render as JSON. Browser/Node-compatible; same output as runtime.log.

Parameters

- `args (...unknown[])` — Values to print, joined by single spaces and terminated with a newline. Primitives (string/number/boolean/null/undefined) print raw; objects and arrays render as JSON via `JSON.stringify`, falling back to `String()` for functions and circular references.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their `String()` form.

```
console.log("user", { id: 1, name: "ada" }); // user {"id":1,"name":"ada"}
```

console.warn

```
warn(args: ...unknown[]): void
```

Like `console.log` but writes to stderr.

Parameters

- `args (...unknown[])` — Values to print; same space-joined / JSON formatting as `console.log` but routed to stderr.

Returns: void — output is written to stderr as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their `String()` form.

```
console.warn("retrying in", 5, "seconds");
```

crypto

Hashing, JWT, age encryption — anything that produces a digest, signature, or ciphertext.

crypto.encrypt.decrypt

```
decrypt(ciphertext: string | Uint8Array | ArrayBuffer, identities: string | string[]): Uint8Array
```

Open a payload with one of the supplied identities. Routes to age or PGP based on the identity / ciphertext format. age: binary or armored auto-detected. Wrong identity throws.

Parameters

- `ciphertext` (*string* | *Uint8Array* | *ArrayBuffer*) — The encrypted payload; age armor and PGP armor are auto-detected.
- `identities` (*string* | *string[]*) — One or more private keys. AGE-SECRET-KEY-1... identities use the age backend; -----BEGIN PGP PRIVATE KEY BLOCK----- entries (or an armored PGP message) use the PGP backend.

Returns: *Uint8Array* — the decrypted plaintext bytes (decode with `new TextDecoder().decode(...)` for a string).

Throws: Throws if ciphertext is empty or an unsupported type, no identities are given, an identity is actually a public key, or no supplied identity matches the ciphertext's recipients.

```
const pt = crypto.encrypt.decrypt(ct, privateKey);
const text = new TextDecoder().decode(pt);
```

crypto.encrypt.detectBackend

```
detectBackend(input: string): { backend: string; kind?: string }
```

Classify a recipient / identity string. Returns { backend: 'age'|'pgp'|'unknown', kind?: 'public'|'private' }. Pure prefix matching; no parsing or I/O.

Parameters

- `input` (*string*) — A recipient or identity string to classify (age bech32, age SSH public key, or PGP armored block).

Returns: { backend: "age" | "pgp" | "unknown", kind?: "public" | "private" } — kind is present only when the backend was identified.

Throws: Never throws; unrecognised input returns { backend: "unknown" }.

```
const info = crypto.encrypt.detectBackend("age1abc..."); // { backend: "age",
kind: "public" }
```

crypto.encrypt.encrypt

```
encrypt(data: string | Uint8Array | ArrayBuffer, recipients: string | string[],
opts?: { armored?: boolean }): Uint8Array
```

Seal data to recipients. age public keys (age1...) → age backend (opts.armored for ASCII); PGP public-key blocks → PGP backend (always armored). Auto-dispatched on key format. Multi-recipient: any listed identity decrypts.

Parameters

- `data` (*string* | *Uint8Array* | *ArrayBuffer*) — Plaintext to encrypt.
- `recipients` (*string* | *string[]*) — One or more recipient public keys. age1... bech32 keys use the age backend; -----BEGIN PGP PUBLIC KEY BLOCK----- entries use the PGP backend. Backends cannot be mixed in one call.

- `opts` (*{ armored?: boolean }, optional*) — armored (age backend only) wraps the output in age ASCII armor; defaults to false (binary). Ignored on the PGP path, which is always armored.

Returns: Uint8Array — the ciphertext (binary age, ASCII-armored age when `opts.armored`, or armored PGP message bytes).

Throws: Throws if data is an unsupported type, no recipients are given, a recipient is actually a private key, or a recipient string fails to parse.

```
const ct = crypto.encrypt.encrypt("secret", publicKey, { armored: true });
```

crypto.encrypt.keygen

```
keygen(): { publicKey: string; privateKey: string }
```

Generate a fresh age X25519 keypair. Returns `{ publicKey: 'age1...', privateKey: 'AGE-SECRET-KEY-1...' }`.

Returns: `{ publicKey: string, privateKey: string }` — `publicKey` is the shareable age1... recipient; `privateKey` is the secret AGE-SECRET-KEY-1... identity.

Throws: Throws if the system RNG fails to generate an identity.

```
const { publicKey, privateKey } = crypto.encrypt.keygen();
```

crypto.encrypt.keygenPgp

```
keygenPgp(opts?: { name?: string, email?: string }): { publicKey: string;
privateKey: string }
```

Generate a PGP keypair (RSA 2048). `opts.name` / `opts.email` populate the user ID. Returns armored `{ publicKey, privateKey }` blocks. `encrypt/decrypt` auto-route to PGP when they see these.

Parameters

- `opts` (*{ name?: string, email?: string }, optional*) — name and email populate the primary user ID; both default to empty (fine for throwaway keys).

Returns: `{ publicKey: string, privateKey: string }` — both ASCII-armored PGP key blocks (-----BEGIN PGP PUBLIC/PRIVATE KEY BLOCK-----).

Throws: Throws if entity generation or armor serialisation fails.

```
const { publicKey, privateKey } = crypto.encrypt.keygenPgp({ name: "Test",
email: "t@example.com" });
```

crypto.encrypt.rekey

```
rekey(ciphertext: string | Uint8Array | ArrayBuffer, oldIdentities: string | string[], newRecipients: string | string[], opts?: { armored?: boolean }): Uint8Array
```

Re-encrypt for a new recipient set without exposing plaintext to JS. Output format defaults to match the input; `opts.armored` forces. Internal decrypt+encrypt loop.

Parameters

- `ciphertext` (*string* | *Uint8Array* | *ArrayBuffer*) — The existing age ciphertext to re-key (armor auto-detected).
- `oldIdentities` (*string* | *string*[]) — Current AGE-SECRET-KEY-1... private keys used to decrypt the existing ciphertext.
- `newRecipients` (*string* | *string*[]) — New age1... public keys to encrypt the result for.
- `opts` (*{ armored?: boolean }, optional*) — `armored` forces the output armor state; default preserves the input's armor state.

Returns: *Uint8Array* — the re-encrypted ciphertext for the new recipient set.

Throws: Throws if ciphertext is empty, either key list is empty, a key is the wrong kind (public vs private), or the decrypt step fails (no matching identity / malformed input). age backend only.

```
const rotated = crypto.encrypt.rekey(ct, oldKey, newPublicKey);
```

`crypto.hash.blake3`

```
blake3(input: string): string
```

BLAKE3 hex digest (32-byte output, lukechampine.com/blake3).

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: *string* — 64-char lowercase hex digest (256-bit output).

Throws: Throws a *TypeError* if input is missing, null, or undefined.

```
const d = crypto.hash.blake3("hello");
```

`crypto.hash.crc32`

```
crc32(input: string): string
```

CRC-32 (IEEE polynomial), zero-padded to 8 hex chars.

Parameters

- `input` (*string*) — Data to checksum, interpreted as a UTF-8 byte sequence.

Returns: string — 8-char zero-padded lowercase hex checksum.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const c = crypto.hash.crc32("hello"); // "3610a686"
```

crypto.hash.md5

```
md5(input: string): string
```

MD5 hex digest of a UTF-8 input. Avoid for security purposes — exposed for compatibility with legacy fingerprints.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 32-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.md5("hello"); // "5d41402abc4b2a76b9719d911017c592"
```

crypto.hash.sha1

```
sha1(input: string): string
```

SHA-1 hex digest of a UTF-8 input. Avoid for security purposes.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 40-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha1("hello");
```

crypto.hash.sha256

```
sha256(input: string): string
```

SHA-256 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 64-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha256("hello");
```

crypto.hash.sha384

```
sha384(input: string): string
```

SHA-384 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 96-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha384("hello");
```

crypto.hash.sha3_256

```
sha3_256(input: string): string
```

SHA-3 256-bit hex digest. The underscore in the name matches recon's binding.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 64-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha3_256("hello");
```

crypto.hash.sha3_512

```
sha3_512(input: string): string
```

SHA-3 512-bit hex digest.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 128-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha3_512("hello");
```

crypto.hash.sha512

```
sha512(input: string): string
```

SHA-512 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 128-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha512("hello");
```

crypto.jwt.sign

```
sign(claims: Record<string, unknown>, secret: string, opts?: { algorithm?: string }): string
```

Sign a claims object. secret is raw bytes for HS*; PEM-encoded private key for RS*/PS*/ES*/EdDSA; or a JWK JSON object (kty picks the key type) for any algorithm. opts.algorithm defaults to HS256.

Parameters

- `claims` (*Record<string, unknown>*) — Claims payload. Passed through to MapClaims; RFC 7519 reserved claims (exp/nbf/iat/iss/aud/sub) are honoured. Reserved claims are NOT synthesised — set iat/exp explicitly if you want them.
- `secret` (*string*) — Key material: raw HMAC bytes for HS256/384/512, a PEM-encoded private key (-----BEGIN ...) for RS*/PS*/ES*/EdDSA, or a JWK JSON object ({ "kty": ... }).
- `opts` (*{ algorithm?: string }, optional*) — algorithm names the RFC 7518 alg (HS256/HS384/HS512, RS256/384/512, PS256/384/512, ES256/384/512, EdDSA); case-insensitive. Defaults to HS256.

Returns: string — the signed compact-serialisation JWT (header.payload.signature).

Throws: Throws if claims is null, secret is empty, the algorithm is unsupported, or the secret shape mismatches the algorithm (e.g. HMAC algo with a PEM key, or asymmetric algo with plain bytes).

```
const tok = crypto.jwt.sign({ sub: "u1" }, "topsecret", { algorithm: "HS256" });
```

crypto.jwt.validate

```
validate(token: string, secret: string, opts?: { algorithm?: string, audience?: string, issuer?: string }): { valid: boolean; claims?: object; reason?: string }
```

Verify signature + standard claims (exp/nbf/iat) + optional aud/iss. secret accepts raw bytes / PEM public key / JWK. Set `opts.algorithm` for the algo-confusion guard. Resolves { valid:true, claims } or { valid:false, reason }.

Parameters

- `token` (*string*) — The compact-serialisation JWT to verify.
- `secret` (*string*) — Verification key: raw HMAC bytes, a PEM-encoded public key (or certificate), or a JWK JSON object.
- `opts` (*{ algorithm?: string, audience?: string, issuer?: string }, optional*) — algorithm pins the accepted alg (algorithm-confusion guard); when unset, any supported alg is accepted. audience / issuer enforce the aud / iss claims when set.

Returns: { valid: true, claims: object } on success, or { valid: false, reason: string } on any verification / claim failure.

Throws: Throws only on structural / wiring errors (empty token or secret, malformed token, or a secret shape that mismatches the token's algorithm). Cryptographic and claim failures resolve as { valid: false, reason } instead of throwing.

```
const r = crypto.jwt.validate(tok, "topsecret", { algorithm: "HS256" });
if (r.valid) runtime.log(r.claims.sub);
```

crypto.jwt.view

```
view(token: string): { header: object; payload: object; signature: string }
```

Decode header + payload WITHOUT verifying the signature. Useful for inspection / debugging auth flows. Malformed input throws.

Parameters

- `token` (*string*) — A compact-serialisation JWT (three dot-separated base64url segments).

Returns: { header: object, payload: object, signature: string } — decoded header and payload (object key order preserved) plus the raw signature segment.

Throws: Throws if the token is not three dot-separated segments or a segment fails base64url / JSON decoding.

```
const { header, payload } = crypto.jwt.view(tok);
```

db

Database / KV / directory clients: SQLite, PostgreSQL, MySQL/MariaDB, SQL Server, Redis, Valkey, memcached, LDAP, dict.

db.clickhouse.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?: string, database?: string, secure?: boolean }): Promise<Record<string, unknown>>
```

Connect to ClickHouse via the pure-Go clickhouse-go v2 driver. Arg is a clickhouse:// URL DSN string or an options object { host, port, user, password, database, secure }. Returns the shared SQL handle { exec, query, queryValue, begin, prepare, close }. Uses ? placeholders; default native port 9000 (set secure:true for TLS). Pings on open.

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string*, *secure?: boolean* }) — A clickhouse:// URL DSN string used verbatim, OR an options object assembled into one (defaults: host localhost, port 9000 native protocol, empty database). secure:true appends secure=true to the URL for TLS (typically port 9440).

Returns: Promise resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise` (first column of the first row, or null); `begin() → Promise`; `prepare(sql) → Promise`; `close() → Promise`. ClickHouse uses ? positional placeholders (or @name).

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.clickhouse.open({ host: "localhost", database: "metrics",
secure: false });
const rows = await db.query("SELECT name, value FROM stats WHERE host = ?",
"web1");
await db.close();
```

db.dict.define

```
define(host: string, word: string, opts?: { database?: string, port?: string,
timeout?: number }): Promise<{ word: string; found: boolean; definitions: { db:
string; dbName: string; text: string }[] >
```

RFC 2229 DICT word lookup. `define(host, word, opts?) -> { word, found, definitions: [{ db, dbName, text }] }`. found:false on no match (not an error). One-shot: connect, query, QUIT.

Parameters

- `host` (*string*) — The DICT server hostname.
- `word` (*string*) — The word to look up.
- `opts` ({ *database?: string*, *port?: string*, *timeout?: number* }, *optional*) — database selects a specific dictionary (default "*" = all); port is the DICT port (default "2628"); timeout is the dial/read deadline in milliseconds (default 10000).

Returns: `Promise<{ word: string, found: boolean, definitions: { db: string, dbName: string, text: string } [] >` — definitions carries one entry per matching dictionary (db is the dictionary code, dbName its human name, text the definition body). A word with no definitions resolves with found:false and an empty list.

Throws: Throws if host or word is empty, on dial/banner failure, or on an unexpected DICT status code (e.g. 550 invalid database). A 552 "no match" is NOT an error — it resolves with found:false.

```
const r = await db.dict.define("dict.org", "serendipity");
if (r.found) runtime.log(r.definitions[0].text);
```

db.dict.match

```
match(host: string, word: string, opts?: { strategy?: string, database?: string,
port?: string, timeout?: number }): Promise<{ word: string; matches: { db:
string; word: string }[] }>
```

RFC 2229 word match. match(host, word, opts?) -> { word, matches: [{ db, word }] }. opts.strategy (default prefix), opts.database (default *), opts.port (default 2628). One-shot: connect, query, QUIT.

Parameters

- `host` (*string*) — The DICT server hostname.
- `word` (*string*) — The word (or pattern) to match.
- `opts` (*{ strategy?: string, database?: string, port?: string, timeout?: number }, optional*) — strategy is the match strategy (default "prefix"); database selects a specific dictionary (default "*" = all); port is the DICT port (default "2628"); timeout is the dial/read deadline in milliseconds (default 10000).

Returns: Promise<{ word: string, matches: { db: string, word: string }[] }> — matches carries one entry per matched word (db is the dictionary it was found in, word the matched headword). No matches resolves with an empty matches list.

Throws: Throws if host or word is empty, on dial/banner failure, or on an unexpected DICT status code. A 552 "no match" is NOT an error — it resolves with an empty matches list.

```
const r = await db.dict.match("dict.org", "seren", { strategy: "prefix" });
runtime.log(r.matches.map(m => m.word));
```

db.ldap.open

```
open(url: string, opts?: { bindDN?: string, password?: string }):
Promise<Record<string, unknown>>
```

Dial LDAP (ldap://host:port or ldaps://...), anonymous bind by default (or opts.bindDN/password). Returns { rootDSE, search, close }. search(baseDN, filter, attrs?) -> entries; rootDSE -> server metadata.

Parameters

- `url` (*string*) — An LDAP URL: ldap://host:port (or ldaps://... for TLS, e.g. ldap://localhost:389).
- `opts` (*{ bindDN?: string, password?: string }, optional*) — When bindDN is set, the connection binds with bindDN/password instead of doing an anonymous bind.

Returns: Promise resolving to { rootDSE, search, close }: rootDSE() → Promise reads the server's Root DSE (an ordered { dn, : string[] } object advertising naming contexts, supported controls, vendor, etc.; an empty object when the server returns no entry); search(baseDN, filter, attrs?) → Promise<object[]> runs a whole-subtree search and returns one ordered { dn, : string[] } object per entry (multi-valued attributes stay arrays; filter defaults to (objectClass=*); attrs is an optional array of attribute names); close() → Promise. A directory-inspection (read) binding, not a write/modify surface.

Throws: open throws if url is empty, the dial fails, or (when bindDN is set) the bind fails (the connection is closed on bind failure). rootDSE / search throw on the underlying LDAP search error.

```
const dir = await db.ldap.open("ldap://localhost:389");
const meta = await dir.rootDSE();
const people = await dir.search("ou=people,dc=example,dc=com", "(uid=alice)",
["cn", "mail"]);
await dir.close();
```

db.memcached.open

```
open(addr: string): Promise<Record<string, unknown>>
```

Connect to memcached (host:port). Returns { get, set, delete }. get -> string or null (miss); delete -> bool (existed). set(key, value, expirySeconds?). No ping on open; the pool is lazy.

Parameters

- `addr` (*string*) — A memcached server address, host:port (e.g. localhost:11211).

Returns: Promise resolving to { get, set, delete }: get(key) → Promise<string | null> (null on a cache miss); set(key, value, expirySeconds?) → Promise (value stored as bytes; expirySeconds 0 or omitted means never expire); delete(key) → Promise (true if the key existed, false on a miss). gomemcache pools connections lazily, so there is no ping-on-open and no close method (the pool is GC'd with the handle).

Throws: open throws if addr is empty. set throws if key is empty or the value cannot be coerced to bytes. get / delete throw on transport errors; a cache miss is data (get → null, delete → false), not an error.

```
const mc = await db.memcached.open("localhost:11211");
await mc.set("session:42", "active", 300);
const v = await mc.get("session:42"); // "active" or null
const existed = await mc.delete("session:42"); // true
```

db.mssql.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string }): Promise<Record<string, unknown>>
```

Connect to Microsoft SQL Server via the pure-Go go-mssqldb driver. Arg is a sqlserver:// URL DSN string or an options object { host, port, user, password, database }. Returns the shared SQL handle.

Uses @p1,@p2,... placeholders. Pings on open.

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string* })
— A sqlserver:// URL DSN string used verbatim, OR an options object assembled into one (defaults: host localhost, port 1433; database goes in the URL query string per the go-mssqldb form).

Returns: Promise resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise` (first column of the first row, or null); `begin() → Promise`; `prepare(sql) → Promise`; `close() → Promise`. SQL Server uses @p1, @p2, ... placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.mssql.open({ host: "localhost", user: "sa", password:
"P@ss", database: "shop" });
const rows = await db.query("SELECT TOP 5 id, name FROM users WHERE region =
@p1", "EU");
await db.close();
```

db.mysql.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string }): Promise<Record<string, unknown>>
```

Connect to MySQL/MariaDB via the pure-Go go-sql-driver. Arg is a go-sql-driver DSN string (user:pass@tcp(host:port)/db) or an options object { host, port, user, password, database }. Returns the shared SQL handle. Uses ? placeholders. Pings on open.

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string* })
— A go-sql-driver DSN string (user:pass@tcp(host:port)/db?params) used verbatim, OR an options object assembled into one (defaults: host localhost, port 3306, empty database). One driver serves both MySQL and MariaDB.

Returns: Promise resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise` (first column of the first row, or null); `begin() → Promise`; `prepare(sql) → Promise`; `close() → Promise`. MySQL uses ? positional placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.mysql.open("app:s3cr3t@tcp(localhost:3306)/shop");
const n = await db.queryValue("SELECT COUNT(*) FROM orders WHERE status = ?",
"paid");
await db.close();
```

db.oracle.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string }): Promise<Record<string, unknown>>
```

Connect to Oracle via the pure-Go go-ora driver (no cgo). Arg is an oracle:// URL DSN string or an options object { host, port, user, password, database } where database is the service name. Returns the shared SQL handle. Uses :1,:2,... bind placeholders; default port 1521. Pings on open.

Parameters

- `dsn` (*string* | { *host?: string, port?: number, user?: string, password?: string, database?: string* }) — An oracle:// URL DSN string used verbatim, OR an options object assembled into one (defaults: host localhost, port 1521). database is the Oracle service name and goes in the URL path. The go-ora driver is pure Go, unlike the OCI-bound godror.

Returns: Promise resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise` (first column of the first row, or null); `begin() → Promise`; `prepare(sql) → Promise`; `close() → Promise`. Oracle uses :1, :2, ... (or :name) bind placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.oracle.open({ host: "localhost", user: "app", password:
"s3cr3t", database: "ORCLPDB1" });
const rows = await db.query("SELECT id, name FROM users WHERE id = :1", 42);
await db.close();
```

db.postgres.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string, sslmode?: string }): Promise<Record<string, unknown>>
```

Connect to PostgreSQL via the pure-Go pgx driver. Arg is a libpq DSN/URL string or an options object { host, port, user, password, database, sslmode }. Returns the shared SQL handle { exec, query, queryValue, begin, prepare, close }. Uses \$1,\$2,... placeholders. Pings on open.

Parameters

- `dsn` (*string* | { *host?: string, port?: number, user?: string, password?: string, database?: string, sslmode?: string* }) — A libpq DSN/URL string used verbatim, OR an options object assembled into a postgres:// URL (defaults: host localhost, port 5432, empty database; sslmode added to

the query string when set). CockroachDB and other Postgres-wire engines connect through the same driver.

Returns: Promise resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row, keyed by column name in column order); `queryValue(sql, ...params) → Promise` (first column of the first row, or null); `begin() → Promise ({ exec, query, queryValue, commit, rollback })`; `prepare(sql) → Promise ({ exec, query, queryValue, close })`; `close() → Promise`. Postgres uses \$1, \$2, ... positional placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails (the pool is closed on ping failure).

```
const db = await db.postgres.open({ host: "localhost", user: "app", password:
"s3cr3t", database: "shop", sslmode: "disable" });
const rows = await db.query("SELECT id, name FROM users WHERE id = $1", 42);
await db.close();
```

db.redis.open

```
open(url: string): Promise<Record<string, unknown>>
```

Connect to Redis (redis://...). Returns { do, ping, close }. `do(cmd, ...args)` runs any RESP command; missing key -> null. Pings on open to surface bad addresses.

Parameters

- `url (string)` — A standard Redis URL: `redis://[:password@]host:port/db` (`rediss://` for TLS), parsed by go-redis's `ParseURL`.

Returns: Promise resolving to { do, ping, close }: `do(cmd, ...args) → Promise` runs an arbitrary RESP command (the first arg is the command name, the rest its arguments) and returns the reply coerced to a JS value — strings, numbers, arrays, or null; a nil reply (missing key) resolves to null rather than throwing. `ping() → Promise ('PONG')`. `close() → Promise`.

Throws: open throws if url is empty, the URL fails to parse, or the open-time ping fails (the client is closed on ping failure). do throws on Redis-level errors (WRONGTYPE, unknown command, etc.); a missing-key nil reply is data, not an error.

```
const r = await db.redis.open("redis://localhost:6379/0");
await r.do("SET", "greeting", "hi");
const v = await r.do("GET", "greeting"); // "hi"
const missing = await r.do("GET", "nope"); // null
await r.close();
```

db.sqlite.open

```
open(path: string): Promise<Record<string, unknown>>
```

Open a SQLite database (':memory:' or a file path; created if absent). Resolves to a handle { exec, query, queryValue, begin, prepare, close }. Connection is Ping-ed before resolving.

Parameters

- `path` (*string*) — ":memory:" for an in-RAM database, or a filesystem path. Missing files are created by the modernc.org/sqlite (pure-Go, no cgo) driver.

Returns: Promise resolving to the shared SQL handle object: `exec(sql, ...params) → Promise<{ rowsAffected: number, lastInsertId: number }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row, keyed by column name in column order); `queryValue(sql, ...params) → Promise` (first column of the first row, or null when no rows match); `begin() → Promise ({ exec, query, queryValue, commit, rollback })`; `prepare(sql) → Promise ({ exec, query, queryValue, close })`; `close() → Promise`. SQLite uses ? positional placeholders. UTF-8 byte columns scan to strings; genuinely binary bytes surface as Uint8Array.

Throws: Throws if path is missing or empty, or if the connection ping fails (the *sql.DB is closed on ping failure rather than leaked). Subsequent exec/query/etc. throw the driver error on bad SQL or bind mismatch.

```
const db = await db.sqlite.open(":memory:");
await db.exec("CREATE TABLE t (id INTEGER PRIMARY KEY, name TEXT)");
await db.exec("INSERT INTO t (name) VALUES (?)", "alice");
const rows = await db.query("SELECT * FROM t WHERE name = ?", "alice");
await db.close();
```

db.valkey.open

```
open(url: string): Promise<Record<string, unknown>>
```

Connect to Valkey (valkey:// or redis://...). Valkey is the RESP-compatible Redis fork, so this is db.redis with the same { do, ping, close } handle; valkey:// / valkeys:// URLs are accepted (normalised to redis:// / rediss://). Pings on open to surface bad addresses.

Parameters

- `url` (*string*) — A connection URL: valkey://[:password@]host:port/db (valkeys:// for TLS), or the equivalent redis:// / rediss:// form. Parsed by go-redis's ParseURL after normalising the valkey scheme.

Returns: Promise resolving to { do, ping, close }: `do(cmd, ...args) → Promise` runs an arbitrary RESP command (the first arg is the command name, the rest its arguments) and returns the reply coerced to a JS value — strings, numbers, arrays, or null; a nil reply (missing key) resolves to null rather than throwing. `ping() → Promise ('PONG')`. `close() → Promise`.

Throws: open throws if url is empty, the URL fails to parse, or the open-time ping fails (the client is closed on ping failure). do throws on RESP-level errors (WRONGTYPE, unknown command, etc.); a missing-key nil reply is data, not an error.

```
const r = await db.valkey.open("valkey://localhost:6379/0");
await r.do("SET", "greeting", "hi");
const v = await r.do("GET", "greeting"); // "hi"
await r.close();
```

fs

Filesystem operations: path manipulation and archive create/extract.

fs.archive.create

```
create(destPath: string, sources: (string | { path: string, name?: string })[]):
Promise<Record<string, unknown>>
```

Create a zip / tar / tar.gz at destPath from a list of paths. Format inferred from extension.

Parameters

- `destPath` (*string*) — Output archive path. Format is inferred from the extension: .zip, .tar, .tar.gz, or .tgz.
- `sources` ((*string* | { *path*: *string*, *name*?: *string* })[]) — Non-empty array of inputs. A bare string uses the disk path as-is and its basename inside the archive; an object overrides the in-archive name via *name*. Directory sources are recursed (the directory's basename becomes the archive subdir). Archive paths always use forward slashes.

Returns: `Promise<{ path: string, format: string, entries: string[], bytes?: number }>` — *path* is *destPath*, *format* is the inferred format ("zip" | "tar" | "tar.gz"), *entries* lists the file paths written (directories excluded), and *bytes* is the final archive size when *stat* succeeds.

Throws: Rejects if *destPath* is empty, *sources* is not an array / is empty / contains a bad entry (object missing 'path', unsupported element type), the format cannot be inferred from the extension, or any disk read / write fails (e.g. a source path does not exist).

```
const r = await fs.archive.create("out.tar.gz", ["dist", { path: "README.md",
name: "docs/README.md" }]);
runtime.log(r.format, r.entries.length);
```

fs.archive.extract

```
extract(archivePath: string, destDir: string, opts?: { overwrite?: boolean }):
Promise<Record<string, unknown>>
```

Extract a zip / tar / tar.gz to *destDir*. *opts.overwrite* controls O_EXCL behaviour.

Parameters

- `archivePath` (*string*) — Path to the archive. Format is inferred from its extension (.zip, .tar, .tar.gz, .tgz).

- `destDir` (*string*) — Destination directory; created (recursively) if absent. All entries are confined to this directory via zip-slip / tar-slip protection.
- `opts` (*{ overwrite?: boolean }, optional*) — `overwrite` (default `false`) clobbers existing files; when `false`, an entry colliding with an existing file fails the call (`O_EXCL`).

Returns: `Promise<{ path: string, format: string, dest: string, entries: string[] }>` — `path` is `archivePath`, `format` is the inferred format, `dest` is `destDir`, and `entries` lists the extracted entry names (regular files only).

Throws: Rejects if `archivePath` or `destDir` is empty, the format cannot be inferred, `destDir` cannot be created, the archive cannot be opened / decoded, an entry escapes `destDir` (absolute path or `'..'` component), or (with `overwrite` `false`) an entry collides with an existing file.

```
const r = await fs.archive.extract("out.tar.gz", "./unpacked", { overwrite: true });
runtime.log(r.entries.length, "files extracted");
```

fs.path.basename

```
basename(path: string, suffix?: string): string
```

Final segment of a path; optional suffix is stripped if it matches.

Parameters

- `path` (*string*) — A forward-slash path. Trailing slashes are stripped before taking the last segment.
- `suffix` (*string, optional*) — Trailing suffix to remove from the result (e.g. an extension). Only stripped when it matches and is not the entire segment; a non-matching or empty suffix is ignored.

Returns: `string` — the last path segment, with suffix removed when it applies.

Throws: Throws a `TypeError` if `path` is missing, null, or undefined. `suffix` is coerced to a string and is optional.

```
const b = fs.path.basename("/var/log/app.log", ".log"); // "app"
```

fs.path.dirname

```
dirname(path: string): string
```

Directory portion of a path. POSIX-style; trailing slashes are stripped.

Parameters

- `path` (*string*) — A forward-slash path. On Windows, normalise separators yourself first.

Returns: `string` — everything up to (not including) the final slash; `."` when the path has no directory component, `"/"` for a rooted single segment.

Throws: Throws a `TypeError` if path is missing, null, or undefined.

```
const d = fs.path.dirname("/var/log/app.log"); // "/var/log"
```

net

Network clients and probes: HTTP, TCP/DNS/TLS/NTP/WHOIS probes, netstatus, email auth, browser-style sessions.

net.browser.open

```
open(...args: unknown[]): Promise<Record<string, unknown>>
```

Open a stateful HTTP session: { setUserAgent, setHeader, get, post, cookies }. Cookie jar + default headers persist across requests (like a browser).

Returns: `Promise<{ setUserAgent(ua: string): void, setHeader(name: string, value: string): void, get(url: string): Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, url: string }>, post(url: string, body?: string): Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, url: string }>, cookies(url: string): Promise<{ name: string, value: string }[]> }>` — a session handle backed by an `http.Client` with an automatic cookie jar (public-suffix scoped). `setUserAgent/setHeader` register default headers replayed on every request; `get/post` return the same result shape as `net.http.request`; `cookies` lists the jar's cookies for a URL.

Throws: `browser.open` rejects only if the cookie jar can't be created. `get/post` reject if the URL is empty or the request fails (transport error); `cookies` rejects if the URL is empty or unparseable. 4xx/5xx responses resolve normally.

```
const b = await net.browser.open();
b.setUserAgent("my-bot/1.0");
await b.post("https://site/login", "user=x&pass=y");
const home = await b.get("https://site/home");
runtime.log(await b.cookies("https://site/"));
```

net.capture.interfaces

```
interfaces(): { name: string; addresses: string[]; up: boolean; loopback: boolean }[]
```

List the host's network interfaces synchronously: `net.capture.interfaces()` → array of { name, addresses: string[], up, loopback }. Pure-Go (no privileges, all platforms).

Returns: { name: string, addresses: string[], up: boolean, loopback: boolean }[] — one entry per interface with its name, assigned addresses (CIDR strings), and up / loopback flags. Synchronous (not a Promise).

Throws: Throws if interface enumeration fails.

```
for (const i of net.capture.interfaces()) runtime.log(i.name, i.up);
```

net.capture.open

```
open(opts: { iface: string, promisc?: boolean, snaplen?: number, filter?: string }, onPacket: (pkt: { ts: number, length: number, captureLength: number, link: string, eth?: { src: string, dst: string, type: string }, vlan?: { id: number, priority: number, drop: boolean, type: string }, arp?: { operation: string, senderMac: string, senderIp: string, targetMac: string, targetIp: string }, ip?: { version: number, src: string, dst: string, protocol: string, ttl: number }, tcp?: { srcPort: number, dstPort: number, seq: number, ack: number, window: number, checksum: number, flags: { syn: boolean, ack: boolean, fin: boolean, rst: boolean, psh: boolean, urg: boolean }, options?: { mss?: number, windowScale?: number, sackPermitted?: boolean, timestamps?: { val: number, ecr: number } } }, udp?: { srcPort: number, dstPort: number, length: number }, icmp?: { type: number, code: number }, dns?: { id: number, qr: boolean, opcode: string, rcode: string, questions: { name: string, type: string }[], answers: { name: string, type: string, data: string }[] }, payload?: Uint8Array, bytes: Uint8Array }) => void): void
```

Live packet capture: `net.capture.open({ iface, promisc?, snaplen?, filter? }, pkt => {...})` → `Promise<{ iface, link, close() }>`. Linux + macOS only (Windows rejects); needs root / `CAP_NET_RAW` (Linux) or `/dev/bpf` access (macOS). `promisc` defaults true. The handler is called per frame with a decoded packet `{ ts, length, captureLength, link, eth?, vlan?, arp?, ip?, tcp?, udp?, icmp?, dns?, payload?, bytes }`. Optional filter is a tcpdump-like expression string (e.g. 'tcp and port 80'), evaluated post-decode in userspace — NOT a kernel BPF program, so it skips the JS callback for non-matching packets but does not avoid the kernel→userspace copy. Supports tcp/udp/icmp/ip/ip6, host/src host/dst host, net/src net/dst net (CIDR), port/src port/dst port, portrange/src portrange/dst portrange, and/or/not + parens, implicit-and between juxtaposed primaries. A malformed expression makes `open` reject. `close()` returns `Promise`. Pure-Go `gopacket` (no `libpcap/cgo`).

Parameters

- `opts` (*{ iface: string, promisc?: boolean, snaplen?: number, filter?: string }*) — `iface` is the interface name to capture on (required); `promisc` enables promiscuous mode (default true); `snaplen` caps the per-packet capture length in bytes (default 262144); `filter` is an optional tcpdump-like expression (e.g. 'tcp and port 80') applied post-decode in userspace — supports tcp/udp/icmp/ip/ip6, host/net (CIDR), port/portrange, src/dst directions, and/or/not + parens.
- `onPacket` (*((pkt: { ts: number, length: number, captureLength: number, link: string, eth?: { src: string, dst: string, type: string }, vlan?: { id: number, priority: number, drop: boolean, type: string }, arp?: { operation: string, senderMac: string, senderIp: string, targetMac: string, targetIp: string }, ip?: { version: number, src: string, dst: string, protocol: string, ttl: number }, tcp?: { srcPort: number, dstPort: number, seq: number, ack: number, window: number, checksum: number, flags: { syn: boolean, ack: boolean, fin: boolean, rst: boolean, psh: boolean, urg: boolean }, options?: { mss?: number, windowScale?: number, sackPermitted?: boolean, timestamps?: { val: number, ecr: number } } }, udp?: { srcPort: number, dstPort: number, length: number }, icmp?: { type: number, code: number }, dns?: { id: number, qr: boolean, opcode: string, rcode: string, questions: { name: string, type: string }[], answers: { name: string, type: string, data: string }[] }, payload?: Uint8Array, bytes: Uint8Array }) => void)*) — Called once per matching frame with the

decoded packet. `ts` is epoch ms; layer keys (`eth/vlan/arp/ip/tcp/udp/icmp/dns`) are present only when that layer decoded; `bytes` is always the raw frame; `payload` is the application-layer bytes when present.

Returns: `Promise<{ iface: string, link: string, close(): Promise }>` — a live-capture handle. `link` is the link-type name; `close()` stops the capture and resolves when the source is torn down. The handler keeps firing until `close()` is called or the source errors.

Throws: Rejects if `iface` is missing, the filter expression is malformed, the platform is unsupported (Windows), or the capture can't be opened (missing `root / CAP_NET_RAW` on Linux, `/dev/bpf` access on macOS). Throws synchronously if `onPacket` is not a function.

```
const cap = await net.capture.open({ iface: "en0", filter: "tcp and port 443" },
  pkt => {
    runtime.log(pkt.ip?.src, "→", pkt.ip?.dst);
  });
await cap.close();
```

net.capture.openFile

```
openFile(path: string, onPacket: (pkt: { ts: number, length: number,
  captureLength: number, link: string, eth?: object, vlan?: object, arp?: object,
  ip?: object, tcp?: object, udp?: object, icmp?: object, dns?: object, payload?:
  Uint8Array, bytes: Uint8Array }) => void, opts?: { filter?: string }): void
```

Read a `.pcap` / `.pcapng` file: `net.capture.openFile(path, pkt => {...}, opts?)` → `Promise`. Calls the handler once per decoded packet (same shape as `capture.open`) and resolves at EOF. Offline; no privileges. `opts` is an optional trailing arg `{ filter? }` — the 2-arg form still works; `filter` is the same tcpdump-like expression string as `capture.open` (post-decode/userspace, not kernel BPF; supports `host/net CIDR + port/portrange`; malformed → rejects).

Parameters

- `path` (*string*) — Path to a `.pcap` or `.pcapng` file; the format is auto-detected from the magic bytes.
- `onPacket` ((*pkt*: { *ts*: number, *length*: number, *captureLength*: number, *link*: string, *eth*?: object, *vlan*?: object, *arp*?: object, *ip*?: object, *tcp*?: object, *udp*?: object, *icmp*?: object, *dns*?: object, *payload*?: *Uint8Array*, *bytes*: *Uint8Array* }) => void) — Called once per decoded packet (same shape as `capture.open`'s handler).
- `opts` ({ *filter*?: string }, optional) — `filter` is the same tcpdump-like expression as `capture.open`, applied post-decode in userspace; omit (2-arg form) to deliver every packet.

Returns: `Promise` — resolves at end-of-file after dispatching every (matching) packet to the handler.

Throws: Rejects if the filter expression is malformed, the file can't be opened or parsed, or the handler throws. Throws synchronously if `onPacket` is not a function.

```
await net.capture.openFile("/tmp/dump.pcap", pkt =>
  runtime.log(pkt.tcp?.dstPort), { filter: "tcp" });
```

net.capture.routes

```
routes(): { destination: string; gateway: string; interface: string; family:
"ip" | "ip6"; metric: number }[]
```

List the host's IP routing table synchronously: `net.capture.routes()` → array of { destination, gateway, interface, family, metric }. Pure-Go, unprivileged: Linux reads `/proc/net/route` + `/proc/net/ipv6_route`; macOS/BSD read the routing socket via `x/net/route`. Windows is unsupported (throws).

Returns: { destination: string, gateway: string, interface: string, family: "ip" | "ip6", metric: number }[] — one entry per route. destination is a CIDR ('0.0.0.0/0' for the default route, ':::0' for the IPv6 default); gateway is the next-hop IP or '' for a directly-connected/link route; interface is the outgoing NIC name (best-effort); metric is 0 when the platform doesn't report one (BSD/macOS). Synchronous (not a Promise).

Throws: Throws if route enumeration fails or the platform is unsupported (Windows).

```
const def = net.capture.routes().find(r => r.destination ===
"0.0.0.0/0");\nruntime.log("default via", def?.gateway, "on", def?.interface);
```

net.capture.toFile

```
toFile(path: string, opts?: { snaplen?: number, linkType?: number }): {
write(bytes: string | Uint8Array, opts?: { ts?: number }): void; close():
Promise<void> }
```

Write raw frames to a .pcap file: `net.capture.toFile(path, { linkType?, snaplen? })` → { `write(bytes, { ts? })`, `close()` }. `write` appends a raw frame (Uint8Array); `ts` (ms) overrides the timestamp. `close()` flushes and returns Promise. Offline; no privileges.

Parameters

- `path` (*string*) — Path of the .pcap file to create (overwritten if it exists).
- `opts` (*{ snaplen?: number, linkType?: number }, optional*) — `snaplen` is the pcap global-header snap length (default 262144); `linkType` is the numeric pcap link-type written into the header (default Ethernet).

Returns: { `write(bytes, opts?): void, close(): Promise` } — a writer handle (returned synchronously, not a Promise). `write` appends one raw frame to the file (`opts.ts` in ms overrides the timestamp, defaulting to now) and returns undefined. `close()` flushes and closes the file, resolving when done.

Throws: Throws synchronously if the file can't be created or the header can't be written, and on a per-frame write error. `write` throws if called after `close`. `close()` rejects on a close error.

```
const w = net.capture.toFile("/tmp/out.pcap");
w.write(frameBytes, { ts: Date.now() });
await w.close();
```

net.email.all

```
all(domain: string): Promise<Record<string, unknown>>
```

Run all five email probes in parallel — five-way handshake aggregate.

Parameters

- `domain` (*string*) — The domain to run all five email-auth probes against.

Returns: `Promise<{ domain: string, spf: object, dmarc: object, mtaSts: object, tlsRpt: object, bimi: object }>` — domain echoes the input; each probe key holds that probe's result (the same shape its individual binding returns) or `{ error: string }` when that single probe failed. A per-probe failure doesn't fail the aggregate.

Throws: Resolves even when individual probes fail (their failure surfaces under `.error`). Always resolves.

```
const a = await net.email.all("example.com"); runtime.log(a.spf, a.dmarc);
```

net.email.bimi

```
bimi(domain: string, opts?: { selector?: string }): Promise<Record<string, unknown>>
```

Probe BIMI: `TXT(_bimi)`; selector defaults to 'default'.

Parameters

- `domain` (*string*) — The domain whose `._bimi` TXT record is queried for a v=BIMI1 record.
- `opts` (*{ selector?: string }, optional*) — selector names the BIMI selector to query (default 'default').

Returns: `Promise<{ present: false, selector: string } | { present: true, selector: string, record: string, tags: Record<string, string>, l: string, a: string }>` — selector echoes the queried selector; when found, `l` is the logo URL tag and `a` is the assertion (VMC) tag. Missing record resolves to `{ present: false, selector }`.

Throws: Rejects only on a DNS lookup error other than `NXDOMAIN`; an absent record resolves to `{ present: false, selector }`.

```
const r = await net.email.bimi("example.com", { selector: "v1" });  
runtime.log(r.present && r.l);
```

net.email.dmarc

```
dmarc(domain: string): Promise<Record<string, unknown>>
```

Query `TXT(_dmarc)` and parse policy / pct / rua / ruf tags.

Parameters

- `domain` (*string*) — The domain whose `_dmarc`.TXT record is queried for a `v=DMARC1` record.

Returns: `Promise<{ present: false } | { present: true, record: string, tags: Record<string, string>, policy: string, subdomain: string, percent: string, rua: string, ruf: string }>` — `tags` is the full parsed tag map; `policy/subdomain/percent/rua/ruf` surface the common `p/sp/pct/rua/ruf` tags. Missing record resolves to `{ present: false }`.

Throws: Rejects only on a DNS lookup error other than `NXDOMAIN`; an absent record resolves to `{ present: false }`.

```
const r = await net.email.dmarc("example.com"); runtime.log(r.present &&
r.policy);
```

net.email.mtaSts

```
mtaSts(domain: string): Promise<Record<string, unknown>>
```

Probe MTA-STS: `TXT(_mta-sts.)` plus the fetched policy file.

Parameters

- `domain` (*string*) — The domain whose `_mta-sts`.TXT record and well-known policy file are probed.

Returns: `Promise<{ present: false } | { present: true, record: string, txt: { v: string, id: string }, policy?: { version?: string, mode?: string, mx?: string[], maxAge?: number | string }, policyError?: string }>` — `txt` carries the versioned id from the TXT marker; `policy` is the parsed well-known file (`mode + mx + maxAge`), or `policyError` holds the fetch/parse error string when the file couldn't be retrieved. Missing TXT resolves to `{ present: false }`.

Throws: Rejects only on a DNS lookup error other than `NXDOMAIN`; an absent record resolves to `{ present: false }`. A policy-file fetch failure is captured in `policyError`, not thrown.

```
const r = await net.email.mtaSts("example.com"); runtime.log(r.present &&
r.policy?.mode);
```

net.email.send

```
send(opts: { to: string | string[], from: string, subject?: string, body?:
string, html?: string, attachments?: { filename: string, contentType?: string,
bytes: Uint8Array | ArrayBuffer }[], headers?: Record<string, string>, server: {
host: string, port?: number, auth?: { username: string, password: string },
tls?: "starttls" | "tls" | "none" }, timeout?: number }): Promise<{ accepted:
string[]; rejected: { address: string; reason: string }[] }>
```

Send an outbound email: `net.email.send({to, from, subject, body, html?, attachments?, headers?, server: {host, port?, auth?, tls?}, timeout?}) → Promise<{accepted: string[], rejected: [{address, reason}]]>`. One TCP connection per call; per-recipient outcome captured. Transport failures throw; per-RCPT rejections surface in the result. TLS modes: `starttls` (default), `tls`, `none`.

Parameters

- `opts` (*{ to: string | string[], from: string, subject?: string, body?: string, html?: string, attachments?: { filename: string, contentType?: string, bytes: Uint8Array | ArrayBuffer }[], headers?: Record<string, string>, server: { host: string, port?: number, auth?: { username: string, password: string }, tls?: "starttls" | "tls" | "none" }, timeout?: number }*) — `to` (string or array) and `from` are required, as is `server.host`. `subject/body/html` shape the message: `body` alone → `text/plain`, `body+html` → `multipart/alternative`, any attachments → `multipart/mixed`. attachments carry raw bytes (`contentType` defaults to `application/octet-stream`). `headers` adds custom headers (CR/LF stripped). `server.port` defaults to 587; `server.auth` enables PLAIN auth (skipped when `tls` is 'none'); `server.tls` picks the transport: 'starttls' (default), implicit 'tls', or 'none'. `timeout` is the dial / connection timeout in ms (default 30000).

Returns: `Promise<{ accepted: string[], rejected: { address: string, reason: string }[] }>` — `accepted` lists recipients the server accepted at RCPT TO; `rejected` pairs each refused address with the server's reason. The DATA body is sent only when at least one recipient was accepted.

Throws: Rejects on missing required fields (`to` / `from` / `server.host`), transport / protocol failures (dial, HELO, STARTTLS unavailable when requested, AUTH, MAIL FROM, DATA), or timeout. Per-recipient RCPT rejections are returned in `rejected`, not thrown.

```
const r = await net.email.send({
  to: "to@example.com", from: "from@example.com", subject: "hi", body: "hello",
  server: { host: "smtp.example.com", port: 587, auth: { username: "u",
    password: "p" } },
});
runtime.log(r.accepted, r.rejected);
```

net.email.spf

```
spf(domain: string): Promise<Record<string, unknown>>
```

Query TXT() for SPF, return record + parsed mechanisms + all-policy.

Parameters

- `domain` (*string*) — The domain whose apex TXT records are queried for an SPF (v=spf1) record.

Returns: `Promise<{ present: false } | { present: true, record: string, mechanisms: string[], allPolicy: string }>` — when found, `record` is the raw SPF string, `mechanisms` is the tokenised list after `v=spf1`, and `allPolicy` summarises the trailing all-style mechanism (pass / fail / softfail / neutral). Missing record resolves to `{ present: false }`.

Throws: Rejects only on a DNS lookup error other than NXDOMAIN; an absent record resolves to `{ present: false }`.

```
const r = await net.email.spf("example.com"); if (r.present)
  runtime.log(r.allPolicy);
```

net.email.tlsRpt

```
tlsRpt(domain: string): Promise<Record<string, unknown>>
```

Probe TLS-RPT: TXT(_smtp._tls.) and parse rua.

Parameters

- `domain` (*string*) — The domain whose _smtp._tls. TXT record is queried for a v=TLSRPTv1 record.

Returns: `Promise<{ present: false } | { present: true, record: string, tags: Record<string, string>, rua: string }>` — `tags` is the parsed tag map; `rua` surfaces the report-URI tag. Missing record resolves to `{ present: false }`.

Throws: Rejects only on a DNS lookup error other than `NXDOMAIN`; an absent record resolves to `{ present: false }`.

```
const r = await net.email.tlsRpt("example.com"); runtime.log(r.present && r.rua);
```

net.http.get

```
get(url: string): Promise<Record<string, unknown>>
```

Perform an HTTP GET with a 5-second default timeout. Returns `{ status, body }`.

Parameters

- `url` (*string*) — Absolute request URL (<http://> or <https://>).

Returns: `Promise<{ status: number, body: string }>` — the HTTP status code and the response body as a string. Redirects are followed by the default client.

Throws: Rejects on transport errors (DNS failure, connection refused, TLS handshake) or if the 5s context deadline is exceeded. 4xx/5xx responses do NOT reject — they surface via status.

```
const r = await net.http.get("https://example.com"); runtime.log(r.status);
```

net.http.post

```
post(url: string, body?: string): Promise<Record<string, unknown>>
```

Perform an HTTP POST with a 5-second default timeout. Returns `{ status, body }`.

Parameters

- `url` (*string*) — Absolute request URL (<http://> or <https://>).
- `body` (*string, optional*) — Request body sent verbatim; omit or pass empty for no body. No Content-Type header is set automatically.

Returns: `Promise<{ status: number, body: string }>` — the HTTP status code and the response body as a string.

Throws: Rejects on transport errors (DNS failure, connection refused, TLS handshake) or if the 5s context deadline is exceeded. 4xx/5xx responses do NOT reject.

```
const r = await net.http.post("https://api.example.com/x", JSON.stringify({ a: 1 }));
```

net.http.request

```
request(method: string, url: string, opts?: { headers?: Record<string, string>, body?: string, timeout?: number, retry?: number, follow?: boolean, username?: string, password?: string }): Promise<Record<string, unknown>>
```

Full HTTP client: method, url, opts {headers, body, timeout, retry, follow, username, password}. Returns {status, ok, headers, body, url}. 4xx/5xx dont throw; retry covers transport errors + 5xx.

Parameters

- `method` (*string*) — HTTP method (GET, POST, PUT, ...); upper-cased internally. Required.
- `url` (*string*) — Absolute request URL. Required.
- `opts` (*{ headers?: Record<string, string>, body?: string, timeout?: number, retry?: number, follow?: boolean, username?: string, password?: string }, optional*) — headers sets request headers; body is the raw request body; timeout is the per-attempt client timeout in ms (default 30000); retry is the number of extra attempts (default 0) applied only to transport errors and 5xx with linear backoff capped at 1s; follow toggles redirect following (default true — false stops at the first 3xx); username/password set HTTP Basic auth.

Returns: `Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, url: string }>` — status is the final status code; ok is status in [200,400); headers is a lower-cased name → value map (last value wins, alphabetically ordered); body is the response text; url is the final URL after redirects.

Throws: Rejects on transport errors (DNS, connection refused, TLS) or context deadline, and after exhausting retries on a persistent transport error / 5xx. A malformed method or URL rejects immediately (not retried). 4xx/5xx that succeed at the transport level resolve normally.

```
const r = await net.http.request("POST", "https://api.example.com", { headers: { "content-type": "application/json" }, body: "{}", retry: 2 });
```

net.icmp.open

```
open(opts?: { network?: "ip4" | "ip6", readBuffer?: number }): void
```

Open a raw ICMP socket: `net.icmp.open(opts?)` → `Promise`. Requires root / `CAP_NET_RAW` (open rejects otherwise). `opts` { `network?: 'ip4'|'ip6'` (default 'ip4'), `readBuffer?` }. `handle.send(opts)` writes a message in one of two modes: Echo mode { `to`, `type?`, `code?`, `id?`, `seq?`, `payload?` } (type defaults to the network's echo request), or raw mode { `to`, `type`, `code?`, `body` } where `body` (`Uint8Array|string`) is

marshalled verbatim (`icmp.RawBody`) for hand-built non-Echo messages such as destination-unreachable — in raw mode type is required and body is mutually exclusive with id/seq/payload. push/callback model — `onMessage(cb)` events carry { address, type, code }; `onClose(cb)/onError(cb)`; `handle.network/local`; `handle.close()`.

Parameters

- `opts` (*{ network?: "ip4" | "ip6", readBuffer?: number }, optional*) — network selects the IP version (default 'ip4'); `readBuffer` is the inbound channel capacity (default 64).

Returns: `Promise<{ network: string, local: string, send(opts: { to: string, type?: number, code?: number, id?: number, seq?: number, payload?: string | Uint8Array, body?: string | Uint8Array }): Promise, onMessage(cb: (ev: { bytes: Uint8Array, text: string, address: string, type: number, code: number }) => void): void, onClose, onError, close(): void }>` — a raw-ICMP handle. `send` writes an ICMP message: omit body for an Echo-shaped body (type defaults to the network's echo request, id/seq/payload optional); provide body for a verbatim raw body (type required, mutually exclusive with id/seq/payload) to hand-build non-Echo messages such as destination-unreachable. `to` is the destination address. `onMessage` fires per received packet with the marshalled body plus { address, type, code } meta.

Throws: Rejects if the raw socket can't be opened — typically because it needs root / `CAP_NET_RAW`. `send` rejects on resolve/marshal/write errors and throws synchronously: after close, if `opts.to` is missing, if a raw body is sent without `opts.type`, or if body is combined with id/seq/payload. Read errors surface via `onError`.

```
const p = await net.icmp.open();
p.onMessage(ev => runtime.log(ev.address, ev.type));
await p.send({ to: "8.8.8.8", id: 1, seq: 1, payload: "ping" });
// raw (non-Echo) body — e.g. a hand-built destination-unreachable:
await p.send({ to: "8.8.8.8", type: 3, code: 1, body: new Uint8Array([0, 0, 0, 0]) });
```

net.netstatus.check

```
check(host: string, opts?: { port?: string, timeout?: number }):
Promise<Record<string, unknown>>
```

Run DNS / TCP / TLS / HTTP against one host concurrently. Returns { reachable, dns, tcp, tls, http } — each sub-probe ok+error; reachable = dns.ok AND tcp.ok. Sub-failures are data, not throws.

Parameters

- `host` (*string*) — The host to check. Required.
- `opts` (*{ port?: string, timeout?: number }, optional*) — port is the TCP/TLS port (default "443"); timeout bounds all four sub-probes in ms (default 10000).

Returns: `Promise<{ host: string, port: string, elapsedMs: number, reachable: boolean, dns: { ok: boolean, ips: string[], error?: string }, tcp: { ok: boolean, latencyMs: number, error?: string }, tls: { ok: boolean, daysRemaining: number, error?: string }, http: { ok: boolean, status: number, error?: string } }>`

— the four sub-probe results plus elapsed time. `reachable` is `dns.ok` AND `tcp.ok`; TLS/HTTP are reported but don't gate it. Each sub-probe carries its own error string instead of failing the call.

Throws: Rejects only if host is empty. Sub-probe failures are captured as data (`ok:false + error`) rather than thrown.

```
const s = await net.netstatus.check("example.com"); runtime.log(s.reachable);
```

net.probe.dns

```
dns(host: string, opts?: { types?: string[] }): Promise<Record<string, unknown>>
```

Look up A / AAAA / MX / TXT / CNAME / NS records. Default: all five.

Parameters

- `host` (*string*) — The hostname to resolve.
- `opts` (*{ types?: string[] }, optional*) — `types` restricts the lookup to a subset (case-insensitive: `'a','aaaa','mx','txt','cname','ns'`); omit to query all.

Returns: `Promise<{ a?: string[], aaaa?: string[], mx?: { preference: number, host: string }[], txt?: string[], cname?: string, ns?: string[] }>` — each key is present only when that record type returned at least one entry, so use `"mx"` in `result` to test presence.

Throws: Resolves with an object omitting record types that errored or were empty; per-type lookup failures are swallowed (not thrown). Always resolves.

```
const r = await net.probe.dns("example.com", { types: ["a", "mx"] });
```

net.probe.ntp

```
ntp(host: string, opts?: { timeout?: number, port?: number | string }):  
Promise<Record<string, unknown>>
```

Query an NTPv4 server (UDP 123) and report offset, RTT, stratum, root delay / dispersion.

Parameters

- `host` (*string*) — The NTP server hostname or IP.
- `opts` (*{ timeout?: number, port?: number | string }, optional*) — `timeout` is the query timeout in ms (default 5000); `port` overrides the default UDP port 123.

Returns: `Promise<{ serverTime: string, offsetMs: number, rttMs: number, stratum: number, referenceTime: string, rootDelayMs: number, rootDispersionMs: number }>` — the server's time and reference time (RFC3339 nanos), clock offset and round-trip in ms, the stratum, and the root delay / dispersion in ms.

Throws: Rejects if the NTP query fails (unreachable server, timeout, malformed response).

```
const r = await net.probe.ntp("pool.ntp.org"); runtime.log(r.offsetMs);
```

net.probe.ping

```
ping(host: string, opts?: { mode?: "tcp" | "icmp" | "udp", count?: number,
  timeout?: number, port?: string }): Promise<{ host: string; ip: string; mode:
  string; sent: number; received: number; lossPercent: number; minMs: number;
  avgMs: number; maxMs: number }>
```

Reachability probe. mode tcp (default; dials host:port), icmp (real ICMP echo, needs raw-socket privileges), or udp (sends a datagram to a closed port and counts ICMP port-unreachable as reachable, needs root / CAP_NET_RAW). Returns { sent, received, lossPercent, minMs, avgMs, maxMs }. Unreachable = received 0, no throw.

Parameters

- `host` (*string*) — The target host. Required.
- `opts` (*{ mode?: "tcp" | "icmp" | "udp", count?: number, timeout?: number, port?: string }, optional*) — mode selects the probe (default 'tcp' — opens count TCP connections; 'icmp' sends real ICMP echo and needs raw-socket privileges; 'udp' sends a datagram to a closed port and counts the ICMP port-unreachable reply as reachable, needs root / CAP_NET_RAW); count is the number of probes (default 4); timeout is the per-probe timeout in ms (default 5000); port is the TCP target port (default "80", tcp mode only).

Returns: Promise<{ host: string, ip: string, mode: string, sent: number, received: number, lossPercent: number, minMs: number, avgMs: number, maxMs: number }> — the resolved IP, the mode used, packets sent/received, loss percentage, and min/avg/max RTT in ms. A fully unreachable host resolves with received:0 and lossPercent:100 rather than rejecting.

Throws: Rejects if host is empty, mode is not one of 'tcp', 'icmp', or 'udp', DNS resolution fails (tcp mode), or the raw ICMP socket can't be opened (icmp/udp modes; typically missing raw-socket privileges). Individual lost packets are counted, not thrown.

```
const p = await net.probe.ping("example.com", { count: 3 });
runtime.log(p.lossPercent);
```

net.probe.smtp

```
smtp(host: string, opts?: { port?: string, timeout?: number, ehloName?: string
  }): Promise<{ host: string; port: string; banner: string; ehloDomain: string;
  extensions: string[]; starttls: boolean; authMechanisms: string[]; sizeLimit:
  number }>
```

SMTP capability probe (no mail sent). EHLO + parse extensions. Returns { banner, ehloDomain, extensions, starttls, authMechanisms, sizeLimit }. Connection failures throw.

Parameters

- `host` (*string*) — The SMTP server host. Required.

- `opts` (*{ port?: string, timeout?: number, ehloName?: string }, optional*) — port is the SMTP port (default "25"); timeout bounds the whole conversation in ms (default 10000); ehloName is the domain sent in EHLO (default "localhost").

Returns: `Promise<{ host: string, port: string, banner: string, ehloDomain: string, extensions: string[], starttls: boolean, authMechanisms: string[], sizeLimit: number }>` — the greeting banner, the server's EHLO greeting line, the raw advertised extension lines, whether STARTTLS is offered, the upper-cased AUTH mechanism names, and the SIZE limit (0 if unadvertised). No mail is sent.

Throws: Rejects if host is empty, the dial fails, or the greeting / EHLO cannot be read. A server that simply omits STARTTLS or AUTH reports them as false / empty — a finding, not an error.

```
const s = await net.probe.smtp("mail.example.com"); runtime.log(s.starttls,
s.authMechanisms);
```

net.probe.tcp

```
tcp(target: string, opts?: { timeout?: number, port?: string }): Promise<{ host:
string; port: number; ip: string; latencyMs: number }>
```

Dial a TCP target and report latency + resolved IP. Default timeout 5s.

Parameters

- `target` (*string*) — host:port to dial; a bare host uses `opts.port` (default 80).
- `opts` (*{ timeout?: number, port?: string }, optional*) — timeout is the dial timeout in ms (default 5000); port is the fallback port when target has no :port (default "80").

Returns: `Promise<{ host: string, port: number, ip: string, latencyMs: number }>` — the parsed host, port, the resolved remote IP, and the connect latency in milliseconds.

Throws: Rejects if the dial fails (refused, unreachable, name resolution failure, or timeout).

```
const r = await net.probe.tcp("example.com:443"); runtime.log(r.latencyMs);
```

net.probe.tls

```
tls(target: string, opts?: { timeout?: number }): Promise<Record<string,
unknown>>
```

Open a TLS connection (InsecureSkipVerify; for probing only) and return the cert chain summary.

Parameters

- `target` (*string*) — host:port to dial; a bare host uses port 443. The host is sent as SNI.
- `opts` (*{ timeout?: number }, optional*) — timeout is the dial timeout in ms (default 5000).

Returns: `Promise<{ cn: string, issuer: string, notBefore: string, notAfter: string, daysRemaining: number, dnsNames: string[], serialNumber: string, fingerprintSha256: string }>` — leaf-certificate fields: common name, issuer CN, validity bounds (RFC3339), days until expiry, SAN DNS names, decimal

serial, and the SHA-256 fingerprint as hex. Verification is skipped, so expired / mismatched certs still report.

Throws: Rejects if the dial / handshake fails, the connection is not TLS, or no peer certificates are presented.

```
const c = await net.probe.tls("example.com:443"); runtime.log(c.daysRemaining);
```

net.probe.traceroute

```
traceroute(host: string, opts?: { protocol?: "icmp" | "udp" | "tcp", port?: number, maxHops?: number, timeout?: number, probes?: number }): Promise<{ ttl: number; address: string | null; rttsMs: number[]; reached: boolean }[]>
```

Trace the network path to a host: `net.probe.traceroute(host, opts?) → Promise<hop[]>`. Sends probes with increasing TTL and reports each responding router. Needs root / CAP_NET_RAW (intermediate hops are seen via ICMP time-exceeded). `opts { protocol?: 'icmp'|'udp'|'tcp' (default 'icmp'), port?: number (udp 33434 / tcp 80), maxHops?: number (30), timeout?: number ms per probe (2000), probes?: number per hop (3) }`. IPv4 only.

Parameters

- `host` (*string*) — The destination host or IP.
- `opts` (*{ protocol?: "icmp" | "udp" | "tcp", port?: number, maxHops?: number, timeout?: number, probes?: number }, optional*) — protocol selects the probe type (icmp echo, udp to an incrementing high port, or tcp SYN via a TTL-limited connect). port is the udp/tcp target (ignored for icmp). maxHops caps the trace. timeout is the per-probe wait in ms. probes is the number of probes per hop.

Returns: One entry per hop (TTL 1..n): ttl is the hop number; address is the responding router/host IP (null if every probe at that TTL timed out); rttsMs are the round-trip times of the probes that answered; reached is true on the hop where the destination itself replied (the array ends there or at maxHops).

Throws: Rejects if the host doesn't resolve, the protocol is unknown, or the raw ICMP socket can't be opened (needs root / CAP_NET_RAW). Per-hop timeouts are normal (address: null), not errors.

```
// needs root / CAP_NET_RAW
const hops = await net.probe.traceroute("1.1.1.1", { protocol: "icmp", maxHops: 20 });
for (const h of hops) runtime.log(h.ttl, h.address ?? "*", h.rttsMs);
```

net.probe.whois

```
whois(domain: string, opts?: { timeout?: number }): Promise<Record<string, unknown>>
```

Two-hop WHOIS via the IANA referral, returning the parsed record plus the raw response text.

Parameters

- `domain` (*string*) — The domain (or IP / ASN) to look up.
- `opts` (*{ timeout?: number }, optional*) — timeout is the wire-level WHOIS client timeout in ms (default 10000).

Returns: `Promise<{ raw: string, domain?: { name: string, punycode: string, whoisServer: string, nameServers: string[], status: string[], dnssec: boolean, createdAt: string, updatedAt: string, expirationDate: string }, registrar?: { name: string } }>` — raw is always the full WHOIS text; domain and registrar are best-effort parsed fields, omitted for TLDs the parser doesn't recognise.

Throws: Rejects if the WHOIS query itself fails (no referral, connection error, timeout). A parse failure is non-fatal — only raw is returned.

```
const w = await net.probe.whois("example.com");
runtime.log(w.domain?.expirationDate);
```

net.probe.wss

```
wss(url: string, opts?: { timeout?: number, ping?: boolean }): Promise<{ url:
string; connected: boolean; subprotocol: string; status: number; handshakeMs:
number; pingMs: number }>
```

WebSocket handshake probe. Opens `ws://wss://` connection, optional ping/pong RTT. Returns `{ connected, subprotocol, status, handshakeMs, pingMs }`. Failed handshake throws.

Parameters

- `url` (*string*) — The WebSocket URL (`ws://` or `wss://`). Required.
- `opts` (*{ timeout?: number, ping?: boolean }, optional*) — timeout bounds the handshake and ping in ms (default 10000); ping toggles the ping/pong RTT measurement (default true).

Returns: `Promise<{ url: string, connected: boolean, subprotocol: string, status: number, handshakeMs: number, pingMs: number }>` — connected is true on a successful upgrade, subprotocol is the negotiated subprotocol (or empty), status is the HTTP status of the 101 upgrade, handshakeMs is the handshake time in ms, and pingMs is the ping/pong RTT (or -1 when the ping was skipped or unanswered). The connection is closed immediately.

Throws: Rejects if url is empty or the handshake fails (non-101, refused, bad URL). A failed ping leaves pingMs at -1 rather than rejecting.

```
const w = await net.probe.wss("wss://echo.websocket.org");
runtime.log(w.handshakeMs);
```

net.raw.open

```
open(opts?: { iface?: string, filter?: string, readBuffer?: number }): Promise<{
link: string; send(spec: object | Uint8Array): Promise<{ bytesSent: number }>;
onPacket(cb: (pkt: any) => void): void; onClose(cb: () => void): void;
onError(cb: (msg: string) => void): void; close(): Promise<void> }>
```

Open a raw IPv4 packet engine: `net.raw.open({ iface?, filter?, readBuffer? })` → Promise. Sends crafted IPv4 packets (TCP flags / UDP / arbitrary IP protocol) via an `IP_HDRINCL` raw socket and receives replies via the capture path. Needs root / `CAP_NET_RAW`; Linux + macOS only (Windows rejects). `iface` defaults to the auto-detected default-route interface; `filter` is a tcpdump-like expression narrowing on packet. The handle: `send(specOrBytes)` → Promise<{ bytesSent }>; `onPacket(cb)` delivers a decoded packet (same shape as `net.capture`); `onClose/onError`; `close()` → Promise. `send spec`: { `dst`, `dstPort?`, `srcPort?`, `src?`, `proto?`: 'tcp'|'udp'|'ip', `protocol?`, `flags?`: string[], `seq?`, `ack?`, `window?`, `ttl?`, `ipId?`, `payload?` }; or pass a `Uint8Array` to send a full IPv4 packet verbatim. Default flags ['SYN'], `ttl` 64, `window` 65535, `src` = egress IP, `srcPort` = random high.

Parameters

- `opts` ({ `iface?`: string, `filter?`: string, `readBuffer?`: number }, optional) — `iface` is the capture/egress interface (auto-detected if omitted); `filter` is a tcpdump-like expression evaluated post-decode; `readBuffer` sizes the inbound channel (default 64).

Returns: A handle: `link` is the capture link type; `send` crafts+fires a packet (structured spec or raw bytes) and resolves { `bytesSent` }; `onPacket` receives decoded reply packets; `close()` tears down the send socket and capture.

Throws: Rejects if the platform is Windows, the raw socket or capture can't be opened (needs root / `CAP_NET_RAW`), the egress interface can't be detected, or the filter is malformed. `send` throws on an invalid spec (missing `dst`, unknown TCP flag, bad port/ttl) and rejects on a write failure.

```
// needs root / CAP_NET_RAW
const h = await net.raw.open({ filter: "tcp and src port 443" });
h.onPacket(p => runtime.log(p.ip.src, p.tcp.flags));
await h.send({ dst: "93.184.216.34", dstPort: 443, flags: ["SYN"] });
await new Promise(r => setTimeout(r, 1000));
await h.close();
```

net.raw.tcp

```
tcp(host: string, port: number, opts?: { flags?: string[], srcPort?: number,
src?: string, seq?: number, ttl?: number, payload?: Uint8Array | string,
timeout?: number, iface?: string }): Promise<{ ts: number; link: string; ip: {
src: string; dst: string; protocol: string; ttl: number }; tcp: { srcPort:
number; dstPort: number; seq: number; ack: number; flags: { syn: boolean; ack:
boolean; fin: boolean; rst: boolean; psh: boolean; urg: boolean } }; payload?:
Uint8Array; bytes: Uint8Array } | null>
```

One-shot raw TCP probe: `net.raw.tcp(host, port, opts?)` → Promise<reply | null>. Sends a single crafted TCP segment (default a SYN) and resolves with the first reply packet correlated by the 4-tuple, or null on timeout. SYN → SYN/ACK means open; RST means closed; null means filtered/no answer. Needs root / `CAP_NET_RAW`; Linux + macOS only.

Parameters

- `host` (string) — Destination host or IPv4 address. Required.
- `port` (number) — Destination TCP port. Required.

- `opts` (*{ flags?: string[], srcPort?: number, src?: string, seq?: number, ttl?: number, payload?: Uint8Array | string, timeout?: number, iface?: string }, optional*) — flags are the TCP flags to set (default ['SYN']); src/srcPort/seq/ttl/payload tune the crafted segment; timeout is the reply wait in ms (default 2000); iface overrides the auto-detected capture interface.

Returns: The decoded reply packet (same shape as `net.capture` packets), or null if no correlated reply arrived within the timeout. A SYN/ACK indicates the port is open; an RST indicates closed.

Throws: Rejects if the platform is Windows, host is empty, port is out of range, DNS resolution fails, or the raw socket / capture can't be opened (needs root / CAP_NET_RAW). A timeout is not an error — it resolves null.

```
// needs root / CAP_NET_RAW
const reply = await net.raw.tcp("scanme.nmap.org", 80, { flags: ["SYN"] });
if (reply && reply.tcp.flags.syn && reply.tcp.flags.ack) runtime.log("open");
else if (reply && reply.tcp.flags.rst) runtime.log("closed");
else runtime.log("filtered / no answer");
```

net.tcp.connect

```
connect(host: string, port: string | number, opts?: { timeout?: number,
readBuffer?: number }): void
```

Open a TCP client socket: `net.tcp.connect(host, port, opts?)` → Promise. Push/callback read model — `handle.onData(cb)/onClose(cb)/onError(cb)` register listeners; `handle.write(data)` sends (string→UTF-8 / Uint8Array); `handle.remote/local` are the peer/local addresses; `handle.close()` shuts down. `opts { timeout?, readBuffer? }`.

Parameters

- `host` (*string*) — The remote host to dial.
- `port` (*string | number*) — The remote port.
- `opts` (*{ timeout?: number, readBuffer?: number }, optional*) — timeout is the dial timeout in ms (default 10000); readBuffer is the inbound channel capacity (default 64).

Returns: Promise<{ remote: string, local: string, write(data: string | Uint8Array): Promise, onData(cb: (ev: { bytes: Uint8Array, text: string }) => void): void, onClose(cb: () => void): void, onError(cb: (err: string) => void): void, close(): void }> — a connected-socket handle. remote/local are the peer/local addresses. write resolves once the bytes are written. onData fires per inbound chunk with both a Uint8Array and a UTF-8 text view; onClose fires when the stream ends; onError forwards non-EOF read/transport errors. close() tears down the connection.

Throws: The returned Promise rejects if the dial fails (refused, unreachable, timeout). After connect, write rejects on a write error and throws synchronously if called after close; transport read errors surface via the onError callback, not as rejections.

```
const sock = await net.tcp.connect("example.com", 80);
sock.onData(ev => runtime.log(ev.text));
await sock.write("GET / HTTP/1.0\r\n\r\n");
```

net.udp.open

```
open(opts: { host?: string, port?: string | number, bind?: string, readBuffer?: number }): void
```

Open a UDP socket: `net.udp.open(opts) → Promise`. Connected mode `{ host, port }` exposes `send(data)`; bound mode `{ bind: '9999' }` exposes `sendTo(data, host, port)` and tags inbound events with `{ address, port }`. Push/callback model — `onMessage(cb)/onClose(cb)/onError(cb)`; `handle.local` is the bound address; `handle.close()` shuts down. `opts` also takes `readBuffer?`.

Parameters

- `opts` (`{ host?: string, port?: string | number, bind?: string, readBuffer?: number }`) — Selects the mode: connected mode needs `{ host, port }` (`net.DialUDP` to that peer); bound mode needs `{ bind }` (e.g. `'9999'`, `net.ListenUDP` on that local address). `readBuffer` is the inbound channel capacity (default 64). Provide exactly one of the two modes.

Returns: `Promise` — connected mode resolves to `{ local: string, send(data: string | Uint8Array): Promise, onMessage, onClose, onError, close(): void }`; bound mode resolves to `{ local: string, sendTo(data: string | Uint8Array, host: string, port: string | number): Promise, send (throws), onMessage, onClose, onError, close(): void }`. `onMessage` fires per datagram with `{ bytes: Uint8Array, text: string }` plus `{ address, port }` in bound mode. `local` is the bound address.

Throws: Rejects if neither `{ bind }` nor `{ host, port }` is supplied, or the dial / listen fails. `send/sendTo` reject on a write error and throw synchronously after close; calling `send` on a bound socket throws (use `sendTo`). Read errors surface via `onError` (a clean close ends silently).

```
const u = await net.udp.open({ host: "1.1.1.1", port: 53 });
u.onMessage(ev => runtime.log(ev.bytes.length));
await u.send(query);
```

runtime

Script-host scaffolding: logging, assertions, time, environment, `runtime.argv`.

runtime.argv

```
argv: string[]
```

Per-script argument vector: `[programName, scriptPath, ...userArgs]`. `argv[0]` is the program name (sercon), `argv[1]` is the running script path, and any args after `--` on the command line start at `argv[2]`.

Returns: `string[]` — the per-run argument vector. `argv[0]` is the program name ("sercon"), `argv[1]` is the running script's path, and entries from index 2 onward are the user arguments passed after `--` on the command line. This is a value (property), not a function.

Throws: Not callable — accessing it never throws; reading an out-of-range index yields undefined per normal array semantics.

```
const target = runtime.argv[2] ?? "default-host";
```

runtime.assert.equal

```
equal(actual: unknown, expected: unknown, msg?: string): void
```

Throw when actual != expected (strict equality on primitives, deep equality on objects). Optional msg appears in the error.

Parameters

- `actual` (*unknown*) — The value produced by the code under test.
- `expected` (*unknown*) — The value to compare against. Primitives use strict equality; objects/arrays use deep structural equality (key order ignored).
- `msg` (*string, optional*) — Optional message prefixed onto the thrown error; defaults to "assert.equal failed".

Returns: void — returns nothing when the values match.

Throws: Throws an Error (": expected , got ") when the values are not equal.

```
runtime.assert.equal(1 + 1, 2, "math still works");
```

runtime.assert.ok

```
ok(cond: unknown, msg?: string): void
```

Throw when cond is falsy. Optional msg appears in the error.

Parameters

- `cond` (*unknown*) — A value tested for truthiness (JS coercion: 0, "", null, undefined, NaN and false are falsy).
- `msg` (*string, optional*) — Optional message used as the thrown error text; defaults to "assert.ok failed".

Returns: void — returns nothing when cond is truthy.

Throws: Throws an Error carrying msg (or "assert.ok failed") when cond is null or coerces to false.

```
runtime.assert.ok(user.id, "user must have an id");
```

runtime.env.get

```
get(name: string): string | undefined
```

Read an environment variable. Returns undefined when unset (not empty string).

Parameters

- `name` (*string*) — Environment variable name to look up.

Returns: string — the variable's value, or undefined when the variable is not set. A variable set to the empty string returns "", which is distinct from undefined.

Throws: Never throws.

```
const home = runtime.env.get("HOME") ?? "/tmp";
```

runtime.log

```
log(args: unknown[]): void
```

Print one space-separated line of the arguments to stdout. Primitives print raw; objects/arrays render as JSON (circular refs fall back to [object Object]). The script-side equivalent of console.log.

Parameters

- `args` (*unknown[]*) — Zero or more values to print. They are joined with single spaces; primitives stringify directly, objects/arrays are JSON-encoded.

Returns: void — writes a single newline-terminated line to stdout.

Throws: Never throws; unserialisable objects degrade to [object Object] rather than erroring.

```
runtime.log("count", 3, { ok: true }); // count 3 {"ok":true}
```

runtime.secrets.available

```
available: boolean
```

True when an OS keystore backend (macOS Keychain, Linux Secret Service, Windows Credential Manager) is plausibly reachable this run. Cheap advisory hint — does not touch the keystore; gate calls on it to self-skip on headless boxes.

Returns: boolean — false on a host with no reachable keystore (e.g. a headless Linux box without a D-Bus session). The authoritative signal is whether get/set/delete throw.

Throws: Never throws.

```
if (!runtime.secrets.available) runtime.log("no keystore — skipping");
```

runtime.secrets.delete

```
delete(name: string, account: string): Promise<boolean>
```

Remove a secret from the OS keystore under prefix + name / account. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the sercon prefix namespace (keystore service is prefix+name).
- `account` (*string*) — Account/user the secret belongs to.

Returns: Promise resolving true when an item was removed, false when there was nothing to remove.

Throws: Rejects if name is missing/empty or account is missing (pass "" for a single-secret name), when the keystore backend is unreachable, or the delete fails ("runtime.secrets.delete: ..."). Bounded by a 10s timeout.

```
const removed = await runtime.secrets.delete("devshop", "tess@example.com");
```

runtime.secrets.get

```
get(name: string, account: string): Promise<string | null>
```

Read a string secret from the OS keystore. The keystore service is the configured prefix + name (default "sercon/"), so reads are confined to sercon's namespace. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the sercon prefix namespace (the keystore service is prefix+name, e.g. "sercon/devshop").
- `account` (*string*) — Account/user the secret belongs to (may be an empty string for a single-secret name).

Returns: Promise resolving to the stored secret string, or null when no such item exists.

Throws: Rejects if name is missing/empty or account is missing (pass "" for a single-secret name), when the keystore backend is unreachable, or the read fails (a clean "runtime.secrets.get: ..." error). Bounded by a 10s timeout (a blocking macOS consent prompt rejects rather than hangs).

```
const pw = await runtime.secrets.get("devshop", "tess@example.com");
```

runtime.secrets.set

```
set(name: string, account: string, secret: string): Promise<void>
```

Store or overwrite a string secret in the OS keystore under prefix + name / account. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the sercon prefix namespace (keystore service is prefix+name).
- `account` (*string*) — Account/user the secret belongs to.
- `secret` (*string*) — The secret value to store.

Returns: Promise resolving when the secret is written.

Throws: Rejects if name is missing/empty, account is missing (pass "" for a single-secret name), or secret is missing; when the keystore backend is unreachable; or the write fails ("runtime.secrets.set: ..."). Bounded by a 10s timeout.

```
await runtime.secrets.set("devshop", "tess@example.com", "hunter2");
```

runtime.time.format

```
format(ms: number, layout: string, tz?: string): string
```

Format a unix-ms timestamp through strftime tokens. Optional IANA tz (e.g. 'Europe/Stockholm'); default is the host's local zone.

Parameters

- `ms` (*number*) — Milliseconds since the Unix epoch (e.g. from `time.nowMs`). Coerced to an integer.
- `layout` (*string*) — strftime-style layout. Supported tokens: %Y %y %m %d %H %M %S %T %F %j %A %a %B %b %z %Z and %% (literal percent). Unknown %X tokens pass through verbatim.
- `tz` (*string, optional*) — IANA timezone name (e.g. "Europe/Stockholm", "UTC"). Defaults to the host's local zone when omitted/null/undefined.

Returns: string — the timestamp rendered in tz with the given layout.

Throws: Throws ("time.format: ...") if tz is not a loadable IANA timezone name.

```
const s = runtime.time.format(runtime.time.nowMs(), "%F %T", "UTC");
```

runtime.time.nowMs

```
nowMs(): number
```

Wall-clock milliseconds since the Unix epoch.

Returns: number — integer milliseconds since 1970-01-01T00:00:00Z (host wall clock).

Throws: Never throws.

```
const t0 = runtime.time.nowMs();
```

runtime.time.sleep

```
sleep(ms: number): Promise<void>
```

Resolve after `ms` milliseconds. Cancellable via the engine timeout.

Parameters

- `ms` (*number*) — Delay in milliseconds. Coerced to an integer; non-positive values resolve effectively immediately.

Returns: Promise — resolves once the delay elapses.

Throws: Rejects if the run is cancelled or hits its timeout before the delay elapses (the underlying context is cancelled).

```
await runtime.time.sleep(250);
```

server

Network servers: HTTP/HTTPS listeners with routing, middleware, static files, WebSocket upgrade.

server.http.listen

```
listen(opts: { port: number; host?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; onError?: (err: unknown, req: Request, res: Response) => unknown }): { address: string; stopped: Promise<void>; close(): Promise<void> }
```

Bind an HTTP listener: `server.http.listen({port, host?, routes, use?, onError?})` → handle with `.address`, `.close()`, `.stopped` Promise. `routes` is a map of stdlib `http.ServeMux` patterns ('GET /users/{id}') to handlers `(req, res) => res.json({...})` or `{use: [...], handler: fn}` for per-route middleware. Optional `onError(err, req, res)` renders a custom response when a handler/middleware throws or rejects (else a stock 500). Handlers can call `res.upgradeWebSocket(opts?)` to hijack the connection and return an AsyncIterable with `.send / .close` — for `await (const msg of ws)` walks frames; `msg` is `{type:'text',text}` or `{type:'binary',bytes:Uint8Array}`. Handlers can also call `res.sse(opts?)` to start a one-way Server-Sent Events stream — returns a handle with `send(data)` (`string` → `data: ;` or `{event,data,id,retry}` with object data JSON-encoded), `close()`, and a `closed` Promise (resolves on close or client disconnect); `opts: {keepAlive?: ms, retry?: ms}`.

Parameters

- `opts` (`{ port: number; host?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise) => unknown)[]; handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res: Response, next: () => Promise) => unknown)[]; onError?: (err: unknown, req: Request, res: Response) => unknown }`) — Listener config. `port` is required; `host` defaults to "0.0.0.0". `routes` maps Go 1.22+ `ServeMux` patterns ('GET /', 'POST /users/{id}', 'GET /assets/{rest...}') to a handler function or a `{use, handler}` object for per-route middleware. `use` is a global middleware chain run before every route. `onError(err, req, res)` is invoked when a handler or middleware throws/rejects — render a response with `res.*` (it may be async); if it doesn't finalize, or itself throws, the stock 500 is sent. Under `sercon serve`, `--port-override` replaces `port`.

Returns: A server handle (returned synchronously): `address` is 'tcp/host:port' (resolved, so a `port:0` ephemeral bind reports its OS-chosen port); `stopped` resolves when the server stops (rejects if `Serve`

fails with a non-close error); `close()` begins a graceful 30s shutdown and resolves with the same stopped Promise.

Throws: Throws synchronously if `opts` is missing, `port` is 0/absent, `routes` is missing, a `use[]` entry or route value is not a function/valid `{use, handler}`, or the bind fails (e.g. address already in use).

```
const srv = server.http.listen({
  port: 8080,
  routes: { "GET /": (req, res) => res.json({ ok: true }) },
});
runtime.log(srv.address);
await srv.close();
```

server.http.static

```
static(opts: { dir: string; stripPrefix?: string; index?: string; etag?: boolean
}): (req: Request, res: Response) => void
```

Static-file mount: `server.http.static({dir, stripPrefix, index?, etag?})` → handler. Assign to a wildcard route (GET /assets/{rest...}). Internally stdlib `http.FileServer` with `stripPrefix`; ETag/Last-Modified/range requests work; no directory listing.

Parameters

- `opts` (*{ dir: string; stripPrefix?: string; index?: string; etag?: boolean }*) — `dir` is the filesystem root to serve. `stripPrefix` is removed from the request path before lookup (set it to the route's static prefix). `index` and `etag` are accepted but currently unused — `http.FileServer` already serves `index.html` and emits ETag/Last-Modified by default.

Returns: A route handler marker (returned synchronously). Assign it as a routes entry, typically under a wildcard pattern like 'GET /assets/{rest...}'. The route compiler unwraps it to a stdlib `http.FileServer` mounted under `http.StripPrefix`.

Throws: The call itself does not throw; an invalid `dir` surfaces as 404s at request time.

```
server.http.listen({
  port: 8080,
  routes: { "GET /assets/{rest...}": server.http.static({ dir: "./public",
stripPrefix: "/assets/" }) },
});
```

server.https.listen

```
listen(opts: { port: number; host?: string; cert: string | "self-signed"; key?:
string; routes: Record<string, ((req: Request, res: Response) => unknown) | {
use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[];
handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res:
Response, next: () => Promise<void>) => unknown)[] }): { address: string;
stopped: Promise<void>; close(): Promise<void> }
```

Like `server.http.listen` plus a `cert` (file path OR inline PEM string) and matching `key`. Set `cert: "self-signed"` to mint an ephemeral in-process dev cert (no key needed) instead — covers localhost / 127.0.0.1 / ::1 plus the listen host. No autocert (own production certs in your supervisor).

Parameters

- `opts` *{{ port: number; host?: string; cert: string | "self-signed"; key?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise) => unknown)[]; handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res: Response, next: () => Promise) => unknown)[] }}* — Same shape as `server.http.listen` plus `cert` and `key`. `cert` is a filesystem path, an inline PEM string (detected by a leading '-----BEGIN'), or the literal "self-signed" to generate an ephemeral P-256 cert in-process (key is then ignored). `key` is required for the path/PEM forms. TLS is pinned to a minimum of TLS 1.2. Self-signed certs fail normal client verification by design — dev only.

Returns: Same handle shape as `server.http.listen`; address is 'tcp/host:port'.

Throws: Throws synchronously on the same conditions as `server.http.listen`, plus if `cert` is missing, `key` is missing for a path/PEM cert, or the key pair fails to load/parse.

```
// Ephemeral dev cert — no openssl, no files:
const srv = server.https.listen({
  port: 8443,
  cert: "self-signed",
  routes: { "GET /": (req, res) => res.text("secure") },
});
```

server.https.static

```
static(opts: { dir: string; stripPrefix?: string; index?: string; etag?: boolean
}): (req: Request, res: Response) => void
```

Like `server.http.static`; same options.

Parameters

- `opts` *{{ dir: string; stripPrefix?: string; index?: string; etag?: boolean }}* — Identical to `server.http.static` — `dir` is the root, `stripPrefix` is removed from the path before lookup, `index/etag` are accepted but unused.

Returns: A route handler marker (returned synchronously); assign it to a wildcard route on an https listener.

Throws: The call itself does not throw; an invalid `dir` surfaces as 404s at request time.

```
server.https.listen({
  port: 8443, cert, key,
  routes: { "GET /assets/{rest...}": server.https.static({ dir: "./public",
    stripPrefix: "/assets/" }) },
});
```

server.icmp.listen

```
listen(opts?: { network?: "ip4" | "ip6" }, handler: (msg: { bytes: Uint8Array; text: string; address: string; type: number; code: number }, reply: (opts?: { to?: string; type?: number; code?: number; id?: number; seq?: number; payload?: string | Uint8Array; body?: string | Uint8Array }) => Promise<void>) => void): { address: string; close(): Promise<void> }
```

Bind a raw ICMP listener: `server.icmp.listen(opts?, (msg, reply) => {...})` → `handle { address: 'icmp/', close() }`. Raw ICMP has no ports — the socket receives ALL host ICMP traffic — and needs `root / CAP_NET_RAW` (synchronous bind throws otherwise). `opts { network?: 'ip4'|'ip6' (default 'ip4') }`. The handler runs once per received packet; `msg` is `{ bytes, text, address, type, code }` (the sender + parsed ICMP header) and `reply(opts?)` sends an ICMP message back to the sender (Echo by default, or a raw body), returning a Promise. Emits a `READY` line under `sercon serve` and joins graceful shutdown.

Parameters

- `opts` (*{ network?: "ip4" | "ip6" }, optional*) — network selects the IP version (default 'ip4'). There is no host/port — raw ICMP binds to all addresses.
- `handler` (*((msg: { bytes: Uint8Array; text: string; address: string; type: number; code: number }, reply: (opts?: { to?: string; type?: number; code?: number; id?: number; seq?: number; payload?: string | Uint8Array; body?: string | Uint8Array }) => Promise) => void)*) — Invoked once per received ICMP packet. `msg` carries the marshalled body (bytes/text), the sender address, and the parsed type/code. `reply(opts?)` sends an ICMP message back to the sender (or `opts.to`): Echo mode `{ type?, code?, id?, seq?, payload? }` or raw mode `{ type, code?, body }` (body marshalled verbatim); it returns a Promise that resolves once written.

Returns: A server handle (returned synchronously): `address` is `'icmp/'`; `close()` closes the socket and resolves. There is no per-connection handle (ICMP is connectionless) — `reply` is bound to the received packet's sender.

Throws: Throws synchronously if the handler is not a function, or if the raw socket can't be opened (typically because it needs `root / CAP_NET_RAW`). `reply` rejects on `resolve/marshal/write` errors and throws synchronously for the raw-body validation rules (a raw body requires `type`; `body` is mutually exclusive with `id/seq/payload`).

```
// Needs root / CAP_NET_RAW. Reply to every echo request with an echo reply:
const srv = server.icmp.listen({}, (msg, reply) => {
  if (msg.type === 8) reply({ type: 0, payload: msg.bytes });
});
runtime.log(srv.address);
await srv.close();
```

server.smtp.listen

```
listen(opts: { port: number; host?: string; hostname?: string; handlers: {
  onMail: (env: Envelope) => boolean | string | void | Promise<boolean | string |
  void>; onRcpt: (env: Envelope, to: string) => boolean | string | void |
  Promise<boolean | string | void>; onData: (env: Envelope, msg: Message) =>
  boolean | string | void | Promise<boolean | string | void> }; auth?: { user:
  string, pass: string, env: Envelope) => boolean | Promise<boolean>; starttls?: {
  cert: string; key: string }; allowInsecureAuth?: boolean; maxMessageBytes?:
  number; maxRecipients?: number; sessionTimeout?: number }): { address: string;
  stopped: Promise<void>; close(): Promise<void> }
```

Bind an SMTP listener: `server.smtp.listen({port, hostname?, handlers: {onMail, onRcpt, onData}, auth?, starttls?, allowInsecureAuth?, maxMessageBytes?, maxRecipients?, sessionTimeout?})` → handle with `.address`, `.close()`, `.stopped` Promise. Handlers receive (envelope, ...) per stage; return true/undefined to accept, false to reject, a string for a 550 reason, throw for 451 temp-fail. `onData` receives a parsed Message with text/html bodies, attachments, and raw bytes.

Parameters

- opts** (*{ port: number; host?: string; hostname?: string; handlers: { onMail: (env: Envelope) => boolean | string | void | Promise<boolean | string | void>; onRcpt: (env: Envelope, to: string) => boolean | string | void | Promise<boolean | string | void>; onData: (env: Envelope, msg: Message) => boolean | string | void | Promise<boolean | string | void> }; auth?: { user: string, pass: string, env: Envelope) => boolean | Promise; starttls?: { cert: string; key: string }; allowInsecureAuth?: boolean; maxMessageBytes?: number; maxRecipients?: number; sessionTimeout?: number }*) — port is required; host (bind interface) defaults to "0.0.0.0"; hostname (advertised EHLO domain) defaults to the OS hostname. handlers.onMail/onRcpt/onData are all required and run per protocol stage — each returns true/undefined to accept, false for a 550 reject, a string for '550 ', or throws for a 451 temp-fail. auth (optional) enables PLAIN+LOGIN SASL; return truthy to accept. starttls {cert, key} (paths or inline PEM) enables STARTTLS. allowInsecureAuth permits AUTH without TLS. maxMessageBytes defaults to 10 MiB (non-positive values ignored), maxRecipients to 100, sessionTimeout to 30000 (milliseconds).

Returns: A server handle (returned synchronously): address is 'tcp/host:port'; stopped resolves when the server stops (rejects on a non-close Serve error); close() shuts the listener down and resolves with the stopped Promise. The Envelope passed to handlers is { from, recipients, remote, helo, authenticatedUser?, tls?: { version, cipher } }; the Message passed to onData is { from, to, cc, subject, headers, body: { text, html }, attachments: { filename, contentType, bytes }[], raw: Uint8Array }.

Throws: Throws synchronously if opts is missing, port is 0/absent, handlers is missing, any of onMail/onRcpt/onData is absent or not a function, auth is present but not a function, the starttls key pair fails to load, or the bind fails.

```
const srv = server.smtp.listen({
  port: 2525,
  handlers: {
    onMail: (env) => true,
    onRcpt: (env, to) => to.endsWith("@example.com"),
    onData: (env, msg) => { runtime.log(msg.subject); },
  },
});
```

server.tcp.listen

```
listen(opts: { port: number; host?: string; readBuffer?: number }, handler:
(conn: { remote: string; local: string; write(data: string | Uint8Array):
Promise<void>; onData(cb: (msg: { bytes: Uint8Array; text: string }) => void):
void; onClose(cb: () => void): void; onError(cb: (err: unknown) => void): void;
close(): void }) => void): { address: string; close(): Promise<void> }
```

Bind a raw TCP server: `server.tcp.listen({port, host?, readBuffer?}, conn => {...})` → handle { address: 'tcp/host:port', close() }. The connection handler runs once per accepted socket; `conn` is the SAME handle shape as `net.tcp.connect` — `onData(cb)` (`cb` gets {bytes, text}), `onClose(cb)`, `onError(cb)`, `write(data)` (string or `Uint8Array`), `close()`, and `remote/local` addresses. Synchronous bind (throws on bind error); `port:0` binds an OS-chosen ephemeral port. Emits a `READY` line under `sercon` `serve` and joins graceful shutdown.

Parameters

- `opts` ({ *port*: number; *host?*: string; *readBuffer?*: number }) — `port` is the listen port (0 binds an OS-chosen ephemeral port). `host` defaults to all interfaces. `readBuffer` is the per-connection inbound channel capacity (frames buffered before backpressure), default 64.
- `handler` ((*conn*: { *remote*: string; *local*: string; *write*(*data*: string | Uint8Array): Promise; *onData*(*cb*: (*msg*: { *bytes*: Uint8Array; *text*: string }) => void): void; *onClose*(*cb*: () => void): void; *onError*(*cb*: (*err*: unknown) => void): void; *close*(): void }) => void) — Invoked once per accepted connection. `conn` matches `net.tcp.connect`'s `handle`: register `onData/onClose/onError` callbacks, `write()` to send (returns a Promise), `close()` to tear down. `remote/local` are 'host:port' strings.

Returns: A server handle (returned synchronously): `address` is 'tcp/host:port' (the resolved bind address, so `port:0` reports its ephemeral port). `close()` closes the listener and all accepted connections, then resolves.

Throws: Throws synchronously if `opts` is missing, the handler is not a function, or the bind fails (e.g. address already in use).

```
const srv = server.tcp.listen({ port: 0 }, (conn) => {
  conn.onData(msg => conn.write("echo: " + msg.text));
});
runtime.log(srv.address);
await srv.close();
```

server.udp.listen

```
listen(opts: { port: number; host?: string }, handler: (msg: { bytes:
Uint8Array; text: string; address: string; port: number }, reply: (data: string
| Uint8Array) => Promise<void>) => void): { address: string; close():
Promise<void> }
```

Bind a raw UDP server: `server.udp.listen({port, host?}, (msg, reply) => {...})` → handle { address: 'udp/host:port', close() }. The handler runs once per inbound datagram; msg is {bytes, text, address, port} (the sender) and reply(data) (string or Uint8Array) sends a datagram back to that sender, returning a Promise. Synchronous bind (throws on bind error); port:0 binds an OS-chosen ephemeral port. Emits a READY line under `sercon` `serve` and joins graceful shutdown.

Parameters

- `opts` ({ *port: number; host?: string* }) — port is the listen port (0 binds an OS-chosen ephemeral port). host defaults to all interfaces.
- `handler` ((*msg: { bytes: Uint8Array; text: string; address: string; port: number }, reply: (data: string | Uint8Array) => Promise) => void*) — Invoked once per inbound datagram. msg carries the payload (bytes/text) and the sender's address/port. reply(data) sends a datagram back to that sender and returns a Promise that resolves once written (rejects on a write error).

Returns: A server handle (returned synchronously): address is 'udp/host:port' (resolved, so port:0 reports its ephemeral port). close() closes the socket and resolves. There is no per-connection handle — UDP is connectionless, so reply is bound to the originating datagram's sender.

Throws: Throws synchronously if opts is missing, the handler is not a function, the address fails to resolve, or the bind fails.

```
const srv = server.udp.listen({ port: 0 }, (msg, reply) => {
  reply("pong: " + msg.text);
});
runtime.log(srv.address);
await srv.close();
```

services

Subprocess and external-CLI / service wrappers: shell, git, gh, AI providers, agent-browser automation, W3C WebDriver browser control.

services.agentBrowser.auth

Namespace object for the auth vault (session-independent): auth.save, auth.list, auth.show, auth.delete. Passwords are never placed in argv — auth.save sends the password via stdin (--password-stdin).

services.agentBrowser.auth.delete

```
delete(...args: unknown[])
```

services.agentBrowser.auth.list

```
list(...args: unknown[])
```

services.agentBrowser.auth.save

```
save(...args: unknown[])
```

services.agentBrowser.auth.show

```
show(...args: unknown[])
```

services.agentBrowser.available

```
available
```

True when the agent-browser CLI is on PATH. Sync boolean, resolved once per Run. Gate calls on this; every binding throws a clean error when the CLI is absent.

Returns: boolean — true if `agent-browser` is on PATH.

```
if (!services.agentBrowser.available) runtime.log("install agent-browser first");
```

services.agentBrowser.clearDefaultOptions

```
clearDefaultOptions(...args: unknown[])
```

Reset the namespace-level launch defaults to an empty object, removing any values set by `setDefaultOptions`.

Returns: void

Throws: Never throws.

```
services.agentBrowser.clearDefaultOptions();  
runtime.log(JSON.stringify(services.agentBrowser.defaultOptions())); // {}
```

services.agentBrowser.defaultOptions

```
defaultOptions(...args: unknown[])
```

Return a shallow copy of the current namespace-level launch defaults. These are merged (under per-call opts) into every subsequent `launch()`.

Returns: object — a plain-object copy of the current defaults map. Empty object when no defaults have been set.

Throws: Never throws.

```
services.agentBrowser.setDefaultOptions({ headed: false });
runtime.log(JSON.stringify(services.agentBrowser.defaultOptions())); //
{"headed":false}
```

services.agentBrowser.eval

```
eval(url: string, js: string): void
```

One-shot shortcut: launch an ephemeral session, open url, evaluate a JS expression in the page, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `js` (*string*) — JavaScript expression to evaluate in the page context. Required.

Returns: `Promise<{ success: boolean, data: object, error: string|null }>` — agent-browser envelope; `data.result` holds the serialised return value.

Throws: Throws if url or js is missing; if agent-browser is not on PATH; on navigation or evaluation failure.

```
const r = await services.agentBrowser.eval("data:text/html,<title>Hi</title>",
"document.title");
runtime.log(r.data?.result); // "Hi"
```

services.agentBrowser.launch

```
launch(opts?: { session?: string, headed?: boolean, profile?: string, proxy?:
string, userAgent?: string, device?: string, colorScheme?: string,
ignoreHttpsErrors?: boolean, engine?: string, executablePath?: string, enable?:
string, args?: string, timeout?: number }): { session: string; open(url: string,
opts?: object): Promise<any>; back(): Promise<any>; forward(): Promise<any>;
reload(): Promise<any>; wait(selOrMs: string | number): Promise<any>;
connect(target: string): Promise<any>; click(sel: string): Promise<any>;
dblclick(sel: string): Promise<any>; hover(sel: string): Promise<any>;
focus(sel: string): Promise<any>; check(sel: string): Promise<any>; uncheck(sel:
string): Promise<any>; scrollIntoView(sel: string): Promise<any>; fill(sel:
string, text: string): Promise<any>; type(sel: string, text: string):
Promise<any>; press(key: string): Promise<any>; select(sel: string, ...values:
string[]): Promise<any>; scroll(dir: string, px?: number): Promise<any>;
drag(src: string, dst: string): Promise<any>; upload(sel: string, files: string
| string[]): Promise<any>; download(sel: string, path: string): Promise<any>;
keyboard: { type(text: string): Promise<any>; insertText(text: string):
Promise<any> }; mouse: { move(x: number, y: number): Promise<any>; down(button?:
string): Promise<any>; up(button?: string): Promise<any>; wheel(dy: number, dx?:
number): Promise<any> }; get(what: string, sel?: string): Promise<any>;
isVisible(sel: string): Promise<any>; isEnabled(sel: string): Promise<any>;
isChecked(sel: string): Promise<any>; eval(code: string): Promise<any>;
snapshot(opts?: object): Promise<any>; console(opts?: object): Promise<any>;
errors(opts?: object): Promise<any>; highlight(sel: string): Promise<any>;
find(locator: string, value: string, opts: { action: string, text?: string }):
Promise<any>; locator(spec: object | string, value?: string): object; set: {
viewport(w: number, h: number, scale?: number): Promise<any>; device(name:
string): Promise<any>; geo(lat: number, lng: number): Promise<any>; offline(on?:
boolean): Promise<any>; headers(headers: Record<string, string>): Promise<any>;
credentials(user: string, pass: string): Promise<any>; media(scheme?: "dark" |
"light", reducedMotion?: boolean): Promise<any> }; record: { start(path: string,
url?: string): Promise<any>; stop(): Promise<any> }; screenshot(path?: string,
opts?: { selector?: string, full?: boolean, annotate?: boolean, format?: "png" |
"jpeg", quality?: number }): Promise<{ path?: string, size?: number, bytes?:
number[], format: string }>; pdf(path?: string): Promise<{ path?: string, size?:
number, bytes?: number[], format: string }>; network: { route(url: string,
opts?: { abort?: boolean, body?: unknown, resourceType?: string }):
Promise<any>; unroute(url?: string): Promise<any>; requests(opts?: { clear?:
boolean, filter?: string, type?: string, method?: string, status?: string }):
Promise<any>; request(id: string): Promise<any>; har: { start(path?: string):
Promise<any>; stop(path?: string): Promise<any> } }; cookies: { get():
Promise<any>; set(name: string, value: string, opts?: { url?: string, domain?:
string, path?: string, httpOnly?: boolean, secure?: boolean, sameSite?: "Strict"
| "Lax" | "None", expires?: number }): Promise<any>; clear(): Promise<any> };
storage: { local: { get(key?: string): Promise<any>; set(key: string, value:
string): Promise<any>; clear(): Promise<any> }; session: { get(key?: string):
Promise<any>; set(key: string, value: string): Promise<any>; clear():
Promise<any> } }; tabs: { list(): Promise<any>; new(url?: string, opts?: {
label?: string }): Promise<any>; close(ref?: string): Promise<any>; select(ref:
string): Promise<any> }; diff: { snapshot(opts?: { baseline?: string, selector?:
string, compact?: boolean, depth?: number }): Promise<any>; screenshot(opts: {
baseline: string, output?: string, threshold?: number }): Promise<any>;
url(url1: string, url2: string): Promise<any> }; trace: { start(): Promise<any>;
```

```
stop(path?: string): Promise<any> }; profiler: { start(opts?: { categories?:
string }): Promise<any>; stop(path?: string): Promise<any> }; inspect():
Promise<any>; clipboard(op: "read" | "write" | "copy" | "paste", text?: string):
Promise<any>; vitals(url?: string): Promise<any>; pushstate(url: string):
Promise<any>; react: { tree(): Promise<any>; inspect(id: string): Promise<any>;
renders: { start(): Promise<any>; stop(): Promise<any> }; suspense(opts?: {
onlyDynamic?: boolean }): Promise<any> }; stream: { enable(opts?: { port?:
number }): Promise<any>; disable(): Promise<any>; status(): Promise<any> };
chat(message: string, opts?: { model?: string }): Promise<any>; cmd(command:
string, ...args: string[]): Promise<any>; batch(cmds: string[], opts?: { bail?:
boolean }): Promise<any>; auth: { login(name: string): Promise<any> }; close():
Promise<any> }
```

Allocate a browser session and return a handle. Synchronous (no browser starts until the first command). Pass `opts.session` to name the session; otherwise a unique id is generated. Launch flags (headed, profile, proxy, userAgent, device, colorScheme, ignoreHttpsErrors, engine, executablePath, enable, args) are threaded into every call the handle makes. Sessions the script does not `close()` are best-effort closed when the Run ends.

Parameters

- `opts` (*{ session?: string, headed?: boolean, profile?: string, proxy?: string, userAgent?: string, device?: string, colorScheme?: string, ignoreHttpsErrors?: boolean, engine?: string, executablePath?: string, enable?: string, args?: string, timeout?: number }, optional*) — Launch flags captured for the lifetime of the handle and threaded into every subprocess call. `session` names the agent-browser session (auto-generated when omitted). `timeout` is the per-call agent-browser subprocess timeout in milliseconds (default 30000, 0 disables); when a call exceeds the limit the Promise rejects with a clear 'timed out' error instead of hanging forever.

Returns: A handle object with a read-only session string and methods: `open`, `back`, `forward`, `reload`, `wait`, `connect`, `click`, `dblclick`, `hover`, `focus`, `fill`, `type`, `press`, `check`, `uncheck`, `select`, `scroll`, `scrollIntoView`, `drag`, `upload`, `download`, `keyboard`.{type,insertText}, `mouse`.{move,down,up,wheel}, `get`, `isVisible`, `isEnabled`, `isChecked`, `eval`, `snapshot`, `console`, `errors`, `highlight`, `find`, `locator`, `close`. Every async method resolves to an agent-browser envelope { success: boolean, data: object, error: string|null }; drill into `.data` for the actual values.

Throws: `launch()` itself does not throw for a missing CLI (it allocates only); the first method call throws if agent-browser is not on PATH.

```
const b = services.agentBrowser.launch({ headed: false });
await b.open("https://example.com");
const r = await b.get("title");
runtime.log(r.data?.title);
await b.close();
```

services.agentBrowser.pdf

```
pdf(url: string, path?: string): Promise<{ path: string, size: number, format:
string } | { bytes: number[], format: string }>
```

One-shot shortcut: launch an ephemeral session, open url, capture a PDF, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `path` (*string, optional*) — Output file path. When supplied the PDF is written there and the result has { `path`, `size`, `format` }; when omitted the PDF bytes are returned as a `number[]` in { `bytes`, `format` }.

Returns: Promise — path given: { `path`: string, `size`: number, `format`: string }; no path: { `bytes`: number[], `format`: string }.

Throws: Throws if url is missing; if agent-browser is not on PATH; on navigation, capture, or I/O failure.

```
const pdf = await services.agentBrowser.pdf("data:text/html,<h1>Hi</h1>");
runtime.log(new Uint8Array(pdf.bytes).length, pdf.format); // e.g. 5678 pdf
```

services.agentBrowser.screenshot

```
screenshot(url: string, path?: string, opts?: { selector?: string, full?:
boolean, annotate?: boolean, format?: "png" | "jpeg", quality?: number }):
Promise<{ path: string, size: number, format: string } | { bytes: number[],
format: string }>
```

One-shot shortcut: launch an ephemeral session, open url, capture a screenshot, and close.

Equivalent to `launch()+open(url)+screenshot(path?,opts?)+close()` but in a single call.

Parameters

- `url` (*string*) — URL to open (http/https or data: URI). Required.
- `path` (*string, optional*) — Output file path. When supplied the screenshot is written there and the result has { `path`, `size`, `format` }; when omitted the image bytes are returned as a `number[]` (byte-value array) in { `bytes`, `format` }.
- `opts` (*{ selector?: string, full?: boolean, annotate?: boolean, format?: "png" | "jpeg", quality?: number }, optional*) — Capture options. `selector` scopes the capture to an element. `full` captures the full page. `format` defaults to png. `quality` (0–100) applies to jpeg.

Returns: Promise — path given: { `path`: string, `size`: number, `format`: string }; no path: { `bytes`: number[], `format`: string } where bytes is a plain JS `number[]` (byte-value array); wrap with `new Uint8Array(bytes)` to get a typed array.

Throws: Throws if url is missing; if agent-browser is not on PATH; on navigation, capture, or I/O failure.

```
const shot = await services.agentBrowser.screenshot("data:text/html,
<h1>Hi</h1>");
runtime.log(new Uint8Array(shot.bytes).length, shot.format); // e.g. 12345 png
```

services.agentBrowser.setDefaultOptions

```
setDefaultOptions(opts: object): void
```

Replace the namespace-level launch defaults with the supplied object. Merged (under per-call opts) into every subsequent launch(). Affects only the current Run.

Parameters

- `opts` (*object*) — A plain object of launch option key/value pairs. The entire defaults map is replaced (not merged) with this object.

Returns: void

Throws: Never throws.

```
services.agentBrowser.setDefaultOptions({ headed: false, proxy:
"http://proxy:3128" });
const b = services.agentBrowser.launch(); // headed:false + proxy inherited
```

services.agentBrowser.snapshot

```
snapshot(url: string, opts?: { interactive?: boolean, compact?: boolean, depth?:
number, selector?: string }): void
```

One-shot shortcut: launch an ephemeral session, open url, take an accessibility-tree snapshot, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `opts` (*{ interactive?: boolean, compact?: boolean, depth?: number, selector?: string }, optional*) — Snapshot options forwarded to the handle's snapshot() method.

Returns: Promise<{ success: boolean, data: object, error: string|null }> — agent-browser envelope; data contains the accessibility tree.

Throws: Throws if url is missing; if agent-browser is not on PATH; on navigation or snapshot failure.

```
const snap = await services.agentBrowser.snapshot("data:text/html,<h1>Hi</h1>",
{ compact: true });
runtime.log(JSON.stringify(snap.data).slice(0, 100));
```

services.agentBrowser.version

```
version(...args: unknown[]): Promise<string>
```

The agent-browser CLI version string.

Returns: Promise — the version reported by `agent-browser --version`.

Throws: Throws if the agent-browser CLI is not on PATH.

```
runtime.log(await services.agentBrowser.version());
```

services.ai.providers

```
providers(): string[]
```

Which of claude / codex / copilot / gemini are on PATH, in preference order.

Returns: string[] — the subset of supported AI CLIs found on PATH, in preference order (claude, codex, copilot, gemini); an empty array when none are installed. Synchronous (not a Promise).

Throws: Never throws.

```
const ps = services.ai.providers(); // e.g. ["claude", "gemini"]
```

services.ai.send

```
send(opts: { prompt: string, provider?: "claude" | "codex" | "copilot" |  
"gemini", system?: string, context?: string, timeout?: number }):  
Promise<Record<string, unknown>>
```

Run a one-shot prompt through a provider. opts { prompt (required), provider?, system?, context?, timeout? }. Returns { provider, output, exitCode }. Non-zero exit is data; no provider throws.

Parameters

- `opts` (*{ prompt: string, provider?: "claude" | "codex" | "copilot" | "gemini", system?: string, context?: string, timeout?: number }*) — prompt is required. provider names the CLI to use; when omitted, the first provider on PATH (in preference order) is chosen. system and context are prepended to the prompt as "System: ..." / "Context: ..." blocks (a portable substitute for each CLI's own flags). timeout in ms (default 120000).

Returns: Promise<{ provider: string, output: string, exitCode: number }> — provider is the CLI that ran; output is its trimmed stdout (or stderr when stdout is empty on a non-zero exit); exitCode is 0 on success.

Throws: Throws if opts.prompt is missing/empty, no provider is on PATH (when provider is unset), the named provider is unknown, the CLI fails to spawn, or on timeout / context cancellation. A non-zero exit code does NOT throw — it resolves with that exitCode.

```
const r = await services.ai.send({ prompt: "Say hi", provider: "claude" });  
runtime.log(r.output);
```

services.exec.http

```
http(method: string, url: string, opts?: { headers?: Record<string, string>,  
body?: string, timeout?: number, follow?: boolean, insecure?: boolean, backend?:  
"auto" | "recon" | "curl" }): Promise<Record<string, unknown>>
```

Make an HTTP request by shelling out to recon (preferred) or curl (fallback). 4xx/5xx resolve as status; transport errors and timeouts throw. opts.backend = 'auto' | 'recon' | 'curl'.

Parameters

- `method` (*string*) — HTTP verb (GET, POST, PUT, DELETE, PATCH, HEAD); lower-case input is uppercased before forwarding.
- `url` (*string*) — Target URL; must be fully qualified (the backend requires a scheme + host).
- `opts` (*{ headers?: Record<string, string>, body?: string, timeout?: number, follow?: boolean, insecure?: boolean, backend?: "auto" | "recon" | "curl" }, optional*) — headers emits one -H "Name: Value" per entry. body is written to a temp file and sent via --data-binary so CR/LF stay intact. timeout in ms (default 30000). follow toggles -L to follow 3xx redirects. insecure toggles -k to skip TLS verification. backend picks the tool: 'auto' (default) prefers recon then curl; 'recon' or 'curl' require that specific binary on PATH.

Returns: Promise<{ status: number, headers: Record<string, string>, body: string, durationMs: number, backend: "recon" | "curl" }> — status is the final HTTP status code; headers have lower-cased names (last response block on a redirect chain); body is the UTF-8 decoded response body; backend is whichever tool ran.

Throws: Throws if method or url is missing/empty; if the requested backend (or, for 'auto', neither recon nor curl) is on PATH; on transport errors (DNS failure, connection refused, TLS handshake); on timeout or context cancellation; or if the response headers can't be parsed. HTTP 4xx/5xx do NOT throw — they resolve with that status.

```
const r = await services.exec.http("GET", "https://example.com");
runtime.log(r.status, r.backend);
```

services.exec.shell

```
shell(cmd: string | string[], opts?: { timeout?: number, cwd?: string, stdin?:
string, env?: Record<string, string>, pane?: string | Pane, pty?: boolean }):
Promise<Record<string, unknown>>
```

Run a subprocess and wait for it to exit. String cmd → /bin/sh -c (or cmd /C on Windows); array cmd → argv (no shell). Non-zero exits resolve normally; spawn failures and timeouts throw.

Parameters

- `cmd` (*string | string[]*) — A string is passed verbatim to the host shell (/bin/sh -c on Unix, cmd /C on Windows) so quoting, pipes, and redirects work. A string[] is treated as argv: argv[0] is run directly with no shell, so use this form when arguments contain whitespace or shell metacharacters you don't want re-interpreted.
- `opts` (*{ timeout?: number, cwd?: string, stdin?: string, env?: Record<string, string>, pane?: string | Pane, pty?: boolean }, optional*) — timeout in ms (default 30000); on expiry the process tree is killed and the call throws. cwd sets the working directory. stdin is fed to the process's standard input. env entries are merged on top of the inherited environment (they do not replace it). pane (a tui.pane name or Pane handle) streams stdout+stderr live into a TUI pane — in that mode the result's stdout/stderr strings stay empty. pty (default false) runs the command under a pseudo-terminal so it believes it is a terminal and emits color/progress; with a pane the output is rendered there, without a pane it is captured into stdout (stderr stays empty since a pty merges both streams). Unix only — on Windows pty is ignored and the normal pipe path is used.

Returns: `Promise<{ stdout: string, stderr: string, exitCode: number, success: boolean, durationMs: number }>` — `stdout/stderr` are captured (empty when streamed to a pane); `exitCode` is 0 on success; `success` is `exitCode === 0`; `durationMs` is wall-clock spawn-to-exit time.

Throws: Throws if `cmd` is missing, an empty string, an empty array, or a non-string array element; if the host binary is not on `PATH` or fails to start; if the timeout (or context cancellation) fires before exit; or if a named pane is referenced without a prior `tui.layout`. A non-zero exit code does NOT throw — it resolves with `success:false`.

```
const r = await services.exec.shell("echo hi");
if (r.success) runtime.log(r.stdout.trim());
```

services.exec.stream

```
stream(cmd: string | string[], onLine: (line: string, stream: "stdout" |
"stderr") => void, opts?: { cwd?: string, env?: Record<string, string>, stdin?:
string, timeout?: number }): Promise<{ exitCode: number; success: boolean;
durationMs: number }>
```

Run a subprocess and stream its `stdout/stderr` to a callback line by line as output arrives (unlike `exec.shell`, which buffers). String `cmd` → `/bin/sh -c` (or `cmd /C` on Windows); array `cmd` → `argv` (no shell). Resolves `{ exitCode, success, durationMs }` on exit; non-zero exits resolve normally; spawn failures and timeouts reject.

Parameters

- `cmd` (*string | string[]*) — A string is passed to the host shell (`/bin/sh -c` on Unix, `cmd /C` on Windows) so pipes, redirects, and globs work. A `string[]` is treated as `argv`: `argv[0]` is run directly with no shell.
- `onLine` (*(line: string, stream: "stdout" | "stderr") => void*) — Called once per output line as it arrives. `line` has its trailing newline stripped; `stream` is `'stdout'` or `'stderr'`. A final line without a trailing newline is still delivered. Required — a non-function throws synchronously.
- `opts` (*{ cwd?: string, env?: Record<string, string>, stdin?: string, timeout?: number }, optional*) — `cwd` sets the working directory; `env` entries merge on top of the inherited environment; `stdin` is fed to the process. `timeout` is in ms with NO default (0 / absent = run until exit, unlike `exec.shell`'s 30000); when set, the process tree is killed on expiry and the call rejects.

Returns: `Promise<{ exitCode: number, success: boolean, durationMs: number }>` — resolves on process exit. `exitCode` is 0 on success; `success` is `exitCode === 0`; `durationMs` is wall-clock spawn-to-exit time. Output is delivered via `onLine`, not captured into the result.

Throws: Throws synchronously if `cmd` is missing or `onLine` is not a function. The `Promise` rejects if the host binary is not on `PATH` or fails to start, or if the timeout (or context cancellation) fires before exit. A non-zero exit code does NOT reject — it resolves with `success:false`.

```
const r = await services.exec.stream("echo one; echo two", (line, stream) => {
  runtime.log(stream, line);
});
runtime.log("exit", r.exitCode);
```

services.gh.authStatus

```
authStatus(...args: unknown[]): Promise<{ authenticated: boolean; user: string; raw: string }>
```

Probe gh's auth state. Missing gh / unauthenticated resolve with { authenticated: false, ... } — only context cancellation throws.

Returns: Promise<{ authenticated: boolean, user: string, raw: string }> — authenticated is true only when gh api user succeeds; user is the resolved login ("" when not authenticated); raw is the login on success or the underlying gh error / "gh not on PATH" otherwise.

Throws: Throws only on context cancellation. A missing gh binary or an unauthenticated session resolve with authenticated:false rather than throwing.

```
const a = await services.gh.authStatus();
if (a.authenticated) runtime.log("logged in as", a.user);
```

services.gh.prList

```
prList(opts?: { cwd?: string, state?: string, limit?: number, author?: string }): Promise<Record<string, unknown>[]>
```

List pull requests on the cwd's repo (or opts.cwd). Defaults: open state, limit 30. Filters: state / limit / author.

Parameters

- `opts` (*{ cwd?: string, state?: string, limit?: number, author?: string }, optional*) — cwd selects the repo (defaults to the engine's working directory, which gh uses to detect the repo). state filters by PR state ("open" default, "closed", "merged", "all"). limit caps results (default 30; must be positive). author filters to PRs opened by that login.

Returns: Promise<Array<{ number: number, title: string, state: string, author: string, headRefName: string, baseRefName: string, url: string, createdAt: string, updatedAt: string }>> — one object per PR; author is flattened from gh's { login } wrapper to the bare login string; createdAt/updatedAt are ISO 8601 timestamps.

Throws: Throws if limit is <= 0, gh is not on PATH, gh exits non-zero (not authenticated, not a GitHub repo, etc.), the JSON can't be parsed, or on context cancellation.

```
const prs = await services.gh.prList({ state: "open", limit: 5 });
for (const pr of prs) runtime.log(pr.number, pr.title);
```

services.gh.repoView

```
repoView(repo?: string, opts?: { cwd?: string }): Promise<Record<string, unknown>>
```

Repo metadata. With no arg uses cwd's repo; pass 'owner/name' for any repo gh can see. owner + defaultBranch are pre-flattened.

Parameters

- `repo` (*string, optional*) — "owner/name" of any repo gh can access. Omit (or pass opts as the first arg) to view the repo detected from cwd.
- `opts` (*{ cwd?: string }, optional*) — cwd selects the checkout gh uses to detect the current repo when repo is omitted.

Returns: `Promise<{ name: string, owner: string, description: string, url: string, defaultBranch: string, visibility: string }>` — owner is flattened from gh's { login } wrapper to the bare login; defaultBranch is flattened from defaultBranchRef.name (" " if absent); key order matches gh's output.

Throws: Throws if gh is not on PATH, gh exits non-zero (repo not found, not authenticated, etc.), the JSON can't be parsed, or on context cancellation.

```
const r = await services.gh.repoView("cli/cli");
runtime.log(r.owner, r.defaultBranch);
```

services.git.add

```
add(paths: string | string[], opts?: { cwd?: string }): Promise<{ paths:
string[] }>
```

Stage one path (string) or several (string[]).

Parameters

- `paths` (*string | string[]*) — Path or paths to stage. Passed after a `--` separator so paths that look like flags (-foo) are handled literally.
- `opts` (*{ cwd?: string }, optional*) — cwd selects the checkout; defaults to the engine's working directory.

Returns: `Promise<{ paths: string[] }>` — the list of paths that were staged.

Throws: Throws if paths is missing, an empty string, or contains a non-string array element; if git is not on PATH; or if git add exits non-zero (e.g. a pathspec matching nothing).

```
await services.git.add(["src/a.ts", "src/b.ts"]);
```

services.git.branch

```
branch(opts?: { cwd?: string }): Promise<{ current: string; detached: boolean;
all: string[] }>
```

Current branch (empty when HEAD is detached) plus the list of local branches.

Parameters

- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout to inspect; defaults to the engine's working directory.

Returns: `Promise<{ current: string, detached: boolean, all: string[] }>` — `current` is the checked-out branch name (" " when detached); `detached` is true on a detached HEAD; `all` lists every local branch (refs/heads) by short name.

Throws: Throws if git is not on PATH, the directory is not a git repository, or the underlying git command fails. A detached HEAD is reported via `detached:true`, not a throw.

```
const b = await services.git.branch();
runtime.log(b.detached ? "(detached)" : b.current);
```

services.git.commit

```
commit(message: string, opts?: { cwd?: string, allowEmpty?: boolean }):
Promise<{ sha: string }>
```

Create a commit; returns the post-commit HEAD SHA. `opts.allowEmpty` toggles `--allow-empty`.

Parameters

- `message` (*string*) — Commit message (passed as a single `-m` argument).
- `opts` (*{ cwd?: string, allowEmpty?: boolean }, optional*) — `cwd` selects the checkout. `allowEmpty` adds `--allow-empty` so a commit succeeds with no staged changes (release markers, etc.); defaults to false.

Returns: `Promise<{ sha: string }>` — `sha` is the full SHA of the newly created HEAD commit.

Throws: Throws if `message` is missing or blank, git is not on PATH, or git commit exits non-zero (e.g. nothing staged and `allowEmpty` is false).

```
const c = await services.git.commit("chore: bump", { allowEmpty: true });
```

services.git.diffStat

```
diffStat(opts?: { cwd?: string, revRange?: string }): Promise<{ files: number;
insertions: number; deletions: number }>
```

Aggregate `{ files, insertions, deletions }` from `git diff --shortstat`. Default `revRange` `HEAD~1..HEAD`.

Parameters

- `opts` (*{ cwd?: string, revRange?: string }, optional*) — `cwd` selects the checkout. `revRange` is the diff range (default `"HEAD~1..HEAD"`, the last commit).

Returns: `Promise<{ files: number, insertions: number, deletions: number }>` — counters parsed from `git diff --shortstat`. An empty diff returns all zeros.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git diff exits non-zero (e.g. a bad revRange).

```
const d = await services.git.diffStat();
runtime.log(d.files, d.insertions, d.deletions);
```

services.git.isClean

```
isClean(opts?: { cwd?: string }): Promise<boolean>
```

True iff `git status --porcelain` is empty.

Parameters

- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: Promise — true when the working tree has no staged, unstaged, or untracked changes.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git status exits non-zero.

```
if (await services.git.isClean()) runtime.log("clean");
```

services.git.log

```
log(opts?: { cwd?: string, limit?: number, revRange?: string }): Promise<{ sha: string; shortSha: string; author: string; email: string; timestamp: number; subject: string }[]>
```

Recent commits as { sha, shortSha, author, email, timestamp, subject }. `opts.limit` / `opts.revRange`.

Parameters

- `opts` (*{ cwd?: string, limit?: number, revRange?: string }, optional*) — `cwd` selects the checkout. `limit` caps the number of commits (default 50; must be positive). `revRange` selects the range/ref to walk (default "HEAD").

Returns: Promise<Array<{ sha: string, shortSha: string, author: string, email: string, timestamp: number, subject: string }>> — newest first; timestamp is the author Unix epoch seconds; subject is the commit's first line.

Throws: Throws if limit is ≤ 0 , git is not on PATH, the directory is not a git repository, or git log exits non-zero (e.g. an unknown revRange).

```
const log = await services.git.log({ limit: 5 });
runtime.log(log[0].subject);
```

services.git.revParse

```
revParse(rev: string, opts?: { cwd?: string }): Promise<string>
```

Full 40-char SHA for the given rev. Invalid refs throw.

Parameters

- `rev` (*string*) — Any revision git understands (branch, tag, HEAD, short SHA, HEAD~2, etc.).
- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: Promise — the full 40-character commit SHA the rev resolves to.

Throws: Throws if rev is missing or empty, git is not on PATH, the directory is not a git repository, or the rev cannot be resolved (git's own error message is included).

```
const sha = await services.git.revParse("HEAD");
```

services.git.runText

```
runText(args: string | string[], opts?: { cwd?: string }): Promise<{ stdout: string; stderr: string; exitCode: number }>
```

Escape hatch: run any `git <args>`, get { stdout, stderr, exitCode } — exitCode is data, not a throw.

Parameters

- `args` (*string | string[]*) — git arguments (without the leading "git"), e.g. ["tag", "--list"]. A bare string is treated as a single argument.
- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: Promise<{ stdout: string, stderr: string, exitCode: number }> — captured streams plus git's exit code, so callers can react to any exit status without try/catch.

Throws: Throws if args is missing, an empty string, or an empty array; if git is not on PATH; or on a spawn failure / context cancellation. A non-zero exit code does NOT throw — it is returned in exitCode.

```
const r = await services.git.runText(["tag", "--list"]);
if (r.exitCode === 0) runtime.log(r.stdout);
```

services.git.status

```
status(opts?: { cwd?: string }): Promise<{ path: string; indexStatus: string;
workingStatus: string }[]>
```

Parsed `git status --porcelain` entries: { path, indexStatus, workingStatus }.

Parameters

- `opts` (`{ cwd?: string }`, *optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: `Promise<Array<{ path: string, indexStatus: string, workingStatus: string }>>` — one entry per changed path; `indexStatus` / `workingStatus` are the porcelain v1 X / Y status characters (e.g. "M", "A", "?"). An empty array means a clean tree.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git status exits non-zero.

```
for (const e of await services.git.status())
  runtime.log(e.indexStatus + e.workingStatus, e.path);
```

services.webdriver.available

```
available
```

True when a W3C WebDriver binary (chromedriver or geckodriver) is on PATH. Sync boolean, resolved once per Run. Gate calls on this before using probe or connect.

Returns: boolean — true if chromedriver or geckodriver is found on PATH.

```
if (!services.webdriver.available) {
  runtime.log("install chromedriver or geckodriver first");
}
```

services.webdriver.connect

```

connect(opts?: { browser?: "chrome" | "firefox", headless?: boolean, url?:
string, args?: string[], capabilities?: object, commandTimeout?: number }):
Promise<{ get(url: string): Promise<{ ok: true }>; url(): Promise<string>;
title(): Promise<string>; back(): Promise<{ ok: true }>; forward(): Promise<{
ok: true }>; refresh(): Promise<{ ok: true }>; find(by: string, value: string):
Promise<{ click(): Promise<{ ok: true }>; sendKeys(text: string): Promise<{ ok:
true }>; clear(): Promise<{ ok: true }>; submit(): Promise<{ ok: true }>;
text(): Promise<string>; getAttribute(name: string): Promise<string>;
cssValue(name: string): Promise<string>; tagName(): Promise<string>;
isDisplayed(): Promise<boolean>; isEnabled(): Promise<boolean>; isSelected():
Promise<boolean>; find(by: string, value: string): Promise<any>; findAll(by:
string, value: string): Promise<any[]>; screenshot(path?: string): Promise<{
path?: string; size?: number; bytes?: number[]; format: "png" }> }>; findAll(by:
string, value: string): Promise<Array<{ click(): Promise<{ ok: true }>;
sendKeys(text: string): Promise<{ ok: true }>; clear(): Promise<{ ok: true }>;
submit(): Promise<{ ok: true }>; text(): Promise<string>; getAttribute(name:
string): Promise<string>; cssValue(name: string): Promise<string>; tagName():
Promise<string>; isDisplayed(): Promise<boolean>; isEnabled(): Promise<boolean>;
isSelected(): Promise<boolean>; find(by: string, value: string): Promise<any>;
findAll(by: string, value: string): Promise<any[]>; screenshot(path?: string):
Promise<{ path?: string; size?: number; bytes?: number[]; format: "png" }> }>;
source(): Promise<string>; screenshot(path?: string): Promise<{ path?: string;
size?: number; bytes?: number[]; format: "png" }>; executeScript(js: string,
args?: unknown[]): Promise<unknown>; executeScriptAsync(js: string, args?:
unknown[]): Promise<unknown>; cookies(): Promise<object[]>; setCookie(c: { name:
string; value: string; path?: string; domain?: string; secure?: boolean;
httpOnly?: boolean; expiry?: number }): Promise<{ ok: true }>;
deleteCookie(name: string): Promise<{ ok: true }>; deleteAllCookies(): Promise<{
ok: true }>; setImplicitWait(ms: number): Promise<{ ok: true }>; waitFor(by:
string, value: string, opts?: { timeout?: number; visible?: boolean }):
Promise<any>; windowHandles(): Promise<string[]>; currentWindow():
Promise<string>; switchToWindow(handle: string): Promise<{ ok: true }>;
newWindow(type?: "tab" | "window"): Promise<{ handle: string; type: string }>;
closeWindow(): Promise<string[]>; switchToFrame(target: number | object):
Promise<{ ok: true }>; switchToParentFrame(): Promise<{ ok: true }>;
switchToDefaultContent(): Promise<{ ok: true }>; acceptAlert(): Promise<{ ok:
true }>; dismissAlert(): Promise<{ ok: true }>; alertText(): Promise<string>;
sendAlertText(text: string): Promise<{ ok: true }>; maximize(): Promise<{ x:
number; y: number; width: number; height: number }>; minimize(): Promise<{ x:
number; y: number; width: number; height: number }>; fullscreen(): Promise<{ x:
number; y: number; width: number; height: number }>; setWindowRect(rect: {
width?: number; height?: number; x?: number; y?: number }): Promise<{ x: number;
y: number; width: number; height: number }>; getWindowRect(): Promise<{ x:
number; y: number; width: number; height: number }>; hover(el: object):
Promise<{ ok: true }>; dragAndDrop(src: object, dst: object): Promise<{ ok: true
}>; keyChord(...keys: string[]): Promise<{ ok: true }>; performActions(sequence:
unknown[]): Promise<{ ok: true }>; releaseActions(): Promise<{ ok: true }>;
quit(): Promise<{ closed: true }> }>

```

Connect to a running WebDriver server (opts.url) or start an installed local chromedriver/geckodriver and dial it. Returns a session handle whose methods drive the browser. Sessions are quit on Run end

if the script does not call `quit()` explicitly.

Parameters

- `opts` (*{ browser?: "chrome" | "firefox", headless?: boolean, url?: string, args?: string[], capabilities?: object, commandTimeout?: number }*, optional) — browser selects the driver binary (default 'chrome'). headless defaults to true. url, if given, dials an already-running driver at that base URL instead of starting one. args appends extra browser flags. capabilities is an escape hatch for raw W3C capability overrides merged last. commandTimeout (ms, default 30000) bounds each low-level WebDriver request so a driver blocked behind an open alert or an unreachable endpoint can't hang the call; 0 or negative disables the per-command deadline. quit()/Run-end also cancel any in-flight command.

Returns: Promise resolving to a session handle with methods: `get(url)`, `url()`, `title()`, `back()`, `forward()`, `refresh()`, `find(by, value) → element handle`, `findAll(by, value) → element handle[]`, `source()`, `screenshot(path?)`, `executeScript(js, args?)`, `executeScriptAsync(js, args?)`, `cookies()`, `setCookie(c)`, `deleteCookie(name)`, `deleteAllCookies()`, `setImplicitWait(ms)`, `waitFor(by, value, opts?)`, `quit()`. Phase 2 session methods: `windowHandles()`, `currentWindow()`, `switchToWindow(handle)`, `newWindow(type?)`, `closeWindow()`, `switchToFrame(indexOrElement)`, `switchToParentFrame()`, `switchToDefaultContent()`, `acceptAlert()`, `dismissAlert()`, `alertText()`, `sendAlertText(text)`, `maximize()`, `minimize()`, `fullscreen()`, `setWindowRect({width?,height?,x?,y?})`, `getWindowRect()`, `hover(el)`, `dragAndDrop(src,dst)`, `keyChord(...keys)`, `performActions(sequence)`, `releaseActions()`. Element handles also expose `hover()` and `dragTo(target)` and carry an `elementId` string. `executeScript/executeScriptAsync` return element handles when the script returns an element (or a top-level array of elements). Locator strategies: `css`, `xpath`, `id`, `name`, `tag`, `className`, `linkText`, `partialLinkText`.

Throws: Throws if no url is given and the driver binary for the selected browser is not on PATH; if dialing the driver fails (driver not running, wrong port); or if the browser launch fails. Subsequent session method calls throw if the session is already closed.

```
if (!services.webdriver.available) {
  runtime.log("no driver on PATH – skip");
} else {
  const d = await services.webdriver.connect({ browser: "chrome", headless: true });
  try {
    await d.get("data:text/html,<title>hi</title><h1 id=h>Hello</h1>");
    runtime.log("title:", await d.title());
    const el = await d.find("id", "h");
    runtime.log("text:", await el.text());
    await d.executeScript("return 1+2", []);
  } finally {
    await d.quit();
  }
}
```

services.webdriver.probe

```
probe(opts: { url: string }): Promise<{ ready: boolean; status?: number; error?: string }>
```

Check whether a WebDriver endpoint responds at `opts.url/status`. Returns `{ ready, status }` on HTTP success or `{ ready: false, error }` on transport failure. Does not throw on network errors.

Parameters

- `opts` (*{ url: string }*) — url is required — the base URL of a running WebDriver server (e.g. `'http://127.0.0.1:9515'`).

Returns: `Promise<{ ready, status? } | { ready: false, error }>` — `ready` is true when the endpoint returns HTTP 200; `status` is the HTTP status code when a response was received; `error` is the transport error message when the request failed entirely.

Throws: Throws if `opts.url` is missing or empty. Transport failures resolve with `{ ready: false, error }` rather than throwing.

```
const r = await services.webdriver.probe({ url: "http://127.0.0.1:9515" });
runtime.log("ready:", r.ready);
```

text

String / regex / charset / data manipulation — all text-shaped transforms.

text.charset.decode

```
decode(input: string | Uint8Array | ArrayBuffer, charset: string):
Promise<string>
```

Decode bytes in a named charset to a UTF-8 string.

Parameters

- `input` (*string | Uint8Array | ArrayBuffer*) — Bytes encoded in `charset`.
- `charset` (*string*) — WHATWG/HTML5 encoding name or alias (UTF-8, ISO-8859-1, Windows-1252, Shift_JIS, GBK, ...).

Returns: `Promise` — the bytes decoded to a UTF-8 string.

Throws: Rejects if `input` is an unsupported type, the `charset` name is unknown, or the bytes are invalid for that encoding.

```
const s = await text.charset.decode(bytes, "Windows-1252");
```

text.charset.detect

```
detect(input: string | Uint8Array | ArrayBuffer): Promise<Record<string,
unknown>>
```

Detect the most-likely charset of a byte sequence (saintfish/chardet). Returns top guess + candidates.

Parameters

- `input` (*string* | *Uint8Array* | *ArrayBuffer*) — Bytes to sniff. A string is taken as its raw UTF-8 bytes.

Returns: `Promise<{ charset: string, confidence: number, language?: string, candidates: { charset: string, confidence: number, language?: string }[] }>` — the top match plus all candidates; confidence is chardet's 0–100 score. `language` is present only when chardet reports one.

Throws: Rejects if input is empty, an unsupported type, or chardet finds no candidates.

```
const r = await text.charset.detect(bytes);
runtime.log(r.charset, r.confidence);
```

`text.charset.encode`

```
encode(input: string, charset: string): Promise<Uint8Array>
```

Encode a UTF-8 string to bytes in the named charset.

Parameters

- `input` (*string*) — UTF-8 string to encode.
- `charset` (*string*) — Target WHATWG/HTML5 encoding name or alias.

Returns: `Promise` — the string encoded as bytes in the target charset.

Throws: Rejects if the charset name is unknown, or a character has no representation in the target encoding (no lossy fallback — characters are not silently dropped).

```
const bytes = await text.charset.encode("café", "ISO-8859-1");
```

`text.diff.compare`

```
compare(a: string | Uint8Array | ArrayBuffer, b: string | Uint8Array |
ArrayBuffer, opts?: { context?: number, fromFile?: string, toFile?: string }):
Promise<Record<string, unknown>>
```

Unified-diff two text inputs. `opts`: `context` (default 3), `fromFile` / `toFile` (default 'a' / 'b'). Binary inputs return `{ binary: true }` with an empty diff.

Parameters

- `a` (*string* | *Uint8Array* | *ArrayBuffer*) — The 'from' / left side. A string is taken as its UTF-8 bytes.
- `b` (*string* | *Uint8Array* | *ArrayBuffer*) — The 'to' / right side.
- `opts` (*{ context?: number, fromFile?: string, toFile?: string }, optional*) — `context` is the number of unchanged lines around each hunk (default 3); `fromFile` / `toFile` are the header labels (default 'a' / 'b').

Returns: Promise<{ identical: boolean, binary: boolean, added: number, removed: number, diff: string, format: "unified" }> — diff holds the unified-diff text (empty when identical or binary); added/removed count body +/- lines excluding file headers; binary is true when either input has a NUL byte in its first 8 KB.

Throws: Rejects if either input is an unsupported type, or the unified-diff generation fails.

```
const d = await text.diff.compare("a\n", "b\n");
runtime.log(d.diff, d.added, d.removed);
```

text.jq.query

```
query(data: unknown, filter: string): Promise<unknown>
```

Run a jq filter over data and return the first emitted value (or null).

Parameters

- `data` (*unknown*) — Input value — any JSON-like JS value (object/array/scalar). Passed to gojq directly; integers of any width are normalised so queries don't blow up.
- `filter` (*string*) — A jq filter expression, e.g. `".users[0].name"` or `".items | length"`.

Returns: Promise — the first value the filter emits, or null when it emits nothing (e.g. an optional path like `.a.b?` that misses).

Throws: Rejects if data is undefined, the filter string is empty, the filter fails to parse, or evaluation produces an error.

```
const name = await text.jq.query(obj, ".users[0].name");
```

text.jq.queryAll

```
queryAll(data: unknown, filter: string): Promise<unknown[]>
```

Run a jq filter and drain the iterator into an array.

Parameters

- `data` (*unknown*) — Input value — any JSON-like JS value.
- `filter` (*string*) — A jq filter expression. Use this when the filter explodes a stream, e.g. `".[]"` or `".users[].id"`.

Returns: Promise<unknown[]> — every value the filter emits, in order (empty array when it emits nothing).

Throws: Rejects if data is undefined, the filter string is empty, the filter fails to parse, or evaluation produces an error.

```
const ids = await text.jq.queryAll(obj, ".users[].id");
```

text.preg.match

```
match(pattern: string, subject: string): { match: string; groups: string[];
index: number } | null
```

First hit of `/pattern/flags` against `subject`, or `null`. Returns `{ match, groups, index }`; optional groups that didn't match surface as empty strings.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (forward-slash delimiter only). Engine is Go RE2. Flags: `i` (case-insensitive), `m` (multiline `^/$`), `s` (dotall). `u/U/x` are rejected.
- `subject` (*string*) — The string to search.

Returns: `{ match, groups, index }` for the first match, where `match` is the full match, `groups` holds the numbered submatches after group 0 (unmatched optionals as `""`), and `index` is the byte offset; `null` when there is no match.

Throws: Throws if the pattern is empty, lacks a leading/closing `/`, uses an unsupported flag (`u/U/x`), or fails RE2 compilation (e.g. backreferences/lookaround, which RE2 does not support).

```
const m = text.preg.match("/(\\d+)/", "x42"); // { match: "42", groups: ["42"],
index: 1 }
```

text.preg.matchAll

```
matchAll(pattern: string, subject: string): { match: string; groups: string[];
index: number }[]
```

Every hit of `/pattern/flags` against `subject`, as an array of `{ match, groups, index }` objects.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (RE2 engine; `i/m/s` flags).
- `subject` (*string*) — The string to search.

Returns: Array of `{ match, groups, index }` objects, one per non-overlapping match; empty array when there is none.

Throws: Throws on the same conditions as `preg.match` (bad delimiter, unsupported flag, RE2 compile failure).

```
const all = text.preg.matchAll("/\\d+"/, "1 22 333"); // 3 matches
```

text.preg.replace

```
replace(pattern: string, replacement: string, subject: string): string
```

Substitute every match of `/pattern/flags` in `subject`. Replacement uses Go's `$1` / `${1}` backref syntax — PHP's `\1` form is NOT translated.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (RE2 engine; `i/m/s` flags).
- `replacement` (*string*) — Replacement template using Go's RE2 syntax: `$1 / ${name}` reference groups; use `${1}` to disambiguate. PHP's `\1` backrefs are NOT supported.
- `subject` (*string*) — The string to transform.

Returns: `string` — subject with every match replaced (all occurrences).

Throws: Throws on the same conditions as `preg.match` (bad delimiter, unsupported flag, RE2 compile failure).

```
text.preg.replace("/(\\w+)@/", "${1}_at_", "a@b"); // "a_at_b"
```

`text.preg2.match`

```
match(pattern: string, subject: string): { match: string; groups: string[];  
index: number } | null
```

First hit of `/pattern/flags` via `regexp2` (PCRE). Supports lookahead/lookbehind/backreferences. Same `{ match, groups, index }` shape as `preg`. No linear-time guarantee.

Parameters

- `pattern` (*string*) — A `/regex/flags` delimited pattern run on the .NET-flavoured `regexp2` engine (lookaround, backreferences, possessive quantifiers). Flags: `i`, `m`, `s`, and `x` (ignore-pattern-whitespace, unavailable in `preg`). `u/U` are rejected.
- `subject` (*string*) — The string to search.

Returns: `{ match, groups, index }` for the first match (same shape as `preg.match`); `null` when there is no match.

Throws: Throws if the pattern is empty, lacks a leading/closing `/`, uses an unsupported flag (`u/U`), fails `regexp2` compilation, or errors during matching. Backtracking engine: guard untrusted input with a timeout.

```
const m = text.preg2.match("/(?!<=)\\w+/", "a@host"); // { match: "host", ... }
```

`text.preg2.matchAll`

```
matchAll(pattern: string, subject: string): { match: string; groups: string[];  
index: number }[]
```

Every hit of `/pattern/flags` via `regexp2` (PCRE), as an array of `{ match, groups, index }`.

Parameters

- `pattern` (*string*) — A `/regex/flags` delimited pattern on the `regexp2` engine (`i/m/s/x` flags).
- `subject` (*string*) — The string to search.

Returns: Array of { match, groups, index } objects, one per match; empty array when there is none.

Throws: Throws on the same conditions as `preg2.match`. Backtracking engine: guard untrusted input with a timeout.

```
const all = text.preg2.matchAll("/\\w+/", "a b c"); // 3 matches
```

text.preg2.replace

```
replace(pattern: string, replacement: string, subject: string): string
```

Substitute every match of `/pattern/flags` via `regexp2`. Replacement uses .NET `$1` / `${1}` syntax. Backtracking engine — keep a timeout around untrusted input.

Parameters

- `pattern` (*string*) — A `/regex/flags` delimited pattern on the `regexp2` engine (i/m/s/x flags).
- `replacement` (*string*) — Replacement template using .NET substitution syntax: `$1` / `${1}` reference groups, `$$` is a literal dollar.
- `subject` (*string*) — The string to transform.

Returns: string — subject with every match replaced (startAt 0, count -1).

Throws: Throws on the same conditions as `preg2.match`, or if replacement fails. Backtracking engine: guard untrusted input with a timeout.

```
text.preg2.replace("/(\\w)\\1/", "X", "aabb"); // backref-aware: "XX"
```

text.str.base64Decode

```
base64Decode(input: string): string
```

Standard base64; URL-safe input is accepted via auto-detect.

Parameters

- `input` (*string*) — Standard-alphabet base64 string (with padding).

Returns: string — the decoded bytes interpreted as a UTF-8 string.

Throws: Throws a `TypeError` if input is missing/null/undefined; throws if the input is not valid standard base64.

```
text.str.base64Decode("aGk="); // "hi"
```

text.str.base64Encode

```
base64Encode(input: string): string
```

Standard base64 (with padding).

Parameters

- `input` (*string*) — UTF-8 string to encode (encoded as its raw bytes).

Returns: string — RFC 4648 standard base64 with `=` padding.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.base64Encode("hi"); // "aGk="
```

`text.str.br2nl`

```
br2nl(input: string): string
```

Inverse of `nl2br`:

,
,
→ `'\n'`.

Parameters

- `input` (*string*) — Source text. Any case-insensitive
,
, or
variant is matched.

Returns: string — input with each
variant replaced by a single `\n`.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.br2nl("a<br/>b"); // "a\nb"
```

`text.str.htmlEntityDecode`

```
htmlEntityDecode(input: string): string
```

Decode named and numeric HTML entities to their UTF-8 equivalents.

Parameters

- `input` (*string*) — Text containing HTML entities (named like `&` or numeric like `' / '`).

Returns: string — input with recognised entities decoded to their UTF-8 characters.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.htmlEntityDecode("a & b"); // "a & b"
```

text.str.lpad

```
lpad(input: string, len: number, padChar?: string): string
```

Shortcut for `pad(side: 'left')`.

Parameters

- `input` (*string*) — The string to pad on the left.
- `len` (*number*) — Target rune length.
- `padChar` (*string, optional*) — Pad string; defaults to a single space.

Returns: string — input left-padded to len runes.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.lpad("7", 3, "0"); // "007"
```

text.str.ltrim

```
ltrim(input: string, mask?: string): string
```

Like `trim`, left side only.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset of characters to strip from the left. Defaults to the whitespace set `"\t\n\r\v\f"`.

Returns: string — input with leading mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.ltrim("--x", "-"); // "x"
```

text.str.nl2br

```
nl2br(input: string, xhtml?: boolean): string
```

Replace newlines with
(or
when `xhtml=true`).

Parameters

- `input` (*string*) — Source text. CRLF is normalised to LF first, so each line break yields one tag.
- `xhtml` (*boolean, optional*) — When truthy, emit the self-closing instead of `<br />`. Defaults to false.

Returns: string — input with each `\n` replaced by the chosen tag followed by the original newline.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.nl2br("a\nb"); // "a<br>\nb"
```

text.str.normalizeNewlines

```
normalizeNewlines(input: string, style?: "lf" | "crlf" | "cr"): string
```

Canonicalise any mix of `\r\n`, `\r`, `\n` to the requested style ('lf' | 'crlf' | 'cr').

Parameters

- `input` (*string*) — Text with any mix of CRLF, CR, and LF line endings.
- `style` (*"lf" | "crlf" | "cr", optional*) — Target line-ending style. Defaults to 'lf'.

Returns: string — input with every line ending rewritten to the requested style.

Throws: Throws a `TypeError` if style is not one of 'lf', 'crlf', or 'cr'.

```
text.str.normalizeNewlines("a\r\nb", "lf"); // "a\nb"
```

text.str.pad

```
pad(input: string, len: number, padChar?: string, side?: "right" | "left" | "both"): string
```

Pad to `len` with `padChar` (default ' '). `side` is 'right' (default), 'left', or 'both'.

Parameters

- `input` (*string*) — The string to pad. Returned unchanged when its rune length already meets or exceeds `len`.
- `len` (*number*) — Target rune length. Measured in runes, not bytes.
- `padChar` (*string, optional*) — Pad string; truncated to fit the needed width. Defaults to a single space.
- `side` (*"right" | "left" | "both", optional*) — Which side(s) to pad. 'both' splits the deficit with the extra rune going right. Defaults to 'right'; any unknown value is treated as 'right'.

Returns: string — input padded to `len` runes on the chosen side(s).

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.pad("7", 3, "0", "left"); // "007"
```

text.str.printf

```
printf(format: string, args?: ...unknown): void
```

sprintf + write to stdout.

Parameters

- `format` (*string*) — A Go fmt format string (same verbs as sprintf).
- `args` (*...unknown, optional*) — Values substituted into the verbs.

Returns: void — writes the formatted text directly to process stdout; returns nothing.

Throws: Throws a `TypeError` if called with no arguments (format string required).

```
text.str.printf("%d items\n", 3);
```

text.str.reverse

```
reverse(input: string): string
```

Rune-aware reversal — `reverse('café')` is `'éfac'`.

Parameters

- `input` (*string*) — The string to reverse. Reversed by Unicode code point, not byte, so multi-byte runes stay intact.

Returns: string — the input reversed rune-by-rune.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.reverse("café"); // "éfac"
```

text.str.rpad

```
rpad(input: string, len: number, padChar?: string): string
```

Shortcut for `pad(side: 'right')`.

Parameters

- `input` (*string*) — The string to pad on the right.
- `len` (*number*) — Target rune length.
- `padChar` (*string, optional*) — Pad string; defaults to a single space.

Returns: string — input right-padded to len runes.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.rpad("7", 3, "."); // "7.."
```

text.str.rtrim

```
rtrim(input: string, mask?: string): string
```

Like trim, right side only.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset of characters to strip from the right. Defaults to the whitespace set `"\t\n\r\v\f"`.

Returns: string — input with trailing mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.rtrim("x...", "."); // "x"
```

text.str.sprintf

```
sprintf(format: string, args?: ...unknown): string
```

Go's fmt verbs (%s, %d, %x, %.2f, %v, %t, %q, ...) — not PHP's.

Parameters

- `format` (*string*) — A Go fmt format string. Uses Go verbs: %s string, %d integer, %f / %.2f float, %x hex, %v default, %t bool, %q quoted, %% literal percent.
- `args` (*...unknown, optional*) — Values substituted into the verbs (passed through `.Export()`, so JS numbers arrive as Go int64/float64).

Returns: string — the formatted result.

Throws: Throws a `TypeError` if called with no arguments (format string required). Verb/arg mismatches are not thrown — Go renders `%!verb(...)` error markers inline.

```
text.str.sprintf("%s=%d", "n", 5); // "n=5"
```

text.str.stripHtml

```
stripHtml(input: string): string
```

Remove HTML tags and decode common entities.

Parameters

- `input` (*string*) — HTML source. Anything matching `<...>` is removed.

Returns: string — the input with all `<...>` tag spans deleted.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.stripHtml("<b>hi</b>"); // "hi"
```

text.str.trim

```
trim(input: string, mask?: string): string
```

Strip whitespace (or any char in the optional mask string) from both ends.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset: any character in this string is trimmed (PHP-style, not a prefix). Defaults to the whitespace set `"\t\n\r\v\f"`.

Returns: string — input with leading and trailing mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.trim(" hi "); // "hi"
```

text.str.urlDecode

```
urlDecode(input: string): string
```

Inverse of `urlEncode`.

Parameters

- `input` (*string*) — A form-encoded string (`+` decodes to space, `%XX` to bytes).

Returns: string — the decoded value.

Throws: Throws a `TypeError` if input is missing/null/undefined; throws on malformed percent-escapes.

```
text.str.urlDecode("a+b%26c"); // "a b&c"
```

text.str.urlEncode

```
urlEncode(input: string): string
```

Form-encoding (`'+'` for space). For path segments use `encodeURIComponent` (provided by `goja`).

Parameters

- `input` (*string*) — String to percent-encode. Uses `application/x-www-form-urlencoded` rules: space becomes `+`.

Returns: string — the form-encoded value.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.urlEncode("a b&c"); // "a+b%26c"
```

tui

Multi-pane terminal UI: layout, pane, write, focus.

tui.layout

```
layout(tree: { name: string; title?: string; weight?: number } | { rows:
object[]; weight?: number } | { cols: object[]; weight?: number }): void
```

Declare the pane layout for this Run. Tree nodes: { name, title?, weight?, autoscroll? } (leaf), { rows: [...], weight? } (vertical split), { cols: [...], weight? } (horizontal split). The root node also accepts { mouse?: boolean }. Throws on duplicate names, empty rows/cols, unknown keys, or under --watch.

Parameters

- `tree` (*{ name: string; title?: string; weight?: number } | { rows: object[]; weight?: number } | { cols: object[]; weight?: number }*) — The root layout node. Exactly one of name / rows / cols must be set per node. A leaf (name) becomes a bordered pane addressable via `tui.pane(name)`; name must be a non-empty string and unique across the whole tree. rows stacks children top-to-bottom; cols places them side-by-side; both arrays must be non-empty. weight (positive integer, default 1) sets the child's proportional share of its parent's space. title (string, leaf only) seeds the pane's border caption. Any other key is rejected. The tree is realised over the full terminal as a tview Flex when stdout is a TTY; otherwise it falls back to prefixed-line output. autoscroll (boolean, leaf only, default true) controls whether the pane follows the tail as new lines arrive; set false to keep it pinned at the top. mouse (boolean, root only, default false) enables mouse support: the wheel scrolls the pane under the cursor (without changing which pane is focused) and a left-click focuses the pane under the cursor — at the cost of the terminal's native click-drag text selection (use Shift/Option+drag to select while mouse mode is on). wrap (string, leaf only, default "char") sets line wrapping: "char" wraps mid-word, "word" wraps at word boundaries, "off" disables wrapping (long lines scroll horizontally). color (boolean, leaf only, default true) renders subprocess ANSI as pane colors; set false to strip ANSI and show plain text. wrap and color affect TTY rendering only (the non-TTY fallback always emits plain prefixed lines).

Returns: void — installs the layout and brings up the UI (TTY) or the fallback line writer (non-TTY); the controller is torn down automatically at Run end.

Throws: Throws if called under --watch; if layout was already called this Run; if the tree argument is missing/null/undefined; if any node violates the structure rules (not an object, more than one of name/rows/cols, missing all three, empty rows/cols, unknown key, non-string or empty name, duplicate name, title on a non-leaf, or non-positive/non-integer weight) — the error includes the tree path (e.g. "rows[1].cols[0]"); or if the terminal screen / fallback writer fails to start.

```
tui.layout({ cols: [{ name: "log", title: "Log" }, { name: "out", weight: 2 }]
});
tui.pane("log").writeln("started");
```

tui.onKey

```
onKey(handler: (key: { name: string; rune: string; ctrl: boolean; alt: boolean; shift: boolean }) => void): () => void
```

Register a callback invoked on every keypress (TTY mode) except Ctrl-C. Returns an unsubscribe function. No-op in non-TTY (fallback) mode.

Parameters

- `handler` *((key: { name: string; rune: string; ctrl: boolean; alt: boolean; shift: boolean }) => void)*
 - Called for each keypress with a key descriptor. name is the tcell key name ("Enter", "Up", "Tab", "Ctrl-A", "F1", ...) or "Rune" for a printable character (rune holds it). For Ctrl+letter combinations the name carries the combo (e.g. "Ctrl-A") and the ctrl flag may be false, so check name for control keys. Built-in navigation (Tab focus, PgUp/PgDn/arrows/Home/End scroll) still runs in addition to the handler (coexist model). Ctrl-C always aborts the script and is never delivered.

Returns: An unsubscribe function; call it to stop receiving keys. In non-TTY mode onKey registers nothing and returns a no-op unsubscribe. Note: registering onKey alone does not keep the Run alive — keep an outstanding await (e.g. waitKey or a sleep) if the script should stay open.

Throws: Throws if called before tui.layout, or if the argument is not a function.

```
tui.layout({ rows: [{ name: "log" }] });  
const off = tui.onKey((k) => { if (k.name === "Rune" && k.rune === "q") off();  
});
```

tui.pane

```
pane(name: string): { write(text: string): void; writeln(text: string): void;  
clear(): void; title(text: string): void }
```

Return a Pane handle for a declared pane. Throws if the name wasn't in the layout. Handle methods: write(text), writeln(text), clear(), title(text). services.exec.shell({pane}) streams subprocess I/O into a pane.

Parameters

- `name` *(string)* — The leaf name declared in the tui.layout tree.

Returns: A Pane handle. write(text) appends text (subprocess ANSI SGR colors are translated to pane colors); writeln(text) appends text followed by a newline; clear() empties the pane (no-op in the non-TTY fallback); title(text) updates the pane's border caption (no-op in fallback). All methods return undefined and are safe to call from any callback. The handle can also be passed as the pane option to services.exec.shell to stream a subprocess's stdout/stderr live into the pane.

Throws: Throws if tui.layout has not been called yet this Run, or if name was not declared as a leaf in the layout (the message lists the available pane names).

```
const p = tui.pane("out");
p.title("Output");
p.writeln("hello");
```

tui.waitKey

```
waitKey(...args: unknown[]): Promise<{ name: string; rune: string; ctrl:
boolean; alt: boolean; shift: boolean }>
```

Resolve with the next keypress (TTY mode). One-shot — await again for the next key. Rejects in non-TTY (fallback) mode.

Returns: A Promise resolving to the next key descriptor (same shape as `onKey`'s argument). Concurrent `waitKey` calls resolve FIFO — one keypress resolves the oldest pending call. While a `waitKey` is pending the TUI stays open (this is the idiomatic "press any key to close" hold). Ctrl-C aborts and is never delivered.

Throws: Rejects if called before `tui.layout`, or in non-TTY mode (no interactive terminal), or if the TUI is closed while waiting.

```
tui.layout({ rows: [{ name: "log" }] });
tui.pane("log").writeln("Done. Press any key to close.");
await tui.waitKey();
```

This manual covers sercon v0.49.1. Whenever you add, remove, or change a flag, a binding, or the script API, update this file alongside the help screen (`--help`), the examples walkthrough (`--examples`), and the `CHANGELOG.md` .