

sercon

User Manual

Version 0.5.17

2026-05-26

Repository · <https://github.com/codedeviate/sercon>

License · MIT

sercon manual

Long-form reference for the `pkg/scriptengine` library, the `sercon` CLI, the built-in script API, and the JavaScript runtime surface that `scriptengine` exposes via `goja` and `goja_nodejs`. Examples are runnable — save snippets as `.ts` files and execute with `sercon file.ts` unless otherwise noted.

Table of contents

1. [Overview](#)
 2. [Quickstart](#)
 3. [Library API — `pkg/scriptengine`](#)
 4. [CLI — `sercon`](#)
 5. [Built-in script `api`](#)
 6. [JavaScript runtime built-ins \(`goja`\)](#)
 7. [Async runtime additions \(`goja\_nodejs`\)](#)
 8. [TypeScript support](#)
 9. [Top-level `await`](#)
 10. [Module resolution](#)
 11. [Timeouts and cancellation](#)
 12. [Error semantics](#)
 13. [Type generation \(`.d.ts`\)](#)
 14. [Limitations and gotchas](#)
-

1. Overview

`sercon` is a Go program that runs TypeScript files. It's built from two parts:

- `pkg/scriptengine` — a library you embed in your own Go code. You register Go-callable bindings, then call `Run` / `RunFile` to execute a `.ts` script that talks to them.
- `cmd/sercon` — a thin CLI built on the library, useful for ad-hoc test scripts and as a working example of every binding kind.

The runtime is pure Go: `goja` is the JS engine, `esbuild` (used as a Go library) is the TS→JS transpiler, and `goja_nodejs` provides Promises, `setTimeout`, `console`, and CommonJS `require`. No cgo, no Node.

2. Quickstart

Install:

```
go install github.com/codedeviate/sercon/cmd/sercon@latest
```

Run one of the bundled examples:

```
sercon examples/scripts/smoke.ts examples/scripts/async.ts
```

...or run them all via `make demo`. `examples/scripts/` ships a runnable `.ts` (or `.tsx`) file per feature area — `hash.ts`, `strings.ts`, `path-and-time.ts`, `default-export.ts`, `tsx-demo.ts`, and the original `smoke.ts` / `async.ts` / `hang.ts` (`hang.ts` is the timeout demo and intentionally exits non-zero).

Generate a declaration file for editor autocomplete:

```
sercon --emit-dts api.d.ts
```

Embed in your own Go program:

```
eng := scriptengine.New(scriptengine.Options{
    Timeout:    5 * time.Second,
    ScriptRoot: "./scripts",
})
eng.Register("upper", strings.ToUpper)
val, err := eng.RunFile(ctx, "scripts/main.ts")
```

3. Library API — `pkg/scriptengine`

Engine, Options, New

```
type Options struct {
    Timeout      time.Duration // wall-clock limit per Run; 0 disables
    ScriptRoot   string        // base for require/import resolution
    DisableConsole bool          // turn off the goja_nodejs console module
}

func New(opts Options) *Engine
```

`Engine` is the host of all registered bindings. Construct it once, then call `Run` / `RunFile` as many times as needed — each call gets a *fresh* `*goja.Runtime` so script state never leaks between invocations.

`ScriptRoot` defaults to the working directory at `New` time if left empty. Both `Timeout` and `ScriptRoot` can be overridden on a per-Run basis only by constructing a fresh `Engine`; see [Out-of-scope](#) for the per-Run override on the roadmap.

Registering bindings

Function	Use when...
<code>Register(name, value)</code>	The binding is a pure function, struct, or primitive that doesn't need access to the runtime.
<code>RegisterNamespace(name, members map[string]any)</code>	You want to expose <code>name.x</code> , <code>name.y</code> without defining a Go struct.
<code>RegisterConstructor(name, ctor)</code>	The binding is a Go function whose return value should be treated as a class in the emitted <code>.d.ts</code> . (Runtime semantics today are identical to <code>Register</code> ; the <code>d.ts</code> emitter is the only differentiator.)
<code>RegisterFactory(name, factory func(vm, loop) any)</code>	The binding needs the per-Run <code>*goja.Runtime</code> or <code>*eventloop.EventLoop</code> in scope (typical for Promise-returning I/O).
<code>RegisterNamespaceFactory(name, factory)</code>	Same, but the factory returns the full member map.

Example — synchronous function binding:

```
eng.Register("greet", func(name string) string { return "hi " + name })
```

Example — namespace with a sync member:

```
eng.RegisterNamespace("math2", map[string]any{
    "square": func(n float64) float64 { return n * n },
    "pi":     3.14159,
})
```

Example — async (Promise-returning) binding via `PromisifyAsync`:

```

eng.RegisterFactory("httpGet", func(vm *goja.Runtime, loop
*eventloop.EventLoop) any {
    return scriptengine.PromisifyAsync(vm, loop,
        func(ctx context.Context, call goja.FunctionCall) (map[string]any,
error) {
            url := call.Argument(0).String()
            resp, err := http.Get(url)
            if err != nil {
                return nil, err
            }
            defer resp.Body.Close()
            body, _ := io.ReadAll(resp.Body)
            return map[string]any{"status": resp.StatusCode, "body":
string(body)}, nil
        })
})

```

Run, RunFile

```

func (e *Engine) Run(ctx context.Context, name, source string, opts
...RunOption) (goja.Value, error)
func (e *Engine) RunFile(ctx context.Context, path string, opts
...RunOption) (goja.Value, error)

```

`Run` executes `source` as an entry-script TS. `name` is used in stack traces. The returned value is currently always `undefined` — top-level expression capture is on the backlog.

Both methods respect `Options.Timeout` and `ctx` cancellation, whichever fires first; the resulting error is either `ErrScriptTimeout`, `ctx.Err()`, or the underlying JS exception.

Per-Run options

```

type RunOption func(*runConfig)

func WithScriptRoot(dir string) RunOption

```

Reuse a single Engine across many scripts that each live in their own directory by passing `WithScriptRoot(dir)` to `Run` / `RunFile`. The override applies only to this call; `Options.ScriptRoot` is used for every other Run.

```

_, _ = eng.Run(ctx, "main.ts", source,
scriptengine.WithScriptRoot("/path/to/run42"))

```

Reset

```
func (e *Engine) Reset()
```

Clears every registered binding. Useful when a long-lived Engine is reused across unrelated batches of scripts that each want a clean global namespace. **Not safe to call concurrently with Run / RunFile.**

WriteTypes

```
func (e *Engine) WriteTypes(w io.Writer) error
```

Emits a `.d.ts` describing the registered surface. See section [13. Type generation](#) for what the mapping looks like.

SetDocs and SetMemberDocs

```
func (e *Engine) SetDocs(path, doc string)
func (e *Engine) SetMemberDocs(namespace string, docs map[string]string)
```

Attach JSDoc strings to registered bindings so the emitted `.d.ts` grows readable editor hover. `SetDocs` takes a dotted path — either a bare top-level name (`"greet"`, `"http"`) or `"namespace.member"` for namespace members (`"http.get"`, `"exec.shell"`); `SetMemberDocs` is the convenience for documenting many members of a namespace at once (the namespace itself can still be documented separately via `SetDocs`).

Multi-line docs are written as `\n`-separated strings; the emitter expands them into a standard `* -`prefixed JSDoc block. Single-line docs collapse to `/** ... */`. Calling `SetDocs` with an empty string removes any previous doc for that path. Bindings without a doc entry get no JSDoc block at all — the emitter doesn't insert empty `/** */` placeholders.

`SetDocs` may be called before or after the matching `Register...` call; docs are looked up by name at emit time. Like the registration methods, `SetDocs` must not race with a `Run` / `WriteTypes` / `Reset` on the same engine.

```

eng.Register("greet", func(name string) string { return "hi " + name })
eng.SetDocs("greet", "Greet someone by name.")

eng.RegisterNamespace("math2", map[string]any{
    "pi":      3.14159,
    "squared": func(n float64) float64 { return n * n },
})
eng.SetDocs("math2", "Tiny math helpers.")
eng.SetMemberDocs("math2", map[string]string{
    "pi":      "Circumference / diameter ratio.",
    "squared": "Multiply a number by itself.",
})

```

PromisifyAsync[T] and AsyncBinding

```

type AsyncBinding struct {
    Func      func(goja.FunctionCall) goja.Value
    TSReturnType string
}

func PromisifyAsync[T any](vm *goja.Runtime, loop *eventloop.EventLoop,
    work func(ctx context.Context, call goja.FunctionCall) (T, error),
) AsyncBinding

```

`PromisifyAsync` turns blocking Go work into a JS Promise. It launches the work in a goroutine, parks a `SetTimeout` sentinel so `loop.Run` doesn't return early, and schedules the resolve/reject back onto the event loop. **Required** for any Promise-returning binding — `RunOnLoop` alone is not counted as a live job by the event loop.

`PromisifyAsync` returns an `AsyncBinding` carrier rather than the bare callback. The engine unwraps it to `.Func` at `vm.Set` time so goja's special-case detection of `func(goja.FunctionCall) goja.Value` host-callbacks still fires; the `.d.ts` emitter reads `.TSReturnType` to emit `Promise<T>` (mapped from the generic `T` via reflect at construction time) instead of the previous `unknown`. Hosts assigning the result into a `map[string]any` for `RegisterNamespace` / `RegisterNamespaceFactory` get the unwrap recursively too — no manual work required.

Version

```

import "github.com/codedeviate/sercon/pkg/scriptengine"

fmt.Println(scriptengine.Version) // "0.1.0"

```

Bumped in lockstep with the git tag.

4. CLI — sercon

```
sercon [flags] <script.ts> [script.ts ...]
sercon [flags] -                # read entry script from stdin
sercon --examples | --help | --version
```

Flag	Purpose
<code>-timeout DURATION</code>	Wall-clock limit per script (default <code>10s</code> ; <code>0</code> disables).
<code>-root DIR</code>	Override the script root for <code>require</code> / <code>import</code> resolution.
<code>-emit-dts PATH</code>	Write the example bindings' <code>.d.ts</code> to <code>PATH</code> and exit.
<code>-v</code>	Verbose: trace the rewritten entry-script JS and each module resolution to stderr; also print duration on failure.
<code>-h</code> , <code>--help</code>	In-depth colorized help.
<code>--examples</code>	In-depth colorized walkthrough of every feature.
<code>--version</code>	Print the engine version (plus goja/esbuild build-info versions).
<code>--watch</code>	Re-run on file change under the script root. Debounced 150 ms. Ctrl-C exits.

Each positional argument is either a path to a `.ts` / `.tsx` file or `-` to read an entry script from standard input:

```
echo 'api.log(1 + 2);' | sercon -
sercon prelude.ts -          # prelude then stdin
```

Exit codes are distinct per failure type so shells can react sensibly. When several scripts run, the highest applicable code wins:

Code	Meaning
<code>0</code>	every script passed.
<code>1</code>	CLI usage error (unknown flag, missing script argument, ...).
<code>2</code>	at least one script failed to transpile — never ran.
<code>3</code>	at least one script timed out (<code>-timeout</code>) or was context-cancelled.
<code>4</code>	at least one script ran and threw a JS exception.

`-v` writes lines prefixed with `[sercon]` to stderr. The traces include the full rewritten entry-script JS (the form goja actually runs, after the ESM→CJS rewrite + async IIFE wrapper) and

every module-resolution event, so debugging an unexpected resolve target or a mis-rewritten import is straightforward.

The CLI registers a single `api` namespace; see the next section.

`--watch` : re-run on file change

`sercon --watch <script.ts>` runs the script once and then blocks, re-running it every time a watched file changes under the script root. Useful for iterating on a script body or on the binding surface during development.

```
sercon --watch examples/scripts/smoke.ts
```

What's watched:

- The script root (resolved the same way as the one-shot mode — `-root DIR` if set, else `dirname` of the first script argument).
- Recursively. Symlinks aren't followed.
- New directories that appear during the session are picked up automatically.

What's filtered:

- **File extensions:** `.ts` / `.tsx` / `.js` / `.jsx` / `.json` trigger a re-run; `.d.ts` files match via suffix too. Anything else (Markdown, images, editor swap files, build artefacts) is ignored.
- **Directories:** hidden directories (`.git`, `.vscode`, anything starting with `.`) and `node_modules` are excluded — both because they generate floods of irrelevant events and because recursing into them inflates the watcher's directory count for no gain.

Re-runs are **debounced 150 ms**. Editors typically fire several events per save (write → rename → chmod); collapsing them into a single re-run keeps the loop responsive without thrashing. The debounce window resets on every event, so a rapid burst of saves becomes one re-run after the burst settles.

Each run is delimited by a line like `--- sercon re-run @ HH:MM:SS ---` so the output is visually distinct from the previous run.

`Ctrl-C` (SIGINT) or `SIGTERM` exit cleanly. Per-script failures inside a watch session log as `FAIL` but don't terminate the loop — a syntax error you're trying to fix shouldn't kick you out of watch mode.

The watcher reuses the same `Engine` across runs. Each `Run` already gets a fresh `*goja.Runtime`, so re-running is just re-invoking `runOne` on the existing engine — there's no per-iteration setup cost beyond the transpile + module-resolution work.

5. Built-in script `api`

The CLI exposes the following globals to every script:

```

declare const api: {
  log(...args: unknown[]): void;

  assert: {
    equal(actual: unknown, expected: unknown, msg?: string): void;
    ok(cond: unknown, msg?: string): void;
  };

  http: {
    get(url: string): Promise<{ status: number; body: string }>;
    post(url: string, body?: string): Promise<{ status: number; body:
string }>;
    request(
      method: string,
      url: string,
      opts?: {
        headers?: Record<string, string>;
        body?: string;
        timeout?: number;           // ms; default 30000
        retry?: number;           // re-attempts on transport error / 5xx;
default 0
        follow?: boolean;         // follow redirects; default true
        username?: string;        // basic auth
        password?: string;
      },
    ): Promise<{
      status: number;
      ok: boolean;                // status in [200, 400)
      headers: Record<string, string>; // lower-cased names
      body: string;
      url: string;                // final URL after redirects
    }>;
  };

  time: {
    nowMs(): number;
    sleep(ms: number): Promise<void>;
    format(unixMs: number, fmt: string, tz?: string): string;
  };

  env: {
    get(name: string): string | undefined;
  };

  str: {
    trim(s: string, mask?: string): string;
    ltrim(s: string, mask?: string): string;
    rtrim(s: string, mask?: string): string;
  };

```

```

reverse(s: string): string;
stripHtml(s: string): string;
nl2br(s: string, xhtml?: boolean): string;
br2nl(s: string): string;
base64Encode(s: string): string;
base64Decode(s: string): string;
urlEncode(s: string): string;
urlDecode(s: string): string;
htmlEntityDecode(s: string): string;
pad(s: string, len: number, padChar?: string, side?: "left" | "right" |
"both"): string;
  lpad(s: string, len: number, padChar?: string): string;
  rpad(s: string, len: number, padChar?: string): string;
sprintf(fmt: string, ...args: unknown[]): string;
printf(fmt: string, ...args: unknown[]): void;
normalizeNewlines(s: string, style: "lf" | "crlf" | "cr"): string;
};

path: {
  dirname(p: string): string;
  basename(p: string, suffix?: string): string;
};

hash: {
  md5(data: string): string;
  sha1(data: string): string;
  sha256(data: string): string;
  sha384(data: string): string;
  sha512(data: string): string;
  sha3_256(data: string): string;
  sha3_512(data: string): string;
  blake3(data: string): string;
  crc32(data: string): string;
};

compression: {
  algos(): string[];
  compress(
    algo: "gzip" | "deflate" | "zlib" | "bzip2" | "zstd" | "brotli" |
"lz4" | "xz" | "snappy",
    data: string | ArrayBuffer | Uint8Array,
  ): Promise<Uint8Array>;
  decompress(
    algo: "gzip" | "deflate" | "zlib" | "bzip2" | "zstd" | "brotli" |
"lz4" | "xz" | "snappy",
    data: string | ArrayBuffer | Uint8Array,
  ): Promise<Uint8Array>;
};

```

```

barcode: {
  formats(): string[];
  decodableFormats(): string[];
  encode(
    format: "qr" | "datamatrix" | "aztec" | "pdf417"
      | "code128" | "code39" | "codabar"
      | "ean13" | "ean8" | "upca",
    data: string,
    opts?: { width?: number; height?: number; quietZone?: boolean |
number },
  ): Promise<Uint8Array>;
  decode(
    data: Uint8Array | ArrayBuffer | string,          // PNG / JPEG / WebP
image bytes
    format?: "qr" | "datamatrix" | "aztec"
      | "code128" | "code39" | "code93" | "codabar"
      | "ean13" | "ean8" | "upca" | "upce" | "itf",
  ): Promise<{ format: string; text: string }>;
};

text: {
  detect(data: string | ArrayBuffer | Uint8Array): Promise<{
    charset: string;
    confidence: number;
    language?: string;
    candidates: Array<{ charset: string; confidence: number; language?:
string }>;
  }>;
  decode(data: string | ArrayBuffer | Uint8Array, charset: string):
Promise<string>;
  encode(text: string, charset: string): Promise<Uint8Array>;
};

checkdigit: {
  algos(): string[];
  validate(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    input: string,
  ): boolean;
  compute(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    partial: string,
  ): string;
  inspect(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    input: string,
  ): {
    algo: string;
    input: string;
  };
};

```

```

        valid: boolean;
        given: string;
        computed: string;
    };
};

archive: {
    create(
        destPath: string,
        sources: Array<string | { path: string; name?: string }>,
    ): Promise<{ path: string; format: string; entries: string[]; bytes?:
number }>;
    extract(
        archivePath: string,
        destDir: string,
        opts?: { overwrite?: boolean },
    ): Promise<{ path: string; format: string; dest: string; entries:
string[] }>;
};

diff: {
    compare(
        a: string | ArrayBuffer | Uint8Array,
        b: string | ArrayBuffer | Uint8Array,
        opts?: { context?: number; fromFile?: string; toFile?: string },
    ): Promise<{
        identical: boolean;
        binary: boolean;
        added: number;
        removed: number;
        diff: string;
        format: "unified";
    }>;
};

jq: {
    query(data: unknown, filter: string): Promise<unknown>;
    queryAll(data: unknown, filter: string): Promise<unknown[]>;
};

net: {
    tcp(target: string, opts?: { timeout?: number; port?: string }):
Promise<{
        host: string; port: number; ip: string; latencyMs: number;
    }>;
    dns(host: string, opts?: { types?: ("a" | "aaaa" | "mx" | "txt" |
"cname" | "ns")[] }): Promise<{
        a?: string[]; aaaa?: string[];
        mx?: { preference: number; host: string }[];
    }>;
};

```

```

    txt?: string[]; cname?: string; ns?: string[];
  }>;
  tls(target: string, opts?: { timeout?: number }): Promise<{
    cn: string; issuer: string;
    notBefore: string; notAfter: string;
    daysRemaining: number;
    dnsNames: string[];
    serialNumber: string;
    fingerprintSha256: string;
  }>;
  ntp(host: string, opts?: { timeout?: number; port?: number }):
Promise<{
    serverTime: string;
    offsetMs: number;
    rttMs: number;
    stratum: number;
    referenceTime: string;
    rootDelayMs: number;
    rootDispersionMs: number;
  }>;
  whois(domain: string, opts?: { timeout?: number }): Promise<{
    raw: string;
    domain?: {
      name: string; punycode?: string;
      whoisServer?: string;
      nameServers?: string[]; status?: string[];
      dnssec?: boolean;
      createdAt?: string; updatedAt?: string; expirationDate?:
string;
    };
    registrar?: { name?: string };
  }>;
};

exec: {
  shell(
    cmd: string | string[],
    opts?: {
      cwd?: string;
      env?: Record<string, string>;
      timeout?: number; // ms; default 30000
      stdin?: string;
    },
  ): Promise<{
    stdout: string;
    stderr: string;
    exitCode: number;
    success: boolean;
    durationMs: number;
  }>;
};

```

```

    }>;
    http(
      method: string,
      url: string,
      opts?: {
        headers?: Record<string, string>;
        body?: string;
        timeout?: number;          // ms; default 30000
        follow?: boolean;          // -L
        insecure?: boolean;        // -k
        backend?: "auto" | "recon" | "curl"; // default "auto"
      },
    ): Promise<{
      status: number;
      headers: Record<string, string>;
      body: string;
      durationMs: number;
      backend: "recon" | "curl";
    }>;
  };

  git: {
    branch(opts?: { cwd?: string }): Promise<{
      current: string;          // branch name; empty on detached HEAD
      detached: boolean;
      all: string[];            // local branches
    }>;
    isClean(opts?: { cwd?: string }): Promise<boolean>;
    revParse(rev: string, opts?: { cwd?: string }): Promise<string>;
    status(opts?: { cwd?: string }): Promise<Array<{
      path: string;
      indexStatus: string;      // single-char porcelain code
      workingStatus: string;
    }>>;
    add(paths: string | string[], opts?: { cwd?: string }): Promise<{
      paths: string[] }>;
    commit(message: string, opts?: { cwd?: string; allowEmpty?: boolean }):
    Promise<{ sha: string }>;
    log(opts?: { cwd?: string; limit?: number; revRange?: string }):
    Promise<Array<{
      sha: string;
      shortSha: string;
      author: string;
      email: string;
      timestamp: number;        // unix seconds
      subject: string;
    }>>;
    diffStat(opts?: { cwd?: string; revRange?: string }): Promise<{
      files: number;

```



```

        insertions: number;
        deletions: number;
    }>;
    runText(args: string[], opts?: { cwd?: string }): Promise<{
        stdout: string;
        stderr: string;
        exitCode: number;           // surfaced as data – does NOT throw
    }>;
};

preg: {
    match(pattern: string, subject: string): {
        match: string;
        groups: string[];         // empty string for an optional group that
didn't match
        index: number;
    } | null;
    matchAll(pattern: string, subject: string): Array<{
        match: string;
        groups: string[];
        index: number;
    }>;
    replace(pattern: string, replacement: string, subject: string): string;
};

preg2: {                               // PCRE engine
(lookaround, backrefs)
    match(pattern: string, subject: string): {
        match: string;
        groups: string[];
        index: number;
    } | null;
    matchAll(pattern: string, subject: string): Array<{
        match: string;
        groups: string[];
        index: number;
    }>;
    replace(pattern: string, replacement: string, subject: string): string;
};

jwt: {
    sign(
        claims: Record<string, unknown>,
        secret: string,                               // raw bytes for HS*; PEM-
encoded private key for RS*/PS*/ES*/EdDSA
        opts?: {
            algorithm?:
                | "HS256" | "HS384" | "HS512"           // HMAC
                | "RS256" | "RS384" | "RS512"           // RSA-PKCS1v1.5

```

```

        | "PS256" | "PS384" | "PS512"           // RSA-PSS
        | "ES256" | "ES384" | "ES512"         // ECDSA
        | "EdDSA";                             // Ed25519
    },                                          // default "HS256"
): string;
view(token: string): {
    header: Record<string, unknown>;
    payload: Record<string, unknown>;
    signature: string;                        // raw base64url, no
padding
};
validate(
    token: string,
    secret: string,                          // raw bytes for HS*; PEM-
encoded public key (or certificate) for asymmetric
    opts?: {
        algorithm?:
            | "HS256" | "HS384" | "HS512"
            | "RS256" | "RS384" | "RS512"
            | "PS256" | "PS384" | "PS512"
            | "ES256" | "ES384" | "ES512"
            | "EdDSA";
        audience?: string;
        issuer?: string;
    },
): { valid: true; claims: Record<string, unknown> }
    | { valid: false; reason: string };
};

encrypt: {
    keygen(): { publicKey: string; privateKey: string };
// age X25519
    keygenPgp(opts?: { name?: string; email?: string }):
// PGP keypair (armored blocks)
    { publicKey: string; privateKey: string };
    encrypt(
        data: string | Uint8Array | ArrayBuffer,
        recipients: string | string[],          // age1... public keys
        opts?: { armored?: boolean },          // default false → binary age
format
    ): Uint8Array;
    decrypt(
        ciphertext: string | Uint8Array | ArrayBuffer,
        identities: string | string[],          // AGE-SECRET-KEY-1...
private keys
    ): Uint8Array;                          // auto-detects binary vs
armored
    rekey(
        ciphertext: string | Uint8Array | ArrayBuffer,

```

```

        oldIdentities: string | string[],
        newRecipients: string | string[],
        opts?: { armored?: boolean },           // default: preserve input
format
    ): Uint8Array;
    detectBackend(input: string):
        | { backend: "age";      kind: "public" | "private" }
        | { backend: "pgp";      kind: "public" | "private" }
        | { backend: "unknown" };
    };

    sqlite: {
        open(path: string): Promise<{           // ":memory:" or a
filesystem path
            exec(sql: string, ...params: unknown[]): Promise<{
                rowsAffected: number;
                lastInsertId: number;
            }>;
            query(sql: string, ...params: unknown[]):
Promise<Array<Record<string, unknown>>>;
            queryValue(sql: string, ...params: unknown[]): Promise<unknown>; //
first column of first row, or null
            begin(): Promise<{                   // transaction handle
                exec(sql: string, ...params: unknown[]): Promise<{ rowsAffected:
number; lastInsertId: number }>;
                query(sql: string, ...params: unknown[]):
Promise<Array<Record<string, unknown>>>;
                queryValue(sql: string, ...params: unknown[]): Promise<unknown>;
                commit(): Promise<void>;
                rollback(): Promise<void>;
            }>;
            prepare(sql: string): Promise<{      // compiled-once statement
handle
                exec(...params: unknown[]): Promise<{ rowsAffected: number;
lastInsertId: number }>;
                query(...params: unknown[]): Promise<Array<Record<string,
unknown>>>;
                queryValue(...params: unknown[]): Promise<unknown>;
                close(): Promise<void>;
            }>;
            close(): Promise<void>;
        }>;
    };

    gh: {
        authStatus(): Promise<{
            authenticated: boolean;
            user: string;           // login or empty
            raw: string;           // login on success, stderr / probe note

```

```

on failure
  >;
  prList(opts?: {
    cwd?: string;
    state?: "open" | "closed" | "merged" | "all";
    limit?: number;           // default 30
    author?: string;
  }): Promise<Array<{
    number: number;
    title: string;
    state: string;
    author: string;           // login (flattened from gh's wrapper)
    headRefName: string;
    baseRefName: string;
    url: string;
    createdAt: string;
    updatedAt: string;
  }>>;
  repoView(
    repo?: string,           // "owner/name"; omit to use cwd's repo
    opts?: { cwd?: string },
  ): Promise<{
    name: string;
    owner: string;           // login (flattened)
    description: string;
    url: string;
    defaultBranch: string;   // flattened from defaultBranchRef.name
    visibility: string;
  }>;
};

email: {
  spf(domain: string): Promise<
    | { present: false }
    | {
      present: true;
      record: string;
      mechanisms: string[];
      allPolicy: "pass" | "fail" | "softfail" | "neutral" | "";
    }
  >;
  dmarc(domain: string): Promise<
    | { present: false }
    | {
      present: true;
      record: string;
      tags: Record<string, string>;
      policy?: string;
      subdomain?: string;
    }
  >;
};

```

```

        percent?: string;
        rua?: string;
        ruf?: string;
    }
>;
mtaSts(domain: string): Promise<
    | { present: false }
    | {
        present: true;
        record: string;
        txt: { v: string; id: string };
        policy?: {
            version?: string;
            mode?: "enforce" | "testing" | "none" | string;
            mx?: string[];
            maxAge?: number | string;
        };
        policyError?: string;
    }
>;
tlsRpt(domain: string): Promise<
    | { present: false }
    | {
        present: true;
        record: string;
        tags: Record<string, string>;
        rua?: string;
    }
>;
bimi(domain: string, opts?: { selector?: string }): Promise<
    | { present: false; selector: string }
    | {
        present: true;
        selector: string;
        record: string;
        tags: Record<string, string>;
        l?: string;
        a?: string;
    }
>;
all(domain: string): Promise<{
    domain: string;
    spf:    unknown;
    dmarc:  unknown;
    mtaSts: unknown;
    tlsRpt: unknown;
    bimi:   unknown;
}>;

```

```
};  
};
```

Examples:

```
api.log("hello", 1 + 2);  
api.assert.equal(api.time.nowMs() > 0, true);  
  
const r = await api.http.get("https://example.com");  
api.assert.ok(r.status === 200, `expected 200, got ${r.status}`);  
  
await api.time.sleep(50);  
const home = api.env.get("HOME") ?? "(none)";  
  
api.log(api.hash.sha256("abc"));  
// → ba7816bf8f01cfea414140de5dae2223b00361a396177a9cb410ff61f20015ad
```

`api.http.request(method, url, opts?)` is the full-featured client beyond `get` / `post` : per-call `headers`, `body`, `timeout` (default 30s), `retry` (re-attempts on transport errors and 5xx — never 4xx, which are deterministic; linear backoff capped at 1s), `follow` (redirect control; default true), and `basic-auth username / password`. The result is `{ status, ok, headers, body, url }` — `ok` is `status` in [200,400), `headers` is a lower-cased name→value map, `url` is the final URL after redirects. HTTP 4xx / 5xx don't throw (surfaced via `status / ok`); transport errors and context deadline do.

HTTP bindings use `net/http` with a 5-second default per-request timeout and surface real `Promise<...>` values through the event loop. They are *not* mockable from JS — they go to the real network.

Compression bindings (`api.compression.*`) cover nine pure-Go algorithms behind a uniform interface. Inputs are either strings (interpreted as UTF-8 byte sequences) or any `ArrayBuffer` / `Uint8Array`; outputs are `Uint8Array` — goja's representation of a Go `[]byte` return is a `Uint8Array`, so scripts can read `.length` and iterate directly without an `ArrayBuffer` view. Every algorithm round-trips:

Algorithm	Library
<code>gzip / deflate / zlib</code>	<code>stdlib compress/*</code>
<code>bzip2</code>	<code>stdlib compress/bzip2</code> for read; github.com/dsnet/compress/bzip2 for write
<code>zstd</code>	github.com/klauspost/compress/zstd
<code>brotli</code>	github.com/andybalholm/brotli
<code>lz4</code>	github.com/pierrec/lz4/v4

Algorithm	Library
xz	github.com/ulikunitz/xz
snappy	github.com/golang/snappy

`api.compression.algos()` returns the supported names so scripts can iterate without hard-coding the list. Unknown algorithm names throw a clear error rather than silently returning empty bytes.

Barcode encoders (`api.barcode.*`) cover ten symbologies through one `encode(format, data, opts?)` entry point, all backed by the pure-Go github.com/boombuler/barcode toolkit. Output is a PNG payload as `Uint8Array`; size defaults to `256x256` for 2D codes (QR / DataMatrix / Aztec) and `400x120` for linear ones, overridable via `opts.width / opts.height`. `opts.quietZone` (`true` for a default margin, or a pixel count) pads the rendered bars with a white border — required for EAN / UPC to decode (see the decode quirks below).

- 2D: `qr` (medium error correction, auto mode), `datamatrix`, `aztec` (33% ECC, auto layers), `pdf417` (security level 5).
- Linear: `code128`, `code39` (Mod-43 checksum, ASCII-only), `codabar`.
- Retail: `ean13`, `ean8`, `upca` — content must be numeric and the right length; `boombuler/ean.Encode` dispatches by length so all three variants share the same encoder.

`api.barcode.formats()` returns the supported encode names;

`api.barcode.decodableFormats()` returns the (slightly different) decode set.

Decoding (`api.barcode.decode(data, format?)`) accepts PNG / JPEG / WebP image bytes and reaches into github.com/makiuchi-d/gozxing. Returns `{ format, text }` — `format` is the snake-case label the encoder accepts, so round-trips line up. WebP support comes via a blank-import of `golang.org/x/image/webp` that registers the decoder with `stdlib's image.Decode`.

- **With a format hint:** only that reader runs. Fast, and a mismatch surfaces a clear error rather than a silent miss.
- **Without a hint:** every reader in a priority list runs (2D formats first; the more-constrained 1D readers — UPC/EAN — before the permissive ones — Code 128 / Codabar). First success wins; if no reader matches, the call throws with "no barcode recognised".

Decode quirks worth knowing — these are properties of the underlying libraries, not sercon bugs:

- **pdf417 isn't decodable.** `gozxing` v0.1.1 has no PDF417 reader; pass it as a hint and the call throws "unsupported format". The encoder still emits PDF417 just fine.
- **EAN / UPC need a quiet zone.** Per ISO/IEC 15420 + 15418, retail symbologies require white margin on both sides of the bars. `boombuler's` encoder emits them edge-to-edge, so a bare `barcode.encode("ean13", x) → barcode.decode(...)` round-trip fails on

the decode step. Pass `opts.quietZone` to `encode` to fix it: `true` adds a default margin (10% of the width, floored at 10px), or pass a pixel count for an explicit margin. Real-world EAN scanners need that margin too, so it's good practice for any EAN/UPC you'll actually print.

- **Code 39 decoded text ends with the Mod-43 checksum char.** boombuler's encoder appends it; gozxing returns it.
- **Codabar wraps payload in start/stop chars** like `A...A` on encode; gozxing strips them on decode.

Charset bindings (`api.text.*`) cover detection plus byte-side round-tripping:

- **`detect(data)`** — runs `saintfish/chardet` over the input and returns the top guess (`charset` , `confidence` on a 0–100 scale, optional `language` hint) plus the full candidate list. Input is bytes (string, `ArrayBuffer` , or `Uint8Array`).
- **`encode(text, charset)`** — converts a UTF-8 string to bytes in the target encoding. Characters with no representation in the target encoding cause the encoder to error out rather than silently lose them; pre-process the input yourself if you want lossy behaviour.
- **`decode(data, charset)`** — the inverse: bytes-in-charset → UTF-8 string.

Charset names follow the WHATWG Encoding Living Standard aliases that

`golang.org/x/text/encoding/htmlindex.Get` accepts — `UTF-8` , `ISO-8859-1` , `Windows-1252` , `Shift_JIS` , `GBK` , `GB18030` , `Big5` , `EUC-JP` , `EUC-KR` , etc., plus all documented aliases. Unknown names throw a clear error.

Check-digit bindings (`api.checkdigit.*`) verify and compute the trailing check digit of common numeric codes. All three members are synchronous — the algorithms are small bits of modular arithmetic, no I/O or library dependency. Supported algorithms (returned by `algorithms()`):

Algorithm	Input length	Notes
<code>luhn</code>	any ≥ 2	Credit cards, IMEI, social-security-style IDs. Mod-10 with right-to-left doubling.
<code>isbn10</code>	10	Last position may be <code>X</code> (= 10). Mod-11 with weighted sum.
<code>isbn13</code>	13	Alias of <code>ean13</code> — same algorithm; the 978 / 979 prefix is the only thing that distinguishes ISBN-13 syntactically.
<code>ean13</code>	13	Mod-10, weights <code>1, 3, 1, 3, ...</code> from position 0.
<code>ean8</code>	8	Mod-10, weights <code>3, 1, 3, 1, ...</code> .
<code>upca</code>	12	Mod-10, weights <code>3, 1, 3, 1, ...</code> .

`validate(algo, input)` returns a plain boolean — non-digit characters, wrong length, and unknown algorithm names all surface as `false` so scripts can run quick presence checks without try/catch. `compute(algo, partial)` takes the input *without* its check digit and

returns just that digit (or "X" for ISBN-10 position 10). `inspect(algo, input)` returns the union of both views — useful when you want to display the diagnosis to a human.

Archive bindings (`api.archive.*`) handle the three formats `stdlib` provides without external deps: `.zip`, `.tar`, and `.tar.gz` (also `.tgz`). Format is inferred from the destination's extension on `create` and from the source path on `extract`.

- **`create(destPath, sources)`** — `sources` is an array of either bare paths (the basename is used as the in-archive name) or `{ path, name }` objects (override the in-archive name). Directory sources are recursed; the directory's basename becomes the archive subdir and the rest of the tree is added relative to it. Returns the count + list of entries and the output file size.
- **`extract(archivePath, destDir, opts?)`** — unpacks into `destDir`. `opts.overwrite` (default `false`) controls whether existing files at the destination are clobbered; on `false` a collision errors out (`os.ErrExist`-shaped). Every entry name is run through a zip-slip / tar-slip guard that rejects absolute paths, `..` segments, and anything that would resolve outside `destDir`.

Diff bindings (`api.diff.compare(a, b, opts?)`) produce a unified diff between two text inputs via github.com/pmezard/go-difflib. Inputs are strings (UTF-8) or any byte sequence (`ArrayBuffer` / `Uint8Array`). The result reports:

- `identical` — byte-equal inputs short-circuit; `diff` is empty.
- `binary` — either input has a NUL byte in its first 8 KB. A unified diff of binary content isn't useful so `diff` is empty.
- `added` / `removed` — body-only `+` / `-` line counts (file headers are excluded so the numbers match `git diff --shortstat`).
- `diff` — the unified diff text, or empty when `identical` or `binary` is `true`.
- `format` — always `"unified"` today; reserved for future alternatives (`"side-by-side"`, `"summary"`, ...).

`opts.context` (default 3) is the context-line count; `opts.fromFile` / `opts.toFile` (default `"a"` / `"b"`) set the labels in the diff headers.

JSON-query bindings (`api.jq.*`) run jq filters over JS data via github.com/itchyny/gojq, a pure-Go re-implementation. The filter syntax is the same as command-line jq.

- **`query(data, filter)`** — runs the filter and returns the first emitted value (or `null` when the filter emits nothing).
- **`queryAll(data, filter)`** — drains the result iterator and returns every emitted value as an array.

Data passes through goja's `Export` as `map[string]any` / `[]any` trees, which `gojq`'s runtime accepts directly. There's a small normalisation pass on the way in that converts all sized integer types (`int64`, `int32`, etc.) to plain `int` and `float32` to `float64` — `gojq`'s arithmetic dispatch only knows the two canonical numeric kinds, and goja exports every JS-side integer as `int64`. Without normalisation a query as simple as `.users[].age` would panic in `gojq`.

Parse errors and runtime errors (type mismatches, division by zero, ...) both surface as JS exceptions. Use jq's optional access operator `?` (e.g. `.does.not.exist?`) to suppress missing-path errors and get `null` back instead.

Subprocess bindings (`api.exec.*`) wrap Go's `os/exec`. `shell` is the only entry today; `http` (recon-with-curl-fallback) and `git / gh` wrappers ride on top of it in later 0.4.x cuts.

- **`shell(cmd, opts?)`** — runs a subprocess and waits for it to exit. `cmd` is either a **string** (passed verbatim to `/bin/sh -c` on Unix or `cmd /C` on Windows so pipes, redirects, and globs work as typed) or a **string[]** (treated as `argv`, no shell — use this when arguments could be re-interpreted by the shell). Returns `{ stdout, stderr, exitCode, success, durationMs }`.

Options: `cwd` sets the working directory; `env` is merged on top of the parent process environment (not a replacement); `timeout` is in milliseconds and defaults to **30 000 ms**; `stdin` is piped to the child's standard input.

Process-start failures (host binary not on PATH, permission denied) throw. Context deadline (`timeout`) and engine cancellation throw too — the subprocess never got to choose its own exit code, so it isn't reasonable to surface a fake one. Non-zero exits do **not** throw: `success: false` + the real `exitCode` is what callers want for the usual "ran the linter, expected to be told it failed" flow.

- **`http(method, url, opts?)`** — curl-compatible HTTP client routed through `recon` (preferred) with `curl` as a fallback. Returns `{ status, headers, body, durationMs, backend }`. The choice between the two binaries is normally invisible to scripts — both implement the same curl-style flag set we use here (`-X`, `-H`, `-D`, `--data-binary`, `-L`, `-k`, `-s`). Pick a specific backend with `opts.backend: "recon" | "curl"` when you need to.

Implementation notes: response headers are dumped to a temp file via `-D <path>` and parsed back, rather than relying on `-i`'s body-stream format (recon's `-i` is verbose-debug style with `<` prefixes, incompatible with curl's wire-format `-i`). Request bodies are materialised to a temp file and passed via `--data-binary @<path>` so CR / LF are preserved regardless of backend. Header names in the returned map are lower-cased. On redirect chains (`opts.follow: true`), the *last* response wins — the intermediate 3xx blocks are skipped.

HTTP 4xx / 5xx do **not** throw — they're a normal HTTP outcome and callers branch on `status`. Process-start failures, transport errors (DNS, connection refused, TLS handshake), and context deadline / cancellation throw. The error message includes the backend's `stderr` when present.

Git bindings (`api.git.*`) shell out to the host `git` binary; no pure-Go alternative (`go-git` is heavier and would drift from the user's installed git). Every binding accepts an `opts.cwd` so a single engine can work across multiple checkouts.

- **branch(opts?)** — Reports the current branch (empty string when HEAD is detached) plus every local branch from `for-each-ref refs/heads`. `detached` is true exactly when `symbolic-ref --short HEAD` exited non-zero, which is git's own signal for the state.
- **isClean(opts?)** — Convenience boolean over `git status --porcelain`. True when the porcelain output is empty.
- **revParse(rev, opts?)** — Returns the 40-char SHA for `rev`. Invalid refs throw with git's stderr message included.
- **status(opts?)** — Parses `--porcelain v1` output into `{ path, indexStatus, workingStatus }`. The two single-char status codes are the raw porcelain codes (M, A, D, R, ?, ...) so scripts can dispatch on them without re-parsing strings.
- **add(paths, opts?)** — Stages one path (string) or a list (string[]). `--` is inserted between `add` and the paths so values that look like flags work too. Returns the resolved paths array.
- **commit(message, opts?)** — Creates a new commit; the post-commit HEAD SHA is returned. `allowEmpty:true` toggles `--allow-empty`. An empty message throws before spawning.
- **log(opts?)** — Returns `limit` (default 50) most-recent commits in `revRange` (default HEAD). Format string `%H%x09%h%x09%an%x09%ae%x09%at%x09%s` tab-separates SHA, short SHA, author name, email, unix timestamp, and subject. Newlines and tabs in subjects are already collapsed by git so the parse stays one-line-per-commit.
- **diffStat(opts?)** — Aggregates `git diff --shortstat` into `{ files, insertions, deletions }`. Default range is `HEAD~1..HEAD`. A pure-add or pure-delete diff returns zero for the missing side instead of throwing.
- **runText(args, opts?)** — Escape hatch. Surfaces `{ stdout, stderr, exitCode }` for any `git <args>` invocation. Non-zero exits do **not** throw — that's the whole point of having a generic wrapper; callers branch on `exitCode`. Spawn failures and context cancellation still throw.

GitHub-CLI bindings (`api.gh.*`) wrap the `gh` binary. They respect whatever authentication state `gh auth` is already in; we don't try to swap accounts or manage tokens. Every call uses `gh --json` so the result is structured rather than parsed out of human-readable text.

- **authStatus()** — Probe whether `gh` is installed *and* authenticated. Under the hood this runs `gh api user --jq .login` (machine-friendly: just the login on success, a clear error otherwise) rather than `gh auth status` (multi-line human report). Missing-`gh` and unauthenticated-session both resolve with `{ authenticated: false, ... }` rather than throwing — the whole point of a status probe is that the script branches on it. Only context cancellation throws.
- **prList(opts?)** — Lists pull requests on the repo identified by `opts.cwd` (or the engine's working directory). Defaults: open state, limit 30. `gh`'s `author: { login, ... }` wrapper is flattened to just the login string so scripts can compare directly.
- **repoView(repo?, opts?)** — Returns metadata about a repo. With no argument it asks `gh` about the `cwd`'s repo (works from inside a checkout); pass `"owner/name"` to look up any repo `gh` has access to. Two convenience flattenings are applied: `owner` is a

login string (not `{login, ...}`), and `defaultBranch` is the branch name (not `defaultBranchRef.name`). Empty repos resolve with `defaultBranch: ""` rather than `undefined`.

Regex bindings (`api.preg.*`) accept PHP-style `/pattern/flags` delimited input but run on Go's `stdlib regexp` (RE2). The "preg" naming is a deliberate homage; the semantics are RE2's. That means **no backreferences inside patterns, no lookahead / lookbehind**, and **no possessive quantifiers** — scripts that need those should reach for goja's native `RegExp` (ECMAScript-flavoured) or wait for a dedicated PCRE-compatible binding (`regexp2` is on the backlog).

Supported flags: `i` (case-insensitive), `m` (multiline — `^` and `$` match line boundaries), `s` (dotall — `.` matches newline). The PHP flags `u` (Unicode), `U` (default-ungreedy), and `x` (extended whitespace-ignoring) all surface a clear error rather than silently dropping; unknown flags do too. RE2 is UTF-8 by default, so `u` is unnecessary in any case — the error message says so.

- **`match(pattern, subject)`** — Returns `{ match, groups, index }` for the first hit, or `null` when nothing matches. `groups` is every numbered submatch after group 0, in order; an optional group that didn't match surfaces as an empty string (not `undefined`) so the result type stays stable.
- **`matchAll(pattern, subject)`** — Same shape, drained: returns an array of match objects, one per hit.
- **`replace(pattern, replacement, subject)`** — Substitutes via Go's `$1 / ${1}` backreference syntax. PHP's `\1` form is **not** translated — that would conflict with `\\` escapes that are already legitimate in the replacement. Use `$1 / ${1}` as in Go.

PCRE regex bindings (`api.preg2.*`) are the heavier-engine sibling of `api.preg`. They run on github.com/dlclark/regexp2, a port of .NET's regex engine, which **does** support lookahead, lookbehind, backreferences inside patterns, and possessive quantifiers — everything RE2 (and therefore `api.preg`) can't do. The API is identical (`match / matchAll / replace`, the same `/pattern/flags` delimited syntax, the same `{ match, groups, index }` shape) so switching engines is a one-word change. The flag set adds `x` (ignore-pattern-whitespace), which RE2 couldn't honour; `u / U` still error (`regexp2` is Unicode-aware by default and has no global-ungreedy switch). `replace` uses .NET / `regexp2` `$1 / ${1}` substitution syntax.

The trade-off is the reason `api.preg` exists at all: `regexp2` **backtracks**, so it has no linear-time guarantee. A pathological pattern against adversarial input can blow up exponentially. Use `api.preg` (RE2) by default; reach for `api.preg2` only when you need a PCRE feature, and keep a timeout around untrusted input.

JWT bindings (`api.jwt.*`) wrap [golang-jwt/jwt/v5](https://github.com/dlclark/jwt/v5). The full RFC 7518 algorithm matrix is supported: HMAC (`HS256 / HS384 / HS512`), RSA-PKCS1v1.5 (`RS256 / RS384 / RS512`), RSA-PSS (`PS256 / PS384 / PS512`), ECDSA (`ES256 / ES384 / ES512`), and Ed25519 (EdDSA). Unsupported algorithm names — including the special "none" value, garbage like

"HS999" , and protocol names that aren't JWT signing algos — throw at `sign / validate` time with a named-algorithm error rather than silently falling through.

Key shape. The `secret` parameter takes three forms, picked by shape: HMAC algorithms use the byte string directly; asymmetric algorithms expect a PEM-encoded key (private for `sign` , public — or a certificate the public key can be pulled from — for `validate`); and **any algorithm** accepts a JWK (JSON Web Key) object — a JSON string carrying a "kty" member. JWK is detected by the leading `{ + "kty"` and parsed via [lestrrat-go/jwx/v2/jwk](https://lestrrat-go.com/jwx/v2/jwk) ; the `kty` (`oct` → HMAC bytes, `RSA` , `EC` , `OKP`) determines the key type, so a JWK works with whatever matching `opts.algorithm` you pass. Useful when keys come from a JWKS endpoint or a config file in JWK form. PEM detection is the literal `-----BEGIN` prefix. Both directions of the PEM/bytes cross-check are enforced (JWK input bypasses them — the `kty` is authoritative):

- **PEM secret + HMAC algorithm** throws with a named-algorithm hint ("looks like a PEM-encoded key — set `opts.algorithm` to `RS256 / ES256 / EdDSA / etc.`"). Without this guard, a script that forgets to set `opts.algorithm` after switching to a PEM key would silently sign an HS256 token using the PEM bytes themselves — that's exactly the algorithm-confusion attack pattern in reverse.
- **Plain bytes + asymmetric algorithm** throws with the inverse hint ("needs a PEM-encoded private/public key but secret is plain bytes"). Without this guard the user would get a cryptic "x509: malformed certificate" deep inside `jwt-go`.

The cross-check fires at the binding boundary before `jwt-go` is called, so it's a thrown error rather than a `valid: false` resolution — these are structural input problems, not validation failures.

- **`sign(claims, secret, opts?)`** — Produces a compact-serialised signed JWT. `claims` is passed straight through to `jwt-go`'s `MapClaims` , which recognises the RFC 7519 reserved claims (`exp` , `nbf` , `iat` , `aud` , `iss` , `sub` , `jti`) and lets you layer arbitrary application claims alongside them. Missing reserved claims aren't synthesised — scripts that want `iat` set should compute it explicitly (e.g. `Math.floor(api.time.nowMs() / 1000)`).
- **`view(token)`** — Decodes the header + payload **without verifying the signature**. Useful for debugging auth flows or surfacing `aud` / `iss` to the user before deciding whether to trust the token. `signature` comes back as the raw base64url-encoded segment (no padding); decode it yourself if you need the bytes. Malformed input — wrong segment count, invalid base64, invalid JSON — throws.
- **`validate(token, secret, opts?)`** — Verifies the signature using `secret` and the algorithm declared in the token's header. Standard-claim validation (`exp` , `nbf` , `iat`) is delegated to `jwt-go`. `opts.audience` / `opts.issuer` are cross-checked against `aud` / `iss` when set. The contract is **resolve, don't throw, on validation failure**: `{ valid: false, reason: "..." }` for bad signature, expired, audience mismatch, issuer mismatch, *and* algorithm mismatch when `opts.algorithm` is set. Only structural input

errors (not a JWT at all, empty secret, empty token, PEM/algorithm cross-check failure) throw — pattern-matching on `valid: false` shouldn't accidentally accept a garbage string.

Setting `opts.algorithm` is the **algorithm-confusion guard**: an attacker who has the public key bytes can forge an HS256 token by using those bytes as the HMAC secret (the classic JWT exploit). When `opts.algorithm` is set, the parser's `WithValidMethods` whitelist refuses any token whose header declares a different algorithm. Always set `opts.algorithm` when validating asymmetric tokens; the binding doesn't infer it from the key shape because that would make the wrong default silent.

Encryption (`api.encrypt.*`) supports **two backends** behind one API: age (filippo.io/age, the default) and PGP ([ProtonMail/go-crypto/openpgp](https://protonmail.com/go-crypto/openpgp)). `encrypt` / `decrypt` auto-dispatch on the key/ciphertext format — the same classifier `detectBackend` exposes — so scripts use one API regardless of backend. The age round-trip core (`keygen` / `encrypt` / `decrypt`) landed in v0.5.5; armoured output (`opts.armored`) in v0.5.6; `rekey` in v0.5.7; `detectBackend` in v0.5.8; the PGP backend in v0.5.16.

age identities use the X25519 flavour: `age1...` recipients (safe to share) and `AGE-SECRET-KEY-1...` identities (kept secret). PGP keys are ASCII-armored blocks (`-----BEGIN PGP PUBLIC/PRIVATE KEY BLOCK-----`), generated by `keygenPgp`.

- **keygenPgp(opts?)** — Generate a PGP keypair (RSA 2048). `opts.name` / `opts.email` populate the primary user ID; both optional. Returns armored `{ publicKey, privateKey }` blocks. When `encrypt` sees a PGP public-key block as a recipient (or `decrypt` sees a PGP private-key identity / a `-----BEGIN PGP MESSAGE-----` ciphertext), it routes through `openpgp` instead of age. PGP output is always armored; `opts.armored` is ignored on that path. age and PGP recipient sets can't be mixed in one call — the formats are incompatible.

All three members are synchronous: encryption is pure CPU work with a small API surface, matching the call shape of the other crypto bindings (`api.jwt`, `api.hash`).

- **keygen()** — Generate a fresh X25519 identity. Returns `{ publicKey, privateKey }` as the bech32 strings age writes to disk. Two consecutive calls produce different keys; entropy comes from `crypto/rand` via age. Destructure to use one or the other separately.
- **encrypt(data, recipients, opts?)** — Seal `data` to one or more recipients. Input accepts `string` / `Uint8Array` / `ArrayBuffer`; `recipients` accepts a single string or an array. Default output is the binary age format as `Uint8Array`. Passing `opts.armored: true` wraps it in age's ASCII armor — the `-----BEGIN AGE ENCRYPTED FILE-----` banner + base64 body that's safe to embed in JSON / YAML / email. Either form decrypts via the same `decrypt` call (auto-detected by the leading bytes). Multi-recipient encryption uses age's native multi-recipient header — any one of the listed identities can decrypt the resulting ciphertext, no re-encryption per reader needed.

- **decrypt(ciphertext, identities)** — Open an age-encrypted payload with one of the supplied identities. **Auto-detects binary vs armored input** by sniffing the leading bytes for age's armor banner (whitespace and BOM bytes ahead of the banner are tolerated, so ciphertext pasted from JSON / email containers still detects correctly). age then walks the identities and uses the first that matches a stanza in the header. Returns `Uint8Array`; let scripts decode to a string via `api.text.decode(bytes, "utf-8")` if appropriate (goja doesn't ship `TextDecoder`).
- **rekey(ciphertext, oldIdentities, newRecipients, opts?)** — Re-encrypt a payload for a fresh recipient set without ever exposing the plaintext to the caller. The decrypt+encrypt loop is internal; the plaintext lives only in this function's stack until the encrypt step finishes. Output format defaults to **match the input** — armored in / armored out, binary in / binary out — which is what you almost always want when key-rotating a payload that lives in a fixed location (file, vault row, JSON field). `opts.armored: true | false` forces the output regardless of input. Internally consistent with `encrypt` and `decrypt`: same cross-checks (private as recipient, public as identity), same auto-detect on the input side, same multi-recipient handling on the output side.
- **detectBackend(input)** — Classify a recipient or identity string by encryption backend without parsing or I/O. Returns `{ backend: "age" | "pgp" | "unknown", kind?: "public" | "private" }`. Age recognises three input forms: bech32 X25519 (age1..., classified `kind: "public"`), AGE-SECRET-KEY-1... (`kind: "private"`), and SSH public keys (ssh-rsa / ssh-ed25519, classified `kind: "public"` since age accepts them as recipients). PGP recognises armored block markers (-----BEGIN PGP PUBLIC KEY BLOCK----- / -----BEGIN PGP PRIVATE KEY BLOCK-----). Anything else returns `{ backend: "unknown" }` — false-negative on unfamiliar input is the safer default than guessing. v0.5.8 ships the classifier only; sercon's actual `encrypt` / `decrypt` paths still handle age recipients only. Useful for scripts that need to dispatch on the format ("is this an age recipient or a PGP block?") before deciding which encrypt path — e.g., a config loader that routes some recipients through `api.encrypt.encrypt` and shells out to `gpg --encrypt` for others.

Cross-checks fire at the binding boundary so the common JS-side mistakes throw with named-key hints rather than cryptic bech32 errors deep inside age:

- **Private key as recipient** ("looks like a private key (AGE-SECRET-KEY-...); recipients are public keys (age1...)") — catches a script that mixed up the two halves on the encrypt call.
- **Public key as identity** ("looks like a public key (age1...); identities are private keys (AGE-SECRET-KEY-1...)") — inverse guard on the decrypt call.

Decryption with an identity that wasn't on the recipient list throws with age's own wording forwarded through: `"identity did not match any of the recipients: incorrect identity for recipient block"`. Scripts that want to silently try multiple identities should pass them all in one call (age tries each in turn internally) rather than catching this error.

SQLite (`api.sqlite.*`) is sercon's first **stateful-handle** binding — the shape future protocol bindings (redis / ldap / ...) will reuse. The namespace exposes only `open`; the object it resolves

to carries the real surface. The driver is modernc.org/sqlite, the pure-Go SQLite translation — no cgo, so the project's cgo-free rule is preserved.

- **open(path)** — Connect to a database. `":memory:"` is an in-RAM database that evaporates when the handle is closed; any other string is a filesystem path (created if absent). The connection is `Ping`-ed before `open` resolves, so a bad path surfaces here rather than at the first query. Resolves to a handle object: `{ exec, query, queryValue, close }`.
- **exec(sql, ...params)** — Run a non-query statement (DDL, INSERT, UPDATE, DELETE). Returns `{ rowsAffected, lastInsertId }` — both are driver-reported and may be zero for statements that don't populate them (e.g. `CREATE TABLE`).
- **query(sql, ...params)** — Run a SELECT and return one object per row, keyed by column name. Column types map as: `INTEGER` → number, `REAL` → number, `TEXT` → string, `NULL` → `null`, `BLOB` → `Uint8Array` (unless the bytes are valid UTF-8, in which case they promote to a string — SQLite stores TEXT and BLOB the same way at the storage layer, so this heuristic recovers the common TEXT case while keeping genuinely-binary BLOBs as bytes).
- **queryValue(sql, ...params)** — Run a statement expected to produce a scalar and return the first column of the first row, or `null` when no rows match. Rows beyond the first are discarded. Ideal for `SELECT count(*)`, `PRAGMA user_version`, single-column lookups.
- **begin()** — Open a transaction; resolves to a nested handle with the same `exec / query / queryValue` surface plus `commit()` and `rollback()`. The transaction reuses the exact same code paths as the top-level handle (via an internal executor interface that both `*sql.DB` and `*sql.Tx` satisfy) — only the error prefix differs (`sqlite.tx.exec:` vs `sqlite.exec:`). A transaction holds a pooled connection until it commits or rolls back, so **every begin() must reach one of them** — there's no auto-rollback on handle close in this cut. Once finalized, the transaction is spent: further calls throw `sql: transaction has already been committed or rolled back`, and so do double-commit / commit-after-rollback. A constraint violation inside the transaction throws (with the `sqlite.tx.exec:` prefix) but doesn't auto-roll-back — the script decides whether to retry or roll back.
- **prepare(sql)** — Compile a statement once and resolve to a handle whose `exec / query / queryValue` run it repeatedly with fresh bind params. Unlike the handle / transaction methods, these take **only the params** — no SQL string — since the SQL was fixed at `prepare()` time. Worthwhile for batch-insert or repeated-lookup loops where the per-call parse + plan cost would otherwise dominate. Invalid SQL throws at `prepare()`, not the first `exec()`. Like every handle here it must be `close()`d; a leaked statement pins driver resources. Statements are bound to the database handle, not a transaction — using a prepared statement inside a transaction is out of scope for now.
- **close()** — Release the connection. There's no finalizer — scripts that forget to close leak a connection until the process exits. The documented pattern is `open / use / close`.

Parameters bind positionally as `?` placeholders, in argument order. `goja` exports JS numbers as `int64` / `float64`, strings as `string`, `null` as `nil`, and `Uint8Array` as `[]byte` — all of which the driver accepts directly, so no per-type coercion is needed on the script side. Every handle and transaction method is async (Promise-returning); `open` and `begin` are too.

Hash bindings interpret the input as a UTF-8 byte sequence and return lowercase hex. SHA-3 functions are `sha3_256` / `sha3_512` (the IETF spec uses the underscore so the JS name matches recon's). BLAKE3 uses the upstream lukechampion.com/blake3 reference implementation with a 32-byte output. `crc32` is the IEEE polynomial, zero-padded to 8 hex chars.

String utilities (`api.str.*`) follow PHP-/recon-style semantics where they differ from JS:

- `trim` / `ltrim` / `rtrim` accept an optional mask string; **any character in the mask** is stripped. Default mask is whitespace (`" \t\n\r\v\f"`).
- `reverse` is rune-aware, so `reverse("café")` is `"éfac"`. Grapheme clusters are not handled — that's a separate problem.
- `nl2br` emits `<br>` by default; pass `true` as the second arg for XHTML-style `<br/>`.
- `urlencode` / `urldecode` use the form encoding (`+` for space). For path-segment encoding, use `encodeURIComponent` instead — that one is provided by `goja` natively.
- `pad` / `lpad` / `rpadd` follow recon's `str_pad`: `pad` defaults to side `"right"`; `"left"` and `"both"` are also accepted.
- `sprintf` / `printf` are thin wrappers over Go's `fmt`, so the verb set is Go's (`%s`, `%d`, `%x`, `%.2f`, `%v`, `%t`, `%q`, etc.) — not PHP's. `printf` writes to `stdout`.
- `normalizeNewlines` canonicalises any mix of `\r\n`, `\r`, and `\n` to the requested style.

Path utilities (`api.path.*`) are POSIX (forward-slash). On Windows either pass forward-slash paths or convert separators yourself.

Time formatting (`api.time.format(unixMs, fmt, tz?)`) takes Unix milliseconds (symmetric with `api.time.nowMs()`) and a small strftime token set: `%Y %y %m %d %H %M %S %F %T %j %A %a %B %b %z %Z %%`. Day and month names are in English. Pass a third-argument IANA zone name (e.g. `"UTC"`, `"America/New_York"`) to render in that zone; if omitted, the host's `time.Local` is used. Unknown `%X` tokens are emitted verbatim so a typo is visible rather than silently dropped.

Protocol probes (`api.net.*`) are `stdlib`-backed and hit the real network. All three return Promises (via `PromisifyAsync`) and accept an optional `{ timeout: <ms> }` second arg; the default is 5 seconds.

- `tcp(target, opts?)` — dials TCP, measures round-trip from resolution to handshake. `target` is `"host:port"` or just `"host"` (with `opts.port` overriding the default `80`). Resolves to the remote IP, port, and `latencyMs`.
- `dns(host, opts?)` — looks up A, AAAA, MX, TXT, CNAME, and NS records via `net.Resolver`. Pass `opts.types` to scope the query; empty record sets are omitted from the result object so scripts can probe membership with `if ("mx" in result)`.

- **tls(target, opts?)** — connects with `InsecureSkipVerify` so even expired or hostname-mismatched certs come back inspectable. Returns the leaf cert's CN, issuer CN, validity window, `daysRemaining`, `dnsNames`, serial number, and a SHA-256 fingerprint (lowercase hex). Default port is `443`. Hosts that care about validity should re-validate the cert themselves.
- **ntp(host, opts?)** — queries an NTPv4 server (UDP 123 by default) via github.com/beevik/ntp. Returns the server time (ISO 8601), local-clock offset, round-trip time, stratum, reference time, and the root delay / dispersion values from the server's response. All durations are in milliseconds with sub-millisecond precision.
- **whois(domain, opts?)** — two-hop WHOIS lookup (IANA → registrar's whois server) via github.com/likexian/whois with parsing through github.com/likexian/whois-parser. `raw` is always populated; `domain` and `registrar` are best-effort parsed views (some TLDs the parser doesn't understand will return only `raw`). The whois library doesn't accept a `context.Context` — `opts.timeout` is plumbed through its own per-client setting and the engine's `Options.Timeout` watchdog catches the outer call.

Email-authentication probes (`api.email.*`) read DNS records and surface the published policy. They all return `{ present: false }` when the relevant record is absent (NXDOMAIN or no TXT record matching the marker prefix), so scripts can write a single presence-check pattern across the family.

- **spf(domain)** — queries `TXT(<domain>)` and returns the first record starting with `v=spf1`. `mechanisms` is the tokenised sequence after the version prefix; `allPolicy` summarises the trailing `all` qualifier ("`pass`" for `all` / `+all`, "`fail`" for `-all`, "`softfail`" for `~all`, "`neutral`" for `?all`, empty string when no `all` is present).
- **dmARC(domain)** — queries `TXT(_dmarc.<domain>)` and parses the `v=DMARC1; key=val; ...` form into a flat tag map (keys are case-folded; values retain internal whitespace and commas). `policy` / `subdomain` / `percent` / `rua` / `ruf` are surfaced separately because those are the fields most scripts care about.
- **mtaSts(domain)** — looks up `TXT(_mta-sts.<domain>)` and, on success, fetches the policy file at `https://mta-sts.<domain>/.well-known/mta-sts.txt`. The TXT carries a versioned `id` for change detection; the policy file is where the actual `mode` + `mx` list live. RFC 8461 format (line-based, `key: value`, `mx:` repeats) is parsed inline. Policy-fetch failures (TLS errors, 4xx, timeout) don't fail the binding — the TXT part is still returned, and the fetch error surfaces under `policyError` so scripts can decide what to do.
- **tlsRpt(domain)** — looks up `TXT(_smtp._tls.<domain>)` and parses the `v=TLSRPTv1; rua=...` form into a tag map. `rua` is surfaced separately because that's the actionable bit.
- **bimi(domain, opts?)** — looks up `<selector>._bimi.<domain>`, selector defaulting to `default`. Surfaces the logo URL (`l`) and VMC URL (`a`) from the tag map.

- `all(domain)` — runs every probe above in parallel via goroutines and returns a single aggregate object keyed by probe name. Per-probe failures don't fail the aggregate — they surface under `<probe>.error` so a partial result is still useful (e.g. SPF + DMARC found, MTA-STS policy fetch timed out).

6. JavaScript runtime built-ins (goja)

`scriptengine` runs on goja, which implements **ES5.1 + a large subset of ES6+**. The following are present and behave per spec:

Globals

- `Object`, `Array`, `String`, `Number`, `Boolean`, `BigInt`
- `Math`, `Date`, `JSON`, `RegExp`
- `Error`, `TypeError`, `RangeError`, `SyntaxError`, `ReferenceError`, `EvalError`, `URIError`
- `Promise` (provided via goja's native implementation; resolution microtasks are pumped by the event loop)
- `Symbol` (partial: well-known symbols `Symbol.iterator`, `Symbol.asyncIterator`, etc., are supported)
- `Proxy`, `Reflect` (partial)

```
api.log(Math.PI.toFixed(4), Math.max(1, 9, 3));
api.log(new Date().toISOString());
api.log(JSON.stringify({ a: 1, b: [2, 3] }));
api.log("abc".repeat(3), "abc".padStart(6, "_"));
```

Collections

- `Map`, `Set`, `WeakMap`, `WeakSet`

```
const counts = new Map<string, number>();
for (const ch of "banana") counts.set(ch, (counts.get(ch) ?? 0) + 1);
api.log([...counts]); // [{"b",1}, {"a",3}, {"n",2}]
```

Typed arrays

- `ArrayBuffer`, `DataView`
- `Int8Array`, `Uint8Array`, `Uint8ClampedArray`, `Int16Array`, `Uint16Array`, `Int32Array`, `Uint32Array`, `Float32Array`, `Float64Array`

```
const buf = new ArrayBuffer(4);
new DataView(buf).setUint32(0, 0xCAFEBAFE);
api.log(new Uint8Array(buf)[0].toString(16)); // ca
```

Iteration and generators

- `for...of`, `for...in`, `spread`, `destructuring`
- Generator functions (`function*`)
- Async generators (`async function*`) — supported via goja's event loop

```
function* fib() {
  let [a, b] = [0, 1];
  while (true) {
    yield a;
    [a, b] = [b, a + b];
  }
}
const it = fib();
const first10 = Array.from({ length: 10 }, () => it.next().value);
api.log(first10);
```

Not (yet) provided

- `fetch` — use `api.http.*` from the CLI, or register your own binding.
- `Worker`, `threading` primitives.
- DOM globals (`window`, `document`).
- Native ES modules (`import` / `export` are *transpiled* to CommonJS at load time — see [section 8](#)).

7. Async runtime additions (goja_nodejs)

The event loop bundles a small Node-compatible runtime on top of goja:

Timers

```
setTimeout(() => api.log("after 50ms"), 50);
setInterval(() => api.log("tick"), 100);
setImmediate(() => api.log("next tick"));
clearTimeout(t); clearInterval(i); clearImmediate(im);
```

The event loop holds the script alive while *any* timer is pending. If your script ends with only `setTimeout` callbacks scheduled, the run will wait for them to fire before returning.

console

```
console.log("info");
console.error("oops");
console.warn("careful");
console.info("fyi");
console.debug("trace");
```

Disable with `Options.DisableConsole = true` if you want a clean global namespace.

require

CommonJS `require`, with TS-aware extension fallback added by `sercon` (see [section 10](#)):

```
const helpers = require("./helpers");
helpers.doThing();
```

Both `require("./foo")` and `import { x } from "./foo"` resolve to the same module instance per Run; esbuild rewrites `import` to `require` at transpile time.

8. TypeScript support

TS is transpiled by esbuild in-process. There is no `tsc`, no `tsconfig.json`, no type-checking — types are erased and the resulting JS is what executes. Practical implications:

- All TS syntax esbuild supports is allowed: type annotations, `interface`, `type`, generics, enums (compiled to objects), namespaces, decorators (legacy & TC39 stage 3).
- `import type { X } from "y"` is stripped entirely (no runtime require).
- TS errors are *not* surfaced as build errors — only esbuild parse errors are. Run a separate `tsc --noEmit` in CI if you want type enforcement.
- `tsconfig.json` is **not** consulted.

`.tsx` and `.jsx` are supported via esbuild's `LoaderTSX` / `LoaderJSX`, including for required modules resolved by the engine's extension fallback. JSX has no React runtime in scope by default — either set the factory at the file level with an `@jsx` pragma:

```
/** @jsx h */
function h(tag: string, props: any, ...children: any[]) {
  return { tag, props: props ?? {}, children };
}
export const Box = (label: string) => <div className="box">{label}</div>;
```

...or provide a `React.createElement`-compatible binding from Go and rely on esbuild's default pragma. ESM `export default <value>` also works through the entry-script rewriter's

`__esModule ? .default : m` unwrap, so `import answer from "./mod"` resolves to the default export.

9. Top-level `await`

Top-level `await` works in entry scripts:

```
const r = await api.http.get("https://example.com");
api.log(r.status);
```

How it works under the hood: esbuild rejects top-level `await` with the `cjs` output format, so the engine transpiles the *entry* script with `format=esm`, then rewrites `import` statements to `require()` calls and wraps the rest of the body in:

```
;(async () => { /* your transpiled body */ })().then(__resolve, __reject);
```

`__resolve` and `__reject` are bindings the engine sets on the runtime to capture the final settlement. *Required-module* code (anything loaded through `require`) is transpiled with `format=cjs` and does **not** support top-level `await` — wrap in `(async () => ... )()` yourself if you need it inside a module.

10. Module resolution

The engine plugs a custom source loader into `goja_nodejs/require` that adds:

1. **Direct path:** if the requested file exists on disk, use it. `.ts` / `.tsx` files are transpiled on the fly.
2. **Extension fallback** when the request has no extension: try `.ts`, `.tsx`, `.js`, `.cjs`, `.mjs`, `.json` in that order.
3. **`.js` → `.ts` swap:** if a path ending in `.js` / `.cjs` / `.mjs` is requested but the literal file doesn't exist, the same path ending in `.ts` (or `.tsx`) is tried. Handles `package.json` `main` fields that point at compiled output where only the TypeScript source is on disk.
4. **`package.json` source preference:** when reading a `package.json`, if a `source` field is present and points at an existing `.ts` / `.tsx` file, `main` is rewritten to that path before the registry consumes it. Matches the convention used by parcel, microbundle, and similar bundlers to mark the TS source of truth.
5. **Directory index:** if the resolved path is a directory, try `index.ts`, `index.tsx`, `index.js`.

The base directory follows Node's CommonJS rules — relative imports resolve against the requiring module's directory, courtesy of `goja_nodejs`.

```
// All resolve to ./helpers/assert.ts:
import { check } from "./helpers/assert";
import { check } from "./helpers/assert.ts";
const { check } = require("./helpers/assert");
const { check } = require("./helpers/assert.js");
```

JSON imports work out of the box via either the ESM-style default import (esbuild rewrites it to a `require`) or a direct `require`:

```
import data from "./data.json";
const r = require("./data.json");
```

`node_modules` lookup *works* via the upstream registry, but the source loader doesn't currently transpile through that path — keep TS helpers under `ScriptRoot`.

11. Timeouts and cancellation

Two mechanisms interrupt a running script:

- `Options.Timeout` — wall-clock limit per `Run`. Exceeding it returns `scriptengine.ErrScriptTimeout`.
- `ctx` passed to `Run / RunFile` — cancelling the context returns `ctx.Err()` (typically `context.Canceled` or `context.DeadlineExceeded`).

Both call `vm.Interrupt(...)` and `loop.Terminate()` from a watcher goroutine. This works for **sync** JS — including `while(true){}` — because the goja VM checks the interrupt flag between bytecode steps. A host Go binding that's blocked in cgo or a long syscall isn't interruptible mid-call; if you write such a binding, plumb the engine `ctx` through to it yourself.

```
eng := scriptengine.New(scriptengine.Options{Timeout: 200 *
time.Millisecond})
_, err := eng.Run(ctx, "loop.ts", `while (true) {}`)
// err is scriptengine.ErrScriptTimeout, returned ~200–300ms after Run.
```

12. Error semantics

- A Go binding returning `(T, error)` automatically throws as a JS exception when `err != nil`. The exception's `.message` is `err.Error()`.

```
try { mayFail(); } catch (e) { api.log(String(e)); }
```

- A Go binding that `panic`s with a `*goja.Object` (e.g. `vm.NewGoError(...)`) does the same.

- A script throwing a JS error out of the top level surfaces from `Run` as a Go `error` whose message is the JS error's `.message`.
- Timeout errors satisfy `errors.Is(err, scriptengine.ErrScriptTimeout)`. Cancellation errors satisfy `errors.Is(err, context.Canceled)` or `context.DeadlineExceeded`.

13. Type generation (.d.ts)

`Engine.WriteTypes(w)` walks the registered bindings and emits a TypeScript declaration file. The mapping table (abbreviated):

Go type	TS type
<code>string</code>	<code>string</code>
<code>any numeric / float64</code>	<code>number</code>
<code>bool</code>	<code>boolean</code>
<code>[]T</code>	<code>T[] ; []byte → Uint8Array</code>
<code>map[string]T</code>	<code>Record<string, T></code>
<code>*T</code>	<code>T</code> (transparent unwrap)
<code>error</code> (single return)	omitted (collapses to <code>void</code>)
<code>(T, error)</code> (tuple return)	<code>T</code>
<code>interface{}</code> / <code>any</code>	<code>unknown</code>
<code>goja.FunctionCall</code> parameter	folded into <code>(...args: unknown[])</code>
<code>goja.Value</code> return	<code>unknown</code>
<code>scriptengine.AsyncBinding</code> value (from <code>PromisifyAsync[T]</code>)	<code>Promise<T></code>
self-referential types	<code>unknown</code> (cycle detection bails after depth 4)

```
sercon --emit-dts api.d.ts
```

In your editor, place the resulting `.d.ts` next to your scripts (or reference it via `/// <reference path="..." />`) for autocomplete.

JSDoc comments on declarations

The emitter pulls JSDoc strings from the engine's doc map (set via `Engine.SetDocs` / `Engine.SetMemberDocs`) and renders them above each declaration. Single-line docs collapse to `/** ... */`; multi-line docs expand to a standard `/*`-prefixed block. Bindings without a doc entry emit no JSDoc — the emitter never inserts empty placeholders. Docs are looked up by the same dotted path the binding was registered under, so namespace members get hover text too:

```
// AUTO-GENERATED by scriptengine.Engine.WriteTypes – do not edit by hand.

/** sercon's built-in script surface. ... */
declare const api: {
  hash: {
    /** SHA-256 hex digest of a UTF-8 input. */
    sha256(...args: unknown[]): string;
    /** BLAKE3 hex digest (32-byte output, lukechampine.com/blake3). */
    blake3(...args: unknown[]): string;
    // ...
  };
  // ...
};
```

The CLI's `api.*` surface ships with a curated doc map; see `cmd/sercon/api_docs.go` for the source of truth and add new entries there when adding new bindings (the lockstep rule keeps `--examples` / `MANUAL` / `api.d.ts` in sync, and the doc map is now the seventh artifact in that chain).

14. Limitations and gotchas

- **No native ES modules at runtime.** All `import` is rewritten to `require` by esbuild. Side-effect-only imports run the module body exactly once per Run.
- **No top-level `await` inside required modules.** Wrap with `(async () => ... )()` inside the module if you need it.
- **No `tsconfig.json` honoured.** Module resolution and type checking are independent from any project-level TS config.
- **Per-Run runtime.** Side effects on `globalThis` do not persist between calls to `Run` on the same `Engine`. The require *compile* cache is shared; module *exports* are not.
- **Multi-line ESM imports may stress the entry rewriter.** The current rewriter is line-based and handles the common forms; especially gnarly inputs (comments inside the spec list, very unusual whitespace) fall back to executing as-is, which may throw.
- **`RegisterConstructor` is d.ts-only today.** At runtime it behaves like `Register`. True new-able constructor semantics are on the roadmap.
- **HTTP bindings are real network calls.** `api.http.*` uses `net/http` with a 5s timeout. They are not mockable from JS.

See [OUT-OF-SCOPE.md](#) for the active backlog of deferred ideas.

This manual covers sercon v0.5.17. Whenever you add, remove, or change a flag, a binding, or the script API, update this file alongside the help screen (`--help`), the examples walkthrough (`--examples`), and the `CHANGELOG.md` .