

sercon

User Manual

Version 0.8.0

2026-05-26

Repository · <https://github.com/codedeviate/sercon>

License · MIT

sercon manual

Long-form reference for the `pkg/scriptengine` library, the `sercon` CLI, the built-in script API, and the JavaScript runtime surface that `scriptengine` exposes via `goja` and `goja_nodejs`. Examples are runnable — save snippets as `.ts` files and execute with `sercon file.ts` unless otherwise noted.

Table of contents

1. [Overview](#)
 2. [Quickstart](#)
 3. [Library API — `pkg/scriptengine`](#)
 4. [CLI — `sercon`](#)
 5. [Built-in script `api`](#)
 6. [JavaScript runtime built-ins \(`goja`\)](#)
 7. [Async runtime additions \(`goja\_nodejs`\)](#)
 8. [TypeScript support](#)
 9. [Top-level `await`](#)
 10. [Module resolution](#)
 11. [Timeouts and cancellation](#)
 12. [Error semantics](#)
 13. [Type generation \(`.d.ts`\)](#)
 14. [Limitations and gotchas](#)
-

1. Overview

`sercon` is a Go program that runs TypeScript files. It's built from two parts:

- `pkg/scriptengine` — a library you embed in your own Go code. You register Go-callable bindings, then call `Run` / `RunFile` to execute a `.ts` script that talks to them.
- `cmd/sercon` — a thin CLI built on the library, useful for ad-hoc test scripts and as a working example of every binding kind.

The runtime is pure Go: `goja` is the JS engine, `esbuild` (used as a Go library) is the TS→JS transpiler, and `goja_nodejs` provides Promises, `setTimeout`, `console`, and CommonJS `require`. No cgo, no Node.

2. Quickstart

Install:

```
go install github.com/codedeviate/sercon/cmd/sercon@latest
```

Run one of the bundled examples:

```
sercon examples/scripts/smoke.ts examples/scripts/async.ts
```

...or run them all via `make demo`. `examples/scripts/` ships a runnable `.ts` (or `.tsx`) file per feature area — `hash.ts`, `strings.ts`, `path-and-time.ts`, `default-export.ts`, `tsx-demo.ts`, and the original `smoke.ts` / `async.ts` / `hang.ts` (`hang.ts` is the timeout demo and intentionally exits non-zero).

Generate a declaration file for editor autocomplete:

```
sercon --emit-dts api.d.ts
```

Embed in your own Go program:

```
eng := scriptengine.New(scriptengine.Options{
    Timeout:    5 * time.Second,
    ScriptRoot: "./scripts",
})
eng.Register("upper", strings.ToUpper)
val, err := eng.RunFile(ctx, "scripts/main.ts")
```

3. Library API — `pkg/scriptengine`

Engine, Options, New

```
type Options struct {
    Timeout      time.Duration // wall-clock limit per Run; 0 disables
    ScriptRoot   string        // base for require/import resolution
    DisableConsole bool          // turn off the goja_nodejs console module
    Verbose      io.Writer     // diagnostic traces, prefixed "[sercon] "
    ModuleLoader func(candidatePath string) (source string, found bool,
err error)
}

func New(opts Options) *Engine
```

`ModuleLoader`, when non-nil, is consulted for every require/import candidate path **before** the filesystem — the hook for embedders that want to serve modules from somewhere other than disk (an in-memory FS, a network source, an embedded bundle). goja probes several candidates per specifier (`./x`, `./x.ts`, `./x/index.ts`, ...); the engine also tries the bare path plus the usual extension fallbacks (`.ts` / `.tsx` / `.js` / `.mjs` / `.cjs` / `.json`) against the loader, so a loader can match on a plain suffix. Return `(source, true, nil)` to serve the module (transpiled when the path ends in `.ts` / `.tsx`, exactly like a disk read); `("", false, nil)` to fall through to the filesystem; `("", false, err)` to abort resolution.

```
eng := scriptengine.New(scriptengine.Options{
    ModuleLoader: func(path string) (string, bool, error) {
        if strings.HasSuffix(path, "greeting.ts") {
            return `export const hi = (n: string) => "hi " + n;`, true, nil
        }
        return "", false, nil // fall through to disk
    },
})
```

`Engine` is the host of all registered bindings. Construct it once, then call `Run` / `RunFile` as many times as needed — each call gets a *fresh* `*goja.Runtime` so script state never leaks between invocations.

`ScriptRoot` defaults to the working directory at `New` time if left empty. Both `Timeout` and `ScriptRoot` can be overridden on a per-Run basis only by constructing a fresh `Engine`; see [Out-of-scope](#) for the per-Run override on the roadmap.

Registering bindings

Function	Use when...
<code>Register(name, value)</code>	The binding is a pure function, struct, or primitive that doesn't need access to the runtime.
<code>RegisterNamespace(name, members map[string]any)</code>	You want to expose <code>name.x</code> , <code>name.y</code> without defining a Go struct.
<code>RegisterConstructor(name, ctor)</code>	The binding is a Go function whose return value should be treated as a class in the emitted <code>.d.ts</code> . (Runtime semantics today are identical to <code>Register</code> ; the <code>d.ts</code> emitter is the only differentiator.)
<code>RegisterFactory(name, factory func(vm, loop) any)</code>	The binding needs the per-Run <code>*goja.Runtime</code> or <code>*eventloop.EventLoop</code> in scope (typical for Promise-returning I/O).

Function	Use when...
<code>RegisterNamespaceFactory(name, factory)</code>	Same, but the factory returns the full member map.

Example — synchronous function binding:

```
eng.Register("greet", func(name string) string { return "hi " + name })
```

Example — namespace with a sync member:

```
eng.RegisterNamespace("math2", map[string]any{
    "square": func(n float64) float64 { return n * n },
    "pi":     3.14159,
})
```

Example — async (Promise-returning) binding via `PromisifyAsync`:

```
eng.RegisterFactory("httpGet", func(vm *goja.Runtime, loop
*eventloop.EventLoop) any {
    return scriptengine.PromisifyAsync(vm, loop,
        func(ctx context.Context, call goja.FunctionCall) (map[string]any,
error) {
            url := call.Argument(0).String()
            resp, err := http.Get(url)
            if err != nil {
                return nil, err
            }
            defer resp.Body.Close()
            body, _ := io.ReadAll(resp.Body)
            return map[string]any{"status": resp.StatusCode, "body":
string(body)}, nil
        })
})
```

Run, RunFile

```
func (e *Engine) Run(ctx context.Context, name, source string, opts
...RunOption) (goja.Value, error)
func (e *Engine) RunFile(ctx context.Context, path string, opts
...RunOption) (goja.Value, error)
```

`Run` executes `source` as an entry-script TS. `name` is used in stack traces. The returned value is currently always `undefined` — top-level expression capture is on the backlog.

Both methods respect `Options.Timeout` and `ctx` cancellation, whichever fires first; the resulting error is either `ErrScriptTimeout`, `ctx.Err()`, or the underlying JS exception.

Per-Run options

```
type RunOption func(*runConfig)

func WithScriptRoot(dir string) RunOption
```

Reuse a single Engine across many scripts that each live in their own directory by passing `WithScriptRoot(dir)` to `Run` / `RunFile`. The override applies only to this call; `Options.ScriptRoot` is used for every other `Run`.

```
_, _ = eng.Run(ctx, "main.ts", source,
scriptengine.WithScriptRoot("/path/to/run42"))
```

```
func WithArgs(args []string) RunOption
```

Set the per-script argument vector. The script reads it as `Sercon.argv`, a global with Node/Bun layout: `argv[0]` is `Options.ProgramName` (defaults to `filepath.Base(os.Args[0])`); the CLI sets it to `"sercon"`, `argv[1]` is the running script's path, and the `WithArgs` values follow from index 2. The global is always present — with no args it is just `[programName, scriptPath]`.

```
_, _ = eng.Run(ctx, "main.ts", source, scriptengine.WithArgs([]string{"--
port", "8080"}))
// inside the script: Sercon.argv === [ProgramName, "/abs/main.ts", "--
port", "8080"]
```

Reset

```
func (e *Engine) Reset()
```

Clears every registered binding. Useful when a long-lived Engine is reused across unrelated batches of scripts that each want a clean global namespace. **Not safe to call concurrently with `Run` / `RunFile`.**

WriteTypes

```
func (e *Engine) WriteTypes(w io.Writer) error
```

Emits a `.d.ts` describing the registered surface. See section [13. Type generation](#) for what the mapping looks like.

SetDocs and SetMemberDocs

```
func (e *Engine) SetDocs(path, doc string)
func (e *Engine) SetMemberDocs(namespace string, docs map[string]string)
```

Attach JSDoc strings to registered bindings so the emitted `.d.ts` grows readable editor hover. `SetDocs` takes a dotted path — either a bare top-level name (`"greet"`, `"http"`) or `"namespace.member"` for namespace members (`"http.get"`, `"exec.shell"`); `SetMemberDocs` is the convenience for documenting many members of a namespace at once (the namespace itself can still be documented separately via `SetDocs`).

Multi-line docs are written as `\n`-separated strings; the emitter expands them into a standard `*`-prefixed JSDoc block. Single-line docs collapse to `/** ... */`. Calling `SetDocs` with an empty string removes any previous doc for that path. Bindings without a doc entry get no JSDoc block at all — the emitter doesn't insert empty `/** */` placeholders.

`SetDocs` may be called before or after the matching `Register...` call; docs are looked up by name at emit time. Like the registration methods, `SetDocs` must not race with a `Run` / `WriteTypes` / `Reset` on the same engine.

```
eng.Register("greet", func(name string) string { return "hi " + name })
eng.SetDocs("greet", "Greet someone by name.")

eng.RegisterNamespace("math2", map[string]any{
    "pi":      3.14159,
    "squared": func(n float64) float64 { return n * n },
})
eng.SetDocs("math2", "Tiny math helpers.")
eng.SetMemberDocs("math2", map[string]string{
    "pi":      "Circumference / diameter ratio.",
    "squared": "Multiply a number by itself.",
})
```

PromisifyAsync[T] and AsyncBinding

```
type AsyncBinding struct {
    Func          func(goja.FunctionCall) goja.Value
    TSReturnType string
}

func PromisifyAsync[T any](vm *goja.Runtime, loop *eventloop.EventLoop,
    work func(ctx context.Context, call goja.FunctionCall) (T, error),
) AsyncBinding
```

`PromisifyAsync` turns blocking Go work into a JS Promise. It launches the work in a goroutine, parks a `SetTimeout` sentinel so `loop.Run` doesn't return early, and schedules the

resolve/reject back onto the event loop. **Required** for any Promise-returning binding — `RunOnLoop` alone is not counted as a live job by the event loop.

`PromisifyAsync` returns an `AsyncBinding` carrier rather than the bare callback. The engine unwraps it to `.Func` at `vm.Set` time so goja's special-case detection of `func(goja.FunctionCall) goja.Value` host-callbacks still fires; the `.d.ts` emitter reads `.TSReturnType` to emit `Promise<T>` (mapped from the generic `T` via `reflect` at construction time) instead of the previous `unknown`. Hosts assigning the result into a `map[string]any` for `RegisterNamespace` / `RegisterNamespaceFactory` get the unwrap recursively too — no manual work required.

Version

```
import "github.com/codedeviate/sercon/pkg/scriptengine"

fmt.Println(scriptengine.Version) // "0.1.0"
```

Bumped in lockstep with the git tag.

4. CLI — `sercon`

```
sercon [flags] <script.ts> [script.ts ...]
sercon [flags] - # read entry script from stdin
sercon --examples | --help | --version
```

Flag	Purpose
<code>-timeout DURATION</code>	Wall-clock limit per script (default <code>10s</code> ; <code>0</code> disables).
<code>-root DIR</code>	Override the script root for <code>require</code> / <code>import</code> resolution.
<code>-emit-dts PATH</code>	Write the example bindings' <code>.d.ts</code> to <code>PATH</code> and exit.
<code>-v</code>	Verbose: trace the rewritten entry-script JS and each module resolution to stderr; also print duration on failure.
<code>-h</code> , <code>--help</code>	In-depth colored help.
<code>--examples</code>	In-depth colored walkthrough of every feature.
<code>--version</code>	Print the engine version (plus goja/esbuild build-info versions).
<code>--watch</code>	Re-run on file change under the script root. Debounced 150 ms. Ctrl-C exits.

Each positional argument is either a path to a `.ts` / `.tsx` file or `-` to read an entry script from standard input:

```
echo 'api.runtime.log(1 + 2);' | sercon -
sercon prelude.ts -                # prelude then stdin
```

Passing arguments: `--` and `Sercon.argv`

Everything after a standalone `--` is the script argument vector, exposed to every script as the `Sercon` runtime global:

```
sercon run.ts -- --port 8080 hello
```

`Sercon.argv` uses the Node/Bun layout `[programName, scriptPath, ...userArgs]` — `argv[0]` is `"sercon"`, `argv[1]` is the path of the script currently executing, and the post-`--` arguments start at index 2 (so `Sercon.argv.slice(2)` is just your args). All scripts in one invocation share the same `--` tail, but each sees its own path at `argv[1]`. The global is always present; with no `--` it is just `[programName, scriptPath]`.

```
const args = Sercon.argv.slice(2); // ["--port", "8080", "hello"]
```

Exit codes are distinct per failure type so shells can react sensibly. When several scripts run, the highest applicable code wins:

Code	Meaning
0	every script passed.
1	CLI usage error (unknown flag, missing script argument, ...).
2	at least one script failed to transpile — never ran.
3	at least one script timed out (<code>-timeout</code>) or was context-cancelled.
4	at least one script ran and threw a JS exception.

`-v` writes lines prefixed with `[sercon]` to `stderr`. The traces include the full rewritten entry-script JS (the form goja actually runs, after the ESM→CJS rewrite + async IIFE wrapper) and every module-resolution event, so debugging an unexpected resolve target or a mis-rewritten import is straightforward.

The CLI registers a single `api` namespace; see the next section.

`--watch`: re-run on file change

`sercon --watch <script.ts>` runs the script once and then blocks, re-running it every time a watched file changes under the script root. Useful for iterating on a script body or on the

binding surface during development.

```
sercon --watch examples/scripts/smoke.ts
```

What's watched:

- The script root (resolved the same way as the one-shot mode — `-root DIR` if set, else `dirname` of the first script argument).
- Recursively. Symlinks aren't followed.
- New directories that appear during the session are picked up automatically.

What's filtered:

- **File extensions:** `.ts` / `.tsx` / `.js` / `.jsx` / `.json` trigger a re-run; `.d.ts` files match via suffix too. Anything else (Markdown, images, editor swap files, build artefacts) is ignored.
- **Directories:** hidden directories (`.git`, `.vscode`, anything starting with `.`) and `node_modules` are excluded — both because they generate floods of irrelevant events and because recursing into them inflates the watcher's directory count for no gain.

Re-runs are **debounced 150 ms**. Editors typically fire several events per save (write → rename → chmod); collapsing them into a single re-run keeps the loop responsive without thrashing. The debounce window resets on every event, so a rapid burst of saves becomes one re-run after the burst settles.

Each run is delimited by a line like `--- sercon re-run @ HH:MM:SS ---` so the output is visually distinct from the previous run.

`Ctrl-C` (SIGINT) or `SIGTERM` exit cleanly. Per-script failures inside a watch session log as `FAIL` but don't terminate the loop — a syntax error you're trying to fix shouldn't kick you out of watch mode.

The watcher reuses the same `Engine` across runs. Each `Run` already gets a fresh `*goja.Runtime`, so re-running is just re-invoking `runOne` on the existing engine — there's no per-iteration setup cost beyond the transpile + module-resolution work.

Module-graph invalidation. Watch mode tracks each entry script's import graph (its own file plus every module it resolves, captured via `Engine.SetResolveHook`) during the run. On a file change it re-runs **only the entries whose graph includes the changed file** — editing a helper that just one of three entry scripts imports re-runs that one entry, not all three. An entry's own file counts (editing it re-runs it); stdin entries and any entry whose graph isn't known yet re-run unconditionally (conservative). Before each re-run the module cache is busted (`Engine.ResetModuleCache`) so an edited import's new source actually takes effect — the registry otherwise caches compiled bytecode across runs.

5. Built-in script `api`

Migrating from 0.7

`api.*` was reorganised in 0.8.0 from 32 flat groups into 9 category buckets. Every script written for 0.7 or earlier needs path updates. The mapping:

Old path	New path
<code>api.log</code>	<code>api.runtime.log</code>
<code>api.assert.equal</code>	<code>api.runtime.assert.equal</code>
<code>api.assert.ok</code>	<code>api.runtime.assert.ok</code>
<code>api.time.*</code>	<code>api.runtime.time.*</code>
<code>api.env.get</code>	<code>api.runtime.env.get</code>
<code>api.hash.*</code>	<code>api.crypto.hash.*</code>
<code>api.jwt.*</code>	<code>api.crypto.jwt.*</code>
<code>api.encrypt.*</code>	<code>api.crypto.encrypt.*</code>
<code>api.str.*</code>	<code>api.text.str.*</code>
<code>api.preg.*</code>	<code>api.text.preg.*</code>
<code>api.preg2.*</code>	<code>api.text.preg2.*</code>
<code>api.text.detect</code>	<code>api.text.charset.detect</code>
<code>api.text.encode</code>	<code>api.text.charset.encode</code>
<code>api.text.decode</code>	<code>api.text.charset.decode</code>
<code>api.jq.*</code>	<code>api.text.jq.*</code>
<code>api.diff.compare</code>	<code>api.text.diff.compare</code>
<code>api.compression.*</code>	<code>api.format.compression.*</code>
<code>api.barcode.*</code>	<code>api.format.barcode.*</code>
<code>api.checkdigit.*</code>	<code>api.format.checkdigit.*</code>
<code>api.path.*</code>	<code>api.fs.path.*</code>
<code>api.archive.*</code>	<code>api.fs.archive.*</code>
<code>api.http.*</code>	<code>api.net.http.*</code>
<code>api.net.tcp</code>	<code>api.net.probe.tcp</code>
<code>api.net.dns</code>	<code>api.net.probe.dns</code>
<code>api.net.tls</code>	<code>api.net.probe.tls</code>
<code>api.net.ntp</code>	<code>api.net.probe.ntp</code>
<code>api.net.whois</code>	<code>api.net.probe.whois</code>

Old path	New path
<code>api.net.ping</code>	<code>api.net.probe.ping</code>
<code>api.net.smtp</code>	<code>api.net.probe.smtp</code>
<code>api.net.wss</code>	<code>api.net.probe.wss</code>
<code>api.netstatus.*</code>	<code>api.net.netstatus.*</code>
<code>api.email.*</code>	<code>api.net.email.*</code>
<code>api.browser.*</code>	<code>api.net.browser.*</code>
<code>api.sqlite.*</code>	<code>api.db.sqlite.*</code>
<code>api.redis.*</code>	<code>api.db.redis.*</code>
<code>api.memcached.*</code>	<code>api.db.memcached.*</code>
<code>api.ldap.*</code>	<code>api.db.ldap.*</code>
<code>api.dict.*</code>	<code>api.db.dict.*</code>
<code>api.exec.shell</code>	<code>api.tools.exec.shell</code>
<code>api.exec.http</code>	<code>api.tools.exec.http</code>
<code>api.git.*</code>	<code>api.tools.git.*</code>
<code>api.gh.*</code>	<code>api.tools.gh.*</code>
<code>api.ai.*</code>	<code>api.tools.ai.*</code>
<code>api.tui.layout</code>	<code>api.ui.tui.layout</code>
<code>api.tui.pane</code>	<code>api.ui.tui.pane</code>

Every leaf function name stays the same; only the addressing changes. A mechanical sed pass over your scripts handles it — see the `CHANGELOG.md` 0.8.0 "Changed" entry for the canonical sed command.

The CLI exposes the following globals to every script:

```

declare const api: {
  log(...args: unknown[]): void;

  assert: {
    equal(actual: unknown, expected: unknown, msg?: string): void;
    ok(cond: unknown, msg?: string): void;
  };

  http: {
    get(url: string): Promise<{ status: number; body: string }>;
    post(url: string, body?: string): Promise<{ status: number; body:
string }>;
    request(
      method: string,
      url: string,
      opts?: {
        headers?: Record<string, string>;
        body?: string;
        timeout?: number;           // ms; default 30000
        retry?: number;           // re-attempts on transport error / 5xx;
default 0
        follow?: boolean;         // follow redirects; default true
        username?: string;        // basic auth
        password?: string;
      },
    ): Promise<{
      status: number;
      ok: boolean;                // status in [200, 400)
      headers: Record<string, string>; // lower-cased names
      body: string;
      url: string;                // final URL after redirects
    }>;
  };

  time: {
    nowMs(): number;
    sleep(ms: number): Promise<void>;
    format(unixMs: number, fmt: string, tz?: string): string;
  };

  env: {
    get(name: string): string | undefined;
  };

  str: {
    trim(s: string, mask?: string): string;
    ltrim(s: string, mask?: string): string;
    rtrim(s: string, mask?: string): string;
  };

```

```

    reverse(s: string): string;
    stripHtml(s: string): string;
    nl2br(s: string, xhtml?: boolean): string;
    br2nl(s: string): string;
    base64Encode(s: string): string;
    base64Decode(s: string): string;
    urlEncode(s: string): string;
    urlDecode(s: string): string;
    htmlEntityDecode(s: string): string;
    pad(s: string, len: number, padChar?: string, side?: "left" | "right" |
"both"): string;
    lpad(s: string, len: number, padChar?: string): string;
    rpad(s: string, len: number, padChar?: string): string;
    sprintf(fmt: string, ...args: unknown[]): string;
    printf(fmt: string, ...args: unknown[]): void;
    normalizeNewlines(s: string, style: "lf" | "crlf" | "cr"): string;
};

path: {
    dirname(p: string): string;
    basename(p: string, suffix?: string): string;
};

hash: {
    md5(data: string): string;
    sha1(data: string): string;
    sha256(data: string): string;
    sha384(data: string): string;
    sha512(data: string): string;
    sha3_256(data: string): string;
    sha3_512(data: string): string;
    blake3(data: string): string;
    crc32(data: string): string;
};

compression: {
    algos(): string[];
    compress(
        algo: "gzip" | "deflate" | "zlib" | "bzip2" | "zstd" | "brotli" |
"lz4" | "xz" | "snappy",
        data: string | ArrayBuffer | Uint8Array,
    ): Promise<Uint8Array>;
    decompress(
        algo: "gzip" | "deflate" | "zlib" | "bzip2" | "zstd" | "brotli" |
"lz4" | "xz" | "snappy",
        data: string | ArrayBuffer | Uint8Array,
    ): Promise<Uint8Array>;
};

```

```

barcode: {
  formats(): string[];
  decodableFormats(): string[];
  encode(
    format: "qr" | "datamatrix" | "aztec" | "pdf417"
      | "code128" | "code39" | "codabar"
      | "ean13" | "ean8" | "upca",
    data: string,
    opts?: { width?: number; height?: number; quietZone?: boolean |
number },
  ): Promise<Uint8Array>;
  decode(
    data: Uint8Array | ArrayBuffer | string,          // PNG / JPEG / WebP
image bytes
    format?: "qr" | "datamatrix" | "aztec"
      | "code128" | "code39" | "code93" | "codabar"
      | "ean13" | "ean8" | "upca" | "upce" | "itf",
  ): Promise<{ format: string; text: string }>;
};

text: {
  detect(data: string | ArrayBuffer | Uint8Array): Promise<{
    charset: string;
    confidence: number;
    language?: string;
    candidates: Array<{ charset: string; confidence: number; language?:
string }>;
  }>;
  decode(data: string | ArrayBuffer | Uint8Array, charset: string):
Promise<string>;
  encode(text: string, charset: string): Promise<Uint8Array>;
};

checkdigit: {
  algos(): string[];
  validate(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    input: string,
  ): boolean;
  compute(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    partial: string,
  ): string;
  inspect(
    algo: "luhn" | "isbn10" | "isbn13" | "ean13" | "ean8" | "upca",
    input: string,
  ): {
    algo: string;
    input: string;
  };
};

```

```

        valid: boolean;
        given: string;
        computed: string;
    };
};

archive: {
    create(
        destPath: string,
        sources: Array<string | { path: string; name?: string }>,
    ): Promise<{ path: string; format: string; entries: string[]; bytes?:
number }>;
    extract(
        archivePath: string,
        destDir: string,
        opts?: { overwrite?: boolean },
    ): Promise<{ path: string; format: string; dest: string; entries:
string[] }>;
};

diff: {
    compare(
        a: string | ArrayBuffer | Uint8Array,
        b: string | ArrayBuffer | Uint8Array,
        opts?: { context?: number; fromFile?: string; toFile?: string },
    ): Promise<{
        identical: boolean;
        binary: boolean;
        added: number;
        removed: number;
        diff: string;
        format: "unified";
    }>;
};

jq: {
    query(data: unknown, filter: string): Promise<unknown>;
    queryAll(data: unknown, filter: string): Promise<unknown[]>;
};

net: {
    tcp(target: string, opts?: { timeout?: number; port?: string }):
Promise<{
        host: string; port: number; ip: string; latencyMs: number;
    }>;
    dns(host: string, opts?: { types?: ("a" | "aaaa" | "mx" | "txt" |
"cname" | "ns")[] }): Promise<{
        a?: string[]; aaaa?: string[];
        mx?: { preference: number; host: string }[];
    }>;
};

```



```

    txt?: string[]; cname?: string; ns?: string[];
  }>;
  tls(target: string, opts?: { timeout?: number }): Promise<{
    cn: string; issuer: string;
    notBefore: string; notAfter: string;
    daysRemaining: number;
    dnsNames: string[];
    serialNumber: string;
    fingerprintSha256: string;
  }>;
  ntp(host: string, opts?: { timeout?: number; port?: number }):
Promise<{
    serverTime: string;
    offsetMs: number;
    rttMs: number;
    stratum: number;
    referenceTime: string;
    rootDelayMs: number;
    rootDispersionMs: number;
  }>;
  whois(domain: string, opts?: { timeout?: number }): Promise<{
    raw: string;
    domain?: {
      name: string; punycode?: string;
      whoisServer?: string;
      nameServers?: string[]; status?: string[];
      dnssec?: boolean;
      createDate?: string; updatedAt?: string; expirationDate?:
string;
    };
    registrar?: { name?: string };
  }>;
  ping(host: string, opts?: {
    count?: number;           // default 4
    timeout?: number;         // ms; default 5000
    mode?: "tcp" | "icmp";    // default "tcp"
    port?: string;            // tcp mode; default "80"
  }): Promise<{
    host: string; ip: string; mode: string;
    sent: number; received: number; lossPercent: number;
    minMs: number; avgMs: number; maxMs: number;
  }>;
  smtp(host: string, opts?: {
    port?: string;            // default "25"
    timeout?: number;         // ms; default 10000
    ehloName?: string;        // default "localhost"
  }): Promise<{
    host: string; port: string; banner: string; ehloDomain: string;
    extensions: string[]; starttls: boolean;

```

```

    authMechanisms: string[]; sizeLimit: number;
  }>;
  wss(url: string, opts?: {
    timeout?: number;          // ms; default 10000
    ping?: boolean;           // measure ping/pong RTT; default true
  }): Promise<{
    url: string; connected: boolean; subprotocol: string;
    status: number; handshakeMs: number; pingMs: number; // pingMs -1 if
skipped/failed
  }>;
};

  netstatus: {
    check(host: string, opts?: { port?: string; timeout?: number }):
Promise<{
      host: string; port: string; elapsedMs: number; reachable: boolean;
      dns: { ok: boolean; ips: string[]; error?: string };
      tcp: { ok: boolean; latencyMs: number; error?: string };
      tls: { ok: boolean; daysRemaining: number; error?: string };
      http: { ok: boolean; status: number; error?: string };
    }>;
  };

  redis: {
    open(url: string): Promise<{ //
redis://[:pass@]host:port/db
      do(command: string, ...args: unknown[]): Promise<unknown>; //
arbitrary RESP command; null on missing key
      ping(): Promise<string>;
      close(): Promise<void>;
    }>;
  };

  memcached: {
    open(addr: string): Promise<{ // host:port
      get(key: string): Promise<string | null>; // null on cache miss
      set(key: string, value: string | Uint8Array, expirySeconds?: number):
Promise<void>;
      delete(key: string): Promise<boolean>; // true if the key
existed
    }>;
  };

  ldap: {
    open(url: string, opts?: { bindDN?: string; password?: string }):
Promise<{
      rootDSE(): Promise<Record<string, unknown>>; // server
metadata entry
      search(baseDN: string, filter: string, attrs?: string[]):

```

```

Promise<Array<Record<string, unknown>>>;
    close(): Promise<void>;
    }>;
};

dict: {
    define(host: string, word: string, opts?: { database?: string; port?:
string; timeout?: number }): Promise<{
        word: string; found: boolean;
        definitions: Array<{ db: string; dbName: string; text: string }>;
    }>;
    match(host: string, word: string, opts?: { database?: string;
strategy?: string; port?: string; timeout?: number }): Promise<{
        word: string; matches: Array<{ db: string; word: string }>;
    }>;
};

ai: {
    providers(): string[]; // detected CLIs on PATH,
preference order
    send(opts: {
        prompt: string;
        provider?: string; // default: first
detected
        system?: string;
        context?: string;
        timeout?: number; // ms; default 120000
    }): Promise<{ provider: string; output: string; exitCode: number }>;
};

browser: {
    open(): Promise<{ // stateful HTTP session
        setUserAgent(ua: string): void;
        setHeader(name: string, value: string): void;
        get(url: string): Promise<{ status: number; ok: boolean; headers:
Record<string, string>; body: string; url: string }>;
        post(url: string, body?: string): Promise<{ status: number; ok:
boolean; headers: Record<string, string>; body: string; url: string }>;
        cookies(url: string): Promise<Array<{ name: string; value: string
}>>;
    }>;
};

exec: {
    shell(
        cmd: string | string[],
        opts?: {
            cwd?: string;
            env?: Record<string, string>;

```

```

        timeout?: number; // ms; default 30000
        stdin?: string;
    },
): Promise<{
    stdout: string;
    stderr: string;
    exitCode: number;
    success: boolean;
    durationMs: number;
}>;
http(
    method: string,
    url: string,
    opts?: {
        headers?: Record<string, string>;
        body?: string;
        timeout?: number; // ms; default 30000
        follow?: boolean; // -L
        insecure?: boolean; // -k
        backend?: "auto" | "recon" | "curl"; // default "auto"
    },
): Promise<{
    status: number;
    headers: Record<string, string>;
    body: string;
    durationMs: number;
    backend: "recon" | "curl";
}>;
};

git: {
    branch(opts?: { cwd?: string }): Promise<{
        current: string; // branch name; empty on detached HEAD
        detached: boolean;
        all: string[]; // local branches
    }>;
    isClean(opts?: { cwd?: string }): Promise<boolean>;
    revParse(rev: string, opts?: { cwd?: string }): Promise<string>;
    status(opts?: { cwd?: string }): Promise<Array<{
        path: string;
        indexStatus: string; // single-char porcelain code
        workingStatus: string;
    }>>;
    add(paths: string | string[], opts?: { cwd?: string }): Promise<{
paths: string[] }>;
        commit(message: string, opts?: { cwd?: string; allowEmpty?: boolean }):
Promise<{ sha: string }>;
        log(opts?: { cwd?: string; limit?: number; revRange?: string }):
Promise<Array<{

```

```

    sha: string;
    shortSha: string;
    author: string;
    email: string;
    timestamp: number;          // unix seconds
    subject: string;
  }>;
  diffStat(opts?: { cwd?: string; revRange?: string }): Promise<{
    files: number;
    insertions: number;
    deletions: number;
  }>;
  runText(args: string[], opts?: { cwd?: string }): Promise<{
    stdout: string;
    stderr: string;
    exitCode: number;          // surfaced as data – does NOT throw
  }>;
};

preg: {
  match(pattern: string, subject: string): {
    match: string;
    groups: string[];        // empty string for an optional group that
    didn't match
    index: number;
  } | null;
  matchAll(pattern: string, subject: string): Array<{
    match: string;
    groups: string[];
    index: number;
  }>;
  replace(pattern: string, replacement: string, subject: string): string;
};

preg2: {                                // PCRE engine
  (lookaround, backrefs)
  match(pattern: string, subject: string): {
    match: string;
    groups: string[];
    index: number;
  } | null;
  matchAll(pattern: string, subject: string): Array<{
    match: string;
    groups: string[];
    index: number;
  }>;
  replace(pattern: string, replacement: string, subject: string): string;
};

```

```

jwt: {
  sign(
    claims: Record<string, unknown>,
    secret: string,                                // raw bytes for HS*; PEM-
encoded private key for RS*/PS*/ES*/EdDSA
    opts?: {
      algorithm?:
        | "HS256" | "HS384" | "HS512"            // HMAC
        | "RS256" | "RS384" | "RS512"            // RSA-PKCS1v1.5
        | "PS256" | "PS384" | "PS512"            // RSA-PSS
        | "ES256" | "ES384" | "ES512"            // ECDSA
        | "EdDSA";                                // Ed25519
    },                                              // default "HS256"
  ): string;
  view(token: string): {
    header: Record<string, unknown>;
    payload: Record<string, unknown>;
    signature: string;                            // raw base64url, no
padding
  };
  validate(
    token: string,
    secret: string,                                // raw bytes for HS*; PEM-
encoded public key (or certificate) for asymmetric
    opts?: {
      algorithm?:
        | "HS256" | "HS384" | "HS512"
        | "RS256" | "RS384" | "RS512"
        | "PS256" | "PS384" | "PS512"
        | "ES256" | "ES384" | "ES512"
        | "EdDSA";
      audience?: string;
      issuer?: string;
    },
  ): { valid: true; claims: Record<string, unknown> }
    | { valid: false; reason: string };
};

encrypt: {
  keygen(): { publicKey: string; privateKey: string };
// age X25519
  keygenPgp(opts?: { name?: string; email?: string }):
// PGP keypair (armored blocks)
  { publicKey: string; privateKey: string };
  encrypt(
    data: string | Uint8Array | ArrayBuffer,
    recipients: string | string[],                // age1... public keys
    opts?: { armored?: boolean },                // default false → binary age
format

```

```

    ): Uint8Array;
    decrypt(
        ciphertext: string | Uint8Array | ArrayBuffer,
        identities: string | string[],           // AGE-SECRET-KEY-1...
private keys
    ): Uint8Array;                             // auto-detects binary vs
armored
    rekey(
        ciphertext: string | Uint8Array | ArrayBuffer,
        oldIdentities: string | string[],
        newRecipients: string | string[],
        opts?: { armored?: boolean },           // default: preserve input
format
    ): Uint8Array;
    detectBackend(input: string):
        | { backend: "age";      kind: "public" | "private" }
        | { backend: "pgp";      kind: "public" | "private" }
        | { backend: "unknown" };
    };

    sqlite: {
        open(path: string): Promise<{           // ":memory:" or a
filesystem path
            exec(sql: string, ...params: unknown[]): Promise<{
                rowsAffected: number;
                lastInsertId: number;
            }>;
            query(sql: string, ...params: unknown[]):
Promise<Array<Record<string, unknown>>>;
            queryValue(sql: string, ...params: unknown[]): Promise<unknown>; //
first column of first row, or null
            begin(): Promise<{                   // transaction handle
                exec(sql: string, ...params: unknown[]): Promise<{ rowsAffected:
number; lastInsertId: number }>;
                query(sql: string, ...params: unknown[]):
Promise<Array<Record<string, unknown>>>;
                queryValue(sql: string, ...params: unknown[]): Promise<unknown>;
                commit(): Promise<void>;
                rollback(): Promise<void>;
            }>;
            prepare(sql: string): Promise<{      // compiled-once statement
handle
                exec(...params: unknown[]): Promise<{ rowsAffected: number;
lastInsertId: number }>;
                query(...params: unknown[]): Promise<Array<Record<string,
unknown>>>;
                queryValue(...params: unknown[]): Promise<unknown>;
                close(): Promise<void>;
            }>;
    };

```

```

        close(): Promise<void>;
    }>;
};

gh: {
    authStatus(): Promise<{
        authenticated: boolean;
        user: string;           // login or empty
        raw: string;            // login on success, stderr / probe note
on failure
    }>;
    prList(opts?: {
        cwd?: string;
        state?: "open" | "closed" | "merged" | "all";
        limit?: number;         // default 30
        author?: string;
    }): Promise<Array<{
        number: number;
        title: string;
        state: string;
        author: string;         // login (flattened from gh's wrapper)
        headRefName: string;
        baseRefName: string;
        url: string;
        createdAt: string;
        updatedAt: string;
    }>>;
    repoView(
        repo?: string,          // "owner/name"; omit to use cwd's repo
        opts?: { cwd?: string },
    ): Promise<{
        name: string;
        owner: string;          // login (flattened)
        description: string;
        url: string;
        defaultBranch: string;   // flattened from defaultBranchRef.name
        visibility: string;
    }>;
};

email: {
    spf(domain: string): Promise<
        | { present: false }
        | {
            present: true;
            record: string;
            mechanisms: string[];
            allPolicy: "pass" | "fail" | "softfail" | "neutral" | "";
        }
    >

```



```

>;
dmARC(domain: string): Promise<
  | { present: false }
  | {
    present: true;
    record: string;
    tags: Record<string, string>;
    policy?: string;
    subdomain?: string;
    percent?: string;
    rua?: string;
    ruf?: string;
  }
>;
mtaSts(domain: string): Promise<
  | { present: false }
  | {
    present: true;
    record: string;
    txt: { v: string; id: string };
    policy?: {
      version?: string;
      mode?: "enforce" | "testing" | "none" | string;
      mx?: string[];
      maxAge?: number | string;
    };
    policyError?: string;
  }
>;
tlsRpt(domain: string): Promise<
  | { present: false }
  | {
    present: true;
    record: string;
    tags: Record<string, string>;
    rua?: string;
  }
>;
bimi(domain: string, opts?: { selector?: string }): Promise<
  | { present: false; selector: string }
  | {
    present: true;
    selector: string;
    record: string;
    tags: Record<string, string>;
    l?: string;
    a?: string;
  }
>;

```

```

    all(domain: string): Promise<{
      domain: string;
      spf:    unknown;
      dmarc:  unknown;
      mtaSts: unknown;
      tlsRpt: unknown;
      bimi:   unknown;
    }>;
  };
};
};

```

Examples:

```

api.runtime.log("hello", 1 + 2);
api.runtime.assert.equal(api.runtime.time.nowMs() > 0, true);

const r = await api.net.http.get("https://example.com");
api.runtime.assert.ok(r.status === 200, `expected 200, got ${r.status}`);

await api.runtime.time.sleep(50);
const home = api.runtime.env.get("HOME") ?? "(none)";

api.runtime.log(api.crypto.hash.sha256("abc"));
// → ba7816bf8f01cfea414140de5dae2223b00361a396177a9cb410ff61f20015ad

```

`api.net.http.request(method, url, opts?)` is the full-featured client beyond `get` / `post`: per-call headers, body, timeout (default 30s), retry (re-attempts on transport errors and 5xx — never 4xx, which are deterministic; linear backoff capped at 1s), follow (redirect control; default true), and basic-auth username / password. The result is `{ status, ok, headers, body, url }` — `ok` is status in [200,400), `headers` is a lower-cased name→value map, `url` is the final URL after redirects. HTTP 4xx / 5xx don't throw (surfaced via `status` / `ok`); transport errors and context deadline do.

HTTP bindings use `net/http` with a 5-second default per-request timeout and surface real `Promise<...>` values through the event loop. They are *not* mockable from JS — they go to the real network.

Compression bindings (`api.format.compression.*`) cover nine pure-Go algorithms behind a uniform interface. Inputs are either strings (interpreted as UTF-8 byte sequences) or any `ArrayBuffer` / `Uint8Array`; outputs are `Uint8Array` — goja's representation of a Go `[]byte` return is a `Uint8Array`, so scripts can read `.length` and iterate directly without an `ArrayBuffer` view. Every algorithm round-trips:

Algorithm	Library
gzip / deflate / zlib	stdlib compress/*

Algorithm	Library
bzip2	stdlib compress/bzip2 for read; github.com/dsnet/compress/bzip2 for write
zstd	github.com/klauspost/compress/zstd
brotli	github.com/andybalholm/brotli
lz4	github.com/pierrec/lz4/v4
xz	github.com/ulikunitz/xz
snappy	github.com/golang/snappy

`api.format.compression.algos()` returns the supported names so scripts can iterate without hard-coding the list. Unknown algorithm names throw a clear error rather than silently returning empty bytes.

Barcode encoders (`api.format.barcode.*`) cover ten symbologies through one `encode(format, data, opts?)` entry point, all backed by the pure-Go github.com/boombuler/barcode toolkit. Output is a PNG payload as `Uint8Array`; size defaults to 256x256 for 2D codes (QR / DataMatrix / Aztec) and 400x120 for linear ones, overridable via `opts.width / opts.height`. `opts.quietZone` (`true` for a default margin, or a pixel count) pads the rendered bars with a white border — required for EAN / UPC to decode (see the decode quirks below).

- 2D: `qr` (medium error correction, auto mode), `datamatrix`, `aztec` (33% ECC, auto layers), `pdf417` (security level 5).
- Linear: `code128`, `code39` (Mod-43 checksum, ASCII-only), `codabar`.
- Retail: `ean13`, `ean8`, `upca` — content must be numeric and the right length; `boombuler/ean.Encode` dispatches by length so all three variants share the same encoder.

`api.format.barcode.formats()` returns the supported encode names;

`api.format.barcode.decodableFormats()` returns the (slightly different) decode set.

Decoding (`api.format.barcode.decode(data, format?)`) accepts PNG / JPEG / WebP image bytes and reaches into github.com/makiuchi-d/gozxing. Returns `{ format, text }` — `format` is the snake-case label the encoder accepts, so round-trips line up. WebP support comes via a blank-import of `golang.org/x/image/webp` that registers the decoder with `stdlib's image.Decode`.

- **With a format hint:** only that reader runs. Fast, and a mismatch surfaces a clear error rather than a silent miss.
- **Without a hint:** every reader in a priority list runs (2D formats first; the more-constrained 1D readers — UPC/EAN — before the permissive ones — Code 128 / Codabar). First success wins; if no reader matches, the call throws with "no barcode recognised".

Decode quirks worth knowing — these are properties of the underlying libraries, not sercon bugs:

- **pdf417 isn't decodable.** gozxing v0.1.1 has no PDF417 reader; pass it as a hint and the call throws "unsupported format". The encoder still emits PDF417 just fine.
- **EAN / UPC need a quiet zone.** Per ISO/IEC 15420 + 15418, retail symbologies require white margin on both sides of the bars. boombuler's encoder emits them edge-to-edge, so a bare `barcode.encode("ean13", x) → barcode.decode(...)` round-trip fails on the decode step. Pass `opts.quietZone` to encode to fix it: `true` adds a default margin (10% of the width, floored at 10px), or pass a pixel count for an explicit margin. Real-world EAN scanners need that margin too, so it's good practice for any EAN/UPC you'll actually print.
- **Code 39 decoded text ends with the Mod-43 checksum char.** boombuler's encoder appends it; gozxing returns it.
- **Codabar wraps payload in start/stop chars** like `A...A` on encode; gozxing strips them on decode.

Charset bindings (`api.text.*`) cover detection plus byte-side round-tripping:

- **detect(data)** — runs `saintfish/chardet` over the input and returns the top guess (`charset`, `confidence` on a 0–100 scale, optional `language` hint) plus the full candidate list. Input is bytes (string, `ArrayBuffer`, or `Uint8Array`).
- **encode(text, charset)** — converts a UTF-8 string to bytes in the target encoding. Characters with no representation in the target encoding cause the encoder to error out rather than silently lose them; pre-process the input yourself if you want lossy behaviour.
- **decode(data, charset)** — the inverse: bytes-in-charset → UTF-8 string.

Charset names follow the WHATWG Encoding Living Standard aliases that

`golang.org/x/text/encoding/htmlindex.Get` accepts — `UTF-8`, `ISO-8859-1`, `Windows-1252`, `Shift_JIS`, `GBK`, `GB18030`, `Big5`, `EUC-JP`, `EUC-KR`, etc., plus all documented aliases. Unknown names throw a clear error.

Check-digit bindings (`api.format.checkdigit.*`) verify and compute the trailing check digit of common numeric codes. All three members are synchronous — the algorithms are small bits of modular arithmetic, no I/O or library dependency. Supported algorithms (returned by `algorithms()`):

Algorithm	Input length	Notes
<code>luhn</code>	any ≥ 2	Credit cards, IMEI, social-security-style IDs. Mod-10 with right-to-left doubling.
<code>isbn10</code>	10	Last position may be <code>X</code> (= 10). Mod-11 with weighted sum.
<code>isbn13</code>	13	Alias of <code>ean13</code> — same algorithm; the 978 / 979 prefix is the only thing that distinguishes ISBN-13 syntactically.

Algorithm	Input length	Notes
ean13	13	Mod-10, weights 1, 3, 1, 3, ... from position 0.
ean8	8	Mod-10, weights 3, 1, 3, 1, ...
upca	12	Mod-10, weights 3, 1, 3, 1, ...

`validate(algo, input)` returns a plain boolean — non-digit characters, wrong length, and unknown algorithm names all surface as `false` so scripts can run quick presence checks without try/catch. `compute(algo, partial)` takes the input *without* its check digit and returns just that digit (or "X" for ISBN-10 position 10). `inspect(algo, input)` returns the union of both views — useful when you want to display the diagnosis to a human.

Archive bindings (`api.fs.archive.*`) handle the three formats `stdlib` provides without external deps: `.zip`, `.tar`, and `.tar.gz` (also `.tgz`). Format is inferred from the destination's extension on `create` and from the source path on `extract`.

- **`create(destPath, sources)`** — `sources` is an array of either bare paths (the basename is used as the in-archive name) or `{ path, name }` objects (override the in-archive name). Directory sources are recursed; the directory's basename becomes the archive subdir and the rest of the tree is added relative to it. Returns the count + list of entries and the output file size.
- **`extract(archivePath, destDir, opts?)`** — unpacks into `destDir`. `opts.overwrite` (default `false`) controls whether existing files at the destination are clobbered; on `false` a collision errors out (`os.ErrExist` -shaped). Every entry name is run through a zip-slip / tar-slip guard that rejects absolute paths, `..` segments, and anything that would resolve outside `destDir`.

Diff bindings (`api.text.diff.compare(a, b, opts?)`) produce a unified diff between two text inputs via github.com/pmezard/go-difflib. Inputs are strings (UTF-8) or any byte sequence (`ArrayBuffer` / `Uint8Array`). The result reports:

- `identical` — byte-equal inputs short-circuit; `diff` is empty.
- `binary` — either input has a NUL byte in its first 8 KB. A unified diff of binary content isn't useful so `diff` is empty.
- `added` / `removed` — body-only `+` / `-` line counts (file headers are excluded so the numbers match `git diff --shortstat`).
- `diff` — the unified diff text, or empty when `identical` or `binary` is `true`.
- `format` — always "unified" today; reserved for future alternatives (`"side-by-side"`, `"summary"`, ...).

`opts.context` (default 3) is the context-line count; `opts.fromFile` / `opts.toFile` (default `"a"` / `"b"`) set the labels in the diff headers.

JSON-query bindings (`api.text.jq.*`) run jq filters over JS data via github.com/itchyny/gojq, a pure-Go re-implementation. The filter syntax is the same as

command-line jq.

- **query(data, filter)** — runs the filter and returns the first emitted value (or `null` when the filter emits nothing).
- **queryAll(data, filter)** — drains the result iterator and returns every emitted value as an array.

Data passes through goja's `Export` as `map[string]any / []any` trees, which gojq's runtime accepts directly. There's a small normalisation pass on the way in that converts all sized integer types (`int64`, `int32`, etc.) to plain `int` and `float32` to `float64` — gojq's arithmetic dispatch only knows the two canonical numeric kinds, and goja exports every JS-side integer as `int64`. Without normalisation a query as simple as `.users[].age` would panic in gojq.

Parse errors and runtime errors (type mismatches, division by zero, ...) both surface as JS exceptions. Use jq's optional access operator `?` (e.g. `.does.not.exist?`) to suppress missing-path errors and get `null` back instead.

Aggregate connectivity (`api.net.netstatus.check(host, opts?)`) is an orchestration layer over the lower-level probes — it runs DNS, TCP, TLS, and HTTP against one host **concurrently** (fan-out via a `WaitGroup`) and folds them into a single status object. Each sub-probe carries its own `ok` flag and, on failure, an `error` string — individual failures are data, not a thrown error, so the result is always a complete snapshot. `reachable` is the AND of `dns.ok` and `tcp.ok` (name resolves + connection opens); TLS / HTTP are reported but don't gate it (a plain-HTTP host or one with an expired cert is still "reachable"). Only a missing host argument throws. No new library — it composes `net / crypto/tls / net/http` directly.

Redis (`api.db.redis.open(url)`) is a stateful-handle binding over [redis/go-redis/v9](#) (the official client). `open` parses a `redis://[:password@]host:port/db` URL and PINGS to surface a bad address up front, then resolves to `{ do, ping, close }`. `do(command, ...args)` runs an arbitrary RESP command — `do("SET", "k", "v")`, `do("LRANGE", "l", "0", "-1")` — so the binding stays small instead of mirroring hundreds of methods. Replies map the obvious way (bulk strings → string, integers → number, arrays → array); a nil reply (missing key) becomes JS `null` rather than throwing, so `await r.do("GET", "absent")` is `null`. Redis-level errors (WRONGTYPE, unknown command) throw. Offline-testable via `alicebob/miniredis`.

Memcached (`api.db.memcached.open(addr)`) is a stateful-handle binding over [bradfitz/gomemcache](#) (the de facto standard client). `addr` is `host:port`; `gomemcache` pools connections lazily so there's no PING-on-open (the first op surfaces a bad address) and no close. Resolves to `{ get, set, delete }`: `get` returns the stored string or `null` on a cache miss; `set(key, value, expirySeconds?)` stores bytes (0 / omitted = never expire); `delete` returns `true` if the key existed, `false` on a miss. Server errors throw.

LDAP (`api.db.ldap.open(url)`) is a stateful-handle binding over [go-ldap/v3](#). Dials `ldap://host:port` (or `ldaps://`), does an anonymous bind by default (or a simple bind with `opts.bindDN / opts.password`), and resolves to `{ rootDSE, search, close }`. `rootDSE()` reads the server's anonymous metadata entry (naming contexts, supported

controls, vendor); `search(baseDN, filter, attrs?)` runs a subtree search and returns `{ dn, <attr>: [values] }` per entry (attributes stay arrays since LDAP is multi-valued). A read / inspection surface — no modify / add / delete.

DICT (`api.db.dict.*`) is an RFC 2229 dictionary-server client. No maintained pure-Go DICT library exists, so the protocol is hand-rolled over `net/textproto` (a simple line-based status-code protocol, much like SMTP). `define(host, word, opts?)` returns `{ word, found, definitions: [{ db, dbName, text }] }`; a word with no entry resolves `found: false` (not an error). `match(host, word, opts?)` returns words matching under `opts.strategy` (default `prefix`). Both are one-shot (connect → query → QUIT); `opts.database` (default `*` = all), `opts.port` (default 2628).

AI agents (`api.tools.ai.*`) shell out to a coding-assistant CLI. `providers()` lists which of `claude` / `codex` / `copilot` / `gemini` are on PATH (preference order); `send(opts)` runs a one-shot prompt through one of them — `opts.provider` picks explicitly, else the first detected. `opts.system` / `opts.context` are prepended to the prompt (the portable common denominator across CLIs with different system-prompt flags). The `argv` builder is a pure function (`buildAIArgv`) so the provider invocations are unit-testable without the CLIs installed. Returns `{ provider, output, exitCode }`; a non-zero exit is surfaced as data (the model CLI may exit non-zero with a useful message), while a missing provider and context deadline throw. This is the options-object shape rather than a builder chain — the idiomatic JS equivalent, and it sidesteps threading a mutable builder handle through `goja`. **Library:** `os/exec` (`stdlib`).

Browser sessions (`api.net.browser.open()`) give a stateful HTTP client — an automatic cookie jar (`net/http/cookiejar` with the `golang.org/x/net/publicsuffix` list, so cookies scope correctly across subdomains) plus default headers replayed on every request. A second stateful-handle binding in the same shape as `api.db.sqlite: open()` resolves to `{ setUserAgent, setHeader, get, post, cookies }`. A login POST followed by a GET replays the session cookie without the script touching it; `cookies(url)` inspects the jar (handy for asserting a login set what you expect). `get` / `post` return the same `{ status, ok, headers, body, url }` shape as `api.net.http.request`. No explicit close — the session is GC'd with the handle.

Subprocess bindings (`api.tools.exec.*`) wrap Go's `os/exec`. `shell` is the only entry today; `http` (recon-with-curl-fallback) and `git` / `gh` wrappers ride on top of it in later 0.4.x cuts.

- **`shell(cmd, opts?)`** — runs a subprocess and waits for it to exit. `cmd` is either a **string** (passed verbatim to `/bin/sh -c` on Unix or `cmd /C` on Windows so pipes, redirects, and globs work as typed) or a **string[]** (treated as `argv`, no shell — use this when arguments could be re-interpreted by the shell). Returns `{ stdout, stderr, exitCode, success, durationMs }`.

Options: `cwd` sets the working directory; `env` is merged on top of the parent process environment (not a replacement); `timeout` is in milliseconds and defaults to **30 000 ms**;

`stdin` is piped to the child's standard input.

Process-start failures (host binary not on PATH, permission denied) throw. Context deadline (`timeout`) and engine cancellation throw too — the subprocess never got to choose its own exit code, so it isn't reasonable to surface a fake one. Non-zero exits do **not** throw: `success: false` + the real `exitCode` is what callers want for the usual "ran the linter, expected to be told it failed" flow.

- **`http(method, url, opts?)`** — curl-compatible HTTP client routed through `recon` (preferred) with `curl` as a fallback. Returns `{ status, headers, body, durationMs, backend }`. The choice between the two binaries is normally invisible to scripts — both implement the same curl-style flag set we use here (`-X`, `-H`, `-D`, `--data-binary`, `-L`, `-k`, `-s`). Pick a specific backend with `opts.backend: "recon" | "curl"` when you need to.

Implementation notes: response headers are dumped to a temp file via `-D <path>` and parsed back, rather than relying on `-i`'s body-stream format (recon's `-i` is verbose-debug style with `<` prefixes, incompatible with curl's wire-format `-i`). Request bodies are materialised to a temp file and passed via `--data-binary @<path>` so CR / LF are preserved regardless of backend. Header names in the returned map are lower-cased. On redirect chains (`opts.follow: true`), the *last* response wins — the intermediate 3xx blocks are skipped.

HTTP 4xx / 5xx do **not** throw — they're a normal HTTP outcome and callers branch on `status`. Process-start failures, transport errors (DNS, connection refused, TLS handshake), and context deadline / cancellation throw. The error message includes the backend's `stderr` when present.

Git bindings (`api.tools.git.*`) shell out to the host `git` binary; no pure-Go alternative (`go-git` is heavier and would drift from the user's installed git). Every binding accepts an `opts.cwd` so a single engine can work across multiple checkouts.

- **`branch(opts?)`** — Reports the current branch (empty string when HEAD is detached) plus every local branch from `for-each-ref refs/heads`. `detached` is true exactly when `symbolic-ref --short HEAD` exited non-zero, which is git's own signal for the state.
- **`isClean(opts?)`** — Convenience boolean over `git status --porcelain`. True when the porcelain output is empty.
- **`revParse(rev, opts?)`** — Returns the 40-char SHA for `rev`. Invalid refs throw with git's `stderr` message included.
- **`status(opts?)`** — Parses `--porcelain v1` output into `{ path, indexStatus, workingStatus }`. The two single-char status codes are the raw porcelain codes (`M`, `A`, `D`, `R`, `?`, ...) so scripts can dispatch on them without re-parsing strings.
- **`add(paths, opts?)`** — Stages one path (string) or a list (string[]). `--` is inserted between `add` and the paths so values that look like flags work too. Returns the resolved paths array.

- `commit(message, opts?)` — Creates a new commit; the post-commit HEAD SHA is returned. `allowEmpty:true` toggles `--allow-empty`. An empty message throws before spawning.
- `log(opts?)` — Returns `limit` (default 50) most-recent commits in `revRange` (default HEAD). Format string `%H%x09%h%x09%an%x09%ae%x09%at%x09%s` tab-separates SHA, short SHA, author name, email, unix timestamp, and subject. Newlines and tabs in subjects are already collapsed by git so the parse stays one-line-per-commit.
- `diffStat(opts?)` — Aggregates `git diff --shortstat` into `{ files, insertions, deletions }`. Default range is `HEAD~1..HEAD`. A pure-add or pure-delete diff returns zero for the missing side instead of throwing.
- `runText(args, opts?)` — Escape hatch. Surfaces `{ stdout, stderr, exitCode }` for any `git <args>` invocation. Non-zero exits do **not** throw — that's the whole point of having a generic wrapper; callers branch on `exitCode`. Spawn failures and context cancellation still throw.

GitHub-CLI bindings (`api.tools.gh.*`) wrap the `gh` binary. They respect whatever authentication state `gh auth` is already in; we don't try to swap accounts or manage tokens. Every call uses `gh --json` so the result is structured rather than parsed out of human-readable text.

- `authStatus()` — Probe whether `gh` is installed *and* authenticated. Under the hood this runs `gh api user --jq .login` (machine-friendly: just the login on success, a clear error otherwise) rather than `gh auth status` (multi-line human report). Missing-`gh` and unauthenticated-session both resolve with `{ authenticated: false, ... }` rather than throwing — the whole point of a status probe is that the script branches on it. Only context cancellation throws.
- `prList(opts?)` — Lists pull requests on the repo identified by `opts.cwd` (or the engine's working directory). Defaults: open state, limit 30. `gh`'s `author: { login, ... }` wrapper is flattened to just the login string so scripts can compare directly.
- `repoView(repo?, opts?)` — Returns metadata about a repo. With no argument it asks `gh` about the `cwd`'s repo (works from inside a checkout); pass `"owner/name"` to look up any repo `gh` has access to. Two convenience flattenings are applied: `owner` is a login string (not `{ login, ... }`), and `defaultBranch` is the branch name (not `defaultBranchRef.name`). Empty repos resolve with `defaultBranch: ""` rather than `undefined`.

Regex bindings (`api.text.preg.*`) accept PHP-style `/pattern/flags` delimited input but run on Go's `stdlib regexp` (RE2). The "preg" naming is a deliberate homage; the semantics are RE2's. That means **no backreferences inside patterns, no lookahead / lookbehind, and no possessive quantifiers** — scripts that need those should reach for goja's native `RegExp` (ECMAScript-flavoured) or wait for a dedicated PCRE-compatible binding (`regexp2` is on the backlog).

Supported flags: `i` (case-insensitive), `m` (multiline — `^` and `$` match line boundaries), `s` (dotall — `.` matches newline). The PHP flags `u` (Unicode), `U` (default-ungreedy), and `x` (extended whitespace-ignoring) all surface a clear error rather than silently dropping; unknown

flags do too. RE2 is UTF-8 by default, so `u` is unnecessary in any case — the error message says so.

- **`match(pattern, subject)`** — Returns `{ match, groups, index }` for the first hit, or `null` when nothing matches. `groups` is every numbered submatch after group 0, in order; an optional group that didn't match surfaces as an empty string (not `undefined`) so the result type stays stable.
- **`matchAll(pattern, subject)`** — Same shape, drained: returns an array of match objects, one per hit.
- **`replace(pattern, replacement, subject)`** — Substitutes via Go's `$1 / ${1}` backreference syntax. PHP's `\1` form is **not** translated — that would conflict with `\\` escapes that are already legitimate in the replacement. Use `$1 / ${1}` as in Go.

PCRE regex bindings (`api.text.preg2.*`) are the heavier-engine sibling of `api.text.preg`. They run on github.com/dlclark/regexp2, a port of .NET's regex engine, which **does** support lookahead, lookbehind, backreferences inside patterns, and possessive quantifiers — everything RE2 (and therefore `api.text.preg`) can't do. The API is identical (`match / matchAll / replace`, the same `/pattern/flags` delimited syntax, the same `{ match, groups, index }` shape) so switching engines is a one-word change. The flag set adds `x` (ignore-pattern-whitespace), which RE2 couldn't honour; `u / U` still error (regexp2 is Unicode-aware by default and has no global-ungreedy switch). `replace` uses .NET / regexp2 `$1 / ${1}` substitution syntax.

The trade-off is the reason `api.text.preg` exists at all: regexp2 **backtracks**, so it has no linear-time guarantee. A pathological pattern against adversarial input can blow up exponentially. Use `api.text.preg` (RE2) by default; reach for `api.text.preg2` only when you need a PCRE feature, and keep a timeout around untrusted input.

JWT bindings (`api.crypto.jwt.*`) wrap [golang-jwt/v5](https://github.com/dgrijalva/jwt-go). The full RFC 7518 algorithm matrix is supported: HMAC (`HS256 / HS384 / HS512`), RSA-PKCS1v1.5 (`RS256 / RS384 / RS512`), RSA-PSS (`PS256 / PS384 / PS512`), ECDSA (`ES256 / ES384 / ES512`), and Ed25519 (`EdDSA`). Unsupported algorithm names — including the special `"none"` value, garbage like `"HS999"`, and protocol names that aren't JWT signing algos — throw at `sign / validate` time with a named-algorithm error rather than silently falling through.

Key shape. The `secret` parameter takes three forms, picked by shape: HMAC algorithms use the byte string directly; asymmetric algorithms expect a PEM-encoded key (private for `sign`, public — or a certificate the public key can be pulled from — for `validate`); and **any algorithm** accepts a JWK (JSON Web Key) object — a JSON string carrying a `"kty"` member. JWK is detected by the leading `{ + "kty"` and parsed via [lestrrat-go/jwx/v2/jwk](https://github.com/lestrrat-go/jwx/v2/jwk); the `kty` (`oct` → HMAC bytes, `RSA`, `EC`, `OKP`) determines the key type, so a JWK works with whatever matching `opts.algorithm` you pass. Useful when keys come from a JWKS endpoint or a config file in JWK form. PEM detection is the literal `-----BEGIN` prefix. Both directions of the PEM/bytes cross-check are enforced (JWK input bypasses them — the `kty` is authoritative):

- **PEM secret + HMAC algorithm** throws with a named-algorithm hint ("looks like a PEM-encoded key — set `opts.algorithm` to RS256 / ES256 / EdDSA / etc."). Without this guard, a script that forgets to set `opts.algorithm` after switching to a PEM key would silently sign an HS256 token using the PEM bytes themselves — that's exactly the algorithm-confusion attack pattern in reverse.
- **Plain bytes + asymmetric algorithm** throws with the inverse hint ("needs a PEM-encoded private/public key but secret is plain bytes"). Without this guard the user would get a cryptic "x509: malformed certificate" deep inside `jwt-go`.

The cross-check fires at the binding boundary before `jwt-go` is called, so it's a thrown error rather than a `valid: false` resolution — these are structural input problems, not validation failures.

- **`sign(claims, secret, opts?)`** — Produces a compact-serialised signed JWT. `claims` is passed straight through to `jwt-go`'s `MapClaims`, which recognises the RFC 7519 reserved claims (`exp`, `nbfi`, `iat`, `aud`, `iss`, `sub`, `jti`) and lets you layer arbitrary application claims alongside them. Missing reserved claims aren't synthesised — scripts that want `iat` set should compute it explicitly (e.g. `Math.floor(api.runtime.time.nowMs() / 1000)`).
- **`view(token)`** — Decodes the header + payload **without verifying the signature**. Useful for debugging auth flows or surfacing `aud` / `iss` to the user before deciding whether to trust the token. `signature` comes back as the raw base64url-encoded segment (no padding); decode it yourself if you need the bytes. Malformed input — wrong segment count, invalid base64, invalid JSON — throws.
- **`validate(token, secret, opts?)`** — Verifies the signature using `secret` and the algorithm declared in the token's header. Standard-claim validation (`exp`, `nbfi`, `iat`) is delegated to `jwt-go`. `opts.audience` / `opts.issuer` are cross-checked against `aud` / `iss` when set. The contract is **resolve, don't throw, on validation failure**: `{ valid: false, reason: "..." }` for bad signature, expired, audience mismatch, issuer mismatch, *and* algorithm mismatch when `opts.algorithm` is set. Only structural input errors (not a JWT at all, empty secret, empty token, PEM/algorithm cross-check failure) throw — pattern-matching on `valid: false` shouldn't accidentally accept a garbage string.

Setting `opts.algorithm` is the **algorithm-confusion guard**: an attacker who has the public key bytes can forge an HS256 token by using those bytes as the HMAC secret (the classic JWT exploit). When `opts.algorithm` is set, the parser's `WithValidMethods` whitelist refuses any token whose header declares a different algorithm. Always set `opts.algorithm` when validating asymmetric tokens; the binding doesn't infer it from the key shape because that would make the wrong default silent.

Encryption (`api.crypto.encrypt.*`) supports **two backends** behind one API: age (filippo.io/age, the default) and PGP ([ProtonMail/go-crypto/openpgp](https://protonmail.com/go-crypto/openpgp)). `encrypt` / `decrypt` auto-dispatch on the key/ciphertext format — the same classifier `detectBackend`

exposes — so scripts use one API regardless of backend. The age round-trip core (keygen / encrypt / decrypt) landed in v0.5.5; armoured output (`opts.armored`) in v0.5.6; rekey in v0.5.7; detectBackend in v0.5.8; the PGP backend in v0.5.16.

age identities use the X25519 flavour: `age1...` recipients (safe to share) and `AGE-SECRET-KEY-1...` identities (kept secret). PGP keys are ASCII-armored blocks (`-----BEGIN PGP PUBLIC/PRIVATE KEY BLOCK-----`), generated by `keygenPgp` .

- **keygenPgp(opts?)** — Generate a PGP keypair (RSA 2048). `opts.name` / `opts.email` populate the primary user ID; both optional. Returns armored { `publicKey`, `privateKey` } blocks. When `encrypt` sees a PGP public-key block as a recipient (or `decrypt` sees a PGP private-key identity / a `-----BEGIN PGP MESSAGE-----` ciphertext), it routes through `openpgp` instead of `age`. PGP output is always armored; `opts.armored` is ignored on that path. `age` and PGP recipient sets can't be mixed in one call — the formats are incompatible.

All three members are synchronous: encryption is pure CPU work with a small API surface, matching the call shape of the other crypto bindings (`api.crypto.jwt` , `api.crypto.hash`).

- **keygen()** — Generate a fresh X25519 identity. Returns { `publicKey`, `privateKey` } as the bech32 strings `age` writes to disk. Two consecutive calls produce different keys; entropy comes from `crypto/rand` via `age`. Destructure to use one or the other separately.
- **encrypt(data, recipients, opts?)** — Seal `data` to one or more recipients. Input accepts `string` / `Uint8Array` / `ArrayBuffer`; `recipients` accepts a single string or an array. Default output is the binary age format as `Uint8Array` . Passing `opts.armored: true` wraps it in `age`'s ASCII armor — the `-----BEGIN AGE ENCRYPTED FILE-----` banner + base64 body that's safe to embed in JSON / YAML / email. Either form decrypts via the same `decrypt` call (auto-detected by the leading bytes). Multi-recipient encryption uses `age`'s native multi-recipient header — any one of the listed identities can decrypt the resulting ciphertext, no re-encryption per reader needed.
- **decrypt(ciphertext, identities)** — Open an `age`-encrypted payload with one of the supplied identities. **Auto-detects binary vs armored input** by sniffing the leading bytes for `age`'s armor banner (whitespace and BOM bytes ahead of the banner are tolerated, so ciphertext pasted from JSON / email containers still detects correctly). `age` then walks the identities and uses the first that matches a stanza in the header. Returns `Uint8Array` ; let scripts decode to a string via `api.text.charset.decode(bytes, "utf-8")` if appropriate (goja doesn't ship `TextDecoder`).
- **rekey(ciphertext, oldIdentities, newRecipients, opts?)** — Re-encrypt a payload for a fresh recipient set without ever exposing the plaintext to the caller. The `decrypt+encrypt` loop is internal; the plaintext lives only in this function's stack until the `encrypt` step finishes. Output format defaults to **match the input** — armored in / armored out, binary in / binary out — which is what you almost always want when key-rotating a payload that lives in a fixed location (file, vault row, JSON field). `opts.armored: true` | `false` forces the output regardless of input. Internally consistent with `encrypt` and

decrypt : same cross-checks (private as recipient, public as identity), same auto-detect on the input side, same multi-recipient handling on the output side.

- **detectBackend(input)** — Classify a recipient or identity string by encryption backend without parsing or I/O. Returns { backend: "age" | "pgp" | "unknown", kind?: "public" | "private" }. Age recognises three input forms: bech32 X25519 (age1..., classified kind: "public"), AGE-SECRET-KEY-1... (kind: "private"), and SSH public keys (ssh-rsa / ssh-ed25519, classified kind: "public" since age accepts them as recipients). PGP recognises armored block markers (-----BEGIN PGP PUBLIC KEY BLOCK----- / -----BEGIN PGP PRIVATE KEY BLOCK-----). Anything else returns { backend: "unknown" } — false-negative on unfamiliar input is the safer default than guessing. v0.5.8 ships the classifier only; sercon's actual encrypt / decrypt paths still handle age recipients only. Useful for scripts that need to dispatch on the format ("is this an age recipient or a PGP block?") before deciding which encrypt path — e.g., a config loader that routes some recipients through api.crypto.encrypt.encrypt and shells out to gpg --encrypt for others.

Cross-checks fire at the binding boundary so the common JS-side mistakes throw with named-key hints rather than cryptic bech32 errors deep inside age:

- **Private key as recipient** ("looks like a private key (AGE-SECRET-KEY-...); recipients are public keys (age1...)") — catches a script that mixed up the two halves on the encrypt call.
- **Public key as identity** ("looks like a public key (age1...); identities are private keys (AGE-SECRET-KEY-1...)") — inverse guard on the decrypt call.

Decryption with an identity that wasn't on the recipient list throws with age's own wording forwarded through: "identity did not match any of the recipients: incorrect identity for recipient block". Scripts that want to silently try multiple identities should pass them all in one call (age tries each in turn internally) rather than catching this error.

SQLite (api.db.sqlite.*) is sercon's first **stateful-handle** binding — the shape future protocol bindings (redis / ldap / ...) will reuse. The namespace exposes only open ; the object it resolves to carries the real surface. The driver is modernc.org/sqlite, the pure-Go SQLite translation — no cgo, so the project's cgo-free rule is preserved.

- **open(path)** — Connect to a database. ":memory:" is an in-RAM database that evaporates when the handle is closed; any other string is a filesystem path (created if absent). The connection is Ping -ed before open resolves, so a bad path surfaces here rather than at the first query. Resolves to a handle object: { exec, query, queryValue, close }.
- **exec(sql, ...params)** — Run a non-query statement (DDL, INSERT, UPDATE, DELETE). Returns { rowsAffected, lastInsertId } — both are driver-reported and may be zero for statements that don't populate them (e.g. CREATE TABLE).
- **query(sql, ...params)** — Run a SELECT and return one object per row, keyed by column name. Column types map as: INTEGER → number, REAL → number, TEXT → string, NULL → null, BLOB → Uint8Array (unless the bytes are valid UTF-8, in which

case they promote to a string — SQLite stores TEXT and BLOB the same way at the storage layer, so this heuristic recovers the common TEXT case while keeping genuinely-binary BLOBs as bytes).

- **queryValue(sql, ...params)** — Run a statement expected to produce a scalar and return the first column of the first row, or `null` when no rows match. Rows beyond the first are discarded. Ideal for `SELECT count(*)`, `PRAGMA user_version`, single-column lookups.
- **begin()** — Open a transaction; resolves to a nested handle with the same `exec` / `query` / `queryValue` surface plus `commit()` and `rollback()`. The transaction reuses the exact same code paths as the top-level handle (via an internal executor interface that both `*sql.DB` and `*sql.Tx` satisfy) — only the error prefix differs (`sqlite.tx.exec:` vs `sqlite.exec:`). A transaction holds a pooled connection until it commits or rolls back, so **every begin() must reach one of them** — there's no auto-rollback on handle close in this cut. Once finalized, the transaction is spent: further calls throw `sql: transaction has already been committed or rolled back`, and so do double-commit / commit-after-rollback. A constraint violation inside the transaction throws (with the `sqlite.tx.exec:` prefix) but doesn't auto-roll-back — the script decides whether to retry or roll back.
- **prepare(sql)** — Compile a statement once and resolve to a handle whose `exec` / `query` / `queryValue` run it repeatedly with fresh bind params. Unlike the handle / transaction methods, these take **only the params** — no SQL string — since the SQL was fixed at `prepare()` time. Worthwhile for batch-insert or repeated-lookup loops where the per-call parse + plan cost would otherwise dominate. Invalid SQL throws at `prepare()`, not the first `exec()`. Like every handle here it must be `close()`; a leaked statement pins driver resources. Statements are bound to the database handle, not a transaction — using a prepared statement inside a transaction is out of scope for now.
- **close()** — Release the connection. There's no finalizer — scripts that forget to close leak a connection until the process exits. The documented pattern is `open` / `use` / `close`.

Parameters bind positionally as `?` placeholders, in argument order. `goja` exports JS numbers as `int64` / `float64`, strings as `string`, `null` as `nil`, and `Uint8Array` as `[]byte` — all of which the driver accepts directly, so no per-type coercion is needed on the script side. Every handle and transaction method is `async` (Promise-returning); `open` and `begin` are too.

Hash bindings interpret the input as a UTF-8 byte sequence and return lowercase hex. SHA-3 functions are `sha3_256` / `sha3_512` (the IETF spec uses the underscore so the JS name matches recon's). BLAKE3 uses the upstream `lukechampine.com/blake3` reference implementation with a 32-byte output. `crc32` is the IEEE polynomial, zero-padded to 8 hex chars.

String utilities (`api.text.str.*`) follow PHP-/recon-style semantics where they differ from JS:

- `trim` / `ltrim` / `rtrim` accept an optional mask string; **any character in the mask** is stripped. Default mask is whitespace (`" \t\n\r\v\f"`).

- `reverse` is rune-aware, so `reverse("café")` is `"éfac"`. Grapheme clusters are not handled — that's a separate problem.
- `nl2br` emits `<br>` by default; pass `true` as the second arg for XHTML-style `<br/>`.
- `urlEncode` / `urlDecode` use the form encoding (`+` for space). For path-segment encoding, use `encodeURIComponent` instead — that one is provided by goja natively.
- `pad` / `lpad` / `rpadd` follow `recon's str_pad`: `pad` defaults to side `"right"`; `"left"` and `"both"` are also accepted.
- `sprintf` / `printf` are thin wrappers over Go's `fmt`, so the verb set is Go's (`%s`, `%d`, `%x`, `%.2f`, `%v`, `%t`, `%q`, etc.) — not PHP's. `printf` writes to `stdout`.
- `normalizeNewlines` canonicalises any mix of `\r\n`, `\r`, and `\n` to the requested style.

Path utilities (`api.fs.path.*`) are POSIX (forward-slash). On Windows either pass forward-slash paths or convert separators yourself.

Time formatting (`api.runtime.time.format(unixMs, fmt, tz?)`) takes Unix milliseconds (symmetric with `api.runtime.time.nowMs()`) and a small strftime token set: `%Y %y %m %d %H %M %S %F %T %j %A %a %B %b %z %Z %%`. Day and month names are in English. Pass a third-argument IANA zone name (e.g. `"UTC"`, `"America/New_York"`) to render in that zone; if omitted, the host's `time.Local` is used. Unknown `%X` tokens are emitted verbatim so a typo is visible rather than silently dropped.

Protocol probes (`api.net.*`) are `stdlib`-backed and hit the real network. All three return Promises (via `PromisifyAsync`) and accept an optional `{ timeout: <ms> }` second arg; the default is 5 seconds.

- **`tcp(target, opts?)`** — dials TCP, measures round-trip from resolution to handshake. `target` is `"host:port"` or just `"host"` (with `opts.port` overriding the default `80`). Resolves to the remote IP, port, and `latencyMs`.
- **`dns(host, opts?)`** — looks up A, AAAA, MX, TXT, CNAME, and NS records via `net.Resolver`. Pass `opts.types` to scope the query; empty record sets are omitted from the result object so scripts can probe membership with `if ("mx" in result)`.
- **`tls(target, opts?)`** — connects with `InsecureSkipVerify` so even expired or hostname-mismatched certs come back inspectable. Returns the leaf cert's CN, issuer CN, validity window, `daysRemaining`, `dnsNames`, serial number, and a SHA-256 fingerprint (lowercase hex). Default port is `443`. Hosts that care about validity should re-validate the cert themselves.
- **`ntp(host, opts?)`** — queries an NTPv4 server (UDP 123 by default) via github.com/beevik/ntp. Returns the server time (ISO 8601), local-clock offset, round-trip time, stratum, reference time, and the root delay / dispersion values from the server's response. All durations are in milliseconds with sub-millisecond precision.
- **`whois(domain, opts?)`** — two-hop WHOIS lookup (IANA → registrar's whois server) via github.com/likexian/whois with parsing through github.com/likexian/whois-parser. `raw` is always populated; `domain` and `registrar` are best-effort parsed views (some TLDs the parser doesn't understand will return only `raw`). The whois library

doesn't accept a `context.Context` — `opts.timeout` is plumbed through its own per-client setting and the engine's `Options.Timeout` watchdog catches the outer call.

- **ping(host, opts?)** — Reachability probe in two modes. "tcp" (default) dials `host:port` `count` times and measures connect RTT — no privileges, works in containers / CI, and "can I open a connection" is what most liveness checks actually want. "icmp" does real ICMP echo via [pro-bing](#) but needs raw-socket privileges (root / CAP_NET_RAW), so it's opt-in. Returns `{ host, ip, mode, sent, received, lossPercent, minMs, avgMs, maxMs }`. An unreachable host resolves with `received: 0 / lossPercent: 100` — "down" is a normal outcome, not a throw; only DNS-resolution failure and bad arguments throw.
- **smtp(host, opts?)** — SMTP capability probe (not a send pipeline). Opens the connection, reads the banner, sends EHLO, and reports the advertised extensions: `starttls` (boolean), `authMechanisms` (e.g. `["PLAIN", "LOGIN"]`), `sizeLimit`, and the raw `extensions` list. Hand-rolled over `net/textproto` (`net/smtp` doesn't expose the banner or full extension list as data). No mail is sent — it QUITs after EHLO. Connection / protocol failures throw; a server that doesn't advertise STARTTLS / AUTH reports them as `false / []` (a finding, not an error).
- **wss(url, opts?)** — WebSocket handshake probe via [coder/websocket](#). Opens the `ws:// / wss://` connection, optionally measures a ping/pong RTT (`opts.ping`, default true), and closes. Returns `{ url, connected, subprotocol, status, handshakeMs, pingMs }` — `status` is the 101 upgrade response code, `pingMs` is `-1` when the ping was skipped or the server didn't pong. A failed handshake (non-101, refused, bad URL) throws; there's no useful partial result for a failed upgrade. Not a streaming client — the connection doesn't outlive the call.

Email-authentication probes (`api.net.email.*`) read DNS records and surface the published policy. They all return `{ present: false }` when the relevant record is absent (NXDOMAIN or no TXT record matching the marker prefix), so scripts can write a single presence-check pattern across the family.

- **spf(domain)** — queries `TXT(<domain>)` and returns the first record starting with `v=spf1`. `mechanisms` is the tokenised sequence after the version prefix; `allPolicy` summarises the trailing `all` qualifier ("pass" for `all / +all`, "fail" for `-all`, "softfail" for `~all`, "neutral" for `?all`, empty string when no `all` is present).
- **dmARC(domain)** — queries `TXT(_dmarc.<domain>)` and parses the `v=DMARC1; key=val; ...` form into a flat tag map (keys are case-folded; values retain internal whitespace and commas). `policy / subdomain / percent / rua / ruf` are surfaced separately because those are the fields most scripts care about.
- **mtaSts(domain)** — looks up `TXT(_mta-sts.<domain>)` and, on success, fetches the policy file at `https://mta-sts.<domain>/.well-known/mta-sts.txt`. The TXT carries a versioned `id` for change detection; the policy file is where the actual `mode + mx` list live. RFC 8461 format (line-based, `key: value`, `mx:` repeats) is parsed inline. Policy-fetch failures (TLS errors, 4xx, timeout) don't fail the binding — the TXT part is still

returned, and the fetch error surfaces under `policyError` so scripts can decide what to do.

- **`tlsRpt(domain)`** — looks up `TXT(_smtp._tls.<domain>)` and parses the `v=TLSRPTv1; rua=...` form into a tag map. `rua` is surfaced separately because that's the actionable bit.
- **`bimi(domain, opts?)`** — looks up `<selector>._bimi.<domain>`, selector defaulting to `default`. Surfaces the logo URL (`l`) and VMC URL (`a`) from the tag map.
- **`all(domain)`** — runs every probe above in parallel via goroutines and returns a single aggregate object keyed by probe name. Per-probe failures don't fail the aggregate — they surface under `<probe>.error` so a partial result is still useful (e.g. SPF + DMARC found, MTA-STS policy fetch timed out).

`api.ui.tui.*` — multi-pane TUI

Scripts that orchestrate multiple subprocesses (e.g. `brew`, `npm`, `cargo` updates running in parallel) can declare a multi-pane terminal layout and route each subprocess's stdout/stderr into its own pane. Activation is **script-driven**: the first call to `api.ui.tui.layout(...)` enters TUI mode. If stdout is not a TTY (CI, pipes, `make demo`), the same calls fall back to prefixed plain-text lines (`[paneName] line`) so the same script runs in both contexts.

Layout

```
api.ui.tui.layout({
  rows: [
    { name: "log", title: "Orchestrator", weight: 1 },
    { cols: [
      { name: "brew" },
      { name: "npm" },
    ],
      weight: 2,
    },
  ],
});
```

Each layout node is one of:

- `{ name: string, title?: string, weight?: number }` — a leaf pane.
- `{ rows: LayoutNode[], weight?: number }` — a vertical split.
- `{ cols: LayoutNode[], weight?: number }` — a horizontal split.

`weight` defaults to 1. Pane names must be unique across the whole tree. `layout` may be called **once** per Run; a second call throws.

Pane handles

```
const log = api.ui.tui.pane("log");
log.writeln("Updating Homebrew...");
log.title("Done");
log.clear();
log.write("partial line ");
```

`api.ui.tui.pane(name)` returns a handle with `write`, `writeln`, `clear`, `title`. Methods are synchronous from the script's perspective — they enqueue to the TUI goroutine.

Subprocess routing

`api.tools.exec.shell(cmd, opts)` accepts `opts.pane` (a Pane handle or a pane name string):

```
await api.tools.exec.shell("brew update && brew upgrade", { pane: "brew"
});
```

When set, the subprocess's stdout **and** stderr stream into the pane line by line. The returned `{ exitCode, durationMs, success }` is the same as without `pane`: except `stdout` and `stderr` are empty (data was streamed, not captured). ANSI colors are translated to `tview` color tags and rendered natively; `\r` without `\n` overwrites the current line so progress spinners render cleanly.

Keybindings (TTY mode only)

Key	Action
Tab / Shift-Tab	Cycle focus through panes
PgUp / PgDn	Scroll focused pane one page
↑ / ↓	Scroll focused pane one line
Home / End	Jump to top / bottom
Ctrl-C	Abort the script

The focused pane has a yellow border; the bottom status bar shows its name and the active keys.

Limitations (v1)

- `api.ui.tui` is **incompatible with** `--watch`: calling `api.ui.tui.layout()` under `--watch` throws. Use one or the other.
- No mouse, no runtime layout mutation, no input panes (`api.ui.tui.input`), no snapshot-to-normal-screen on exit. The alt screen is restored at script end and the usual `PASS / FAIL` line prints.

6. JavaScript runtime built-ins (goja)

`scriptengine` runs on goja, which implements **ES5.1 + a large subset of ES6+**. The following are present and behave per spec:

Globals

- `Object`, `Array`, `String`, `Number`, `Boolean`, `BigInt`
- `Math`, `Date`, `JSON`, `RegExp`
- `Error`, `TypeError`, `RangeError`, `SyntaxError`, `ReferenceError`, `EvalError`, `URIError`
- `Promise` (provided via goja's native implementation; resolution microtasks are pumped by the event loop)
- `Symbol` (partial: well-known symbols `Symbol.iterator`, `Symbol.asyncIterator`, etc., are supported)
- `Proxy`, `Reflect` (partial)

```
api.runtime.log(Math.PI.toFixed(4), Math.max(1, 9, 3));
api.runtime.log(new Date().toISOString());
api.runtime.log(JSON.stringify({ a: 1, b: [2, 3] }));
api.runtime.log("abc".repeat(3), "abc".padStart(6, "_"));
```

Collections

- `Map`, `Set`, `WeakMap`, `WeakSet`

```
const counts = new Map<string, number>();
for (const ch of "banana") counts.set(ch, (counts.get(ch) ?? 0) + 1);
api.runtime.log([...counts]); // [{"b",1}, {"a",3}, {"n",2}]
```

Typed arrays

- `ArrayBuffer`, `DataView`
- `Int8Array`, `Uint8Array`, `Uint8ClampedArray`, `Int16Array`, `Uint16Array`, `Int32Array`, `Uint32Array`, `Float32Array`, `Float64Array`

```
const buf = new ArrayBuffer(4);
new DataView(buf).setUint32(0, 0xCAFEBAFE);
api.runtime.log(new Uint8Array(buf)[0].toString(16)); // ca
```

Iteration and generators

- `for...of`, `for...in`, `spread`, `destructuring`

- Generator functions (`function*`)
- Async generators (`async function*`) — supported via goja's event loop

```
function* fib() {
  let [a, b] = [0, 1];
  while (true) {
    yield a;
    [a, b] = [b, a + b];
  }
}
const it = fib();
const first10 = Array.from({ length: 10 }, () => it.next().value);
api.runtime.log(first10);
```

Not (yet) provided

- `fetch` — use `api.net.http.*` from the CLI, or register your own binding.
- `Worker`, threading primitives.
- DOM globals (`window`, `document`).
- Native ES modules (`import` / `export` are *transpiled* to CommonJS at load time — see [section 8](#)).

7. Async runtime additions (goja_nodejs)

The event loop bundles a small Node-compatible runtime on top of goja:

Timers

```
setTimeout(() => api.runtime.log("after 50ms"), 50);
setInterval(() => api.runtime.log("tick"), 100);
setImmediate(() => api.runtime.log("next tick"));
clearTimeout(t); clearInterval(i); clearImmediate(im);
```

The event loop holds the script alive while *any* timer is pending. If your script ends with only `setTimeout` callbacks scheduled, the run will wait for them to fire before returning.

console

```
console.log("info");
console.error("oops");
console.warn("careful");
console.info("fyi");
console.debug("trace");
```

Disable with `Options.DisableConsole = true` if you want a clean global namespace.

`require`

CommonJS `require`, with TS-aware extension fallback added by `sercon` (see [section 10](#)):

```
const helpers = require("./helpers");
helpers.doThing();
```

Both `require("./foo")` and `import { x } from "./foo"` resolve to the same module instance per Run; esbuild rewrites `import` to `require` at transpile time.

8. TypeScript support

TS is transpiled by esbuild in-process. There is no `tsc`, no `tsconfig.json`, no type-checking — types are erased and the resulting JS is what executes. Practical implications:

- All TS syntax esbuild supports is allowed: type annotations, `interface`, `type`, generics, enums (compiled to objects), namespaces, decorators (legacy & TC39 stage 3).
- `import type { X } from "y"` is stripped entirely (no runtime require).
- TS errors are *not* surfaced as build errors — only esbuild parse errors are. Run a separate `tsc --noEmit` in CI if you want type enforcement.
- `tsconfig.json` is **not** consulted.

`.tsx` and `.jsx` are supported via esbuild's `LoaderTSX` / `LoaderJSX`, including for required modules resolved by the engine's extension fallback. JSX has no React runtime in scope by default — either set the factory at the file level with an `@jsx` pragma:

```
/** @jsx h */
function h(tag: string, props: any, ...children: any[]) {
  return { tag, props: props ?? {}, children };
}
export const Box = (label: string) => <div className="box">{label}</div>;
```

...or provide a `React.createElement`-compatible binding from Go and rely on esbuild's default pragma. ESM `export default <value>` also works through the entry-script rewriter's `__esModule ? .default : m` unwrap, so `import answer from "./mod"` resolves to the default export.

9. Top-level `await`

Top-level `await` works in entry scripts:

```
const r = await api.net.http.get("https://example.com");
api.runtime.log(r.status);
```

How it works under the hood: esbuild rejects top-level await with the `cjs` output format, so the engine transpiles the *entry* script with `format=esm`, then rewrites `import` statements to `require()` calls and wraps the rest of the body in:

```
;(async () => { /* your transpiled body */ })().then(__resolve, __reject);
```

`__resolve` and `__reject` are bindings the engine sets on the runtime to capture the final settlement. *Required-module* code (anything loaded through `require`) is transpiled with `format=cjs` and does **not** support top-level await — wrap in `(async () => ...)` yourself if you need it inside a module.

10. Module resolution

The engine plugs a custom source loader into `goja_nodejs/require` that adds:

1. **Direct path**: if the requested file exists on disk, use it. `.ts` / `.tsx` files are transpiled on the fly.
2. **Extension fallback** when the request has no extension: try `.ts`, `.tsx`, `.js`, `.cjs`, `.mjs`, `.json` in that order.
3. **.js → .ts swap**: if a path ending in `.js` / `.cjs` / `.mjs` is requested but the literal file doesn't exist, the same path ending in `.ts` (or `.tsx`) is tried. Handles `package.json` `main` fields that point at compiled output where only the TypeScript source is on disk.
4. **package.json source preference**: when reading a `package.json`, if a `source` field is present and points at an existing `.ts` / `.tsx` file, `main` is rewritten to that path before the registry consumes it. Matches the convention used by parcel, microbundle, and similar bundlers to mark the TS source of truth.
5. **Directory index**: if the resolved path is a directory, try `index.ts`, `index.tsx`, `index.js`.

The base directory follows Node's CommonJS rules — relative imports resolve against the requiring module's directory, courtesy of `goja_nodejs`.

```
// All resolve to ./helpers/assert.ts:
import { check } from "./helpers/assert";
import { check } from "./helpers/assert.ts";
const { check } = require("./helpers/assert");
const { check } = require("./helpers/assert.js");
```

JSON imports work out of the box via either the ESM-style default import (esbuild rewrites it to a `require`) or a direct `require`:

```
import data from "./data.json";
const r = require("./data.json");
```

`node_modules` lookup *works* via the upstream registry, but the source loader doesn't currently transpile through that path — keep TS helpers under `ScriptRoot`.

11. Timeouts and cancellation

Two mechanisms interrupt a running script:

- `Options.Timeout` — wall-clock limit per `Run`. Exceeding it returns `scriptengine.ErrScriptTimeout`.
- `ctx` passed to `Run / RunFile` — cancelling the context returns `ctx.Err()` (typically `context.Canceled` or `context.DeadlineExceeded`).

Both call `vm.Interrupt(...)` and `loop.Terminate()` from a watcher goroutine. This works for **sync** JS — including `while(true){}` — because the goja VM checks the interrupt flag between bytecode steps. A host Go binding that's blocked in cgo or a long syscall isn't interruptible mid-call; if you write such a binding, plumb the engine `ctx` through to it yourself.

```
eng := scriptengine.New(scriptengine.Options{Timeout: 200 *
time.Millisecond})
_, err := eng.Run(ctx, "loop.ts", `while (true) {}`)
// err is scriptengine.ErrScriptTimeout, returned ~200–300ms after Run.
```

12. Error semantics

- A Go binding returning `(T, error)` automatically throws as a JS exception when `err != nil`. The exception's `.message` is `err.Error()`.

```
try { mayFail(); } catch (e) { api.runtime.log(String(e)); }
```

- A Go binding that `panic`s with a `*goja.Object` (e.g. `vm.NewGoError(...)`) does the same.
- A script throwing a JS error out of the top level surfaces from `Run` as a Go `error` whose message is the JS error's `.message`.
- Timeout errors satisfy `errors.Is(err, scriptengine.ErrScriptTimeout)`. Cancellation errors satisfy `errors.Is(err, context.Canceled)` or `context.DeadlineExceeded`.

13. Type generation (.d.ts)

`Engine.WriteTypes(w)` walks the registered bindings and emits a TypeScript declaration file. The mapping table (abbreviated):

Go type	TS type
<code>string</code>	<code>string</code>
<code>any numeric / float64</code>	<code>number</code>
<code>bool</code>	<code>boolean</code>
<code>[]T</code>	<code>T[] ; []byte → Uint8Array</code>
<code>map[string]T</code>	<code>Record<string, T></code>
<code>*T</code>	<code>T</code> (transparent unwrap)
<code>error</code> (single return)	omitted (collapses to <code>void</code>)
<code>(T, error)</code> (tuple return)	<code>T</code>
<code>interface{}</code> / <code>any</code>	<code>unknown</code>
<code>goja.FunctionCall</code> parameter	folded into <code>(...args: unknown[])</code>
<code>goja.Value</code> return	<code>unknown</code>
<code>scriptengine.AsyncBinding</code> value (from <code>PromisifyAsync[T]</code>)	<code>Promise<T></code>
self-referential types	<code>unknown</code> (cycle detection bails after depth 4)

```
sercon --emit-dts api.d.ts
```

In your editor, place the resulting `.d.ts` next to your scripts (or reference it via `/// <reference path="..." />`) for autocomplete.

JSDoc comments on declarations

The emitter pulls JSDoc strings from the engine's doc map (set via `Engine.SetDocs / Engine.SetMemberDocs`) and renders them above each declaration. Single-line docs collapse to `/** ... */`; multi-line docs expand to a standard `*`-prefixed block. Bindings without a doc entry emit no JSDoc — the emitter never inserts empty placeholders. Docs are looked up by the same dotted path the binding was registered under, so namespace members get hover text too:


```
// AUTO-GENERATED by scriptengine.Engine.WriteTypes – do not edit by hand.

/** sercon's built-in script surface. ... */
declare const api: {
  hash: {
    /** SHA-256 hex digest of a UTF-8 input. */
    sha256(...args: unknown[]): string;
    /** BLAKE3 hex digest (32-byte output, lukechampine.com/blake3). */
    blake3(...args: unknown[]): string;
    // ...
  };
  // ...
};
```

The CLI's `api.*` surface ships with a curated doc map; see `cmd/sercon/api_docs.go` for the source of truth and add new entries there when adding new bindings (the lockstep rule keeps `--examples / MANUAL / api.d.ts` in sync, and the doc map is now the seventh artifact in that chain).

14. Limitations and gotchas

- **No native ES modules at runtime.** All `import` is rewritten to `require` by esbuild. Side-effect-only imports run the module body exactly once per Run.
- **No top-level await inside required modules.** Wrap with `(async () => ... )()` inside the module if you need it.
- **No `tsconfig.json` honoured.** Module resolution and type checking are independent from any project-level TS config.
- **Per-Run runtime.** Side effects on `globalThis` do not persist between calls to `Run` on the same `Engine`. The `require` *compile* cache is shared; module *exports* are not.
- **Multi-line ESM imports may stress the entry rewriter.** The current rewriter is line-based and handles the common forms; especially gnarly inputs (comments inside the spec list, very unusual whitespace) fall back to executing as-is, which may throw.
- **RegisterConstructor is d.ts-only today.** At runtime it behaves like `Register`. True `new`-able constructor semantics are on the roadmap.
- **HTTP bindings are real network calls.** `api.net.http.*` uses `net/http` with a 5s timeout. They are not mockable from JS.

See [OUT-OF-SCOPE.md](#) for the active backlog of deferred ideas.

This manual covers sercon v0.8.0. Whenever you add, remove, or change a flag, a binding, or the script API, update this file alongside the help screen (`--help`), the examples walkthrough (`--examples`), and the `CHANGELOG.md`.