

sercon User Manual

Embeddable TypeScript script engine — reference and
guide

0.66.0
2026-06-24
Thomas Bjork

Contents

sercon manual	1
1. Overview	1
2. Quickstart	1
3. Library API — pkg/scriptengine	1
Engine, Options, New	1
Registering bindings	2
Run, RunFile	3
Reset	4
WriteTypes	4
SetDocs and SetMemberDocs	4
PromisifyAsync[T] and AsyncBinding	5
Version	5
4. CLI — sercon	6
Passing arguments: -- and runtime.argv	7
--watch : re-run on file change	7
Editor autocomplete (sercon init)	8
Shebang lines and executable scripts (sercon run)	9
sercon serve : long-running scripts with production niceties	9
5. Reserved globals (script surface)	10
console (browser/Node compatibility)	11
Migration from v0.8.0	11
runtime	11
crypto	13
text	14
codec	15
fs	18
net	19
db	26
services	27
tui	46
image	48
web	50
6. Servers	51
6.1 Foundation: HoldRun + LoopCallable	52
6.2 server.http.listen / server.https.listen	52
6.3 Static-file serving	55
6.4 Middleware	56
6.5 WebSocket upgrade	56
6.6 Lifecycle	58
6.7 SMTP	58
6.8 Raw TCP / UDP	62
7. JavaScript runtime built-ins (goja)	63
Globals	63
Collections	64
Typed arrays	64
Iteration and generators	64

Not (yet) provided	64
8. Async runtime additions (goja_nodejs)	65
Timers	65
console	65
require	65
9. TypeScript support	65
10. Top-level await	66
11. Module resolution	66
12. Timeouts and cancellation	67
13. Error semantics	67
14. Type generation (.d.ts)	68
JSDoc comments on declarations	68
15. Limitations and gotchas	69
Bundled libraries	69
paymentproviders	69
16. Binding reference (generated)	70
codec	70
console	78
crypto	80
db	87
fs	93
image	97
net	99
runtime	118
server	125
services	131
text	151
tui	162
web	164

sercon manual

Long-form reference for the `pkg/scriptengine` library, the `sercon` CLI, the built-in script API, and the JavaScript runtime surface that `scriptengine` exposes via `goja` and `goja_nodejs`. Examples are runnable — save snippets as `.ts` files and execute with `sercon file.ts` unless otherwise noted.

1. Overview

`sercon` is a Go program that runs TypeScript files. It's built from two parts:

- `pkg/scriptengine` — a library you embed in your own Go code. You register Go-callable bindings, then call `Run` / `RunFile` to execute a `.ts` script that talks to them.
- `cmd/sercon` — a thin CLI built on the library, useful for ad-hoc test scripts and as a working example of every binding kind.

The runtime is pure Go: `goja` is the JS engine, `esbuild` (used as a Go library) is the TS→JS transpiler, and `goja_nodejs` provides Promises, `setTimeout`, `console`, and CommonJS `require`. No cgo, no Node.

2. Quickstart

Install:

```
go install github.com/codedeviate/sercon/cmd/sercon@latest
```

Run one of the bundled examples:

```
sercon examples/scripts/smoke.ts examples/scripts/async.ts
```

...or run them all via `make demo`. `examples/scripts/` ships a runnable `.ts` (or `.tsx`) file per feature area — `hash.ts`, `strings.ts`, `path-and-time.ts`, `default-export.ts`, `tsx-demo.ts`, and the original `smoke.ts` / `async.ts` / `hang.ts` (`hang.ts` is the timeout demo and intentionally exits non-zero).

Generate a declaration file for editor autocomplete:

```
sercon --emit-dts sercon.d.ts
```

Embed in your own Go program:

```
eng := scriptengine.New(scriptengine.Options{
    Timeout:    5 * time.Second,
    ScriptRoot: "./scripts",
})
eng.Register("upper", strings.ToUpper)
val, err := eng.RunFile(ctx, "scripts/main.ts")
```

3. Library API — `pkg/scriptengine`

`Engine`, `Options`, `New`

```

type Options struct {
    Timeout      time.Duration // wall-clock limit per Run; 0 disables
    ScriptRoot    string         // base for require/import resolution
    DisableConsole bool          // turn off the goja_nodejs console module
    Verbose       io.Writer      // diagnostic traces, prefixed "[sercon] "
    ModuleLoader  func(candidatePath string) (source string, found bool, err
error)
}

func New(opts Options) *Engine

```

`ModuleLoader`, when non-nil, is consulted for every require/import candidate path **before** the filesystem — the hook for embedders that want to serve modules from somewhere other than disk (an in-memory FS, a network source, an embedded bundle). goja probes several candidates per specifier (`./x`, `./x.ts`, `./x/index.ts`, ...); the engine also tries the bare path plus the usual extension fallbacks (`.ts` / `.tsx` / `.js` / `.mjs` / `.cjs` / `.json`) against the loader, so a loader can match on a plain suffix. Return `(source, true, nil)` to serve the module (transpiled when the path ends in `.ts` / `.tsx`, exactly like a disk read); `("", false, nil)` to fall through to the filesystem; `("", false, err)` to abort resolution.

```

eng := scriptengine.New(scriptengine.Options{
    ModuleLoader: func(path string) (string, bool, error) {
        if strings.HasSuffix(path, "greeting.ts") {
            return `export const hi = (n: string) => "hi " + n;`, true, nil
        }
        return "", false, nil // fall through to disk
    },
})

```

`Engine` is the host of all registered bindings. Construct it once, then call `Run` / `RunFile` as many times as needed — each call gets a *fresh* `*goja.Runtime` so script state never leaks between invocations.

`ScriptRoot` defaults to the working directory at `New` time if left empty. Both `Timeout` and `ScriptRoot` can be overridden on a per-Run basis only by constructing a fresh `Engine` see Out-of-scope for the per-Run override on the roadmap.

Registering bindings

Function	Use when...
<code>Register(name, value)</code>	The binding is a pure function, struct, or primitive that doesn't need access to the runtime.
<code>RegisterNamespace(name, members map[string]any)</code>	You want to expose <code>name.x</code> , <code>name.y</code> without defining a Go struct.
<code>RegisterConstructor(name, ctor)</code>	The binding is a Go function whose return value should be treated as a class in the emitted <code>.d.ts</code> . (Runtime semantics today are identical to <code>Register</code> the <code>d.ts</code> emitter is the only differentiator.)

<code>RegisterFactory(name, factory func(vm, loop) any)</code>	The binding needs the per-Run <code>*goja.Runtime</code> or <code>*eventloop.EventLoop</code> in scope (typical for Promise-returning I/O).
<code>RegisterNamespaceFactory(name, factory)</code>	Same, but the factory returns the full member map.

Example — synchronous function binding:

```
eng.Register("greet", func(name string) string { return "hi " + name })
```

Example — namespace with a sync member:

```
eng.RegisterNamespace("math2", map[string]any{
    "square": func(n float64) float64 { return n * n },
    "pi":     3.14159,
})
```

Example — async (Promise-returning) binding via `PromisifyAsync`:

```
eng.RegisterFactory("httpGet", func(vm *goja.Runtime, loop *eventloop.EventLoop)
any {
    return scriptengine.PromisifyAsync(vm, loop,
        func(ctx context.Context, call goja.FunctionCall) (map[string]any, error)
        {
            url := call.Argument(0).String()
            resp, err := http.Get(url)
            if err != nil {
                return nil, err
            }
            defer resp.Body.Close()
            body, _ := io.ReadAll(resp.Body)
            return map[string]any{"status": resp.StatusCode, "body":
string(body)}, nil
        })
})
```

Run, RunFile

```
func (e *Engine) Run(ctx context.Context, name, source string, opts ...RunOption)
(goja.Value, error)
func (e *Engine) RunFile(ctx context.Context, path string, opts ...RunOption)
(goja.Value, error)
```

`Run` executes `source` as an entry-script TS. `name` is used in stack traces. The returned value is currently always `undefined` — top-level expression capture is on the backlog.

Both methods respect `Options.Timeout` and `ctx` cancellation, whichever fires first; the resulting error is either `ErrScriptTimeout`, `ctx.Err()`, or the underlying JS exception.

Per-Run options

```
type RunOption func(*runConfig)
```

```
func WithScriptRoot(dir string) RunOption
```

Reuse a single Engine across many scripts that each live in their own directory by passing `WithScriptRoot(dir)` to `Run / RunFile`. The override applies only to this call; `Options.ScriptRoot` is used for every other `Run`.

```
_, _ = eng.Run(ctx, "main.ts", source, scriptengine.WithScriptRoot("/path/to/run42"))
```

```
func WithArgs(args []string) RunOption
```

Set the per-script argument vector. The script reads it as `runtime.argv`, a global with Node/Bun layout: `argv[0]` is `Options.ProgramName` (defaults to `filepath.Base(os.Args[0])` the CLI sets it to `"sercon"`), `argv[1]` is the running script's path, and the `WithArgs` values follow from index 2. The global is always present — with no args it is just `[programName, scriptPath]`.

```
_, _ = eng.Run(ctx, "main.ts", source, scriptengine.WithArgs([]string{"--port", "8080"}))
// inside the script: runtime.argv == [ProgramName, "/abs/main.ts", "--port", "8080"]
```

Reset

```
func (e *Engine) Reset()
```

Clears every registered binding. Useful when a long-lived Engine is reused across unrelated batches of scripts that each want a clean global namespace. **Not safe to call concurrently with `Run / RunFile`.**

WriteTypes

```
func (e *Engine) WriteTypes(w io.Writer) error
```

Emits a `.d.ts` describing the registered surface. See section 14. Type generation for what the mapping looks like.

SetDocs and SetMemberDocs

```
func (e *Engine) SetDocs(path, doc string)
func (e *Engine) SetMemberDocs(namespace string, docs map[string]string)
```

Attach JSDoc strings to registered bindings so the emitted `.d.ts` grows readable editor hover. `SetDocs` takes a dotted path — either a bare top-level name (`"greet"`, `"http"`) or `"namespace.member"` for namespace members (`"http.get"`, `"exec.shell"`); `SetMemberDocs` is the convenience for documenting many members of a namespace at once (the namespace itself can still be documented separately via `SetDocs`).

Multi-line docs are written as `\n`-separated strings; the emitter expands them into a standard `*`-prefixed JSDoc block. Single-line docs collapse to `/** ... */`. Calling `SetDocs` with an empty string removes any previous doc for that path. Bindings without a doc entry get no JSDoc block at all — the emitter doesn't insert empty `/** */` placeholders.

`SetDocs` may be called before or after the matching `Register...` call; docs are looked up by name at emit time. Like the registration methods, `SetDocs` must not race with a `Run` / `WriteTypes` / `Reset` on the same engine.

```
eng.Register("greet", func(name string) string { return "hi " + name })
eng.SetDocs("greet", "Greet someone by name.")

eng.RegisterNamespace("math2", map[string]any{
    "pi":      3.14159,
    "squared": func(n float64) float64 { return n * n },
})
eng.SetDocs("math2", "Tiny math helpers.")
eng.SetMemberDocs("math2", map[string]string{
    "pi":      "Circumference / diameter ratio.",
    "squared": "Multiply a number by itself.",
})
```

PromisifyAsync[T] and AsyncBinding

```
type AsyncBinding struct {
    Func      func(goja.FunctionCall) goja.Value
    TSReturnType string
}

func PromisifyAsync[T any](vm *goja.Runtime, loop *eventloop.EventLoop,
    work func(ctx context.Context, call goja.FunctionCall) (T, error),
) AsyncBinding
```

`PromisifyAsync` turns blocking Go work into a JS Promise. It launches the work in a goroutine, parks a `SetTimeout` sentinel so `loop.Run` doesn't return early, and schedules the resolve/reject back onto the event loop. **Required** for any Promise-returning binding — `RunOnLoop` alone is not counted as a live job by the event loop.

`PromisifyAsync` returns an `AsyncBinding` carrier rather than the bare callback. The engine unwraps it to `.Func` at `vm.Set` time so `goja`'s special-case detection of `func(goja.FunctionCall) goja.Value` host-callbacks still fires; the `.d.ts` emitter reads `.TSReturnType` to emit `Promise<T>` (mapped from the generic `T` via `reflect` at construction time) instead of the previous `unknown`. Hosts assigning the result into a `map[string]any` for `RegisterNamespace` / `RegisterNamespaceFactory` get the unwrap recursively too — no manual work required.

Version

```
import "github.com/codedeviate/sercon/pkg/scriptengine"

fmt.Println(scriptengine.Version) // "0.1.0"
```

Bumped in lockstep with the git tag.

4. CLI — sercon

```
sercon [flags] <script.ts> [script.ts ...]
sercon [flags] -                # read entry script from stdin
sercon --examples | --help | --version
```

Flag	Purpose
-timeout DURATION	Wall-clock limit per script (default 10s 0 disables).
-root DIR	Override the script root for require / import resolution.
-emit-dts PATH	Write the example bindings' .d.ts to PATH and exit.
-v	Verbose: trace the rewritten entry-script JS and each module resolution to stderr; also print duration on failure.
-h, --help	In-depth colored help.
--examples	In-depth colored walkthrough of every feature.
--no-pager	Don't page --help / --examples . By default, when stdout is a terminal they pipe through \$PAGER (falling back to less with LESS=FRX , color preserved); a pipe/redirect, --no-pager , or PAGER=cat renders directly.
--version	Print the engine version (plus goja/esbuild build-info versions).
--doctor	Check external tool requirements (git, gh, AI providers, agent-browser, chromedriver/geckodriver, typst, recon/curl, clipboard/image): installed?, version, purpose — and validate the chromedriver↔Chrome major-version match; then exit. Prints a category-grouped table. Exits 0 normally (missing tools are optional), 5 on a detected compatibility conflict. Same report as the services.doctor binding.
--watch	Re-run on file change under the script root. Debounced 150 ms. Ctrl-C exits.
--secrets-prefix=P	Namespace prefix for runtime.secrets keystore items. Overrides \$SERCON_SECRETS_PREFIX default sercon/ .
--env-file=PATH	Load KEY=VALUE pairs from a .env file into the environment before running (repeatable). Parses KEY=VALUE , # comments, blank lines, an optional leading export , and optional surrounding quotes — no shell expansion. A variable already in the real environment always wins; among multiple files, a later file overrides an earlier one. Replaces the set -a; source .env; set +a ritual and makes shebang scripts self-sufficient.

Each positional argument is either a path to a .ts / .tsx file or - to read an entry script from standard input:

```
echo 'runtime.log(1 + 2);' | sercon -
sercon prelude.ts -                # prelude then stdin
```

Passing arguments: `--` and `runtime.argv`

Everything after a standalone `--` is the script argument vector, exposed to every script via `runtime.argv`:

```
sercon run.ts -- --port 8080 hello
```

`runtime.argv` uses the Node/Bun layout `[programName, scriptPath, ...userArgs]` — `argv[0]` is `"sercon"`, `argv[1]` is the path of the script currently executing, and the post-`--` arguments start at index 2 (so `runtime.argv.slice(2)` is just your args). All scripts in one invocation share the same `--` tail, but each sees its own path at `argv[1]`. The global is always present; with no `--` it is just `[programName, scriptPath]`.

```
const args = runtime.argv.slice(2); // ["--port", "8080", "hello"]
```

Exit codes are distinct per failure type so shells can react sensibly. When several scripts run, the highest applicable code wins:

Code	Meaning
0	every script passed.
1	CLI usage error (unknown flag, missing script argument, ...).
2	at least one script failed to transpile — never ran.
3	at least one script timed out (<code>-timeout</code>) or was context-cancelled.
4	at least one script ran and threw a JS exception.
5	<code>--doctor</code> detected a compatibility conflict (e.g. chromedriver↔Chrome major mismatch).

`-v` writes lines prefixed with `[sercon]` to stderr. The traces include the full rewritten entry-script JS (the form goja actually runs, after the ESM→CJS rewrite + async IIFE wrapper) and every module-resolution event, so debugging an unexpected resolve target or a mis-rewritten import is straightforward.

The CLI registers twelve reserved top-level globals; see the next section.

`--watch` : re-run on file change

`sercon --watch <script.ts>` runs the script once and then blocks, re-running it every time a watched file changes under the script root. Useful for iterating on a script body or on the binding surface during development.

```
sercon --watch examples/scripts/smoke.ts
```

What's watched:

- The script root (resolved the same way as the one-shot mode — `-root DIR` if set, else `dirname` of the first script argument).
- Recursively. Symlinks aren't followed.
- New directories that appear during the session are picked up automatically.

What's filtered:

- **File extensions:** `.ts` / `.tsx` / `.js` / `.jsx` / `.json` trigger a re-run; `.d.ts` files match via suffix too. Anything else (Markdown, images, editor swap files, build artefacts) is ignored.
- **Directories:** hidden directories (`.git` , `.vscode` , anything starting with `.`) and `node_modules` are excluded — both because they generate floods of irrelevant events and because recursing into them inflates the watcher’s directory count for no gain.

Re-runs are **debounced 150 ms**. Editors typically fire several events per save (write → rename → chmod); collapsing them into a single re-run keeps the loop responsive without thrashing. The debounce window resets on every event, so a rapid burst of saves becomes one re-run after the burst settles.

Each run is delimited by a line like `--- sercon re-run @ HH:MM:SS ---` so the output is visually distinct from the previous run.

`Ctrl-C` (SIGINT) or `SIGTERM` exit cleanly. Per-script failures inside a watch session log as `FAIL` but don’t terminate the loop — a syntax error you’re trying to fix shouldn’t kick you out of watch mode.

The watcher reuses the same `Engine` across runs. Each `Run` already gets a fresh `*goja.Runtime` , so re-running is just re-invoking `runOne` on the existing engine — there’s no per-iteration setup cost beyond the transpile + module-resolution work.

Module-graph invalidation. Watch mode tracks each entry script’s import graph (its own file plus every module it resolves, captured via `Engine.SetResolveHook`) during the run. On a file change it re-runs **only the entries whose graph includes the changed file** — editing a helper that just one of three entry scripts imports re-runs that one entry, not all three. An entry’s own file counts (editing it re-runs it); stdin entries and any entry whose graph isn’t known yet re-run unconditionally (conservative). Before each re-run the module cache is busted (`Engine.ResetModuleCache`) so an edited import’s new source actually takes effect — the registry otherwise caches compiled bytecode across runs.

Editor autocomplete (`sercon init`)

```
sercon init [dir]
```

Drops two files into `dir` (default the current directory) so any editor backed by the TypeScript language server (VSCode, Zed, Neovim+coc, Sublime LSP, ...) gives completion and hover docs for the reserved globals with no plugin or manual config:

- **`sercon.d.ts`** — the binding declarations (same content as `sercon -emit-dts`): ambient declare const blocks for `runtime` , `crypto` , `text` , `codec` , `fs` , `net` , `db` , `server` , `services` , `tui` , and `console` , each with JSDoc.
- **`jsconfig.json`** — points the language server at `sercon.d.ts` and sets `moduleResolution: "Bundler"` (so the extensionless relative imports sercon scripts use resolve) and `types: []` (sercon isn’t Node, so no stray `@types/*`).

Existing files are left untouched unless `--force` is given. The `examples/scripts/` directory ships with both files already, so the bundled demos autocomplete out of the box. For a single file without a config, the no-setup fallback is a leading

```
/// <reference path="./sercon.d.ts" /> .
```

Shebang lines and executable scripts (`sercon run`)

A script may begin with a `#!/` shebang line. It is stripped before transpile (the line is blanked in place, so transpile/syntax error line numbers still match the source), which means a `.ts` / `.tsx` / `.js` file can be made directly executable:

```
#!/usr/bin/env sercon
runtime.log("hello from an executable script");
```

```
chmod +x hello.ts
./hello.ts
```

The kernel launches a shebang script as `sercon <script> <args...>`, and in the default mode every positional is treated as a *separate script* — so positional arguments to a shebang script would be mistaken for additional script paths. For an executable script that takes arguments, use the **run subcommand** in the shebang:

```
sercon run [flags] <script.ts> [args...]
```

```
#!/usr/bin/env -S sercon run
runtime.log("args:", JSON.stringify(runtime.argv.slice(2)));
```

`sercon run` executes exactly one script and hands every token after the script path to it as `runtime.argv[2:]` (Node/Bun layout: `[program, script, ...args]`) — no standalone `--` separator is needed (and a shebang line can't inject one). The `env -S` form splits the single shebang argument so the kernel runs `sercon run /abs/path/script.ts arg1 arg2 ...`. Flags (`-timeout`, `-root`, `-v`) are accepted before the script path; everything from the script path onward is positional and becomes script args. `env -S` is supported by GNU coreutils, macOS, and the BSDs.

```
sercon run script.ts --port 8080 alice    # argv[2:] = ["--port", "8080", "alice"]
```

`sercon serve` : long-running scripts with production niceties

```
sercon serve [flags] <script.ts> [-- args...]
```

`sercon serve` is a sibling subcommand that wraps the same engine but adds the conveniences a process supervisor (systemd, docker, launchd) usually wants from a long-lived script. Use it whenever a script binds an HTTP/HTTPS listener (or any future `server.*` protocol) and is expected to keep running. Vanilla `sercon script.ts` also keeps the loop alive while listeners are bound, but gives you none of the deltas below.

Flag	Purpose
<code>--shutdown-timeout DURATION</code>	Graceful-shutdown deadline on SIGTERM/SIGINT (default 30s). After the deadline the engine hard-cancels.
<code>--port-override N</code>	If non-zero, replaces every literal <code>port:</code> in <code>server.*.listen({...})</code> calls with <code>N</code> . Useful for spinning up the same script on a different port without editing

	source. Only literal port values are rewritten — computed expressions (<code>{port: getPort()}</code>) are left alone.
<code>-timeout DURATION</code>	Per-script wall-clock limit. Defaults to <code>0</code> (disabled) under <code>serve</code> so listeners don't get cut off; opt back in if you want one.
<code>-root DIR</code>	Script root for <code>require</code> / <code>import</code> resolution (same semantics as vanilla <code>sercon</code>).
<code>-v</code>	Verbose engine tracing to <code>stderr</code> .

Behavioural deltas vs vanilla `sercon` :

- **Access log to `stderr`.** Each request emits one line in the format `ts remote method path status dur_us`, e.g.
`2026-05-29T10:23:14Z 127.0.0.1:54321 GET /users/42 200 1843µs`. Websocket upgrades log only the upgrade line; per-frame traffic is suppressed.
- **Readiness signal to `stdout`.** When each listener calls back into the binding with `bound`, the binding prints `READY listening on tcp/0.0.0.0:8080` to `stdout`. Multiple listeners produce one `READY` line each. Supervisors that read-line on startup (`systemd-notify`, `docker-compose` healthchecks) can wait on it.
- **Graceful shutdown.** `SIGTERM` / `SIGINT` triggers `srv.close()` on every active listener concurrently, then waits up to `--shutdown-timeout` for the script to drain. Clean exit is exit code `0`. Past the deadline, `loop.Terminate()` fires and the process exits with the usual classification.
- **No `--watch`.** Re-running while a listener is bound would race port rebinding; `sercon serve --watch ...` exits with a usage error. Use vanilla `sercon --watch script.ts` for the hot-reload development loop.

```
sercon serve examples/scripts/server-http.ts
sercon serve --port-override 9090 server.ts
sercon serve --shutdown-timeout 10s server.ts
```

5. Reserved globals (script surface)

`sercon` scripts get twelve reserved top-level globals. Use them directly — there is no enclosing namespace. The full per-binding reference is the generated `examples/scripts/sercon.d.ts` this section is the at-a-glance prose.

Reserved names: `runtime`, `crypto`, `text`, `codec`, `fs`, `net`, `db`, `server`, `services`, `tui`, `image`, `web`. User code can shadow these with a local `let` / `const` / `var` per normal JavaScript scoping — `sercon` does not intervene at runtime. `server` (added in v0.10.0) covers inbound listeners — see section 6 for the long-form treatment.

Deterministic key order. Objects returned from bindings enumerate their keys in a stable order (declaration order for fixed-shape results; source / column / insertion order for dynamic ones), so `JSON.stringify(result)` is byte-stable run-to-run — safe to hash for canonical serialization (payment signing, webhook signatures). The lone exception is `text.jq`, whose underlying engine discards key order internally.

console (browser/Node compatibility)

Beyond the twelve reserved globals, sercon provides a `console` global so scripts pasted from a browser or Node run unchanged:

Method	Stream
<code>console.log</code> / <code>console.info</code> / <code>console.debug</code>	<code>stdout</code>
<code>console.warn</code> / <code>console.error</code>	<code>stderr</code>

Each prints one space-joined line — clean and prefix-free (no timestamp), unlike the goja default console. Primitives print raw; objects and arrays render as JSON

(`console.log({a:1})` → `{ "a": 1 }`), with circular references falling back to `[object Object]` rather than throwing. The CLI disables the engine’s built-in console so this stream-correct shim is the only `console` a script sees. `runtime.log` is the native equivalent of `console.log` (`stdout`) and formats the same way. Library embedders get the `goja_nodejs` console by default instead; toggle it with `Options.DisableConsole`.

Migration from v0.8.0

v0.8.0	v0.9.0
<code>api.runtime.log</code>	<code>runtime.log</code>
<code>api.crypto.hash.sha256</code>	<code>crypto.hash.sha256</code>
<code>api.text.str.trim</code>	<code>text.str.trim</code>
<code>api.format.barcode.encode</code>	<code>codec.barcode.encode</code>
<code>api.fs.path.basename</code>	<code>fs.path.basename</code>
<code>api.net.http.get</code>	<code>net.http.get</code>
<code>api.db.sqlite.open</code>	<code>db.sqlite.open</code>
<code>api.tools.exec.shell</code>	<code>services.exec.shell</code>
<code>api.tools.git.commit</code>	<code>services.git.commit</code>
<code>api.ui.tui.layout</code>	<code>tui.layout</code>
<code>Sercon.argv</code>	<code>runtime.argv</code>

The `api` global no longer exists; reading it throws a `ReferenceError`. The `Sercon` global is also gone.

A mechanical `sed` pass over your scripts handles the prefix drop and the three renames — see the `CHANGELOG.md` [0.9.0] “Changed” entry for a sketch.

runtime

Script-host scaffolding. Members:

- `runtime.log(...args)` — print a space-joined line (primitives raw, objects/arrays as JSON).
- `runtime.assert.equal(actual, expected, msg?)` / `runtime.assert.ok(cond, msg?)` — throw on mismatch / falsy. `assert.equal` is **deep**: distinct objects/arrays with identical contents compare equal (structural, recursive — key order irrelevant).
- `runtime.time.nowMs()` — current wall-clock in milliseconds since the Unix epoch (host clock; never throws).

- `runtime.time.sleep(ms)` — Promise that resolves after `ms` on the event loop. It is **cancellable**: if the run hits its deadline or is cancelled before the delay elapses, the pending sleep rejects (the underlying context is cancelled) rather than holding the loop open.
- `runtime.time.format(unixMs, layout, tz?)` — render a Unix-ms timestamp through **strftime-style** tokens (not Go's reference layout). Supported tokens: `%Y %y %m %d %H %M %S %T %F %j %A %a %B %b %z %Z` and `%%` (a literal percent); unknown `%X` tokens pass through verbatim. `tz` is an optional IANA zone name (e.g. "Europe/Stockholm", "UTC"); it defaults to the host's local zone and **throws** if the name is not loadable. For example `runtime.time.format(runtime.time.nowMs(), "%F %T", "UTC")` yields "2026-06-24 13:45:09".
- `runtime.env.get(name)` — process environment variable; returns `undefined` when unset (never throws).
- `runtime.setDeadline(ms)` / `runtime.clearDeadline()` / `runtime.getDeadline()` — control the running script's own wall-clock kill deadline at runtime. This is the **same** deadline the `-timeout` flag sets, *not* the JS global `setTimeout` (which merely schedules a callback and does not move the run's kill timer).
 - `runtime.setDeadline(ms)` — move the kill deadline to **now + ms**, replacing any prior deadline. `ms <= 0` disables it (identical to `clearDeadline()`). Use it to extend a long-running task or to add a deadline to a run started with `-timeout 0`. Applies immediately to the in-flight run; never throws (a non-numeric argument coerces to `0`).
 - `runtime.clearDeadline()` — remove the deadline entirely (equivalent to `-timeout 0 / setDeadline(0)`); the script then runs without a timeout.
 - `runtime.getDeadline()` — milliseconds remaining until the kill deadline (`number`, `>= 0`), or `null` when no deadline is active (disabled, or started with `-timeout 0`).
- `runtime.argv` — `[programName, scriptPath, ...userArgs]`. User args come from the CLI args after a `--` separator. The layout mirrors Node/Bun: `argv[0]` is "sercon", `argv[1]` is the path of the currently-executing script (so each script in a multi-arg invocation sees its own path), and `argv.slice(2)` is just your args. The property is always present; with no `--` it is just `[programName, scriptPath]`.
- **runtime.secrets** — read/write string credentials in the OS keystore (macOS Keychain, Linux Secret Service / libsecret, Windows Credential Manager). Pure-Go (no cgo) via `zalando/go-keyring`.
 - `runtime.secrets.available` — `boolean`, advisory: is a keystore backend reachable this run? Gate on it to self-skip on headless boxes.
 - `runtime.secrets.get(name, account)` — `Promise<string | null>` `null` when absent.
 - `runtime.secrets.set(name, account, secret)` — `Promise<void>`.
 - `runtime.secrets.delete(name, account)` — `Promise<boolean>` (`true` if removed).

All operations are confined to a **prefix namespace**: the keystore service is `PREFIX + name`, so a script can neither read nor clobber secrets outside it. `PREFIX` resolves from

`--secrets-prefix > SERCON_SECRETS_PREFIX > the default sercon/`. Store a secret once with your OS tools, e.g. macOS

`security add-generic-password -s 'sercon/devshop' -a tess -w`, Linux

`secret-tool store --label='sercon devshop' service 'sercon/devshop' account tess`.

Operations are async and bounded by a 10s timeout (a blocking macOS consent prompt rejects rather than hangs).

- **runtime.clipboard** — read/write the host OS system clipboard text. An external-CLI fallback (no pure-Go, no-cgo library covers every platform).
 - `runtime.clipboard.available` — `boolean`, cheap advisory (no clipboard access): `true` when a backend is on PATH. Gate on it to self-skip on headless boxes.
 - `runtime.clipboard.read()` — `Promise<string>` the clipboard text (`""` when empty).
 - `runtime.clipboard.write(text)` — `Promise<void>` replace the clipboard text.
 - `runtime.clipboard.imageAvailable` — `boolean`, cheap advisory: `true` when a PNG image backend is usable on this host. Gated separately from `available` (text).
 - `runtime.clipboard.readImage()` — `Promise<Uint8Array | null>` the clipboard image as PNG bytes, or `null` when no image is present.
 - `runtime.clipboard.writeImage(png)` — `Promise<void>` set the clipboard image from PNG bytes (validated by signature; non-PNG rejects).

Backends: macOS `pbcopy` / `pbpaste`, Linux `wl-clipboard` (Wayland) or `xclip` / `xsel` (X11), Windows `clip` + PowerShell `Get-Clipboard`. When none is installed the calls throw a clean error and `available` is `false` the binary itself stays fully functional. Image is PNG-only and feature-detected separately (`imageAvailable`); macOS image read needs `pngpaste`, Linux needs `wl-clipboard` or `xclip` (not `xsel`). RTF/HTML/non-PNG are out of scope.

```
runtime.log("hello");
runtime.assert.equal(2 + 2, 4);
const args = runtime.argv.slice(2); // post-`--` user args
```

crypto

Hashing, JWT, and age/PGP encryption. Three sub-namespaces:

crypto.hash.*

`md5`, `sha1`, `sha256`, `sha384`, `sha512`, `sha3_256`, `sha3_512`, `blake3`, `crc32`. Each takes a single `string` (interpreted as a UTF-8 byte sequence) and returns a lowercase hex digest. Digest widths: MD5 32, SHA-1 40, SHA-256/SHA3-256/BLAKE3 64, SHA-384 96, SHA-512/SHA3-512 128. `crc32` is the IEEE polynomial, zero-padded to 8 hex chars. All are pure and deterministic. A missing/ `null` / `undefined` argument throws a `TypeError`. MD5 and SHA-1 are exposed for legacy fingerprinting only — do not use them for security. BLAKE3 is lukechampion.com/blake3 (256-bit output); the `sha3_*` underscore naming matches `recon`'s binding.

crypto.jwt.*

`sign(claims, secret, opts?)`, `view(token)`, `validate(token, secret, opts?)`.

The `secret` argument is polymorphic and the same across `sign/validate`: raw bytes for HMAC algorithms (`HS256` / `HS384` / `HS512`); a PEM-encoded key (`-----BEGIN ...`) for the asymmetric families (`RS*`, `PS*`, `ES*`, `EdDSA`); or a JWK JSON object (whose `kt` selects the key type). `sign` uses the private key, `validate` the public key (a certificate is accepted). `opts.algorithm` is case-insensitive and defaults to `HS256`.

`sign` does **not** synthesise reserved claims — set `iat` / `exp` / `nbf` yourself if you want them; everything in `claims` is passed straight through as `MapClaims`. `view` decodes the header and payload **without verifying** the signature (handy for debugging an auth flow), preserving

object key order and returning the raw signature segment alongside. `validate` verifies the signature plus the standard time claims (`exp` / `nbf` — `iat` is not checked, per RFC 7519 §4.1.6) and, when `opts.audience` / `opts.issuer` are set, those claims too; pinning `opts.algorithm` activates the algorithm-confusion guard. It resolves `{ valid: true, claims }` or `{ valid: false, reason }` — cryptographic and claim failures do *not* throw; only structural/wiring errors (empty token/secret, malformed token, key shape mismatching the token's algorithm) throw.

`crypto.encrypt.*`

```
keygen(), keygenPgp(opts?), encrypt(data, recipients, opts?),  
decrypt(ciphertext, identities),  
rekey(ciphertext, oldIdentities, newRecipients, opts?), detectBackend(input).
```

Two interoperable backends, auto-dispatched on key format — you never name the backend explicitly. `age (filippo.io/age)` keys are bech32: `age1...` recipients (public) and `AGE-SECRET-KEY-1...` identities (private); `age` also accepts SSH public keys as recipients. PGP keys are ASCII-armored `-----BEGIN PGP PUBLIC/PRIVATE KEY BLOCK-----` blocks. `keygen` mints a fresh `age X25519` keypair; `keygenPgp` mints an `RSA-2048` PGP keypair (`opts.name` / `opts.email` populate the user ID). `encrypt` / `decrypt` route to `age` or PGP by inspecting the supplied keys (and, for `decrypt`, the ciphertext armor) — the two backends cannot be mixed in a single call.

`encrypt` accepts `string` / `Uint8Array` / `ArrayBuffer` data and one or many recipients; `age` output is binary unless `opts.armor` requests ASCII armor, PGP output is always armored. Multi-recipient ciphertext can be opened by any listed identity. `decrypt` returns the plaintext as a `Uint8Array` (decode with `new TextDecoder().decode(...)`). `rekey` is an `age`-only `decrypt+re-encrypt` that rotates a ciphertext to a new recipient set without exposing plaintext to JS; output armor defaults to the input's and `opts.armor` overrides. `detectBackend` is pure prefix matching (no parsing, no I/O) and never throws — it returns `{ backend: "age" | "pgp" | "unknown", kind?: "public" | "private" }`. The `encrypt/decrypt/rekey` calls throw on type/shape errors: wrong key kind (public where private is expected or vice versa), empty key lists, an unsupported data type, or no identity matching the ciphertext.

`text`

String, regex, charset, and data manipulation. The `str.*` and `preg.*` families are **synchronous**; `charset.*`, `jq.*`, and `diff.*` return **Promises** (await them). Members:

- `text.str.*` — synchronous string helpers: `trim`, `ltrim`, `rtrim`, `reverse`, `stripHtml`, `nl2br`, `br2nl`, `base64Encode`, `base64Decode`, `urlEncode`, `urlDecode`, `htmlEntityDecode`, `pad` / `lpad` / `rpadd`, `sprintf`, `printf`, `normalizeNewlines`.
 - `trim` / `ltrim` / `rtrim` take a PHP-style **cutset** mask, not a prefix: every character in the mask is stripped (default mask is the whitespace set `" \t\n\r\v\f"`). `reverse` and the `pad` family count **runes**, not bytes, so multi-byte text stays intact (`reverse("café")` → `"éfac"`). `pad`'s side is `"right"` (default), `"left"`, or `"both"` `"both"` splits the deficit with the extra rune going right.
 - `base64Encode` / `base64Decode` are the **standard** RFC 4648 alphabet with padding — URL-safe input (`-` / `_`) is *not* auto-detected and throws; translate it to the standard alphabet first.

- `urlencode / urlDecode` use `application/x-www-form-urlencoded` rules (space $\leftrightarrow$ +); for path segments use `goja's built-in encodeURIComponent`. `sprintf / printf` use **Go's** `fmt` verbs (`%s`, `%d`, `%.2f`, `%v`, `%q`, ...), not PHP's; `printf` writes straight to `stdout` and returns nothing. `normalizeNewlines` rewrites any mix of CRLF/CR/LF to `"lf"` (default), `"crlf"`, or `"cr"`.
- `text.preg.*` and `text.preg2.*` — two distinct regex engines sharing a PHP-style `/pattern/flags` syntax (forward-slash delimiter only) and a `{ match, groups, index }` result shape. **preg** runs Go's **RE2**: linear-time, no catastrophic backtracking, but **no** lookahead or backreferences; flags `i / m / s` only (`u / U / x` are rejected). **preg2** runs the .NET -flavoured `regexp2` engine: lookahead, lookbehind, backreferences, and the extra `x` flag — at the cost of a backtracking engine, so guard untrusted input with a timeout. `match` returns the first hit or `null` `matchAll` returns an array; `replace` substitutes every match. Replacement templates use the engine's `$1 / ${1}` group syntax — PHP's `\1` backref form is *not* translated.
- `text.charset.detect(data) / decode(data, charset) / encode(text, charset)` — **async** character-set detection and conversion over a `string`, `Uint8Array`, or `ArrayBuffer`. `detect` (`saintfish/chardet`) resolves `{ charset, confidence, language?, candidates }` where `confidence` is a 0–100 score and `candidates` ranks every guess. `decode / encode` map between a named WHATWG/HTML5 encoding (UTF-8, ISO-8859-1, Windows-1252, Shift_JIS, GBK, ...) and UTF-8; `encode` has **no lossy fallback** — a character with no representation in the target encoding rejects rather than being silently dropped.
- `text.jq.query(json, expr) / queryAll(json, expr)` — **async** jq filtering (`itchyny/gojq`) against any JSON-like JS value. `query` resolves the **first** value the filter emits (or `null` when it emits nothing, e.g. an optional path `.a.b?` that misses); `queryAll` drains the whole stream into an array.
- `text.diff.compare(a, b, opts?)` — **async** unified diff between two text or byte inputs. `opts.context` (default 3) sets the surrounding lines; `opts.fromFile / opts.toFile` (default `"a" / "b"`) label the headers. Resolves `{ identical, binary, added, removed, diff, format: "unified" }` inputs with a NUL byte in their first 8 KB are reported as `binary: true` with an empty `diff`.

codec

Binary-format codecs (was `format` in v0.8.0). Members:

- `codec.compression.{algos, compress, decompress}` — `gzip`, `deflate`, `zlib`, `bzip2`, `zstd`, `brotli`, `lz4`, `xz`, `snappy`.
- `codec.barcode.{formats, decodableFormats, encode, decode}` — `QR`, `DataMatrix`, `Aztec`, `PDF417`, `Code128`, `Code39`, `Codabar`, `EAN-13/8`, `UPC-A`, `ITF`. `encode` returns PNG bytes; `decode` reads image bytes.
- `codec.checkdigit.{algos, validate, compute, inspect}` — `Luhn`, `ISBN-10/13`, `EAN-13/8`, `UPC-A`.
- `codec.php.{serialize, unserialize, varExport, parseVarExport, varDump, parseVarDump}` — read and write PHP data dumps.
- `codec.perl.{dumper, parseDumper}` — read and write Perl `Data::Dumper` output.
- `codec.xml.encode(value, opts?) / codec.xml.decode(xml)` — `value` $\leftrightarrow$ XML. Attributes are `@`-prefixed keys, text is `#text`, other keys are child elements, and an array value

becomes repeated sibling elements; a text-only element collapses to a bare string and an empty/self-closing element to `null`. The value must be a single-key object naming the root (or pass `opts.rootName`). `opts.indent` pretty-prints and `opts.declaration` prepends the XML declaration. Object key order and namespace prefixes are preserved. Decoded values are **strings** (XML is untyped — a number round-trips as its string form); surrounding whitespace is trimmed and mixed text/element ordering is not preserved. Cycles, mismatched tags, multiple roots, and malformed XML throw.

codec.compression — compress / decompress

Nine algorithms, exposed by name (case-insensitive). `codec.compression.algos()` returns them in registration order: `gzip`, `deflate`, `zlib`, `bzip2`, `zstd`, `brotli`, `lz4`, `xz`, `snappy`.

- `codec.compression.compress(algo, data): Promise<Uint8Array>` — both `compress` and `decompress` are **async** (they run off the loop via `PromisifyAsync`). `data` is a string (taken as its UTF-8 bytes), a `Uint8Array`, or an `ArrayBuffer`.
- `codec.compression.decompress(algo, data): Promise<Uint8Array>` — the inverse; pass the **same** algorithm name you compressed with. The frame format is algorithm-specific and not self-describing, so there is no auto-detection.

Notes on the implementations: `bzip2` is compressed via `dsnet/compress` (Go's standard `compress/bzip2` is decompression-only); `snappy` uses its one-shot block form (no streaming frame header), while `lz4` uses the LZ4 frame format (`pierrec/lz4` writer/reader); `gzip` / `zlib` / `deflate` are `stdlib`. Outputs surface to JS as `Uint8Array` decode bytes back to text with `new TextDecoder().decode(bytes)`. Round-trips are byte-faithful for the same algorithm; cross-algorithm or truncated input throws.

codec.barcode — encode / decode

- `codec.barcode.encode(format, data, opts?): Promise<Uint8Array>` — **async**; renders `data` into PNG bytes. Encodable formats (`codec.barcode.formats()`): `qr`, `datamatrix`, `aztec`, `pdf417`, `code128`, `code39`, `codabar`, `ean13`, `ean8`, `upca`.
- `codec.barcode.decode(data, format?): Promise<{ format, text }>` — **async**; reads PNG / JPEG / WebP image bytes via `goxing`. Decodable formats (`codec.barcode.decodableFormats()`): the encodable set minus `pdf417` (`goxing` has no PDF417 reader) plus `code93`, `upce`, and `itf`. With no `format` hint every reader is tried in priority order and the first hit wins; a hint runs only that reader.

Dimensions and the quiet zone. `opts.width` / `opts.height` default to 256×256 for the 2D codes (`qr`, `datamatrix`, `aztec`) and 400×120 for the 1D symbologies. `opts.quietZone` pads a white margin around the rendered bars: `true` uses 10 % of the width (minimum 10 px), a number uses that many pixels per side, and `false` / `0` / absent adds none. **EAN/UPC barcodes need a quiet zone to be decodable** — encode them with `quietZone: true` (or supply margin in the source image) before round-tripping through `decode`. Invalid payloads for a symbology (e.g. wrong EAN digit count) throw at encode time.

codec.checkdigit — validate / compute / inspect

Six algorithms (`codec.checkdigit.algos()`): `luhn`, `isbn10`, `isbn13`, `ean13`, `ean8`, `upca`. `isbn13` is an alias for `ean13` (identical check-digit math). Algorithm names are case-insensitive and trimmed.

- `codec.checkdigit.validate(algo, input): boolean` — does `input` (the full number *including* its trailing check digit) pass? Returns `false` — never throws — for the wrong length, non-digit characters, or an unknown algorithm.
- `codec.checkdigit.compute(algo, partial): string` — the single missing check digit for `partial` (the number *without* the check digit). Returns `'0' – '9'`, or `'X'` for ISBN-10 when the value is 10. **Throws** on an unknown algorithm, wrong length, or a non-digit. Fixed-length algorithms expect exactly length – 1 digits (12 for `ean13`, 9 for `isbn10`, etc.).
- `codec.checkdigit.inspect(algo, input): { algo, input, valid, given, computed }` — a diagnostic that combines the two: `given` is the input's last character, `computed` is the recalculated check digit (empty when the input is too short to split), and `valid` is `given === computed` (case-insensitive). Never throws; malformed input yields `valid: false` with an empty `computed`. ISBN-10 accepts a trailing `X` as the check digit.

codec.php — PHP dump formats

Three PHP textual formats, each with an encoder and a matching parser:

- `codec.php.serialize(value, opts?): string` — PHP `serialize()`. Emits the canonical wire form (`s: , i: , d: , b: , N; , a: , 0:`).
- `codec.php.unserialize(text, opts?): value` — inverse of `serialize`.
- `codec.php.varExport(value, opts?): string` — PHP `var_export()`: valid PHP source (`array ( ... ), \Cls::__set_state( ... )`).
- `codec.php.parseVarExport(text, opts?): value` — read a `var_export` literal back to a value.
- `codec.php.varDump(value, opts?): string` — PHP `var_dump()`: human-readable debug output with type/length annotations.
- `codec.php.parseVarDump(text, opts?): value` — **best-effort** read of `var_dump` output.

Array heuristic. A JS array, or a JS object whose keys are exactly the contiguous integer sequence `0..n-1`, maps to a PHP list array. Any other object maps to a PHP associative array (string/int keys preserved).

The `__class` sentinel. An object carrying a `__class` string key (configurable via `opts.classKey`, default `"__class"`) round-trips as a PHP *object* — `serialize` emits `0:...`, `var_export` emits `\Cls::__set_state(...)`, `var_dump` emits `object(Cls)#...`. The remaining keys become the object's properties. Decoding restores the `__class` key. Without the sentinel a value is treated as an array.

Shared references and cycles. `serialize/unserialize` model PHP's `r: / R:` reference markers: when the same object appears more than once, the decoder resolves both occurrences to the *same* JS object (a DAG is preserved). A true cycle (an object reachable from itself) **throws** on encode — there is no JS value that can represent it losslessly here.

`var_dump` parse is lossy. PHP's `var_dump` is a debug view, not a serialization format; `parseVarDump` is best-effort and **throws** when it hits something it cannot faithfully reconstruct: the `*RECURSION*` marker, a length-truncated string (its declared byte length disagrees with the bytes present), or visibility-annotated property names (`[ "x": "Cls":private ]`).

codec.perl — Perl `Data::Dumper`

- `codec.perl.dumper(value, opts?): string` — emit a `Data::Dumper`-style dump (`$VAR1 = ... ;`) with normalized indentation.

- `codec.perl.parseDumper(text, opts?): value` — read `Data::Dumper` output back to a value.

The JSON bool convention. Perl has no native boolean, so JSON modules represent one as a blessed scalar reference (e.g. `bless( do{\(my $o = 1)}, 'JSON::XS::Boolean' )`). `dumper` emits JS booleans in this form, and `parseDumper` decodes a blessed scalar ref in the JSON-bool family (`JSON::XS::Boolean`, `JSON::PP::Boolean`, ...) back to a JS boolean. A bare `1 / 0` stays a number. The class used on encode is `opts.perlBoolClass` (default `"JSON::XS::Boolean"`). Blessed hashes carry the `__class` sentinel, mirroring the PHP side. Cycles **throw**.

opts and shared caveats

All `codec.php.* / codec.perl.*` functions accept an optional `opts` bag:

- `classKey` (default `"__class"`) — the key used as the class sentinel.
- `perlBoolClass` (default `"JSON::XS::Boolean"`) — class for Perl booleans.
- `indent` — override the default 2-space indentation step (`varExport`, `dumper`).

Caveats shared with JSON: strings are UTF-8 and lengths are reported in **bytes** (multibyte chars count as more than one); JS has a single number type, so `int` and `float` collapse the same way `JSON.stringify` does (e.g. `1.0` may re-emit as `1`). Decoded objects preserve a stable key order, so re-encoding is byte-stable — safe for canonical-JSON / payment hashing.

fs

Filesystem operations:

- `fs.path.dirname(p) / fs.path.basename(p, suffix?)` — POSIX-style path splitting (forward slashes). `dirname` returns everything up to (not including) the final slash — `."` when there is no directory component, `"/` for a rooted single segment; trailing slashes are stripped first. `basename` returns the final segment and, when `suffix` is supplied and the segment ends with it (and is not equal to it), strips that suffix — handy for dropping a known extension. Both **throw a TypeError** if `p` is missing, `null`, or `undefined`. These are pure string transforms — they never touch disk. On Windows, normalise separators yourself before calling.
- `fs.archive.create(destPath, sources) / fs.archive.extract(archivePath, destDir, opts?)` — both **async** (return Promises). The archive format is **inferred from the file extension**: `.zip`, `.tar`, `.tar.gz`, and `.tgz` (a `.tgz` resolves to the `tar.gz` format); an unrecognised extension rejects. `create` takes a non-empty array of sources — a bare string uses the disk path as-is with its basename inside the archive, while an `{ path, name }` object overrides the in-archive name; directory sources are recursed (their basename becomes the archive subdirectory) and archive paths always use forward slashes. It resolves to `{ path, format, entries, bytes? }` (`entries` lists the regular files written; `bytes` is the final archive size when stat succeeds). `extract` creates `destDir` recursively if absent, applies **zip-slip / tar-slip protection** (entries that resolve outside `destDir` via an absolute path or a `..` component reject the whole call), and honours `opts.overwrite` — the default `false` opens targets with `O_EXCL` so an entry colliding with an existing file fails; pass `true` to clobber. It resolves to `{ path, format, dest, entries }`.

General file read/write (all async; paths used as given, no sandboxing — matching `image.save` and the WebDriver screenshot writers):

- `fs.writeText(path, text) → { path, bytes }` — UTF-8, truncating.
- `fs.writeBytes(path, data: Uint8Array) → { path, bytes }`.
- `fs.readText(path) → string` `fs.readBytes(path) → Uint8Array`.
- `fs.mkdir(path) → { path }` — creates parents (`mkdir -p`), idempotent.
- `fs.exists(path) → boolean` — never throws for a missing path.
- `fs.remove(path) → { path }` — file or directory tree; no error if absent.
- `fs.stat(path) → { size, isDir, modifiedMs }`.

`writeText / writeBytes` **do not create parent directories** — call `fs.mkdir` first (Node-like); reads and `stat` reject when the target is absent. See `examples/scripts/fs-report.ts` for a per-step screenshot report built on these.

net

Network clients and probes:

- `net.http.{get, post, request}` — fetch-style HTTP client; `request` accepts headers, body, timeout, retry, follow, and basic auth.
- `net.probe.{tcp, dns, tls, ntp, whois, ping, traceroute, smtp, wss}` — one-shot reachability / handshake probes. Each returns a structured result; failures surface as `Error` objects with details.
- `net.probe.traceroute(host, opts?)` — trace the network path: send probes with increasing TTL and report each responding router. **Needs root / CAP_NET_RAW** (intermediate hops appear only via ICMP `time-exceeded`). `opts.protocol` is `"icmp"` (default), `"udp"` (to an incrementing high port), or `"tcp"` (SYN via a TTL-limited connect); `port` sets the udp/tcp target; `maxHops` (30), `timeout ms` per probe (2000), and `probes per hop` (3) bound the trace. Resolves to one `{ ttl, address, rttMs, reached }` per hop (`address` is `null` for a timed-out hop). IPv4 only.
- `net.probe.ping` also accepts `mode: "udp"` — a UDP datagram to a (closed) port whose ICMP `port-unreachable` proves reachability (needs root / `CAP_NET_RAW`), alongside the existing `"tcp"` (default) and `"icmp"` modes.
- `net.netstatus.check(host)` — combined reachability summary.
- `net.email.*` — `spf`, `dmARC`, `mtaSts`, `tlsRpt`, `bimi`, and `all(domain)` which runs every probe in parallel and aggregates; plus `send({...})` — outbound SMTP sender (see §6.7).
- `net.browser.open()` — open a stateful HTTP session (cookie jar + replayed default headers); **not** an OS-browser launcher.
- `net.tcp.connect(host, port, opts?)` — open a TCP client socket.
- `net.udp.open(opts)` — open a UDP socket (connected or bound).
- `net.icmp.open(opts?)` — open a raw ICMP socket (needs privileges).
- `net.capture.{interfaces, open, openFile, toFile}` — packet capture (live + pcap file I/O), pure-Go gopacket.
- `net.raw.{open, tcp}` — craft & send raw IPv4 packets (TCP flags / UDP / arbitrary IP protocol) and receive replies (needs privileges).
- `net.load.http(opts)` — authorized HTTP load / resilience self-test (worker-pool, latency percentiles + error rate; public hosts need `confirm:true`).

HTTP client (`net.http`)

Three entry points, all returning a `Promise`. **A 4xx/5xx response never rejects** — it resolves with the status set; only transport-level failures (DNS, connection refused, TLS handshake) or a deadline reject.

- `net.http.get(url) / net.http.post(url, body?)` — quick one-liners with a **5-second** default timeout. Resolve to `{ status, body, bodyBytes }`: `body` is the response decoded as UTF-8, `bodyBytes` is the raw, undecoded `Uint8Array` (pair with `text.charset.decode` for ISO-8859-1 and other non-UTF-8 payloads). `post` sends `body` verbatim and sets **no** `Content-Type` — add one via `request` if the server needs it.
- `net.http.request(method, url, opts?)` — the full client. `opts` is `{ headers?, body?, timeout? (ms, default 30000), retry? (extra attempts, default 0), follow? (default true), username?, password? }`. `retry` applies **only** to transport errors and 5xx, with linear backoff capped at 1 s; a malformed method/URL rejects immediately and is not retried. `follow: false` stops at the first 3xx. Resolves to `{ status, ok, headers, body, bodyBytes, url }` — `ok` is `status` in `[200, 400)` `headers` is a lower-cased name→value map (last value wins); `url` is the final URL after redirects.

One-shot probes (`net.probe`)

Each `net.probe.*` helper performs a single handshake/lookup and resolves with a structured result. Connection-level failures reject; *findings* (an expired cert, a missing STARTTLS, an unreachable host) are returned as data.

- `net.probe.tcp(target, opts?)` — dial and report `{ host, port, ip, latencyMs }`. `opts.port` (default "80") supplies a port when `target` is a bare host; `opts.timeout ms` (default 5000).
- `net.probe.dns(host, opts?)` — resolve A / AAAA / MX / TXT / CNAME / NS. `opts.types` restricts the set (case-insensitive). **Always resolves** — per-type failures and empties are simply omitted from the result, so test presence with "mx" in `result`. `mx` entries are `{ preference, host }`.
- `net.probe.tls(target, opts?)` — open a TLS connection with **verification skipped** (probe-only) and return the leaf cert summary `{ cn, issuer, notBefore, notAfter, daysRemaining, dnsNames, serialNumber, fingerprintSha256 }`. The host is sent as SNI; expired or name-mismatched certs still report (that's the point). A bare host uses port 443.
- `net.probe.ntp(host, opts?)` — query an NTPv4 server (UDP 123) and report `{ serverTime, offsetMs, rttMs, stratum, referenceTime, rootDelayMs, rootDispersionMs }`.
- `net.probe.whois(domain, opts?)` — two-hop WHOIS via the IANA referral. `raw` is always the full response text; `domain / registrar` are best-effort parsed and omitted for TLDs the parser doesn't recognise. A parse failure is non-fatal (only `raw` comes back); a wire failure rejects.
- `net.probe.smtp(host, opts?)` — SMTP **capability** probe (EHLO only, no mail). Returns `{ banner, ehloDomain, extensions, starttls, authMechanisms, sizeLimit }`. A server that omits STARTTLS/AUTH reports `false / []` — a finding, not an error. `opts.port` default "25", `opts.ehloName` default "localhost".

- `net.probe.wss(url, opts?)` — WebSocket handshake probe. Returns `{ connected, subprotocol, status, handshakeMs, pingMs }` and closes immediately. `opts.ping` (default `true`) toggles the RTT measurement; `pingMs` is `-1` when skipped or unanswered.

`net.netstatus.check(host, opts?)` rolls DNS / TCP / TLS / HTTP into **one concurrent** sweep, resolving to `{ host, port, elapsedMs, reachable, dns, tcp, tls, http }`. Each sub-probe is `{ ok, ..., error? }` — sub-failures are **captured as data**, never thrown; only an empty `host` rejects. `reachable` is `dns.ok && tcp.ok` (TLS/HTTP are reported but don't gate it). `opts.port` defaults to `"443"`.

Email-authentication probes (`net.email`)

DNS-based deliverability / anti-spoofing posture checks. Each takes a `domain` and resolves to a discriminated union — `{ present: false }` when the record is absent, or `{ present: true, record, ... }` with the parsed fields when found. **They reject only on a DNS error other than NXDOMAIN**; a missing record is the normal `{ present: false }` outcome, not a throw.

- `net.email.spf(domain)` → `{ record, mechanisms, allPolicy }` from the apex TXT (`v=spf1`); `allPolicy` summarises the trailing `all` mechanism (`pass` / `fail` / `softfail` / `neutral`).
- `net.email.dmarc(domain)` → the `_dmarc.<domain> v=DMARC1` record with `{ tags, policy, subdomain, percent, rua, ruf }`.
- `net.email.mtaSts(domain)` → the `_mta-sts.<domain> TXT` **plus** the fetched well-known policy file (`{ mode, mx, maxAge }`); a policy-fetch failure surfaces in `policyError`, not as a throw.
- `net.email.tlsRpt(domain)` → the `_smtp._tls.<domain> v=TLSRPTv1` record with `{ tags, rua }`.
- `net.email.bimi(domain, opts?)` → the `<selector>._bimi.<domain>` record with `{ l, a }` (logo URL + VMC assertion); `opts.selector` defaults to `"default"` and is echoed back in the result.
- `net.email.all(domain)` runs all five **in parallel** and aggregates under `{ domain, spf, dmarc, mtaSts, tlsRpt, bimi }` a single probe's failure surfaces as `{ error }` under its key rather than failing the aggregate, so this call always resolves.

Stateful HTTP session (`net.browser`)

`net.browser.open()` (no arguments) opens a **stateful HTTP session** — an `http.Client` with an automatic, public-suffix-scoped **cookie jar** and a set of default headers replayed on every request, the way a browser carries a session across page loads. It does **not** launch the OS web browser. It resolves to a handle:

- `setUserAgent(ua)` / `setHeader(name, value)` — register default headers replayed on subsequent requests (synchronous).
- `get(url)` / `post(url, body?)` — issue a request through the session; resolve to the same `{ status, ok, headers, body, url }` shape as `net.http.request`. The jar is updated automatically, so a `login post` followed by a `get` replays the session cookie with no manual handling.
- `cookies(url)` — list the jar's cookies scoped to `url` as `{ name, value }[]`, e.g. to confirm a login set the expected cookie.

`open()` rejects only if the cookie jar can't be created; `get` / `post` / `cookies` reject on an empty/unparseable URL or a transport error (4xx/5xx resolve normally).

Raw sockets (`net.tcp` / `net.udp` / `net.icmp`) are long-lived, bidirectional client sockets with a *push* / *callback* read model — unlike the one-shot `net.probe.*` helpers, they stay open and deliver inbound data through callbacks until you `close()` them. Each constructor returns a `Promise` for a handle object:

- `net.tcp.connect(host, port, opts?)` — dial a TCP peer. `opts` { `timeout?` (ms, default 10000), `readBuffer?` (inbound channel capacity, default 64) }. The handle has `write(data)` (`string` → UTF-8 bytes, `Uint8Array` → raw bytes; returns a `Promise`), the `remote` / `local` address strings, plus the shared callbacks below. Inbound chunks arrive via `onData`.
- `net.udp.open(opts)` — two modes selected by `opts`. **Connected** { `host`, `port`, `readBuffer?` } resolves the peer up front and exposes `send(data)`. **Bound** { `bind`: "127.0.0.1:0", `readBuffer?` } listens on a local address and exposes `sendTo(data, host, port)` its inbound events also carry `address` / `port` of the sender. Either way the handle has `local` (the bound address — read the chosen port back from it when you bind to `:0`) and delivers datagrams via `onMessage`.
- `net.icmp.open(opts?)` — open a raw ICMP socket. **Requires root** / `CAP_NET_RAW`; without the privilege `open()` rejects with an error naming that requirement. `opts` { `network?`: "ip4" | "ip6" (default "ip4"), `readBuffer?` }. `send(opts)` writes a message in one of two modes: **Echo mode** { `to`, `type?`, `code?`, `id?`, `seq?`, `payload?` } builds an Echo-shaped body (`type` defaults to the network's echo request), or **raw mode** { `to`, `type`, `code?`, `body` } marshals `body` (`Uint8Array` | `string`) verbatim — for hand-built non-Echo messages such as a crafted destination-unreachable. In raw mode `type` is required and `body` is mutually exclusive with `id` / `seq` / `payload`. The handle has `network` / `local` inbound events carry `address` / `type` / `code` and arrive via `onMessage`.

All three handles share the same callback surface and lifecycle:

- `onData(cb)` (TCP) / `onMessage(cb)` (UDP, ICMP) — register a listener for inbound data. The event object carries `bytes` (a `Uint8Array`) and `text` (the bytes decoded as UTF-8); UDP-bound / ICMP events add the sender metadata noted above.
- `onClose(cb)` — fires once when the socket closes (locally or remotely).
- `onError(cb)` — fires on a read / connection error (benign teardown closes are reported through `onClose`, not `onError`).
- `close()` — shut the socket down; returns a `Promise`.

The socket keeps the engine's event loop alive while open (via the internal `HoldRun` refcount), so a script that only listens won't exit until it `close()`s. `sercon` scripts can't *bind* a listening TCP/UDP server socket yet — these are client sockets — so a TCP example needs a peer to dial; a UDP loopback pair is fully self-contained.

Packet capture (`net.capture`)

Packet capture and pcap file I/O, powered by **pure-Go gopacket** (no `libpcap`, no `cgo`). Five bindings:

- `net.capture.interfaces()` — **synchronous**; returns an array of `{ name, addresses: string[], up, loopback }` for the host's network interfaces. Pure-Go (`net.Interfaces`); no privileges, all platforms.
- `net.capture.routes()` — **synchronous**; returns an array of `{ destination, gateway, interface, family, metric }` for the host's IP routing table. `destination` is a CIDR (`0.0.0.0/0` / `::/0` are the default routes); `gateway` is the next-hop IP or "" for a directly-connected route; `family` is "ip" or "ip6" `metric` is 0 where the platform doesn't report one. Pure-Go, unprivileged: Linux parses `/proc/net/route` + `/proc/net/ipv6_route` macOS/BSD read the routing socket via `golang.org/x/net/route`. **Windows is unsupported** (throws).
- `net.capture.open({ iface, promisc?, snaplen?, filter? }, pkt => {...})` — start a **live** capture on `iface`. Resolves to a handle `{ iface, link, close() }` the per-frame handler receives a decoded packet object (see below). `promisc` defaults `true` `snaplen` defaults `262144`. `filter` is an optional tcpdump-like expression string (see **Filtering** below).
Linux + macOS only — Linux uses `AF_PACKET`, macOS uses `BPF` (`/dev/bpf`); **Windows is unsupported** and `open()` rejects there. Requires **root** / **CAP_NET_RAW (Linux)** or **/dev/bpf read access (macOS)**; without the privilege `open()` rejects naming the requirement. `close()` returns `Promise<void>`.
- `net.capture.openFile(path, pkt => {...}, opts?)` — read a `.pcap` / `.pcapng` file (format auto-detected from the magic bytes), calling the handler once per decoded packet. Returns a `Promise` that resolves at EOF. Offline; no privileges. `opts` is an optional trailing argument `{ filter? }` (the two-argument form still works); `filter` is the same tcpdump-like expression string as `open` (see **Filtering** below).
- `net.capture.toFile(path, { linkType?, snaplen? })` — open/create a `.pcap` file and return `{ write(bytes, { ts? }), close() }`. `write` appends a raw frame (a `Uint8Array`); `opts.ts` (ms) overrides the per-frame timestamp (defaults to now). `close()` flushes and returns `Promise<void>`. `linkType` defaults to Ethernet, `snaplen` to `262144`. Offline; no privileges.

The **decoded packet object** passed to the `open` / `openFile` handlers:

```
{
  ts, length, captureLength, link,      // always present
  eth?: { src, dst, type },
  vlan?: { id, priority, drop, type }, // 802.1Q tag; `type` is the inner
ethertype
  arp?: { operation, senderMac, senderIp, targetMac, targetIp }, // operation:
"request"|"reply"
  ip?: { version, src, dst, protocol, ttl },
  tcp?: { srcPort, dstPort, seq, ack, window, checksum,
         flags: { syn, ack, fin, rst, psh, urg },
         options?: { mss?, windowScale?, sackPermitted?, timestamps?: { val,
ecr } } },
  udp?: { srcPort, dstPort, length },
  icmp?: { type, code },
  dns?: { id, qr, opcode, rcode,
         questions: { name, type }[],
         answers: { name, type, data }[] }, // data rendered by type (A/
AAAA→IP, CNAME→name, ...)
  payload?: Uint8Array,                // application-layer bytes, if any
```

```
bytes:    Uint8Array,           // the full raw frame (always present)
}
```

Layer keys (`eth` / `vlan` / `arp` / `ip` / `tcp` / `udp` / `icmp` / `dns` / `payload`) are present **only when that layer decodes** — a truncated or unrecognised frame still yields the always-present `ts` / `length` / `captureLength` / `link` / `bytes` . `tcp.options` is present only when the segment carries recognised options.

Filtering. Both `open` and `openFile` accept an optional `filter` string — a **subset of tcpdump syntax**. Supported grammar:

- protocols: `tcp` , `udp` , `icmp` , `ip` , `ip6`
- hosts: `host X` , `src host X` , `dst host X` (IPv4 or IPv6 address);
- networks: `net X/Y` , `src net X/Y` , `dst net X/Y` (IPv4 or IPv6 CIDR prefix, e.g. `net 10.0.0.0/8` or `net 2001:db8::/32`);
- ports: `port N` , `src port N` , `dst port N`
- port ranges: `portrange A-B` , `src portrange A-B` , `dst portrange A-B` (inclusive, e.g. `portrange 8000-8100`);
- boolean composition: `and` , `or` , `not` , and parentheses;
- **implicit-and** between juxtaposed primaries — `tcp port 80` is exactly `tcp and port 80` .

Example: `net.capture.open({ iface: "en0", filter: "tcp and port 80" }, pkt => {...})` , or `net.capture.openFile("cap.pcap", pkt => {...}, { filter: "udp or icmp" })` .

The filter is a **userspace, post-decode** predicate — it is **not** compiled to a kernel BPF program. It runs on each frame *after* `gopacket` has decoded it, so it saves the JS-callback dispatch and object-conversion cost for non-matching packets but does **not** avoid the kernel→userspace copy (a real kernel filter would). A malformed expression makes `open` / `openFile` **reject**.

Limitations. No live capture on Windows. The `filter` grammar is the `tcpdump` subset above; for anything beyond it, filter inside the callback on the decoded fields. Common-layer decode only (Ethernet / IPv4 / IPv6 / TCP / UDP / ICMP); exotic protocols surface only as raw `bytes` . At high packet rates frames backpressure and drop at the kernel, exactly like `tcpdump` .

The `pcap` file API round-trips: write raw frames with `toFile` , read them back decoded with `openFile` , no privileges required (see `examples/scripts/capture-file.ts`).

Raw packets (`net.raw`)

Craft and send raw IPv4 packets and receive the replies — for SYN scans, RST/port-state inference, custom-TTL probes, and exotic-packet testing. **Linux + macOS only; needs root / CAP_NET_RAW** (Windows rejects). Two bindings:

- `net.raw.open({ iface?, filter?, readBuffer? })` — open a raw packet engine.
`send(spec | Uint8Array)` crafts and fires a packet: a structured `spec`

```
{ dst, dstPort?, srcPort?, src?, proto?: "tcp"|"udp"|"ip", protocol?, flags?:
  string[], seq?, ack?, window?, ttl?, ipId?, payload? }
```

 (full source-IP / TTL / IP-ID control; `sercon` computes checksums), or a `Uint8Array` holding a complete IPv4 packet sent verbatim. `onPacket(cb)` delivers decoded replies (the **same decoded packet object** as `net.capture`); `onClose` / `onError` / `close()` (returns a `Promise`) round out the handle. `iface` defaults to the auto-detected default-route

interface — set it explicitly on multi-homed hosts; `filter` is the same tcpdump-like expression as `net.capture`. Default `flags ["SYN"]`, `ttl 64`, `window 65535`, `src` = the egress interface IP, `srcPort` = a random high port.

- `net.raw.tcp(host, port, opts?)` — one-shot: send a single crafted TCP segment (default a SYN) and resolve with the first reply correlated by the 4-tuple, or `null` on timeout. SYN → SYN/ACK = open, RST = closed, null = filtered/no answer. `opts`
`{ flags?, srcPort?, src?, seq?, ttl?, payload?, timeout? (ms, default 2000), iface? }`

> **Kernel-RST caveat:** because the host kernel does not own the crafted > connection, it may answer the peer's SYN/ACK with its own RST. The probe > still captures the reply; only the target's connection state is perturbed. > To probe silently, add a host firewall rule dropping the outbound RST > (`iptables -A OUTPUT -p tcp --tcp-flags RST RST -j DROP`, or the `pf` > equivalent). `sercon` does not modify your firewall.

HTTP load / resilience self-test (`net.load`)

`net.load.http(opts)` is an **authorized HTTP load / resilience self-test harness**: a worker pool drives a target you operate at a chosen concurrency, and the call resolves with a latency/error report. It is a **defensive** tool — for staging / CI / lab self-tests — not a flooding weapon: it is a plain HTTP client loop with **no raw packets, source spoofing, amplification, or randomized-target fan-out**.

> **Authorized-use guardrail.** A target whose host resolves to loopback, > a private/ULA range, link-local, the unspecified address, or `localhost` > runs unconditionally (it's your own infra). Any **public** host is > **refused unless you pass `confirm: true`** — throwing > `net.load.http: refusing to load-test public host <h> without confirm:true` (authorized self-testing only) .>

`confirm: true` is the script author asserting they are authorized to > load-test that host. Concurrency is hard-capped at **1000** (values above > throw); the default is a conservative **10**.

```
net.load.http({
  url: string,           // http/https target (required)
  method?: string,      // default "GET"
  headers?: Record<string, string>,
  body?: string,
  concurrency?: number, // parallel workers; default 10, capped at 1000
  requests?: number,    // total requests to send ...
  duration?: number,    // ... OR run for this many ms (exactly one required)
  rps?: number,         // optional client-side rate cap (req/s; 0 =
unlimited)
  timeout?: number,     // per-request ms; default 10000
  confirm?: boolean,    // required true to target a public host
}): Promise<LoadReport>
```

- Provide **exactly one** of `requests` (> 0) or `duration` (> 0 ms); giving neither or both throws.
- A 4xx/5xx is a normal **recorded** response (it counts in `statusCounts`), not a thrown error; only transport failures (no response) count as `failed`. `errorRate` is `(failed + 5xx) / sent`.

The report:

```

interface LoadReport {
    target: string;
    method: string;
    concurrency: number;
    durationMs: number;           // wall-clock of the run
    sent: number;                 // requests started
    completed: number;           // got an HTTP response (any status)
    failed: number;               // transport errors / timeouts (no response)
    rps: number;                  // achieved throughput (completed / durationSec)
    errorRate: number;           // (failed + 5xx) / sent, 0..1
    latency: {                    // milliseconds, over completed requests
        min: number; mean: number; p50: number; p90: number;
        p95: number; p99: number; max: number;
    };
    statusCounts: Record<string, number>; // e.g. { "200": 990, "503": 10 }
    errors: Record<string, number>;       // transport error kind → count
}

```

Latency percentiles are nearest-rank over the completed-request latencies. Transport errors are bucketed by kind (`timeout` , `refused` , `dns` , `reset` , `canceled` , `error`). The run honours the engine's cancellation (Run end / `--timeout` aborts in-flight requests). See `examples/scripts/load.ts` .

db

Database / KV / directory clients:

- `db.sqlite.open(path)` — embedded SQLite (cgo-free driver).
- `db.postgres.open(dsn | opts)` — PostgreSQL via pure-Go `pgx` .
- `db.mysql.open(dsn | opts)` — MySQL / MariaDB via pure-Go `go-sql-driver` .
- `db.mssql.open(dsn | opts)` — Microsoft SQL Server via pure-Go `go-mssqldb` .
- `db.clickhouse.open(dsn | opts)` — ClickHouse via pure-Go `clickhouse-go v2` (`opts.secure` for TLS).
- `db.oracle.open(dsn | opts)` — Oracle via pure-Go `go-ora` (`opts.database` is the service name).
- `db.redis.open(addr, opts?)` — Redis client.
- `db.valkey.open(url, opts?)` — Valkey client (the RESP-compatible Redis fork; same { `do` , `ping` , `close` } handle, accepts `valkey://` / `valkeys://` as well as `redis://`).
- `db.memcached.open(addr)` — Memcached client.
- `db.ldap.open(url, opts?)` — LDAP client (search, bind).
- `db.dict.{define, match}` — local dictionary lookup / fuzzy match.

SQL engines (`sqlite` / `postgres` / `mysql` / `mssql`) share one handle: `open()` resolves to { `exec` , `query` , `queryValue` , `begin` , `prepare` , `close` } (transactions via `begin()` → { `exec` , `query` , `queryValue` , `commit` , `rollback` } prepared statements via `prepare(sql)` → { `exec` , `query` , `queryValue` , `close` }). The server engines accept either a driver **DSN string** or a connection **options object** ({ `host` , `port` , `user` , `password` , `database` } , plus `sslmode` for postgres); credentials in the assembled URL DSNs are percent-escaped. The connection is pinged on `open()` so a bad DSN or unreachable server fails there, not at the first query. Bind parameters are positional and passed after the SQL string — write your engine's placeholder syntax: `?` (sqlite, mysql), `$1` (postgres), `@p1` (mssql). All four drivers are pure

Go (no cgo). Scripts must `close()` the handle (and commit/rollback transactions, close prepared statements) — there is no GC finalizer.

Integration testing. The `db.*` tests above are unit-level. To exercise the networked bindings against real servers, point the `SERCON_TEST_*` env vars at a running database and run the Integration tests; unset vars skip, so `make test` / CI stay green without any server. `make test-integration` does it all against the dbplayground fleet: it brings the fleet up, runs the tests (postgres / mysql / mariadb / clickhouse / redis / valkey / memcached / ldap), and tears it down. Requires Docker + Compose v2.

services

Subprocess and external-CLI / service wrappers (was `tools` in v0.8.0). Members:

- `services.exec.shell(cmd, opts?)` — run a shell command; captures stdout/stderr or streams into a `tui` pane when `opts.pane` is set.
- `services.exec.stream(cmd, onLine, opts?)` — like `exec.shell` but streams the subprocess's output to `onLine(line, stream)` **line by line as it arrives** instead of buffering it. `cmd` and `opts` (`cwd` / `env` / `stdin` / `timeout`) match `exec.shell`, except `timeout` has **no default** (0 / absent = run until exit), since streaming targets long-running output. `stream` is `"stdout"` or `"stderr"`. Returns a `Promise` resolving to `{ exitCode, success, durationMs }` on exit; a non-zero exit resolves with `success: false`, while spawn failures and timeouts reject. Useful for processing large or incremental output without holding it all in memory.
- `services.exec.http(method, url, opts?)` — shell-level `curl` -style HTTP (separate from `net.http`, intended for raw protocol use).
- `services.git.*` — git porcelain wrappers (`clone`, `status`, `commit`, `log`, ...) invoking the local `git` binary.
- `services.gh.{authStatus, prList, repoView}` — thin `gh` CLI wrappers.
- `services.ai.{providers, send}` — provider-agnostic AI chat client.
- `services.agentBrowser.*` — headless-Chrome automation via the `agent-browser` CLI. See the subsection below.
- `services.webdriver.*` — W3C WebDriver client (via `tebeka/selenium`). See the subsection below.
- `services.typst.*` — Typst typesetting via the external `typst` CLI. See the subsection below.

services.agentBrowser

Bridge to the `agent-browser` headless-Chrome CLI. Every call shells out to `agent-browser --json --session \<name\> \<command\>` and returns the parsed JSON response.

Availability gate. `services.agentBrowser.available` is a sync boolean — check it before calling anything else. Every binding throws a clean error if the CLI is absent.

Launch model. `services.agentBrowser.launch(opts?)` is **synchronous**; no browser starts until the first command. It returns a handle whose I/O methods are all async (Promises). Sessions the script doesn't `close()` are best-effort closed when the Run ends.

Launch options (all optional):

Key	CLI flag	Description
-----	----------	-------------

session	--session	Name the session (auto-generated when omitted).
headed	--headed	Run with visible browser window.
profile	--profile	Browser profile directory.
proxy	--proxy	Proxy URL.
userAgent	--user-agent	Custom user-agent string.
device	--device	Emulate a device (e.g. "iPhone 12").
colorScheme	--color-scheme	"light" or "dark".
ignoreHttpsErrors	--ignore-https-errors	Skip TLS verification.
engine	--engine	Browser engine to use.
executablePath	--executable-path	Path to the browser binary.
enable	--enable	Enable a browser feature flag.
args	--args	Extra browser arguments.
timeout	(<i>sercon</i>)	Per-call subprocess timeout in ms (default 30 000, 0 disables). When the limit is reached the Promise rejects with a clear "timed out" message instead of hanging. Session <code>close()</code> is independently bounded at 10 s regardless of this value.

Response envelope. Every async method returns a parsed JSON object with the envelope `{ success: boolean, data: {...}, error: string | null }`. Drill into `.data` for the actual values (e.g. `data.title`, `data.value`, `data.visible`).

Handle method groups:

- **Navigation:** `open(url)`, `back()`, `forward()`, `reload()`, `wait(msOrSelector)`, `connect(portOrUrl)`
- **Interaction:** `click(sel)`, `dblclick(sel)`, `hover(sel)`, `focus(sel)`, `fill(sel, text)`, `type(sel, text)`, `press(key)`, `check(sel)`, `uncheck(sel)`, `select(sel, ...values)`, `scroll(dir, px?)`, `scrollIntoView(sel)`, `drag(src, dst)`, `upload(sel, files)`, `download(sel, path)`, `keyboard.{type, insertText}`, `mouse.{move, down, up, wheel}`
- **Inspection:** `get(what, sel?)`, `isVisible(sel)`, `isEnabled(sel)`, `isChecked(sel)`, `eval(jsCode)`, `snapshot(opts?)`, `console(opts?)`, `errors(opts?)`, `highlight(sel)`
- **Locators:** `find(locator, value, { action, text? })`, `locator(spec)` — returns a chainable handle bound to a `(locator, value) spec`.
- **Settings (Phase 2):** `set.viewport(w,h,scale?)`, `set.device(name)`, `set.geo(lat,lng)`, `set.offline(on?)`, `set.headers(obj)`, `set.credentials(user,pass)`, `set.media(scheme?,reducedMotion?)`
- **Recording (Phase 2):** `record.start(path.webm, url?)`, `record.stop()`
- **Capture (Phase 2):** `screenshot(path?,opts?)`, `pdf(path?)`
- **Network (Phase 3):** `network.route(url, opts?)`, `network.unroute(url?)`, `network.requests(opts?)`, `network.request(id)`, `network.har.start(path?)`, `network.har.stop(path?)`
- **Cookies (Phase 3):** `cookies.get()`, `cookies.set(name, value, opts?)`, `cookies.clear()`

- **Storage (Phase 3):** `storage.local.get(key?)`, `storage.local.set(key, value)`, `storage.local.clear()` and matching `storage.session.*`
- **Tabs (Phase 3):** `tabs.list()`, `tabs.new(url?, opts?)`, `tabs.close(ref?)`, `tabs.select(ref)`
- **Diff (Phase 3):** `diff.snapshot(opts?)`, `diff.screenshot(opts)`, `diff.url(url1, url2)`
- **Debug/Perf (Phase 4):** `trace.start()`, `trace.stop(path?)`, `profiler.start(opts?)`, `profiler.stop(path?)`, `inspect()`, `clipboard(op, text?)`, `vitals(url?)`, `pushstate(url)`
- **React DevTools (Phase 4):** `react.tree()`, `react.inspect(id)`, `react.renderers.start()`, `react.renderers.stop()`, `react.suspense(opts?)` — requires `launch({ enable: "react-devtools" })`
- **Streaming (Phase 4):** `stream.enable(opts?)`, `stream.disable()`, `stream.status()`
- **AI (Phase 4):** `chat(message, opts?)` — requires an AI gateway configured in agent-browser
- **Escape hatch (Phase 4):** `cmd(command, ...args)`, `batch(cmds, opts?)`
- **Auth login (Phase 4):** `auth.login(name)` — uses a saved auth profile on the current session
- **Lifecycle:** `session` (read-only name), `close()`

Argument vocabularies (the agent-browser surface). The string arguments to `get`, the `is*` checks, and `find` are passed through to the agent-browser CLI verbatim, so the accepted values are agent-browser's, not sercon's. The set below is current for **agent-browser 0.27.1** (the version this manual was written against — see the authoritative-reference note at the end of this section):

- **get(what, sel?)** — `what` is one of `text`, `html`, `value`, `attr` (the attribute name goes in the `sel` slot, e.g. `get("attr", "href")`), `title`, `url`, `count`, `box` (bounding box), `styles`, or `cdp-url`. When `sel` is a plain CSS selector it scopes the query; a snapshot `@ref` (see below) targets a specific element; omit it to query the page or the active element. The result is in the `.data` field of the envelope (e.g. `(await b.get("title")).data.title`).
- **isVisible(sel) / isEnabled(sel) / isChecked(sel)** — map to agent-browser's `is visible|enabled|checked <selector>`. `sel` is required. The boolean lands in `.data` (e.g. `.data.visible`).
- **find(locator, value, { action, text? })** — `locator` is one of `role`, `text`, `label`, `placeholder`, `alt`, `title`, `testid`, `first`, `last`, or `nth` `value` is the locator's argument (the role/text/etc. to match). `action` is **required** and is an act-style verb — `click`, `dblclick`, `hover`, `focus`, `check`, `uncheck`, `fill`, or `type` (the last two consume text). agent-browser's `find` cannot return a handle without acting, so there is no read-only `find` for read-only matching use `snapshot()`. The chainable `locator(spec)` handle exposes the same verbs as methods.
- **snapshot(opts?)** — `opts` is `{ interactive?: boolean, compact?: boolean, depth?: number, selector?: string }`, mapping to agent-browser's `-i / -c / -d <n> / -s <sel>`. The snapshot is an accessibility tree whose interactive nodes carry refs like `@e2` **those refs are valid selectors** for `get`, `click`, `fill`, etc. on the same session — the snapshot → act-by-ref loop is agent-browser's intended workflow for AI agents.
- **console(opts?) / errors(opts?)** — `opts` is `{ clear?: boolean }` (passes `--clear`); returns captured console logs / page errors.

Reaching the rest of agent-browser. sercon binds the commonly-used subset; agent-browser ships more (sessions, dashboard, confirmation prompts, iOS-simulator and cloud-provider backends, init scripts, ...). Any subcommand that isn't a first-class binding is reachable through the `cmd(command, ...args)` escape hatch — `cmd` maps 1:1 onto an `agent-browser \<command\> [args]` invocation. For the **authoritative, version-matched** command and flag reference — including data shapes returned in `.data`, which evolve with the CLI — run `agent-browser --help`, or the agent-oriented guide `agent-browser skills get core --full`. Treat the tables in this manual as a convenience snapshot, not a contract: agent-browser owns its own surface.

Capture (Phase 2). `screenshot` and `pdf` are path-first with opt-in in-memory bytes:

- `b.screenshot(path?, opts?)` — when `path` is given the image is written to that file and the result is `{ path: string, size: number, format: string }`. When omitted the raw bytes are returned as `{ bytes: number[], format: string }` — a plain JS `number[]` (byte-value array). Convert to a typed array with `new Uint8Array(bytes)`.
- `b.pdf(path?)` — same path-first model; format is always `"pdf"`.
- `opts` for `screenshot`:
`{ selector?, full?, annotate?, format?: "png"|"jpeg", quality? }`.

Defaults bag (Phase 2). A per-Run namespace-level defaults object is merged under per-call `launch()` `opts` so common flags don't need to be repeated:

```
services.agentBrowser.setDefaultOptions({ headed: false, proxy: "http://proxy:3128" });
const b = services.agentBrowser.launch(); // inherits headed + proxy
```

- `setDefaultOptions(obj)` — replace the defaults map entirely.
- `defaultOptions()` — return a shallow copy of the current defaults.
- `clearDefaultOptions()` — reset to `{}`.

Defaults are scoped to the current Run; they do not persist across `Run` calls.

Flat one-shot shortcuts (Phase 2). Launch an ephemeral session, open a URL, perform one action, and close — without managing a handle:

Shortcut	Equivalent handle sequence
<code>agentBrowser.screenshot(url, path?, opts?)</code>	<code>launch()+open(url)+screenshot(path?,opts?)+close()</code>
<code>agentBrowser.pdf(url, path?)</code>	<code>launch()+open(url)+pdf(path?)+close()</code>
<code>agentBrowser.snapshot(url, opts?)</code>	<code>launch()+open(url)+snapshot(opts?)+close()</code>
<code>agentBrowser.eval(url, js)</code>	<code>launch()+open(url)+eval(js)+close()</code>

Example (Phase 1 + Phase 2):

```
if (!services.agentBrowser.available) {
  runtime.log("install agent-browser first");
} else {
  const html = "data:text/html," + encodeURIComponent(
    "<title>sercon demo</title><h1 id=hi>Hello</h1><input id=box>"
  );
}
```

```

const b = services.agentBrowser.launch({ headed: false });
try {
  await b.open(html);

  const titleRes = await b.get("title");
  runtime.log("title:", titleRes.data?.title); // "sercon demo"

  await b.fill("#box", "typed by sercon");
  const valRes = await b.get("value", "#box");
  runtime.log("value:", valRes.data?.value); // "typed by sercon"

  const visRes = await b.isVisible("#hi");
  runtime.log("visible:", visRes.data?.visible); // true

  const snap = await b.snapshot({ compact: true });
  runtime.log("snap keys:", Object.keys(snap).join(", "));

  // Phase 2 – settings + capture
  await b.set.viewport(1280, 800);
  const shot = await b.screenshot({ full: true }); // bytes path
  runtime.log("bytes:", new Uint8Array(shot.bytes).length, shot.format);

  const file = await b.screenshot("/tmp/demo.png"); // file path
  runtime.log("file:", JSON.stringify(file)); // { path, size, format }
} finally {
  await b.close();
}

// One-shot shortcut
const r = await services.agentBrowser.eval(html, "document.title");
runtime.log("title:", r.data?.result); // "sercon demo"
}

```

Phase 3: network, cookies, storage, tabs, diff.

All Phase-3 groups are handle methods (operate on an open session) and return the standard { success, data, error } envelope.

Network interception (network.*).

```

// Block all XHR requests matching a glob pattern.
await b.network.route("**/api/*", { abort: true });

// Stub a JSON response body instead of blocking.
await b.network.route("**/data.json", { body: { mock: true } });

// Remove a specific route (or all routes when url is omitted).
await b.network.unroute("**/api/*");

// Inspect captured requests.
const reqs = await b.network.requests({ clear: true, method: "POST" });
runtime.log(JSON.stringify(reqs.data));

// Get a single request by id.
const req = await b.network.request("req-1");

```

```
// Record a HAR trace.
await b.network.har.start("/tmp/trace.har");
// ... navigate and interact ...
await b.network.har.stop();
```

Route opts: abort?: boolean, body?: unknown (JSON-serialised), resourceType?: string (e.g. "xhr").

Requests opts: clear?: boolean, filter?: string, type?: string, method?: string, status?: string.

Cookie management (cookies.*).

```
// Read the current cookie jar.
const jar = await b.cookies.get();
runtime.log(jar.data?.cookies); // array of cookie objects

// Set a cookie (requires a real HTTP origin – not available on data: URLs).
await b.cookies.set("session_id", "abc123", {
  domain: ".example.com",
  httpOnly: true,
  secure: true,
  sameSite: "Lax",
  expires: Math.floor(Date.now() / 1000) + 3600, // Unix seconds
});

// Clear all cookies.
await b.cookies.clear();
```

cookies.set opts: url?, domain?, path?, httpOnly?, secure?, sameSite?: "Strict" | "Lax" | "None", expires? (Unix seconds).

> **Note:** cookies.set and storage.* require a real HTTP origin. > They will fail on data: URLs — wrap in try-catch when testing with > data: URLs.

Web storage (storage.*).

```
// localStorage round-trip (requires a real HTTP origin).
await b.storage.local.set("theme", "dark");
const val = await b.storage.local.get("theme");
runtime.log(val.data); // { theme: "dark" } or similar
await b.storage.local.clear();

// sessionStorage works identically under storage.session.
await b.storage.session.set("token", "xyz");
const tok = await b.storage.session.get("token");
await b.storage.session.clear();
```

get(key?) — omit key to get all entries; pass a key to get one.

Tab management (tabs.*).

```

// Open a new tab (optionally at a URL, with an optional label).
await b.tabs.new("https://docs.example.com", { label: "docs" });

// List all open tabs. Refs are t1/t2/... or the user-supplied label.
const list = await b.tabs.list();
runtime.log(list.data?.tabs); // [{ tabId, title, url, active, label, type }]

// Switch the active tab.
await b.tabs.select("docs"); // by label
await b.tabs.select("t1"); // by ref

// Close a specific tab (or the current tab when ref is omitted).
await b.tabs.close("docs");

```

Page diffing (diff.*).

```

// Snapshot the current DOM state (baseline snapshot when no -b given).
const snap = await b.diff.snapshot();

// Compare against a saved baseline snapshot.
const delta = await b.diff.snapshot({ baseline: "/tmp/base.json", compact:
true });
runtime.log(delta.data);

// Pixel-level screenshot comparison.
const px = await b.diff.screenshot({
  baseline: "/tmp/base.png",
  output: "/tmp/diff.png",
  threshold: 0.1, // 0-1; pixels above this are flagged
});

// Compare two URLs side-by-side without an open session.
const cmp = await b.diff.url("https://staging.example.com", "https://prod.example.
com");

```

diff.snapshot opts: baseline?: string (path to a prior snapshot JSON),
 selector?: string (scope to a subtree), compact?: boolean, depth?: number.

Phase 4: debug/perf, React DevTools, streaming, AI chat, escape hatch, auth vault.

All Phase-4 groups are handle methods (operate on an open session) unless noted. They return the standard { success, data, error } envelope; batch returns an array of per-command envelopes.

Debug / perf (trace.*, profiler.*, inspect, clipboard, vitals, pushstate).

```

// Chrome DevTools tracing.
await b.trace.start();
// ... navigate and interact ...
await b.trace.stop("/tmp/trace.json");

// V8 CPU profiling (optional category filter).
await b.profiler.start({ categories: "v8,blink" });
await b.profiler.stop("/tmp/profile.json");

```

```
// Open the Chrome DevTools inspector and get the DevTools URL.
const devtools = await b.inspect();
runtime.log("DevTools:", devtools.data?.url);

// Clipboard read/write.
await b.clipboard("write", "copied text");
const clip = await b.clipboard("read");
runtime.log(clip.data?.text);

// Core Web Vitals for the current page (or a given URL).
const v = await b.vitals();
runtime.log("LCP:", v.data?.lcp);

// SPA client-side navigation without a full page reload.
await b.pushstate("/app/dashboard");
```

`profiler.start` opts: { categories?: string } — a comma-separated list of V8/Blink profiling categories.

React DevTools (`react.*`).

Requires the session to have been launched with `launch({ enable: "react-devtools" })`. `agent-browser` returns a clear error (surfaced as a throw) when `react-devtools` was not enabled at launch time.

```
const b = services.agentBrowser.launch({ enable: "react-devtools" });
await b.open("https://my-react-app.example.com");

// Get the component tree.
const tree = await b.react.tree();
runtime.log(JSON.stringify(tree.data).slice(0, 200));

// Inspect a specific fiber node by id (id from tree.data).
const node = await b.react.inspect("fiber-123");

// Measure re-renders.
await b.react.renders.start();
// ... interact ...
const renders = await b.react.renders.stop();
runtime.log(JSON.stringify(renders.data));

// List Suspense boundaries (optionally only dynamic ones).
const suspense = await b.react.suspense({ onlyDynamic: true });
```

`react.suspense` opts: { onlyDynamic?: boolean } — when true, only boundaries that have actually suspended are included.

Live streaming (`stream.*`).

```
// Enable a browser-event stream on a chosen port.
await b.stream.enable({ port: 9229 });

const status = await b.stream.status();
```

```
runtime.log("streaming:", status.data?.enabled);

await b.stream.disable();
```

`stream.enable` `opts: { port?: number }` — selects the streaming port (agent-browser picks a default when omitted).

AI chat (`chat`).

Requires an AI gateway to be configured in agent-browser; errors cleanly if not configured.

```
// Single-shot natural-language instruction.
const r = await b.chat("Click the Login button");
runtime.log("chat ok:", r.success);

// Optionally specify the model.
const r2 = await b.chat("Fill out the form with test data", { model: "gpt-4o" });
```

Escape hatch (`cmd` / `batch`).

`cmd` is a generic passthrough for any agent-browser subcommand. `batch` runs multiple commands in a single round-trip and returns an **array** of per-command result envelopes (not the usual single envelope).

```
// cmd: call any agent-browser subcommand.
const title = await b.cmd("get", "title");
runtime.log("title:", title.data?.title);

// batch: multiple commands, one round-trip.
const results = await b.batch(["get title", "get url"], { bail: false });
for (const r of results) {
  runtime.log(r.success, JSON.stringify(r.data));
}
```

`batch` `opts: { bail?: boolean }` — stop at the first failing command and return results up to that point.

Auth vault — namespace (`auth.save` / `list` / `show` / `delete`) + `handle` (`auth.login`).

Vault CRUD is session-independent and lives on the namespace (`services.agentBrowser.auth.*`). `login` fills a form on a live page and is a handle method (`b.auth.login(name)`).

Passwords are never placed in argv. `auth.save` feeds the password through the subprocess stdin using `--password-stdin`. This means the password never appears in `ps` output or shell history.

```
// Save a login profile (password is sent via stdin, not argv).
await services.agentBrowser.auth.save("prod", {
  url: "https://app.example.com/login",
  username: "admin",
  password: "s3cret", // sent via stdin
  usernameSelector: "#user",
  passwordSelector: "#pass",
});
```

```

    submitSelector: "#submit",
  });

  // List all saved profiles (returns an envelope; data contains profile names).
  const list = await services.agentBrowser.auth.list();
  runtime.log(JSON.stringify(list.data));

  // Show details of a saved profile.
  const info = await services.agentBrowser.auth.show("prod");

  // Delete a profile.
  await services.agentBrowser.auth.delete("prod");

  // Log in using a saved profile (handle method — requires an open session).
  const b = services.agentBrowser.launch();
  await b.open("https://app.example.com/login");
  await b.auth.login("prod"); // fills the form and submits
  runtime.log("logged in");
  await b.close();

```

`auth.save` opts:

```
{ url, username, password, usernameSelector?, passwordSelector?, submitSelector? }.
```

All string opts except `password` map to the corresponding agent-browser flags; `password` is stripped from the option object and fed via stdin.

services.webdriver

W3C WebDriver client backed by github.com/tebeka/selenium (pure Go — no cgo). Drives a real browser via an installed `chromedriver` or `geckodriver` no driver binary is bundled or downloaded.

Availability gate. `services.webdriver.available` is a sync boolean — true when `chromedriver` or `geckodriver` is found on PATH. Gate all calls on this value.

Connect-or-start model. `connect(opts?)` is async. When `opts.url` is given, it dials the running driver endpoint at that URL. When `opts.url` is omitted, it locates the driver binary on PATH, picks an ephemeral port, starts the driver as a subprocess, and dials it. The resolved session handle has the same surface in both cases. Sessions quit (and any started subprocess is stopped) when the Run ends; call `session.quit()` explicitly for deterministic teardown.

Probe. `probe({ url })` checks whether a WebDriver endpoint is ready without opening a session. Returns `{ ready: boolean, status?: number }` on HTTP success or `{ ready: false, error: string }` on transport failure; never throws on network errors.

Connect options (all optional):

Key	Type	Default	Description
<code>browser</code>	<code>"chrome"</code> <code>"firefox"</code>	<code>"chrome"</code>	Selects the driver binary (<code>chromedriver</code> or <code>geckodriver</code>).
<code>headless</code>	<code>boolean</code>	<code>true</code>	Run the browser in headless mode.

<code>url</code>	<code>string</code>	—	Dial an already-running WebDriver server instead of starting one.
<code>args</code>	<code>string[]</code>	—	Extra browser command-line flags appended to the default set.
<code>capabilities</code>	<code>object</code>	—	Raw W3C capability overrides merged last (escape hatch).
<code>commandTimeout</code>	<code>number</code>	30000	Per low-level WebDriver request timeout (ms). Bounds each command so a driver blocked behind an open alert or an unreachable endpoint can't hang the call. 0 /negative disables it. <code>quit()</code> (and Run-end cleanup) also cancels any in-flight command.

Browser flags (`args`) and capabilities (the WebDriver surface). Unlike `agentBrowser`, this binding speaks the W3C WebDriver protocol directly (via `tebeka/selenium`) rather than wrapping a CLI, so the “external surface” is the W3C capability set and the browser’s own command-line flags:

- **`args`** are the browser’s launch flags, appended *after* the headless flag `sercon` adds for you (`--headless=new` for Chrome, `-headless` for Firefox when `headless` is true). They land in the vendor capability block (`goog:chromeOptions.args` / `moz:firefoxOptions.args`). Common Chrome flags: `--no-sandbox`, `--disable-gpu`, `--disable-dev-shm-usage`, `--window-size=1280,800`, `--lang=en-US`, `--proxy-server=host:port`. Firefox takes `-width 1280`, `-height 800`, `-private`, etc.
- **`capabilities`** is an object of top-level **W3C standard capabilities**, merged into the session request *last* (so it overrides `sercon`’s defaults). Standard keys: `browserName`, `browserVersion`, `platformName`, `acceptInsecureCerts` (boolean — accept self-signed/expired TLS), `pageLoadStrategy` ("normal" | "eager" | "none"), `unhandledPromptBehavior` ("dismiss" | "accept" | "ignore" | ...), `timeouts` ({ `script`, `pageLoad`, `implicit` } in ms), and `proxy` ({ `proxyType`, `httpProxy`, `sslProxy`, ... }). Vendor blocks go here too — `goog:chromeOptions` and `moz:firefoxOptions` (each { `args?`, `binary?`, `prefs?`, ... }). Because the merge is by top-level key, supplying `goog:chromeOptions` wholesale *replaces* the one built from `args` / `headless` — set `args` for additive flags, `capabilities` only when you need full control. See the W3C WebDriver capabilities spec for the complete list.

Session handle methods (all async):

Method	Returns	Description
<code>get(url)</code>	{ <code>ok</code> : true }	Navigate to a URL.
<code>url()</code>	<code>string</code>	Current URL.
<code>title()</code>	<code>string</code>	Current page title.
<code>back()</code> / <code>forward()</code> / <code>refresh()</code>	{ <code>ok</code> : true }	Browser history controls.

<code>find(by, value)</code>	element handle	Find the first matching element.
<code>findAll(by, value)</code>	element handle[]	Find all matching elements.
<code>source()</code>	string	Current page source HTML.
<code>screenshot(path?)</code>	{ path?, size?, bytes?, format }	Capture the viewport. No path → bytes: number[] with path → { path, size, format }.
<code>executeScript(js, args?)</code>	unknown	Run a synchronous script in the page context.
<code>executeScriptAsync(js, args?)</code>	unknown	Run an async script (callback-resolved).
<code>cookies()</code>	object[]	Get all cookies.
<code>setCookie(c)</code>	{ ok: true }	Set a cookie { name, value, path?, domain?, secure?, (httpOnly?, expiry? }).
<code>deleteCookie(name)</code>	{ ok: true }	Delete a named cookie.
<code>deleteAllCookies()</code>	{ ok: true }	Clear all cookies.
<code>setImplicitWait(ms)</code>	{ ok: true }	Set the implicit element-wait timeout.
<code>waitFor(by, value, opts?)</code>	element handle	Poll in the active frame until an element appears (opts: timeout? ms, visible?: boolean, enabled?: boolean — wait past a disabled→enabled flip).
<code>clickWhenReady(by, value, opts?)</code>	{ ok: true }	Wait in the active frame (default visible + enabled) then issue a native (trusted) click that fires React handlers (opts: timeout?, visible?, enabled?, poll? ms).
<code>quit()</code>	{ closed: true }	Close the session (idempotent; also called by Run-end cleanup).

Locator strategies (by argument): `css`, `xpath`, `id`, `name`, `tag`, `className`, `linkText`, `partialLinkText`.

Element handle methods (all async):

Method	Description
--------	-------------

<code>click()</code>	Click the element.
<code>sendKeys(text)</code>	Type text into the element.
<code>clear()</code>	Clear the element's value.
<code>submit()</code>	Submit the form.
<code>text()</code>	Inner text content.
<code>getAttribute(name)</code>	Get the named HTML attribute ; returns <code>null</code> when the attribute is absent (W3C semantics). This is the <i>attribute</i> , not the live DOM <i>property</i> — an <code><input></code> 's typed text lives in its <code>value</code> property, not a <code>value</code> attribute, so read it with <code>executeScript</code> (e.g. <code>executeScript("return document.querySelector('#box').value")</code>). Chrome's legacy fallback returns the property; Firefox correctly returns <code>null</code> .
<code>cssValue(name)</code>	Get a computed CSS property value.
<code>tagName()</code>	Tag name in lower case.
<code>isDisplayed()</code> / <code>isEnabled()</code> / <code>isSelected()</code>	Visibility / state checks.
<code>find(by, value)</code> / <code>findAll(by, value)</code>	Sub-tree element search.
<code>screenshot(path?)</code>	Screenshot the element (scrolls into view).

Example:

```

if (!services.webdriver.available) {
  runtime.log("no chromedriver/geckodriver on PATH – skip");
} else {
  const d = await services.webdriver.connect({ browser: "chrome", headless:
true });
  try {
    await d.get("data:text/html," + encodeURIComponent(
      "<title>wd demo</title><h1 id=hi>Hello</h1><input id=box>"));
    runtime.log("title:", await d.title());
    const h1 = await d.find("id", "hi");
    runtime.log("text:", await h1.text(), "visible:", await h1.isDisplayed());
    const box = await d.find("css", "#box");
    await box.sendKeys("typed by sercon");
    runtime.log("id attr:", await box.getAttribute("id")); // "box"
    // Live input value is a DOM property, not an attribute – read via script
    // so it works the same on Chrome and Firefox.
    runtime.log("value:", await d.executeScript("return
document.querySelector('#box').value", []));
    runtime.log("eval:", await d.executeScript("return 6*7", []));
    const shot = await d.screenshot();
    runtime.log("bytes:", new Uint8Array(shot.bytes).length, shot.format);
  } finally {
    await d.quit();
  }
}

```

```

    }
}

```

Phase 2 (v0.41.0). The session handle gains the following additional async methods.

Phase 2 session methods:

Method	Returns	Description
<code>windowHandles()</code>	<code>string[]</code>	All open window/tab handles.
<code>currentWindow()</code>	<code>string</code>	Handle of the currently focused window/tab.
<code>switchToWindow(handle)</code>	<code>{ ok: true }</code>	Focus a window/tab by its handle.
<code>newWindow(type?)</code>	<code>{ handle, type }</code>	Open a new tab ("tab" , default) or window ("window"). Returns the new handle.
<code>closeWindow()</code>	<code>string[]</code>	Close the current window/tab; returns the remaining handles and auto-switches to a survivor.
<code>switchToFrame(target)</code>	<code>{ ok: true }</code>	Switch into a frame by index (number) or by passing an element handle for the <code><iframe></code> . A name/id string is not accepted — locate the iframe via <code>find</code> and pass the element handle.
<code>switchToParentFrame()</code>	<code>{ ok: true }</code>	Switch to the parent frame.
<code>switchToDefaultContent()</code>	<code>{ ok: true }</code>	Switch back to the top-level browsing context.
<code>acceptAlert()</code>	<code>{ ok: true }</code>	Accept (OK/confirm) the current alert/confirm/prompt dialog.
<code>dismissAlert()</code>	<code>{ ok: true }</code>	Dismiss (cancel) the current dialog.
<code>alertText()</code>	<code>string</code>	Get the message text of the current dialog.
<code>sendAlertText(text)</code>	<code>{ ok: true }</code>	Type into a prompt dialog.
<code>maximize()</code>	<code>{ x, y, width, height }</code>	Maximize the browser window; returns the new rect.
<code>minimize()</code>	<code>{ x, y, width, height }</code>	Minimize the browser window; returns the new rect.
<code>fullscreen()</code>	<code>{ x, y, width, height }</code>	Enter fullscreen; returns the new rect.

<code>setWindowRect({width?,height?,x,y,width,height})</code>	<code>{ x, y, width, height }</code>	Set position and/or size.
<code>getWindowRect()</code>	<code>{ x, y, width, height }</code>	Get the current window rect.
<code>hover(e1)</code>	<code>{ ok: true }</code>	Move the mouse pointer over an element (W3C actions).
<code>dragAndDrop(src, dst)</code>	<code>{ ok: true }</code>	Drag from <code>src</code> element to <code>dst</code> element.
<code>keyChord(...keys)</code>	<code>{ ok: true }</code>	Send a key chord (e.g. <code>keyChord("Control", "a")</code>).
<code>performActions(sequence)</code>	<code>{ ok: true }</code>	Send a raw W3C actions array (escape hatch).
<code>releaseActions()</code>	<code>{ ok: true }</code>	Release any held action state (buttons/keys).

Element handle additions (Phase 2): Every element handle now carries an `elementId` string property (the raw W3C element reference id). Two new async methods are available: `e1.hover()` moves the mouse over the element; `e1.dragTo(target)` drags the element to a target element handle.

executeScript / executeScriptAsync — element handle return values: When a script returns a single W3C element reference, the binding wraps it in an element handle automatically. A top-level array of element references is also wrapped (each entry becomes an element handle). Nested references (e.g. an object with element-reference values) are returned as plain objects — return them flat or unwrap manually.

Caveat — alert unhandledPromptBehavior : chromedriver's default `unhandledPromptBehavior` is "dismiss and notify", which may auto-dismiss dialogs before your script can interact with them. To inspect alerts manually, pass

`capabilities: { unhandledPromptBehavior: "ignore" }` in `connect()`. With "ignore", the alert is kept open but the renderer is blocked by the modal, so any WebDriver command that executes script in the page (including `executeScript`) will hang until the alert is dismissed. Trigger alerts via `<body onload="alert(...)">` or a synchronous inline script so that `get()` completes with the dialog already open — then `alertText()` and `acceptAlert()` work without racing against the modal state.

Caveat — file upload: For local `<input type="file">`, use `e1.sendKeys("/absolute/path/to/file")` — this works with both Chrome and Firefox. Remote-grid file upload (the `/session/{id}/file` endpoint) is not supported by the underlying `tebeka/selenium` library and is out of scope.

Phase 2 example:

```
if (!services.webdriver.available) {
  runtime.log("no chromedriver/geckodriver on PATH — skip");
} else {
  const d = await services.webdriver.connect({
    browser: "chrome",
    headless: true,
    capabilities: { unhandledPromptBehavior: "ignore" },
  });
```

```

try {
  // Windows / tabs
  const nw = await d.newWindow("tab");
  await d.switchToWindow(nw.handle);
  runtime.log("open tabs:", (await d.windowHandles()).length);
  await d.closeWindow();

  // Window rect
  await d.setWindowRect({ width: 1024, height: 768 });
  const r = await d.getWindowRect();
  runtime.log("rect:", r.width + "x" + r.height);

  // Frames
  await d.get("data:text/html," + encodeURIComponent(
    "<title>outer</title><iframe srcdoc='<p id=inner>hi</p>'></iframe>"));
  await d.switchToFrame(0);
  runtime.log("frame text:", await (await d.find("id", "inner")).text());
  await d.switchToParentFrame();
  await d.switchToDefaultContent();

  // Alerts (requires unhandledPromptBehavior: "ignore")
  // Navigate to a page that fires alert() synchronously on load; get()
  // blocks until the load event, so the dialog is open when it returns.
  try {
    await d.get("data:text/html," + encodeURIComponent(
      "<body onload=\"alert('hello')\"><title>a</title>"));
  } catch (_) {
    // chromedriver may surface an unexpected-alert-open error – the alert
    // is still live, carry on.
  }
  runtime.log("alert:", await d.alertText());
  await d.acceptAlert();

  // Hover + drag
  await d.get("about:blank");
  await d.executeScript(`
    document.body.innerHTML =
      '<div id=a
style="position:absolute;top:50px;left:50px;width:80px;height:80px"></div>' +
      '<div id=b
style="position:absolute;top:50px;left:200px;width:80px;height:80px"></div>';
    return 1;`, []);
  const a = await d.find("id", "a");
  const b = await d.find("id", "b");
  await a.hover();
  await d.dragAndDrop(a, b);

  // executeScript returning an element handle
  const el = await d.executeScript("return document.getElementById('b')", []);
  runtime.log("script element tag:", await el.tagName());
} finally {
  await d.quit();
}
}

```

Frames / iframes

WebDriver — nested + cross-origin. `switchToFrame(target)` accepts a frame index (number), an iframe **element handle** from `find()`, or a **CSS-selector string** (it finds the iframe in the current context and switches). After a switch, element queries are scoped to that frame; `switchToParentFrame()` goes up one level and `switchToDefaultContent()` returns to the top document. `frameChain([t1, t2, ...])` switches from the top document through each level in one call — the ergonomic way to reach a deeply nested frame. This uses the W3C `/frame` protocol, so cross-origin frames (e.g. a Klarna Checkout widget's nested payment iframe) work the same as same-origin ones.

```
await d.frameChain(["#klarna-checkout-iframe", "#klarna-fullscreen-iframe"]);
const confirm = await d.find("css", "button.kco-confirm"); // scoped to the inner
frame
await d.switchToDefaultContent();
```

WebDriver — waiting & reliable clicks. `find` and all element interaction follow the **active frame** set via `switchToFrame` / `frameChain` (W3C frame-scoping — there is no separate frame-context to manage). For buttons that render or enable asynchronously inside a frame (a common pattern in checkout widgets, e.g. a Klarna Checkout “Pay order” that flips disabled→enabled), combine a switch with `clickWhenReady`:

- `waitFor(by, value, opts?)` polls in the active frame until the element is present, and — when requested — `visible` and/or `enabled`. Opts: `timeout?` (ms, default 10000), `visible?: boolean`, `enabled?: boolean` (waits past a disabled→enabled transition). Resolves to the element handle.
- `clickWhenReady(by, value, opts?)` waits — by default both `visible` **and** `enabled` — in the active frame, then issues a **native (trusted)** W3C Element-Click that fires React/synthetic handlers. Opts: `timeout?`, `visible?`, `enabled?` (default `true` for both), and `poll?` (inter-attempt interval in ms). Resolves `{ ok: true }`.

```
await d.frameChain(["#klarna-checkout-iframe", "#payment-frame"]);
await d.clickWhenReady("css", "button.pay", { timeout: 8000 });
```

What made such nested-checkout flows flaky was **timing** (the button rendering or enabling late), not frame context — `find` already follows the chain. `clickWhenReady` is the fix: a wait (visible+enabled, in the active frame) followed by a trusted click.

CDP trusted clicks (Chrome only)

Some elements cannot be activated by the standard WebDriver click: when a button lives inside a **nested cross-origin iframe**, the W3C Element Click hit-tests to the parent `<iframe>` element and fails with `element click intercepted` (synthetic `executeScript().click()` is untrusted and ignored by frameworks like React). The classic case is completing a Klarna Checkout “Pay order”.

`session.cdpClick(by, value, opts?)` solves this via chromedriver's CDP passthrough. It locates the element across the **whole pierced frame tree** (no `switchToFrame` needed), reads its true viewport coordinates with `DOM.getContentQuads`, scrolls it into view, and dispatches a **trusted** `Input.dispatchEvent` `press/release`. Opts: `timeout` (ms, default 10000), `poll` (ms, default 50), `button` (`"left"` / `"right"` / `"middle"`), `scrollIntoView` (default `true`), `offsetX` / `offsetY`. Resolves to `{ clicked: true, x, y }`. Throws on firefox.

`session.cdp(command, params?)` is a raw escape hatch that sends a single CDP command (e.g. `Input.dispatchMouseEvent`, `Browser.getVersion`) and returns its result. Chrome-only.

```
const d = await services.webdriver.connect({ browser: "chrome" });
await d.get(checkoutUrl);
await d.cdpClick("xpath", '//button[normalize-space="Pay order"]', { timeout:
8000 });
```

Out-of-process iframes (OOPIFs). A cross-site iframe (different eTLD+1 — e.g. a Klarna Checkout hosted on another domain) runs in a separate renderer process, which chromedriver’s page-session CDP passthrough cannot reach. `cdpClick` therefore opens a **browser-level CDP connection** (a direct DevTools WebSocket, the way Puppeteer/Playwright do), attaches to the page and every iframe target, locates the element in its owning session, and dispatches the trusted `Input.dispatchMouseEvent` there — browser-level input is hit-tested across OOPIF boundaries. (A different *port* is still same-site and stays in-process; only a different site becomes an OOPIF.)

`session.targets()` lists the browser’s CDP targets (including OOPIF iframe targets); `session.attach(target)` (a target object or `targetId`) returns a `{ targetId, sessionId, cdp(method, params?), detach() }` handle for running CDP against that specific target. These require a CDP-capable session (local chromedriver, or a Grid advertising `se:cdp`); `cdpClick` falls back to the page-session path when browser-level CDP is unavailable.

agentBrowser — single level. `b.frame(target)` switches the session’s frame context to an iframe (CSS selector or `@ref` from `snapshot`); `b.frame("main")` returns to the top document. Subsequent `click` / `fill` / `get` / `snapshot` operate inside the switched frame. It is **single-level only**: `agent-browser` resolves the selector against the *main* document and cannot descend into nested frames, so sequential `frame()` calls won’t reach an inner frame. For nested or cross-origin nesting, use `WebDriver (frameChain)`. (A Klarna/KCO completion helper is out of scope — it needs the live payment widget.)

services.typst

External-CLI binding to the Typst compiler. Every call shells out to the local `typst` binary; nothing is linked into `sercon` (the `typst-as-lib` Rust embedding can’t be used under our no-cgo constraint, so the external CLI is the only viable route — same opt-in, feature-detected model as `services.git` / `services.gh` / `services.ai` / `services.agentBrowser`).

Availability gate. `services.typst.available` is a sync boolean over `exec.LookPath("typst")` — resolved once per Run. Gate on it; every other `typst` binding throws a clean error when the CLI is absent. Install with `brew install typst` (or from [<https://typst.app>](https://typst.app)).

Surface.

- `services.typst.available` — `boolean`. True when `typst` is on PATH.
- `services.typst.version()` — `Promise<string>`. The trimmed `typst --version` line.
- `services.typst.fonts()` — `Promise<string[]>`. Font families `typst` can see, deduplicated and sorted.
- `services.typst.compile(opts)` — `Promise<{ format, bytes?, path? }>`. Compile a document to PDF/PNG/SVG.

- `services.typst.query(opts)` — `Promise<unknown>`. Query a document for elements matching a selector; returns parsed JSON.

Input. `compile` and `query` take exactly **one** of `input` (a path to a `.typ` file) or `source` (inline Typst markup). Passing both — or neither — throws.

compile output rule. With `no output`, a **PDF** is compiled to a temp file and returned as **bytes** (`{ format: "pdf", bytes: Uint8Array }`) — this bytes-return path is **PDF only**. With an `output path`, the file is written there (`{ format, path }`) and `format` is inferred from the extension (`.pdf` / `.png` / `.svg`); **PNG and SVG require an output path**. Other options: `root` (project root for imports), `inputs` (`sys.inputs` key/value pairs → `--input k=v`), `ppi` (PNG resolution), `fontPaths` (extra `--font-path` dirs), `timeout` (ms, default 60000).

Inline-source caveat. When you pass `source`, it is written to a temporary `main.typ` and compiled in that temp directory, so relative `#import` / `read()` paths resolve against the temp dir, **not** your working directory. For documents that pull in sibling files, pass `input` (or set `root`) instead of `source`.

query. `selector` is required (e.g. `"<label>"`, `"heading"`); `field` extracts a single field per match, `one` returns just the first match.

```
if (!services.typst.available) {
  runtime.log("install typst (brew install typst) first");
} else {
  runtime.log("typst:", await services.typst.version());

  // Inline source → PDF bytes (PDF-only without an output path).
  const pdf = await services.typst.compile({ source: "= Hi\nFrom Typst." });
  runtime.log(pdf.format, "bytes:", pdf.bytes.length);

  // Render to PNG on disk (png/svg require an output path).
  await services.typst.compile({ source: "= Hi", output: "/tmp/out.png", ppi:
144 });

  // Query a labelled metadata value as JSON.
  const v = await services.typst.query({
    source: "#metadata(42) <answer>",
    selector: "<answer>", field: "value", one: true,
  });
  runtime.log("answer:", v); // 42
}
```

services.doctor

`services.doctor(requires?)` reports on every external tool sercon can use and, optionally, asserts a script's prerequisites. It is the script-side form of the `--doctor` CLI flag and returns the same report.

```
const r = await services.doctor(requires?: string[]);
```

Report shape. The promise resolves to:

```
{
  ok: boolean;           // false only when a compatibility conflict was found
```



```

satisfied: boolean; // true when every `requires` entry is met (unmet === [])
unmet: string[];    // requested requirements that are absent or conflicted
tools: {
  name: string;      // binary name, e.g. "chromedriver"
  category: string;  // feature group, e.g. "webdriver"
  purpose: string;   // what sercon uses it for
  installed: boolean;
  version: string | null; // null when not installed or no --version
  ok: boolean;        // false only on a compatibility conflict
  detail?: string;    // e.g. "matches Chrome 149" or the conflict reason
}[];
}

```

Tools probed include `git`, `gh`, the AI providers (`claude`, `codex`, `copilot`, `gemini`), `agent-browser`, `chromedriver` / `geckodriver`, `typst`, the HTTP backends (`recon`, `curl`), and the platform clipboard / image tools (`pbcopy` / `pbpaste` / `osascript` / `pngpaste` on macOS, `wl-copy` / `wl-paste` / `xclip` / `xsel` on Linux, `clip` / `powershell` on Windows).

requires — assert prerequisites. Each entry is either a feature/category name (`"webdriver"`, `"typst"`, `"ai"`, `"git"`, `"gh"`, `"agentBrowser"`, `"clipboard"`, `"image"`, `"http"`) or a specific binary name (`"chromedriver"`, `"pngpaste"`, `"claude"`, ...). A category is met when at least one tool in it is installed and `ok`. Anything unmet lands in `unmet` — **a missing tool is reported, not thrown**, so a script can self-skip an optional feature. An **unknown** requirement name (a typo) throws.

The chromedriver↔Chrome check. When `chromedriver` is installed, sercon locates the installed Chrome/Chromium and compares the two major versions. A mismatch sets that tool's `ok` to `false` (with a `detail` explaining it), the top-level `ok` to `false`, and — for `--doctor` — exits with code `5`. A missing Chrome leaves `ok` `true` and notes that the version couldn't be verified. **Missing is fine; only a conflict fails.**

```

// Assert a hard prerequisite.
const need = await services.doctor(["git"]);
runtime.assert.ok(need.satisfied, "git is required");

// Self-skip an optional feature.
const opt = await services.doctor(["typst", "webdriver"]);
if (!opt.satisfied) runtime.log("skipping; unmet:", JSON.stringify(opt.unmet));

```

tui

Multi-pane terminal UI (was `ui.tui` in v0.8.0). `tui.layout(tree)` declares panes; `tui.pane(name)` returns a pane handle with `write`, `writeln`, `clear`, `title`. `services.exec.shell({pane})` streams subprocess I/O into a pane.

Scripts that orchestrate multiple subprocesses (e.g. `brew`, `npm`, `cargo` updates running in parallel) can declare a multi-pane terminal layout and route each subprocess's stdout/stderr into its own pane. Activation is **script-driven**: the first call to `tui.layout(...)` enters TUI mode. If stdout is not a TTY (CI, pipes, `make demo`), the same calls fall back to prefixed plain-text lines (`[paneName] line`) so the same script runs in both contexts.

Layout

```
tui.layout({
  rows: [
    { name: "log", title: "Orchestrator", weight: 1 },
    { cols: [
      { name: "brew" },
      { name: "npm" },
    ],
      weight: 2,
    },
  ],
});
```

Each layout node is one of:

- { name: string, title?: string, weight?: number } — a leaf pane.
- { rows: LayoutNode[], weight?: number } — a vertical split.
- { cols: LayoutNode[], weight?: number } — a horizontal split.

`weight` defaults to 1. Pane names must be unique across the whole tree. `layout` may be called **once** per Run; a second call throws.

Pane handles

```
const log = tui.pane("log");
log.writeln("Updating Homebrew...");
log.title("Done");
log.clear();
log.write("partial line ");
```

`tui.pane(name)` returns a handle with `write`, `writeln`, `clear`, `title`. Methods are synchronous from the script's perspective — they enqueue to the TUI goroutine.

Subprocess routing

`services.exec.shell(cmd, opts)` accepts `opts.pane` (a Pane handle or a pane name string):

```
await services.exec.shell("brew update && brew upgrade", { pane: "brew" });
```

When set, the subprocess's `stdout` **and** `stderr` stream into the pane line by line. The returned { `exitCode`, `durationMs`, `success` } is the same as without `pane`: except `stdout` and `stderr` are empty (data was streamed, not captured). ANSI colors are translated to `tview` color tags and rendered natively; `\r` without `\n` overwrites the current line so progress spinners render cleanly.

Pass { `pty`: true } to run the command under a pseudo-terminal (Unix only). The child then believes it is a terminal and emits color, progress bars, and spinners, which a `pane` renders (or, without a `pane`, are captured into `stdout`). This is the general alternative to per-tool force-color flags (`FORCE_COLOR=1`, `--color=always`): it works for any tool that gates output on a TTY. A `pty` merges `stdout` and `stderr` onto one stream, so `stderr` is empty in `pty` mode. On Windows `pty` is ignored (pipe fallback, no color).

Keybindings (TTY mode only)

Key	Action
Tab / Shift-Tab	Cycle focus through panes
PgUp / PgDn	Scroll focused pane one page
↑ / ↓	Scroll focused pane one line
Home / End	Jump to top / bottom
Ctrl-C	Abort the script (single press)

The focused pane has a yellow border; the bottom status bar shows its name and the active keys.

Panes follow the tail automatically; scroll up (keys or, when enabled, the mouse wheel) to pause following, and scroll back to the bottom (or press `End`) to resume. Set `{ autoscroll: false }` on a leaf to keep it pinned at the top. Pass `{ mouse: true }` at the layout root to enable mouse support: the wheel scrolls the pane **under the cursor** (without changing which pane is focused), and a left-click focuses the pane under the cursor. This disables the terminal's native click-drag selection while active; hold Shift/Option to select.

Per-leaf `{ wrap: "char" | "word" | "off" }` controls line wrapping (default `"char"` `"off"` lets long lines scroll horizontally), and `{ color: false }` strips a pane's ANSI to plain text instead of rendering it. Both affect TTY rendering only — the non-TTY fallback always emits plain prefixed lines.

`await tui.waitKey()` resolves with the next key (`{ name, rune, ctrl, alt, shift }`); `tui.onKey(handler)` registers a persistent per-keypress callback and returns an unsubscribe function. Both see every key except `Ctrl-C` (which always aborts) and coexist with the built-in navigation keys. Both require a TTY: `waitKey` rejects and `onKey` is a no-op in non-TTY (fallback) mode.

Limitations (v1)

- `tui` is **incompatible with** `--watch`: calling `tui.layout()` under `--watch` throws. Use one or the other.
- No runtime layout mutation, no input panes (`tui.input`), no snapshot-to-normal-screen on exit. The alt screen is restored at script end and the usual `PASS / FAIL` line prints.

image

Image I/O and manipulation, all pure-Go and synchronous. Decode a raster image (PNG/JPEG/GIF/TIFF/BMP/WebP) or rasterize an SVG subset, then chain transforms on the returned **Image handle**. Each transform returns a *fresh* handle (immutable chaining); the source is never mutated.

Module functions

Function	Returns	Notes
<code>image.open(path)</code>	<code>Image</code>	Read + decode a file; format sniffed from magic bytes, not the extension.
<code>image.decode(bytes)</code>	<code>Image</code>	Decode in-memory <code>Uint8Array</code> image bytes.

<code>image.rasterizeSVG(src, { width, height })</code>	Image	Rasterize an SVG (subset) to a raster image at the given pixel size. <code>src</code> is a path or <code>Uint8Array</code> both dimensions required (> 0).
---	-------	--

Handle properties (read-only): `width`, `height`, `format` (the decoded source format string, or "svg" for a rasterized SVG).

Handle methods (each returns a fresh `Image` unless noted):

Method	Effect
<code>resize(w, h, { filter? })</code>	Resize; a <code>0</code> dimension preserves aspect ratio. Filter $\in$ <code>lanczos</code> (default), <code>nearest</code> , <code>linear</code> , <code>box</code> , <code>catmullrom</code> .
<code>fit(w, h)</code>	Scale to fit within <code>w</code> $\times$ <code>h</code> , preserving aspect (no crop).
<code>thumbnail(w, h)</code>	Scale + centre-crop to exactly <code>w</code> $\times$ <code>h</code> .
<code>crop(x, y, w, h)</code>	Crop a rectangle (bounds-checked).
<code>rotate(deg)</code>	Rotate CCW by an arbitrary angle; corners filled transparent.
<code>rotate90()</code> / <code>rotate180()</code> / <code>rotate270()</code>	Lossless quarter-turns (CCW).
<code>flipH()</code> / <code>flipV()</code>	Mirror horizontally / vertically.
<code>brightness(pct)</code> / <code>contrast(pct)</code> / <code>saturation(pct)</code>	Adjust by a percentage in <code>[-100, 100]</code> .
<code>gamma(g)</code>	Gamma correction (<1 darkens, >1 brightens).
<code>sharpen(sigma)</code> / <code>blur(sigma)</code>	Unsharp-mask / Gaussian blur.
<code>grayscale()</code> / <code>invert()</code>	Desaturate / photographic negative.
<code>overlay(other, x, y, opacity?)</code>	Alpha-blend another handle on top.
<code>paste(other, x, y)</code>	Paste another handle (no blend).
<code>bytes(format, { quality? })</code>	Encode $\rightarrow$ <code>Uint8Array</code> .
<code>save(path, { format?, quality? })</code>	Encode + write to disk (format inferred from the extension if not given). Returns <code>void</code> .

Encode formats: `png`, `jpeg`, `gif`, `tiff`, `bmp`, `webp`.

Worth knowing:

- `resize` with a `0` width or `0` height computes the missing dimension to preserve the aspect ratio.
- `quality` (1..100, default 90) applies to **JPEG** only. **WebP encode is lossless** (via `nativewebp`), so the `quality` option is accepted for API symmetry but ignored.
- SVG is rasterize-in only** — there is no SVG output; `rasterizeSVG` renders a *subset* of SVG to a raster image you then encode as PNG/etc.
- GIF decode is first-frame** — animation is not preserved.

```
const im = image.open("avatar.png");
const thumb = im.resize(128, 0).grayscale().blur(0.5);
thumb.save("thumb.webp"); // lossless webp
const png = im.crop(0, 0, 64, 64).bytes("png"); // → Uint8Array
```

Pure-Go stack: `disintegration/imaging` + `golang.org/x/image` (webp/tiff/bmp decode) + `HugoSmits86/nativewebp` (webp encode) + `srwiley/oksvg` + `rasterx` (SVG). No cgo.

web

Fetch & parse web documents. Three families — `web.feed` (RSS/Atom/JSON feeds), `web.sitemap` (sitemaps), and `web.html` (lenient HTML scraping) — each exposing a synchronous `parse(string)` and an async `load(url, opts?)`. `parse` works on a string you already have; `load` GETs the URL first, then parses the body. This closes the HTML-scraping gap that `codec.xml` (strict XML only) left open. All pure-Go, no cgo.

Shared fetch options. Every `load(url, opts?)` reuses the `net.http` option surface — `timeout` (ms number or duration string), `headers`, `follow` (redirects), `userAgent`, and `basic-auth username / password` — and sends a default `sercon-web/<version> User-Agent` unless you override it. A non-2xx response throws.

web.feed

Function	Returns	Notes
<code>web.feed.parse(source)</code>	<code>Feed</code>	Parse RSS, Atom, or JSON-feed text; format auto-detected (<code>feedType</code> reports it).
<code>web.feed.load(url, opts?)</code>	<code>Promise<Feed></code>	Fetch then parse.

`Feed` is normalized to one model regardless of source format — the RSS/Atom field differences (`pubDate` / `updated`, `description` / `summary`) are unified:

```
{ feedType, title, description, link, updated, items[] }. Each item is
{ title, link, published, updated, content, summary, author, guid, categories[],
  raw }
```

The `raw` object is the escape hatch carrying format-specific extras (enclosures, namespaced elements like `media:*` / `dc:*`) that the normalized model drops.

web.sitemap

Function	Returns	Notes
<code>web.sitemap.parse(source)</code>	<code>Sitemap</code>	Parse a <code>urlset</code> or <code>sitemapindex</code> document.
<code>web.sitemap.load(url, opts?)</code>	<code>Promise<Sitemap></code>	Fetch then parse; transparently decompresses <code>.xml.gz</code> .

`Sitemap` is `{ type: "urlset" | "sitemapindex", urls[], sitemaps[], errors[] }`. `urls` entries are `{ loc, lastmod?, changefreq?, priority? }` an index lists child sitemap URLs in `sitemaps`. `load` detects `gzip` by magic bytes on the body (a `Content-Encoding: gzip` response is already unwrapped by the HTTP transport). For an index, pass `{ expand: true }` to fetch every child sitemap (bounded, single-level expansion — a child

that is itself a `sitemapindex` is not recursed) and merge their URLs into `urls` per-child fetch/parse failures are collected in `errors[]` rather than thrown.

`web.html`

Function	Returns	Notes
<code>web.html.parse(source)</code>	<code>Node</code>	Parse HTML leniently — real-world tag soup is accepted, never throws on bad markup.
<code>web.html.load(url, opts?)</code>	<code>Promise<Node></code>	Fetch then parse.

`parse` / `load` return a chainable **Node** rooted at the document. Query it with CSS or XPath, read it with the accessor methods:

Method	Effect
<code>find(selector)</code>	First CSS match (or <code>null</code>).
<code>findAll(selector)</code>	All CSS matches as a <code>Node[]</code> .
<code>xpath(expr)</code>	First XPath match (or <code>null</code>).
<code>xpathAll(expr)</code>	All XPath matches as a <code>Node[]</code> .
<code>text()</code>	Concatenated text content.
<code>html()</code>	Outer HTML of the node.
<code>innerHTML()</code>	Inner HTML (children only).
<code>tag()</code>	Tag name.
<code>attr(name)</code>	Attribute value (or <code>null</code>).
<code>attrs()</code>	All attributes as a <code>Record<string, string></code> .

Sub-queries are scoped to the receiver node — use `./` for relative XPath; a leading `//` is document-wide.

```
const feed = await web.feed.load("https://example.com/feed.xml");
feed.items.forEach(i => console.log(i.title, i.link));

const sm = await web.sitemap.load("https://example.com/sitemap.xml", { expand:
true });
console.log(sm.urls.length, "urls");

const doc = await web.html.load("https://example.com");
doc.findAll("a.product").forEach(a => console.log(a.text(), a.attr("href")));
doc.xpathAll("//h2").forEach(h => console.log(h.text()));
```

Pure-Go stack: `mmcdole/gofeed` (feeds) + `andybalholm/cascadia` (CSS selectors over `golang.org/x/net/html`) + `antchfx/htmlquery` (XPath). No cgo.

6. Servers

The `server` global (added in v0.10.0) hosts inbound listeners — scripts that *accept* connections rather than make them. It ships `server.http` and `server.https` (§6.2), `server.smtp` (§6.7), and the raw `server.tcp` / `server.udp` listeners (§6.8); future cycles will add IMAP, FTP, and POP3 against the same engine foundation.

6.1 Foundation: `HoldRun` + `LoopCallable`

Two engine primitives back every long-lived binding. They live on `*scriptengine.Engine` (CLI use of them is internal; the public contract is “use these if you write a similar binding yourself”).

`Engine.HoldRun(reason string) (release func())` — keeps `loop.Run` from exiting while a binding has resources outstanding. The event loop normally exits as soon as its `jobCount` drops to zero; `setTimeout` / `setInterval` / `setImmediate` are the only things that count. `HoldRun` parks a 24-hour sentinel timer under the hood (same trick `PromisifyAsync` uses for one-shot async work), returns a `release` function that clears it, and increments a refcount so multiple concurrent holders compose. `release` is idempotent. Any sentinels still parked when `Run` returns are cleanup-drained automatically as a safety net — a binding that forgets to call `release` does not leak into the next `Run`.

`scriptengine.NewLoopCallable(loop, fn)` — wraps a captured `goja.Callable` (a JS function reference) so it can be invoked from *any* goroutine. Goja’s runtime is single-threaded; calling a `Callable` directly from a non-loop goroutine corrupts engine state.

`LoopCallable.Call(buildArgs)` enqueues a `loop.RunOnLoop` callback, builds the arg list on the loop (where `vm.ToValue` is valid), invokes the callable, and returns the result to the caller’s goroutine. `LoopCallable.CallOnLoop(vm, args...)` is the already-on-the-loop variant — using `Call` from inside a loop callback deadlocks because the new `RunOnLoop` job can’t run until the current one returns.

You use them together. A typical server binding holds one or two `LoopCallable`s for the user’s handlers, parks one `HoldRun` sentinel per active listener, and releases on close. The CLI bindings under `cmd/sercon/api_server*.go` follow this pattern.

Concurrency model. Because every handler invocation marshals back onto the single goja loop, **handlers serialize**: only one JS handler runs at a time, across all listeners. This is a feature (no JS data races, no shared-state confusion) and a throughput ceiling. For CPU-bound work in a handler, offload to a Go binding that does the work in a goroutine and resolves a `Promise`.

6.2 `server.http.listen` / `server.https.listen`

```
// HTTP
const srv = await server.http.listen({
  port: 8080, // required
  host: "0.0.0.0", // default
  routes: { "GET /": (req, res) => res.text("hi") },
  use: [logger], // optional global middleware
});

// HTTPS — identical plus cert/key (file paths OR inline PEM strings)
const srv2 = await server.https.listen({
  port: 8443,
  cert: "/etc/ssl/server.pem",
  key: "/etc/ssl/server.key",
  routes,
});

// HTTPS dev shortcut — ephemeral self-signed cert, no key, no files
const srv3 = await server.https.listen({
```

```

    port: 8443,
    cert: "self-signed",
    routes,
  });

  srv.address; // "tcp/0.0.0.0:8080"
  await srv.close(); // graceful: stop accepting, drain, release HoldRun
  await srv.stopped; // Promise that resolves when the listener exits

```

Options:

Key	Type	Notes
port	number	Required. Bind port. 0 lets the kernel pick; read the chosen port off <code>srv.address</code> .
host	string	Default "0.0.0.0". Pass "127.0.0.1" for loopback-only.
routes	Record<string, RouteValue>	Required (may be <code>{}</code>). Keys are stdlib <code>http.ServeMux</code> patterns.
use	Middleware[]	Optional global middleware; runs in array order before any per-route middleware.
onError	(err, req, res) => unknown	Optional error handler; renders a custom response when a handler/middleware throws or rejects. See below.
cert	string	HTTPS only. File path, inline PEM, or the literal "self-signed" to mint an ephemeral in-process dev cert.
key	string	HTTPS only. File path or inline PEM. Not needed (ignored) when <code>cert</code> is "self-signed".

Custom error pages (`onError`). When a handler or middleware throws (or returns a rejected Promise) and `onError` is set, it is invoked as `onError(err, req, res)` instead of sending the stock 500 Internal Server Error. `err` is the thrown value (an `Error` for a plain `throw new Error(...)`); render whatever you like via the usual `res.*` terminals (`res.status(500).json({...})`, `res.html(...)`, etc.). `onError` may be `async`. If it finishes without finalizing a response, or itself throws/rejects, `sercon` falls back to the stock 500 — a buggy error handler can never wedge the request. Without `onError`, behaviour is unchanged (stock 500 + a log line); the per-route `try/catch` pattern still works and takes precedence (an error you catch never reaches `onError`).

Self-signed dev cert. `cert: "self-signed"` generates a fresh P-256 certificate in-process (valid 1 year), with SANs covering `localhost`, `127.0.0.1`, `::1`, and the listen `host`. The keypair lives only for the life of the run and is never written to disk. Self-signed certs fail normal client verification by design (`net.probe.tls` skips verification, so it still works against them) — this is a local-dev convenience, not a production path. For real deployments own the cert in your supervisor (`caddy` / `nginx` / `traefik`) in front.

Route patterns are stdlib Go 1.22+ `http.ServeMux` syntax: `"METHOD /path/{param}/{rest...}"`. The leading method is optional (omit to match any).

`{name}` captures one path segment; `{name...}` captures the tail (wildcard). No new routing library — `stdlib` does the matching, including longest-prefix wins. Captured values land in `req.params`.

RouteValue is either a bare handler — `(req, res) => res.text("...")` — or an object `{ use?: Middleware[], handler: HandlerFn }` for per-route middleware (see §6.4).

Request shape (`req`):

```
type Request = {
  method: string; // "GET"
  url: string; // full URL incl. scheme + host
  path: string; // "/users/42"
  query: Record<string, string[]>; // ?a=1&a=2 → {a: ["1","2"]}
  headers: Record<string, string[]>; // lowercase keys
  params: Record<string, string>; // path-pattern captures
  body: string; // UTF-8 decoded request body
  bodyBytes: Uint8Array; // same data as raw bytes
  remote: string; // "1.2.3.4:54321"
  cookies: Record<string, string>; // parsed Cookie header
};
```

Response builder (`res`) — every method returns `res` for chaining; the terminal `.json` / `.text` / `.html` / `.bytes` / `.empty` / `.redirect` sets the body and flips an internal finalized flag:

```
type CookieOpts = {
  domain?: string;
  path?: string;
  maxAge?: number; // seconds; 0 = session; <0 = delete
  expires?: number; // unix ms
  secure?: boolean;
  httpOnly?: boolean;
  sameSite?: "strict" | "lax" | "none";
};

type Response = {
  status(code: number): Response;
  header(name: string, value: string): Response;
  cookie(name: string, value: string, opts?: CookieOpts): Response;
  // Terminals (all return res so they remain chainable but no further
  // write makes sense after one fires; a second terminal throws):
  json(value: unknown): Response; // → application/json
  text(s: string): Response; // → text/plain
  html(s: string): Response; // → text/html
  bytes(b: Uint8Array, ct?: string): Response; // default application/octet-
stream
  empty(): Response; // no body; pair with .status() for
204/304
  redirect(loc: string, code?: number): Response; // default 302
  // Upgrade (WebSocket):
  upgradeWebSocket(opts?: { readBuffer?: number }): Promise<WebSocket>;
  // Stream (Server-Sent Events):
  sse(opts?: { keepAlive?: number; retry?: number }): SSEStream;
};
```

```

type SSEStream = {
  send(data: string | { data: unknown; event?: string; id?: string; retry?:
number }): Promise<void>;
  close(): Promise<void>;
  readonly closed: Promise<void>; // resolves on close OR client disconnect
  readonly remote: string;
};

```

After the handler's returned Promise resolves, the engine checks `res.finalized`:

- **true** — buffered status / headers / body get written to the underlying `http.ResponseWriter`.
- **false** — engine sends `204 No Content`.
- **handler threw** (sync or async) — engine sends `500 Internal Server Error` and logs the error (`stderr` under vanilla `sercon` access log under `sercon serve`). The script keeps running; only that one request is affected.

Calling a terminal twice on the same `res` throws a `TypeError` (the second call sees `finalized === true` and refuses).

Multiple listeners compose naturally. A script can run

```

const [a, b] = await Promise.all([
  server.http.listen({ port: 80, routes }),
  server.https.listen({ port: 443, routes, cert, key }),
]);

```

Each `listen` call holds its own `HoldRun` sentinel; closing one listener does not affect the other.

6.3 Static-file serving

```

"GET /assets/{rest...}": server.http.static({
  dir:           "/var/www/public", // root directory on disk
  stripPrefix:   "/assets/",        // URL prefix to strip before disk lookup
  index?:        "index.html",      // accepted but currently unused (stdlib
default applies)
  etag?:         true,              // accepted but currently unused (stdlib
default applies)
}),

```

`server.http.static({...})` returns an opaque handler the route compiler unwraps and registers directly on the mux. **The route pattern must include a `{rest...}` wildcard** — without it the match only ever produces the literal prefix and no subpath ever resolves on disk. Internally a stdlib `http.FileServer(http.Dir(dir))` wrapped in `http.StripPrefix(stripPrefix, ...)` ETag and range requests work out of the box. Symlinks outside `dir` are blocked (stdlib `http.Dir`'s default). No directory listing customisation yet — `index` and `etag` options are reserved for future tuning; v0.10.0 inherits the stdlib defaults.

6.4 Middleware

Onion model. A middleware is `(req, res, next) => Promise<void>`. `next` is an async function that runs the rest of the chain; awaiting it lets the middleware post-process.

```
const logger = async (req, res, next) => {
  const start = runtime.time.nowMs();
  await next(); // run downstream chain
  runtime.log(req.method, req.path, "→", runtime.time.nowMs() - start, "ms");
};

await server.http.listen({
  port: 8080,
  use: [logger], // global middleware
  routes: {
    "GET /api/secure": { // per-route middleware
      use: [authCheck],
      handler: (req, res) => res.json({ ok: true }),
    },
  },
});
```

Attachment points:

- **Global** — `use`: array in `listen({...})` options. Runs for every request, in declaration order, before any per-route middleware.
- **Per-route** — when a route value is `{ use: [...], handler: fn }`, those middleware run after globals but before the handler.

For a request to `POST /api/secure` with global middleware `[A, B]` and per-route middleware `[C]`, the composition is:

```
A(req, res, () => B(req, res, () => C(req, res, () => handler(req, res))))
```

Each layer can `await next()` then mutate response headers or record timings, exactly like Express / Koa.

Short-circuit: a middleware that does *not* call `await next()` stops the chain. It must terminate `res` itself (e.g. `res.status(401).text("denied")`) — otherwise the engine sends 204 No Content per the unfinalized-response rule.

Throwing: any throw mid-chain (sync or via rejected Promise) bubbles to the engine's 500 path. Wrap with `try / catch` if you want custom error handling.

6.5 WebSocket upgrade

```
"GET /ws": async (req, res) => {
  const ws = await res.upgradeWebSocket({ readBuffer: 64 });
  // ws is both an async iterator AND has .send / .close.
  for await (const msg of ws) {
    if (msg.type === "text") {
      await ws.send("echo:" + msg.text);
    } else {
      await ws.send(msg.bytes); // msg.type === "binary"
    }
  }
}
```

```

    }
    // iterator ends on close; ws.closeCode / ws.closeReason populated
  },

```

Type:

```

type WSMMessage =
  | { type: "text"; text: string }
  | { type: "binary"; bytes: Uint8Array };

type WebSocket = AsyncIterable<WSMMessage> & {
  send(data: string | Uint8Array): Promise<void>;
  close(code?: number, reason?: string): Promise<void>;
  closeCode?: number;
  closeReason?: string;
  remote: string;
};

```

Library: `github.com/coder/websocket` (the modern successor to `nhooyr.io/websocket` `gorilla/websocket` is in maintenance mode). Pure Go, zero deps.

Marshalling. Per upgraded connection, a Go goroutine reads frames into a buffered channel (capacity = `readBuffer`, default 64). Each `next()` on the async iterator pulls one message off the channel and resolves on the loop via `LoopCallable`. `ws.send()` schedules a write back through the connection-owning goroutine.

Backpressure. If the read buffer fills (slow JS consumer), back-pressure propagates into the websocket library's read loop — the client eventually sees `1009 message too big` after the library's own limit (1MB default). Bump `readBuffer` if your workload bursts and you want more in-flight queueing.

`ws.close()` closes the send channel, drains the receive channel, sends a close frame, and ends the async iterator on the next `next()` call. The script's `for await` exits cleanly; the handler's Promise resolves; the connection's goroutine returns. The HTTP server's `HoldRun` is **unaffected** — it stays alive for new connections until `srv.close()`.

`async function*` and `for await (...)` are not natively parsed by `goja`. `esbuild`'s `Supported` flag lowers both during transpile (see CHANGELOG), and every Run installs `Symbol.asyncIterator = Symbol.for("@@asyncIterator")` so the lowered helper and user code agree on the same iteration key.

Server-Sent Events (`res.sse`)

`res.sse(opts?)` starts a one-way text/event-stream response: the handler holds the connection open and pushes events until it (or the client) closes the stream.

```

"GET /events": (req, res) => {
  const stream = res.sse({ keepAlive: 15000 });
  let n = 0;
  const t = setInterval(() => stream.send({ event: "tick", data: { n: n++ } })),
  1000);
  stream.closed.then(() => clearInterval(t)); // stop on client disconnect

```

```

    return stream.closed; // keep the handler alive
  },

```

Options. `keepAlive (ms)` sends : ping comment frames at that interval to defeat idle-proxy timeouts (off by default). `retry (ms)` emits a one-time `retry:` line telling the client its reconnect delay.

`send(data)` . `data` is a string (→ a single `data: line`) or an object `{ event?, data, id?, retry? }`. Object `data` is JSON-encoded; string `data` is passed through, and multi-line string `data` becomes one `data: line` per source line (per the SSE grammar). The returned Promise resolves once the frame is flushed to the socket, and rejects if the stream is already closed.

`close()` ends the stream and resolves once it is torn down. **`closed`** is a Promise that resolves on `close()` **or** when the client disconnects — use it to stop timers/producers.

Mechanics. Unlike `upgradeWebSocket` , SSE does **not** hijack the connection — it keeps writing to the same `http.ResponseWriter` . So the request’s dispatcher goroutine parks until the stream closes (otherwise `net/http` would finish the response). A dedicated pump goroutine owns the writer and flushes each event; the event loop never writes to the socket directly. The listener’s `HoldRun` keeps the loop alive while streams are open, exactly like `WebSocket`.

6.6 Lifecycle

Vanilla `sercon script.ts` . Each `server.*.listen` call parks a `HoldRun` sentinel and the event loop stays alive while the listener is bound. SIGINT terminates abruptly: the engine’s interrupt watcher calls `loop.Terminate()` , the cleanup drain runs (closing each listener), and the process exits. No access log, no readiness signal, no shutdown timeout.

`sercon serve script.ts` . Adds the production niceties documented in §4: structured access log to stderr, a `READY` listening on `tcp/...` line on stdout per listener, and graceful shutdown with `--shutdown-timeout` (default 30s). Clean SIGTERM exits 0 . Choose `serve` whenever the script is supervised (systemd, docker compose, launchd).

```

# Quick local test:
sercon examples/scripts/server-http.ts

# Production-style supervised run:
sercon serve examples/scripts/server-http.ts

```

6.7 SMTP

SMTP is a **stateful, multi-stage transaction**: a client connects, optionally authenticates, declares a sender (`MAIL FROM`), one or more recipients (`RCPT TO`), then streams the message body (`DATA`). `sercon` maps each stage to a JavaScript callback so a script can accept or reject the transaction incrementally — refuse an unknown sender at `MAIL` , refuse an unroutable recipient at `RCPT` , or inspect the parsed message at `DATA` . The inbound listener is `server.smtp.listen({...})` the outbound side is `net.email.send({...})` (covered at the end of this section).

`server.smtp.listen({...})`

Synchronously binds the listening socket (so a port already in use throws immediately), then serves in the background. Returns a handle once the loop schedules it.

```
const srv = await server.smtp.listen({
  port: 2525,
  hostname: "mx.example.com",
  handlers: {
    onMail: (env) => env.from.endsWith("@example.com"),
    onRcpt: (env, rcpt) => isLocalMailbox(rcpt),
    onData: (env, msg) => { store(msg); return true; },
  },
});

// srv.address → "tcp/0.0.0.0:2525"
// srv.close() → Promise<void> (graceful, 30s drain)
// srv.stopped → Promise<void> (resolves when the listener exits)
```

Options:

Field	Type	Default	Notes
port	number	—	Required. TCP port to bind.
host	string	"0.0.0.0"	Bind address.
hostname	string	os.Hostname()	EHLO greeting + Received: identity.
handlers	{onMail, onRcpt, onData}	—	Required. All three functions; see below.
auth	(user, pass, env) => bool Promise<bool>	—	Optional SASL verifier (PLAIN + LOGIN).
starttls	{cert, key}	—	Enables STARTTLS. File paths or inline PEM.
allowInsecureAuth	boolean	false	Permit AUTH on a plaintext connection (dev only).
maxMessageBytes	number	10485760 (10 MB)	DATA size cap; a non-positive value keeps the default.
maxRecipients	number	100	Max RCPT TO per transaction.
sessionTimeout	number	30000	Per-session idle timeout, milliseconds.

Handlers. `onMail(envelope)`, `onRcpt(envelope, recipient)`, and `onData(envelope, message)` are invoked at the matching protocol stage. All three are required. Each may be `async` — a returned Promise is awaited before the SMTP response is written. The return value (or a thrown exception) controls the SMTP reply, uniformly across all three:

Return value	SMTP response	Use case
--------------	---------------	----------

true / undefined	250 OK	Accept the stage
false	550 Command refused	Generic permanent reject
string	550 <string>	Permanent reject with a reason
throw	451 Temporary failure	Transient error (client should retry)

Envelope (passed to all three handlers; recipients grows as RCPT TO accumulates):

```
type Envelope = {
  from: string; // "" for the null sender (DSN bounces)
  recipients: string[]; // accepted RCPT TO so far
  remote: string; // "1.2.3.4:54321"
  helo: string; // EHLO/HELO hostname the client gave
  authenticatedUser?: string; // set if AUTH succeeded
  tls?: { version: string; cipher: string }; // set under STARTTLS
};
```

Message (only onData parsed via jhillyerd/enmime, with the exact original bytes always available as raw):

```
type Message = {
  from: string; // From: header (may differ from
  envelope.from)
  to: string[]; // To: header
  cc: string[]; // Cc: header
  subject: string;
  headers: Record<string, string[]>; // lowercase keys; multi-valued
  body: {
    text: string; // text/plain part; "" if absent
    html: string; // text/html part; "" if absent
  };
  attachments: Array<{
    filename: string;
    contentType: string;
    bytes: Uint8Array;
  }>;
  raw: Uint8Array; // exact DATA bytes (post dot-unstuffing)
};
```

Forwarders and DKIM-style signature verifiers should work from raw everyone else uses the parsed fields. enmime is lenient with malformed messages, so a script needing strict RFC 5322 conformance can re-parse raw itself.

AUTH. Supply an auth callback to enable SASL. Both **PLAIN** and **LOGIN** mechanisms are advertised; the callback receives the decoded (username, password, envelope) and returns a boolean (or a Promise of one). By default AUTH is refused on a plaintext connection — the client must complete STARTTLS first. Set allowInsecureAuth: true to accept credentials over plaintext; this is a **dev-only** escape hatch for trusted local networks and should never be set in production.

STARTTLS. Pass starttls: {cert, key} (file paths or inline PEM, same convention as server.https) to advertise STARTTLS. The upgrade happens mid-session on the same

connection; once active, `envelope.tls` is populated. The listener itself is plaintext at bind time — there is no implicit-TLS (SMTPS / port 465) inbound mode.

Concurrency. Each connection runs on its own Go goroutine, but every handler invocation **serializes on the goja loop** — exactly one handler runs at a time (the same single-threaded model as the HTTP listener in §6.2). A slow `onData` therefore blocks other connections' handlers. For slow work, acknowledge the stage (`return true`) and process in the background (start a Promise without awaiting it).

A short, self-contained round-trip (bind, send to self, assert receipt):

```
const port = 38095;
let captured: { subject: string; from: string } | null = null;

const srv = await server.smtp.listen({
  port,
  hostname: "test.local",
  handlers: {
    onMail: () => true,
    onRcpt: () => true,
    onData: (env, msg) => {
      captured = { subject: msg.subject, from: env.from };
      return true;
    },
  },
});

await net.email.send({
  to:      "alice@test.local",
  from:    "bob@test.local",
  subject: "round-trip demo",
  body:    "hello from the SMTP demo",
  server:  { host: "127.0.0.1", port, tls: "none" },
});

runtime.assert.equal(captured!.subject, "round-trip demo", "subject");
await srv.close();
```

`net.email.send({...})` — outbound

The outbound counterpart lives under `net.email` (next to the `spf` / `dmARC` / ... probes). It opens **one TCP connection per call** — no pooling, no queueing, no automatic retries — composes the MIME body in-tree, and returns a per-recipient outcome.

```
const result = await net.email.send({
  to:      ["alice@example.com", "bob@example.com"], // string OR string[]
  from:    "noreply@my.app",
  subject: "your verification code",
  body:    "your code is 123456",                    // plain text
  html:    "<p>your code is <b>123456</b></p>",          // optional HTML alternative
  attachments: [
    { filename: "qr.png", contentType: "image/png", bytes: qrBytes },
  ],
  headers: { "X-App": "myapp" },                    // optional extra headers
  server: {
```



```

    host: "smtp.example.com",
    port: 587, // default 587 (submission)
    auth: { username: "u", password: "p" }, // optional PLAIN auth
    tls: "starttls", // "starttls" (default) |
    "tls" | "none"
  },
  timeout: 30000, // ms; default 30s
});

// result = { accepted: string[], rejected: Array<{ address, reason }> }

```

MIME layout. `text only` → `text/plain; charset=utf-8` (no multipart). `text + html` → `multipart/alternative`. Any attachments → `multipart/mixed` (each part base64, `Content-Disposition: attachment`). `Date`, `Message-ID`, and `MIME-Version` are auto-generated; headers are merged on top.

TLS modes. `"starttls"` (default) connects plaintext then upgrades, refusing AUTH until TLS is active. `"tls"` is implicit TLS from connect (SMTPS, typically port 465). `"none"` is plaintext throughout, with AUTH disabled.

Outcomes. Transport-level failures (connection refused, TLS handshake, AUTH, `MAIL FROM` rejected) **throw** a JS exception. The `rejected` array captures only individual RCPT TO addresses the server permitted us to attempt but then refused — the transaction continues for the rest, so a partial send still delivers to the accepted recipients.

6.8 Raw TCP / UDP

`server.tcp.listen` and `server.udp.listen` are the inbound counterparts to the `net.tcp.connect` / `net.udp.open` clients (§5 net). They bind a listening socket **synchronously** and return the server handle directly — a port already in use throws immediately — then serve in the background, keeping the event loop alive via the same `HoldRun` model as the HTTP and SMTP listeners. `srv.close()` returns a `Promise<void>` (resolving once the listener is closed and accepted connections are drained), matching the rest of the `server.*` family. Passing `port: 0` binds an OS-chosen ephemeral port; read it back from the returned handle's `address`.

Both return a server handle:

```

srv.address; // "tcp/127.0.0.1:54321" or "udp/127.0.0.1:54321"
srv.close(); // Promise<void> – stop accepting, drain, release HoldRun

```

`server.tcp.listen({...}, conn => {...})`

The second argument is a **connection handler** invoked once per accepted socket. The `conn` it receives is the **same handle shape** as `net.tcp.connect` — there is no separate server-connection type to learn:

```

const srv = await server.tcp.listen({ port: 0 }, (conn) => {
  conn.onData(ev => conn.write(ev.bytes)); // ev = { bytes: Uint8Array, text }
  conn.onClose(() => runtime.log("peer disconnected"));
  conn.onError(e => runtime.log("conn error", String(e)));
  // conn.write(data) – string (UTF-8) or Uint8Array
  // conn.remote / conn.local – peer / local "host:port"
});

```

```
// conn.close() - half/full close this connection
});

const port = Number(srv.address.split(":").pop()); // "tcp/127.0.0.1:PORT"
// ... net.tcp.connect("127.0.0.1", port) to talk to it ...
await srv.close();
```

Options: port (required; 0 for ephemeral), host (default 127.0.0.1), readBuffer (per-connection read chunk size, same meaning as net.tcp.connect).

server.udp.listen({...}, (msg, reply) => {...})

UDP is connectionless, so the handler fires **once per inbound datagram** rather than once per connection. msg describes the datagram and its sender; reply sends a datagram back to that same sender:

```
const srv = await server.udp.listen({ port: 0 }, (msg, reply) => {
  // msg = { bytes: Uint8Array, text, address, port } - the sender
  runtime.log("got", msg.text, "from", msg.address + ":" + msg.port);
  reply("ack:" + msg.text); // string or Uint8Array; returns a Promise
});

const port = Number(srv.address.split(":").pop()); // "udp/127.0.0.1:PORT"
await srv.close();
```

Options: port (required; 0 for ephemeral), host (default 127.0.0.1).

Lifecycle. Identical to §6.6: vanilla sercon script.ts keeps the loop alive while bound and terminates abruptly on SIGINT; sercon serve script.ts adds the access log, a READY listening on tcp/... (or udp/...) line on stdout per listener, and graceful shutdown. As with all server.* bindings, the connection / datagram handlers marshal back onto the single goja loop, so **handlers serialize** (one at a time across all listeners).

- server.icmp.listen(opts?, (msg, reply) => ...) — bind a raw ICMP listener. **Requires root / CAP_NET_RAW** (synchronous bind throws otherwise); raw ICMP has no ports, so the socket receives **all** host ICMP traffic. opts { network?: "ip4" | "ip6" (default "ip4"), readBuffer? }. The handler runs per received packet with msg { bytes, text, address, type, code } and a reply(opts?) that sends an ICMP message back to the sender (or opts.to) — Echo mode { type?, code?, id?, seq?, payload? } or raw mode { type, code?, body }, the same options as net.icmp send. The handle is { address: "icmp/<addr>", close() } it emits a READY line under sercon serve and joins graceful shutdown.

7. JavaScript runtime built-ins (goja)

scriptengine runs on goja, which implements **ES5.1 + a large subset of ES6+**. The following are present and behave per spec:

Globals

- Object, Array, String, Number, Boolean, BigInt
- Math, Date, JSON, RegExp
- Error, TypeError, RangeError, SyntaxError, ReferenceError, EvalError, URIError

- `Promise` (provided via goja's native implementation; resolution microtasks are pumped by the event loop)
- `Symbol` (partial: well-known symbols `Symbol.iterator`, `Symbol.asyncIterator`, etc., are supported)
- `Proxy`, `Reflect` (partial)

```
runtime.log(Math.PI.toFixed(4), Math.max(1, 9, 3));
runtime.log(new Date().toISOString());
runtime.log(JSON.stringify({ a: 1, b: [2, 3] }));
runtime.log("abc".repeat(3), "abc".padStart(6, "_"));
```

Collections

- `Map`, `Set`, `WeakMap`, `WeakSet`

```
const counts = new Map<string, number>();
for (const ch of "banana") counts.set(ch, (counts.get(ch) ?? 0) + 1);
runtime.log([...counts]); // [[ "b", 1 ], [ "a", 3 ], [ "n", 2 ]]
```

Typed arrays

- `ArrayBuffer`, `DataView`
- `Int8Array`, `Uint8Array`, `Uint8ClampedArray`, `Int16Array`, `Uint16Array`, `Int32Array`, `Uint32Array`, `Float32Array`, `Float64Array`

```
const buf = new ArrayBuffer(4);
new DataView(buf).setUint32(0, 0xCAFEBAFE);
runtime.log(new Uint8Array(buf)[0].toString(16)); // ca
```

Iteration and generators

- `for...of`, `for...in`, `spread`, `destructuring`
- Generator functions (`function*`)
- Async generators (`async function*`) — supported via goja's event loop

```
function* fib() {
  let [a, b] = [0, 1];
  while (true) {
    yield a;
    [a, b] = [b, a + b];
  }
}
const it = fib();
const first10 = Array.from({ length: 10 }, () => it.next().value);
runtime.log(first10);
```

Not (yet) provided

- `fetch` — use `net.http.*` from the CLI, or register your own binding.
- `Worker`, threading primitives.
- DOM globals (`window`, `document`).
- Native ES modules (`import` / `export` are *transpiled* to CommonJS at load time — see section 9).

8. Async runtime additions (goja_nodejs)

The event loop bundles a small Node-compatible runtime on top of goja:

Timers

```
setTimeout(() => runtime.log("after 50ms"), 50);
setInterval(() => runtime.log("tick"), 100);
setImmediate(() => runtime.log("next tick"));
clearTimeout(t); clearInterval(i); clearImmediate(im);
```

The event loop holds the script alive while *any* timer is pending. If your script ends with only `setTimeout` callbacks scheduled, the run will wait for them to fire before returning.

console

```
console.log("info");
console.error("oops");
console.warn("careful");
console.info("fyi");
console.debug("trace");
```

Disable with `Options.DisableConsole = true` if you want a clean global namespace.

require

CommonJS `require`, with TS-aware extension fallback added by `sercon` (see section 11):

```
const helpers = require("./helpers");
helpers.doThing();
```

Both `require("./foo")` and `import { x } from "./foo"` resolve to the same module instance per Run; esbuild rewrites `import` to `require` at transpile time.

9. TypeScript support

TS is transpiled by esbuild in-process. There is no `tsc`, no `tsconfig.json`, no type-checking — types are erased and the resulting JS is what executes. Practical implications:

- All TS syntax esbuild supports is allowed: type annotations, `interface`, `type`, generics, enums (compiled to objects), namespaces, decorators (legacy & TC39 stage 3).
- `import type { X } from "y"` is stripped entirely (no runtime require).
- TS errors are *not* surfaced as build errors — only esbuild parse errors are. Run a separate `tsc --noEmit` in CI if you want type enforcement.
- `tsconfig.json` is **not** consulted.
- **Runtime errors report TypeScript positions.** Stack traces in thrown errors point at the original `.ts` line/column, not the transpiled JS — for both the entry script and imported `.ts` / `.tsx` modules, and for both synchronous throws and async (top-level- `await`) rejections. This is done with inline source maps that goja consumes natively; the entry script's ESM→CJS rewrite is line-shift-aware so its frames map correctly too. If a map is ever unavailable the trace falls back to transpiled-JS positions.

`.tsx` and `.jsx` are supported via esbuild's `LoaderTSX` / `LoaderJSX`, including for required modules resolved by the engine's extension fallback. JSX has no React runtime in scope by default — either set the factory at the file level with an `@jsx` pragma:

```

/** @jsx h */
function h(tag: string, props: any, ...children: any[]) {
  return { tag, props: props ?? {}, children };
}
export const Box = (label: string) => <div className="box">{label}</div>;

```

...or provide a `React.createElement`-compatible binding from Go and rely on esbuild's default pragma. ESM `export default <value>` also works through the entry-script rewriter's `__esModule ? .default : m` unwrap, so `import answer from "./mod"` resolves to the default export.

An **entry script's own** `export default <expr>` is captured as the value `Engine.Run` resolves to, and the `sercon` CLI prints that value as JSON to stdout (scripts without a default export resolve to `undefined` and print nothing; a value that doesn't JSON-encode, such as a function, is skipped). This is the supported way for a script to emit a structured result.

10. Top-level await

Top-level `await` works in entry scripts:

```

const r = await net.http.get("https://example.com");
runtime.log(r.status);

```

How it works under the hood: esbuild rejects top-level `await` with the `cjs` output format, so the engine transpiles the *entry* script with `format=esm`, then rewrites `import` statements to `require()` calls and wraps the rest of the body in:

```
;(async () => { /* your transpiled body */ })().then(__resolve, __reject);
```

`__resolve` and `__reject` are bindings the engine sets on the runtime to capture the final settlement. *Required-module* code (anything loaded through `require`) is transpiled with `format=cjs` and does **not** support top-level `await` — wrap in `(async () => ... )()` yourself if you need it inside a module.

11. Module resolution

The engine plugs a custom source loader into `goja_nodejs/require` that adds:

1. Direct path: if the requested file exists on disk, use it. `.ts` / `.tsx` files are transpiled on the fly.
2. **Extension fallback** when the request has no extension: try `.ts`, `.tsx`, `.js`, `.cjs`, `.mjs`, `.json` in that order.
3. **.js → .ts swap**: if a path ending in `.js` / `.cjs` / `.mjs` is requested but the literal file doesn't exist, the same path ending in `.ts` (or `.tsx`) is tried. Handles `package.json` `main` fields that point at compiled output where only the TypeScript source is on disk.
4. **package.json source preference**: when reading a `package.json`, if a `source` field is present and points at an existing `.ts` / `.tsx` file, `main` is rewritten to that path before the registry consumes it. Matches the convention used by parcel, microbundle, and similar bundlers to mark the TS source of truth.
5. **Directory index**: if the resolved path is a directory, try `index.ts`, `index.tsx`, `index.js`.

The base directory follows Node's CommonJS rules — relative imports resolve against the requiring module's directory, courtesy of `goja_nodejs`.

```
// All resolve to ./helpers/assert.ts:
import { check } from "./helpers/assert";
import { check } from "./helpers/assert.ts";
const { check } = require("./helpers/assert");
const { check } = require("./helpers/assert.js");
```

JSON imports work out of the box via either the ESM-style default import (esbuild rewrites it to a `require`) or a direct `require`:

```
import data from "./data.json";
const r = require("./data.json");
```

`node_modules` lookup *works* via the upstream registry, but the source loader doesn't currently transpile through that path — keep TS helpers under `ScriptRoot`.

12. Timeouts and cancellation

Two mechanisms interrupt a running script:

- `Options.Timeout` — wall-clock limit per `Run`. Exceeding it returns `scriptengine.ErrScriptTimeout`.
- `ctx` passed to `Run` / `RunFile` — cancelling the context returns `ctx.Err()` (typically `context.Canceled` or `context.DeadlineExceeded`).

Both call `vm.Interrupt(...)` and `loop.Terminate()` from a watcher goroutine. This works for **sync** JS — including `while(true){}` — because the goja VM checks the interrupt flag between bytecode steps. A host Go binding that's blocked in cgo or a long syscall isn't interruptible mid-call; if you write such a binding, plumb the engine `ctx` through to it yourself.

```
eng := scriptengine.New(scriptengine.Options{Timeout: 200 * time.Millisecond})
_, err := eng.Run(ctx, "loop.ts", `while (true) {}`)
// err is scriptengine.ErrScriptTimeout, returned ~200-300ms after Run.
```

13. Error semantics

- A Go binding returning `(T, error)` automatically throws as a JS exception when `err != nil`. The exception's `.message` is `err.Error()`.

```
try { mayFail(); } catch (e) { runtime.log(String(e)); }
```

- A Go binding that `panic`s with a `*goja.Object` (e.g. `vm.NewGoError(...)`) does the same.
- A script throwing a JS error out of the top level surfaces from `Run` as a Go `error` whose message is the JS error's `.message`.
- Timeout errors satisfy `errors.Is(err, scriptengine.ErrScriptTimeout)`. Cancellation errors satisfy `errors.Is(err, context.Canceled)` or `context.DeadlineExceeded`.

14. Type generation (.d.ts)

`Engine.WriteTypes(w)` walks the registered bindings and emits a TypeScript declaration file. The mapping table (abbreviated):

Go type	TS type
<code>string</code>	<code>string</code>
any numeric / <code>float64</code>	<code>number</code>
<code>bool</code>	<code>boolean</code>
<code>[]T</code>	<code>T[]</code> <code>[]byte</code> <code>→ Uint8Array</code>
<code>map[string]T</code>	<code>Record<string, T></code>
<code>*T</code>	<code>T</code> (transparent unwrap)
<code>error</code> (single return)	omitted (collapses to <code>void</code>)
<code>(T, error)</code> (tuple return)	<code>T</code>
<code>interface{}</code> / any	<code>unknown</code>
<code>goja.FunctionCall</code> parameter	folded into <code>(...args: unknown[])</code>
<code>goja.Value</code> return	<code>unknown</code>
<code>scriptengine.AsyncBinding</code> value (from <code>PromisifyAsync[T]</code>)	<code>Promise<T></code>
self-referential types	<code>unknown</code> (cycle detection bails after depth 4)

```
sercon --emit-dts sercon.d.ts
```

In your editor, place the resulting `.d.ts` next to your scripts (or reference it via `/// <reference path="..." />`) for autocomplete.

JSDoc comments on declarations

The emitter pulls JSDoc strings from the engine's doc map (set via `Engine.SetDocs` / `Engine.SetMemberDocs`) and renders them above each declaration. Single-line docs collapse to `/** ... */` multi-line docs expand to a standard `*`-prefixed block. Bindings without a doc entry emit no JSDoc — the emitter never inserts empty placeholders. Docs are looked up by the same dotted path the binding was registered under, so namespace members get hover text too:

```
// AUTO-GENERATED by scriptengine.Engine.WriteTypes – do not edit by hand.

/** Hashing primitives. */
declare const crypto: {
  hash: {
    /** SHA-256 hex digest of a UTF-8 input. */
    sha256(...args: unknown[]): string;
    /** BLAKE3 hex digest (32-byte output, lukechampine.com/blake3). */
    blake3(...args: unknown[]): string;
    // ...
  };
  // ...
};
```

The CLI's reserved-global surface ships with a curated doc map; see `cmd/sercon/docs.go` for the source of truth and add new entries there when adding new bindings (the lockstep rule keeps `--examples / MANUAL / sercon.d.ts` in sync, and the doc map is now the seventh artifact in that chain).

15. Limitations and gotchas

- **No native ES modules at runtime.** All `import` is rewritten to `require` by esbuild. Side-effect-only imports run the module body exactly once per Run.
- **No top-level await inside required modules.** Wrap with `(async () => ... )()` inside the module if you need it.
- **No `tsconfig.json` honoured.** Module resolution and type checking are independent from any project-level TS config.
- **Per-Run runtime.** Side effects on `globalThis` do not persist between calls to `Run` on the same `Engine`. The require *compile* cache is shared; module *exports* are not.
- **Multi-line ESM imports may stress the entry rewriter.** The current rewriter is line-based and handles the common forms; especially gnarly inputs (comments inside the spec list, very unusual whitespace) fall back to executing as-is, which may throw.
- **RegisterConstructor gives real new semantics.** `new Foo(...)` runs the registered Go constructor; JS arguments are coerced to its parameter types (an argument that can't be coerced becomes the zero value rather than throwing); the result's exported methods/fields are reachable (subject to the `json` -tag field mapper); and a non-nil trailing `error` result throws. A panic inside the constructor propagates like any Go binding (it is not converted to a catchable JS error), and `instanceof Foo` is not guaranteed (the instance is the Go-wrapped object). This is a library-side API — the CLI registers namespaces, not constructors.
- **HTTP bindings are real network calls.** `net.http.*` uses `net/http` with a 5s timeout. They are not mockable from JS.

See OUT-OF-SCOPE.md for the active backlog of deferred ideas.

Bundled libraries

`paymentproviders`

A TypeScript payments library compiled into the sercon binary and importable by scripts — no install needed:

```
import { Kcov3 } from "paymentproviders";
const api = Kcov3.client(); // reads KCO_* from the environment
const order = await api.getPayment(orderId);
await api.capturePayment(orderId, { amount: order.order_amount });
```

KCO v3 (`Kcov3`) wraps Kustom's Order-Management + Checkout APIs: `getPayment`, `acknowledge`, `capturePayment`, `refundPayment`, `cancelPayment`, `releaseRemainingAuthorization`, and (bonus) `createCheckout` / `getCheckout`. `Kcov3.client(overrides?)` reads credentials from the environment (or `.env` via `--env-file`): `KCO_MERCHANT_ID`, `KCO_SHARED_SECRET`, `KCO_ENV` (`test` | `prod`, selecting `api.playground.kustom.co` / `api.kustom.co`) or `KCO_BASE_URL` any can be overridden per call. Auth is HTTP Basic (`merchantId:sharedSecret`); POST mutations carry an auto-generated `Klarna-Idempotency-Key`. A non-2xx response throws a `PaymentError`.

(`provider`, `status`, `parsed body`, `requestId?`). Amounts are integer minor units (öre). Additional providers (Nets, Svea, Qliro, SwedbankPay) are planned as additive namespaces.

Additional providers (same `client()` + lifecycle shape):

- `netstv1` (Nexi/Nets Checkout Payment API v1) — `createPayment` / `getPayment` / `capturePayment` / `refundPayment` / `cancelPayment` env `NETS_SECRET_KEY`, `NETS_ENV` auth = the secret key in the `Authorization` header.
- `sveacheckout2` (Svea Checkout) — `createOrder` / `getOrder` / `capturePayment` / `refundPayment` / `cancelPayment` env `SCO_MERCHANT_ID`, `SCO_SECRET_KEY`, `SCO_ENV` auth = Svea `base64(merchantId:UPPER(SHA512(body+secret+timestamp)))` plus a `Timestamp` header.
- `qlirov2` (Qliro One) — `createOrder` / `getOrder` / `getPayment` env `QLIRO_API_KEY`, `QLIRO_APIPASSWORD`, `QLIRO_ENV` auth = Qliro `base64(SHA256(body+secret))`.
- `swedbankpayv3` / `swedbankpayv2` (SwedbankPay Checkout v3 / v2) — Bearer auth plus HAL/hypermedia: `createPaymentOrder`, `getPaymentOrder` / `getPayment`, the low-level operation(`paymentOrderUrl`, `rel`, `body?`) primitive (resolves the `rel`'s `href` in the payment-order response and POSTs to it), and convenience `capturePayment` / `refundPayment` / `cancelPayment` (the `capture` / `reversal` / `cancel` operations). Env `SWEDBANKPAY_ACCESS_TOKEN` (Bearer), `SWEDBANKPAY_MERCHANT_ID` (the `payee.payeeId`), `SWEDBANKPAY_ENV`. Base URLs `api.externalintegration.payex.com` (test) / `api.payex.com` (prod).

Provider namespaces are version-labelled (e.g. `kcov3`, `netstv1`) so a future API version ships alongside (e.g. `netstv2`) without breaking callers.

16. Binding reference (generated)

The per-function reference below is generated from the structured binding docs (`MemberDoc`) by `sercon --emit-reference`, spliced in by `make reference`. **Do not edit between the markers** — regenerate instead. Coverage grows per release as each namespace adopts the structured doc model (parameters, return shapes, errors, examples); namespaces not yet migrated show their one-line summary and a collapsed signature.

codec

Binary-format codecs: compression, barcodes, check digits.

codec.barcode.decodableFormats

```
decodableFormats(): string[]
```

Available decode formats (`qr` / `datamatrix` / `aztec` / `code128` / `code39` / `code93` / `codabar` / `ean13` / `ean8` / `upca` / `upce` / `itf`). `PDF417` is encode-only.

Returns: `string[]` — the twelve symbology names `barcode.decode` can recognise. `PDF417` is absent (`gozxing` has no `PDF417` decoder).

Throws: Never throws.

```
const fmts = codec.barcode.decodableFormats();
```

codec.barcode.decode

```
decode(data: string | Uint8Array | ArrayBuffer, format?: string):  
Promise<{ format: string, text: string }>
```

Decode a PNG/JPEG/WebP image to { format, text } via goxing. Optional format hint skips the auto-detect walk. EAN/UPC need a quiet zone in the input. Async.

Parameters

- `data` (*string* / *Uint8Array* / *ArrayBuffer*) — Image bytes (PNG / JPEG / WebP). A string is treated as its raw UTF-8 bytes.
- `format` (*string*, *optional*) — Symbology hint (case-insensitive) from `decodableFormats`. When given, only that reader runs; otherwise every decoder is tried in priority order and the first hit wins.

Returns: `Promise<{ format: string, text: string }>` — the detected symbology name and the decoded payload.

Throws: Throws if data is missing/empty or not a string/ArrayBuffer/Uint8Array, the image cannot be decoded, the format hint is unsupported, or no barcode is recognised.

```
const { format, text } = await codec.barcode.decode(png);
```

codec.barcode.encode

```
encode(format: string, data: string, opts?: { width?: number, height?: number,  
quietZone?: boolean | number }): Promise<Uint8Array>
```

Render data into a PNG of the chosen format. `opts.width` / `opts.height` default to 256x256 (2D) or 400x120 (1D). `opts.quietZone` (true or px count) pads a white margin — required for EAN/UPC to decode. Async.

Parameters

- `format` (*string*) — Symbology (case-insensitive): qr / datamatrix / aztec / pdf417 / code128 / code39 / codabar / ean13 / ean8 / upca.
- `data` (*string*) — Payload to encode. EAN/UPC require the exact digit count for the variant; the encoder validates content per symbology.
- `opts` (*{ width?: number, height?: number, quietZone?: boolean | number }, optional*) — `width` / `height` set the output pixel dimensions (default 256x256 for 2D qr/datamatrix/aztec, 400x120 otherwise). `quietZone` pads a white margin: true uses 10% of width (min 10px), a number uses that many pixels per side, false/0/absent adds none. EAN/UPC need a quiet zone to be decodable.

Returns: `Promise<Uint8Array>` — PNG image bytes.

Throws: Throws if the format is unknown, the data is invalid for that symbology, or scaling / PNG encoding fails.

```
const png = await codec.barcode.encode("qr", "https://example.com", { width: 320,  
height: 320 });
```

codec.barcode.formats

```
formats(): string[]
```

Available encode formats (qr / datamatrix / aztec / pdf417 / code128 / code39 / codabar / ean13 / ean8 / upca).

Returns: string[] — the ten symbology names accepted by barcode.encode.

Throws: Never throws.

```
const fmts = codec.barcode.formats(); // ["qr", "datamatrix", ...]
```

codec.checkdigit.algos

```
algos(): string[]
```

Supported algorithms (luhn / isbn10 / isbn13 / ean13 / ean8 / upca).

Returns: string[] — the six supported algorithm names. isbn13 is an alias for ean13 (same check-digit math).

Throws: Never throws.

```
const algos = codec.checkdigit.algos();
```

codec.checkdigit.compute

```
compute(algo: string, partial: string): string
```

Compute the missing trailing check digit for a partial input.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `partial` (*string*) — The number WITHOUT its check digit (whitespace trimmed). Fixed-length algorithms expect exactly length-1 digits (e.g. 12 for ean13, 9 for isbn10).

Returns: string — the single check digit ('0'–'9', or 'X' for isbn10 when the value is 10).

Throws: Throws if the algorithm is unknown, the input is empty / the wrong length, or contains a non-digit.

```
const cd = codec.checkdigit.compute("ean13", "123456789012"); // "8"
```

codec.checkdigit.inspect

```
inspect(algo: string, input: string): { algo: string, input: string, valid: boolean, given: string, computed: string }
```

Diagnostic combining validate + compute: { valid, given, computed, ... }.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `input` (*string*) — The full number including its trailing check digit (whitespace trimmed).

Returns: { algo, input, valid, given, computed } — algo/input echo the normalised arguments; given is the input's last character; computed is the recalculated check digit (empty when the input is too short or malformed to split); valid is true when given equals computed (case-insensitive).

Throws: Never throws; malformed input yields valid:false with an empty computed.

```
const r = codec.checkdigit.inspect("ean13", "1234567890128");
// { algo: "ean13", input: "...", valid: true, given: "8", computed: "8" }
```

codec.checkdigit.validate

```
validate(algo: string, input: string): boolean
```

Return whether the input passes the named algorithm's check digit.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive, trimmed): luhn / isbn10 / isbn13 / ean13 / ean8 / upca.
- `input` (*string*) — The full number including its trailing check digit (whitespace trimmed). ISBN-10 may end in 'X'.

Returns: boolean — true when the input's check digit is valid for the algorithm. Returns false (does not throw) for wrong length, non-digit characters, or an unknown algorithm.

Throws: Never throws; any failure (unknown algorithm included) returns false.

```
const ok = codec.checkdigit.validate("luhn", "4539578763621486"); // true
```

codec.compression.algos

```
algos(): string[]
```

Available compression algorithm names (gzip / deflate / zlib / bzip2 / zstd / brotli / lz4 / xz / snappy).

Returns: string[] — the nine supported algorithm names, lowercase, in registration order.

Throws: Never throws.

```
const algos = codec.compression.algos(); // ["gzip", "deflate", ...]
```

codec.compression.compress

```
compress(algo: string, data: string | Uint8Array | ArrayBuffer):
Promise<Uint8Array>
```

Compress data with the named algorithm. Returns Uint8Array. Async.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive): `gzip` / `deflate` / `zlib` / `bzip2` / `zstd` / `brotli` / `lz4` / `xz` / `snappy`.
- `data` (*string* / *Uint8Array* / *ArrayBuffer*) — Input bytes. Strings are interpreted as their UTF-8 byte sequence.

Returns: `Promise<Uint8Array>` — the compressed bytes.

Throws: Throws if data is undefined/null or an unsupported type, the algorithm name is unknown, or the underlying compressor errors.

```
const packed = await codec.compression.compress("gzip", "hello world");
```

`codec.compression.decompress`

```
decompress(algo: string, data: string | Uint8Array | ArrayBuffer):  
Promise<Uint8Array>
```

Decompress data previously produced by `compress` (same algorithm name required).
Returns `Uint8Array`. Async.

Parameters

- `algo` (*string*) — Algorithm name (case-insensitive), matching the one used to compress: `gzip` / `deflate` / `zlib` / `bzip2` / `zstd` / `brotli` / `lz4` / `xz` / `snappy`.
- `data` (*string* / *Uint8Array* / *ArrayBuffer*) — Compressed input bytes.

Returns: `Promise<Uint8Array>` — the original decompressed bytes.

Throws: Throws if data is undefined/null or an unsupported type, the algorithm name is unknown, or the input is not valid for that algorithm.

```
const raw = await codec.compression.decompress("gzip", packed);  
const text = new TextDecoder().decode(raw);
```

`codec.perl.dumper`

```
dumper(value: unknown, opts?: { classKey?: string, perlBoolClass?: string,  
indent?: string }): string
```

Perl `Data::Dumper`-style dump (`$VAR1 = ... ;`), normalized indentation. JS booleans emit the `JSON::XS::Boolean` blessed-ref form (`opts.perlBoolClass`).

Parameters

- `value` (*unknown*) — Any JSON-like value (see `php.serialize`). Arrays/objects emit as Perl array/hash refs; class-key objects emit as blessed refs.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `perlBoolClass` names the blessed class emitted for JS booleans (default `"JSON::XS::Boolean"`). `indent` overrides the indentation step. `classKey` overrides the class-name sentinel (default `"__class"`).

Returns: `string` — a `Data::Dumper`-style dump (`$VAR1 = ... ;`) with normalized indentation.

Throws: Throws on a circular reference or an unsupported value type.

```
const d = codec.perl.dumper({ ok: true });
```

codec.perl.parseDumper

```
parseDumper(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

Read Data::Dumper output back. Blessed scalar refs in the JSON bool family decode to booleans; bare 1/0 stay numbers; cycles throw.

Parameters

- `input` (*string*) — Data::Dumper output to parse back (a `$VARn = ...`; assignment or a bare value).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded blessed refs (default “`__class`”). The JSON bool family (JSON::XS::Boolean, JSON::PP::Boolean, Types::Serialiser::Boolean) decodes to JS booleans regardless.

Returns: `unknown` — the decoded value. Blessed scalar refs in the JSON-bool family become JS booleans; bare 1/0 stay numbers; other blessed refs carry the `classKey` sentinel.

Throws: Throws on malformed input or a reference that would close a cycle.

```
const v = codec.perl.parseDumper("$VAR1 = [1, 2];"); // [1, 2]
```

codec.php.parseVarDump

```
parseVarDump(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

Best-effort read of `var_dump()` output. Throws on lossy markers (*RECURSION*, truncation, visibility-annotated props).

Parameters

- `input` (*string*) — PHP `var_dump()` output to parse back.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded objects (default “`__class`”).

Returns: `unknown` — the reconstructed value (best-effort; `var_dump` is a lossy format).

Throws: Throws on lossy markers it cannot faithfully reverse: *RECURSION*, truncation, or visibility-annotated (private/protected) properties.

```
const v = codec.php.parseVarDump('int(42)'); // 42
```

codec.php.parseVarExport

```
parseVarExport(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

Read a `var_export()` literal (arrays, scalars, NULL, `\Cls::__set_state`) back to a value.

Parameters

- `input` (*string*) — A PHP `var_export()` literal: array(...), scalars, NULL, or `\Cls::__set_state(array(...))`.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded `__set_state` objects (default “`__class`”).

Returns: unknown — the decoded value; `\Cls::__set_state` objects become plain objects carrying the `classKey` sentinel.

Throws: Throws on input that is not a parseable `var_export()` literal.

```
const v = codec.php.parseVarExport("array (\n 0 => 1,\n)");
```

codec.php.serialize

```
serialize(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP `serialize()`: encode a value to PHP’s canonical serialization string. Objects use the `__class` sentinel; cycles throw.

Parameters

- `value` (*unknown*) — Any JSON-like value: null, boolean, number, string, array, or plain object. An object carrying the class-key sentinel (`opts.classKey`, default “`__class`”) encodes as a PHP object (O:).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` overrides the class-name sentinel property (default “`__class`”). `indent` / `perlBoolClass` are unused by `serialize`.

Returns: string — the PHP `serialize()` string (e.g. `a:1:{...}`).

Throws: Throws on a circular reference or an unsupported value type (function, Symbol, BigInt, Date, Map, Set, RegExp).

```
const s = codec.php.serialize({ a: 1, b: [2, 3] });
```

codec.php.unserialize

```
unserialize(input: string, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): unknown
```

PHP `unserialize()`: decode a `serialize()` string back to a value. `r:/R`: references resolve to shared objects (DAGs); cycles throw.

Parameters

- `input` (*string*) — A PHP `serialize()` string.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `classKey` sets the sentinel property used to tag decoded PHP objects (default “`__class`”).

Returns: unknown — the decoded value. PHP objects become plain objects carrying the classKey sentinel; r:/R: references rebuild as shared object identities (DAGs).

Throws: Throws on malformed input or a reference that would close a cycle.

```
const v = codec.php.unserialize('a:2:{i:0;i:1;i:1;i:2;}'); // [1, 2]
```

codec.php.varDump

```
varDump(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP var_dump(): human-readable debug output. String lengths are byte counts.

Parameters

- `value` (*unknown*) — Any JSON-like value (see `php.serialize`).
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `indent` overrides the default indentation step. `classKey` overrides the class-name sentinel (default “__class”).

Returns: string — `var_dump()`-style output. String lengths in the output are byte counts, matching PHP.

Throws: Throws on a circular reference or an unsupported value type.

```
const dump = codec.php.varDump({ name: "ok" });
```

codec.php.varExport

```
varExport(value: unknown, opts?: { classKey?: string, perlBoolClass?: string, indent?: string }): string
```

PHP var_export(): emit valid PHP code for a value. `opts.indent` overrides the 2-space step.

Parameters

- `value` (*unknown*) — Any JSON-like value (see `php.serialize`). Objects with the class-key sentinel emit as `\Cls::__set_state(...)`.
- `opts` (*{ classKey?: string, perlBoolClass?: string, indent?: string }, optional*) — `indent` overrides the default 2-space indentation step. `classKey` overrides the class-name sentinel (default “__class”).

Returns: string — valid PHP source, the kind `var_export()` prints.

Throws: Throws on a circular reference or an unsupported value type.

```
const code = codec.php.varExport({ x: 1 }, { indent: "  " });
```

codec.xml.decode

```
decode(xml: string): unknown
```


Parse an XML string to a value using the same @-prefix + #text convention as `xml.encode`. Attributes become @-keys, text becomes #text (or a bare string for a text-only element), child elements become keys, and repeated same-name siblings become an array. Empty/self-closing elements decode to null. Namespace prefixes are kept literally; all values are strings (no type coercion). Mismatched tags, multiple roots, and malformed XML throw.

Parameters

- `xml` (*string*) — The XML document to parse.

Returns: A single-key object whose key is the root element name and whose value is the parsed content (key order follows document order; all leaf values are strings).

Throws: Throws on malformed XML, mismatched/mis-nested end tags, multiple root elements, or no root element.

```
const v = codec.xml.decode("<note id=\"5\">hi</note>");
// { note: { "@id": "5", "#text": "hi" } }
```

codec.xml.encode

```
encode(value: unknown, opts?: { rootName?: string, indent?: string, declaration?: boolean }): string
```

Serialize a value to an XML string. Convention: @-prefixed keys are attributes, #text is element text, other keys are child elements, and an array value becomes repeated sibling elements (a scalar key becomes a text-only element, null a self-closing tag). The value must be a single-key object naming the root element, or pass `opts.rootName` to wrap it. Scalars are stringified; object key order is preserved.

Parameters

- `value` (*unknown*) — A single-key object whose one key names the root element — e.g. `{ note: { "@id": "5", "#text": "hi", to: "alice" } }` → `<note id="5">hi<to>alice</to></note>`. Or any value plus `opts.rootName` to wrap it. Cycles throw.
- `opts` (*{ rootName?: string, indent?: string, declaration?: boolean }, optional*) — `rootName` wraps the value under that root element. `indent` pretty-prints with the given unit per level (default compact). `declaration` prepends `<?xml version="1.0" encoding="UTF-8"?>` (default off).

Returns: The XML string.

Throws: Throws if the value has no single root element and no `opts.rootName`, if the root content is an array, if a non-scalar is used as an attribute or #text value, or if the value contains a cycle.

```
const xml = codec.xml.encode({ note: { "@id": "5", "#text": "hi" } });
// <note id="5">hi</note>
```

console

Browser/Node-style console shim: `log/info/debug` to `stdout`, `warn/error` to `stderr`. For porting scripts; `runtime.log` is the native equivalent.

console.debug

```
debug(args: ...unknown[]): void
```

Alias of console.log — stringified arguments, space-joined, to stdout.

Parameters

- `args (...unknown[])` — Values to print; identical formatting and stdout destination as console.log.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.debug("cache hit", key);
```

console.error

```
error(args: ...unknown[]): void
```

Like console.log but writes to stderr.

Parameters

- `args (...unknown[])` — Values to print; same space-joined / JSON formatting as console.log but routed to stderr.

Returns: void — output is written to stderr as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.error("request failed", { status: 500 });
```

console.info

```
info(args: ...unknown[]): void
```

Alias of console.log — stringified arguments, space-joined, to stdout.

Parameters

- `args (...unknown[])` — Values to print; identical formatting and stdout destination as console.log.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their String() form.

```
console.info("listening on", 8080);
```

console.log

```
log(args: ...unknown[]): void
```

Print a space-joined line of the arguments to stdout. Primitives print raw; objects/arrays render as JSON. Browser/Node-compatible; same output as runtime.log.

Parameters

- `args (...unknown[])` — Values to print, joined by single spaces and terminated with a newline. Primitives (string/number/boolean/null/undefined) print raw; objects and arrays render as JSON via `JSON.stringify`, falling back to `String()` for functions and circular references.

Returns: void — output is written to stdout as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their `String()` form.

```
console.log("user", { id: 1, name: "ada" }); // user {"id":1,"name":"ada"}
```

console.warn

```
warn(args: ...unknown[]): void
```

Like `console.log` but writes to stderr.

Parameters

- `args (...unknown[])` — Values to print; same space-joined / JSON formatting as `console.log` but routed to stderr.

Returns: void — output is written to stderr as a side effect.

Throws: Never throws; values JSON cannot serialise degrade to their `String()` form.

```
console.warn("retrying in", 5, "seconds");
```

crypto

Hashing, JWT, age encryption — anything that produces a digest, signature, or ciphertext.

crypto.encrypt.decrypt

```
decrypt(ciphertext: string | Uint8Array | ArrayBuffer, identities: string | string[]): Uint8Array
```

Open a payload with one of the supplied identities. Routes to age or PGP based on the identity / ciphertext format. age: binary or armored auto-detected. Wrong identity throws.

Parameters

- `ciphertext (string | Uint8Array | ArrayBuffer)` — The encrypted payload; age armor and PGP armor are auto-detected.
- `identities (string | string[])` — One or more private keys. AGE-SECRET-KEY-1... identities use the age backend; —BEGIN PGP PRIVATE KEY BLOCK— entries (or an armored PGP message) use the PGP backend.

Returns: `Uint8Array` — the decrypted plaintext bytes (decode with `new TextDecoder().decode(...)` for a string).

Throws: Throws if ciphertext is empty or an unsupported type, no identities are given, an identity is actually a public key, or no supplied identity matches the ciphertext's recipients.

```
const pt = crypto.encrypt.decrypt(ct, privateKey);
const text = new TextDecoder().decode(pt);
```

crypto.encrypt.detectBackend

```
detectBackend(input: string): { backend: string; kind?: string }
```

Classify a recipient / identity string. Returns { backend: 'age'|'pgp'|'unknown', kind?: 'public'|'private' }. Pure prefix matching; no parsing or I/O.

Parameters

- `input` (*string*) — A recipient or identity string to classify (age bech32, age SSH public key, or PGP armored block).

Returns: { backend: "age" | "pgp" | "unknown", kind?: "public" | "private" } — kind is present only when the backend was identified.

Throws: Never throws; unrecognised input returns { backend: "unknown" }.

```
const info = crypto.encrypt.detectBackend("age1abc..."); // { backend: "age",
kind: "public" }
```

crypto.encrypt.encrypt

```
encrypt(data: string | Uint8Array | ArrayBuffer, recipients: string | string[],
opts?: { armored?: boolean }): Uint8Array
```

Seal data to recipients. age public keys (age1...) → age backend (opts.armored for ASCII); PGP public-key blocks → PGP backend (always armored). Auto-dispatched on key format. Multi-recipient: any listed identity decrypts.

Parameters

- `data` (*string | Uint8Array | ArrayBuffer*) — Plaintext to encrypt.
- `recipients` (*string | string[]*) — One or more recipient public keys. age1... bech32 keys use the age backend; `—BEGIN PGP PUBLIC KEY BLOCK—` entries use the PGP backend. Backends cannot be mixed in one call.
- `opts` (*{ armored?: boolean }, optional*) — armored (age backend only) wraps the output in age ASCII armor; defaults to false (binary). Ignored on the PGP path, which is always armored.

Returns: Uint8Array — the ciphertext (binary age, ASCII-armored age when opts.armored, or armored PGP message bytes).

Throws: Throws if data is an unsupported type, no recipients are given, a recipient is actually a private key, or a recipient string fails to parse.

```
const ct = crypto.encrypt.encrypt("secret", publicKey, { armored: true });
```

crypto.encrypt.keygen

```
keygen(): { publicKey: string; privateKey: string }
```

Generate a fresh age X25519 keypair. Returns { publicKey: 'age1...', privateKey: 'AGE-SECRET-KEY-1...' }.

Returns: { publicKey: string, privateKey: string } — publicKey is the shareable age1... recipient; privateKey is the secret AGE-SECRET-KEY-1... identity.

Throws: Throws if the system RNG fails to generate an identity.

```
const { publicKey, privateKey } = crypto.encrypt.keygen();
```

crypto.encrypt.keygenPgp

```
keygenPgp(opts?: { name?: string, email?: string }): { publicKey: string;  
privateKey: string }
```

Generate a PGP keypair (RSA 2048). opts.name / opts.email populate the user ID. Returns armored { publicKey, privateKey } blocks. encrypt/decrypt auto-route to PGP when they see these.

Parameters

- **opts** (*{ name?: string, email?: string }, optional*) — name and email populate the primary user ID; both default to empty (fine for throwaway keys).

Returns: { publicKey: string, privateKey: string } — both ASCII-armored PGP key blocks (—BEGIN PGP PUBLIC/PRIVATE KEY BLOCK—).

Throws: Throws if entity generation or armor serialisation fails.

```
const { publicKey, privateKey } = crypto.encrypt.keygenPgp({ name: "Test", email:  
"t@example.com" });
```

crypto.encrypt.rekey

```
rekey(ciphertext: string | Uint8Array | ArrayBuffer, oldIdentities: string |  
string[], newRecipients: string | string[], opts?: { armored?: boolean }):  
Uint8Array
```

Re-encrypt for a new recipient set without exposing plaintext to JS. Output format defaults to match the input; opts.armored forces. Internal decrypt+encrypt loop.

Parameters

- **ciphertext** (*string | Uint8Array | ArrayBuffer*) — The existing age ciphertext to re-key (armor auto-detected).
- **oldIdentities** (*string | string[]*) — Current AGE-SECRET-KEY-1... private keys used to decrypt the existing ciphertext.
- **newRecipients** (*string | string[]*) — New age1... public keys to encrypt the result for.
- **opts** (*{ armored?: boolean }, optional*) — armored forces the output armor state; default preserves the input's armor state.

Returns: Uint8Array — the re-encrypted ciphertext for the new recipient set.

Throws: Throws if ciphertext is empty, either key list is empty, a key is the wrong kind (public vs private), or the decrypt step fails (no matching identity / malformed input). age backend only.

```
const rotated = crypto.encrypt.rekey(ct, oldKey, newPublicKey);
```

crypto.hash.blake3

```
blake3(input: string): string
```

BLAKE3 hex digest (32-byte output, lukechampion.com/blake3).

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 64-char lowercase hex digest (256-bit output).

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.blake3("hello");
```

crypto.hash.crc32

```
crc32(input: string): string
```

CRC-32 (IEEE polynomial), zero-padded to 8 hex chars.

Parameters

- `input` (*string*) — Data to checksum, interpreted as a UTF-8 byte sequence.

Returns: string — 8-char zero-padded lowercase hex checksum.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const c = crypto.hash.crc32("hello"); // "3610a686"
```

crypto.hash.md5

```
md5(input: string): string
```

MD5 hex digest of a UTF-8 input. Avoid for security purposes — exposed for compatibility with legacy fingerprints.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 32-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.md5("hello"); // "5d41402abc4b2a76b9719d911017c592"
```

crypto.hash.sha1

```
sha1(input: string): string
```

SHA-1 hex digest of a UTF-8 input. Avoid for security purposes.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 40-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha1("hello");
```

crypto.hash.sha256

```
sha256(input: string): string
```

SHA-256 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 64-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha256("hello");
```

crypto.hash.sha384

```
sha384(input: string): string
```

SHA-384 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 96-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha384("hello");
```

crypto.hash.sha3_256

```
sha3_256(input: string): string
```

SHA-3 256-bit hex digest. The underscore in the name matches `recon`'s binding.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 64-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha3_256("hello");
```

crypto.hash.sha3_512

```
sha3_512(input: string): string
```

SHA-3 512-bit hex digest.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 128-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha3_512("hello");
```

crypto.hash.sha512

```
sha512(input: string): string
```

SHA-512 hex digest of a UTF-8 input.

Parameters

- `input` (*string*) — Data to hash, interpreted as a UTF-8 byte sequence.

Returns: string — 128-char lowercase hex digest.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
const d = crypto.hash.sha512("hello");
```

crypto.jwt.sign

```
sign(claims: Record<string, unknown>, secret: string, opts?: { algorithm?: string }): string
```

Sign a claims object. `secret` is raw bytes for HS*; PEM-encoded private key for RS*/PS*/ES*/EdDSA; or a JWK JSON object (`kty` picks the key type) for any algorithm. `opts.algorithm` defaults to HS256.

Parameters

- `claims` (*Record<string, unknown>*) — Claims payload. Passed through to `MapClaims`; RFC 7519 reserved claims (`exp`/`nbf`/`iat`/`iss`/`aud`/`sub`) are honoured. Reserved claims are NOT synthesised — set `iat`/`exp` explicitly if you want them.
- `secret` (*string*) — Key material: raw HMAC bytes for HS256/384/512, a PEM-encoded private key (—BEGIN ...) for RS*/PS*/ES*/EdDSA, or a JWK JSON object (`{“kty”:...}`).

- `opts` (*{ algorithm?: string }, optional*) — algorithm names the RFC 7518 alg (HS256/HS384/HS512, RS256/384/512, PS256/384/512, ES256/384/512, EdDSA); case-insensitive. Defaults to HS256.

Returns: string — the signed compact-serialisation JWT (header.payload.signature).

Throws: Throws if claims is null, secret is empty, the algorithm is unsupported, or the secret shape mismatches the algorithm (e.g. HMAC algo with a PEM key, or asymmetric algo with plain bytes).

```
const tok = crypto.jwt.sign({ sub: "u1" }, "topsecret", { algorithm: "HS256" });
```

crypto.jwt.validate

```
validate(token: string, secret: string, opts?: { algorithm?: string, audience?: string, issuer?: string }): { valid: boolean; claims?: object; reason?: string }
```

Verify signature + standard time claims (exp/nbf; iat is not checked, per RFC 7519) + optional aud/iss. secret accepts raw bytes / PEM public key / JWK. Set `opts.algorithm` for the algo-confusion guard. Resolves `{ valid:true, claims }` or `{ valid:false, reason }`.

Parameters

- `token` (*string*) — The compact-serialisation JWT to verify.
- `secret` (*string*) — Verification key: raw HMAC bytes, a PEM-encoded public key (or certificate), or a JWK JSON object.
- `opts` (*{ algorithm?: string, audience?: string, issuer?: string }, optional*) — algorithm pins the accepted alg (algorithm-confusion guard); when unset, any supported alg is accepted. audience / issuer enforce the aud / iss claims when set.

Returns: `{ valid: true, claims: object }` on success, or `{ valid: false, reason: string }` on any verification / claim failure.

Throws: Throws only on structural / wiring errors (empty token or secret, malformed token, or a secret shape that mismatches the token's algorithm). Cryptographic and claim failures resolve as `{ valid: false, reason }` instead of throwing.

```
const r = crypto.jwt.validate(tok, "topsecret", { algorithm: "HS256" });
if (r.valid) runtime.log(r.claims.sub);
```

crypto.jwt.view

```
view(token: string): { header: object; payload: object; signature: string }
```

Decode header + payload WITHOUT verifying the signature. Useful for inspection / debugging auth flows. Malformed input throws.

Parameters

- `token` (*string*) — A compact-serialisation JWT (three dot-separated base64url segments).

Returns: `{ header: object, payload: object, signature: string }` — decoded header and payload (object key order preserved) plus the raw signature segment.

Throws: Throws if the token is not three dot-separated segments or a segment fails base64url / JSON decoding.

```
const { header, payload } = crypto.jwt.view(tok);
```

db

Database / KV / directory clients: SQLite, PostgreSQL, MySQL/MariaDB, SQL Server, Redis, Valkey, memcached, LDAP, dict.

db.clickhouse.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?: string, database?: string, secure?: boolean }): Promise<Record<string, unknown>>
```

Connect to ClickHouse via the pure-Go clickhouse-go v2 driver. Arg is a clickhouse:

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string*, *secure?: boolean* }) — A clickhouse:

Returns: Promise<handle> resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise<any>` (first column of the first row, or null); `begin() → Promise<tx>`; `prepare(sql) → Promise<stmt>`; `close() → Promise<void>`. ClickHouse uses ? positional placeholders (or @name).

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.clickhouse.open({ host: "localhost", database: "metrics",
secure: false });
const rows = await db.query("SELECT name, value FROM stats WHERE host = ?",
"web1");
await db.close();
```

db.dict.define

```
define(host: string, word: string, opts?: { database?: string, port?: string,
timeout?: number }): Promise<{ word: string; found: boolean; definitions: { db:
string; dbName: string; text: string }[]>
```

RFC 2229 DICT word lookup. `define(host, word, opts?) -> { word, found, definitions: [{ db, dbName, text }] }`. `found:false` on no match (not an error). One-shot: connect, query, QUIT.

Parameters

- `host` (*string*) — The DICT server hostname.
- `word` (*string*) — The word to look up.
- `opts` ({ *database?: string*, *port?: string*, *timeout?: number* }, *optional*) — database selects a specific dictionary (default "*" = all); port is the DICT port (default "2628"); timeout is the dial/read deadline in milliseconds (default 10000).

Returns: Promise<{ word: string, found: boolean, definitions: { db: string, dbName: string, text: string }[] }> — definitions carries one entry per matching dictionary (db is the dictionary code, dbName its human name, text the definition body). A word with no definitions resolves with found:false and an empty list.

Throws: Throws if host or word is empty, on dial/banner failure, or on an unexpected DICT status code (e.g. 550 invalid database). A 552 “no match” is NOT an error — it resolves with found:false.

```
const r = await db.dict.define("dict.org", "serendipity");
if (r.found) runtime.log(r.definitions[0].text);
```

db.dict.match

```
match(host: string, word: string, opts?: { strategy?: string, database?: string,
port?: string, timeout?: number }): Promise<{ word: string; matches: { db: string;
word: string }[] }>
```

RFC 2229 word match. match(host, word, opts?) -> { word, matches: [{ db, word }] }.
opts.strategy (default prefix), opts.database (default *), opts.port (default 2628). One-shot: connect, query, QUIT.

Parameters

- `host` (*string*) — The DICT server hostname.
- `word` (*string*) — The word (or pattern) to match.
- `opts` (*{ strategy?: string, database?: string, port?: string, timeout?: number }, optional*) — strategy is the match strategy (default “prefix”); database selects a specific dictionary (default “*” = all); port is the DICT port (default “2628”); timeout is the dial/read deadline in milliseconds (default 10000).

Returns: Promise<{ word: string, matches: { db: string, word: string }[] }> — matches carries one entry per matched word (db is the dictionary it was found in, word the matched headword). No matches resolves with an empty matches list.

Throws: Throws if host or word is empty, on dial/banner failure, or on an unexpected DICT status code. A 552 “no match” is NOT an error — it resolves with an empty matches list.

```
const r = await db.dict.match("dict.org", "seren", { strategy: "prefix" });
runtime.log(r.matches.map(m => m.word));
```

db.ldap.open

```
open(url: string, opts?: { bindDN?: string, password?: string }):
Promise<Record<string, unknown>>
```

Dial LDAP (ldap:

Parameters

- `url` (*string*) — An LDAP URL: ldap:
- `opts` (*{ bindDN?: string, password?: string }, optional*) — When bindDN is set, the connection binds with bindDN/password instead of doing an anonymous bind.

Returns: Promise<handle> resolving to { rootDSE, search, close }: rootDSE() → Promise<object> reads the server's Root DSE (an ordered { dn, <attr>: string[] } object advertising naming contexts, supported controls, vendor, etc.; an empty object when the server returns no entry); search(baseDN, filter, attrs?) → Promise<object[]> runs a whole-subtree search and returns one ordered { dn, <attr>: string[] } object per entry (multi-valued attributes stay arrays; filter defaults to (objectClass=*); attrs is an optional array of attribute names); close() → Promise<void>. A directory-inspection (read) binding, not a write/modify surface.

Throws: open throws if url is empty, the dial fails, or (when bindDN is set) the bind fails (the connection is closed on bind failure). rootDSE / search throw on the underlying LDAP search error.

```
const dir = await db.ldap.open("ldap://localhost:389");
const meta = await dir.rootDSE();
const people = await dir.search("ou=people,dc=example,dc=com", "(uid=alice)",
["cn", "mail"]);
await dir.close();
```

db.memcached.open

```
open(addr: string): Promise<Record<string, unknown>>
```

Connect to memcached (host:port). Returns { get, set, delete }. get -> string or null (miss); delete -> bool (existed). set(key, value, expirySeconds?). No ping on open; the pool is lazy.

Parameters

- `addr` (*string*) — A memcached server address, host:port (e.g. localhost:11211).

Returns: Promise<handle> resolving to { get, set, delete }: get(key) → Promise<string | null> (null on a cache miss); set(key, value, expirySeconds?) → Promise<void> (value stored as bytes; expirySeconds 0 or omitted means never expire); delete(key) → Promise<boolean> (true if the key existed, false on a miss). gomemcache pools connections lazily, so there is no ping-on-open and no close method (the pool is GC'd with the handle).

Throws: open throws if addr is empty. set throws if key is empty or the value cannot be coerced to bytes. get / delete throw on transport errors; a cache miss is data (get → null, delete → false), not an error.

```
const mc = await db.memcached.open("localhost:11211");
await mc.set("session:42", "active", 300);
const v = await mc.get("session:42"); // "active" or null
const existed = await mc.delete("session:42"); // true
```

db.mssql.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string }): Promise<Record<string, unknown>>
```

Connect to Microsoft SQL Server via the pure-Go go-mssqldb driver. Arg is a sqlserver:

Parameters

- `dsn (string | { host?: string, port?: number, user?: string, password?: string, database?: string })` — A sqlserver:

Returns: Promise<handle> resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise<any>` (first column of the first row, or null); `begin() → Promise<tx>`; `prepare(sql) → Promise<stmt>`; `close() → Promise<void>`. SQL Server uses @p1, @p2, ... placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.mssql.open({ host: "localhost", user: "sa", password: "P@ss",
  database: "shop" });
const rows = await db.query("SELECT TOP 5 id, name FROM users WHERE region = @p1",
  "EU");
await db.close();
```

db.mysql.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
  string, database?: string }): Promise<Record<string, unknown>>
```

Connect to MySQL/MariaDB via the pure-Go go-sql-driver. Arg is a go-sql-driver DSN string (user:pass@tcp(host:port)/db) or an options object { host, port, user, password, database }. Returns the shared SQL handle. Uses ? placeholders. Pings on open.

Parameters

- `dsn (string | { host?: string, port?: number, user?: string, password?: string, database?: string })` — A go-sql-driver DSN string (user:pass@tcp(host:port)/db?params) used verbatim, OR an options object assembled into one (defaults: host localhost, port 3306, empty database). One driver serves both MySQL and MariaDB.

Returns: Promise<handle> resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise<any>` (first column of the first row, or null); `begin() → Promise<tx>`; `prepare(sql) → Promise<stmt>`; `close() → Promise<void>`. MySQL uses ? positional placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.mysql.open("app:s3cr3t@tcp(localhost:3306)/shop");
const n = await db.queryValue("SELECT COUNT(*) FROM orders WHERE status = ?",
  "paid");
await db.close();
```

db.oracle.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
  string, database?: string }): Promise<Record<string, unknown>>
```

Connect to Oracle via the pure-Go go-ora driver (no cgo). Arg is an oracle:

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string* }) — An oracle:

Returns: `Promise<handle>` resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row); `queryValue(sql, ...params) → Promise<any>` (first column of the first row, or null); `begin() → Promise<tx>`; `prepare(sql) → Promise<stmt>`; `close() → Promise<void>`. Oracle uses `:1`, `:2`, ... (or `:name`) bind placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails.

```
const db = await db.oracle.open({ host: "localhost", user: "app", password:
"s3cr3t", database: "ORCLPDB1" });
const rows = await db.query("SELECT id, name FROM users WHERE id = :1", 42);
await db.close();
```

db.postgres.open

```
open(dsn: string | { host?: string, port?: number, user?: string, password?:
string, database?: string, sslmode?: string }): Promise<Record<string, unknown>>
```

Connect to PostgreSQL via the pure-Go pgx driver. Arg is a libpq DSN/URL string or an options object { `host`, `port`, `user`, `password`, `database`, `sslmode` }. Returns the shared SQL handle { `exec`, `query`, `queryValue`, `begin`, `prepare`, `close` }. Uses `$1`, `$2`, ... placeholders. Pings on open.

Parameters

- `dsn` (*string* | { *host?: string*, *port?: number*, *user?: string*, *password?: string*, *database?: string*, *sslmode?: string* }) — A libpq DSN/URL string used verbatim, OR an options object assembled into a postgres:

Returns: `Promise<handle>` resolving to the shared SQL handle: `exec(sql, ...params) → Promise<{ rowsAffected, lastInsertId }>`; `query(sql, ...params) → Promise<object[]>` (one ordered object per row, keyed by column name in column order); `queryValue(sql, ...params) → Promise<any>` (first column of the first row, or null); `begin() → Promise<tx>` ({ `exec`, `query`, `queryValue`, `commit`, `rollback` }); `prepare(sql) → Promise<stmt>` ({ `exec`, `query`, `queryValue`, `close` }); `close() → Promise<void>`. Postgres uses `$1`, `$2`, ... positional placeholders.

Throws: Throws if no argument is given, the DSN string is empty, the argument is neither a string nor an object, or the connection ping fails (the pool is closed on ping failure).

```
const db = await db.postgres.open({ host: "localhost", user: "app", password:
"s3cr3t", database: "shop", sslmode: "disable" });
const rows = await db.query("SELECT id, name FROM users WHERE id = $1", 42);
await db.close();
```

db.redis.open

```
open(url: string): Promise<Record<string, unknown>>
```

Connect to Redis (redis:

Parameters

- `url` (*string*) — A standard Redis URL: redis:

Returns: Promise<handle> resolving to { do, ping, close }: do(cmd, ...args) → Promise<any> runs an arbitrary RESP command (the first arg is the command name, the rest its arguments) and returns the reply coerced to a JS value — strings, numbers, arrays, or null; a nil reply (missing key) resolves to null rather than throwing. ping() → Promise<string> ('PONG'). close() → Promise<void>.

Throws: open throws if url is empty, the URL fails to parse, or the open-time ping fails (the client is closed on ping failure). do throws on Redis-level errors (WRONGTYPE, unknown command, etc.); a missing-key nil reply is data, not an error.

```
const r = await db.redis.open("redis://localhost:6379/0");
await r.do("SET", "greeting", "hi");
const v = await r.do("GET", "greeting"); // "hi"
const missing = await r.do("GET", "nope"); // null
await r.close();
```

db.sqlite.open

```
open(path: string): Promise<Record<string, unknown>>
```

Open a SQLite database (':memory:' or a file path; created if absent). Resolves to a handle { exec, query, queryValue, begin, prepare, close }. Connection is Ping-ed before resolving.

Parameters

- `path` (*string*) — “:memory:” for an in-RAM database, or a filesystem path. Missing files are created by the modernc.org/sqlite (pure-Go, no cgo) driver.

Returns: Promise<handle> resolving to the shared SQL handle object: exec(sql, ...params) → Promise<{ rowsAffected: number, lastInsertId: number }>; query(sql, ...params) → Promise<object[]> (one ordered object per row, keyed by column name in column order); queryValue(sql, ...params) → Promise<any> (first column of the first row, or null when no rows match); begin() → Promise<tx> ({ exec, query, queryValue, commit, rollback }); prepare(sql) → Promise<stmt> ({ exec, query, queryValue, close }); close() → Promise<void>. SQLite uses ? positional placeholders. UTF-8 byte columns scan to strings; genuinely binary bytes surface as Uint8Array.

Throws: Throws if path is missing or empty, or if the connection ping fails (the *sql.DB is closed on ping failure rather than leaked). Subsequent exec/query/etc. throw the driver error on bad SQL or bind mismatch.

```
const db = await db.sqlite.open(":memory:");
await db.exec("CREATE TABLE t (id INTEGER PRIMARY KEY, name TEXT)");
await db.exec("INSERT INTO t (name) VALUES (?)", "alice");
```

```
const rows = await db.query("SELECT * FROM t WHERE name = ?", "alice");
await db.close();
```

db.valkey.open

```
open(url: string): Promise<Record<string, unknown>>
```

Connect to Valkey (valkey:

Parameters

- `url` (*string*) — A connection URL: valkey:

Returns: `Promise<handle>` resolving to `{ do, ping, close }`: `do(cmd, ...args) → Promise<any>` runs an arbitrary RESP command (the first arg is the command name, the rest its arguments) and returns the reply coerced to a JS value — strings, numbers, arrays, or null; a nil reply (missing key) resolves to null rather than throwing. `ping() → Promise<string>` ('PONG'). `close() → Promise<void>`.

Throws: `open` throws if `url` is empty, the URL fails to parse, or the open-time ping fails (the client is closed on ping failure). `do` throws on RESP-level errors (WRONGTYPE, unknown command, etc.); a missing-key nil reply is data, not an error.

```
const r = await db.valkey.open("valkey://localhost:6379/0");
await r.do("SET", "greeting", "hi");
const v = await r.do("GET", "greeting"); // "hi"
await r.close();
```

fs

Filesystem operations: path manipulation and archive create/extract.

fs.archive.create

```
create(destPath: string, sources: (string | { path: string, name?: string })[]):
Promise<{ path: string; format: string; entries: string[]; bytes?: number }>
```

Create a zip / tar / tar.gz at `destPath` from a list of paths. Format inferred from extension.

Parameters

- `destPath` (*string*) — Output archive path. Format is inferred from the extension: `.zip`, `.tar`, `.tar.gz`, or `.tgz`.
- `sources` (*((string | { path: string, name?: string })[])*) — Non-empty array of inputs. A bare string uses the disk path as-is and its basename inside the archive; an object overrides the in-archive name via `name`. Directory sources are recursed (the directory's basename becomes the archive subdir). Archive paths always use forward slashes.

Returns: `Promise<{ path: string, format: string, entries: string[], bytes?: number }>` — `path` is `destPath`, `format` is the inferred format ("`zip`" | "`tar`" | "`tar.gz`"), `entries` lists the file paths written (directories excluded), and `bytes` is the final archive size when `stat` succeeds.

Throws: Rejects if `destPath` is empty, `sources` is not an array / is empty / contains a bad entry (object missing 'path', unsupported element type), the format cannot be inferred from the extension, or any disk read / write fails (e.g. a source path does not exist).


```
const r = await fs.archive.create("out.tar.gz", ["dist", { path: "README.md",
name: "docs/README.md" }]);
runtime.log(r.format, r.entries.length);
```

fs.archive.extract

```
extract(archivePath: string, destDir: string, opts?: { overwrite?: boolean }):
Promise<{ path: string; format: string; dest: string; entries: string[] }>
```

Extract a zip / tar / tar.gz to destDir. opts.overwrite controls O_EXCL behaviour.

Parameters

- `archivePath` (*string*) — Path to the archive. Format is inferred from its extension (.zip, .tar, .tar.gz, .tgz).
- `destDir` (*string*) — Destination directory; created (recursively) if absent. All entries are confined to this directory via zip-slip / tar-slip protection.
- `opts` (*{ overwrite?: boolean }, optional*) — overwrite (default false) clobbers existing files; when false, an entry colliding with an existing file fails the call (O_EXCL).

Returns: Promise<{ path: string, format: string, dest: string, entries: string[] }> — path is archivePath, format is the inferred format, dest is destDir, and entries lists the extracted entry names (regular files only).

Throws: Rejects if archivePath or destDir is empty, the format cannot be inferred, destDir cannot be created, the archive cannot be opened / decoded, an entry escapes destDir (absolute path or '..' component), or (with overwrite false) an entry collides with an existing file.

```
const r = await fs.archive.extract("out.tar.gz", "./unpacked", { overwrite:
true });
runtime.log(r.entries.length, "files extracted");
```

fs.exists

```
exists(path: string): Promise<boolean>
```

Report whether a path exists. Never throws for a missing path.

Parameters

- `path` (*string*) — Path to test.

Returns: Promise resolving to true if the path exists, false if it does not.

Throws: Rejects only on an unexpected stat error (e.g. a permission error on a parent); a missing path resolves to false.

```
if (!(await fs.exists("report"))) await fs.mkdir("report");
```

fs.mkdir

```
mkdir(path: string): Promise<{ path: string }>
```

Create a directory, including any missing parents (`mkdir -p`). Idempotent.

Parameters

- `path` (*string*) — Directory path to create (mode 0755). Existing directories are fine.

Returns: Promise resolving to { `path` }.

Throws: Rejects if path is missing/empty or creation fails (e.g. a path component is an existing file).

```
await fs.mkdir("report/assets");
```

fs.path.basename

```
basename(path: string, suffix?: string): string
```

Final segment of a path; optional suffix is stripped if it matches.

Parameters

- `path` (*string*) — A forward-slash path. Trailing slashes are stripped before taking the last segment.
- `suffix` (*string, optional*) — Trailing suffix to remove from the result (e.g. an extension). Only stripped when it matches and is not the entire segment; a non-matching or empty suffix is ignored.

Returns: string — the last path segment, with suffix removed when it applies.

Throws: Throws a `TypeError` if path is missing, null, or undefined. suffix is coerced to a string and is optional.

```
const b = fs.path.basename("/var/log/app.log", ".log"); // "app"
```

fs.path.dirname

```
dirname(path: string): string
```

Directory portion of a path. POSIX-style; trailing slashes are stripped.

Parameters

- `path` (*string*) — A forward-slash path. On Windows, normalise separators yourself first.

Returns: string — everything up to (not including) the final slash; “.” when the path has no directory component, “/” for a rooted single segment.

Throws: Throws a `TypeError` if path is missing, null, or undefined.

```
const d = fs.path.dirname("/var/log/app.log"); // "/var/log"
```

fs.readBytes

```
readBytes(path: string): Promise<Uint8Array>
```

Read an entire file as bytes.

Parameters

- `path` (*string*) — File to read (CWD-relative or absolute).

Returns: Promise resolving to the file contents as a Uint8Array.

Throws: Rejects if path is missing/empty or the file cannot be read.

```
const bytes = await fs.readBytes("shot.png");
```

fs.readText

```
readText(path: string): Promise<string>
```

Read an entire file as a UTF-8 string.

Parameters

- `path` (*string*) — File to read (CWD-relative or absolute).

Returns: Promise resolving to the file contents decoded as UTF-8.

Throws: Rejects if path is missing/empty or the file cannot be read (absent, permissions).

```
const html = await fs.readText("report/index.html");
```

fs.remove

```
remove(path: string): Promise<{ path: string }>
```

Remove a file or a directory tree (recursive). No error if the path is already absent.

Parameters

- `path` (*string*) — File or directory to remove. Directories are removed recursively.

Returns: Promise resolving to { path }.

Throws: Rejects only if removal fails (e.g. permissions); removing an absent path is a no-op.

```
await fs.remove("report"); // clean slate
```

fs.stat

```
stat(path: string): Promise<{ size: number; isDir: boolean; modifiedMs: number }>
```

File metadata: size, whether it is a directory, and last-modified time.

Parameters

- `path` (*string*) — Path to stat.

Returns: Promise resolving to { size (bytes), isDir, modifiedMs (epoch milliseconds) }.

Throws: Rejects if path is missing/empty or the target does not exist.

```
const st = await fs.stat("report/index.html");
runtime.log(st.size, st.isDir, st.modifiedMs);
```

fs.writeBytes

```
writeBytes(path: string, data: Uint8Array): Promise<{ path: string; bytes:
number }>
```

Write binary data (a Uint8Array) to a file, truncating. Fails if the parent directory does not exist.

Parameters

- `path` (*string*) — Output file path (CWD-relative or absolute).
- `data` (*Uint8Array*) — Bytes to write (mode 0644). Pass a Uint8Array (e.g. `new Uint8Array(shot.bytes)`).

Returns: Promise resolving to { path, bytes } where bytes is the number of bytes written.

Throws: Rejects if path is missing/empty, data is not a Uint8Array, the parent directory does not exist, or the write fails.

```
const shot = await d.screenshot();
await fs.writeBytes("shot.png", new Uint8Array(shot.bytes));
```

fs.writeText

```
writeText(path: string, text: string): Promise<{ path: string; bytes: number }>
```

Write a string to a file (UTF-8, truncating). Fails if the parent directory does not exist — call `fs.mkdir` first.

Parameters

- `path` (*string*) — Output file path (CWD-relative or absolute). Used as given; no sandboxing.
- `text` (*string*) — Content to write as UTF-8. The file is created (mode 0644) or truncated.

Returns: Promise resolving to { path, bytes } where bytes is the number of bytes written.

Throws: Rejects if path is missing/empty, text is not a string, the parent directory does not exist, or the write fails (permissions, etc.).

```
await fs.mkdir("report");
await fs.writeText("report/index.html", "<h1>Hi</h1>");
```

image

Image decode/encode (PNG/JPEG/GIF/TIFF/BMP/WebP, SVG rasterize-in) and a chainable Image handle: `resize/crop/rotate/flip/adjust/filter/compose`.

image.decode

```
decode(data: Uint8Array): { readonly width: number; readonly height: number;
readonly format: string; resize(width: number, height: number, opts?: { filter?:
"lanczos" | "nearest" | "linear" | "box" | "catmullrom" }): unknown; fit(width:
```

```
number, height: number): unknown; thumbnail(width: number, height: number):
unknown; crop(x: number, y: number, w: number, h: number): unknown;
rotate(degrees: number): unknown; rotate90(): unknown; rotate180(): unknown;
rotate270(): unknown; flipH(): unknown; flipV(): unknown; brightness(percent:
number): unknown; contrast(percent: number): unknown; gamma(gamma: number):
unknown; saturation(percent: number): unknown; sharpen(sigma: number): unknown;
blur(sigma: number): unknown; grayscale(): unknown; invert(): unknown;
overlay(other: unknown, x: number, y: number, opacity?: number): unknown;
paste(other: unknown, x: number, y: number): unknown; bytes(format: "png" | "jpeg"
| "gif" | "tiff" | "bmp" | "webp", opts?: { quality?: number }): Uint8Array;
save(path: string, opts?: { format?: string; quality?: number }): void }
```

Decode in-memory image bytes into a chainable Image handle. The format is sniffed from the magic bytes (PNG/JPEG/GIF/TIFF/BMP/WebP); GIF decodes the first frame only.

Parameters

- `data` (*Uint8Array*) — The raw, encoded image bytes (e.g. from net.http, fs, or a clipboard read).

Returns: Image — a handle exposing read-only width/height/format and chainable transform methods.

Throws: Throws a `TypeError` if data is not a `Uint8Array`, or (“image.decode: ...”) if the bytes are not a recognised/decodable image.

```
const im = image.decode(pngBytes);
```

image.open

```
open(path: string): { readonly width: number; readonly height: number; readonly
format: string; resize(width: number, height: number, opts?: { filter?: "lanczos"
| "nearest" | "linear" | "box" | "catmullrom" }): unknown; fit(width: number,
height: number): unknown; thumbnail(width: number, height: number): unknown;
crop(x: number, y: number, w: number, h: number): unknown; rotate(degrees:
number): unknown; rotate90(): unknown; rotate180(): unknown; rotate270(): unknown;
flipH(): unknown; flipV(): unknown; brightness(percent: number): unknown;
contrast(percent: number): unknown; gamma(gamma: number): unknown;
saturation(percent: number): unknown; sharpen(sigma: number): unknown; blur(sigma:
number): unknown; grayscale(): unknown; invert(): unknown; overlay(other: unknown,
x: number, y: number, opacity?: number): unknown; paste(other: unknown, x: number,
y: number): unknown; bytes(format: "png" | "jpeg" | "gif" | "tiff" | "bmp" |
"webp", opts?: { quality?: number }): Uint8Array; save(path: string, opts?:
{ format?: string; quality?: number }): void }
```

Read an image file from disk and decode it into a chainable Image handle. The format is sniffed from the file’s magic bytes (PNG/JPEG/GIF/TIFF/BMP/WebP), not the extension. GIF decodes the first frame only.

Parameters

- `path` (*string*) — Filesystem path to the image file to read and decode.

Returns: Image — a handle exposing read-only width/height/format and chainable transform methods.

Throws: Throws (“image.open: ...”) if the file cannot be read, or (“image.decode: ...”) if the bytes are not a recognised/decodable image.

```
const im = image.open("avatar.jpg");
```

image.rasterizeSVG

```
rasterizeSVG(src: string | Uint8Array, opts: { width: number; height: number }):  
{ readonly width: number; readonly height: number; readonly format: string;  
resize(width: number, height: number, opts?: { filter?: "lanczos" | "nearest" |  
"linear" | "box" | "catmullrom" }): unknown; fit(width: number, height: number):  
unknown; thumbnail(width: number, height: number): unknown; crop(x: number, y:  
number, w: number, h: number): unknown; rotate(degrees: number): unknown;  
rotate90(): unknown; rotate180(): unknown; rotate270(): unknown; flipH(): unknown;  
flipV(): unknown; brightness(percent: number): unknown; contrast(percent: number):  
unknown; gamma(gamma: number): unknown; saturation(percent: number): unknown;  
sharpen(sigma: number): unknown; blur(sigma: number): unknown; grayscale():  
unknown; invert(): unknown; overlay(other: unknown, x: number, y: number,  
opacity?: number): unknown; paste(other: unknown, x: number, y: number): unknown;  
bytes(format: "png" | "jpeg" | "gif" | "tiff" | "bmp" | "webp", opts?: { quality?:  
number }): Uint8Array; save(path: string, opts?: { format?: string; quality?:  
number }): void }
```

Rasterize an SVG (a supported subset) to a raster Image at the requested pixel size. SVG is rasterize-IN only — there is no SVG output; the result is a raster image you then encode as PNG/JPEG/etc. The accepts a path string or in-memory SVG bytes.

Parameters

- `src` (*string* | *Uint8Array*) — An SVG file path (string) or the raw SVG document bytes (Uint8Array).
- `opts` (*{ width: number; height: number }*) — Target raster size in pixels; both width and height are required and must be > 0.

Returns: Image — a raster handle (format reported as “svg”) sized to `opts.width` × `opts.height`.

Throws: Throws a `TypeError` if `src` is neither a string nor `Uint8Array`, or if width/height are missing/non-positive; throws (“image: svg parse: ...”) on a malformed/unsupported SVG.

```
const im = image.rasterizeSVG("logo.svg", { width: 256, height: 256 });
```

net

Network clients and probes: HTTP, TCP/DNS/TLS/NTP/WHOIS probes, netstatus, email auth, browser-style sessions.

net.browser.open

```
open(...args: unknown[]): Promise<{ setUserAgent(ua: string): void,  
setHeader(name: string, value: string): void, get(url: string): Promise<{ status:  
number, ok: boolean, headers: Record<string, string>, body: string, url: string }  
>, post(url: string, body?: string): Promise<{ status: number, ok: boolean,
```

```
headers: Record<string, string>, body: string, url: string }>, cookies(url:
string): Promise<{ name: string, value: string }[]> }>
```

Open a stateful HTTP session: { setUserAgent, setHeader, get, post, cookies }. Cookie jar + default headers persist across requests (like a browser).

Returns: Promise<{ setUserAgent(ua: string): void, setHeader(name: string, value: string): void, get(url: string): Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, url: string }>, post(url: string, body?: string): Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, url: string }>, cookies(url: string): Promise<{ name: string, value: string }[]> }> — a session handle backed by an http.Client with an automatic cookie jar (public-suffix scoped). setUserAgent/setHeader register default headers replayed on every request; get/post return the same result shape as net.http.request; cookies lists the jar's cookies for a URL.

Throws: browser.open rejects only if the cookie jar can't be created. get/post reject if the URL is empty or the request fails (transport error); cookies rejects if the URL is empty or unparseable. 4xx/5xx responses resolve normally.

```
const b = await net.browser.open();
b.setUserAgent("my-bot/1.0");
await b.post("https://site/login", "user=x&pass=y");
const home = await b.get("https://site/home");
runtime.log(await b.cookies("https://site/"));
```

net.capture.interfaces

```
interfaces(): { name: string; addresses: string[]; up: boolean; loopback:
boolean }[]
```

List the host's network interfaces synchronously: net.capture.interfaces() → array of { name, addresses: string[], up, loopback }. Pure-Go (no privileges, all platforms).

Returns: { name: string, addresses: string[], up: boolean, loopback: boolean }[] — one entry per interface with its name, assigned addresses (CIDR strings), and up / loopback flags. Synchronous (not a Promise).

Throws: Throws if interface enumeration fails.

```
for (const i of net.capture.interfaces()) runtime.log(i.name, i.up);
```

net.capture.open

```
open(opts: { iface: string, promisc?: boolean, snaplen?: number, filter?:
string }, onPacket: (pkt: { ts: number, length: number, captureLength: number,
link: string, eth?: { src: string, dst: string, type: string }, vlan?: { id:
number, priority: number, drop: boolean, type: string }, arp?: { operation:
string, senderMac: string, senderIp: string, targetMac: string, targetIp:
string }, ip?: { version: number, src: string, dst: string, protocol: string, ttl:
number }, tcp?: { srcPort: number, dstPort: number, seq: number, ack: number,
window: number, checksum: number, flags: { syn: boolean, ack: boolean, fin:
boolean, rst: boolean, psh: boolean, urg: boolean }, options?: { mss?: number,
```

```

windowScale?: number, sackPermitted?: boolean, timestamps?: { val: number, ecr:
number } } }, udp?: { srcPort: number, dstPort: number, length: number }, icmp?:
{ type: number, code: number }, dns?: { id: number, qr: boolean, opcode: string,
rcode: string, questions: { name: string, type: string }[], answers: { name:
string, type: string, data: string }[] }, payload?: Uint8Array, bytes:
Uint8Array }) => void): Promise<{ iface: string, link: string, close():
Promise<void> }>

```

Live packet capture: `net.capture.open({ iface, promisc?, snaplen?, filter? }, pkt => {...})` → `Promise<{ iface, link, close() }>`. Linux + macOS only (Windows rejects); needs root / `CAP_NET_RAW` (Linux) or `/dev/bpf` access (macOS). `promisc` defaults true. The handler is called per frame with a decoded packet { `ts`, `length`, `captureLength`, `link`, `eth?`, `vlan?`, `arp?`, `ip?`, `tcp?`, `udp?`, `icmp?`, `dns?`, `payload?`, `bytes` }. Optional filter is a tcpdump-like expression string (e.g. 'tcp and port 80'), evaluated post-decode in userspace — NOT a kernel BPF program, so it skips the JS callback for non-matching packets but does not avoid the kernel→userspace copy. Supports tcp/udp/icmp/ip/ip6, host/src host/dst host, net/src net/dst net (CIDR), port/src port/dst port, portrange/src portrange/dst portrange, and/or/not + parens, implicit-and between juxtaposed primaries. A malformed expression makes `open` reject. `close()` returns `Promise<void>`. Pure-Go `gopacket` (no `libpcap/cgo`).

Parameters

- `opts` (*{ iface: string, promisc?: boolean, snaplen?: number, filter?: string }*) — `iface` is the interface name to capture on (required); `promisc` enables promiscuous mode (default true); `snaplen` caps the per-packet capture length in bytes (default 262144); `filter` is an optional tcpdump-like expression (e.g. 'tcp and port 80') applied post-decode in userspace — supports tcp/udp/icmp/ip/ip6, host/net (CIDR), port/portrange, src/dst directions, and/or/not + parens.
- `onPacket` (*((pkt: { ts: number, length: number, captureLength: number, link: string, eth?: { src: string, dst: string, type: string }, vlan?: { id: number, priority: number, drop: boolean, type: string }, arp?: { operation: string, senderMac: string, senderIp: string, targetMac: string, targetIp: string }, ip?: { version: number, src: string, dst: string, protocol: string, ttl: number }, tcp?: { srcPort: number, dstPort: number, seq: number, ack: number, window: number, checksum: number, flags: { syn: boolean, ack: boolean, fin: boolean, rst: boolean, psh: boolean, urg: boolean }, options?: { mss?: number, windowScale?: number, sackPermitted?: boolean, timestamps?: { val: number, ecr: number } } }, udp?: { srcPort: number, dstPort: number, length: number }, icmp?: { type: number, code: number }, dns?: { id: number, qr: boolean, opcode: string, rcode: string, questions: { name: string, type: string }[], answers: { name: string, type: string, data: string }[] }, payload?: Uint8Array, bytes: Uint8Array }) => void)*) — Called once per matching frame with the decoded packet. `ts` is epoch ms; layer keys (`eth/vlan/arp/ip/tcp/udp/icmp/dns`) are present only when that layer decoded; `bytes` is always the raw frame; `payload` is the application-layer bytes when present.

Returns: `Promise<{ iface: string, link: string, close(): Promise<void> }>` — a live-capture handle. `link` is the link-type name; `close()` stops the capture and resolves when the source is torn down. The handler keeps firing until `close()` is called or the source errors.

Throws: Rejects if iface is missing, the filter expression is malformed, the platform is unsupported (Windows), or the capture can't be opened (missing root / CAP_NET_RAW on Linux, /dev/bpf access on macOS). Throws synchronously if onPacket is not a function.

```
const cap = await net.capture.open({ iface: "en0", filter: "tcp and port 443" },
pkt => {
  runtime.log(pkt.ip?.src, "→", pkt.ip?.dst);
});
await cap.close();
```

net.capture.openFile

```
openFile(path: string, onPacket: (pkt: { ts: number, length: number,
captureLength: number, link: string, eth?: object, vlan?: object, arp?: object,
ip?: object, tcp?: object, udp?: object, icmp?: object, dns?: object, payload?:
Uint8Array, bytes: Uint8Array }) => void, opts?: { filter?: string }):
Promise<void>
```

Read a .pcap / .pcapng file: `net.capture.openFile(path, pkt => {...}, opts?) → Promise<void>`. Calls the handler once per decoded packet (same shape as `capture.open`) and resolves at EOF. Offline; no privileges. `opts` is an optional trailing arg `{ filter? }` — the 2-arg form still works; filter is the same tcpdump-like expression string as `capture.open` (post-decode/userspace, not kernel BPF; supports host/net CIDR + port/portrange; malformed → rejects).

Parameters

- `path` (*string*) — Path to a .pcap or .pcapng file; the format is auto-detected from the magic bytes.
- `onPacket` (*((pkt: { ts: number, length: number, captureLength: number, link: string, eth?: object, vlan?: object, arp?: object, ip?: object, tcp?: object, udp?: object, icmp?: object, dns?: object, payload?: Uint8Array, bytes: Uint8Array }) => void)*) — Called once per decoded packet (same shape as `capture.open`'s handler).
- `opts` (*{ filter?: string }, optional*) — filter is the same tcpdump-like expression as `capture.open`, applied post-decode in userspace; omit (2-arg form) to deliver every packet.

Returns: `Promise<void>` — resolves at end-of-file after dispatching every (matching) packet to the handler.

Throws: Rejects if the filter expression is malformed, the file can't be opened or parsed, or the handler throws. Throws synchronously if `onPacket` is not a function.

```
await net.capture.openFile("/tmp/dump.pcap", pkt => runtime.log(pkt.tcp?.dstPort),
{ filter: "tcp" });
```

net.capture.routes

```
routes(): { destination: string; gateway: string; interface: string; family: "ip"
| "ip6"; metric: number }[]
```

List the host's IP routing table synchronously: `net.capture.routes() → array of { destination, gateway, interface, family, metric }`. Pure-Go, unprivileged: Linux reads `/proc/net/route + /`

proc/net/ipv6_route; macOS/BSD read the routing socket via x/net/route. Windows is unsupported (throws).

Returns: { destination: string, gateway: string, interface: string, family: "ip" | "ip6", metric: number }[] — one entry per route. destination is a CIDR ('0.0.0.0/0' for the default route, '::/0' for the IPv6 default); gateway is the next-hop IP or '' for a directly-connected/link route; interface is the outgoing NIC name (best-effort); metric is 0 when the platform doesn't report one (BSD/macOS). Synchronous (not a Promise).

Throws: Throws if route enumeration fails or the platform is unsupported (Windows).

```
const def = net.capture.routes().find(r => r.destination === "0.0.0.0/0");
\nruntime.log("default via", def?.gateway, "on", def?.interface);
```

net.capture.toFile

```
toFile(path: string, opts?: { snaplen?: number, linkType?: number }):
{ write(bytes: string | Uint8Array, opts?: { ts?: number }): void; close():
Promise<void> }
```

Write raw frames to a .pcap file: net.capture.toFile(path, { linkType?, snaplen? }) → { write(bytes, { ts? }), close() }. write appends a raw frame (Uint8Array); ts (ms) overrides the timestamp. close() flushes and returns Promise<void>. Offline; no privileges.

Parameters

- path (*string*) — Path of the .pcap file to create (overwritten if it exists).
- opts ({ snaplen?: number, linkType?: number }, *optional*) — snaplen is the pcap global-header snap length (default 262144); linkType is the numeric pcap link-type written into the header (default Ethernet).

Returns: { write(bytes, opts?): void, close(): Promise<void> } — a writer handle (returned synchronously, not a Promise). write appends one raw frame to the file (opts.ts in ms overrides the timestamp, defaulting to now) and returns undefined. close() flushes and closes the file, resolving when done.

Throws: Throws synchronously if the file can't be created or the header can't be written, and on a per-frame write error. write throws if called after close. close() rejects on a close error.

```
const w = net.capture.toFile("/tmp/out.pcap");
w.write(frameBytes, { ts: Date.now() });
await w.close();
```

net.email.all

```
all(domain: string): Promise<{ domain: string, spf: object, dmarc: object, mtaSts:
object, tlsRpt: object, bimi: object }>
```

Run all five email probes in parallel — five-way handshake aggregate.

Parameters

- domain (*string*) — The domain to run all five email-auth probes against.

Returns: Promise<{ domain: string, spf: object, dmarc: object, mtaSts: object, tlsRpt: object, bimi: object }> — domain echoes the input; each probe key holds that probe's result (the same shape its individual binding returns) or { error: string } when that single probe failed. A per-probe failure doesn't fail the aggregate.

Throws: Resolves even when individual probes fail (their failure surfaces under <probe>.error). Always resolves.

```
const a = await net.email.all("example.com"); runtime.log(a.spf, a.dmarc);
```

net.email.bimi

```
bimi(domain: string, opts?: { selector?: string }): Promise<{ present: false, selector: string } | { present: true, selector: string, record: string, tags: Record<string, string>, l: string, a: string }>
```

Probe BIMI: TXT(<selector>._bimi.<domain>); selector defaults to 'default'.

Parameters

- `domain` (*string*) — The domain whose <selector>._bimi.<domain> TXT record is queried for a v=BIMI1 record.
- `opts` (*{ selector?: string }, optional*) — selector names the BIMI selector to query (default 'default').

Returns: Promise<{ present: false, selector: string } | { present: true, selector: string, record: string, tags: Record<string, string>, l: string, a: string }> — selector echoes the queried selector; when found, l is the logo URL tag and a is the assertion (VMC) tag. Missing record resolves to { present: false, selector }.

Throws: Rejects only on a DNS lookup error other than NXDOMAIN; an absent record resolves to { present: false, selector }.

```
const r = await net.email.bimi("example.com", { selector: "v1" }); runtime.log(r.present && r.l);
```

net.email.dmarc

```
dmarc(domain: string): Promise<{ present: false } | { present: true, record: string, tags: Record<string, string>, policy: string, subdomain: string, percent: string, rua: string, ruf: string }>
```

Query TXT(_dmarc.<domain>) and parse policy / pct / rua / ruf tags.

Parameters

- `domain` (*string*) — The domain whose _dmarc.<domain> TXT record is queried for a v=DMARC1 record.

Returns: Promise<{ present: false } | { present: true, record: string, tags: Record<string, string>, policy: string, subdomain: string, percent: string, rua: string, ruf: string }> — tags is the full parsed tag map; policy/subdomain/percent/rua/ruf surface the common p/sp/pct/rua/ruf tags. Missing record resolves to { present: false }.

Throws: Rejects only on a DNS lookup error other than NXDOMAIN; an absent record resolves to { present: false }.

```
const r = await net.email.dmarc("example.com"); runtime.log(r.present &&
r.policy);
```

net.email.mtaSts

```
mtaSts(domain: string): Promise<{ present: false } | { present: true, record:
string, txt: { v: string, id: string }, policy?: { version?: string, mode?:
string, mx?: string[], maxAge?: number | string }, policyError?: string }>
```

Probe MTA-STS: TXT(_mta-sts.<domain>) plus the fetched policy file.

Parameters

- `domain` (*string*) — The domain whose _mta-sts.<domain> TXT record and well-known policy file are probed.

Returns: Promise<{ present: false } | { present: true, record: string, txt: { v: string, id: string }, policy?: { version?: string, mode?: string, mx?: string[], maxAge?: number | string }, policyError?: string }> — txt carries the versioned id from the TXT marker; policy is the parsed well-known file (mode + mx + maxAge), or policyError holds the fetch/parse error string when the file couldn't be retrieved. Missing TXT resolves to { present: false }.

Throws: Rejects only on a DNS lookup error other than NXDOMAIN; an absent record resolves to { present: false }. A policy-file fetch failure is captured in policyError, not thrown.

```
const r = await net.email.mtaSts("example.com"); runtime.log(r.present &&
r.policy?.mode);
```

net.email.send

```
send(opts: { to: string | string[], from: string, subject?: string, body?: string,
html?: string, attachments?: { filename: string, contentType?: string, bytes:
Uint8Array | ArrayBuffer }[], headers?: Record<string, string>, server: { host:
string, port?: number, auth?: { username: string, password: string }, tls?:
"starttls" | "tls" | "none" }, timeout?: number }): Promise<{ accepted: string[],
rejected: { address: string, reason: string }[] }>
```

Send an outbound email: net.email.send({to, from, subject, body, html?, attachments?, headers?, server: {host, port?, auth?, tls?}, timeout?}) → Promise<{accepted: string[], rejected: [{address, reason}]}>. One TCP connection per call; per-recipient outcome captured. Transport failures throw; per-RCPT rejections surface in the result. TLS modes: starttls (default), tls, none.

Parameters

- `opts` (*{ to: string | string[], from: string, subject?: string, body?: string, html?: string, attachments?: { filename: string, contentType?: string, bytes: Uint8Array | ArrayBuffer }[], headers?: Record<string, string>, server: { host: string, port?: number, auth?: { username: string, password: string }, tls?: "starttls" | "tls" | "none" }, timeout?: number }*) — to (string or array) and from are required, as is server.host. subject/body/html shape the message:

body alone → text/plain, body+html → multipart/alternative, any attachments → multipart/mixed. attachments carry raw bytes (contentType defaults to application/octet-stream). headers adds custom headers (CR/LF stripped). server.port defaults to 587; server.auth enables PLAIN auth (skipped when tls is 'none'); server.tls picks the transport: 'starttls' (default), implicit 'tls', or 'none'. timeout is the dial / connection timeout in ms (default 30000).

Returns: Promise<{ accepted: string[], rejected: { address: string, reason: string }[] }> — accepted lists recipients the server accepted at RCPT TO; rejected pairs each refused address with the server's reason. The DATA body is sent only when at least one recipient was accepted.

Throws: Rejects on missing required fields (to / from / server.host), transport / protocol failures (dial, HELO, STARTTLS unavailable when requested, AUTH, MAIL FROM, DATA), or timeout. Per-recipient RCPT rejections are returned in rejected, not thrown.

```
const r = await net.email.send({
  to: "to@example.com", from: "from@example.com", subject: "hi", body: "hello",
  server: { host: "smtp.example.com", port: 587, auth: { username: "u", password:
    "p" } },
});
runtime.log(r.accepted, r.rejected);
```

net.email.spf

```
spf(domain: string): Promise<{ present: false } | { present: true, record: string,
mechanisms: string[], allPolicy: string }>
```

Query TXT(<domain>) for SPF, return record + parsed mechanisms + all-policy.

Parameters

- `domain` (*string*) — The domain whose apex TXT records are queried for an SPF (v=spf1) record.

Returns: Promise<{ present: false } | { present: true, record: string, mechanisms: string[], allPolicy: string }> — when found, record is the raw SPF string, mechanisms is the tokenised list after v=spf1, and allPolicy summarises the trailing all-style mechanism (pass / fail / softfail / neutral). Missing record resolves to { present: false }.

Throws: Rejects only on a DNS lookup error other than NXDOMAIN; an absent record resolves to { present: false }.

```
const r = await net.email.spf("example.com"); if (r.present)
runtime.log(r.allPolicy);
```

net.email.tlsRpt

```
tlsRpt(domain: string): Promise<{ present: false } | { present: true, record:
string, tags: Record<string, string>, rua: string }>
```

Probe TLS-RPT: TXT(_smtp._tls.<domain>) and parse rua.

Parameters

- `domain` (*string*) — The domain whose `_smtp._tls.<domain>` TXT record is queried for a `v=TLSRPTv1` record.

Returns: `Promise<{ present: false } | { present: true, record: string, tags: Record<string, string>, rua: string }>` — `tags` is the parsed tag map; `rua` surfaces the report-URI tag. Missing record resolves to `{ present: false }`.

Throws: Rejects only on a DNS lookup error other than `NXDOMAIN`; an absent record resolves to `{ present: false }`.

```
const r = await net.email.tlsRpt("example.com"); runtime.log(r.present && r.rua);
```

net.http.get

```
get(url: string): Promise<{ status: number, body: string, bodyBytes: Uint8Array }>
```

Perform an HTTP GET with a 5-second default timeout. Returns `{ status, body, bodyBytes }`.

Parameters

- `url` (*string*) — Absolute request URL (`http://` or `https://`).

Returns: `Promise<{ status: number, body: string, bodyBytes: Uint8Array }>` — the HTTP status code, the response body as a UTF-8 string, and `bodyBytes` (the raw, undecoded bytes; pair with `text.charset.decode` for non-UTF-8 content like ISO-8859-1). Redirects are followed by the default client.

Throws: Rejects on transport errors (DNS failure, connection refused, TLS handshake) or if the 5s context deadline is exceeded. 4xx/5xx responses do NOT reject — they surface via status.

```
const r = await net.http.get("https://example.com"); runtime.log(r.status);
```

net.http.post

```
post(url: string, body?: string): Promise<{ status: number, body: string, bodyBytes: Uint8Array }>
```

Perform an HTTP POST with a 5-second default timeout. Returns `{ status, body, bodyBytes }`.

Parameters

- `url` (*string*) — Absolute request URL (`http://` or `https://`).
- `body` (*string, optional*) — Request body sent verbatim; omit or pass empty for no body. No Content-Type header is set automatically.

Returns: `Promise<{ status: number, body: string, bodyBytes: Uint8Array }>` — the HTTP status code, the response body as a UTF-8 string, and `bodyBytes` (the raw, undecoded bytes; pair with `text.charset.decode` for non-UTF-8 content).

Throws: Rejects on transport errors (DNS failure, connection refused, TLS handshake) or if the 5s context deadline is exceeded. 4xx/5xx responses do NOT reject.

```
const r = await net.http.post("https://api.example.com/x", JSON.stringify({ a: 1 }));
```

net.http.request

```
request(method: string, url: string, opts?: { headers?: Record<string, string>,
body?: string, timeout?: number, retry?: number, follow?: boolean, username?:
string, password?: string }): Promise<{ status: number, ok: boolean, headers:
Record<string, string>, body: string, bodyBytes: Uint8Array, url: string }>
```

Full HTTP client: method, url, opts {headers, body, timeout, retry, follow, username, password}. Returns {status, ok, headers, body, bodyBytes, url}. 4xx/5xx don't throw; retry covers transport errors + 5xx.

Parameters

- `method` (*string*) — HTTP method (GET, POST, PUT, ...); upper-cased internally. Required.
- `url` (*string*) — Absolute request URL. Required.
- `opts` (*{ headers?: Record<string, string>, body?: string, timeout?: number, retry?: number, follow?: boolean, username?: string, password?: string }, optional*) — headers sets request headers; body is the raw request body; timeout is the per-attempt client timeout in ms (default 30000); retry is the number of extra attempts (default 0) applied only to transport errors and 5xx with linear backoff capped at 1s; follow toggles redirect following (default true — false stops at the first 3xx); username/password set HTTP Basic auth.

Returns: Promise<{ status: number, ok: boolean, headers: Record<string, string>, body: string, bodyBytes: Uint8Array, url: string }> — status is the final status code; ok is status in [200,400); headers is a lower-cased name → value map (last value wins, alphabetically ordered); body is the response text (UTF-8); bodyBytes is the raw, undecoded response bytes (pair with `text.charset.decode` for non-UTF-8 content like ISO-8859-1); url is the final URL after redirects.

Throws: Rejects on transport errors (DNS, connection refused, TLS) or context deadline, and after exhausting retries on a persistent transport error / 5xx. A malformed method or URL rejects immediately (not retried). 4xx/5xx that succeed at the transport level resolve normally.

```
const r = await net.http.request("POST", "https://api.example.com", { headers:
{ "content-type": "application/json" }, body: "{}", retry: 2 });
```

net.icmp.open

```
open(opts?: { network?: "ip4" | "ip6", readBuffer?: number }): Promise<{ network:
string, local: string, send(opts: { to: string, type?: number, code?: number, id?:
number, seq?: number, payload?: string | Uint8Array, body?: string |
Uint8Array }): Promise<void>, onMessage(cb: (ev: { bytes: Uint8Array, text:
string, address: string, type: number, code: number }) => void): void, onClose(cb:
() => void): void, onError(cb: (err: string) => void): void, close(): void }>
```

Open a raw ICMP socket: `net.icmp.open(opts?) → Promise<handle>`. Requires root / CAP_NET_RAW (open rejects otherwise). `opts` { network?: 'ip4'|'ip6' (default 'ip4'), readBuffer? }. `handle.send(opts)` writes a message in one of two modes: Echo mode { to,

type?, code?, id?, seq?, payload? } (type defaults to the network's echo request), or raw mode { to, type, code?, body } where body (Uint8Array|string) is marshalled verbatim (icmp.RawBody) for hand-built non-Echo messages such as destination-unreachable — in raw mode type is required and body is mutually exclusive with id/seq/payload. push/callback model — onMessage(cb) events carry { address, type, code }; onClose(cb)/onError(cb); handle.network/local; handle.close().

Parameters

- `opts` ({ `network?: "ip4" | "ip6"`, `readBuffer?: number` }, optional) — network selects the IP version (default 'ip4'); `readBuffer` is the inbound channel capacity (default 64).

Returns: Promise<{ network: string, local: string, send(opts: { to: string, type?: number, code?: number, id?: number, seq?: number, payload?: string | Uint8Array, body?: string | Uint8Array }): Promise<void>, onMessage(cb: (ev: { bytes: Uint8Array, text: string, address: string, type: number, code: number }) => void): void, onClose, onError, close(): void }> — a raw-ICMP handle. `send` writes an ICMP message: omit body for an Echo-shaped body (type defaults to the network's echo request, id/seq/payload optional); provide body for a verbatim raw body (type required, mutually exclusive with id/seq/payload) to hand-build non-Echo messages such as destination-unreachable. `to` is the destination address. `onMessage` fires per received packet with the marshalled body plus { address, type, code } meta.

Throws: Rejects if the raw socket can't be opened — typically because it needs root / CAP_NET_RAW. `send` rejects on resolve/marshal/write errors and throws synchronously: after close, if `opts.to` is missing, if a raw body is sent without `opts.type`, or if body is combined with id/seq/payload. Read errors surface via `onError`.

```
const p = await net.icmp.open();
p.onMessage(ev => runtime.log(ev.address, ev.type));
await p.send({ to: "8.8.8.8", id: 1, seq: 1, payload: "ping" });
// raw (non-Echo) body — e.g. a hand-built destination-unreachable:
await p.send({ to: "8.8.8.8", type: 3, code: 1, body: new Uint8Array([0, 0, 0, 0]) });
```

net.load.http

```
http(opts: { url: string, method?: string, headers?: Record<string, string>,
body?: string, concurrency?: number, requests?: number, duration?: number, rps?:
number, timeout?: number, confirm?: boolean }): Promise<{ target: string, method:
string, concurrency: number, durationMs: number, sent: number, completed: number,
failed: number, rps: number, errorRate: number, latency: { min: number, mean:
number, p50: number, p90: number, p95: number, p99: number, max: number },
statusCounts: Record<string, number>, errors: Record<string, number> }>
```

Authorized HTTP load / resilience self-test: drive a target with a worker pool at a given concurrency for a fixed `requests` count or `duration` (exactly one required), optional client-side `rps` cap, and return a latency/error report. Dual-use guardrail: public targets are refused unless `confirm:true` (loopback/private/localhost hosts are always allowed); concurrency is capped at 1000. Defensive self-testing only — no raw packets, spoofing, or amplification.

Parameters

- `opts` (*{ url: string, method?: string, headers?: Record<string, string>, body?: string, concurrency?: number, requests?: number, duration?: number, rps?: number, timeout?: number, confirm?: boolean }*) — `url` is the http/https target (required). `method` defaults to GET. `headers/body` set the request. `concurrency` is the number of parallel workers (default 10, clamped to [1,1000]). Provide exactly one of `requests` (total requests to send) or `duration` (run time in ms). `rps` caps client-side throughput (req/s, 0 = unlimited). `timeout` is the per-request timeout in ms (default 10000). `confirm:true` asserts you are authorized and is required to target a public host.

Returns: `Promise<LoadReport>` — `durationMs` is the wall-clock of the run; `sent` is requests started; `completed` got an HTTP response (any status); `failed` is transport errors/timeouts; `rps` is achieved throughput (completed/sec); `errorRate` is (failed + 5xx)/sent in [0,1]; `latency` is milliseconds over completed requests; `statusCounts` maps status code → count; `errors` maps transport error kind (timeout/refused/dns/reset/canceled/error) → count. A 4xx/5xx is a recorded response, not a failure.

Throws: Rejects if `url` is missing or not http/https, if neither or both of `requests` / `duration` are given, if `concurrency` exceeds the 1000 cap, or if the target host is public and `confirm:true` is not set.

```
const r = await net.load.http({ url: "http://127.0.0.1:8080/", requests: 200,
concurrency: 10 }); runtime.log(r.rps, r.latency.p95, r.errorRate);
```

net.netstatus.check

```
check(host: string, opts?: { port?: string, timeout?: number }): Promise<{ host:
string, port: string, elapsedMs: number, reachable: boolean, dns: { ok: boolean,
ips: string[], error?: string }, tcp: { ok: boolean, latencyMs: number, error?:
string }, tls: { ok: boolean, daysRemaining: number, error?: string }, http: { ok:
boolean, status: number, error?: string } }>
```

Run DNS / TCP / TLS / HTTP against one host concurrently. Returns { reachable, dns, tcp, tls, http } — each sub-probe ok+error; reachable = dns.ok AND tcp.ok. Sub-failures are data, not throws.

Parameters

- `host` (*string*) — The host to check. Required.
- `opts` (*{ port?: string, timeout?: number }*, optional) — `port` is the TCP/TLS port (default “443”); `timeout` bounds all four sub-probes in ms (default 10000).

Returns: `Promise<{ host: string, port: string, elapsedMs: number, reachable: boolean, dns: { ok: boolean, ips: string[], error?: string }, tcp: { ok: boolean, latencyMs: number, error?: string }, tls: { ok: boolean, daysRemaining: number, error?: string }, http: { ok: boolean, status: number, error?: string } }>` — the four sub-probe results plus elapsed time. reachable is dns.ok AND tcp.ok; TLS/HTTP are reported but don’t gate it. Each sub-probe carries its own error string instead of failing the call.

Throws: Rejects only if host is empty. Sub-probe failures are captured as data (ok:false + error) rather than thrown.

```
const s = await net.netstatus.check("example.com"); runtime.log(s.reachable);
```

net.probe.dns

```
dns(host: string, opts?: { types?: string[] }): Promise<{ a?: string[], aaaa?: string[], mx?: { preference: number, host: string }[], txt?: string[], cname?: string, ns?: string[] }>
```

Look up A / AAAA / MX / TXT / CNAME / NS records. Default: all five.

Parameters

- `host` (*string*) — The hostname to resolve.
- `opts` (*{ types?: string[] }, optional*) — `types` restricts the lookup to a subset (case-insensitive: 'a','aaaa','mx','txt','cname','ns'); omit to query all.

Returns: `Promise<{ a?: string[], aaaa?: string[], mx?: { preference: number, host: string }[], txt?: string[], cname?: string, ns?: string[] }>` — each key is present only when that record type returned at least one entry, so use `"mx"` in `result` to test presence.

Throws: Resolves with an object omitting record types that errored or were empty; per-type lookup failures are swallowed (not thrown). Always resolves.

```
const r = await net.probe.dns("example.com", { types: ["a", "mx"] });
```

net.probe.ntp

```
ntp(host: string, opts?: { timeout?: number, port?: number | string }): Promise<{ serverTime: string, offsetMs: number, rttMs: number, stratum: number, referenceTime: string, rootDelayMs: number, rootDispersionMs: number }>
```

Query an NTPv4 server (UDP 123) and report offset, RTT, stratum, root delay / dispersion.

Parameters

- `host` (*string*) — The NTP server hostname or IP.
- `opts` (*{ timeout?: number, port?: number | string }, optional*) — `timeout` is the query timeout in ms (default 5000); `port` overrides the default UDP port 123.

Returns: `Promise<{ serverTime: string, offsetMs: number, rttMs: number, stratum: number, referenceTime: string, rootDelayMs: number, rootDispersionMs: number }>` — the server's time and reference time (RFC3339 nanos), clock offset and round-trip in ms, the stratum, and the root delay / dispersion in ms.

Throws: Rejects if the NTP query fails (unreachable server, timeout, malformed response).

```
const r = await net.probe.ntp("pool.ntp.org"); runtime.log(r.offsetMs);
```

net.probe.ping

```
ping(host: string, opts?: { mode?: "tcp" | "icmp" | "udp", count?: number, timeout?: number, port?: string }): Promise<{ host: string, ip: string, mode:
```

```
string, sent: number, received: number, lossPercent: number, minMs: number, avgMs: number, maxMs: number }>
```

Reachability probe. mode tcp (default; dials host:port), icmp (real ICMP echo, needs raw-socket privileges), or udp (sends a datagram to a closed port and counts ICMP port-unreachable as reachable, needs root / CAP_NET_RAW). Returns { sent, received, lossPercent, minMs, avgMs, maxMs }. Unreachable = received 0, no throw.

Parameters

- `host` (*string*) — The target host. Required.
- `opts` (*{ mode?: "tcp" | "icmp" | "udp", count?: number, timeout?: number, port?: string, optional }*) — mode selects the probe (default 'tcp' — opens count TCP connections; 'icmp' sends real ICMP echo and needs raw-socket privileges; 'udp' sends a datagram to a closed port and counts the ICMP port-unreachable reply as reachable, needs root / CAP_NET_RAW); count is the number of probes (default 4); timeout is the per-probe timeout in ms (default 5000); port is the TCP target port (default "80", tcp mode only).

Returns: Promise<{ host: string, ip: string, mode: string, sent: number, received: number, lossPercent: number, minMs: number, avgMs: number, maxMs: number }> — the resolved IP, the mode used, packets sent/received, loss percentage, and min/avg/max RTT in ms. A fully unreachable host resolves with received:0 and lossPercent:100 rather than rejecting.

Throws: Rejects if host is empty, mode is not one of 'tcp', 'icmp', or 'udp', DNS resolution fails (tcp mode), or the raw ICMP socket can't be opened (icmp/udp modes; typically missing raw-socket privileges). Individual lost packets are counted, not thrown.

```
const p = await net.probe.ping("example.com", { count: 3 });
runtime.log(p.lossPercent);
```

net.probe.smtp

```
smtp(host: string, opts?: { port?: string, timeout?: number, ehloName?: string }):
Promise<{ host: string, port: string, banner: string, ehloDomain: string,
extensions: string[], starttls: boolean, authMechanisms: string[], sizeLimit:
number }>
```

SMTP capability probe (no mail sent). EHLO + parse extensions. Returns { banner, ehloDomain, extensions, starttls, authMechanisms, sizeLimit }. Connection failures throw.

Parameters

- `host` (*string*) — The SMTP server host. Required.
- `opts` (*{ port?: string, timeout?: number, ehloName?: string, optional }*) — port is the SMTP port (default "25"); timeout bounds the whole conversation in ms (default 10000); ehloName is the domain sent in EHLO (default "localhost").

Returns: Promise<{ host: string, port: string, banner: string, ehloDomain: string, extensions: string[], starttls: boolean, authMechanisms: string[], sizeLimit: number }> — the greeting banner, the server's EHLO greeting line, the raw advertised extension lines, whether STARTTLS is offered, the upper-cased AUTH mechanism names, and the SIZE limit (0 if unadvertised). No mail is sent.

Throws: Rejects if host is empty, the dial fails, or the greeting / EHLO cannot be read. A server that simply omits STARTTLS or AUTH reports them as false / empty — a finding, not an error.

```
const s = await net.probe.smtp("mail.example.com"); runtime.log(s.starttls, s.authMechanisms);
```

net.probe.tcp

```
tcp(target: string, opts?: { timeout?: number, port?: string }): Promise<{ host: string, port: number, ip: string, latencyMs: number }>
```

Dial a TCP target and report latency + resolved IP. Default timeout 5s.

Parameters

- `target` (*string*) — host:port to dial; a bare host uses `opts.port` (default 80).
- `opts` (*{ timeout?: number, port?: string }, optional*) — timeout is the dial timeout in ms (default 5000); port is the fallback port when target has no :port (default “80”).

Returns: `Promise<{ host: string, port: number, ip: string, latencyMs: number }>` — the parsed host, port, the resolved remote IP, and the connect latency in milliseconds.

Throws: Rejects if the dial fails (refused, unreachable, name resolution failure, or timeout).

```
const r = await net.probe.tcp("example.com:443"); runtime.log(r.latencyMs);
```

net.probe.tls

```
tls(target: string, opts?: { timeout?: number }): Promise<{ cn: string, issuer: string, notBefore: string, notAfter: string, daysRemaining: number, dnsNames: string[], serialNumber: string, fingerprintSha256: string }>
```

Open a TLS connection (`InsecureSkipVerify`; for probing only) and return the cert chain summary.

Parameters

- `target` (*string*) — host:port to dial; a bare host uses port 443. The host is sent as SNI.
- `opts` (*{ timeout?: number }, optional*) — timeout is the dial timeout in ms (default 5000).

Returns: `Promise<{ cn: string, issuer: string, notBefore: string, notAfter: string, daysRemaining: number, dnsNames: string[], serialNumber: string, fingerprintSha256: string }>` — leaf-certificate fields: common name, issuer CN, validity bounds (RFC3339), days until expiry, SAN DNS names, decimal serial, and the SHA-256 fingerprint as hex. Verification is skipped, so expired / mismatched certs still report.

Throws: Rejects if the dial / handshake fails, the connection is not TLS, or no peer certificates are presented.

```
const c = await net.probe.tls("example.com:443"); runtime.log(c.daysRemaining);
```

net.probe.traceroute

```
traceroute(host: string, opts?: { protocol?: "icmp" | "udp" | "tcp", port?: number, maxHops?: number, timeout?: number, probes?: number }): Promise<{ ttl: number; address: string | null; rttsMs: number[]; reached: boolean }[]>
```

Trace the network path to a host: `net.probe.traceroute(host, opts?) → Promise<hop[]>`. Sends probes with increasing TTL and reports each responding router. Needs root / CAP_NET_RAW (intermediate hops are seen via ICMP time-exceeded). `opts { protocol?: 'icmp'|'udp'|'tcp' (default 'icmp'), port?: number (udp 33434 / tcp 80), maxHops?: number (30), timeout?: number ms per probe (2000), probes?: number per hop (3) }`. IPv4 only.

Parameters

- `host` (*string*) — The destination host or IP.
- `opts` (*{ protocol?: "icmp" | "udp" | "tcp", port?: number, maxHops?: number, timeout?: number, probes?: number }, optional*) — protocol selects the probe type (icmp echo, udp to an incrementing high port, or tcp SYN via a TTL-limited connect). port is the udp/tcp target (ignored for icmp). maxHops caps the trace. timeout is the per-probe wait in ms. probes is the number of probes per hop.

Returns: One entry per hop (TTL 1..n): `ttl` is the hop number; `address` is the responding router/host IP (null if every probe at that TTL timed out); `rttsMs` are the round-trip times of the probes that answered; `reached` is true on the hop where the destination itself replied (the array ends there or at maxHops).

Throws: Rejects if the host doesn't resolve, the protocol is unknown, or the raw ICMP socket can't be opened (needs root / CAP_NET_RAW). Per-hop timeouts are normal (`address: null`), not errors.

```
// needs root / CAP_NET_RAW
const hops = await net.probe.traceroute("1.1.1.1", { protocol: "icmp", maxHops: 20 });
for (const h of hops) runtime.log(h.ttl, h.address ?? "*", h.rttsMs);
```

net.probe.whois

```
whois(domain: string, opts?: { timeout?: number }): Promise<{ raw: string, domain?: { name: string, punycode: string, whoisServer: string, nameServers: string[], status: string[], dnssec: boolean, createdAt: string, updatedAt: string, expirationDate: string }, registrar?: { name: string } }>
```

Two-hop WHOIS via the IANA referral, returning the parsed record plus the raw response text.

Parameters

- `domain` (*string*) — The domain (or IP / ASN) to look up.
- `opts` (*{ timeout?: number }, optional*) — timeout is the wire-level WHOIS client timeout in ms (default 10000).

Returns: `Promise<{ raw: string, domain?: { name: string, punycode: string, whoisServer: string, nameServers: string[], status: string[], dnssec: boolean, createdAt: string, updatedAt: string, expirationDate: string }, registrar?: { name: string } }>` — raw is always

the full WHOIS text; domain and registrar are best-effort parsed fields, omitted for TLDs the parser doesn't recognise.

Throws: Rejects if the WHOIS query itself fails (no referral, connection error, timeout). A parse failure is non-fatal — only raw is returned.

```
const w = await net.probe.whois("example.com");
runtime.log(w.domain?.expirationDate);
```

net.probe.wss

```
wss(url: string, opts?: { timeout?: number, ping?: boolean }): Promise<{ url:
string, connected: boolean, subprotocol: string, status: number, handshakeMs:
number, pingMs: number }>
```

WebSocket handshake probe. Opens ws:

Parameters

- `url` (*string*) — The WebSocket URL (ws:
- `opts` (*{ timeout?: number, ping?: boolean }, optional*) — timeout bounds the handshake and ping in ms (default 10000); ping toggles the ping/pong RTT measurement (default true).

Returns: `Promise<{ url: string, connected: boolean, subprotocol: string, status: number, handshakeMs: number, pingMs: number }>` — connected is true on a successful upgrade, subprotocol is the negotiated subprotocol (or empty), status is the HTTP status of the 101 upgrade, handshakeMs is the handshake time in ms, and pingMs is the ping/pong RTT (or -1 when the ping was skipped or unanswered). The connection is closed immediately.

Throws: Rejects if url is empty or the handshake fails (non-101, refused, bad URL). A failed ping leaves pingMs at -1 rather than rejecting.

```
const w = await net.probe.wss("wss://echo.websocket.org");
runtime.log(w.handshakeMs);
```

net.raw.open

```
open(opts?: { iface?: string, filter?: string, readBuffer?: number }):
Promise<{ link: string; send(spec: object | Uint8Array): Promise<{ bytesSent:
number }>; onPacket(cb: (pkt: any) => void): void; onClose(cb: () => void): void;
onError(cb: (msg: string) => void): void; close(): Promise<void> }>
```

Open a raw IPv4 packet engine: `net.raw.open({ iface?, filter?, readBuffer? })` → `Promise<handle>`. Sends crafted IPv4 packets (TCP flags / UDP / arbitrary IP protocol) via an `IP_HDRINCL` raw socket and receives replies via the capture path. Needs root / `CAP_NET_RAW`; Linux + macOS only (Windows rejects). `iface` defaults to the auto-detected default-route interface; `filter` is a tcpdump-like expression narrowing `onPacket`. The handle: `send(specOrBytes)` → `Promise<{ bytesSent }>`; `onPacket(cb)` delivers a decoded packet (same shape as `net.capture`); `onClose/onError`; `close()` → `Promise<void>`. send spec: { dst, dstPort?, srcPort?, src?, proto?: 'tcp'|'udp'|'ip', protocol?, flags?: string[], seq?, ack?, window?,

ttl?, ipId?, payload? }; or pass a Uint8Array to send a full IPv4 packet verbatim. Default flags ['SYN'], ttl 64, window 65535, src = egress IP, srcPort = random high.

Parameters

- `opts` (*{ iface?: string, filter?: string, readBuffer?: number }, optional*) — `iface` is the capture/egress interface (auto-detected if omitted); `filter` is a tcpdump-like expression evaluated post-decode; `readBuffer` sizes the inbound channel (default 64).

Returns: A handle: `link` is the capture link type; `send` crafts+fires a packet (structured spec or raw bytes) and resolves { `bytesSent` }; `onPacket` receives decoded reply packets; `close()` tears down the send socket and capture.

Throws: Rejects if the platform is Windows, the raw socket or capture can't be opened (needs root / CAP_NET_RAW), the egress interface can't be detected, or the filter is malformed. `send` throws on an invalid spec (missing dst, unknown TCP flag, bad port/ttl) and rejects on a write failure.

```
// needs root / CAP_NET_RAW
const h = await net.raw.open({ filter: "tcp and src port 443" });
h.onPacket(p => runtime.log(p.ip.src, p.tcp.flags));
await h.send({ dst: "93.184.216.34", dstPort: 443, flags: ["SYN"] });
await new Promise(r => setTimeout(r, 1000));
await h.close();
```

net.raw.tcp

```
tcp(host: string, port: number, opts?: { flags?: string[], srcPort?: number, src?: string, seq?: number, ttl?: number, payload?: Uint8Array | string, timeout?: number, iface?: string }): Promise<{ ts: number; link: string; ip: { src: string; dst: string; protocol: string; ttl: number }; tcp: { srcPort: number; dstPort: number; seq: number; ack: number; flags: { syn: boolean; ack: boolean; fin: boolean; rst: boolean; psh: boolean; urg: boolean }; payload?: Uint8Array; bytes: Uint8Array } | null>
```

One-shot raw TCP probe: `net.raw.tcp(host, port, opts?)` → `Promise<reply | null>`. Sends a single crafted TCP segment (default a SYN) and resolves with the first reply packet correlated by the 4-tuple, or null on timeout. SYN → SYN/ACK means open; RST means closed; null means filtered/no answer. Needs root / CAP_NET_RAW; Linux + macOS only.

Parameters

- `host` (*string*) — Destination host or IPv4 address. Required.
- `port` (*number*) — Destination TCP port. Required.
- `opts` (*{ flags?: string[], srcPort?: number, src?: string, seq?: number, ttl?: number, payload?: Uint8Array | string, timeout?: number, iface?: string }, optional*) — `flags` are the TCP flags to set (default ['SYN']); `src/srcPort/seq/ttl/payload` tune the crafted segment; `timeout` is the reply wait in ms (default 2000); `iface` overrides the auto-detected capture interface.

Returns: The decoded reply packet (same shape as `net.capture` packets), or null if no correlated reply arrived within the timeout. A SYN/ACK indicates the port is open; an RST indicates closed.

Throws: Rejects if the platform is Windows, host is empty, port is out of range, DNS resolution fails, or the raw socket / capture can't be opened (needs root / CAP_NET_RAW). A timeout is not an error — it resolves null.

```
// needs root / CAP_NET_RAW
const reply = await net.raw.tcp("scanme.nmap.org", 80, { flags: ["SYN"] });
if (reply && reply.tcp.flags.syn && reply.tcp.flags.ack) runtime.log("open");
else if (reply && reply.tcp.flags.rst) runtime.log("closed");
else runtime.log("filtered / no answer");
```

net.tcp.connect

```
connect(host: string, port: string | number, opts?: { timeout?: number,
readBuffer?: number }): Promise<{ remote: string, local: string, write(data:
string | Uint8Array): Promise<void>, onData(cb: (ev: { bytes: Uint8Array, text:
string }) => void): void, onClose(cb: () => void): void, onError(cb: (err: string)
=> void): void, close(): void }>
```

Open a TCP client socket: `net.tcp.connect(host, port, opts?) → Promise<handle>`. Push/callback read model — `handle.onData(cb)/onClose(cb)/onError(cb)` register listeners; `handle.write(data)` sends (string→UTF-8 / Uint8Array); `handle.remote/local` are the peer/local addresses; `handle.close()` shuts down. `opts { timeout?, readBuffer? }`.

Parameters

- `host` (*string*) — The remote host to dial.
- `port` (*string* / *number*) — The remote port.
- `opts` (*{ timeout?: number, readBuffer?: number }, optional*) — `timeout` is the dial timeout in ms (default 10000); `readBuffer` is the inbound channel capacity (default 64).

Returns: `Promise<{ remote: string, local: string, write(data: string | Uint8Array): Promise<void>, onData(cb: (ev: { bytes: Uint8Array, text: string }) => void): void, onClose(cb: () => void): void, onError(cb: (err: string) => void): void, close(): void }>` — a connected-socket handle. `remote/local` are the peer/local addresses. `write` resolves once the bytes are written. `onData` fires per inbound chunk with both a `Uint8Array` and a UTF-8 text view; `onClose` fires when the stream ends; `onError` forwards non-EOF read/transport errors. `close()` tears down the connection.

Throws: The returned Promise rejects if the dial fails (refused, unreachable, timeout). After connect, write rejects on a write error and throws synchronously if called after close; transport read errors surface via the `onError` callback, not as rejections.

```
const sock = await net.tcp.connect("example.com", 80);
sock.onData(ev => runtime.log(ev.text));
await sock.write("GET / HTTP/1.0\r\n\r\n");
```

net.udp.open

```
open(opts: { host?: string, port?: string | number, bind?: string, readBuffer?:
number }): Promise<{ local: string, send(data: string | Uint8Array):
Promise<void>, sendTo(data: string | Uint8Array, host: string, port: string |
number): Promise<void>, onMessage(cb: (ev: { bytes: Uint8Array, text: string,
```



```
address?: string, port?: number }) => void): void, onClose(cb: () => void): void,
onError(cb: (err: string) => void): void, close(): void }>
```

Open a UDP socket: `net.udp.open(opts) → Promise<handle>`. Connected mode `{ host, port }` exposes `send(data)`; bound mode `{ bind: '9999' }` exposes `sendTo(data, host, port)` and tags inbound events with `{ address, port }`. Push/callback model — `onMessage(cb)/onClose(cb)/onError(cb)`; `handle.local` is the bound address; `handle.close()` shuts down. `opts` also takes `readBuffer?`.

Parameters

- `opts` (`{ host?: string, port?: string | number, bind?: string, readBuffer?: number }`) — Selects the mode: connected mode needs `{ host, port }` (`net.DialUDP` to that peer); bound mode needs `{ bind }` (e.g. `'9999'`, `net.ListenUDP` on that local address). `readBuffer` is the inbound channel capacity (default 64). Provide exactly one of the two modes.

Returns: `Promise<handle>` — connected mode resolves to `{ local: string, send(data: string | Uint8Array): Promise<void>, onMessage, onClose, onError, close(): void }`; bound mode resolves to `{ local: string, sendTo(data: string | Uint8Array, host: string, port: string | number): Promise<void>, send (throws), onMessage, onClose, onError, close(): void }`. `onMessage` fires per datagram with `{ bytes: Uint8Array, text: string }` plus `{ address, port }` in bound mode. `local` is the bound address.

Throws: Rejects if neither `{ bind }` nor `{ host, port }` is supplied, or the dial / listen fails. `send`/`sendTo` reject on a write error and throw synchronously after `close`; calling `send` on a bound socket throws (use `sendTo`). Read errors surface via `onError` (a clean close ends silently).

```
const u = await net.udp.open({ host: "1.1.1.1", port: 53 });
u.onMessage(ev => runtime.log(ev.bytes.length));
await u.send(query);
```

runtime

Script-host scaffolding: logging, assertions, time, environment, `runtime.argv`.

runtime.argv

```
argv: string[]
```

Per-script argument vector: `[programName, scriptPath, ...userArgs]`. `argv[0]` is the program name (sercon), `argv[1]` is the running script path, and any args after `--` on the command line start at `argv[2]`.

Returns: `string[]` — the per-run argument vector. `argv[0]` is the program name (“sercon”), `argv[1]` is the running script’s path, and entries from index 2 onward are the user arguments passed after `--` on the command line. This is a value (property), not a function.

Throws: Not callable — accessing it never throws; reading an out-of-range index yields undefined per normal array semantics.

```
const target = runtime.argv[2] ?? "default-host";
```

runtime.assert.equal

```
equal(actual: unknown, expected: unknown, msg?: string): void
```

Throw when actual !== expected (strict equality on primitives, deep equality on objects). Optional msg appears in the error.

Parameters

- `actual` (*unknown*) — The value produced by the code under test.
- `expected` (*unknown*) — The value to compare against. Primitives use strict equality; objects/arrays use deep structural equality (key order ignored).
- `msg` (*string, optional*) — Optional message prefixed onto the thrown error; defaults to “assert.equal failed”.

Returns: void — returns nothing when the values match.

Throws: Throws an Error (“<msg>: expected <expected>, got <actual>”) when the values are not equal.

```
runtime.assert.equal(1 + 1, 2, "math still works");
```

runtime.assert.ok

```
ok(cond: unknown, msg?: string): void
```

Throw when cond is falsy. Optional msg appears in the error.

Parameters

- `cond` (*unknown*) — A value tested for truthiness (JS coercion: 0, “”, null, undefined, NaN and false are falsy).
- `msg` (*string, optional*) — Optional message used as the thrown error text; defaults to “assert.ok failed”.

Returns: void — returns nothing when cond is truthy.

Throws: Throws an Error carrying msg (or “assert.ok failed”) when cond is null or coerces to false.

```
runtime.assert.ok(user.id, "user must have an id");
```

runtime.clearDeadline

```
clearDeadline(): void
```

Remove the running script’s wall-clock kill deadline entirely (equivalent to -timeout 0 / setDeadline(0)). The script then runs without a timeout.

Returns: void.

Throws: Does not throw.

```
runtime.clearDeadline(); // run without a timeout
```

runtime.clipboard.available

available: boolean

True when a host clipboard backend is on PATH (macOS pbcopy/pbpaste; Linux wl-clipboard or xclip/xsel; Windows clip + PowerShell). Cheap, side-effect-free advisory — does not touch the clipboard; gate calls on it to self-skip on headless boxes.

Returns: boolean — false when no clipboard CLI is installed (e.g. a headless server). The authoritative signal is whether read/write throw.

Throws: Never throws.

```
if (!runtime.clipboard.available) runtime.log("no clipboard – skipping");
```

runtime.clipboard.imageAvailable

imageAvailable: boolean

True when a PNG image clipboard backend is usable on this host (macOS pngpaste; Linux wl-clipboard or xclip — not xsel; Windows PowerShell). Cheap advisory; does not touch the clipboard.

Returns: boolean — false when no image backend is installed (e.g. macOS without pngpaste). Gate readImage/writeImage on it.

Throws: Never throws.

```
if (runtime.clipboard.imageAvailable) { /* ... */ }
```

runtime.clipboard.read

read(...args: unknown[]): Promise<string>

Read the host OS system clipboard as UTF-8 text. Async (shells out to the platform clipboard tool). An empty clipboard resolves with "".

Returns: Promise resolving to the clipboard text ("" when the clipboard is empty or holds no text).

Throws: Rejects when no clipboard backend is on PATH (a clean “runtime.clipboard: no clipboard backend ...” message) or the underlying command fails / times out (5s).

```
const text = await runtime.clipboard.read();
```

runtime.clipboard.readImage

readImage(...args: unknown[]): Promise<Uint8Array | null>

Read the host clipboard image as PNG bytes. Async (shells out). Resolves null when the clipboard holds no image.

Returns: Promise resolving to the clipboard image as PNG bytes, or null when no image is present.

Throws: Rejects when no image backend is available (see `imageAvailable`) or the underlying command fails / times out (5s).

```
const png = await runtime.clipboard.readImage();
```

runtime.clipboard.write

```
write(text: string): Promise<void>
```

Replace the host OS system clipboard with the given text. Async (shells out to the platform clipboard tool); text is passed via stdin (no shell-injection risk).

Parameters

- `text` (*string*) — The text to place on the clipboard (non-string values are `String()`-coerced).

Returns: Promise resolving when the clipboard has been set.

Throws: Rejects when no clipboard backend is on PATH or the underlying command fails / times out (5s).

```
await runtime.clipboard.write("copied from sercon");
```

runtime.clipboard.writeImage

```
writeImage(png: Uint8Array): Promise<void>
```

Set the host clipboard image from PNG bytes. Async (shells out). The input must be a PNG (validated by signature) — other formats reject.

Parameters

- `png` (*Uint8Array*) — PNG image bytes (must begin with the PNG signature).

Returns: Promise resolving when the clipboard image is set.

Throws: Rejects when the data is not a PNG, no image backend is available, or the command fails / times out (5s).

```
await runtime.clipboard.writeImage(pngBytes);
```

runtime.env.get

```
get(name: string): string | undefined
```

Read an environment variable. Returns undefined when unset (not empty string).

Parameters

- `name` (*string*) — Environment variable name to look up.

Returns: string — the variable's value, or undefined when the variable is not set. A variable set to the empty string returns "", which is distinct from undefined.

Throws: Never throws.

```
const home = runtime.env.get("HOME") ?? "/tmp";
```

runtime.getDeadline

```
getDeadline(): number | null
```

Return the milliseconds remaining until the running script's kill deadline, or null when no deadline is active (disabled or started with -timeout 0).

Returns: number — ms remaining (≥ 0) — or null when there is no active deadline.

Throws: Does not throw.

```
const left = runtime.getDeadline(); // e.g. 9871, or null
```

runtime.log

```
log(args: unknown[]): void
```

Print one space-separated line of the arguments to stdout. Primitives print raw; objects/arrays render as JSON (circular refs fall back to [object Object]). The script-side equivalent of console.log.

Parameters

- `args` (*unknown[]*) — Zero or more values to print. They are joined with single spaces; primitives stringify directly, objects/arrays are JSON-encoded.

Returns: void — writes a single newline-terminated line to stdout.

Throws: Never throws; unserialisable objects degrade to [object Object] rather than erroring.

```
runtime.log("count", 3, { ok: true }); // count 3 {"ok":true}
```

runtime.secrets.available

```
available: boolean
```

True when an OS keystore backend (macOS Keychain, Linux Secret Service, Windows Credential Manager) is plausibly reachable this run. Cheap advisory hint — does not touch the keystore; gate calls on it to self-skip on headless boxes.

Returns: boolean — false on a host with no reachable keystore (e.g. a headless Linux box without a D-Bus session). The authoritative signal is whether get/set/delete throw.

Throws: Never throws.

```
if (!runtime.secrets.available) runtime.log("no keystore - skipping");
```

runtime.secrets.delete

```
delete(name: string, account: string): Promise<boolean>
```

Remove a secret from the OS keystore under prefix + name / account. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the sercon prefix namespace (keystore service is prefix+name).
- `account` (*string*) — Account/user the secret belongs to.

Returns: Promise resolving true when an item was removed, false when there was nothing to remove.

Throws: Rejects if name is missing/empty or account is missing (pass "" for a single-secret name), when the keystore backend is unreachable, or the delete fails ("runtime.secrets.delete: ..."). Bounded by a 10s timeout.

```
const removed = await runtime.secrets.delete("devshop", "tess@example.com");
```

runtime.secrets.get

```
get(name: string, account: string): Promise<string | null>
```

Read a string secret from the OS keystore. The keystore service is the configured prefix + name (default "sercon/"), so reads are confined to sercon's namespace. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the sercon prefix namespace (the keystore service is prefix+name, e.g. "sercon/devshop").
- `account` (*string*) — Account/user the secret belongs to (may be an empty string for a single-secret name).

Returns: Promise resolving to the stored secret string, or null when no such item exists.

Throws: Rejects if name is missing/empty or account is missing (pass "" for a single-secret name), when the keystore backend is unreachable, or the read fails (a clean "runtime.secrets.get: ..." error). Bounded by a 10s timeout (a blocking macOS consent prompt rejects rather than hangs).

```
const pw = await runtime.secrets.get("devshop", "tess@example.com");
```

runtime.secrets.set

```
set(name: string, account: string, secret: string): Promise<void>
```

Store or overwrite a string secret in the OS keystore under prefix + name / account. Async (keystore I/O).

Parameters

- `name` (*string*) — Secret name within the `sercon` prefix namespace (keystore service is `prefix+name`).
- `account` (*string*) — Account/user the secret belongs to.
- `secret` (*string*) — The secret value to store.

Returns: Promise resolving when the secret is written.

Throws: Rejects if `name` is missing/empty, `account` is missing (pass `""` for a single-secret name), or `secret` is missing; when the keystore backend is unreachable; or the write fails (`"runtime.secrets.set: ..."`). Bounded by a 10s timeout.

```
await runtime.secrets.set("devshop", "tess@example.com", "hunter2");
```

runtime.setDeadline

```
setDeadline(ms: number): void
```

Set the running script's wall-clock kill deadline to `now + ms` (replacing any prior deadline; `ms<=0` disables it). This is the same deadline the `-timeout` flag sets — NOT the JS global `setTimeout` (which schedules a callback). Use it to extend a long task or add a deadline to a `-timeout 0` run.

Parameters

- `ms` (*number*) — Milliseconds from now until the run is killed; `<= 0` disables the deadline (same as `clearDeadline()`).

Returns: `void` — applies immediately to the in-flight run.

Throws: Does not throw; a non-numeric argument coerces to 0 (disable).

```
runtime.setDeadline(30000); // give this run 30s from now
```

runtime.time.format

```
format(ms: number, layout: string, tz?: string): string
```

Format a unix-ms timestamp through `strftime` tokens. Optional IANA `tz` (e.g. `'Europe/Stockholm'`); default is the host's local zone.

Parameters

- `ms` (*number*) — Milliseconds since the Unix epoch (e.g. from `time.nowMs`). Coerced to an integer.
- `layout` (*string*) — `strftime`-style layout. Supported tokens: `%Y %y %m %d %H %M %S %T %F %j %A %a %B %b %z %Z` and `%%` (literal percent). Unknown `%X` tokens pass through verbatim.
- `tz` (*string, optional*) — IANA timezone name (e.g. `"Europe/Stockholm"`, `"UTC"`). Defaults to the host's local zone when omitted/null/undefined.

Returns: `string` — the timestamp rendered in `tz` with the given layout.

Throws: Throws (`"time.format: ..."`) if `tz` is not a loadable IANA timezone name.

```
const s = runtime.time.format(runtime.time.nowMs(), "%F %T", "UTC");
```

runtime.time.nowMs

```
nowMs(): number
```

Wall-clock milliseconds since the Unix epoch.

Returns: number — integer milliseconds since 1970-01-01T00:00:00Z (host wall clock).

Throws: Never throws.

```
const t0 = runtime.time.nowMs();
```

runtime.time.sleep

```
sleep(ms: number): Promise<void>
```

Resolve after `ms` milliseconds. Cancellable via the engine timeout.

Parameters

- `ms` (*number*) — Delay in milliseconds. Coerced to an integer; non-positive values resolve effectively immediately.

Returns: `Promise<void>` — resolves once the delay elapses.

Throws: Rejects if the run is cancelled or hits its timeout before the delay elapses (the underlying context is cancelled).

```
await runtime.time.sleep(250);
```

server

Network servers: HTTP/HTTPS listeners with routing, middleware, static files, WebSocket upgrade.

server.http.listen

```
listen(opts: { port: number; host?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; handler: (req: Request, res: Response) => unknown } >; use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; onError?: (err: unknown, req: Request, res: Response) => unknown }): { address: string; stopped: Promise<void>; close(): Promise<void> }
```

Bind an HTTP listener: `server.http.listen({port, host?, routes, use?, onError?})` → handle with `.address`, `.close()`, `.stopped` Promise. `routes` is a map of stdlib `http.ServeMux` patterns ('GET /users/{id}') to handlers (`(req, res) => res.json({...})` or `{use: [...], handler: fn}` for per-route middleware. Optional `onError(err, req, res)` renders a custom response when a handler/middleware throws or rejects (else a stock 500). Handlers can call `res.upgradeWebSocket(opts?)` to hijack the connection and return an `AsyncIterable<WSMessage>` with `.send` / `.close` — `for await (const msg of ws)` walks frames; `msg` is `{type:'text',text}` or `{type:'binary',bytes:Uint8Array}`. Handlers can also call

`res.sse(opts?)` to start a one-way Server-Sent Events stream — returns a handle with `send(data)` (`string` → `data`: or `{event,data,id,retry}` with object data JSON-encoded), `close()`, and a `closed` Promise (resolves on close or client disconnect); `opts`: `{keepAlive?: ms, retry?: ms}`.

Parameters

- `opts` (`{ port: number; host?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; onError?: (err: unknown, req: Request, res: Response) => unknown }`) — Listener config. `port` is required; `host` defaults to “0.0.0.0”. `routes` maps Go 1.22+ ServeMux patterns (‘GET /’, ‘POST /users/{id}’, ‘GET /assets/{rest...}’) to a handler function or a `{use, handler}` object for per-route middleware. `use` is a global middleware chain run before every route. `onError(err, req, res)` is invoked when a handler or middleware throws/rejects — render a response with `res.*` (it may be async); if it doesn’t finalize, or itself throws, the stock 500 is sent. Under `sercon serve`, `-port-` override replaces `port`.

Returns: A server handle (returned synchronously): address is ‘tcp/host:port’ (resolved, so a `port:0` ephemeral bind reports its OS-chosen port); `stopped` resolves when the server stops (rejects if Serve fails with a non-close error); `close()` begins a graceful 30s shutdown and resolves with the same `stopped` Promise.

Throws: Throws synchronously if `opts` is missing, `port` is 0/absent, `routes` is missing, a `use[]` entry or route value is not a function/valid `{use, handler}`, or the bind fails (e.g. address already in use).

```
const srv = server.http.listen({
  port: 8080,
  routes: { "GET /": (req, res) => res.json({ ok: true }) },
});
runtime.log(srv.address);
await srv.close();
```

server.http.static

```
static(opts: { dir: string; stripPrefix?: string; index?: string; etag?:
boolean }): (req: Request, res: Response) => void
```

Static-file mount: `server.http.static({dir, stripPrefix, index?, etag?})` → handler. Assign to a wildcard route (GET /assets/{rest...}). Internally `stdlib http.FileServer` with `stripPrefix`; ETag/Last-Modified/range requests work; no directory listing.

Parameters

- `opts` (`{ dir: string; stripPrefix?: string; index?: string; etag?: boolean }`) — `dir` is the filesystem root to serve. `stripPrefix` is removed from the request path before lookup (set it to the route’s static prefix). `index` and `etag` are accepted but currently unused — `http.FileServer` already serves `index.html` and emits ETag/Last-Modified by default.

Returns: A route handler marker (returned synchronously). Assign it as a routes entry, typically under a wildcard pattern like 'GET /assets/{rest...}'. The route compiler unwraps it to a stdlib http.FileServer mounted under http.StripPrefix.

Throws: The call itself does not throw; an invalid dir surfaces as 404s at request time.

```
server.http.listen({
  port: 8080,
  routes: { "GET /assets/{rest...}": server.http.static({ dir: "./public",
stripPrefix: "/assets/" }) },
});
```

server.https.listen

```
listen(opts: { port: number; host?: string; cert: string | "self-signed"; key?:
string; routes: Record<string, ((req: Request, res: Response) => unknown) |
{ use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[];
handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res:
Response, next: () => Promise<void>) => unknown)[] }): { address: string; stopped:
Promise<void>; close(): Promise<void> }
```

Like server.http.listen plus a cert (file path OR inline PEM string) and matching key. Set cert: "self-signed" to mint an ephemeral in-process dev cert (no key needed) instead — covers localhost / 127.0.0.1 / ::1 plus the listen host. No autocert (own production certs in your supervisor).

Parameters

- *opts* (*{ port: number; host?: string; cert: string | "self-signed"; key?: string; routes: Record<string, ((req: Request, res: Response) => unknown) | { use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[]; handler: (req: Request, res: Response) => unknown }>; use?: ((req: Request, res: Response, next: () => Promise<void>) => unknown)[] }*) — Same shape as server.http.listen plus cert and key. cert is a filesystem path, an inline PEM string (detected by a leading '—BEGIN'), or the literal "self-signed" to generate an ephemeral P-256 cert in-process (key is then ignored). key is required for the path/PEM forms. TLS is pinned to a minimum of TLS 1.2. Self-signed certs fail normal client verification by design — dev only.

Returns: Same handle shape as server.http.listen; address is 'tcp/host:port'.

Throws: Throws synchronously on the same conditions as server.http.listen, plus if cert is missing, key is missing for a path/PEM cert, or the key pair fails to load/parse.

```
// Ephemeral dev cert — no openssl, no files:
const srv = server.https.listen({
  port: 8443,
  cert: "self-signed",
  routes: { "GET /": (req, res) => res.text("secure") },
});
```

server.https.static

```
static(opts: { dir: string; stripPrefix?: string; index?: string; etag?: boolean }): (req: Request, res: Response) => void
```

Like `server.http.static`; same options.

Parameters

- `opts` (*{ dir: string; stripPrefix?: string; index?: string; etag?: boolean }*) — Identical to `server.http.static` — `dir` is the root, `stripPrefix` is removed from the path before lookup, `index/etag` are accepted but unused.

Returns: A route handler marker (returned synchronously); assign it to a wildcard route on an https listener.

Throws: The call itself does not throw; an invalid `dir` surfaces as 404s at request time.

```
server.https.listen({
  port: 8443, cert, key,
  routes: { "GET /assets/{rest...}": server.https.static({ dir: "./public",
    stripPrefix: "/assets/" }) },
});
```

`server.icmp.listen`

```
listen(opts?: { network?: "ip4" | "ip6" }, handler: (msg: { bytes: Uint8Array;
  text: string; address: string; type: number; code: number }, reply: (opts?: { to?:
  string; type?: number; code?: number; id?: number; seq?: number; payload?: string
  | Uint8Array; body?: string | Uint8Array }) => Promise<void>) => void): { address:
  string; close(): Promise<void> }
```

Bind a raw ICMP listener: `server.icmp.listen(opts?, (msg, reply) => {...})` → handle { address: 'icmp/<addr>', close() }. Raw ICMP has no ports — the socket receives ALL host ICMP traffic — and needs `root / CAP_NET_RAW` (synchronous bind throws otherwise). `opts` { `network?: 'ip4'|'ip6'` (default 'ip4') }. The handler runs once per received packet; `msg` is { `bytes`, `text`, `address`, `type`, `code` } (the sender + parsed ICMP header) and `reply(opts?)` sends an ICMP message back to the sender (Echo by default, or a raw body), returning a Promise. Emits a READY line under `sercon serve` and joins graceful shutdown.

Parameters

- `opts` (*{ network?: "ip4" | "ip6" }, optional*) — network selects the IP version (default 'ip4'). There is no host/port — raw ICMP binds to all addresses.
- `handler` (*((msg: { bytes: Uint8Array; text: string; address: string; type: number; code: number }, reply: (opts?: { to?: string; type?: number; code?: number; id?: number; seq?: number; payload?: string | Uint8Array; body?: string | Uint8Array }) => Promise<void>) => void)*) — Invoked once per received ICMP packet. `msg` carries the marshalled body (bytes/text), the sender address, and the parsed type/code. `reply(opts?)` sends an ICMP message back to the sender (or `opts.to`): Echo mode { `type?`, `code?`, `id?`, `seq?`, `payload?` } or raw mode { `type`, `code?`, `body` } (body marshalled verbatim); it returns a Promise that resolves once written.

Returns: A server handle (returned synchronously): address is 'icmp/<local-addr>'; close() closes the socket and resolves. There is no per-connection handle (ICMP is connectionless) — reply is bound to the received packet's sender.

Throws: Throws synchronously if the handler is not a function, or if the raw socket can't be opened (typically because it needs root / CAP_NET_RAW). reply rejects on resolve/marshal/write errors and throws synchronously for the raw-body validation rules (a raw body requires type; body is mutually exclusive with id/seq/payload).

```
// Needs root / CAP_NET_RAW. Reply to every echo request with an echo reply:
const srv = server.icmp.listen({}, (msg, reply) => {
  if (msg.type === 8) reply({ type: 0, payload: msg.bytes });
});
runtime.log(srv.address);
await srv.close();
```

server.smtp.listen

```
listen(opts: { port: number; host?: string; hostname?: string; handlers: { onMail:
(env: Envelope) => boolean | string | void | Promise<boolean | string | void>;
onRcpt: (env: Envelope, to: string) => boolean | string | void | Promise<boolean |
string | void>; onData: (env: Envelope, msg: Message) => boolean | string | void |
Promise<boolean | string | void> }; auth?: { user: string, pass: string, env:
Envelope) => boolean | Promise<boolean>; starttls?: { cert: string; key: string };
allowInsecureAuth?: boolean; maxMessageBytes?: number; maxRecipients?: number;
sessionTimeout?: number }): { address: string; stopped: Promise<void>; close():
Promise<void> }
```

Bind an SMTP listener: server.smtp.listen({port, hostname?, handlers: {onMail, onRcpt, onData}, auth?, starttls?, allowInsecureAuth?, maxMessageBytes?, maxRecipients?, sessionTimeout?}) → handle with .address, .close(), .stopped Promise. Handlers receive (envelope, ...) per stage; return true/undefined to accept, false to reject, a string for a 550 reason, throw for 451 temp-fail. onData receives a parsed Message with text/html bodies, attachments, and raw bytes.

Parameters

- opts (*{ port: number; host?: string; hostname?: string; handlers: { onMail: (env: Envelope) => boolean | string | void | Promise<boolean | string | void>; onRcpt: (env: Envelope, to: string) => boolean | string | void | Promise<boolean | string | void>; onData: (env: Envelope, msg: Message) => boolean | string | void | Promise<boolean | string | void> }; auth?: { user: string, pass: string, env: Envelope) => boolean | Promise<boolean>; starttls?: { cert: string; key: string }; allowInsecureAuth?: boolean; maxMessageBytes?: number; maxRecipients?: number; sessionTimeout?: number }*) — port is required; host (bind interface) defaults to "0.0.0.0"; hostname (advertised EHLO domain) defaults to the OS hostname. handlers.onMail/onRcpt/onData are all required and run per protocol stage — each returns true/undefined to accept, false for a 550 reject, a string for '550 <string>', or throws for a 451 temp-fail. auth (optional) enables PLAIN+LOGIN SASL; return truthy to accept. starttls {cert, key} (paths or inline PEM) enables STARTTLS. allowInsecureAuth permits AUTH without TLS. maxMessageBytes defaults to 10 MiB (non-positive values ignored), maxRecipients to 100, sessionTimeout to 30000 (milliseconds).

Returns: A server handle (returned synchronously): address is ‘tcp/host:port’; stopped resolves when the server stops (rejects on a non-close Serve error); close() shuts the listener down and resolves with the stopped Promise. The Envelope passed to handlers is { from, recipients, remote, helo, authenticatedUser?, tls?: { version, cipher } }; the Message passed to onData is { from, to, cc, subject, headers, body: { text, html }, attachments: { filename, contentType, bytes }[], raw: Uint8Array }.

Throws: Throws synchronously if opts is missing, port is 0/absent, handlers is missing, any of onMail/onRcpt/onData is absent or not a function, auth is present but not a function, the starttls key pair fails to load, or the bind fails.

```
const srv = server.smtp.listen({
  port: 2525,
  handlers: {
    onMail: (env) => true,
    onRcpt: (env, to) => to.endsWith("@example.com"),
    onData: (env, msg) => { runtime.log(msg.subject); },
  },
});
```

server.tcp.listen

```
listen(opts: { port: number; host?: string; readBuffer?: number }, handler: (conn:
{ remote: string; local: string; write(data: string | Uint8Array): Promise<void>;
onData(cb: (msg: { bytes: Uint8Array; text: string }) => void): void; onClose(cb:
() => void): void; onError(cb: (err: unknown) => void): void; close(): void }) =>
void): { address: string; close(): Promise<void> }
```

Bind a raw TCP server: server.tcp.listen({port, host?, readBuffer?}, conn => {...}) → handle { address: ‘tcp/host:port’, close() }. The connection handler runs once per accepted socket; conn is the SAME handle shape as net.tcp.connect — onData(cb) (cb gets {bytes, text}), onClose(cb), onError(cb), write(data) (string or Uint8Array), close(), and remote/local addresses. Synchronous bind (throws on bind error); port:0 binds an OS-chosen ephemeral port. Emits a READY line under `sercon serve` and joins graceful shutdown.

Parameters

- `opts` ({ *port*: number; *host*?: string; *readBuffer*?: number }) — port is the listen port (0 binds an OS-chosen ephemeral port). host defaults to all interfaces. readBuffer is the per-connection inbound channel capacity (frames buffered before backpressure), default 64.
- `handler` ((conn: { *remote*: string; *local*: string; *write*(data: string | Uint8Array): Promise<void>; *onData*(cb: (msg: { bytes: Uint8Array; text: string }) => void): void; *onClose*(cb: () => void): void; *onError*(cb: (err: unknown) => void): void; *close*(): void }) => void) — Invoked once per accepted connection. conn matches net.tcp.connect’s handle: register onData/onClose/onError callbacks, write() to send (returns a Promise), close() to tear down. remote/local are ‘host:port’ strings.

Returns: A server handle (returned synchronously): address is ‘tcp/host:port’ (the resolved bind address, so port:0 reports its ephemeral port). close() closes the listener and all accepted connections, then resolves.

Throws: Throws synchronously if `opts` is missing, the handler is not a function, or the bind fails (e.g. address already in use).

```
const srv = server.tcp.listen({ port: 0 }, (conn) => {
  conn.onData((msg) => conn.write("echo: " + msg.text));
});
runtime.log(srv.address);
await srv.close();
```

server.udp.listen

```
listen(opts: { port: number; host?: string }, handler: (msg: { bytes: Uint8Array;
text: string; address: string; port: number }, reply: (data: string | Uint8Array)
=> Promise<void>) => void): { address: string; close(): Promise<void> }
```

Bind a raw UDP server: `server.udp.listen({port, host?}, (msg, reply) => {...})` → handle { address: 'udp/host:port', close() }. The handler runs once per inbound datagram; `msg` is {bytes, text, address, port} (the sender) and `reply(data)` (string or `Uint8Array`) sends a datagram back to that sender, returning a `Promise`. Synchronous bind (throws on bind error); `port:0` binds an OS-chosen ephemeral port. Emits a `READY` line under `sercon serve` and joins graceful shutdown.

Parameters

- `opts` (*{ port: number; host?: string }*) — `port` is the listen port (0 binds an OS-chosen ephemeral port). `host` defaults to all interfaces.
- `handler` (*(msg: { bytes: Uint8Array; text: string; address: string; port: number }, reply: (data: string | Uint8Array) => Promise<void>) => void)* — Invoked once per inbound datagram. `msg` carries the payload (bytes/text) and the sender's address/port. `reply(data)` sends a datagram back to that sender and returns a `Promise` that resolves once written (rejects on a write error).

Returns: A server handle (returned synchronously): `address` is 'udp/host:port' (resolved, so `port:0` reports its ephemeral port). `close()` closes the socket and resolves. There is no per-connection handle — UDP is connectionless, so `reply` is bound to the originating datagram's sender.

Throws: Throws synchronously if `opts` is missing, the handler is not a function, the address fails to resolve, or the bind fails.

```
const srv = server.udp.listen({ port: 0 }, (msg, reply) => {
  reply("pong: " + msg.text);
});
runtime.log(srv.address);
await srv.close();
```

services

Subprocess and external-CLI / service wrappers: shell, git, gh, AI providers, agent-browser automation, W3C WebDriver browser control, Typst typesetting, external-requirement diagnostics (doctor).

services.agentBrowser.auth

Namespace object for the auth vault (session-independent): `auth.save`, `auth.list`, `auth.show`, `auth.delete`. Passwords are never placed in `argv` — `auth.save` sends the password via `stdin` (`-password-stdin`).

services.agentBrowser.auth.delete

```
delete(...args: unknown[])
```

services.agentBrowser.auth.list

```
list(...args: unknown[])
```

services.agentBrowser.auth.save

```
save(...args: unknown[])
```

services.agentBrowser.auth.show

```
show(...args: unknown[])
```

services.agentBrowser.available

```
available
```

True when the agent-browser CLI is on `PATH`. Sync boolean, resolved once per Run. Gate calls on this; every binding throws a clean error when the CLI is absent.

Returns: boolean — true if `agent-browser` is on `PATH`.

```
if (!services.agentBrowser.available) runtime.log("install agent-browser first");
```

services.agentBrowser.clearDefaultOptions

```
clearDefaultOptions(...args: unknown[])
```

Reset the namespace-level launch defaults to an empty object, removing any values set by `setDefaultOptions`.

Returns: void

Throws: Never throws.

```
services.agentBrowser.clearDefaultOptions();  
runtime.log(JSON.stringify(services.agentBrowser.defaultOptions())); // {}
```

services.agentBrowser.defaultOptions

```
defaultOptions(...args: unknown[])
```

Return a shallow copy of the current namespace-level launch defaults. These are merged (under per-call `opts`) into every subsequent `launch()`.

Returns: object — a plain-object copy of the current defaults map. Empty object when no defaults have been set.

Throws: Never throws.

```
services.agentBrowser.setDefaultOptions({ headed: false });
runtime.log(JSON.stringify(services.agentBrowser.defaultOptions())); //
{"headed":false}
```

services.agentBrowser.eval

```
eval(url: string, js: string): void
```

One-shot shortcut: launch an ephemeral session, open url, evaluate a JS expression in the page, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `js` (*string*) — JavaScript expression to evaluate in the page context. Required.

Returns: `Promise<{ success: boolean, data: object, error: string|null }>` — agent-browser envelope; `data.result` holds the serialised return value.

Throws: Throws if url or js is missing; if agent-browser is not on PATH; on navigation or evaluation failure.

```
const r = await services.agentBrowser.eval("data:text/html,<title>Hi</title>",
"document.title");
runtime.log(r.data?.result); // "Hi"
```

services.agentBrowser.launch

```
launch(opts?: { session?: string, headed?: boolean, profile?: string, proxy?:
string, userAgent?: string, device?: string, colorScheme?: string,
ignoreHttpsErrors?: boolean, engine?: string, executablePath?: string, enable?:
string, args?: string, timeout?: number }): { session: string; open(url: string,
opts?: object): Promise<any>; back(): Promise<any>; forward(): Promise<any>;
reload(): Promise<any>; wait(selOrMs: string | number): Promise<any>;
connect(target: string): Promise<any>; frame(target: string): Promise<any>;
click(sel: string): Promise<any>; dblclick(sel: string): Promise<any>; hover(sel:
string): Promise<any>; focus(sel: string): Promise<any>; check(sel: string):
Promise<any>; uncheck(sel: string): Promise<any>; scrollIntoView(sel: string):
Promise<any>; fill(sel: string, text: string): Promise<any>; type(sel: string,
text: string): Promise<any>; press(key: string): Promise<any>; select(sel:
string, ...values: string[]): Promise<any>; scroll(dir: string, px?: number):
Promise<any>; drag(src: string, dst: string): Promise<any>; upload(sel: string,
files: string | string[]): Promise<any>; download(sel: string, path: string):
Promise<any>; keyboard: { type(text: string): Promise<any>; insertText(text:
string): Promise<any> }; mouse: { move(x: number, y: number): Promise<any>;
down(button?: string): Promise<any>; up(button?: string): Promise<any>; wheel(dy:
number, dx?: number): Promise<any> }; get(what: string, sel?: string):
Promise<any>; isVisible(sel: string): Promise<any>; isEnabled(sel: string):
Promise<any>; isChecked(sel: string): Promise<any>; eval(code: string):
Promise<any>; snapshot(opts?: object): Promise<any>; console(opts?: object):
```



```

Promise<any>; errors(opts?: object): Promise<any>; highlight(sel: string):
Promise<any>; find(locator: string, value: string, opts: { action: string, text?:
string }): Promise<any>; locator(spec: object | string, value?: string): object;
set: { viewport(w: number, h: number, scale?: number): Promise<any>; device(name:
string): Promise<any>; geo(lat: number, lng: number): Promise<any>; offline(on?:
boolean): Promise<any>; headers(headers: Record<string, string>): Promise<any>;
credentials(user: string, pass: string): Promise<any>; media(scheme?: "dark" |
"light", reducedMotion?: boolean): Promise<any> }; record: { start(path: string,
url?: string): Promise<any>; stop(): Promise<any> }; screenshot(path?: string,
opts?: { selector?: string, full?: boolean, annotate?: boolean, format?: "png" |
"jpeg", quality?: number }): Promise<{ path?: string, size?: number, bytes?:
number[], format: string }>; pdf(path?: string): Promise<{ path?: string, size?:
number, bytes?: number[], format: string }>; network: { route(url: string, opts?:
{ abort?: boolean, body?: unknown, resourceType?: string }): Promise<any>;
unroute(url?: string): Promise<any>; requests(opts?: { clear?: boolean, filter?:
string, type?: string, method?: string, status?: string }): Promise<any>;
request(id: string): Promise<any>; har: { start(path?: string): Promise<any>;
stop(path?: string): Promise<any> } }; cookies: { get(): Promise<any>; set(name:
string, value: string, opts?: { url?: string, domain?: string, path?: string,
httpOnly?: boolean, secure?: boolean, sameSite?: "Strict" | "Lax" | "None",
expires?: number }): Promise<any>; clear(): Promise<any> }; storage: { local:
{ get(key?: string): Promise<any>; set(key: string, value: string): Promise<any>;
clear(): Promise<any> }; session: { get(key?: string): Promise<any>; set(key:
string, value: string): Promise<any>; clear(): Promise<any> } }; tabs: { list():
Promise<any>; new(url?: string, opts?: { label?: string }): Promise<any>;
close(ref?: string): Promise<any>; select(ref: string): Promise<any> }; diff:
{ snapshot(opts?: { baseline?: string, selector?: string, compact?: boolean,
depth?: number }): Promise<any>; screenshot(opts: { baseline: string, output?:
string, threshold?: number }): Promise<any>; url(url1: string, url2: string):
Promise<any> }; trace: { start(): Promise<any>; stop(path?: string):
Promise<any> }; profiler: { start(opts?: { categories?: string }): Promise<any>;
stop(path?: string): Promise<any>; inspect(): Promise<any>; clipboard(op: "read"
| "write" | "copy" | "paste", text?: string): Promise<any>; vitals(url?: string):
Promise<any>; pushstate(url: string): Promise<any>; react: { tree(): Promise<any>;
inspect(id: string): Promise<any>; renders: { start(): Promise<any>; stop():
Promise<any> }; suspense(opts?: { onlyDynamic?: boolean }): Promise<any> };
stream: { enable(opts?: { port?: number }): Promise<any>; disable(): Promise<any>;
status(): Promise<any> }; chat(message: string, opts?: { model?: string }):
Promise<any>; cmd(command: string, ...args: string[]): Promise<any>; batch(cmds:
string[], opts?: { bail?: boolean }): Promise<any>; auth: { login(name: string):
Promise<any> }; close(): Promise<any> }

```

Allocate a browser session and return a handle. Synchronous (no browser starts until the first command). Pass `opts.session` to name the session; otherwise a unique id is generated. Launch flags (`headed`, `profile`, `proxy`, `userAgent`, `device`, `colorScheme`, `ignoreHttpsErrors`, `engine`, `executablePath`, `enable`, `args`) are threaded into every call the handle makes. Sessions the script does not `close()` are best-effort closed when the Run ends.

Parameters

- `opts` (*{ session?: string, headed?: boolean, profile?: string, proxy?: string, userAgent?: string, device?: string, colorScheme?: string, ignoreHttpsErrors?: boolean, engine?: string, executablePath?: string, enable?: string, args?: string, timeout?: number }, optional*) — Launch flags captured for the lifetime of the handle and threaded into every subprocess

call. session names the agent-browser session (auto-generated when omitted). timeout is the per-call agent-browser subprocess timeout in milliseconds (default 30000, 0 disables); when a call exceeds the limit the Promise rejects with a clear 'timed out' error instead of hanging forever.

Returns: A handle object with a read-only session string and methods: open, back, forward, reload, wait, connect, click, dblclick, hover, focus, fill, type, press, check, uncheck, select, scroll, scrollIntoView, drag, upload, download, keyboard.{type,insertText}, mouse.{move,down,up,wheel}, get, isVisible, isEnabled, isChecked, eval, snapshot, console, errors, highlight, find, locator, close. Every async method resolves to an agent-browser envelope { success: boolean, data: object, error: string|null }; drill into .data for the actual values.

Throws: launch() itself does not throw for a missing CLI (it allocates only); the first method call throws if agent-browser is not on PATH.

```
const b = services.agentBrowser.launch({ headed: false });
await b.open("https://example.com");
const r = await b.get("title");
runtime.log(r.data?.title);
await b.close();
```

services.agentBrowser.pdf

```
pdf(url: string, path?: string): Promise<{ path: string, size: number, format: string } | { bytes: number[], format: string }>
```

One-shot shortcut: launch an ephemeral session, open url, capture a PDF, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `path` (*string, optional*) — Output file path. When supplied the PDF is written there and the result has { path, size, format }; when omitted the PDF bytes are returned as a number[] in { bytes, format }.

Returns: Promise — path given: { path: string, size: number, format: string }; no path: { bytes: number[], format: string }.

Throws: Throws if url is missing; if agent-browser is not on PATH; on navigation, capture, or I/O failure.

```
const pdf = await services.agentBrowser.pdf("data:text/html,<h1>Hi</h1>");
runtime.log(new Uint8Array(pdf.bytes).length, pdf.format); // e.g. 5678 pdf
```

services.agentBrowser.screenshot

```
screenshot(url: string, path?: string, opts?: { selector?: string, full?: boolean,
annotate?: boolean, format?: "png" | "jpeg", quality?: number }): Promise<{ path: string, size: number, format: string } | { bytes: number[], format: string }>
```

One-shot shortcut: launch an ephemeral session, open url, capture a screenshot, and close. Equivalent to launch()+open(url)+screenshot(path?,opts?)+close() but in a single call.

Parameters

- `url` (*string*) — URL to open (http/https or data: URI). Required.
- `path` (*string, optional*) — Output file path. When supplied the screenshot is written there and the result has { `path`, `size`, `format` }; when omitted the image bytes are returned as a `number[]` (byte-value array) in { `bytes`, `format` }.
- `opts` (*{ selector?: string, full?: boolean, annotate?: boolean, format?: "png" | "jpeg", quality?: number }, optional*) — Capture options. `selector` scopes the capture to an element. `full` captures the full page. `format` defaults to `png`. `quality` (0–100) applies to `jpeg`.

Returns: Promise — `path` given: { `path`: string, `size`: number, `format`: string }; no `path`: { `bytes`: number[], `format`: string } where `bytes` is a plain JS `number[]` (byte-value array); wrap with `new Uint8Array(bytes)` to get a typed array.

Throws: Throws if `url` is missing; if `agent-browser` is not on `PATH`; on navigation, capture, or I/O failure.

```
const shot = await services.agentBrowser.screenshot("data:text/html,<h1>Hi</h1>");
runtime.log(new Uint8Array(shot.bytes).length, shot.format); // e.g. 12345 png
```

services.agentBrowser.setDefaultOptions

```
setDefaultOptions(opts: object): void
```

Replace the namespace-level launch defaults with the supplied object. Merged (under per-call `opts`) into every subsequent `launch()`. Affects only the current Run.

Parameters

- `opts` (*object*) — A plain object of launch option key/value pairs. The entire defaults map is replaced (not merged) with this object.

Returns: void

Throws: Never throws.

```
services.agentBrowser.setDefaultOptions({ headed: false, proxy: "http://proxy:3128" });
const b = services.agentBrowser.launch(); // headed:false + proxy inherited
```

services.agentBrowser.snapshot

```
snapshot(url: string, opts?: { interactive?: boolean, compact?: boolean, depth?: number, selector?: string }): void
```

One-shot shortcut: launch an ephemeral session, open `url`, take an accessibility-tree snapshot, and close.

Parameters

- `url` (*string*) — URL to open. Required.
- `opts` (*{ interactive?: boolean, compact?: boolean, depth?: number, selector?: string }, optional*) — Snapshot options forwarded to the handle's `snapshot()` method.

Returns: Promise<{ success: boolean, data: object, error: string|null }> — agent-browser envelope; data contains the accessibility tree.

Throws: Throws if url is missing; if agent-browser is not on PATH; on navigation or snapshot failure.

```
const snap = await services.agentBrowser.snapshot("data:text/html,<h1>Hi</h1>",
{ compact: true });
runtime.log(JSON.stringify(snap.data).slice(0, 100));
```

services.agentBrowser.version

```
version(...args: unknown[]): Promise<string>
```

The agent-browser CLI version string.

Returns: Promise<string> — the version reported by `agent-browser --version`.

Throws: Throws if the agent-browser CLI is not on PATH.

```
runtime.log(await services.agentBrowser.version());
```

services.ai.providers

```
providers(): string[]
```

Which of claude / codex / copilot / gemini are on PATH, in preference order.

Returns: string[] — the subset of supported AI CLIs found on PATH, in preference order (claude, codex, copilot, gemini); an empty array when none are installed. Synchronous (not a Promise).

Throws: Never throws.

```
const ps = services.ai.providers(); // e.g. ["claude", "gemini"]
```

services.ai.send

```
send(opts: { prompt: string, provider?: "claude" | "codex" | "copilot" | "gemini",
system?: string, context?: string, timeout?: number }): Promise<Record<string,
unknown>>
```

Run a one-shot prompt through a provider. opts { prompt (required), provider?, system?, context?, timeout? }. Returns { provider, output, exitCode }. Non-zero exit is data; no provider throws.

Parameters

- opts (*{ prompt: string, provider?: "claude" | "codex" | "copilot" | "gemini", system?: string, context?: string, timeout?: number }*) — prompt is required. provider names the CLI to use; when omitted, the first provider on PATH (in preference order) is chosen. system and context are prepended to the prompt as "System: ..." / "Context: ..." blocks (a portable substitute for each CLI's own flags). timeout in ms (default 120000).

Returns: Promise<{ provider: string, output: string, exitCode: number }> — provider is the CLI that ran; output is its trimmed stdout (or stderr when stdout is empty on a non-zero exit); exitCode is 0 on success.

Throws: Throws if opts.prompt is missing/empty, no provider is on PATH (when provider is unset), the named provider is unknown, the CLI fails to spawn, or on timeout / context cancellation. A non-zero exit code does NOT throw — it resolves with that exitCode.

```
const r = await services.ai.send({ prompt: "Say hi", provider: "claude" });
runtime.log(r.output);
```

services.doctor

```
doctor(requires?: string[]): Promise<{ ok: boolean; satisfied: boolean; unmet:
string[]; tools: { name: string; category: string; purpose: string; installed:
boolean; version: string | null; ok: boolean; detail?: string }[] }>
```

Report on every external tool sercon can use (git, gh, AI providers, agent-browser, chromedriver/geckodriver, typst, recon/curl, clipboard/image tools): installed?, version, purpose — and validate the chromedriver↔Chrome major-version match. Optionally assert a script's prerequisites via `requires`.

Parameters

- `requires` (*string[], optional*) — Feature/category names (e.g. “webdriver”, “typst”, “ai”, “git”, “gh”, “agentBrowser”, “clipboard”, “image”, “http”) OR specific binaries (e.g. “chromedriver”, “pngpaste”, “claude”) to assert. Each unmet requirement (absent or in a compatibility conflict) lands in the returned `unmet` array — a missing optional tool is reported, not thrown. An unrecognized name throws (catches typos).

Returns: Promise<{ ok, satisfied, unmet, tools }> — ok is false only when a compatibility conflict was detected (e.g. chromedriver↔Chrome major mismatch); satisfied is true when every entry in requires is met (unmet is empty); unmet lists the requested requirements that are absent or conflicted; tools is the full report, one entry per probed tool ({ name, category, purpose, installed, version (string|null), ok, optional detail }). A missing tool is installed:false, ok:true (optional, not a failure).

Throws: Throws if a requires entry matches neither a known category/feature nor a known binary name (an unknown requirement). Missing tools do NOT throw — they are reported via installed:false and, when requested, listed in unmet.

```
const r = await services.doctor(["git"]);
runtime.log("git satisfied:", r.satisfied, "of", r.tools.length, "tools");
const opt = await services.doctor(["typst", "webdriver"]);
runtime.log("unmet optional features:", JSON.stringify(opt.unmet));
```

services.exec.http

```
http(method: string, url: string, opts?: { headers?: Record<string, string>,
body?: string, timeout?: number, follow?: boolean, insecure?: boolean, backend?:
"auto" | "recon" | "curl" }): Promise<Record<string, unknown>>
```

Make an HTTP request by shelling out to recon (preferred) or curl (fallback). 4xx/5xx resolve as status; transport errors and timeouts throw. `opts.backend = 'auto' | 'recon' | 'curl'`.

Parameters

- `method` (*string*) — HTTP verb (GET, POST, PUT, DELETE, PATCH, HEAD); lower-case input is uppercased before forwarding.
- `url` (*string*) — Target URL; must be fully qualified (the backend requires a scheme + host).
- `opts` (*{ headers?: Record<string, string>, body?: string, timeout?: number, follow?: boolean, insecure?: boolean, backend?: "auto" | "recon" | "curl" }, optional*) — headers emits one `-H "Name: Value"` per entry. body is written to a temp file and sent via `-data-binary` so CR/LF stay intact. timeout in ms (default 30000). follow toggles `-L` to follow 3xx redirects. insecure toggles `-k` to skip TLS verification. backend picks the tool: 'auto' (default) prefers recon then curl; 'recon' or 'curl' require that specific binary on PATH.

Returns: `Promise<{ status: number, headers: Record<string, string>, body: string, durationMs: number, backend: "recon" | "curl" }>` — status is the final HTTP status code; headers have lower-cased names (last response block on a redirect chain); body is the UTF-8 decoded response body; backend is whichever tool ran.

Throws: Throws if method or url is missing/empty; if the requested backend (or, for 'auto', neither recon nor curl) is on PATH; on transport errors (DNS failure, connection refused, TLS handshake); on timeout or context cancellation; or if the response headers can't be parsed. HTTP 4xx/5xx do NOT throw — they resolve with that status.

```
const r = await services.exec.http("GET", "https://example.com");
runtime.log(r.status, r.backend);
```

services.exec.shell

```
shell(cmd: string | string[], opts?: { timeout?: number, cwd?: string, stdin?:
string, env?: Record<string, string>, pane?: string | Pane, pty?: boolean }):
Promise<Record<string, unknown>>
```

Run a subprocess and wait for it to exit. String `cmd` → `/bin/sh -c` (or `cmd /C` on Windows); array `cmd` → `argv` (no shell). Non-zero exits resolve normally; spawn failures and timeouts throw.

Parameters

- `cmd` (*string | string[]*) — A string is passed verbatim to the host shell (`/bin/sh -c` on Unix, `cmd /C` on Windows) so quoting, pipes, and redirects work. A `string[]` is treated as `argv`: `argv[0]` is run directly with no shell, so use this form when arguments contain whitespace or shell metacharacters you don't want re-interpreted.
- `opts` (*{ timeout?: number, cwd?: string, stdin?: string, env?: Record<string, string>, pane?: string | Pane, pty?: boolean }, optional*) — timeout in ms (default 30000); on expiry the process tree is killed and the call throws. `cwd` sets the working directory. `stdin` is fed to the process's standard input. `env` entries are merged on top of the inherited environment (they do not replace it). `pane` (a `tui.pane` name or `Pane` handle) streams `stdout+stderr` live into a TUI pane — in that mode the result's `stdout/stderr` strings stay empty. `pty` (default false) runs the command under a pseudo-terminal so it believes it is a terminal and emits

color/progress; with a pane the output is rendered there, without a pane it is captured into stdout (stderr stays empty since a pty merges both streams). Unix only — on Windows pty is ignored and the normal pipe path is used.

Returns: `Promise<{ stdout: string, stderr: string, exitCode: number, success: boolean, durationMs: number }>` — stdout/stderr are captured (empty when streamed to a pane); exitCode is 0 on success; success is `exitCode === 0`; durationMs is wall-clock spawn-to-exit time.

Throws: Throws if cmd is missing, an empty string, an empty array, or a non-string array element; if the host binary is not on PATH or fails to start; if the timeout (or context cancellation) fires before exit; or if a named pane is referenced without a prior `tui.layout`. A non-zero exit code does NOT throw — it resolves with `success:false`.

```
const r = await services.exec.shell("echo hi");
if (r.success) runtime.log(r.stdout.trim());
```

services.exec.stream

```
stream(cmd: string | string[], onLine: (line: string, stream: "stdout" | "stderr")
=> void, opts?: { cwd?: string, env?: Record<string, string>, stdin?: string,
timeout?: number }): Promise<{ exitCode: number; success: boolean; durationMs:
number }>
```

Run a subprocess and stream its stdout/stderr to a callback line by line as output arrives (unlike `exec.shell`, which buffers). String `cmd` → `/bin/sh -c` (or `cmd /C` on Windows); array `cmd` → `argv` (no shell). Resolves `{ exitCode, success, durationMs }` on exit; non-zero exits resolve normally; spawn failures and timeouts reject.

Parameters

- `cmd` (*string* / *string[]*) — A string is passed to the host shell (`/bin/sh -c` on Unix, `cmd /C` on Windows) so pipes, redirects, and globs work. A `string[]` is treated as `argv`: `argv[0]` is run directly with no shell.
- `onLine` (*(line: string, stream: "stdout" | "stderr") => void*) — Called once per output line as it arrives. `line` has its trailing newline stripped; `stream` is 'stdout' or 'stderr'. A final line without a trailing newline is still delivered. Required — a non-function throws synchronously.
- `opts` (*{ cwd?: string, env?: Record<string, string>, stdin?: string, timeout?: number }*, optional) — `cwd` sets the working directory; `env` entries merge on top of the inherited environment; `stdin` is fed to the process. `timeout` is in ms with NO default (0 / absent = run until exit, unlike `exec.shell`'s 30000); when set, the process tree is killed on expiry and the call rejects.

Returns: `Promise<{ exitCode: number, success: boolean, durationMs: number }>` — resolves on process exit. `exitCode` is 0 on success; success is `exitCode === 0`; durationMs is wall-clock spawn-to-exit time. Output is delivered via `onLine`, not captured into the result.

Throws: Throws synchronously if `cmd` is missing or `onLine` is not a function. The Promise rejects if the host binary is not on PATH or fails to start, or if the timeout (or context cancellation) fires before exit. A non-zero exit code does NOT reject — it resolves with `success:false`.

```
const r = await services.exec.stream("echo one; echo two", (line, stream) => {
  runtime.log(stream, line);
});
runtime.log("exit", r.exitCode);
```

services.gh.authStatus

```
authStatus(...args: unknown[]): Promise<{ authenticated: boolean; user: string;
raw: string }>
```

Probe gh’s auth state. Missing gh / unauthenticated resolve with { authenticated: false, ... } — only context cancellation throws.

Returns: Promise<{ authenticated: boolean, user: string, raw: string }> — authenticated is true only when gh api user succeeds; user is the resolved login (“” when not authenticated); raw is the login on success or the underlying gh error / “gh not on PATH” otherwise.

Throws: Throws only on context cancellation. A missing gh binary or an unauthenticated session resolve with authenticated:false rather than throwing.

```
const a = await services.gh.authStatus();
if (a.authenticated) runtime.log("logged in as", a.user);
```

services.gh.prList

```
prList(opts?: { cwd?: string, state?: string, limit?: number, author?: string }):
Promise<Record<string, unknown>[]>
```

List pull requests on the cwd’s repo (or opts.cwd). Defaults: open state, limit 30. Filters: state / limit / author.

Parameters

- `opts` (*{ cwd?: string, state?: string, limit?: number, author?: string }, optional*) — `cwd` selects the repo (defaults to the engine’s working directory, which gh uses to detect the repo). `state` filters by PR state (“open” default, “closed”, “merged”, “all”). `limit` caps results (default 30; must be positive). `author` filters to PRs opened by that login.

Returns: Promise<Array<{ number: number, title: string, state: string, author: string, headRefName: string, baseRefName: string, url: string, createdAt: string, updatedAt: string }>> — one object per PR; author is flattened from gh’s { login } wrapper to the bare login string; createdAt/updatedAt are ISO 8601 timestamps.

Throws: Throws if limit is <= 0, gh is not on PATH, gh exits non-zero (not authenticated, not a GitHub repo, etc.), the JSON can’t be parsed, or on context cancellation.

```
const prs = await services.gh.prList({ state: "open", limit: 5 });
for (const pr of prs) runtime.log(pr.number, pr.title);
```

services.gh.repoView


```
repoView(repo?: string, opts?: { cwd?: string }): Promise<Record<string, unknown>>
```

Repo metadata. With no arg uses cwd's repo; pass 'owner/name' for any repo gh can see. owner + defaultBranch are pre-flattened.

Parameters

- `repo` (*string, optional*) — "owner/name" of any repo gh can access. Omit (or pass opts as the first arg) to view the repo detected from cwd.
- `opts` (*{ cwd?: string }, optional*) — cwd selects the checkout gh uses to detect the current repo when repo is omitted.

Returns: Promise<{ name: string, owner: string, description: string, url: string, defaultBranch: string, visibility: string }> — owner is flattened from gh's { login } wrapper to the bare login; defaultBranch is flattened from defaultBranchRef.name (" " if absent); key order matches gh's output.

Throws: Throws if gh is not on PATH, gh exits non-zero (repo not found, not authenticated, etc.), the JSON can't be parsed, or on context cancellation.

```
const r = await services.gh.repoView("cli/cli");
runtime.log(r.owner, r.defaultBranch);
```

services.git.add

```
add(paths: string | string[], opts?: { cwd?: string }): Promise<{ paths:
string[] }>
```

Stage one path (string) or several (string[]).

Parameters

- `paths` (*string / string[]*) — Path or paths to stage. Passed after a `--` separator so paths that look like flags (-foo) are handled literally.
- `opts` (*{ cwd?: string }, optional*) — cwd selects the checkout; defaults to the engine's working directory.

Returns: Promise<{ paths: string[] }> — the list of paths that were staged.

Throws: Throws if paths is missing, an empty string, or contains a non-string array element; if git is not on PATH; or if git add exits non-zero (e.g. a pathspec matching nothing).

```
await services.git.add(["src/a.ts", "src/b.ts"]);
```

services.git.branch

```
branch(opts?: { cwd?: string }): Promise<{ current: string; detached: boolean;
all: string[] }>
```

Current branch (empty when HEAD is detached) plus the list of local branches.

Parameters

- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout to inspect; defaults to the engine's working directory.

Returns: `Promise<{ current: string, detached: boolean, all: string[] }>` — `current` is the checked-out branch name ("" when detached); `detached` is true on a detached HEAD; all lists every local branch (`refs/heads`) by short name.

Throws: Throws if git is not on PATH, the directory is not a git repository, or the underlying git command fails. A detached HEAD is reported via `detached:true`, not a throw.

```
const b = await services.git.branch();
runtime.log(b.detached ? "(detached)" : b.current);
```

services.git.commit

```
commit(message: string, opts?: { cwd?: string, allowEmpty?: boolean }):
Promise<{ sha: string }>
```

Create a commit; returns the post-commit HEAD SHA. `opts.allowEmpty` toggles `–allow-empty`.

Parameters

- `message` (*string*) — Commit message (passed as a single `-m` argument).
- `opts` (*{ cwd?: string, allowEmpty?: boolean }, optional*) — `cwd` selects the checkout. `allowEmpty` adds `–allow-empty` so a commit succeeds with no staged changes (release markers, etc.); defaults to false.

Returns: `Promise<{ sha: string }>` — `sha` is the full SHA of the newly created HEAD commit.

Throws: Throws if message is missing or blank, git is not on PATH, or git commit exits non-zero (e.g. nothing staged and `allowEmpty` is false).

```
const c = await services.git.commit("chore: bump", { allowEmpty: true });
```

services.git.diffStat

```
diffStat(opts?: { cwd?: string, revRange?: string }): Promise<{ files: number;
insertions: number; deletions: number }>
```

Aggregate `{ files, insertions, deletions }` from `git diff --shortstat`. Default `revRange` HEAD 1..HEAD.

Parameters

- `opts` (*{ cwd?: string, revRange?: string }, optional*) — `cwd` selects the checkout. `revRange` is the diff range (default "HEAD 1..HEAD", the last commit).

Returns: `Promise<{ files: number, insertions: number, deletions: number }>` — counters parsed from `git diff –shortstat`. An empty diff returns all zeros.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git diff exits non-zero (e.g. a bad `revRange`).

```
const d = await services.git.diffStat();
runtime.log(d.files, d.insertions, d.deletions);
```

services.git.isClean

```
isClean(opts?: { cwd?: string }): Promise<boolean>
```

True iff `git status --porcelain` is empty.

Parameters

- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: `Promise<boolean>` — true when the working tree has no staged, unstaged, or untracked changes.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git status exits non-zero.

```
if (await services.git.isClean()) runtime.log("clean");
```

services.git.log

```
log(opts?: { cwd?: string, limit?: number, revRange?: string }): Promise<{ sha: string; shortSha: string; author: string; email: string; timestamp: number; subject: string }[]>
```

Recent commits as `{ sha, shortSha, author, email, timestamp, subject }`. `opts.limit / opts.revRange`.

Parameters

- `opts` (*{ cwd?: string, limit?: number, revRange?: string }, optional*) — `cwd` selects the checkout. `limit` caps the number of commits (default 50; must be positive). `revRange` selects the range/ref to walk (default "HEAD").

Returns: `Promise<Array<{ sha: string, shortSha: string, author: string, email: string, timestamp: number, subject: string }>>` — newest first; timestamp is the author Unix epoch seconds; subject is the commit's first line.

Throws: Throws if limit is `<= 0`, git is not on PATH, the directory is not a git repository, or git log exits non-zero (e.g. an unknown `revRange`).

```
const log = await services.git.log({ limit: 5 });
runtime.log(log[0].subject);
```

services.git.revParse

```
revParse(rev: string, opts?: { cwd?: string }): Promise<string>
```

Full 40-char SHA for the given rev. Invalid refs throw.

Parameters

- `rev` (*string*) — Any revision git understands (branch, tag, HEAD, short SHA, HEAD 2, etc.).
- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: `Promise<string>` — the full 40-character commit SHA the `rev` resolves to.

Throws: Throws if `rev` is missing or empty, git is not on PATH, the directory is not a git repository, or the `rev` cannot be resolved (git's own error message is included).

```
const sha = await services.git.revParse("HEAD");
```

services.git.runText

```
runText(args: string | string[], opts?: { cwd?: string }): Promise<{ stdout: string; stderr: string; exitCode: number }>
```

Escape hatch: run any `git <args>`, get `{ stdout, stderr, exitCode }` — `exitCode` is data, not a throw.

Parameters

- `args` (*string | string[]*) — git arguments (without the leading “git”), e.g. [“tag”, “-list”]. A bare string is treated as a single argument.
- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: `Promise<{ stdout: string, stderr: string, exitCode: number }>` — captured streams plus git's exit code, so callers can react to any exit status without try/catch.

Throws: Throws if `args` is missing, an empty string, or an empty array; if git is not on PATH; or on a spawn failure / context cancellation. A non-zero exit code does NOT throw — it is returned in `exitCode`.

```
const r = await services.git.runText(["tag", "--list"]);
if (r.exitCode === 0) runtime.log(r.stdout);
```

services.git.status

```
status(opts?: { cwd?: string }): Promise<{ path: string; indexStatus: string; workingStatus: string }[]>
```

Parsed `git status --porcelain` entries: `{ path, indexStatus, workingStatus }`.

Parameters

- `opts` (*{ cwd?: string }, optional*) — `cwd` selects the checkout; defaults to the engine's working directory.

Returns: `Promise<Array<{ path: string, indexStatus: string, workingStatus: string }>>` — one entry per changed path; `indexStatus` / `workingStatus` are the porcelain v1 X / Y status characters (e.g. “M”, “A”, “?”). An empty array means a clean tree.

Throws: Throws if git is not on PATH, the directory is not a git repository, or git status exits non-zero.

```
for (const e of await services.git.status())
  runtime.log(e.indexStatus + e.workingStatus, e.path);
```

services.typst.available

```
available: boolean
```

True when the typst CLI is on PATH. Sync boolean, resolved once per Run. Gate calls on this — every other typst binding throws a clean error when the CLI is absent.

Returns: boolean — true if `typst` is found on PATH.

Throws: Never throws.

```
if (!services.typst.available) {
  runtime.log("install typst (brew install typst) first");
}
```

services.typst.compile

```
compile(opts: { input?: string, source?: string, output?: string, format?: "pdf" |
"png" | "svg", root?: string, inputs?: Record<string, string>, ppi?: number,
fontPaths?: string[], timeout?: number }): Promise<{ format: string; bytes?:
Uint8Array; path?: string }>
```

Compile a Typst document to PDF/PNG/SVG. Provide exactly one of `input` (a .typ path) or `source` (inline Typst). With no `output`, a PDF is compiled to a temp file and returned as bytes (PDF only); with an `output` path the result is written there and `format` is inferred from the extension (png/svg require an output path).

Parameters

- `opts` (*{ input?: string, source?: string, output?: string, format?: "pdf" | "png" | "svg", root?: string, inputs?: Record<string, string>, ppi?: number, fontPaths?: string[], timeout?: number }*) — Provide exactly one of `input` (a path to a .typ file) or `source` (inline Typst markup) — passing both, or neither, throws. `output` is the destination path; when omitted a PDF is produced in a temp dir and returned as bytes (PDF only — png/svg require an output path). `format` (pdf/png/svg) is inferred from `output`'s extension when omitted, defaulting to pdf when there is no output. `root` sets the project root for absolute imports. `inputs` are sys.inputs key/value pairs passed as `–input k=v` (deterministic order). `ppi` sets PNG resolution (png only). `fontPaths` adds `–font-path` search dirs. `timeout` in ms (default 60000). Inline source caveat: `source` is written to a temp `main.typ`, so relative imports/reads resolve against that temp dir, not your cwd — use `input` (or `root`) when the document imports or reads sibling files.

Returns: `Promise<{ format, bytes?, path? }>` — `format` echoes the chosen format. With no `output`: `{ format, bytes }` where `bytes` is the compiled PDF as a `Uint8Array`. With an `output` path: `{ format, path }` where `path` is the written file.

Throws: Throws if neither or both of `input/source` are given; if `format` is not pdf/png/svg; if png/svg is requested without an output path; if the output extension can't be mapped to a

format; if typst is not on PATH; if typst exits non-zero (the trimmed stderr is included); or on timeout / context cancellation.

```
// Inline source → PDF bytes.
const pdf = await services.typst.compile({ source: "= Hi\nBody." });
runtime.log(pdf.format, pdf.bytes.length);

// .typ file → PNG on disk.
await services.typst.compile({ input: "report.typ", output: "/tmp/report.png",
  ppi: 144 });
```

services.typst.fonts

```
fonts(...args: unknown[]): Promise<string[]>
```

List the font families typst can see (from `typst fonts`), de-duplicated and sorted.

Returns: `Promise<string[]>` — sorted, unique font family names available to typst.

Throws: Throws if typst is not on PATH or on timeout / context cancellation.

```
const families = await services.typst.fonts();
runtime.log("fonts:", families.length);
```

services.typst.query

```
query(opts: { selector: string, input?: string, source?: string, field?: string,
  one?: boolean, root?: string, inputs?: Record<string, string>, timeout?:
  number }): Promise<unknown>
```

Query a compiled Typst document for elements matching a selector and return the result as parsed JSON. Provide exactly one of `input` or `source` `selector` is required.

Parameters

- `opts` (*{ selector: string, input?: string, source?: string, field?: string, one?: boolean, root?: string, inputs?: Record<string, string>, timeout?: number }*) — `selector` is required — a Typst query selector (e.g. `<label>`, `<heading>`, `<figure>`). Provide exactly one of `input` (a .typ path) or `source` (inline Typst). `field` extracts a single field from each match (`–field`). `one` returns just the first match instead of an array (`–one`). `root` sets the project root; `inputs` are `–input k=v` sys.inputs pairs; `timeout` in ms (default 60000). Same inline-source caveat as `compile`: `source` is written to a temp file, so relative imports/reads resolve there.

Returns: `Promise<unknown>` — the parsed JSON typst emits: an array of matched elements (or matched field values when `field` is set), or a single value when `one` is set.

Throws: Throws if `selector` is missing; if neither or both of `input/source` are given; if typst is not on PATH; if typst exits non-zero (the trimmed stderr is included); if the output isn't valid JSON; or on timeout / context cancellation.

```
const v = await services.typst.query({
  source: "#metadata(42) <answer>",
  selector: "<answer>", field: "value", one: true,
```

```
});
runtime.log(v); // 42
```

services.typst.version

```
version(...args: unknown[]): Promise<string>
```

The typst CLI version string (from `typst --version`).

Returns: `Promise<string>` — the trimmed version line reported by `typst --version` (e.g. “typst 0.12.0 (...”).

Throws: Throws if typst is not on PATH or on timeout / context cancellation.

```
runtime.log(await services.typst.version());
```

services.webdriver.available

```
available
```

True when a W3C WebDriver binary (chromedriver or geckodriver) is on PATH. Sync boolean, resolved once per Run. Gate calls on this before using probe or connect.

Returns: boolean — true if chromedriver or geckodriver is found on PATH.

```
if (!services.webdriver.available) {
  runtime.log("install chromedriver or geckodriver first");
}
```

services.webdriver.connect

```
connect(opts?: { browser?: "chrome" | "firefox", headless?: boolean, url?: string,
args?: string[], capabilities?: object, commandTimeout?: number }):
Promise<{ get(url: string): Promise<{ ok: true }>; url(): Promise<string>;
title(): Promise<string>; back(): Promise<{ ok: true }>; forward(): Promise<{ ok:
true }>; refresh(): Promise<{ ok: true }>; find(by: string, value: string):
Promise<{ click(): Promise<{ ok: true }>; sendKeys(text: string): Promise<{ ok:
true }>; clear(): Promise<{ ok: true }>; submit(): Promise<{ ok: true }>; text():
Promise<string>; getAttribute(name: string): Promise<string>; cssValue(name:
string): Promise<string>; tagName(): Promise<string>; isDisplayed():
Promise<boolean>; isEnabled(): Promise<boolean>; isSelected(): Promise<boolean>;
find(by: string, value: string): Promise<any>; findAll(by: string, value: string):
Promise<any[]>; screenshot(path?: string): Promise<{ path?: string; size?: number;
bytes?: number[]; format: "png" }> }>; findAll(by: string, value: string):
Promise<Array<{ click(): Promise<{ ok: true }>; sendKeys(text: string):
Promise<{ ok: true }>; clear(): Promise<{ ok: true }>; submit(): Promise<{ ok:
true }>; text(): Promise<string>; getAttribute(name: string): Promise<string>;
cssValue(name: string): Promise<string>; tagName(): Promise<string>;
isDisplayed(): Promise<boolean>; isEnabled(): Promise<boolean>; isSelected():
Promise<boolean>; find(by: string, value: string): Promise<any>; findAll(by:
string, value: string): Promise<any[]>; screenshot(path?: string):
Promise<{ path?: string; size?: number; bytes?: number[]; format: "png" }> }>;
source(): Promise<string>; screenshot(path?: string): Promise<{ path?: string;
```

```

size?: number; bytes?: number[]; format: "png" }>; executeScript(js: string,
args?: unknown[]): Promise<unknown>; executeScriptAsync(js: string, args?:
unknown[]): Promise<unknown>; cookies(): Promise<object[]>; setCookie(c: { name:
string; value: string; path?: string; domain?: string; secure?: boolean;
httpOnly?: boolean; expiry?: number }): Promise<{ ok: true }>; deleteCookie(name:
string): Promise<{ ok: true }>; deleteAllCookies(): Promise<{ ok: true }>;
setImplicitWait(ms: number): Promise<{ ok: true }>; waitFor(by: string, value:
string, opts?: { timeout?: number; visible?: boolean; enabled?: boolean }):
Promise<any>; clickWhenReady(by: string, value: string, opts?: { timeout?: number;
visible?: boolean; enabled?: boolean; poll?: number }): Promise<{ ok: true }>;
windowHandles(): Promise<string[]>; currentWindow(): Promise<string>;
switchToWindow(handle: string): Promise<{ ok: true }>; newWindow(type?: "tab" |
"window"): Promise<{ handle: string; type: string }>; closeWindow():
Promise<string[]>; switchToFrame(target: number | string | object): Promise<{ ok:
true }>; frameChain(targets: (number | string | object)[]): Promise<{ ok: true }>;
switchToParentFrame(): Promise<{ ok: true }>; switchToDefaultContent():
Promise<{ ok: true }>; acceptAlert(): Promise<{ ok: true }>; dismissAlert():
Promise<{ ok: true }>; alertText(): Promise<string>; sendAlertText(text: string):
Promise<{ ok: true }>; maximize(): Promise<{ x: number; y: number; width: number;
height: number }>; minimize(): Promise<{ x: number; y: number; width: number;
height: number }>; fullscreen(): Promise<{ x: number; y: number; width: number;
height: number }>; setWindowRect(rect: { width?: number; height?: number; x?:
number; y?: number }): Promise<{ x: number; y: number; width: number; height:
number }>; getWindowRect(): Promise<{ x: number; y: number; width: number; height:
number }>; hover(el: object): Promise<{ ok: true }>; dragAndDrop(src: object, dst:
object): Promise<{ ok: true }>; keyChord(...keys: string[]): Promise<{ ok: true }
>; performActions(sequence: unknown[]): Promise<{ ok: true }>; releaseActions():
Promise<{ ok: true }>; cdp(command: string, params?: object): Promise<unknown>;
cdpClick(by: string, value: string, opts?: { timeout?: number; poll?: number;
button?: "left" | "right" | "middle"; scrollIntoView?: boolean; offsetX?: number;
offsetY?: number }): Promise<{ clicked: true; x: number; y: number; targetId?:
string }>; targets(): Promise<Array<{ targetId: string; type: string; url: string;
title: string }>>; attach(target: string | { targetId: string }):
Promise<{ targetId: string; sessionId: string; cdp(method: string, params?:
object): Promise<unknown>; detach(): Promise<{ detached: true }> }>; quit():
Promise<{ closed: true }> }>

```

Connect to a running WebDriver server (opts.url) or start an installed local chromedriver/geckodriver and dial it. Returns a session handle whose methods drive the browser. Sessions are quit on Run end if the script does not call quit() explicitly.

Parameters

- **opts** ({ *browser?*: "chrome" | "firefox", *headless?*: boolean, *url?*: string, *args?*: string[], *capabilities?*: object, *commandTimeout?*: number }, optional) — browser selects the driver binary (default 'chrome'). headless defaults to true. url, if given, dials an already-running driver at that base URL instead of starting one. args appends extra browser flags. capabilities is an escape hatch for raw W3C capability overrides merged last. commandTimeout (ms, default 30000) bounds each low-level WebDriver request so a driver blocked behind an open alert or an unreachable endpoint can't hang the call; 0 or negative disables the per-command deadline. quit()/Run-end also cancel any in-flight command.

Returns: Promise resolving to a session handle with methods: `get(url)`, `url()`, `title()`, `back()`, `forward()`, `refresh()`, `find(by, value) → element handle`, `findAll(by, value) → element handle[]`, `source()`, `screenshot(path?)`, `executeScript(js, args?)`, `executeScriptAsync(js, args?)`, `cookies()`, `setCookie(c)`, `deleteCookie(name)`, `deleteAllCookies()`, `setImplicitWait(ms)`, `waitFor(by, value, opts?)` (`opts.enabled` waits past a disabled→enabled flip; `opts.visible` waits for visibility), `clickWhenReady(by, value, opts?)` (waits — by default visible+enabled — in the active frame, then issues a native trusted click that fires React handlers; `opts.poll` sets the inter-attempt interval), `quit()`. `find/clickWhenReady` operate in the active frame set via `switchToFrame/frameChain`. Phase 2 session methods: `windowHandles()`, `currentWindow()`, `switchToWindow(handle)`, `newWindow(type?)`, `closeWindow()`, `switchToFrame(selectorOrIndexOrElement)`, `frameChain([...targets])` (nested: switch from the top document through each level), `switchToParentFrame()`, `switchToDefaultContent()`, `acceptAlert()`, `dismissAlert()`, `alertText()`, `sendAlertText(text)`, `maximize()`, `minimize()`, `fullscreen()`, `setWindowRect({width?,height?,x?,y?})`, `getWindowRect()`, `hover(el)`, `dragAndDrop(src,dst)`, `keyChord(...keys)`, `performActions(sequence)`, `releaseActions()`. `cdp(command, params?)` sends one raw Chrome DevTools Protocol command (Chrome-only) and returns its result; `cdpClick(by, value, opts?)` locates the element across the whole frame tree (including nested cross-origin iframes), scrolls it into view, and dispatches a trusted CDP mouse click at its true viewport coordinates — the way to activate a button the W3C Element Click hit-tests to a parent iframe and intercepts (e.g. a Klarna Checkout “Pay order”). `cdp/cdpClick` throw on firefox. `cdpClick` now crosses out-of-process iframe (OOPIF) boundaries by dispatching input over a browser-level CDP connection; `targets()` lists the browser’s CDP targets (incl. cross-site OOPIF iframe: targets) and `attach(target)` returns a session handle whose `cdp(method, params?)` runs a CDP command against that target (e.g. an OOPIF) and `detach()` releases it. `targets()/attach()` require a CDP-capable session (local chromedriver, or a Grid advertising `se:cdp`). Element handles also expose `hover()` and `dragTo(target)` and carry an `elementId` string. `executeScript/executeScriptAsync` return element handles when the script returns an element (or a top-level array of elements). Locator strategies: `css`, `xpath`, `id`, `name`, `tag`, `className`, `linkText`, `partialLinkText`.

Throws: Throws if no url is given and the driver binary for the selected browser is not on PATH; if dialing the driver fails (driver not running, wrong port); or if the browser launch fails. Subsequent session method calls throw if the session is already closed.

```
if (!services.webdriver.available) {
  runtime.log("no driver on PATH – skip");
} else {
  const d = await services.webdriver.connect({ browser: "chrome", headless:
true });
  try {
    await d.get("data:text/html,<title>hi</title><h1 id=h>Hello</h1>");
    runtime.log("title:", await d.title());
    const el = await d.find("id", "h");
    runtime.log("text:", await el.text());
    await d.executeScript("return 1+2", []);
  } finally {
    await d.quit();
  }
}
```

services.webdriver.probe

```
probe(opts: { url: string }): Promise<{ ready: boolean; status?: number; error?: string }>
```

Check whether a WebDriver endpoint responds at `opts.url/status`. Returns `{ ready, status }` on HTTP success or `{ ready: false, error }` on transport failure. Does not throw on network errors.

Parameters

- `opts` (*{ url: string }*) — `url` is required — the base URL of a running WebDriver server (e.g. `'http://127.0.0.1:9515'`).

Returns: `Promise<{ ready, status? } | { ready: false, error }>` — `ready` is true when the endpoint returns HTTP 200; `status` is the HTTP status code when a response was received; `error` is the transport error message when the request failed entirely.

Throws: Throws if `opts.url` is missing or empty. Transport failures resolve with `{ ready: false, error }` rather than throwing.

```
const r = await services.webdriver.probe({ url: "http://127.0.0.1:9515" });
runtime.log("ready:", r.ready);
```

text

String / regex / charset / data manipulation — all text-shaped transforms.

text.charset.decode

```
decode(input: string | Uint8Array | ArrayBuffer, charset: string): Promise<string>
```

Decode bytes in a named charset to a UTF-8 string.

Parameters

- `input` (*string | Uint8Array | ArrayBuffer*) — Bytes encoded in `charset`.
- `charset` (*string*) — WHATWG/HTML5 encoding name or alias (UTF-8, ISO-8859-1, Windows-1252, Shift_JIS, GBK, ...).

Returns: `Promise<string>` — the bytes decoded to a UTF-8 string.

Throws: Rejects if `input` is an unsupported type, the charset name is unknown, or the bytes are invalid for that encoding.

```
const s = await text.charset.decode(bytes, "Windows-1252");
```

text.charset.detect

```
detect(input: string | Uint8Array | ArrayBuffer): Promise<{ charset: string;
confidence: number; language?: string; candidates: { charset: string; confidence:
number; language?: string }[] }>
```

Detect the most-likely charset of a byte sequence (saintfish/chardet). Resolves to the top guess plus the full candidate list, each scored 0–100.

Parameters

- `input` (*string* | *Uint8Array* | *ArrayBuffer*) — Bytes to sniff. A string is taken as its raw UTF-8 bytes.

Returns: `Promise<{ charset, confidence, language?, candidates }>` — the top match plus all candidates; confidence is chardet's 0–100 score. language is present only when chardet reports one.

Throws: Rejects if input is empty, an unsupported type, or chardet finds no candidates.

```
const r = await text.charset.detect(bytes);
runtime.log(r.charset, r.confidence);
```

text.charset.encode

```
encode(input: string, charset: string): Promise<Uint8Array>
```

Encode a UTF-8 string to bytes in the named charset.

Parameters

- `input` (*string*) — UTF-8 string to encode.
- `charset` (*string*) — Target WHATWG/HTML5 encoding name or alias.

Returns: `Promise<Uint8Array>` — the string encoded as bytes in the target charset.

Throws: Rejects if the charset name is unknown, or a character has no representation in the target encoding (no lossy fallback — characters are not silently dropped).

```
const bytes = await text.charset.encode("café", "ISO-8859-1");
```

text.diff.compare

```
compare(a: string | Uint8Array | ArrayBuffer, b: string | Uint8Array |
ArrayBuffer, opts?: { context?: number; fromFile?: string; toFile?: string }):
Promise<{ identical: boolean; binary: boolean; added: number; removed: number;
diff: string; format: "unified" }>
```

Unified-diff two text inputs. `opts`: `context` (default 3), `fromFile` / `toFile` (default 'a' / 'b'). Binary inputs return `{ binary: true }` with an empty diff.

Parameters

- `a` (*string* | *Uint8Array* | *ArrayBuffer*) — The 'from' / left side. A string is taken as its UTF-8 bytes.
- `b` (*string* | *Uint8Array* | *ArrayBuffer*) — The 'to' / right side.
- `opts` (*{ context?: number; fromFile?: string; toFile?: string }, optional*) — `context` is the number of unchanged lines around each hunk (default 3); `fromFile` / `toFile` are the header labels (default 'a' / 'b').

Returns: `Promise<{ identical, binary, added, removed, diff, format }>` — `diff` holds the unified-diff text (empty when identical or binary); `added/removed` count body +/- lines excluding file headers; `binary` is true when either input has a NUL byte in its first 8 KB.

Throws: Rejects if either input is an unsupported type, or the unified-diff generation fails.

```
const d = await text.diff.compare("a\n", "b\n");
runtime.log(d.diff, d.added, d.removed);
```

text.jq.query

```
query(data: unknown, filter: string): Promise<unknown>
```

Run a jq filter over data and return the first emitted value (or null).

Parameters

- `data` (*unknown*) — Input value — any JSON-like JS value (object/array/scalar). Passed to gojq directly; integers of any width are normalised so queries don't blow up.
- `filter` (*string*) — A jq filter expression, e.g. `".users[0].name"` or `".items | length"`.

Returns: `Promise<unknown>` — the first value the filter emits, or null when it emits nothing (e.g. an optional path like `.a.b?` that misses).

Throws: Rejects if data is undefined, the filter string is empty, the filter fails to parse, or evaluation produces an error.

```
const name = await text.jq.query(obj, ".users[0].name");
```

text.jq.queryAll

```
queryAll(data: unknown, filter: string): Promise<unknown[]>
```

Run a jq filter and drain the iterator into an array.

Parameters

- `data` (*unknown*) — Input value — any JSON-like JS value.
- `filter` (*string*) — A jq filter expression. Use this when the filter explodes a stream, e.g. `"[]"` or `".users[].id"`.

Returns: `Promise<unknown[]>` — every value the filter emits, in order (empty array when it emits nothing).

Throws: Rejects if data is undefined, the filter string is empty, the filter fails to parse, or evaluation produces an error.

```
const ids = await text.jq.queryAll(obj, ".users[].id");
```

text.preg.match

```
match(pattern: string, subject: string): { match: string; groups: string[]; index: number } | null
```

First hit of `/pattern/flags` against subject, or null. Returns `{ match, groups, index }`; optional groups that didn't match surface as empty strings.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (forward-slash delimiter only). Engine is Go RE2. Flags: `i` (case-insensitive), `m` (multiline `^/$`), `s` (dotall). `u/U/x` are rejected.
- `subject` (*string*) — The string to search.

Returns: `{ match, groups, index }` for the first match, where `match` is the full match, `groups` holds the numbered submatches after group 0 (unmatched optionals as `""`), and `index` is the byte offset; null when there is no match.

Throws: Throws if the pattern is empty, lacks a leading/closing `/`, uses an unsupported flag (`u/U/x`), or fails RE2 compilation (e.g. backreferences/lookaround, which RE2 does not support).

```
const m = text.preg.match("/(\\d+)/", "x42"); // { match: "42", groups: ["42"],
index: 1 }
```

text.preg.matchAll

```
matchAll(pattern: string, subject: string): { match: string; groups: string[];
index: number }[]
```

Every hit of `/pattern/flags` against `subject`, as an array of `{ match, groups, index }` objects.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (RE2 engine; `i/m/s` flags).
- `subject` (*string*) — The string to search.

Returns: Array of `{ match, groups, index }` objects, one per non-overlapping match; empty array when there is none.

Throws: Throws on the same conditions as `preg.match` (bad delimiter, unsupported flag, RE2 compile failure).

```
const all = text.preg.matchAll("/\\d+"/, "1 22 333"); // 3 matches
```

text.preg.replace

```
replace(pattern: string, replacement: string, subject: string): string
```

Substitute every match of `/pattern/flags` in `subject`. Replacement uses Go's `$1 / ${1}` backref syntax — PHP's `\1` form is NOT translated.

Parameters

- `pattern` (*string*) — A PHP-style `/regex/flags` delimited pattern (RE2 engine; `i/m/s` flags).
- `replacement` (*string*) — Replacement template using Go's RE2 syntax: `$1 / ${name}` reference groups; use `${1}` to disambiguate. PHP's `\1` backrefs are NOT supported.
- `subject` (*string*) — The string to transform.

Returns: `string` — `subject` with every match replaced (all occurrences).

Throws: Throws on the same conditions as `preg.match` (bad delimiter, unsupported flag, RE2 compile failure).

```
text.preg.replace("/(\\w+)@/", "${1}_at_", "a@b"); // "a_at_b"
```

text.preg2.match

```
match(pattern: string, subject: string): { match: string; groups: string[]; index: number } | null
```

First hit of /pattern/flags via regexp2 (PCRE). Supports lookahead/lookbehind/backreferences. Same { match, groups, index } shape as preg. No linear-time guarantee.

Parameters

- `pattern` (*string*) — A /regex/flags delimited pattern run on the .NET-flavoured regexp2 engine (lookaround, backreferences, possessive quantifiers). Flags: i, m, s, and x (ignore-pattern-whitespace, unavailable in preg). u/U are rejected.
- `subject` (*string*) — The string to search.

Returns: { match, groups, index } for the first match (same shape as preg.match); null when there is no match.

Throws: Throws if the pattern is empty, lacks a leading/closing /, uses an unsupported flag (u/U), fails regexp2 compilation, or errors during matching. Backtracking engine: guard untrusted input with a timeout.

```
const m = text.preg2.match("/(?<=@)\\w+/", "a@host"); // { match: "host", ... }
```

text.preg2.matchAll

```
matchAll(pattern: string, subject: string): { match: string; groups: string[]; index: number }[]
```

Every hit of /pattern/flags via regexp2 (PCRE), as an array of { match, groups, index }.

Parameters

- `pattern` (*string*) — A /regex/flags delimited pattern on the regexp2 engine (i/m/s/x flags).
- `subject` (*string*) — The string to search.

Returns: Array of { match, groups, index } objects, one per match; empty array when there is none.

Throws: Throws on the same conditions as preg2.match. Backtracking engine: guard untrusted input with a timeout.

```
const all = text.preg2.matchAll("/\\w+/", "a b c"); // 3 matches
```

text.preg2.replace

```
replace(pattern: string, replacement: string, subject: string): string
```

Substitute every match of /pattern/flags via regexp2. Replacement uses .NET \$1 / \${1} syntax. Backtracking engine — keep a timeout around untrusted input.

Parameters

- `pattern` (*string*) — A `/regex/flags` delimited pattern on the `regexp2` engine (i/m/s/x flags).
- `replacement` (*string*) — Replacement template using .NET substitution syntax: `$1` / `${1}` reference groups, `$$` is a literal dollar.
- `subject` (*string*) — The string to transform.

Returns: string — subject with every match replaced (startAt 0, count -1).

Throws: Throws on the same conditions as `preg2.match`, or if replacement fails.

Backtracking engine: guard untrusted input with a timeout.

```
text.preg2.replace("/(\\w)\\1/", "X", "aabb"); // backref-aware: "XX"
```

text.str.base64Decode

```
base64Decode(input: string): string
```

Decode standard (RFC 4648) base64 with padding. The standard alphabet only — URL-safe input (containing `-` or `_`) is NOT auto-detected and will throw; pre-translate it to the standard alphabet first.

Parameters

- `input` (*string*) — Standard-alphabet base64 string with `=` padding. URL-safe characters (`-` / `_`) are not accepted.

Returns: string — the decoded bytes interpreted as a UTF-8 string.

Throws: Throws a `TypeError` if input is missing/null/undefined; throws if the input is not valid standard base64 (including URL-safe alphabet input).

```
text.str.base64Decode("aGk="); // "hi"
```

text.str.base64Encode

```
base64Encode(input: string): string
```

Standard base64 (with padding).

Parameters

- `input` (*string*) — UTF-8 string to encode (encoded as its raw bytes).

Returns: string — RFC 4648 standard base64 with `=` padding.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.base64Encode("hi"); // "aGk="
```

text.str.br2nl

```
br2nl(input: string): string
```

Inverse of `nl2br`: `<br>`, `<br/>`, `<br />` → `'\n'`.

Parameters

- `input` (*string*) — Source text. Any case-insensitive `<br>`, `<br/>`, or `<br />` variant is matched.

Returns: string — input with each `<br>` variant replaced by a single `\n`.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.br2nl("a<br/>b"); // "a\nb"
```

text.str.htmlEntityDecode

```
htmlEntityDecode(input: string): string
```

Decode named and numeric HTML entities to their UTF-8 equivalents.

Parameters

- `input` (*string*) — Text containing HTML entities (named like `&` or numeric like `' / '`).

Returns: string — input with recognised entities decoded to their UTF-8 characters.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.htmlEntityDecode("a & b"); // "a & b"
```

text.str.lpad

```
lpad(input: string, len: number, padChar?: string): string
```

Shortcut for `pad(side: 'left')`.

Parameters

- `input` (*string*) — The string to pad on the left.
- `len` (*number*) — Target rune length.
- `padChar` (*string, optional*) — Pad string; defaults to a single space.

Returns: string — input left-padded to len runes.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.lpad("7", 3, "0"); // "007"
```

text.str.ltrim

```
ltrim(input: string, mask?: string): string
```

Like `trim`, left side only.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset of characters to strip from the left. Defaults to the whitespace set `"\t\n\r\v\f"`.

Returns: string — input with leading mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.ltrim("--x", "-"); // "x"
```

text.str.nl2br

```
nl2br(input: string, xhtml?: boolean): string
```

Replace newlines with `<br>` (or `<br/>` when `xhtml=true`).

Parameters

- `input` (*string*) — Source text. CRLF is normalised to LF first, so each line break yields one tag.
- `xhtml` (*boolean, optional*) — When truthy, emit the self-closing `<br/>` instead of `<br>`. Defaults to false.

Returns: string — input with each `\n` replaced by the chosen `<br>` tag followed by the original newline.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.nl2br("a\nb"); // "a<br>\nb"
```

text.str.normalizeNewlines

```
normalizeNewlines(input: string, style?: "lf" | "crlf" | "cr"): string
```

Canonicalise any mix of `\r\n`, `\r`, `\n` to the requested style (`'lf'` | `'crlf'` | `'cr'`).

Parameters

- `input` (*string*) — Text with any mix of CRLF, CR, and LF line endings.
- `style` (`"lf"` | `"crlf"` | `"cr"`, *optional*) — Target line-ending style. Defaults to `'lf'`.

Returns: string — input with every line ending rewritten to the requested style.

Throws: Throws a `TypeError` if style is not one of `'lf'`, `'crlf'`, or `'cr'`.

```
text.str.normalizeNewlines("a\r\nb", "lf"); // "a\nb"
```

text.str.pad

```
pad(input: string, len: number, padChar?: string, side?: "right" | "left" | "both"): string
```

Pad to `len` with `padChar` (default `' '`). `side` is `'right'` (default), `'left'`, or `'both'`.

Parameters

- `input` (*string*) — The string to pad. Returned unchanged when its rune length already meets or exceeds `len`.
- `len` (*number*) — Target rune length. Measured in runes, not bytes.
- `padChar` (*string, optional*) — Pad string; truncated to fit the needed width. Defaults to a single space.

- `side` (*“right”* / *“left”* / *“both”*, *optional*) — Which side(s) to pad. ‘both’ splits the deficit with the extra rune going right. Defaults to ‘right’; any unknown value is treated as ‘right’.

Returns: string — input padded to len runes on the chosen side(s).

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.pad("7", 3, "0", "left"); // "007"
```

text.str.printf

```
printf(format: string, args?: ...unknown): void
```

sprintf + write to stdout.

Parameters

- `format` (*string*) — A Go fmt format string (same verbs as `sprintf`).
- `args` (*...unknown*, *optional*) — Values substituted into the verbs.

Returns: void — writes the formatted text directly to process stdout; returns nothing.

Throws: Throws a `TypeError` if called with no arguments (format string required).

```
text.str.printf("%d items\n", 3);
```

text.str.reverse

```
reverse(input: string): string
```

Rune-aware reversal — `reverse('café')` is `'éfac'`.

Parameters

- `input` (*string*) — The string to reverse. Reversed by Unicode code point, not byte, so multi-byte runes stay intact.

Returns: string — the input reversed rune-by-rune.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.reverse("café"); // "éfac"
```

text.str.rpad

```
rpadd(input: string, len: number, padChar?: string): string
```

Shortcut for `pad(side: 'right')`.

Parameters

- `input` (*string*) — The string to pad on the right.
- `len` (*number*) — Target rune length.
- `padChar` (*string*, *optional*) — Pad string; defaults to a single space.

Returns: string — input right-padded to len runes.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.pad("7", 3, "."); // "7.."
```

text.str.trim

```
rtrim(input: string, mask?: string): string
```

Like trim, right side only.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset of characters to strip from the right. Defaults to the whitespace set “\t\n\r\v\f”.

Returns: string — input with trailing mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.rtrim("x...", "."); // "x"
```

text.str.sprintf

```
sprintf(format: string, args?: ...unknown): string
```

Go’s fmt verbs (%s, %d, %x, %.2f, %v, %t, %q, ...) — not PHP’s.

Parameters

- `format` (*string*) — A Go fmt format string. Uses Go verbs: %s string, %d integer, %f / %.2f float, %x hex, %v default, %t bool, %q quoted, %% literal percent.
- `args` (*...unknown, optional*) — Values substituted into the verbs (passed through `.Export()`, so JS numbers arrive as Go `int64/float64`).

Returns: string — the formatted result.

Throws: Throws a `TypeError` if called with no arguments (format string required). Verb/arg mismatches are not thrown — Go renders `%!verb(...)` error markers inline.

```
text.str.sprintf("%s=%d", "n", 5); // "n=5"
```

text.str.stripHtml

```
stripHtml(input: string): string
```

Remove HTML tags and decode common entities.

Parameters

- `input` (*string*) — HTML source. Anything matching `<...>` is removed.

Returns: string — the input with all `<...>` tag spans deleted.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.stripHtml("<b>hi</b>"); // "hi"
```

text.str.trim

```
trim(input: string, mask?: string): string
```

Strip whitespace (or any char in the optional mask string) from both ends.

Parameters

- `input` (*string*) — The string to trim.
- `mask` (*string, optional*) — Cutset: any character in this string is trimmed (PHP-style, not a prefix). Defaults to the whitespace set “\t\n\r\v\f”.

Returns: string — input with leading and trailing mask characters removed.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.trim("  hi  "); // "hi"
```

text.str.urlDecode

```
urlDecode(input: string): string
```

Inverse of `urlEncode`.

Parameters

- `input` (*string*) — A form-encoded string (+ decodes to space, %XX to bytes).

Returns: string — the decoded value.

Throws: Throws a `TypeError` if input is missing/null/undefined; throws on malformed percent-escapes.

```
text.str.urlDecode("a+b%26c"); // "a b&c"
```

text.str.urlEncode

```
urlEncode(input: string): string
```

Form-encoding (‘+’ for space). For path segments use `encodeURIComponent` (provided by goja).

Parameters

- `input` (*string*) — String to percent-encode. Uses application/x-www-form-urlencoded rules: space becomes +.

Returns: string — the form-encoded value.

Throws: Throws a `TypeError` if input is missing, null, or undefined.

```
text.str.urlEncode("a b&c"); // "a+b%26c"
```

tui

Multi-pane terminal UI: layout, pane, write, focus.

tui.layout

```
layout(tree: { name: string; title?: string; weight?: number } | { rows: object[]; weight?: number } | { cols: object[]; weight?: number }): void
```

Declare the pane layout for this Run. Tree nodes: { name, title?, weight?, autoscroll? } (leaf), { rows: [...], weight? } (vertical split), { cols: [...], weight? } (horizontal split). The root node also accepts { mouse?: boolean }. Throws on duplicate names, empty rows/cols, unknown keys, or under `-watch`.

Parameters

- `tree` (*{ name: string; title?: string; weight?: number } | { rows: object[]; weight?: number } | { cols: object[]; weight?: number }*) — The root layout node. Exactly one of name / rows / cols must be set per node. A leaf (name) becomes a bordered pane addressable via `tui.pane(name)`; name must be a non-empty string and unique across the whole tree. rows stacks children top-to-bottom; cols places them side-by-side; both arrays must be non-empty. weight (positive integer, default 1) sets the child's proportional share of its parent's space. title (string, leaf only) seeds the pane's border caption. Any other key is rejected. The tree is realised over the full terminal as a `tview Flex` when `stdout` is a TTY; otherwise it falls back to prefixed-line output. autoscroll (boolean, leaf only, default true) controls whether the pane follows the tail as new lines arrive; set false to keep it pinned at the top. mouse (boolean, root only, default false) enables mouse support: the wheel scrolls the pane under the cursor (without changing which pane is focused) and a left-click focuses the pane under the cursor — at the cost of the terminal's native click-drag text selection (use Shift/Option+drag to select while mouse mode is on). wrap (string, leaf only, default "char") sets line wrapping: "char" wraps mid-word, "word" wraps at word boundaries, "off" disables wrapping (long lines scroll horizontally). color (boolean, leaf only, default true) renders subprocess ANSI as pane colors; set false to strip ANSI and show plain text. wrap and color affect TTY rendering only (the non-TTY fallback always emits plain prefixed lines).

Returns: void — installs the layout and brings up the UI (TTY) or the fallback line writer (non-TTY); the controller is torn down automatically at Run end.

Throws: Throws if called under `-watch`; if layout was already called this Run; if the tree argument is missing/null/undefined; if any node violates the structure rules (not an object, more than one of name/rows/cols, missing all three, empty rows/cols, unknown key, non-string or empty name, duplicate name, title on a non-leaf, or non-positive/non-integer weight) — the error includes the tree path (e.g. "rows[1].cols[0]"); or if the terminal screen / fallback writer fails to start.

```
tui.layout({ cols: [{ name: "log", title: "Log" }, { name: "out", weight: 2 }] });
tui.pane("log").writeln("started");
```

tui.onKey

```
onKey(handler: (key: { name: string; rune: string; ctrl: boolean; alt: boolean; shift: boolean }) => void): () => void
```

Register a callback invoked on every keypress (TTY mode) except Ctrl-C. Returns an unsubscribe function. No-op in non-TTY (fallback) mode.

Parameters

- `handler` *((key: { name: string; rune: string; ctrl: boolean; alt: boolean; shift: boolean }) => void)* — Called for each keypress with a key descriptor. `name` is the tcell key name (“Enter”, “Up”, “Tab”, “Ctrl-A”, “F1”, ...) or “Rune” for a printable character (rune holds it). For Ctrl+letter combinations the name carries the combo (e.g. “Ctrl-A”) and the `ctrl` flag may be false, so check name for control keys. Built-in navigation (Tab focus, PgUp/PgDn/arrows/Home/End scroll) still runs in addition to the handler (coexist model). Ctrl-C always aborts the script and is never delivered.

Returns: An unsubscribe function; call it to stop receiving keys. In non-TTY mode `onKey` registers nothing and returns a no-op unsubscribe. Note: registering `onKey` alone does not keep the Run alive — keep an outstanding await (e.g. `waitKey` or a `sleep`) if the script should stay open.

Throws: Throws if called before `tui.layout`, or if the argument is not a function.

```
tui.layout({ rows: [{ name: "log" }] });  
const off = tui.onKey((k) => { if (k.name === "Rune" && k.rune === "q") off(); });
```

tui.pane

```
pane(name: string): { write(text: string): void; writeln(text: string): void;  
clear(): void; title(text: string): void }
```

Return a Pane handle for a declared pane. Throws if the name wasn’t in the layout. Handle methods: `write(text)`, `writeln(text)`, `clear()`, `title(text)`. `services.exec.shell({pane})` streams subprocess I/O into a pane.

Parameters

- `name` *(string)* — The leaf name declared in the `tui.layout` tree.

Returns: A Pane handle. `write(text)` appends text (subprocess ANSI SGR colors are translated to pane colors); `writeln(text)` appends text followed by a newline; `clear()` empties the pane (no-op in the non-TTY fallback); `title(text)` updates the pane’s border caption (no-op in fallback). All methods return undefined and are safe to call from any callback. The handle can also be passed as the pane option to `services.exec.shell` to stream a subprocess’s stdout/stderr live into the pane.

Throws: Throws if `tui.layout` has not been called yet this Run, or if `name` was not declared as a leaf in the layout (the message lists the available pane names).

```
const p = tui.pane("out");  
p.title("Output");  
p.writeln("hello");
```

tui.waitKey

```
waitKey(...args: unknown[]): Promise<{ name: string; rune: string; ctrl: boolean; alt: boolean; shift: boolean }>
```

Resolve with the next keypress (TTY mode). One-shot — await again for the next key. Rejects in non-TTY (fallback) mode.

Returns: A Promise resolving to the next key descriptor (same shape as onKey's argument). Concurrent waitKey calls resolve FIFO — one keypress resolves the oldest pending call. While a waitKey is pending the TUI stays open (this is the idiomatic “press any key to close” hold). Ctrl-C aborts and is never delivered.

Throws: Rejects if called before tui.layout, or in non-TTY mode (no interactive terminal), or if the TUI is closed while waiting.

```
tui.layout({ rows: [{ name: "log" }] });
tui.pane("log").writeln("Done. Press any key to close.");
await tui.waitKey();
```

web

Fetch & parse web documents: RSS/Atom/JSON feeds (web.feed), sitemaps incl. gzip + index expand (web.sitemap), and lenient HTML scraping with CSS + XPath (web.html). Each offers parse(string) and async load(url, opts?).

web.feed.load

```
load(url: string, opts?: { timeout?: number | string; headers?: Record<string, string>; follow?: boolean; userAgent?: string; username?: string; password?: string }): Promise<{ feedType: string; title: string; description: string; link: string; updated: string | null; items: Array<{ title: string; link: string; published: string | null; updated: string | null; content: string; summary: string; author: string; guid: string; categories: string[]; raw: Record<string, unknown> }> }>
```

Fetch a URL and parse it as a feed (see feed.parse). Reuses the net.http option surface and sends a default User-Agent unless overridden. Throws on a non-2xx response.

Parameters

- `url` (*string*) — The feed URL to GET.
- `opts` (*{ timeout?: number | string; headers?: Record<string, string>; follow?: boolean; userAgent?: string; username?: string; password?: string }, optional*) — Fetch options: timeout (ms or duration string), headers, follow redirects, userAgent, basic-auth username/password.

Returns: A promise resolving to the normalized feed object.

Throws: Throws on transport failure or a non-2xx response, and on malformed feed content.

```
const f = await web.feed.load("https://example.com/feed.xml");
```

web.feed.parse

```
parse(source: string): { feedType: string; title: string; description: string;
link: string; updated: string | null; items: Array<{ title: string; link: string;
published: string | null; updated: string | null; content: string; summary:
string; author: string; guid: string; categories: string[]; raw: Record<string,
unknown> }> }
```

Parse RSS, Atom, or JSON-feed text into a normalized feed model. Format is auto-detected; feedType reports it. RSS/Atom field differences are unified (pubDate/updated, description/summary). Each item carries a `raw` escape hatch with format-specific extras (enclosures, namespaced elements like `media:/dc:`).

Parameters

- `source` (*string*) — The feed document text (RSS/Atom XML or JSON Feed).

Returns: The normalized feed object.

Throws: Throws on empty/malformed/undetectable feed input.

```
const f = web.feed.parse(xml); f.items[0].title;
```

web.html.load

```
load(url: string, opts?: { timeout?: number | string; headers?: Record<string,
string>; follow?: boolean; userAgent?: string; username?: string; password?:
string }): Promise<{ find(selector: string): unknown; findAll(selector: string):
unknown[]; xpath(expr: string): unknown; xpathAll(expr: string): unknown[];
text(): string; html(): string; innerHTML(): string; tag(): string; attr(name:
string): string | null; attrs(): Record<string, string>; }>
```

Fetch a URL and parse the response as lenient HTML (see `html.parse`). Reuses the `net.http` option surface with a default User-Agent. Throws on a non-2xx response.

Parameters

- `url` (*string*) — The page URL to GET.
- `opts` (*{ timeout?: number | string; headers?: Record<string, string>; follow?: boolean; userAgent?: string; username?: string; password?: string }, optional*) — Fetch options: timeout (ms or duration string), headers, follow redirects, userAgent, basic-auth username/password.

Returns: A promise resolving to the document Node handle.

Throws: Throws on transport failure or a non-2xx response.

```
const doc = await web.html.load("https://example.com"); doc.findAll("a").map(a =>
a.attr("href"));
```

web.html.parse

```
parse(source: string): { find(selector: string): unknown; findAll(selector:
string): unknown[]; xpath(expr: string): unknown; xpathAll(expr: string):
```



```
unknown[]; text(): string; html(): string; innerHTML(): string; tag(): string;
attr(name: string): string | null; attrs(): Record<string, string>; }
```

Parse HTML leniently (real-world tag soup is accepted, never throws on bad markup) into a chainable Node. Query with CSS (find/findAll) or XPath (xpath/xpathAll); read with text/html/innerHTML/tag/attr/attrs. Sub-queries are scoped to the receiver node (use .

Parameters

- `source` (*string*) — The HTML document text.

Returns: A Node handle rooted at the document.

Throws: Does not throw on malformed markup; only on unreadable input.

```
const doc = web.html.parse(html); doc.find("h1").text();
```

web.sitemap.load

```
load(url: string, opts?: { expand?: boolean } & { timeout?: number | string;
headers?: Record<string, string>; follow?: boolean; userAgent?: string; username?:
string; password?: string }): Promise<{ type: "urlset" | "sitemapindex"; urls:
Array<{ loc: string; lastmod?: string; changefreq?: string; priority?: number }>;
sitemaps: string[]; errors: Array<{ url: string; error: string }> >>
```

Fetch a sitemap URL and parse it. Transparently decompresses .xml.gz (gzip magic-byte detection on the body; a Content-Encoding: gzip response is unwrapped by the HTTP transport). For a sitemapindex, pass {expand:true} to fetch all child sitemaps (bounded; per-child errors collected in `errors[]`) and merge their urls into `urls`.

Parameters

- `url` (*string*) — The sitemap URL to GET (may be .xml.gz).
- `opts` (*{ expand?: boolean } & { timeout?: number | string; headers?: Record<string, string>; follow?: boolean; userAgent?: string; username?: string; password?: string }, optional*) — expand:true fetches & merges child sitemaps for an index. Plus the standard fetch options.

Returns: A promise resolving to the sitemap object.

Throws: Throws on transport failure, non-2xx, or a non-sitemap document. Per-child expand failures are captured in `errors[]`, not thrown.

```
const sm = await web.sitemap.load("https://example.com/sitemap.xml", { expand:
true });
```

web.sitemap.parse

```
parse(source: string): { type: "urlset" | "sitemapindex"; urls: Array<{ loc:
string; lastmod?: string; changefreq?: string; priority?: number }>; sitemaps:
string[]; errors: Array<{ url: string; error: string }> }
```

Parse a sitemap document (urlset or sitemapindex) into {type, urls, sitemaps, errors}. urls carry loc/lastmod/changefreq/priority; an index lists child sitemap URLs in `sitemaps`.

Parameters

- `source` (*string*) — The sitemap XML text (decompressed).

Returns: The parsed sitemap object.

Throws: Throws when the document is neither `<urlset>` nor `<sitemapindex>`.

```
const sm = web.sitemap.parse(xml); sm.urls.map(u => u.loc);
```

This manual covers sercon v0.66.0. Whenever you add, remove, or change a flag, a binding, or the script API, update this file alongside the help screen (`--help`), the examples walkthrough (`--examples`), and the `CHANGELOG.md` .